

AI VIETNAM  
All-in-One Course  
(TA Session)

# Logistic Regression

## Exercise



AI VIET NAM  
[@aivietnam.edu.vn](http://aivietnam.edu.vn)

Dinh-Thang Duong – TA

# Outline

- Review
- Titanic Survival Prediction
- Sentiment Analysis
- Question

# Review

## ❖ Introduction

| Positive   | Negative  |
|--|---|
| <br><p>“Delicious food for an affordable price. Nice view. Fast service. We will definitely be visiting again.”</p> | <br><p>“Food was terrible with unprofessional service. I will never revisit this restaurant again.”</p> |

**Problem Statement:** Given a review, classify it into one of the three classes: Positive, Negative or Neutral.

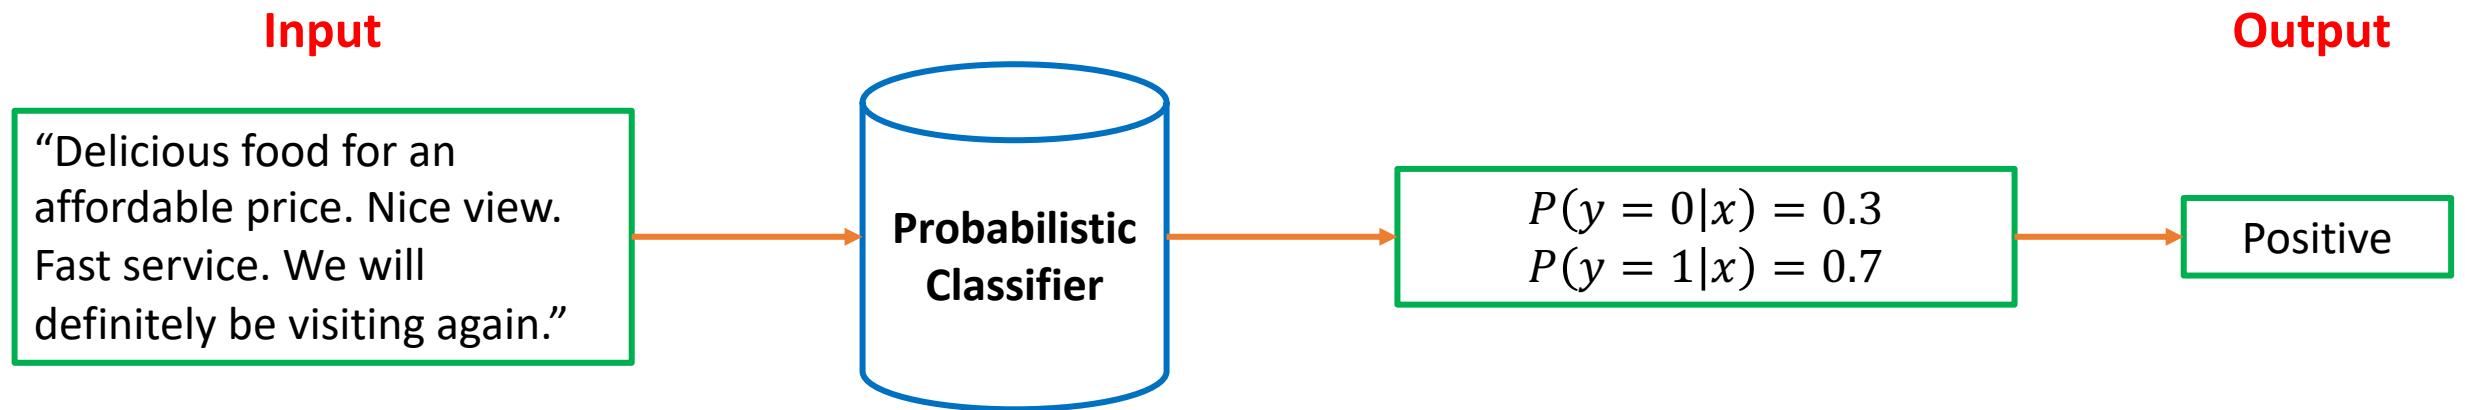
**Input:** “Delicious food for an affordable price. Nice view. Fast service. We will definitely be visiting again.”

**Output:** “Positive”

Sentiment Analysis Problem

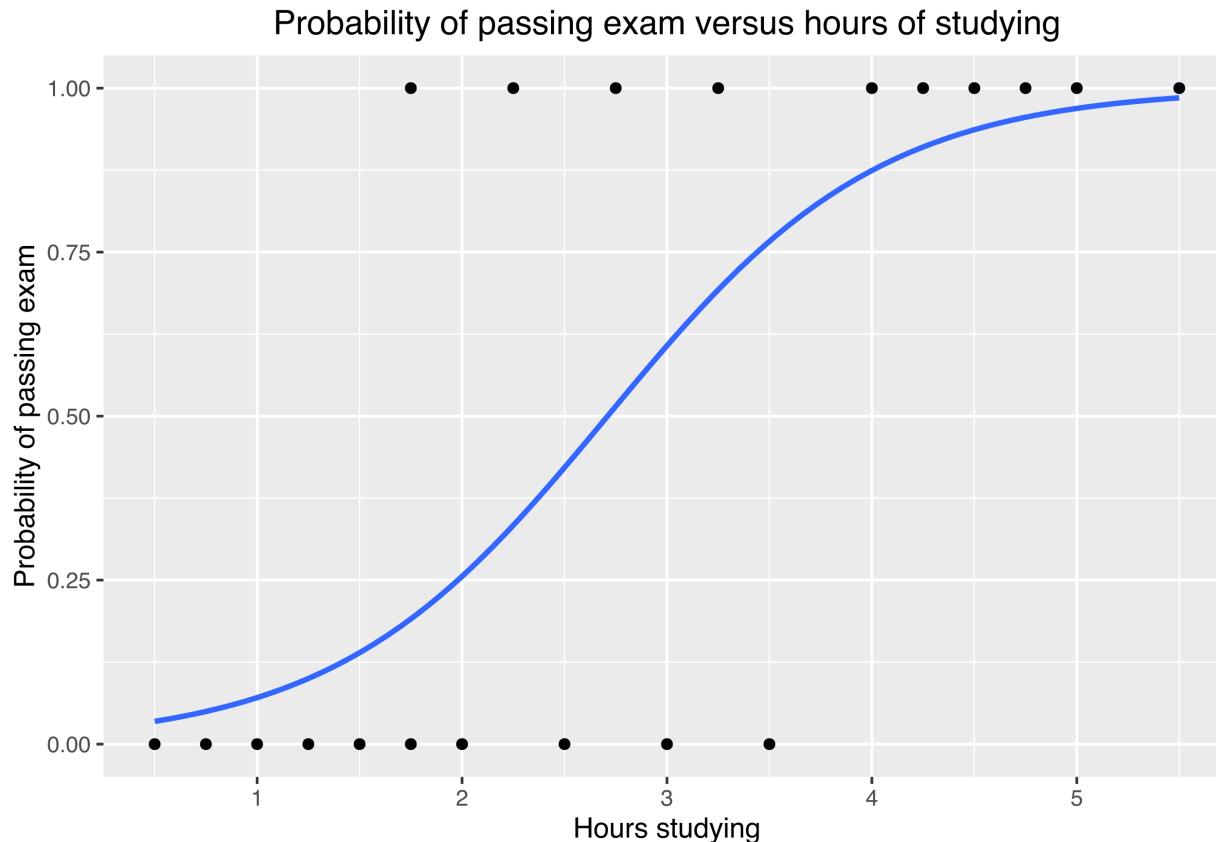
# Review

## ❖ Probabilistic Classifier



# Review

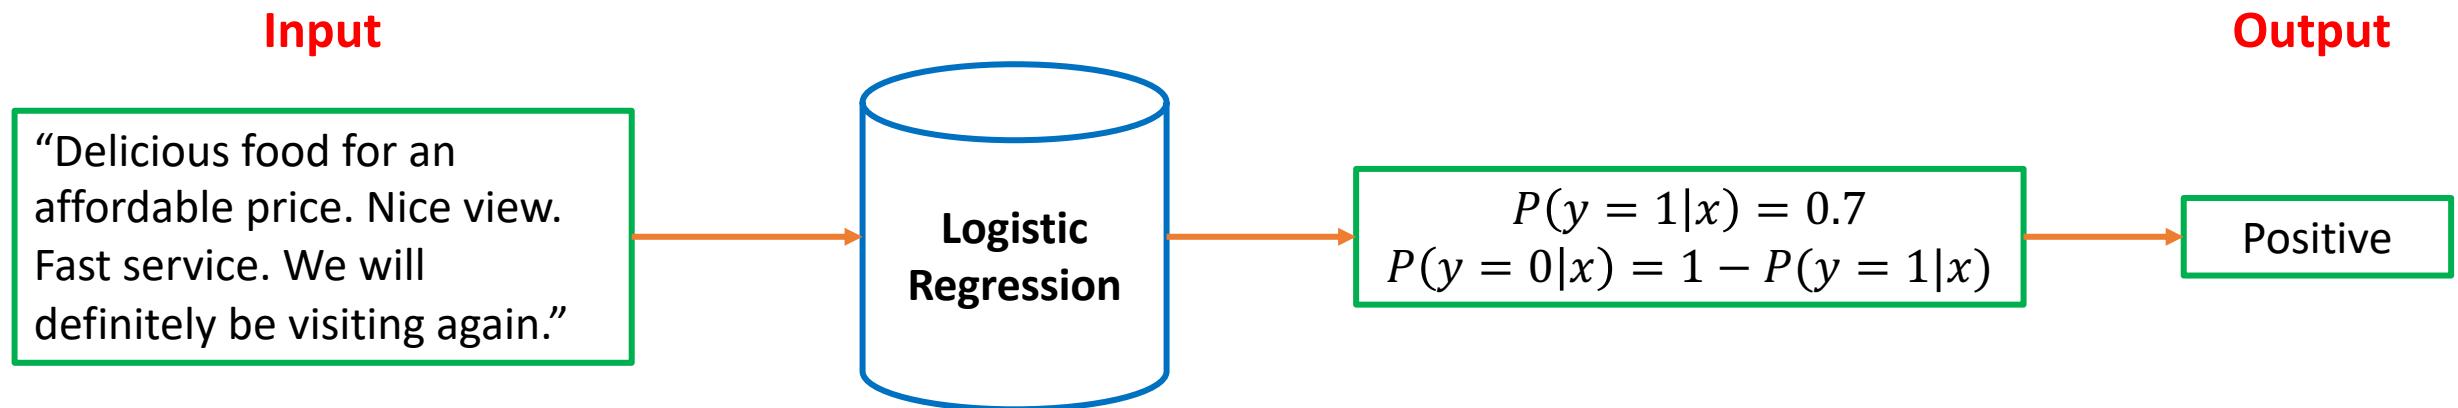
## ❖ Logistic Regression



**Logistic Regression:** A statistical method used for binary classification and probability estimation. This is one of the fundamental tools in ML and statistics that models the relationship between one or more features and binary outcome variable.

# Review

## ❖ Logistic Regression



# Review

## ❖ Logistic Regression

### Raw Text

“Delicious food for an affordable price. Nice view. Fast service. We will definitely be visiting again.”

|       |       |       |     |       |
|-------|-------|-------|-----|-------|
| $x_1$ | $x_2$ | $x_3$ | ... | $x_n$ |
|-------|-------|-------|-----|-------|

Features Vector Representation

### Training Data

| x         | y         |
|-----------|-----------|
| $x^{(1)}$ | $y^{(1)}$ |
| $x^{(2)}$ | $y^{(2)}$ |
| $x^{(3)}$ | $y^{(3)}$ |
| ...       | ...       |
| $x^{(n)}$ | $y^{(n)}$ |

### Features X

|       |       |       |     |       |
|-------|-------|-------|-----|-------|
| $x_1$ | $x_2$ | $x_3$ | ... | $x_n$ |
|-------|-------|-------|-----|-------|



### Weights W

|       |       |       |     |       |
|-------|-------|-------|-----|-------|
| $w_1$ | $w_2$ | $w_3$ | ... | $w_n$ |
|-------|-------|-------|-----|-------|



learning

The bias term



When make a decision:  $z = \sum_{i=1}^n w_i x_i + b = w \cdot x + b$

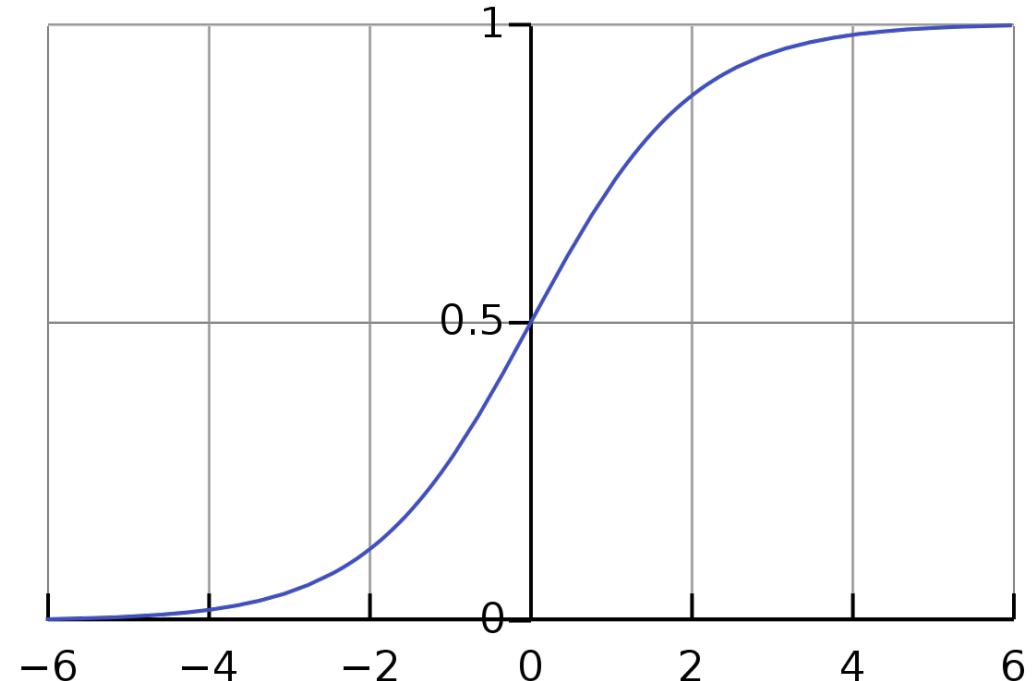
# Review

## ❖ Logistic Regression: Sigmoid function

When make a decision:  $z = \sum_{i=1}^n w_i x_i + b = w \cdot x + b$

The value of this formula ranges from  $(-\infty, +\infty)$ .  
=> Not introduce probability.

To create probability, we will pass z to **sigmoid function**.



$$\sigma(z) = \frac{1}{1 + e^{-z}} = \frac{1}{1 + \exp(-z)}$$

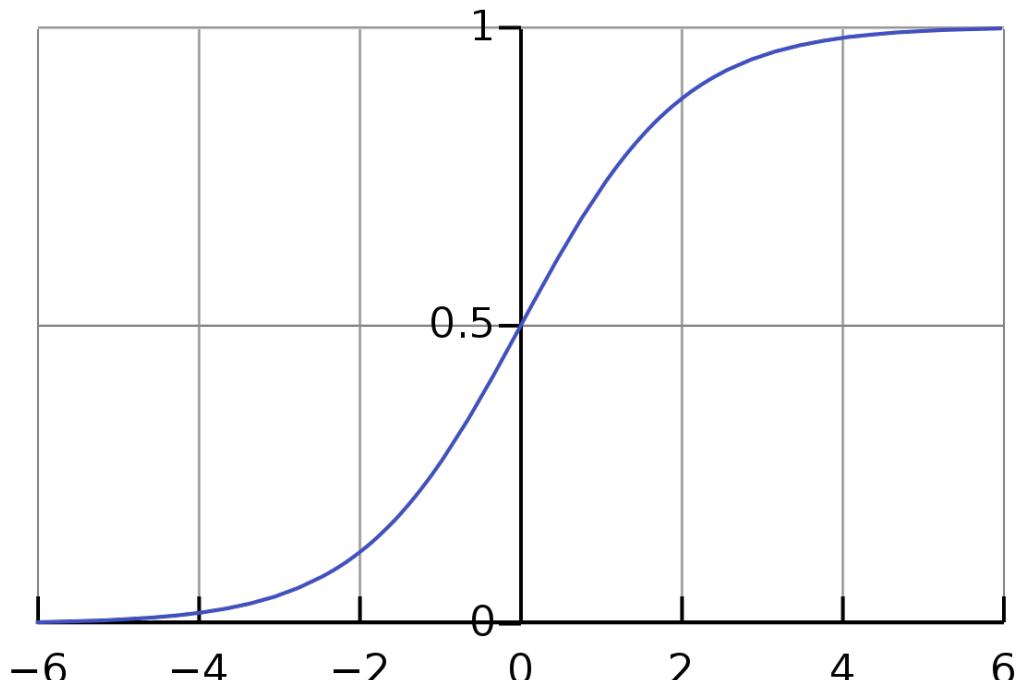
# Review

## ❖ Logistic Regression property

We need to make sure the sum of two cases = 1:

- $z = \sum_{i=1}^n w_i x_i + b = w \cdot x + b$
- $P(y = 1) = \sigma(z) = \frac{1}{1 + \exp(-z)}$
- $P(y = 0) = 1 - \sigma(z) = 1 - \frac{1}{1 + \exp(-z)} = \frac{\exp(-z)}{1 + \exp(-z)}$

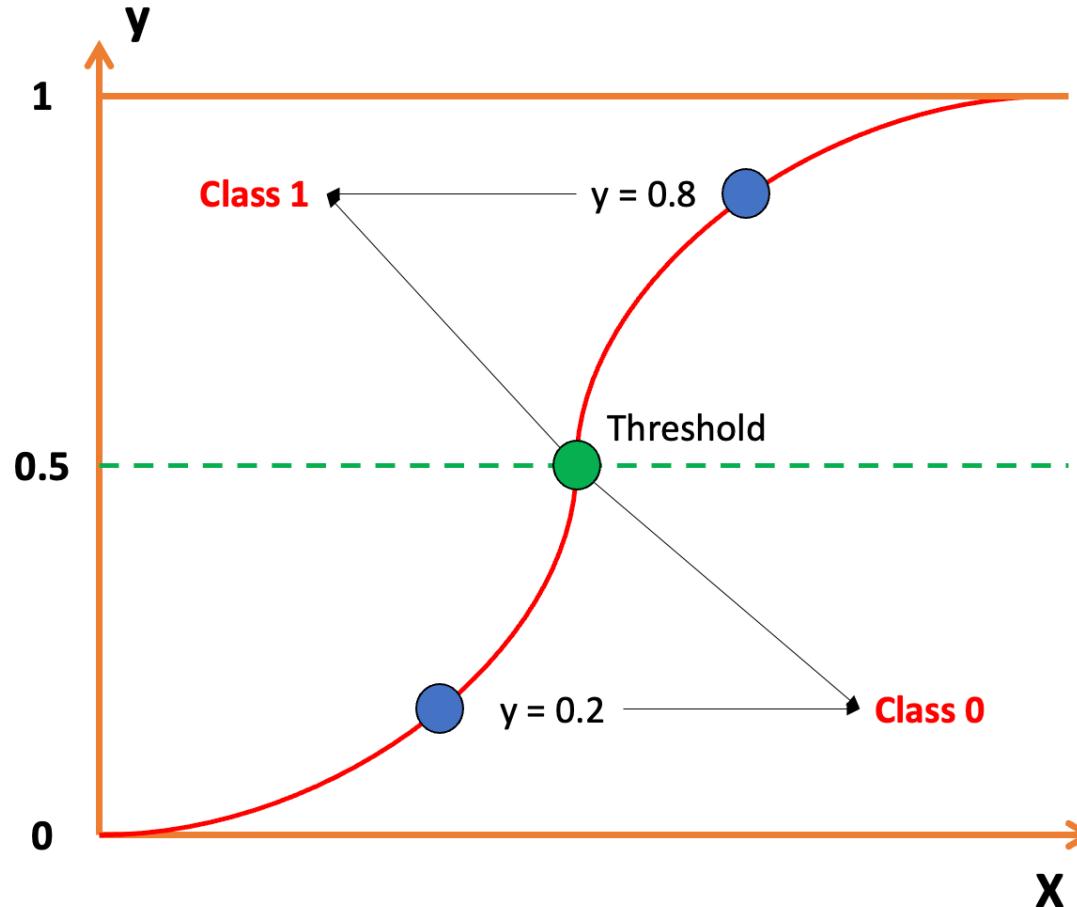
$$1 - \sigma(z) = \frac{\exp(-z)}{1 + \exp(-z)} = \frac{1}{\frac{1}{\exp(-z)} + 1} = \frac{1}{1 + \exp(z)} = \sigma(-z)$$
$$\Rightarrow P(y = 0) = 1 - \sigma(z) = \sigma(-z)$$



$$\sigma(z) = \frac{1}{1 + e^{-z}} = \frac{1}{1 + \exp(-z)}$$

# Review

## ❖ Logistic Regression Prediction



When making a prediction:

$$\text{predict}(x) = \begin{cases} 1 & \text{if } P(y = 1|x) > \text{threshold} \\ 0 & \text{otherwise} \end{cases}$$

# Review

## ❖ Logistic Regression loss function

- $z = \sum_{i=1}^n w_i x_i + b = w \cdot x + b$
- $\hat{y} = \sigma(z) = \frac{1}{1+e^{-z}} = \frac{1}{1+\exp(-z)}$

Now we want to measure how different between  $\hat{y}$  (prediction) and  $y$  (true label) =>  $L(\hat{y}, y)$

**Case 1 (Correct prediction):**  $\hat{y} = 0.7$  and  $y = 1$

$$\begin{aligned} L(\hat{y}, y) &= -\log(p(y|x)) \\ &= -[y \log(\hat{y}) + (1-y) \log(1-\hat{y})] \\ &= -\log(\hat{y}) = -\log(0.7) = 0.36 \end{aligned}$$

- $p(y|x) = \hat{y}^y (1-\hat{y})^{1-y}$  (Bernoulli Distribution)
- $\log(p(y|x)) = \log(\hat{y}^y (1-\hat{y})^{1-y})$   
 $= \log(\hat{y}^y) + \log((1-\hat{y})^{1-y})$   
 $= y \log(\hat{y}) + (1-y) \log(1-\hat{y})$
- $L(\hat{y}, y) = -\log(p(y|x)) = -[y \log(\hat{y}) + (1-y) \log(1-\hat{y})]$   
=> **Cross-entropy loss.**

**Case 2 (Wrong prediction):**  $\hat{y} = 0.7$  and  $y = 0$

$$\begin{aligned} L(\hat{y}, y) &= -\log(p(y|x)) \\ &= -[y \log(\hat{y}) + (1-y) \log(1-\hat{y})] \\ &= -\log(1-\hat{y}) = -\log(0.3) = 1.2 \end{aligned}$$

# Review

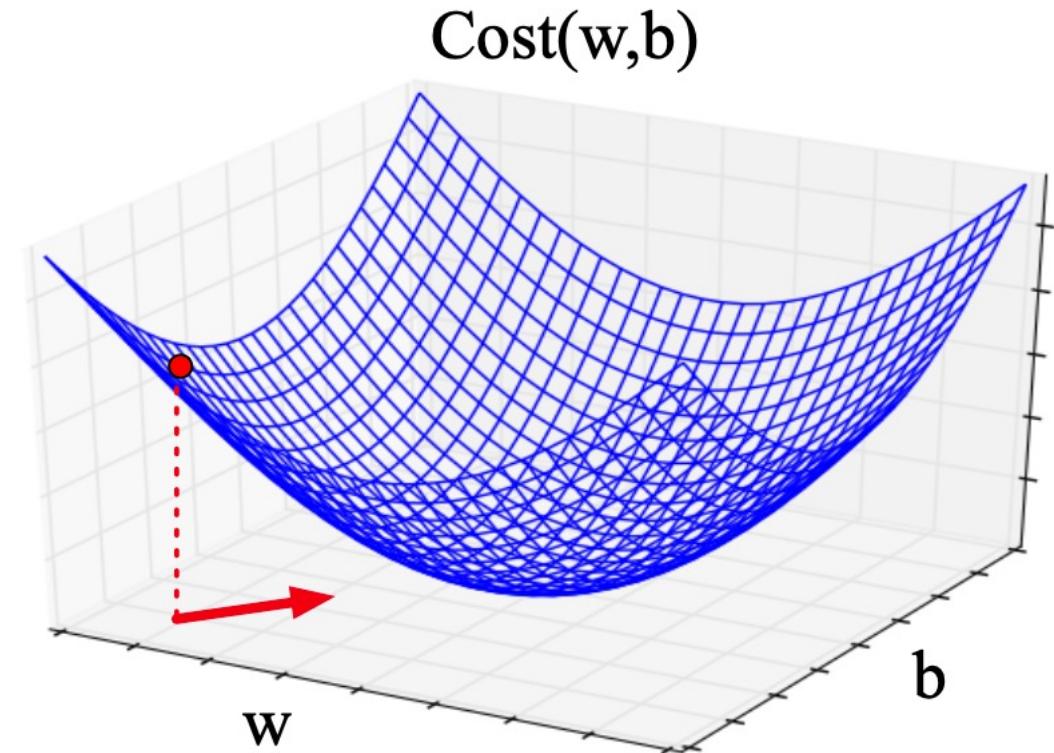
## ❖ Logistic Regression objective

Our objective is to find the optimal weights and bias (theta) that minimize the loss function:

$$\hat{\theta} = \arg\min_{\theta} \frac{1}{n} \sum_{i=1}^n L(\hat{y}^i, y^i)$$

For each step, compute the gradient and update the current theta:

- $\nabla_{\theta_s} L = \frac{1}{N} X^T (\hat{y}^i - y^i)$
- $\theta_{s+1} = \theta_s - \eta \nabla_{\theta_s} L$



Gradient Descent

# Review

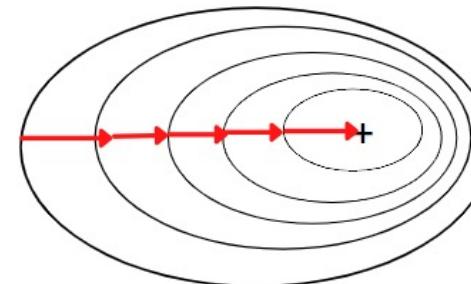
## ❖ Logistic Regression objective

For each step, compute the gradient and update the current theta:

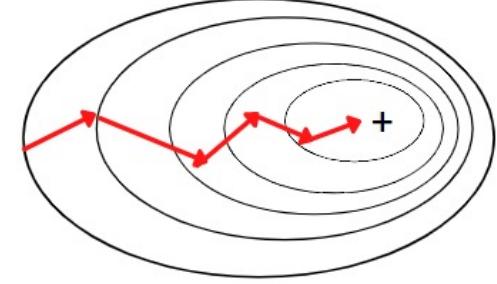
- $\nabla_{\theta_s} L = \frac{1}{N} X^T (\hat{y}^i - y^i)$
- $\theta_{s+1} = \theta_s - \eta \nabla_{\theta_s} L$

- **Stochastic GD:** batch\_size = 1
- **Mini-Batch GD:**  $1 < \text{batch\_size} < N$
- **Batch GD:** batch\_size = N

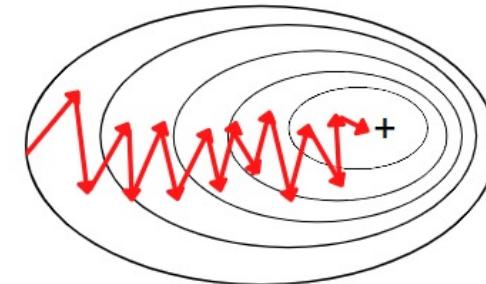
Batch Gradient Descent



Mini-Batch Gradient Descent



Stochastic Gradient Descent



# Review

## ❖ Training Logistic Regression

→ 1) Pick a sample  $(x, y)$  from training data

↓ 2) Compute output  $\hat{y}$

$$\downarrow z = \theta^T x = x^T \theta \quad \hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

3) Compute loss

$$\downarrow L(\hat{y}, y) = (-y \log \hat{y} - (1-y) \log(1-\hat{y}))$$

4) Compute derivative

$$\downarrow \nabla_{\theta} L = x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

$\eta$  is learning rate

```
def sigmoid_function(z):
    return 1 / (1 + np.exp(-z))

def predict(x, theta):
    return sigmoid_function(np.dot(x.T, theta))

def loss_function(y_hat, y):
    return -y*np.log(y_hat) - (1 - y)*np.log(1 - y_hat)

def compute_gradient(x, y_hat, y):
    return x*(y_hat - y)

def update(theta, lr, gradient):
    return theta - lr*gradient

# Given x and y
# compute output
y_hat = predict(x, theta)

# compute loss
loss = loss_function(y_hat, y)

# compute mean of gradient
gradient = compute_gradient(x, y_hat, y)

# update
theta = update(theta, lr, gradient)
```

**Dataset**

| Petal_Length | Petal_Width | Label |
|--------------|-------------|-------|
| 1.4          | 0.2         | 0     |
| 1.5          | 0.2         | 0     |
| 3            | 1.1         | 1     |
| 4.1          | 1.3         | 1     |

→ 1) Pick a sample  $(x, y)$  from training data

↓ 2) Compute output  $\hat{y}$

$$\downarrow z = \theta^T x = x^T \theta \quad \hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

3) Compute loss

$$\downarrow L(\hat{y}, y) = (-y \log \hat{y} - (1-y) \log(1-\hat{y}))$$

4) Compute derivative

$$\downarrow \nabla_{\theta} L = x(\hat{y} - y)$$

5) Update parameters

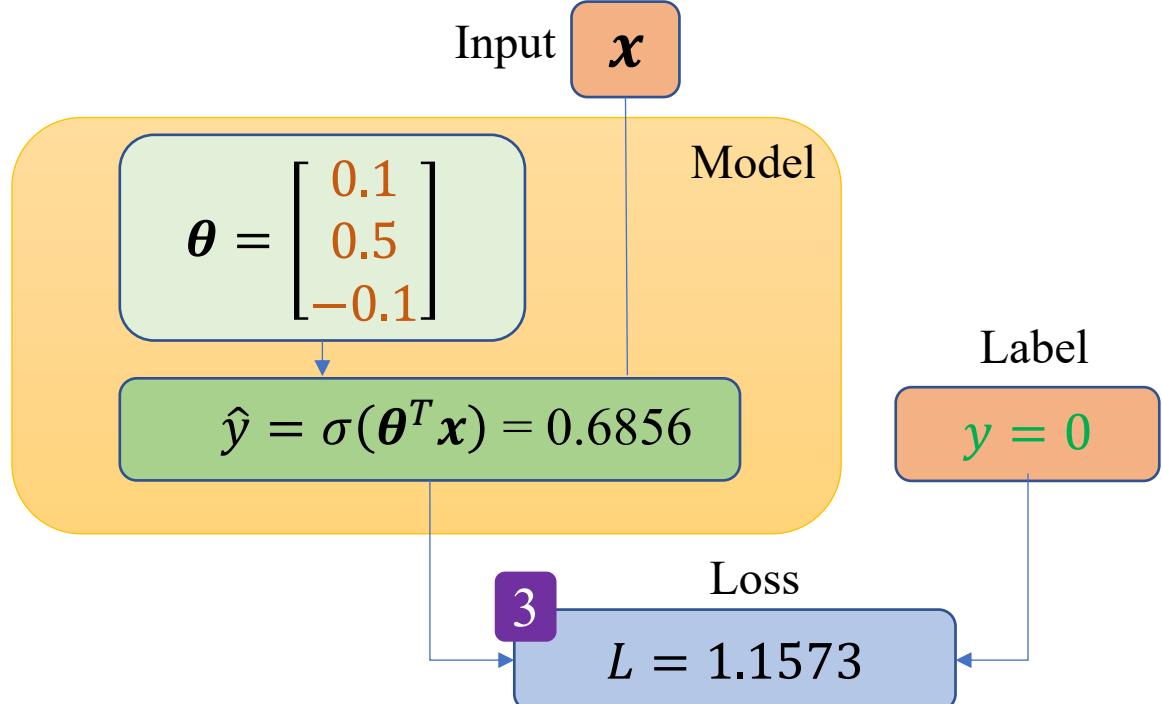
$$\theta = \theta - \eta \nabla_{\theta} L$$

1

$$x = \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1.4 \\ 0.2 \end{bmatrix}$$

Given  $\theta = \begin{bmatrix} b \\ w_1 \\ w_2 \end{bmatrix} = \begin{bmatrix} 0.1 \\ 0.5 \\ -0.1 \end{bmatrix}$

$$\eta = 0.01$$



4

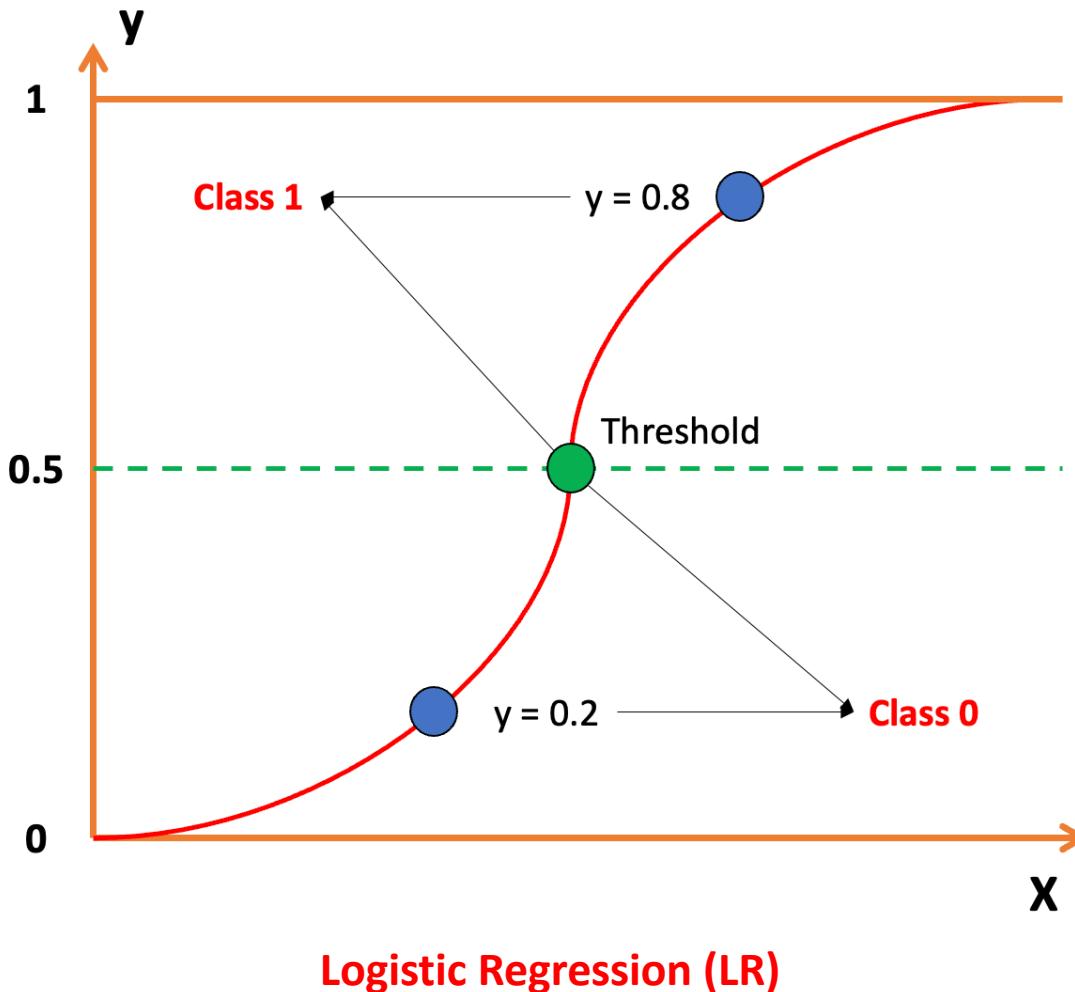
$$\nabla_{\theta} L = x(\hat{y} - y) = \begin{bmatrix} 1 \\ 1.4 \\ 0.2 \end{bmatrix} [0.6856] = \begin{bmatrix} 0.6856 \\ 0.9599 \\ 0.1371 \end{bmatrix} = \begin{bmatrix} L'_b \\ L'_{w_1} \\ L'_{w_2} \end{bmatrix}$$

5

$$\theta - \eta L'_{\theta} = \begin{bmatrix} 0.1 \\ 0.5 \\ -0.1 \end{bmatrix} - \eta \begin{bmatrix} 0.6856 \\ 0.9599 \\ 0.1371 \end{bmatrix} = \begin{bmatrix} 0.093 \\ 0.499 \\ -0.101 \end{bmatrix}$$

# Review

## ❖ Logistic Regression



Sigmoid function

$$z = b_0 + \sum_{i=1}^{n\_features} b_i x_i$$
$$\text{sigmoid}(z) = \frac{1}{1 + e^{-z}}$$

Binary Cross Entropy:

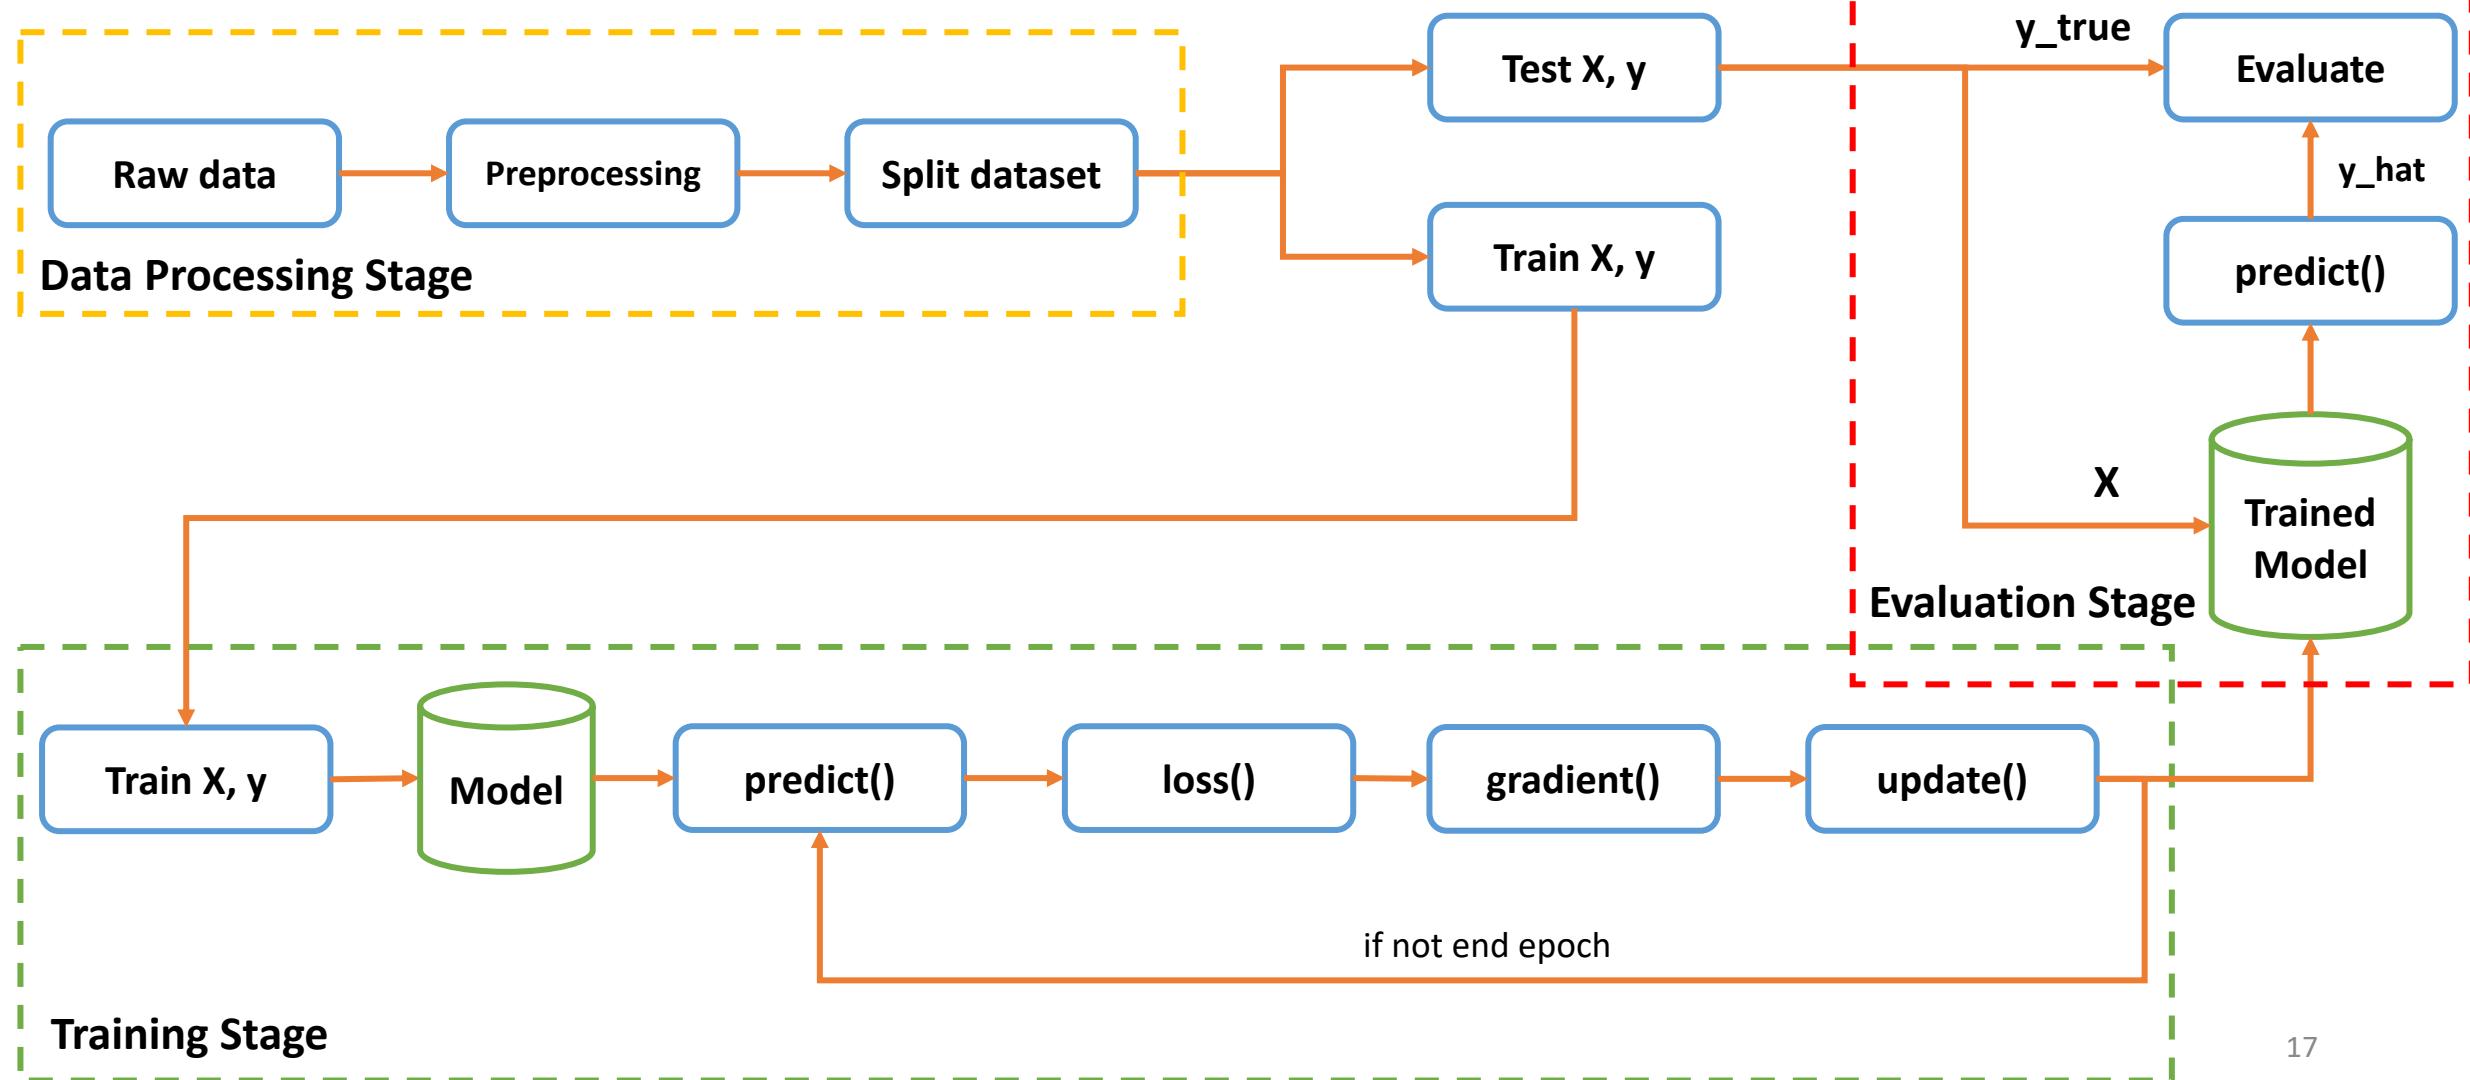
$$L(\hat{y}, y) = -\frac{1}{N} [y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})]$$

Gradient:

$$\nabla_{\theta} L = \frac{1}{N} X^T (\hat{y} - y)$$

# Review

## ❖ Logistic Regression Implementation Pipeline



# Titanic Survival Prediction

## ❖ Introduction

**Description:** Given [Titanic Survival dataset](#), build a Logistic Regression model to determine whether a passenger was survived or not in the Titanic incident.

| PassengerId | Pclass | Sex | Age | SibSp | Parch | Fare    | Embarked | Title | Survived |
|-------------|--------|-----|-----|-------|-------|---------|----------|-------|----------|
| 1           | 3      | 0   | 22  | 1     | 0     | 7.25    | 0        | 0     | 0        |
| 2           | 1      | 1   | 38  | 1     | 0     | 71.2833 | 1        | 1     | 1        |
| 3           | 3      | 1   | 26  | 0     | 0     | 7.925   | 0        | 2     | 1        |
| 4           | 1      | 1   | 35  | 1     | 0     | 53.1    | 0        | 1     | 1        |
| 5           | 3      | 0   | 35  | 0     | 0     | 8.05    | 0        | 0     | 0        |
| 6           | 3      | 0   | 28  | 0     | 0     | 8.4583  | 2        | 0     | 0        |
| 7           | 1      | 0   | 54  | 0     | 0     | 51.8625 | 0        | 0     | 0        |
| 8           | 3      | 0   | 2   | 3     | 1     | 21.075  | 0        | 3     | 0        |
| 9           | 3      | 1   | 27  | 0     | 2     | 11.1333 | 0        | 1     | 1        |

# Titanic Survival Prediction

## ❖ Dataset information

| PassengerId | Pclass | Sex | Age | SibSp | Parch | Fare    | Embarked | Title | Survived |
|-------------|--------|-----|-----|-------|-------|---------|----------|-------|----------|
| 1           | 3      | 0   | 22  | 1     | 0     | 7.25    | 0        | 0     | 0        |
| 2           | 1      | 1   | 38  | 1     | 0     | 71.2833 | 1        | 1     | 1        |
| 3           | 3      | 1   | 26  | 0     | 0     | 7.925   | 0        | 2     | 1        |
| 4           | 1      | 1   | 35  | 1     | 0     | 53.1    | 0        | 1     | 1        |
| 5           | 3      | 0   | 35  | 0     | 0     | 8.05    | 0        | 0     | 0        |
| 6           | 3      | 0   | 28  | 0     | 0     | 8.4583  | 2        | 0     | 0        |
| 7           | 1      | 0   | 54  | 0     | 0     | 51.8625 | 0        | 0     | 0        |
| 8           | 3      | 0   | 2   | 3     | 1     | 21.075  | 0        | 3     | 0        |
| 9           | 3      | 1   | 27  | 0     | 2     | 11.1333 | 0        | 1     | 1        |

- Number of samples: 891
- Number of features: 8
- Number of classes: 2

| Feature     | Description   |
|-------------|---|
| PassengerID | Dataset Index   |
| Pclass      | Ticket class  |
| Sex         | Sex   |
| Age         | Age in years  |
| SibSp       | Number of passenger's siblings / spouses aboard the Titanic |
| Parch       | Number of passenger's parents / children aboard the Titanic |
| Fare        | Passenger fare  |
| Embarked    | Port of Embarkation   |
| Title       | Passenger title   |
| Survived    | Survival  |

# Titanic Survival Prediction

## ❖ Step 1: Import libraries and read dataset

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 from sklearn.model_selection import train_test_split
6 from sklearn.preprocessing import StandardScaler
```

```
1 dataset_path = 'titanic_modified_dataset.csv'
2 df = pd.read_csv(
3     dataset_path,
4     index_col='PassengerId'
5 )
6 df
```

|             | Pclass | Sex | Age  | SibSp | Parch | Fare    | Embarked | Title | Survived |
|-------------|--------|-----|------|-------|-------|---------|----------|-------|----------|
| PassengerId |        |     |      |       |       |         |          |       |          |
| 1           | 3      | 0   | 22.0 | 1     | 0     | 7.2500  | 0        | 0     | 0        |
| 2           | 1      | 1   | 38.0 | 1     | 0     | 71.2833 | 1        | 1     | 1        |
| 3           | 3      | 1   | 26.0 | 0     | 0     | 7.9250  | 0        | 2     | 1        |
| 4           | 1      | 1   | 35.0 | 1     | 0     | 53.1000 | 0        | 1     | 1        |
| 5           | 3      | 0   | 35.0 | 0     | 0     | 8.0500  | 0        | 0     | 0        |
| ...         | ...    | ... | ...  | ...   | ...   | ...     | ...      | ...   | ...      |
| 887         | 2      | 0   | 27.0 | 0     | 0     | 13.0000 | 0        | 5     | 0        |
| 888         | 1      | 1   | 19.0 | 0     | 0     | 30.0000 | 0        | 2     | 1        |
| 889         | 3      | 1   | 28.0 | 1     | 2     | 23.4500 | 0        | 2     | 0        |
| 890         | 1      | 0   | 26.0 | 0     | 0     | 30.0000 | 1        | 0     | 1        |
| 891         | 3      | 0   | 32.0 | 0     | 0     | 7.7500  | 2        | 0     | 0        |

891 rows × 9 columns

# Titanic Survival Prediction

## ❖ Step 2: Get dataset information

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 891 entries, 1 to 891
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Pclass       891 non-null    int64  
 1   Sex          891 non-null    int64  
 2   Age          891 non-null    float64 
 3   SibSp        891 non-null    int64  
 4   Parch        891 non-null    int64  
 5   Fare          891 non-null    float64 
 6   Embarked     891 non-null    int64  
 7   Title         891 non-null    int64  
 8   Survived     891 non-null    int64  
dtypes: float64(2), int64(7)
memory usage: 69.6 KB
```

# Titanic Survival Prediction

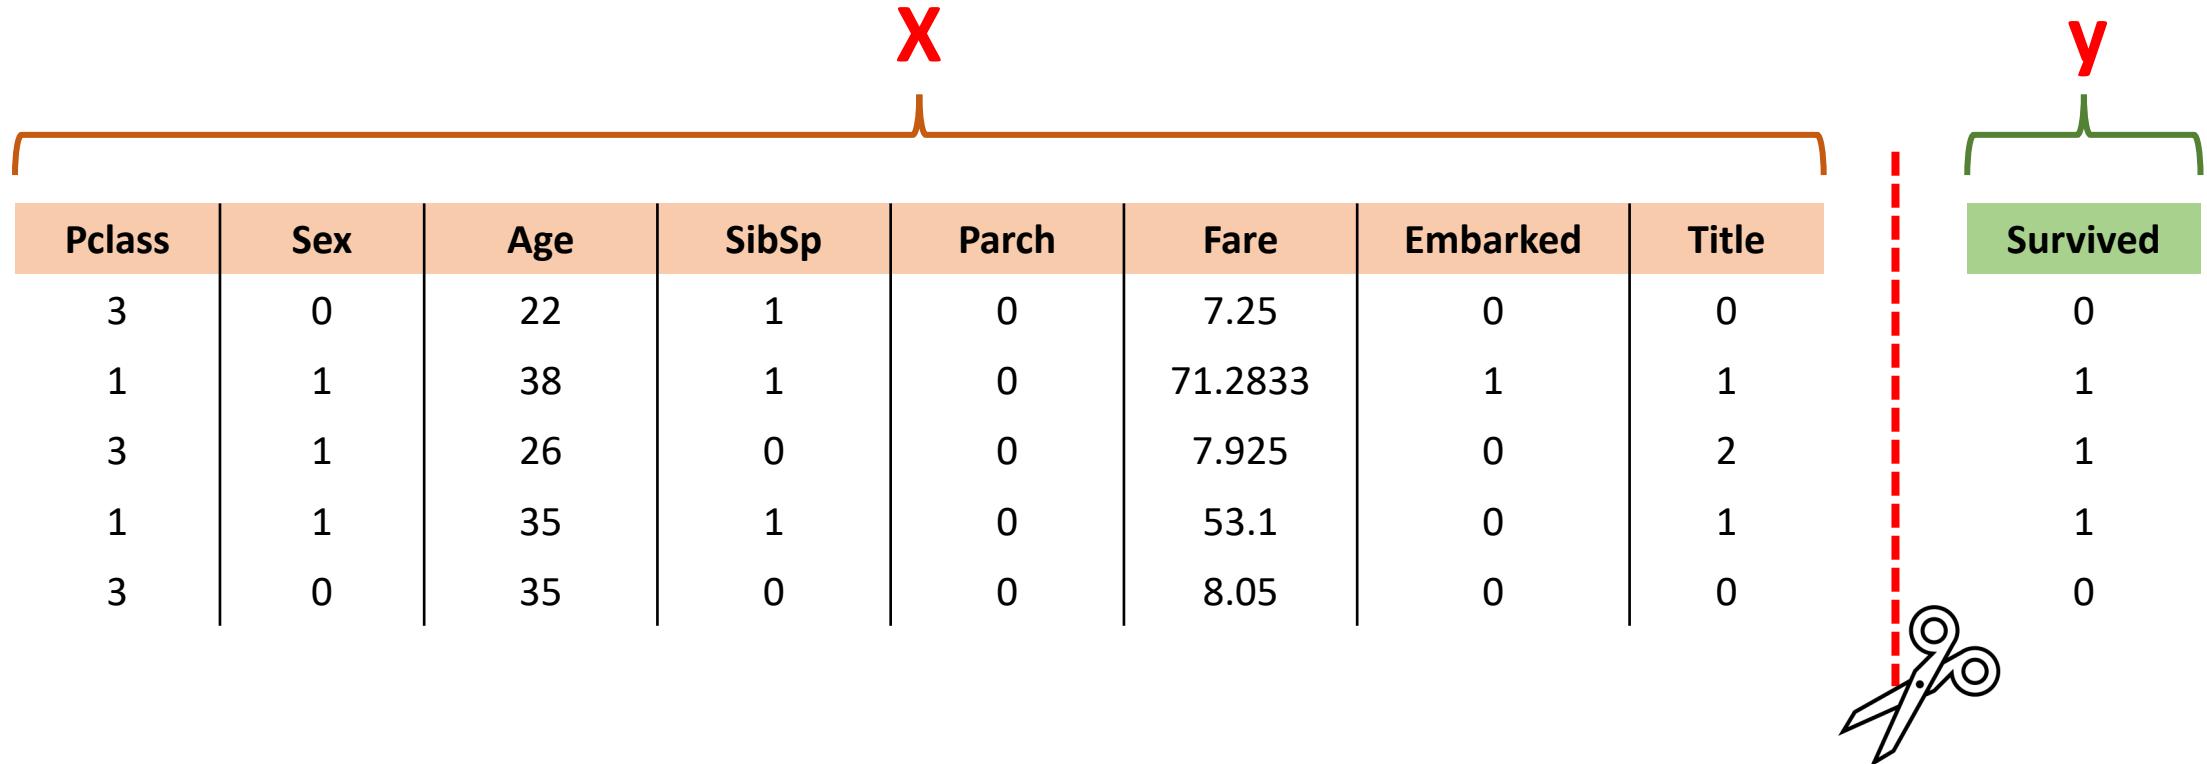
## ❖ Step 2: Get dataset information

```
1 df.describe()
```

|       | Pclass     | Sex        | Age        | SibSp      | Parch      | Fare       | Embarked   | Title      | Survived   |
|-------|------------|------------|------------|------------|------------|------------|------------|------------|------------|
| count | 891.000000 | 891.000000 | 891.000000 | 891.000000 | 891.000000 | 891.000000 | 891.000000 | 891.000000 | 891.000000 |
| mean  | 2.308642   | 0.352413   | 29.361582  | 0.523008   | 0.381594   | 32.204208  | 0.359147   | 0.936027   | 0.383838   |
| std   | 0.836071   | 0.477990   | 13.019697  | 1.102743   | 0.806057   | 49.693429  | 0.638707   | 1.725341   | 0.486592   |
| min   | 1.000000   | 0.000000   | 0.420000   | 0.000000   | 0.000000   | 0.000000   | -1.000000  | 0.000000   | 0.000000   |
| 25%   | 2.000000   | 0.000000   | 22.000000  | 0.000000   | 0.000000   | 7.910400   | 0.000000   | 0.000000   | 0.000000   |
| 50%   | 3.000000   | 0.000000   | 28.000000  | 0.000000   | 0.000000   | 14.454200  | 0.000000   | 0.000000   | 0.000000   |
| 75%   | 3.000000   | 1.000000   | 35.000000  | 1.000000   | 0.000000   | 31.000000  | 1.000000   | 2.000000   | 1.000000   |
| max   | 3.000000   | 1.000000   | 80.000000  | 8.000000   | 6.000000   | 512.329200 | 2.000000   | 16.000000  | 1.000000   |

# Titanic Survival Prediction

## ❖ Step 3: Split X, y



**Split raw dataset to X, y**

- ❖ X: first 8 columns (except PassengerID)
- ❖ y: last column

```
dataset = df.to_numpy().astype(np.float64)
```

```
x, y = dataset[:, :-1], dataset[:, -1]
```

# Titanic Survival Prediction

## ❖ Step 3: Split X, y (add intercepts)

$$z = b_0 + \sum_{i=1}^{n\_features} b_i x_i$$

Currently our X only has n\_features, thus we need to add b (intercept)

```
1 print(X.shape)      #   Column
---  -----
(891, 8)           0 Pclass
                  1 Sex
                  2 Age
                  3 SibSp
                  4 Parch
                  5 Fare
                  6 Embarked
                  7 Title
```

shape[1] == 8 means 8 features

```
1 intercept = np.ones((X.shape[0], 1))
2
3 )
4 X_b = np.concatenate(
5     (intercept, X),
6     axis=1
7 )
```

```
1 intercept      1 X
array([[1.        , 3.        , 0.        ,
       [1.        , 1.        , 1.        ,
       [1.        , 3.        , 1.        ,
       ...,
       [1.        , ..., 1.        ,
       [1.        , 1.        , 0.        ,
       [1.        , 3.        , 0.        ]])]
```

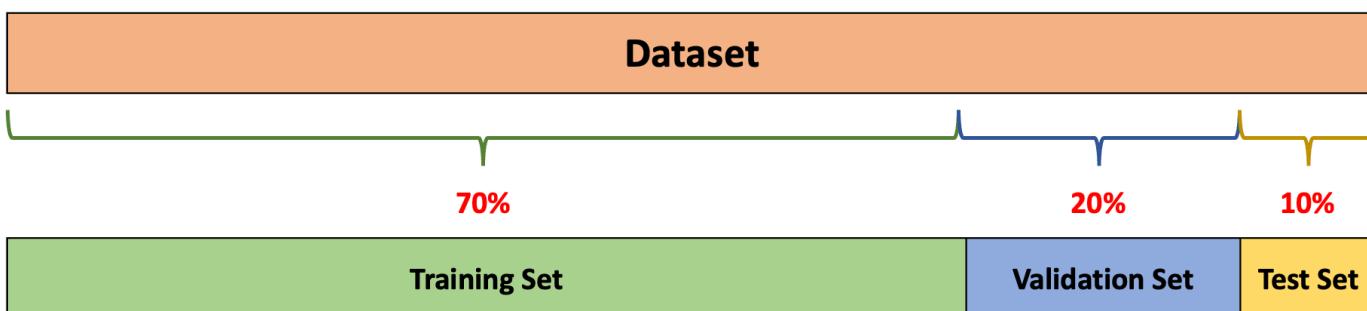
New shape of X after adding intercepts:

```
1 print(X_b.shape)
(891, 9)
```

```
1 X_b
array([[ 1.        ,  3.        ,  ...,  1.        ,
       [ 1.        ,  1.        ,  ...,  3.        ,
       [ 1.        ,  3.        ,  ...,  1.        ,
       ...,
       [ 1.        ,  1.        ,  ...,  3.        ,
       [ 1.        ,  1.        ,  ...,  1.        ,
       [ 1.        ,  3.        ,  ...,  3.        ]])]
```

# Titanic Survival Prediction

## ❖ Step 4: Split train, val, test set



```
1 val_size = 0.2
2 test_size = 0.125
3 random_state = 2
4 is_shuffle = True
5
6 X_train, X_val, y_train, y_val = train_test_split(
7     X_b, y,
8     test_size=val_size,
9     random_state=random_state,
10    shuffle=is_shuffle
11 )
12
13 X_train, X_test, y_train, y_test = train_test_split(
14     X_train, y_train,
15     test_size=test_size,
16     random_state=random_state,
17     shuffle=is_shuffle
18 )
```

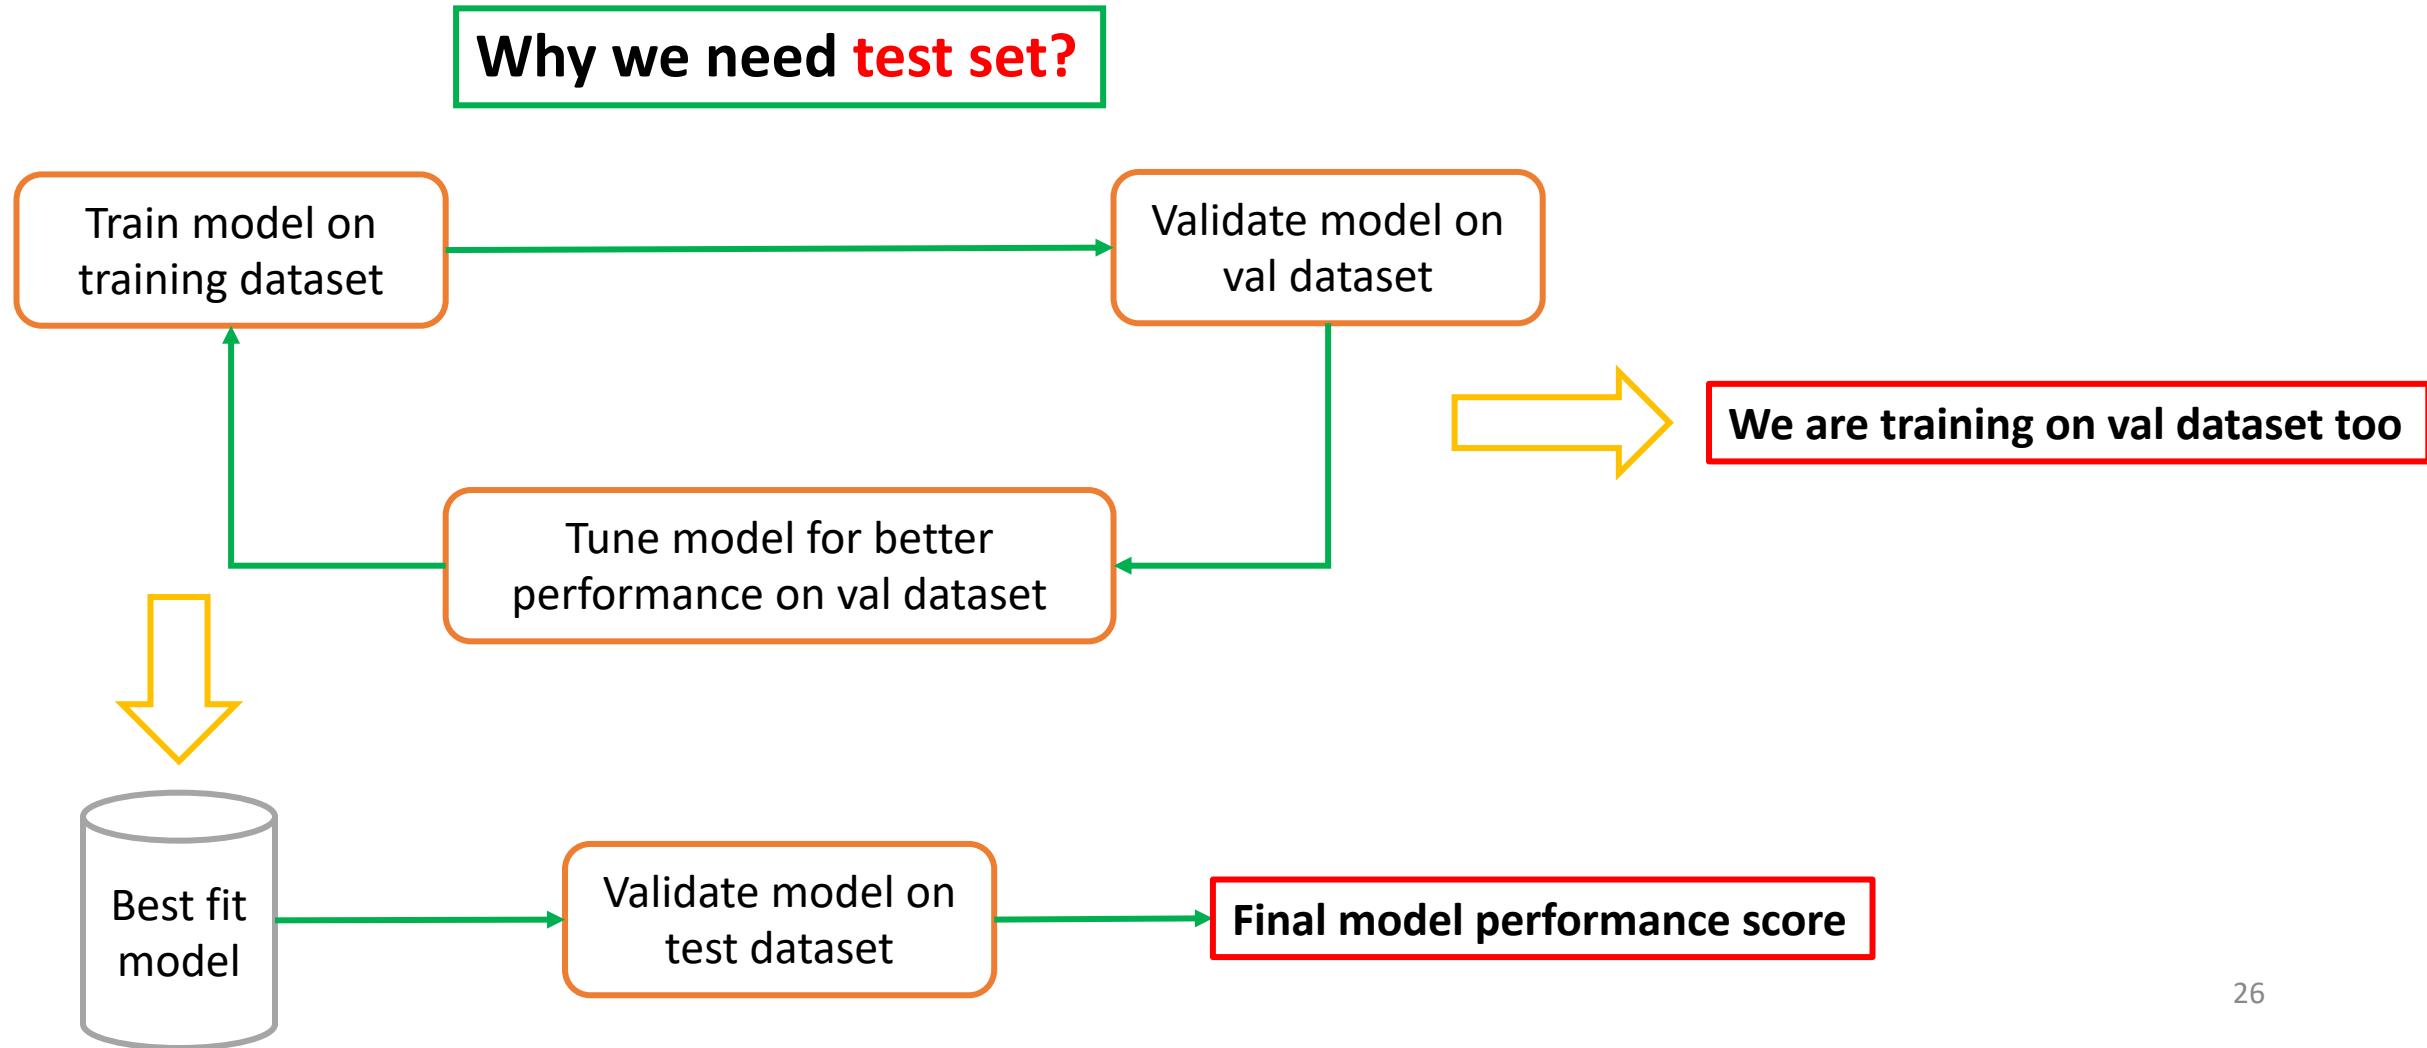
  

```
1 print(f'Number of training samples: {X_train.shape[0]}')
2 print(f'Number of val samples: {X_val.shape[0]}')
3 print(f'Number of test samples: {X_test.shape[0]}')
```

Number of training samples: 623  
Number of val samples: 179  
Number of test samples: 89

# Titanic Survival Prediction

## ❖ Step 4: Split train, val, test set



# Titanic Survival Prediction

## ❖ Step 5: Normalization

Using `sklearn.preprocessing.StandardScaler()` to scale all values in dataset.

$$z = \frac{x_i - \mu}{\sigma}$$

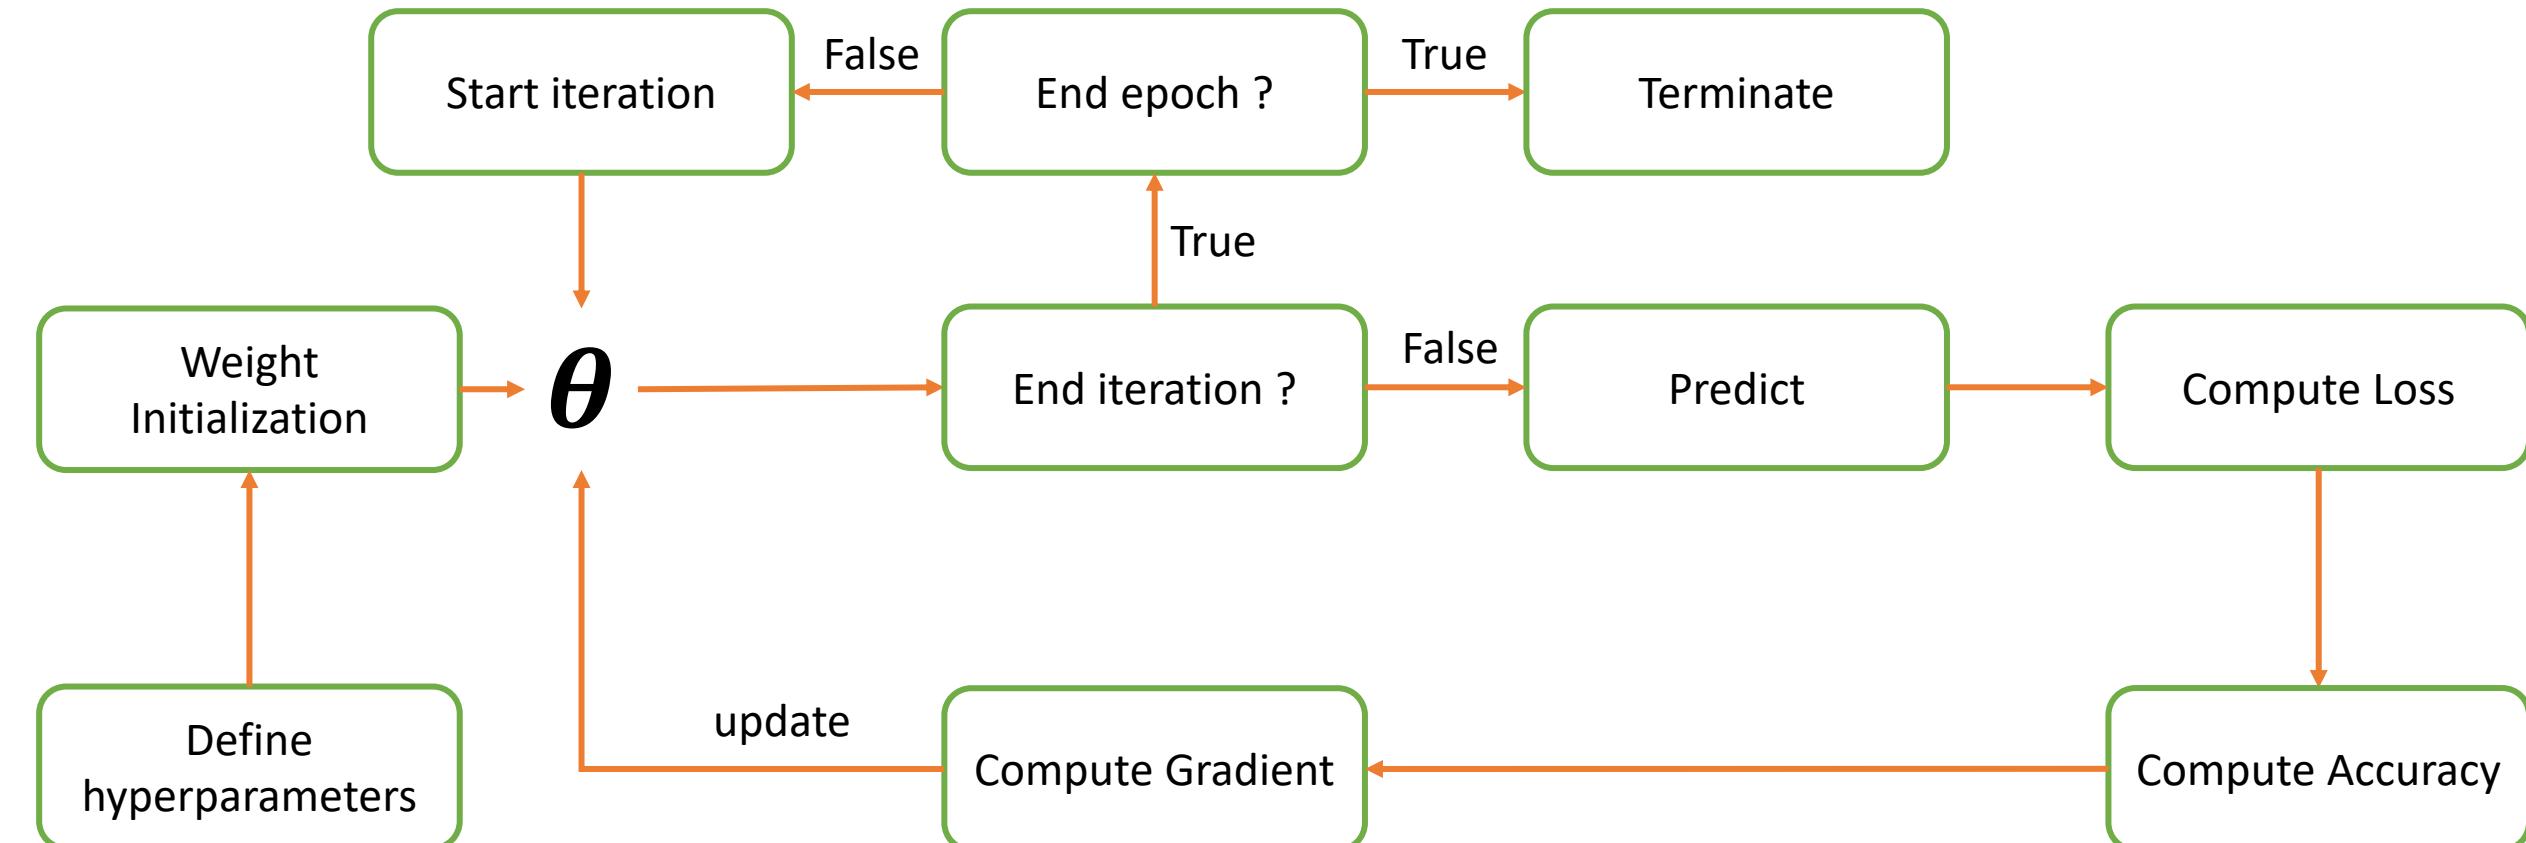
```
1 normalizer = StandardScaler()  
2 X_train = normalizer.fit_transform(X_train)  
3 X_val = normalizer.transform(X_val)  
4 X_test = normalizer.transform(X_test)
```

```
1 X_train  
  
array([[ 0.82313781, -0.75449952, -0.09246773, ..., -0.5944735 ,  
       -0.56180558,  0.          ],  
      [-0.37502774,  1.3253819 , -0.40715477, ..., -0.5944735 ,  
       0.64719226,  0.          ],  
      [ 0.82313781, -0.75449952, -0.80051356, ..., -0.5944735 ,  
       -0.56180558,  0.          ],  
      ...,  
      [-1.57319329, -0.75449952,  3.29041791, ...,  0.91104272,  
       -0.56180558,  0.          ],  
      [-1.57319329, -0.75449952, -0.09246773, ..., -0.5944735 ,  
       -0.56180558,  0.          ],  
      [-1.57319329, -0.75449952, -0.09246773, ..., -0.5944735 ,  
       -0.56180558,  0.          ]])
```

**Note:** We only use the train set to fit the scaler.

# Titanic Survival Prediction

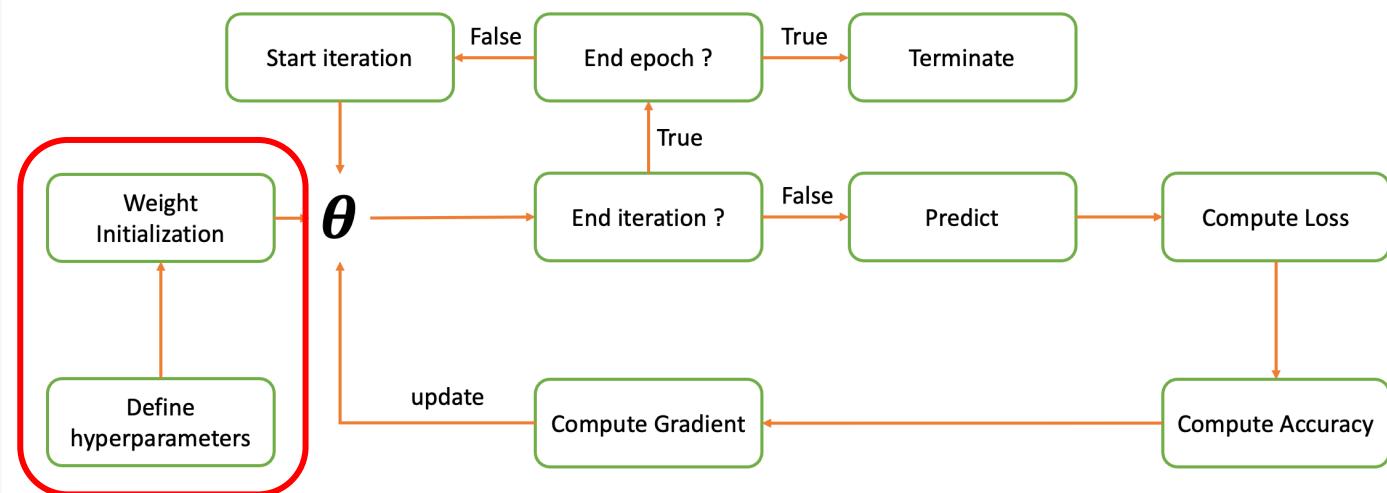
## ❖ Step 6: Training



# Titanic Survival Prediction

## ❖ Step 6: Training

```
1 lr = 0.01
2 epochs = 100
3 batch_size = 16
4
5 np.random.seed(random_state)
6 theta = np.random.uniform(
7     size=x_train.shape[1]
8 )
```



Define essential hyperparameters and initialize weights

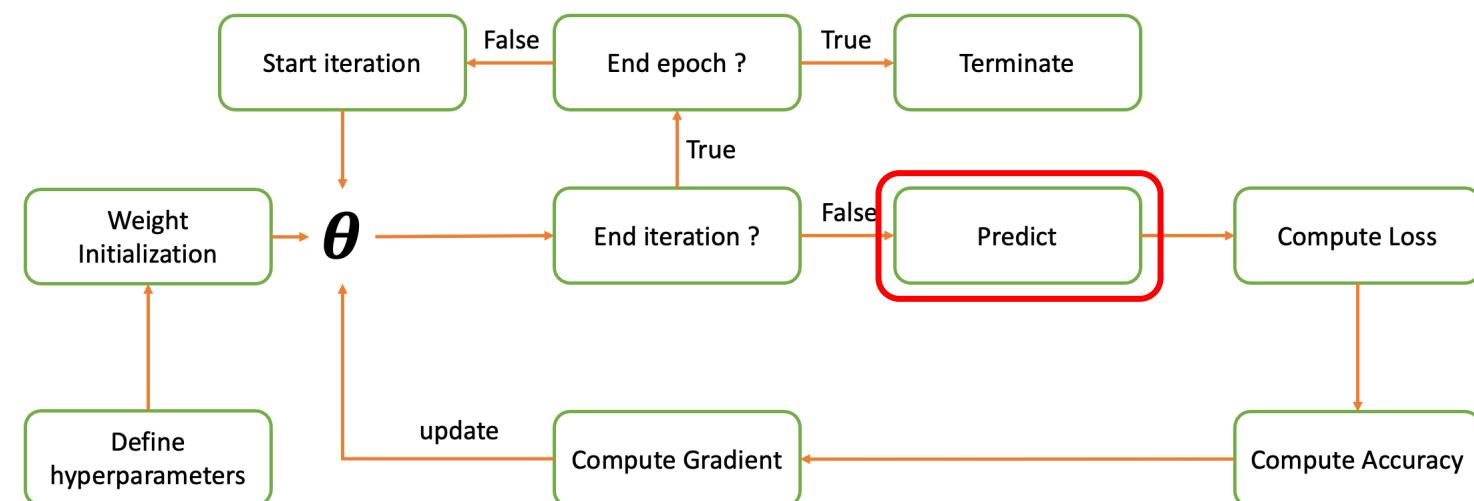
# Titanic Survival Prediction

## ❖ Step 6: Training

### Hypothesis function

- $z = b_0 + \sum_{i=1}^{n\_features} b_i x_i$
- $y_{\text{hat}} = h(z) = \text{sigmoid}(z) = \frac{1}{1+e^{-z}}$

```
1 def sigmoid(z):  
2     return 1 / (1 + np.exp(-z))  
3  
4  
5 def predict(X, theta):  
6     dot_product = np.dot(X, theta)  
7     y_hat = sigmoid(dot_product)  
8  
9     return y_hat
```



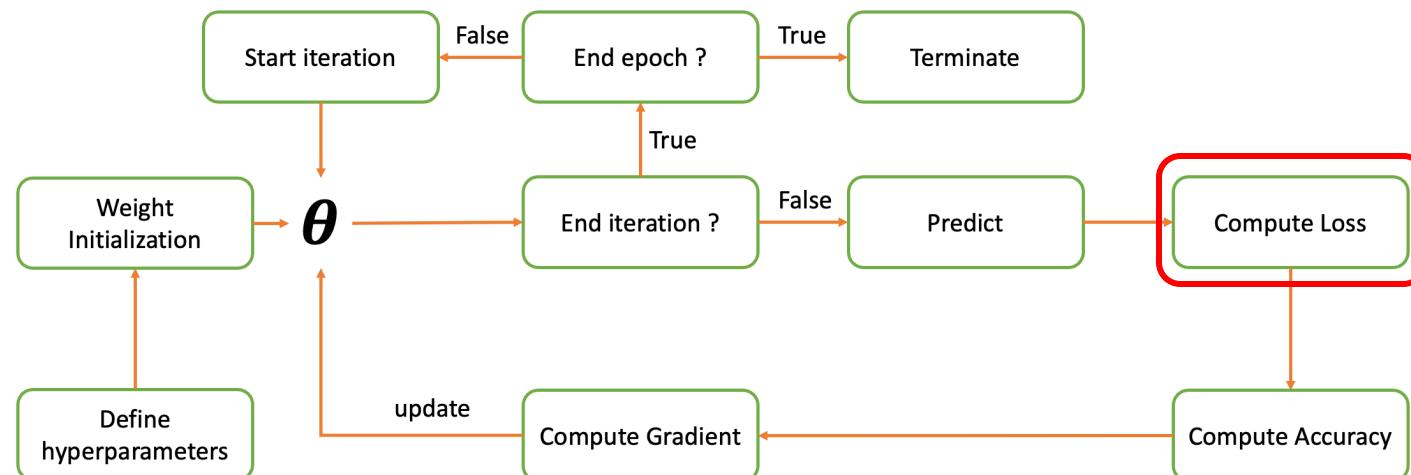
# Titanic Survival Prediction

## ❖ Step 6: Training

### Logistic Regression Loss (Cross-entropy)

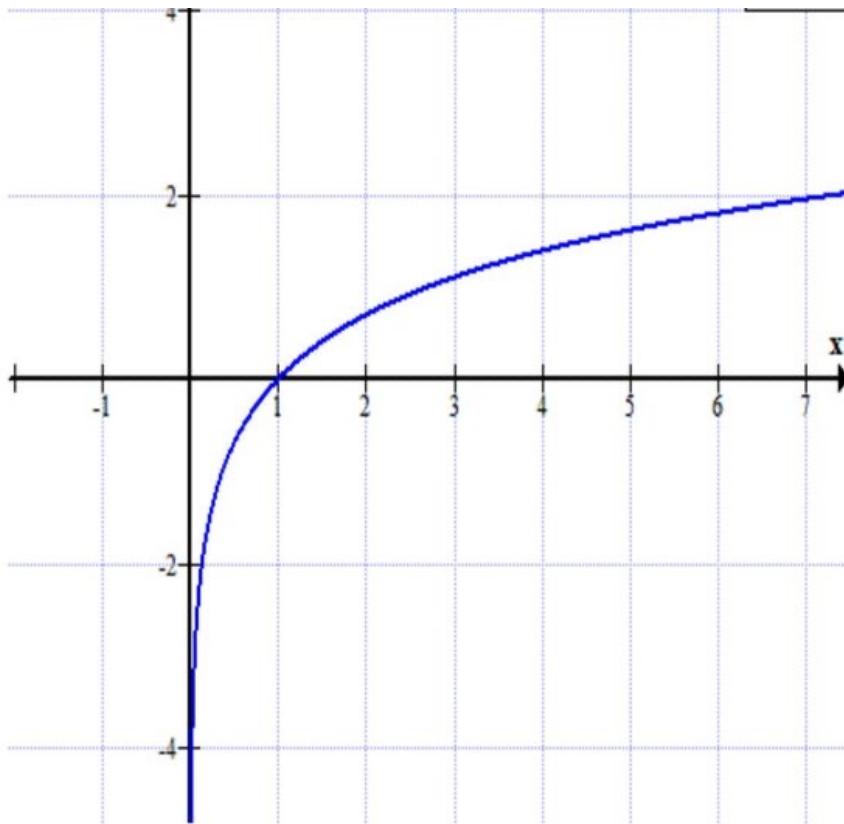
$$L(\hat{y}, y) = -\frac{1}{N} \sum_{i=1}^N y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)$$

```
1 def loss(y_hat, y):
2     y_hat = np.clip(
3         y_hat, 1e-7, 1 - 1e-7
4     )
5
6     return (
7         -y * \
8             np.log(y_hat) - (1 - y) * \
9                 np.log(1 - y_hat)
10    ).mean()
```



# Titanic Survival Prediction

## ❖ Step 6: Training



→ Can cause nan if log(0)

```
1 def loss(y_hat, y):
2     y_hat = np.clip(
3         y_hat, 1e-7, 1 - 1e-7
4     )
5
6     return (
7         -y * \
8             np.log(y_hat) - (1 - y) * \
9                 np.log(1 - y_hat)
10    ).mean()
```

Avoid prediction values too small by limiting the value range

# Titanic Survival Prediction

## ❖ Step 6: Training

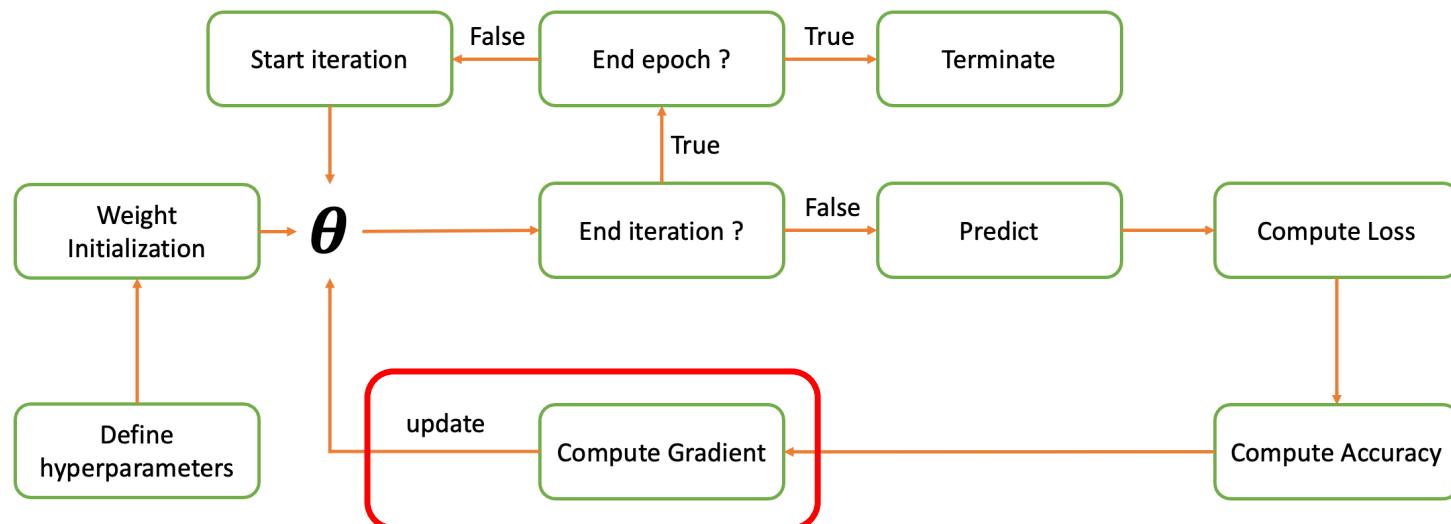
- Gradient computation formula:

$$\nabla_{\theta} L = \frac{1}{N} X^T (\hat{y} - y)$$

- Weights update formula:

$$\theta = \theta - \eta \nabla_{\theta} L$$

```
1 def compute_gradient(X, y, y_hat):  
2     return np.dot(  
3         X.T, (y_hat - y))  
4     ) / y.size  
5  
6 def update_theta(theta, gradient, lr):  
7     return theta - lr * gradient
```



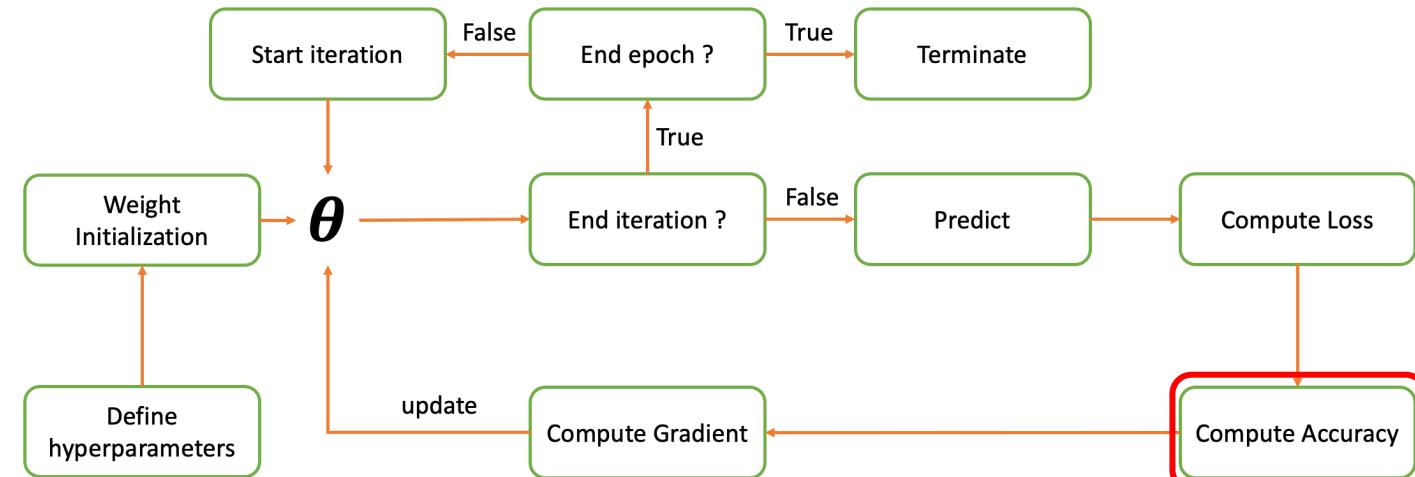
# Titanic Survival Prediction

## ❖ Step 6: Training

### Accuracy Formula

$$\text{accuracy} = \frac{\text{true\_predictions}}{\text{n\_samples}}$$

```
1 def compute_accuracy(X, y, theta):  
2     y_hat = predict(X, theta).round()  
3     acc = (y_hat == y).mean()  
4  
5     return acc
```



# Titanic Survival Prediction

## ❖ Step 6: Training

Line 6: Loop over number of epochs.

Line 7, 8, 9, 10: Declare empty lists to store batch accuracies, losses of train and val sets.

```
6 for epoch in range(epochs):
7     train_batch_losses = []
8     train_batch_accs = []
9     val_batch_losses = []
10    val_batch_accs = []

11
12    for i in range(0, X_train.shape[0], batch_size):
13        X_i = X_train[i:i+batch_size]
14        y_i = y_train[i:i+batch_size]

15
16        y_hat = predict(X_i, theta)
17        train_loss = compute_loss(y_hat, y_i)
18        gradient = compute_gradient(X_i, y_i, y_hat)
19        theta = update_theta(theta, gradient, lr)
20        train_batch_losses.append(train_loss)

21
22        train_acc = compute_accuracy(X_train, y_train, theta)
23        train_batch_accs.append(train_acc)

24
25        y_val_hat = predict(X_val, theta)
26        val_loss = compute_loss(y_val_hat, y_val)
27        val_batch_losses.append(val_loss)

28
29        val_acc = compute_accuracy(X_val, y_val, theta)
30        val_batch_accs.append(val_acc)

31
32        train_batch_loss = sum(train_batch_losses) / len(train_batch_losses)
33        val_batch_loss = sum(val_batch_losses) / len(val_batch_losses)
34        train_batch_acc = sum(train_batch_accs) / len(train_batch_accs)
35        val_batch_acc = sum(val_batch_accs) / len(val_batch_accs)
```

# Titanic Survival Prediction

## ❖ Step 6: Training

Line 12: Loop over number of batches (based on batch size).

Line 13, 14: Get X and y data of current batch.

```
6 for epoch in range(epochs):
7     train_batch_losses = []
8     train_batch_accs = []
9     val_batch_losses = []
10    val_batch_accs = []

11
12    for i in range(0, X_train.shape[0], batch_size):
13        X_i = X_train[i:i+batch_size]
14        y_i = y_train[i:i+batch_size]

15
16        y_hat = predict(X_i, theta)
17        train_loss = compute_loss(y_hat, y_i)
18        gradient = compute_gradient(X_i, y_i, y_hat)
19        theta = update_theta(theta, gradient, lr)
20        train_batch_losses.append(train_loss)

21
22        train_acc = compute_accuracy(X_train, y_train, theta)
23        train_batch_accs.append(train_acc)

24
25        y_val_hat = predict(X_val, theta)
26        val_loss = compute_loss(y_val_hat, y_val)
27        val_batch_losses.append(val_loss)

28
29        val_acc = compute_accuracy(X_val, y_val, theta)
30        val_batch_accs.append(val_acc)

31
32        train_batch_loss = sum(train_batch_losses) / len(train_batch_losses)
33        val_batch_loss = sum(val_batch_losses) / len(val_batch_losses)
34        train_batch_acc = sum(train_batch_accs) / len(train_batch_accs)
35        val_batch_acc = sum(val_batch_accs) / len(val_batch_accs)
```

# Titanic Survival Prediction

## ❖ Step 6: Training

**Line 16:** Do prediction on  $X_i$  and theta.

**Line 17:** Compute the loss of  $y_{\hat{}}$  and  $y_i$ .

**Line 18:** Compute the gradient.

**Line 19:** Update theta using the computed gradient.

```
6 for epoch in range(epochs):
7     train_batch_losses = []
8     train_batch_accs = []
9     val_batch_losses = []
10    val_batch_accs = []

11
12    for i in range(0, X_train.shape[0], batch_size):
13        X_i = X_train[i:i+batch_size]
14        y_i = y_train[i:i+batch_size]

15
16    y_hat = predict(X_i, theta)
17    train_loss = compute_loss(y_hat, y_i)
18    gradient = compute_gradient(X_i, y_i, y_hat)
19    theta = update_theta(theta, gradient, lr)
20    train_batch_losses.append(train_loss)

21
22    train_acc = compute_accuracy(X_train, y_train, theta)
23    train_batch_accs.append(train_acc)

24
25    y_val_hat = predict(X_val, theta)
26    val_loss = compute_loss(y_val_hat, y_val)
27    val_batch_losses.append(val_loss)

28
29    val_acc = compute_accuracy(X_val, y_val, theta)
30    val_batch_accs.append(val_acc)

31
32    train_batch_loss = sum(train_batch_losses) / len(train_batch_losses)
33    val_batch_loss = sum(val_batch_losses) / len(val_batch_losses)
34    train_batch_acc = sum(train_batch_accs) / len(train_batch_accs)
35    val_batch_acc = sum(val_batch_accs) / len(val_batch_accs)
```

# Titanic Survival Prediction

## ❖ Step 6: Training

**From line 20 to line 30:** Compute and store accuracies and losses on train and val sets.

```
6 for epoch in range(epochs):
7     train_batch_losses = []
8     train_batch_accs = []
9     val_batch_losses = []
10    val_batch_accs = []

11
12    for i in range(0, X_train.shape[0], batch_size):
13        X_i = X_train[i:i+batch_size]
14        y_i = y_train[i:i+batch_size]

15
16        y_hat = predict(X_i, theta)
17        train_loss = compute_loss(y_hat, y_i)
18        gradient = compute_gradient(X_i, y_i, y_hat)
19        theta = update_theta(theta, gradient, lr)
20        train_batch_losses.append(train_loss)

21
22        train_acc = compute_accuracy(X_train, y_train, theta)
23        train_batch_accs.append(train_acc)

24
25        y_val_hat = predict(X_val, theta)
26        val_loss = compute_loss(y_val_hat, y_val)
27        val_batch_losses.append(val_loss)

28
29        val_acc = compute_accuracy(X_val, y_val, theta)
30        val_batch_accs.append(val_acc)

31
32        train_batch_loss = sum(train_batch_losses) / len(train_batch_losses)
33        val_batch_loss = sum(val_batch_losses) / len(val_batch_losses)
34        train_batch_acc = sum(train_batch_accs) / len(train_batch_accs)
35        val_batch_acc = sum(val_batch_accs) / len(val_batch_accs)
```

# Titanic Survival Prediction

## ❖ Step 6: Training

**From line 32 to line 35:** Compute and store batch accuracies and losses of train and val sets.

```
6 for epoch in range(epochs):
7     train_batch_losses = []
8     train_batch_accs = []
9     val_batch_losses = []
10    val_batch_accs = []

11
12    for i in range(0, X_train.shape[0], batch_size):
13        X_i = X_train[i:i+batch_size]
14        y_i = y_train[i:i+batch_size]

15
16        y_hat = predict(X_i, theta)
17        train_loss = compute_loss(y_hat, y_i)
18        gradient = compute_gradient(X_i, y_i, y_hat)
19        theta = update_theta(theta, gradient, lr)
20        train_batch_losses.append(train_loss)

21
22        train_acc = compute_accuracy(X_train, y_train, theta)
23        train_batch_accs.append(train_acc)

24
25        y_val_hat = predict(X_val, theta)
26        val_loss = compute_loss(y_val_hat, y_val)
27        val_batch_losses.append(val_loss)

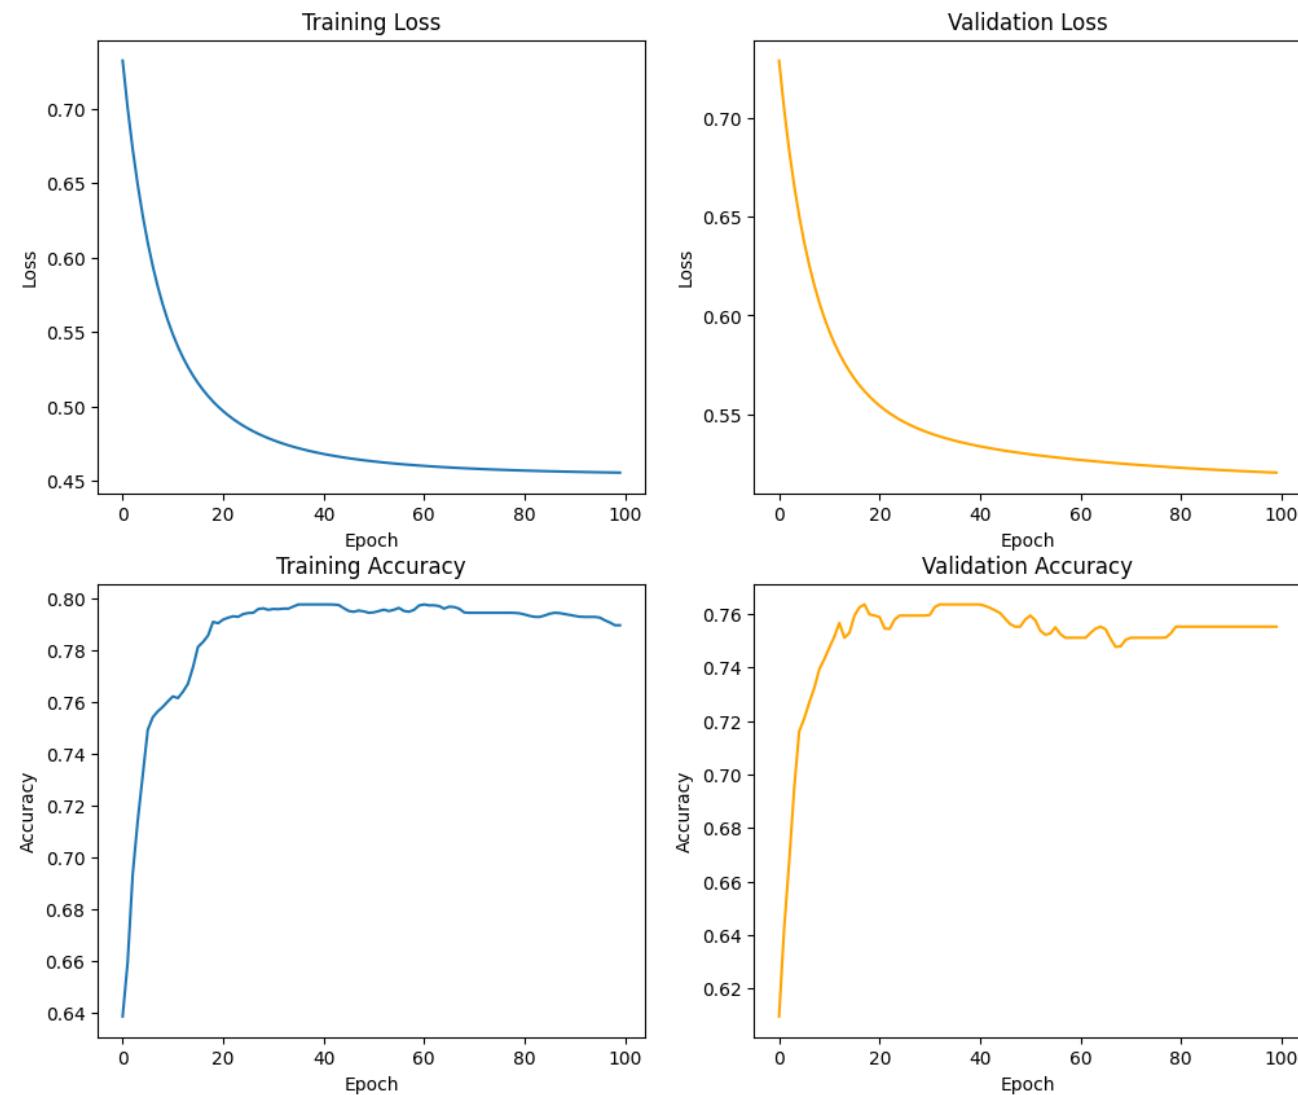
28
29        val_acc = compute_accuracy(X_val, y_val, theta)
30        val_batch_accs.append(val_acc)

31
32        train_batch_loss = sum(train_batch_losses) / len(train_batch_losses)
33        val_batch_loss = sum(val_batch_losses) / len(val_batch_losses)
34        train_batch_acc = sum(train_batch_accs) / len(train_batch_accs)
35        val_batch_acc = sum(val_batch_accs) / len(val_batch_accs)
```

# Titanic Survival Prediction

## ❖ Step 6: Training (Visualization)

```
1 fig, ax = plt.subplots(2, 2, figsize=(12, 10))
2 ax[0, 0].plot(train_losses)
3 ax[0, 0].set(xlabel='Epoch', ylabel='Loss')
4 ax[0, 0].set_title('Training Loss')
5
6 ax[0, 1].plot(val_losses, 'orange')
7 ax[0, 1].set(xlabel='Epoch', ylabel='Loss')
8 ax[0, 1].set_title('Validation Loss')
9
10 ax[1, 0].plot(train_accs)
11 ax[1, 0].set(xlabel='Epoch', ylabel='Accuracy')
12 ax[1, 0].set_title('Training Accuracy')
13
14 ax[1, 1].plot(val_accs, 'orange')
15 ax[1, 1].set(xlabel='Epoch', ylabel='Accuracy')
16 ax[1, 1].set_title('Validation Accuracy')
17
18 plt.show()
```



# Titanic Survival Prediction

## ❖ Step 7: Evaluation

$$\text{accuracy} = \frac{\text{true\_predictions}}{\text{n\_samples}}$$

```
1 def compute_accuracy(X, y, theta):  
2     y_hat = predict(X, theta).round()  
3     acc = (y_hat == y).mean()  
4  
5     return acc
```

```
1 # Val set  
2 val_set_acc = compute_accuracy(X_val, y_val, theta)  
3 print('Evaluation on validation set:')4 print(f'Accuracy: {val_set_acc}')
```

Evaluation on validation set:  
Accuracy: 0.770949720670391

```
1 # Test set  
2 test_set_acc = compute_accuracy(X_test, y_test, theta)  
3 print('Evaluation on test set:')4 print(f'Accuracy: {test_set_acc}')
```

Evaluation on test set:  
Accuracy: 0.7752808988764045

# Sentiment Analysis

## ❖ Introduction

**Description:** Given [Twitter Sentiment Analysis](#), build a Logistic Regression model to determine whether a tweet (text) has a positive sentiment or not.



Positive



Negative

# Sentiment Analysis

## ❖ Step 1: Import libraries

```
1 import pandas as pd
2 import numpy as np
3 import re
4 import nltk
5 import matplotlib.pyplot as plt
6
7 from sklearn.model_selection import train_test_split
8 from sklearn.preprocessing import StandardScaler
9 from nltk.tokenize import TweetTokenizer
10 from collections import defaultdict
```



# Sentiment Analysis

## ❖ Step 2: Read dataset

```
1 dataset_path = 'sentiment_analysis.csv'  
2 df = pd.read_csv(  
3     dataset_path,  
4     index_col='id'  
5 )  
6 df
```

|      | label | tweet   |
|------|-------|---|
|      | id    |   |
| 1    | 0     | #fingerprint #Pregnancy Test https://goo.gl/h1... |
| 2    | 0     | Finally a transparant silicon case ^^ Thanks t... |
| 3    | 0     | We love this! Would you go? #talk #makememorie... |
| 4    | 0     | I'm wired I know I'm George I was made that wa... |
| 5    | 1     | What amazing service! Apple won't even talk to... |
| ...  | ...   | ...   |
| 7916 | 0     | Live out loud #lol #liveoutloud #selfie #smile... |
| 7917 | 0     | We would like to wish you an amazing day! Make... |
| 7918 | 0     | Helping my lovely 90 year old neighbor with he... |
| 7919 | 0     | Finally got my #smart #pocket #wifi stay conne... |
| 7920 | 0     | Apple Barcelona!!! #Apple #Store #BCN #Barcelo... |

7920 rows × 2 columns

# Sentiment Analysis

## ❖ Step 3: Get dataset information

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 7920 entries, 1 to 7920
Data columns (total 2 columns):
 #   Column   Non-Null Count   Dtype  
--- 
 0   label     7920 non-null    int64  
 1   tweet     7920 non-null    object 
dtypes: int64(1), object(1)
memory usage: 185.6+ KB
```

```
1 df.describe()
```

|       | label       |
|-------|-------------|
| count | 7920.000000 |
| mean  | 0.255808    |
| std   | 0.436342    |
| min   | 0.000000    |
| 25%   | 0.000000    |
| 50%   | 0.000000    |
| 75%   | 1.000000    |
| max   | 1.000000    |

# Sentiment Analysis

## ❖ Step 4: Data preprocessing

| label | tweet   |
|-------|---|
| 0     | #fingerprint #Pregnancy Test https://goo.gl/h1... |
| 0     | Finally a transparent silicon case ^^ Thanks t... |
| 0     | We love this! Would you go? #talk #makememorie... |
| 0     | I'm wired I know I'm George I was made that wa... |
| 1     | What amazing service! Apple won't even talk to... |
| ...   | ...   |
| 0     | Live out loud #lol #liveoutloud #selfie #smile... |
| 0     | We would like to wish you an amazing day! Make... |
| 0     | Helping my lovely 90 year old neighbor with he... |
| 0     | Finally got my #smart #pocket #wifi stay conne... |
| 0     | Apple Barcelona!!! #Apple #Store #BCN #Barcelo... |

**Raw Text:** We love this! Would you go? #talk  
#makememories #unplug #relax #iphone #smartphone #wifi  
#connect... http://fb.me/6N3LsUpCu

Data  
Preprocessing

**Vector Representation:** array([1.0, 4768.0, 1425.0])

# Sentiment Analysis

## ❖ Step 4: Data preprocessing

**Raw Text:** We love this! Would you go? #talk  
#makememories #unplug #relax #iphone #smartphone #wifi  
#connect... http://fb.me/6N3LsUpCu

Data  
Preprocessing

**Vector Representation:** `array([1.0, 4768.0, 1425.0])`

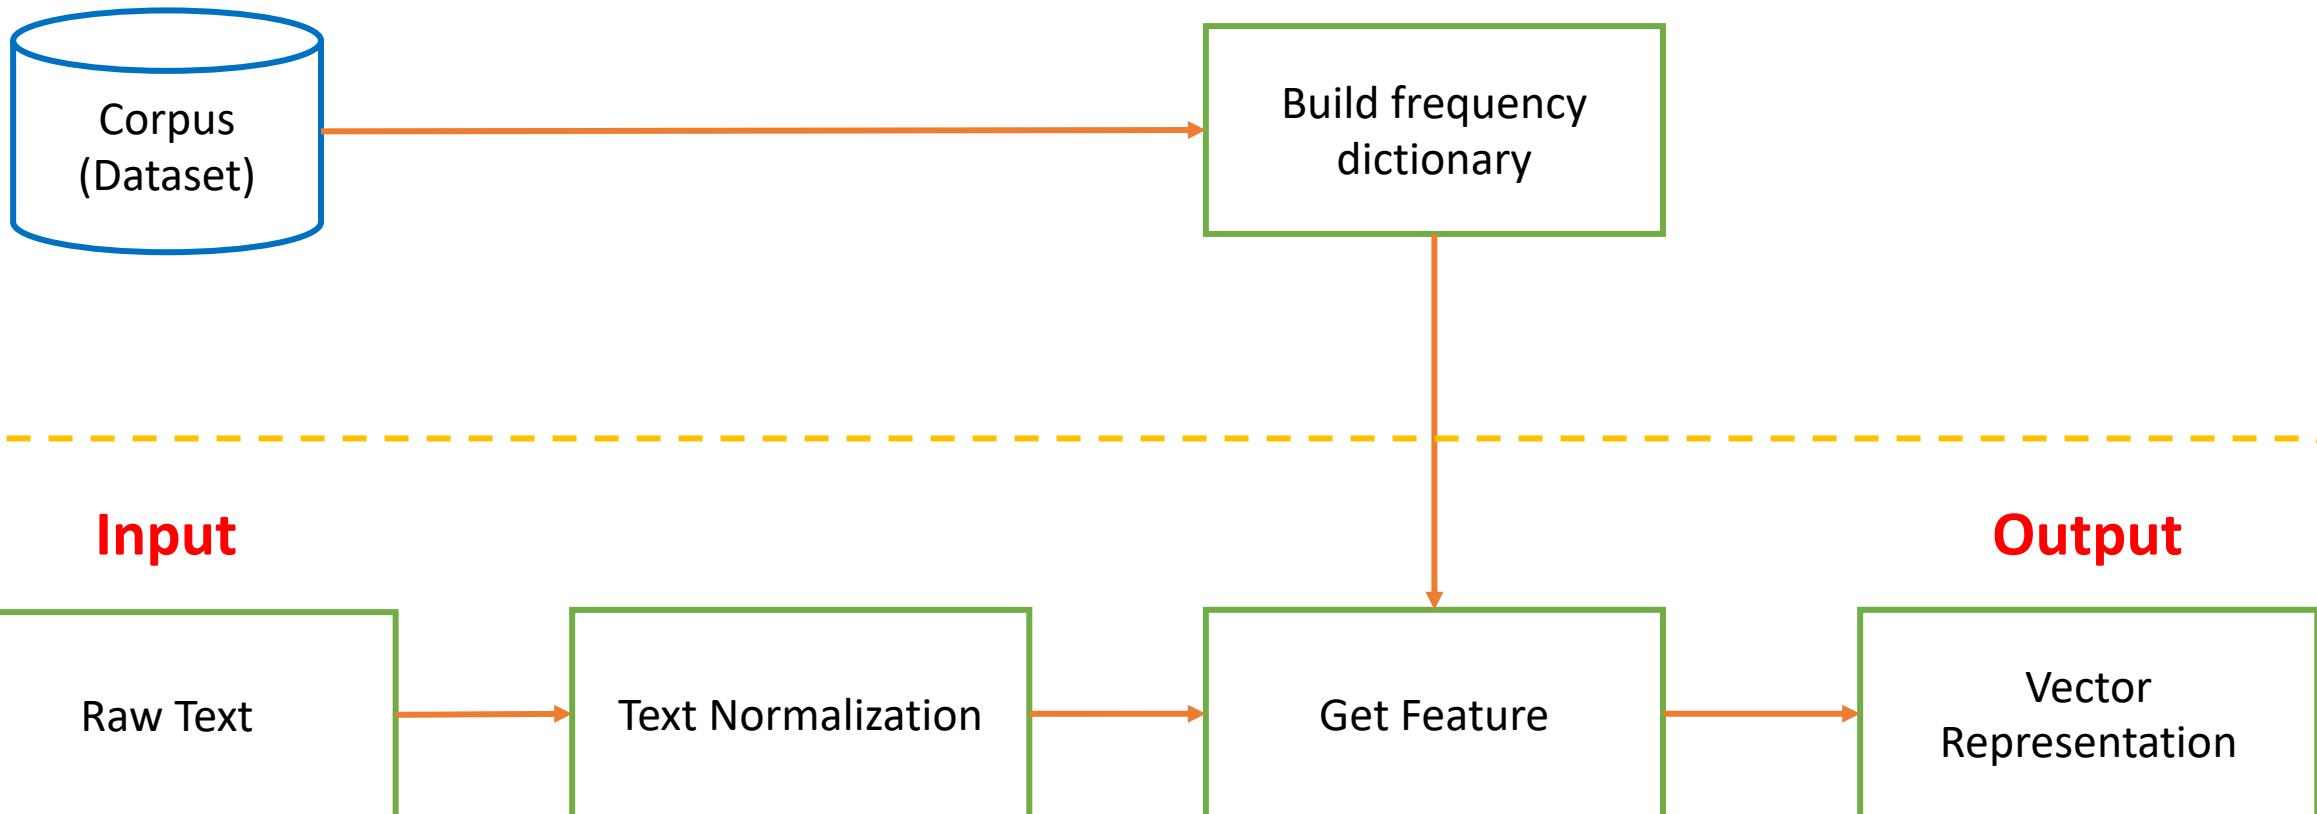
**Vector Representation Meaning**

0                    1                    2

| Bias | Positive Word Count | Negative Word Count |
|------|---------------------|---------------------|
|------|---------------------|---------------------|

# Sentiment Analysis

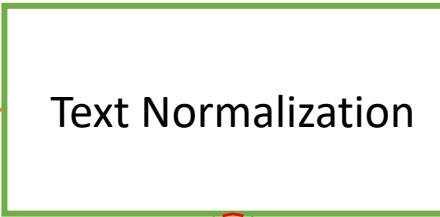
## ❖ Step 4: Data preprocessing



# Sentiment Analysis

## ❖ Step 4: Data preprocessing

**Input**



Raw Text

Text Normalization

Normalized Text

**Output**



Retweet Acronym  
Removal

Hyperlinks  
Removal

Hashtags Removal

Punctuation  
Removal

Tokenization

# Sentiment Analysis

## ❖ Step 4: Data preprocessing

**Retweet Acronym Removal:** Remove "RT" at the start of the string

```
1 text = "RT Hello World"
2 text = re.sub(r'^RT[\s]+', '', text)
3 print(text)
```

Hello World

**Hyperlinks Removal:** Remove all hyperlinks in the string.

```
1 text = "http://fb.me/6N3LsUpCu"
2 text = re.sub(
3     r'https?:\/\/.*[\r\n]*',
4     '',
5     text
6 )
7 print(text)
```

**Hashtag Removal:** Remove all hashtag symbols "#".

```
1 text = "#xinchao, #hello, #konnichiwa"
2 text = re.sub(r'#[\w\s]+', '', text)
3 print(text)
```

xinchao hello konnichiwa

**Tokenization:** Tokenize the string into list of tokens (list of words).

```
1 text = "hello world"
2 tokenizer = TweetTokenizer(
3     preserve_case=False,
4     strip_handles=True,
5     reduce_len=True
6 )
7 text_tokens = tokenizer.tokenize(text)
8 print(text_tokens)
```

[ 'hello', 'world' ]

# Sentiment Analysis

## ❖ Step 4: Data preprocessing

```
1 def text_normalize(text):
2     # Retweet old acronym "RT" removal
3     text = re.sub(r'^RT[\s]+', '', text)
4
5     # Hyperlinks removal
6     text = re.sub(r'https?:\/\/.*[\r\n]*', '', text)
7
8     # Punctuation removal
9     text = re.sub(r'[^w\s]', '', text)
10
11    # Tokenization
12    tokenizer = TweetTokenizer(
13        preserve_case=False,
14        strip_handles=True,
15        reduce_len=True
16    )
17    text_tokens = tokenizer.tokenize(text)
18
19    return text_tokens
```

```
1 text = """We love this! Would you go?
2 #talk #makememories #unplug
3 | #relax #iphone #smartphone #wifi #connect...
4 http://fb.me/6N3LsUpCu
5 """
6 text = text_normalize(text)
7 text
[ 'we',
  'love',
  'this',
  'would',
  'you',
  'go',
  'talk',
  'makememories',
  'unplug',
  'relax',
  'iphone',
  'smartphone',
  'wifi',
  'connect' ]
```

# Sentiment Analysis

## ❖ Step 4: Data preprocessing

**Raw Text:** We love this! Would you go? #talk  
#makememories #unplug #relax #iphone #smartphone #wifi  
#connect... http://fb.me/6N3LsUpCu

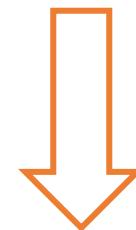
Data  
Preprocessing

**Vector Representation:** `array([1.0, 4768.0, 1425.0])`

### Vector Representation Meaning

0                  1                  2

|      |                     |                     |
|------|---------------------|---------------------|
| Bias | Positive Word Count | Negative Word Count |
|------|---------------------|---------------------|



Build a dictionary containing number of occurrences of each word in positive and negative sentences.

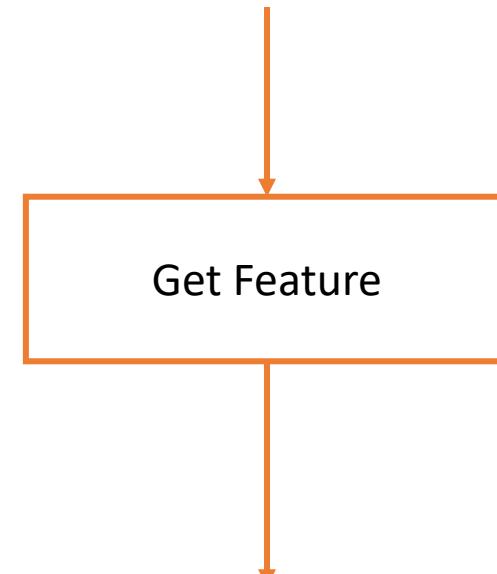
# Sentiment Analysis

## ❖ Step 4: Data preprocessing

Build a dictionary containing number of occurrences of each word in both pos and neg sentences:

```
{  
    ('hello', 'pos'): 10,  
    ('hello', 'neg'): 0,  
    ('world', 'pos'): 30,  
    ('world', 'neg'): 5,  
    ('my', 'pos'): 25,  
    ('check', 'pos'): 50  
}
```

**Input:** [hello, world, my, name, is, Thang]



# Sentiment Analysis

## ❖ Step 4: Data preprocessing

```
1 def get_freqs(df):  
2     freqs = defaultdict(lambda: 0)  
3     for idx, row in df.iterrows():  
4         tweet = row['tweet']  
5         label = row['label']  
6  
7         tokens = text_normalize(tweet)  
8         for token in tokens:  
9             pair = (token, label)  
10            freqs[pair] += 1  
11  
12    return freqs  
13  
14 freqs = get_freqs(df)
```

```
{('fingerprint', 0): 4,  
 ('pregnancy', 0): 1,  
 ('test', 0): 8,  
 ('finally', 0): 168,  
 ('a', 0): 727,  
 ('transparant', 0): 1,  
 ('silicon', 0): 1,  
 ('case', 0): 228,  
 ('thanks', 0): 94,  
 ('to', 0): 876,  
 ('my', 0): 1227,  
 ('uncle', 0): 4,  
 ('yay', 0): 63,  
 ('sony', 0): 701,  
 ('xperia', 0): 54,  
 ('s', 0): 38,  
 ('sonyexperias', 0): 1,  
 ('we', 0): 159,  
 ('love', 0): 385,
```

# Sentiment Analysis

## ❖ Step 4: Data preprocessing

```
1 def get_freqs(df):
2     freqs = defaultdict(lambda: 0)
3     for idx, row in df.iterrows():
4         tweet = row['tweet']
5         label = row['label']
6
7         tokens = text_normalize(tweet)
8         for token in tokens:
9             pair = (token, label)
10            freqs[pair] += 1
11
12    return freqs
13
14 freqs = get_freqs(df)
```

`class collections.defaultdict(default_factory=None, /[, ...])`  
Return a new dictionary-like object. `defaultdict` is a subclass of the built-in `dict` class. It overrides one method and adds one writable instance variable. The remaining functionality is the same as for the `dict` class and is not documented here.

The first argument provides the initial value for the `default_factory` attribute; it defaults to `None`. All remaining arguments are treated the same as if they were passed to the `dict` constructor, including keyword arguments.

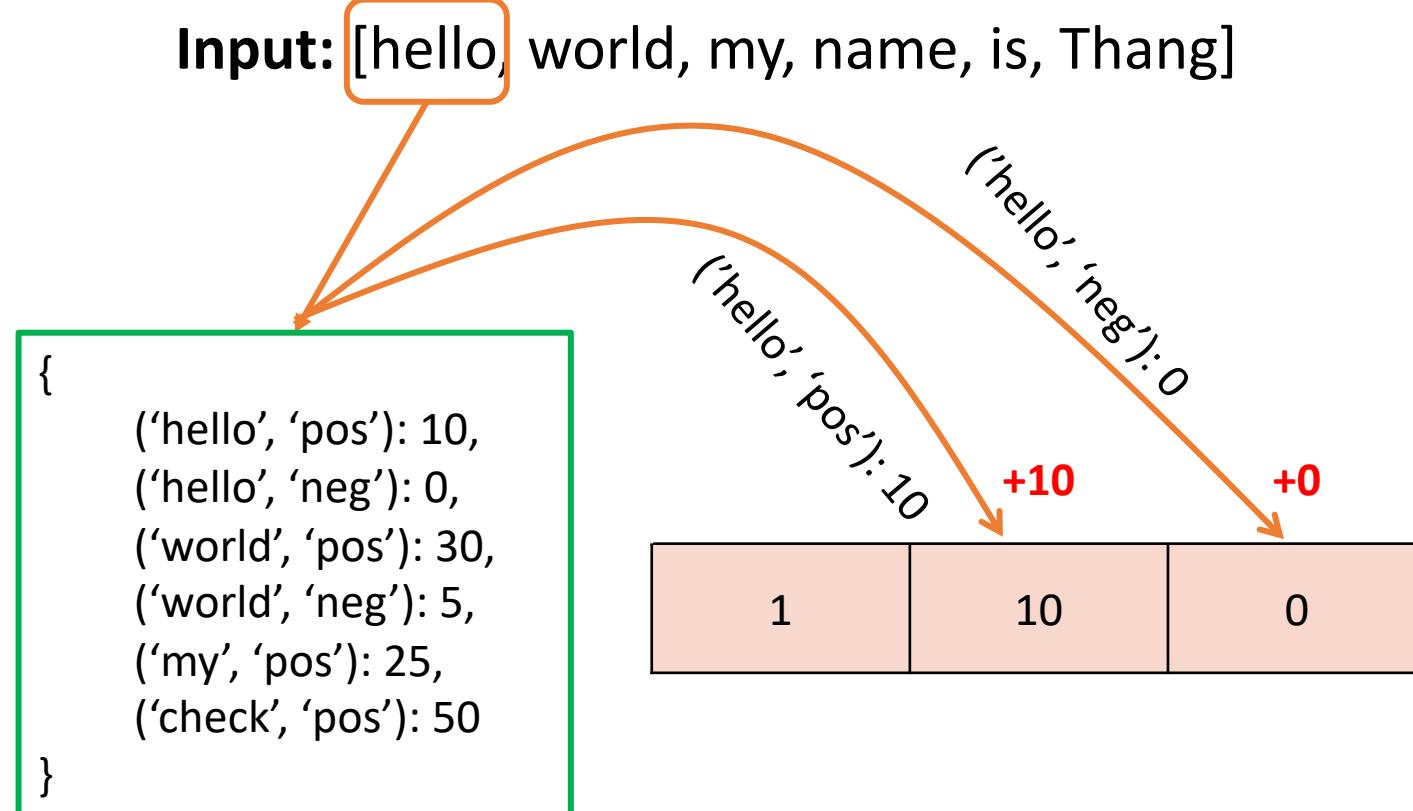
**defaultdict benefit:** Assign a default value for non-existing keys. For e.g., `lambda: 0` will assign value 0 to keys excluded from the dictionary.

# Sentiment Analysis

## ❖ Step 4: Data preprocessing

With frequency dictionary, we can create new vector representation for raw texts as following:

```
1 def get_feature(text, freqs):
2     tokens = text_normalize(text)
3
4     x = np.zeros(3)
5     x[0] = 1
6
7     for token in tokens:
8         x[1] += freqs[(token, 0)]
9         x[2] += freqs[(token, 1)]
10
11    return x
```



# Sentiment Analysis

## ❖ Step 4: Data preprocessing

**Raw Text:** We love this! Would you go? #talk  
#makememories #unplug #relax #iphone #smartphone #wifi  
#connect... http://fb.me/6N3LsUpCu

Data  
Preprocessing

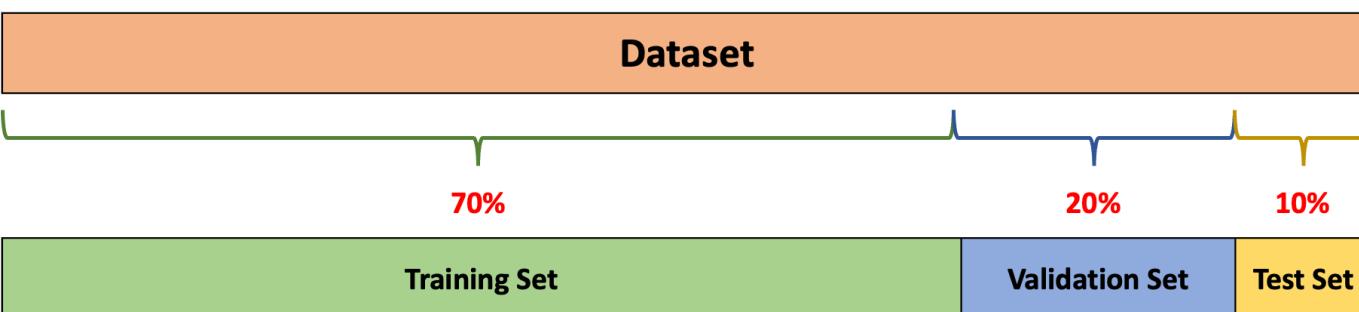
**Vector Representation:** array([1.0, 4768.0, 1425.0])

```
1 x = []
2 y = []
3
4 for idx, row in df.iterrows():
5     tweet = row['tweet']
6     label = row['label']
7
8     x_i = get_feature(tweet, freqs)
9     x.append(x_i)
10    y.append(label)
11
12 x = np.array(X)
13 y = np.array(y).reshape(-1, 1)
14
15 print(f'x shape: {x.shape}')
16 print(f'y shape: {y.shape}')
```

x shape: (7920, 3)  
y shape: (7920, 1)

# Sentiment Analysis

## ❖ Step 5: Split train, val, test set



```
1 val_size = 0.2
2 test_size = 0.125
3 random_state = 2
4 is_shuffle = True
5
6 X_train, X_val, y_train, y_val = train_test_split(
7     X, y,
8     test_size=val_size,
9     random_state=random_state,
10    shuffle=is_shuffle
11 )
12
13 X_train, X_test, y_train, y_test = train_test_split(
14     X_train, y_train,
15     test_size=test_size,
16     random_state=random_state,
17     shuffle=is_shuffle
18 )
```

```
1 print(f'Number of training samples: {X_train.shape[0]}')
2 print(f'Number of val samples: {X_val.shape[0]}')
3 print(f'Number of test samples: {X_test.shape[0]}')
```

Number of training samples: 5544  
Number of val samples: 1584  
Number of test samples: 792

# Sentiment Analysis

## ❖ Step 6: Normalization

Using `sklearn.preprocessing.StandardScaler()` to scale all values in dataset.

$$z = \frac{x_i - \mu}{\sigma}$$

```
1 normalizer = StandardScaler()  
2 X_train[:, 1:] = normalizer.fit_transform(X_train[:, 1:])  
3 X_val[:, 1:] = normalizer.transform(X_val[:, 1:])  
4 X_test[:, 1:] = normalizer.transform(X_test[:, 1:])
```

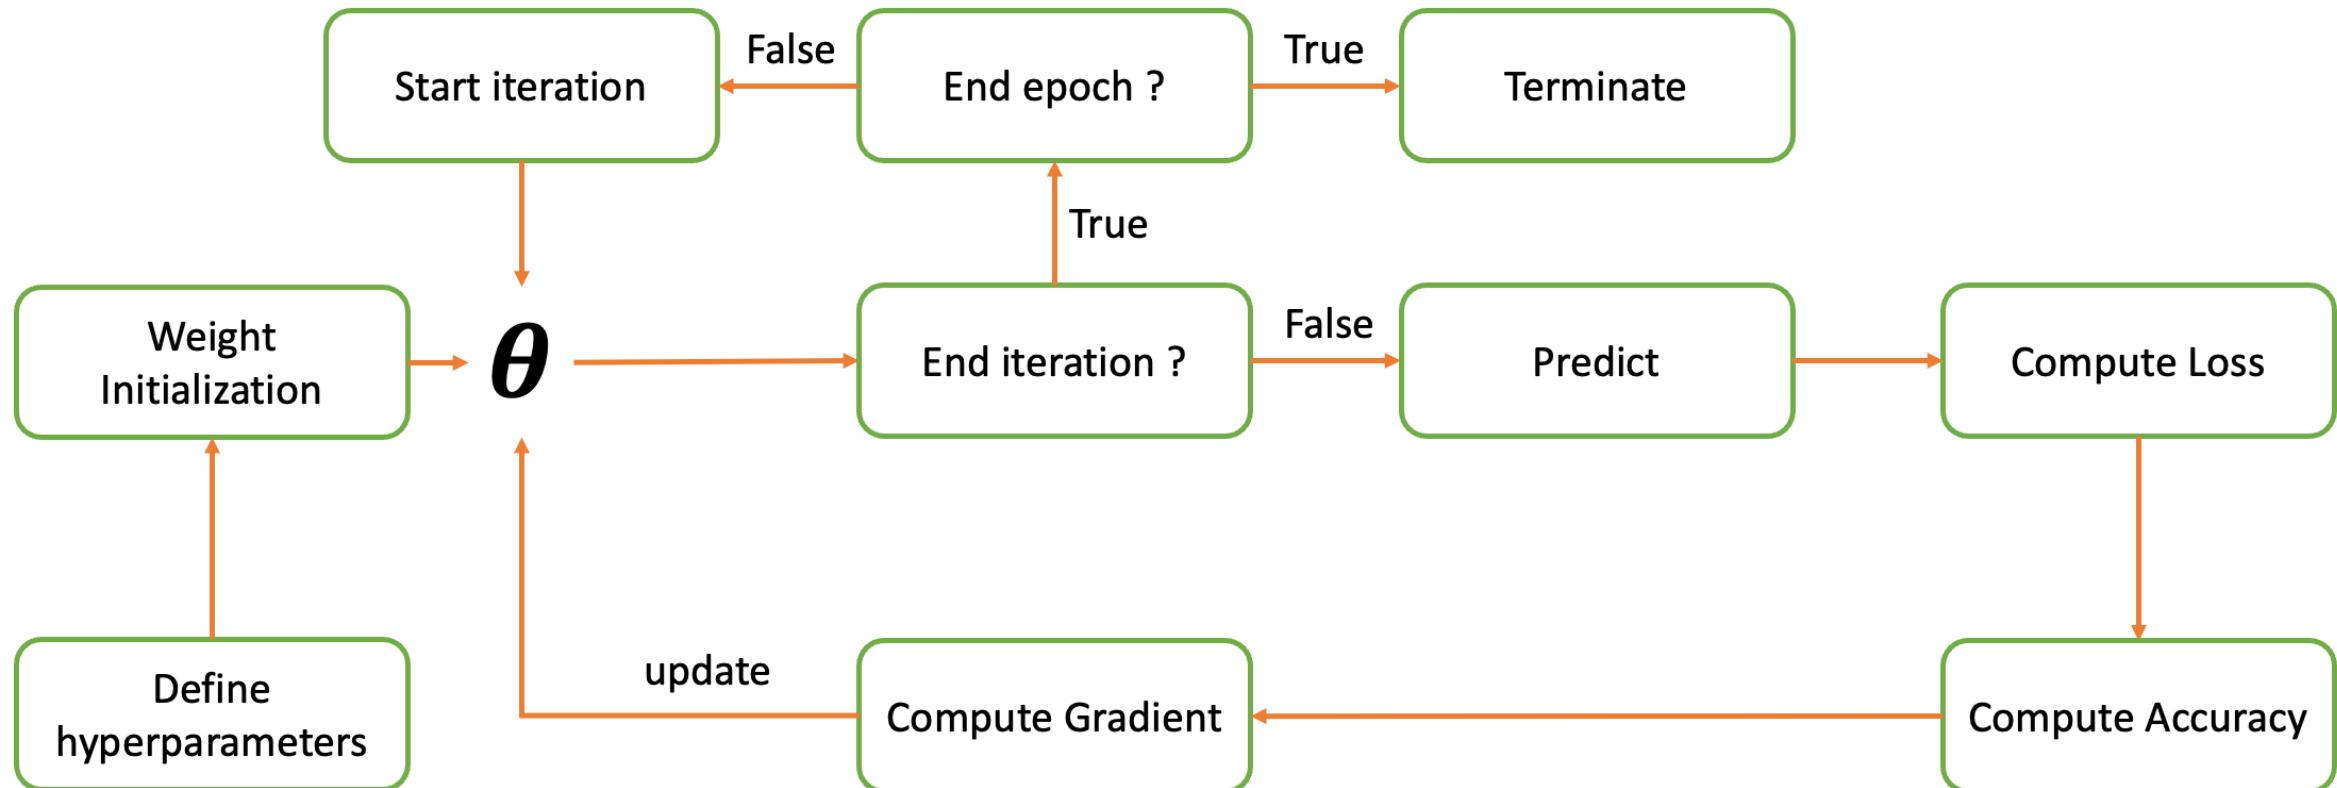
1 X\_train

```
array([[ 1.          , -0.04496429, -0.80123214],  
       [ 1.          ,  0.83177357,  0.04495408],  
       [ 1.          , -0.78519562, -0.93920268],  
       ...,  
       [ 1.          , -0.37919422,  0.26878398],  
       [ 1.          ,  0.33852057, -0.10989055],  
       [ 1.          ,  1.73384255,  1.86239334]])
```

**Note:** We only use the train set to fit the scaler.

# Sentiment Analysis

## ❖ Step 7: Training



# Sentiment Analysis

## ❖ Step 7: Training

```
1 lr = 0.01
2 epochs = 200
3 batch_size = 128
4
5 np.random.seed(random_state)
6 theta = np.random.uniform(
7     size=x_train.shape[1]
8 )
```

Define essential functions and hyperparameters

```
1 def sigmoid(z):
2     return 1 / (1 + np.exp(-z))
3
4 def predict(X, theta):
5     dot_product = np.dot(X, theta)
6     y_hat = sigmoid(dot_product)
7
8     return y_hat
9
10 def compute_loss(y_hat, y):
11     y_hat = np.clip(
12         y_hat, 1e-7, 1 - 1e-7
13     )
14
15     return (-y * np.log(y_hat) - (1 - y) * np.log(1 - y_hat)).mean()
16
17 def compute_gradient(X, y, y_hat):
18     return np.dot(
19         X.T, (y_hat - y)
20     ) / y.size
21
22 def update_theta(theta, gradient, lr):
23     return theta - lr * gradient
24
25 def compute_accuracy(X, y, theta):
26     y_hat = predict(X, theta).round()
27     acc = (y_hat == y).mean()
28
29     return acc
```

# Sentiment Analysis

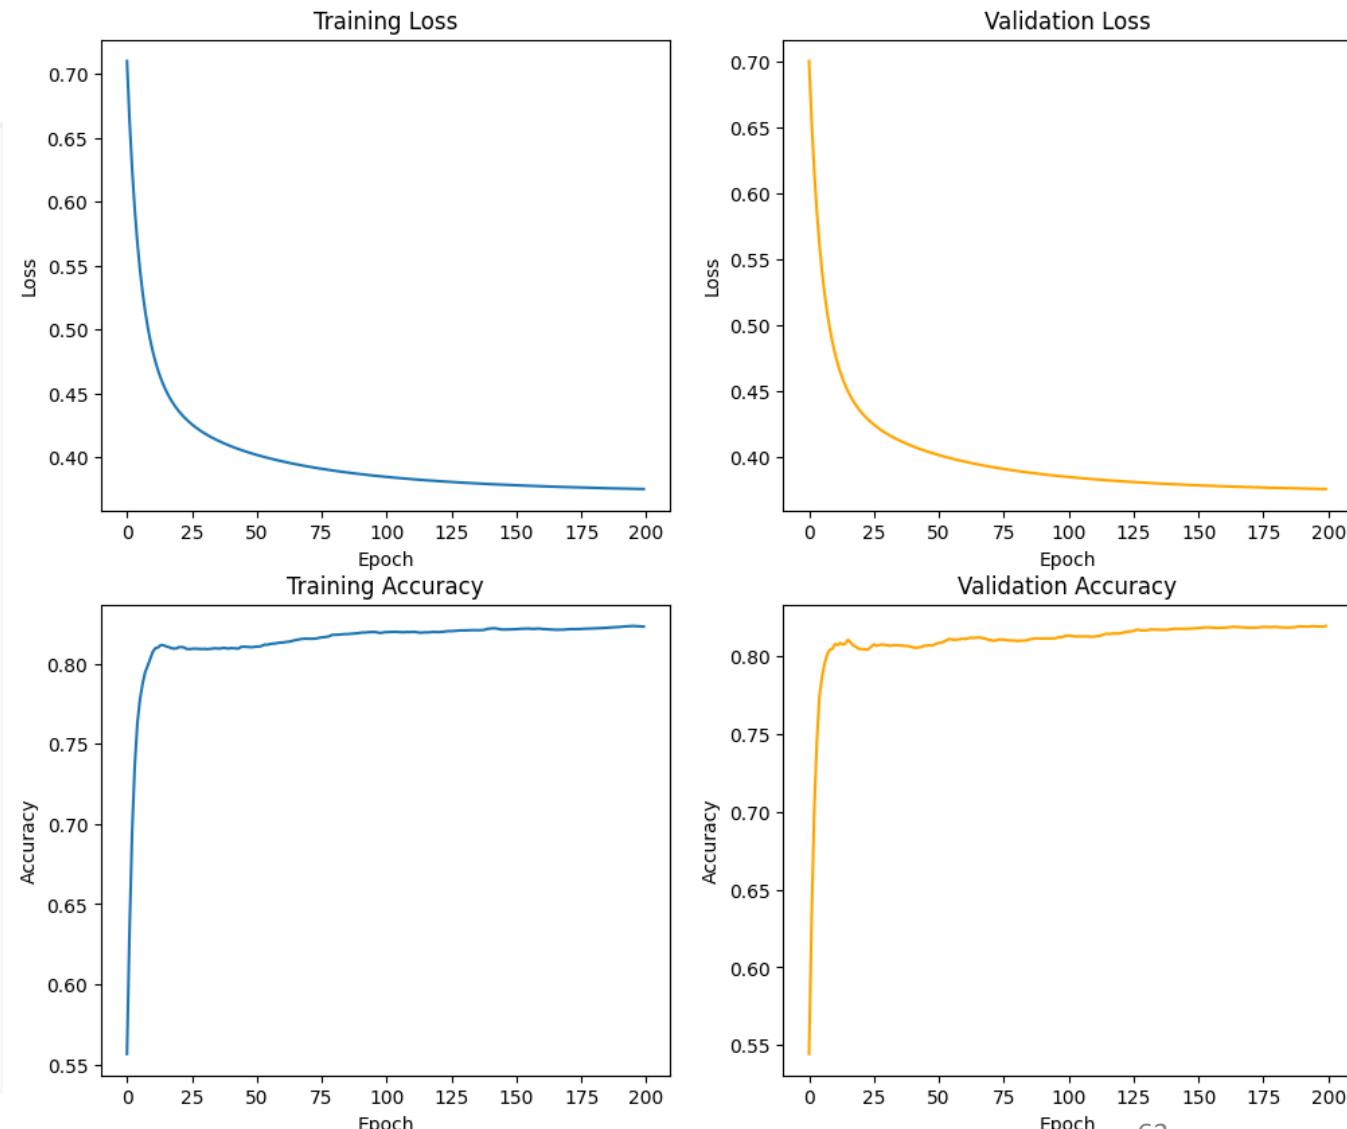
## ❖ Step 7: Training

```
6 for epoch in range(epochs):
7     train_batch_losses = []
8     train_batch_accs = []
9     val_batch_losses = []
10    val_batch_accs = []
11
12    for i in range(0, X_train.shape[0], batch_size):
13        X_i = X_train[i:i+batch_size]
14        y_i = y_train[i:i+batch_size]
15
16        y_hat = predict(X_i, theta)
17        train_loss = compute_loss(y_hat, y_i)
18        gradient = compute_gradient(X_i, y_i, y_hat)
19        theta = update_theta(theta, gradient, lr)
20        train_batch_losses.append(train_loss)
21
22        train_acc = compute_accuracy(X_train, y_train, theta)
23        train_batch_accs.append(train_acc)
24
25        y_val_hat = predict(X_val, theta)
26        val_loss = compute_loss(y_val_hat, y_val)
27        val_batch_losses.append(val_loss)
28
29        val_acc = compute_accuracy(X_val, y_val, theta)
30        val_batch_accs.append(val_acc)
31
32        train_batch_loss = sum(train_batch_losses) / len(train_batch_losses)
33        val_batch_loss = sum(val_batch_losses) / len(val_batch_losses)
34        train_batch_acc = sum(train_batch_accs) / len(train_batch_accs)
35        val_batch_acc = sum(val_batch_accs) / len(val_batch_accs)
```

# Sentiment Analysis

## ❖ Step 7: Training (Visualization)

```
1 fig, ax = plt.subplots(2, 2, figsize=(12, 10))
2 ax[0, 0].plot(train_losses)
3 ax[0, 0].set(xlabel='Epoch', ylabel='Loss')
4 ax[0, 0].set_title('Training Loss')
5
6 ax[0, 1].plot(val_losses, 'orange')
7 ax[0, 1].set(xlabel='Epoch', ylabel='Loss')
8 ax[0, 1].set_title('Validation Loss')
9
10 ax[1, 0].plot(train_accs)
11 ax[1, 0].set(xlabel='Epoch', ylabel='Accuracy')
12 ax[1, 0].set_title('Training Accuracy')
13
14 ax[1, 1].plot(val_accs, 'orange')
15 ax[1, 1].set(xlabel='Epoch', ylabel='Accuracy')
16 ax[1, 1].set_title('Validation Accuracy')
17
18 plt.show()
```



# Sentiment Analysis

## ❖ Step 8: Evaluation

$$\text{accuracy} = \frac{\text{true\_predictions}}{\text{n\_samples}}$$

```
1 def compute_accuracy(X, y, theta):  
2     y_hat = predict(X, theta).round()  
3     acc = (y_hat == y).mean()  
4  
5     return acc
```

```
1 # Val set  
2 val_set_acc = compute_accuracy(X_val, y_val, theta)  
3 print('Evaluation on validation set:')  
4 print(f'Accuracy: {val_set_acc}')
```

Evaluation on validation set:  
Accuracy: 0.821969696969697

```
1 # Test set  
2 test_set_acc = compute_accuracy(X_test, y_test, theta)  
3 print('Evaluation on test set:')  
4 print(f'Accuracy: {test_set_acc}')
```

Evaluation on test set:  
Accuracy: 0.8434343434343434

# Question

