

# Pytorch

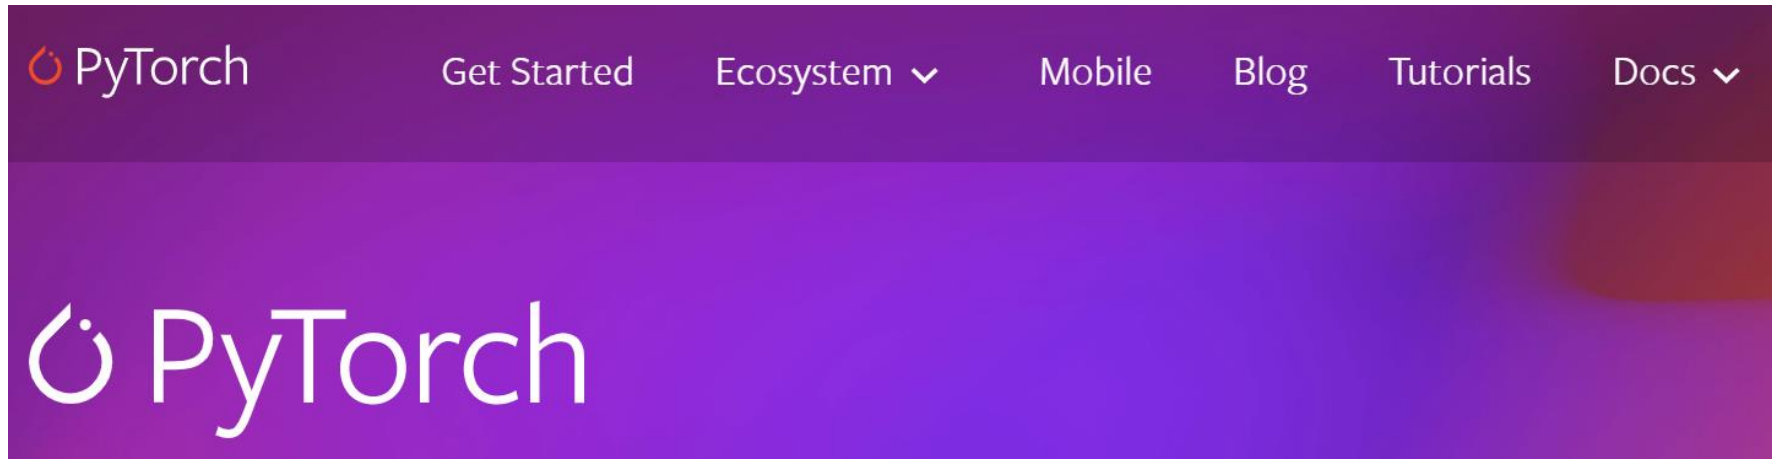
## Deep Learning Framework

**Quang-Vinh Dinh**  
**Ph.D. in Computer Science**

# Outline

- **Introduction to PyTorch**
- **Model Construction**
- **Model Training and Inference**
- **Applying Softmax for Image Data**

# Introduction



PyTorch Build	Stable (2.1.0)		Preview (Nightly)	
Your OS	Linux	Mac	Windows	
Package	Conda	Pip	LibTorch	Source
Language	Python		C++ / Java	
Compute Platform	CUDA 11.8	CUDA 12.1	ROCm 5.6	CPU
Run this Command:	<pre>pip3 install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cu118</pre>			

 Search Tutorials

[PyTorch Recipes](#) [ + ]

[Introduction to PyTorch](#) [ + ]

[Introduction to PyTorch on YouTube](#) [ + ]

[Learning PyTorch](#) [ + ]

[Image and Video](#) [ + ]

[Audio](#) [ + ]

[Text](#) [ + ]

[Backends](#) [ + ]

[Reinforcement Learning](#) [ + ]

[Deploying PyTorch Models in Production](#) [ + ]

[Code Transforms with FX](#) [ + ]

[Frontend APIs](#) [ + ]

[Extending PyTorch](#) [ + ]

[Model Optimization](#) [ + ]

[Parallel and Distributed Training](#) [ + ]

[Mobile](#) [ + ]

[Recommendation Systems](#) [ + ]

[Multimodality](#) [ + ]

# Create a tensor in PyTorch

```
1 # From a List
2 import torch
3
4 data = [1, 2, 3]
5 data = torch.tensor(data)
6 print(data)
```

```
tensor([1, 2, 3])
```

```
1 # From a NumPy Array
2 import numpy as np
3
4 data = np.array([1, 2, 3])
5 data = torch.from_numpy(data)
6 print(data)
```

```
tensor([1, 2, 3], dtype=torch.int32)
```

```
1 # random uniform
2 import torch
3
4 # Creates a 3x3 tensor with
5 # random values between 0 and 1
6 data = torch.rand(3, 3)
7 print(data)
```

```
tensor([[0.0615, 0.2716, 0.2160],
        [0.1595, 0.9780, 0.6053],
        [0.7718, 0.1605, 0.4205]])
```

```
1 # random normal
2 import torch
3
4 # Creates a 3x3 tensor with values
5 # sampled from a standard normal distribution
6 data = torch.randn(3, 3)
7 print(data)
```

```
tensor([[ -1.0891, -0.2750,  0.8105],
        [  1.0836, -1.1107,  0.3683],
        [  2.0638, -1.2075, -0.5610]])
```


```
1 # Using the arange function
2 import torch
3
4 # Creates a tensor [0, 1, 2, 3, 4]
5 data = torch.arange(start=0, end=5, step=1)
6 print(data)
```

```
tensor([0, 1, 2, 3, 4])
```

# Introduction

## ❖ Tensor

### ❖ Broadcasting



```
1 # broadcasting
2 import numpy as np
3
4 # create a tensor
5 tensor1 = torch.tensor([[1, 2],
6                          [3, 4]])
7 tensor2 = torch.tensor([1])
8 print(f'Tensor1:\n {tensor1}')
9 print(f'Tensor2:\n {tensor2}')
10
11 # Addition between two tensors
12 tensor3 = tensor1 + tensor2
13 print(f'Tensor3:\n {tensor3}')
```

Tensor1:  
tensor([[1, 2],  
 [3, 4]])


Tensor2:  
tensor([1])

Tensor3:  
tensor([[2, 3],  
 [4, 5]])

# Introduction

## ❖ Tensor

### ❖ Broadcasting



```
1 # broadcasting
2 import numpy as np
3
4 # create a tensor
5 tensor1 = torch.tensor([[1, 2],
6                          [3, 4]])
7 print(f'Tensor1:\n {tensor1}')
8
9 # Between a tensor and a number
10 tensor2 = tensor1 + 2
11 print(f'Tensor2:\n {tensor2}')
```

```
Tensor1:
  tensor([[1, 2],
          [3, 4]])
Tensor2:
  tensor([[3, 4],
          [5, 6]])
```

# Introduction

## ❖ Tensor

### ❖ Important functions

Squared Difference

$$sd = (x - y)^2$$

x		sd
1		16
2		9
3		4
4		1

$(x - y)^2 =$

```
1 # Compute squared difference
2 import torch
3
4 # Create two tensors
5 x = torch.tensor([1.0, 2.0, 3.0, 4.0])
6 y = 5
7
8 # Compute squared difference
9 squared_diff = (x - y) ** 2
10
11 print(f'x:\n {x}')
12 print(f'y = {y}')
13 print(f'squared_diff:\n {squared_diff}')
```

```
x:
  tensor([1., 2., 3., 4.])
y = 5
squared_diff:
  tensor([16.,  9.,  4.,  1.]')
```

# Introduction

## ❖ Tensor

### ❖ Important functions

#### Mean Squared Error

$$sd = \sum_i (x_i - y_i)^2$$

x	y	sd
1	5	16
2	5	9
3	5	4
4	5	1

$\sum_i (x_i - y_i)^2 = \sum_i ($

```
1 # Compute squared difference
2 import torch
3
4 # Create two tensors
5 x = torch.tensor([1.0, 2.0, 3.0, 4.0])
6 y = torch.tensor([5.0, 5.0, 5.0, 5.0])
7
8 # Compute squared difference
9 loss_fn = torch.nn.MSELoss()
10 mse = loss_fn(x,y)
11
12 print(f'x:\n {x}')
13 print(f'y = {y}')
14 print(f'mse:\n {mse}')
```


```
x:
  tensor([1., 2., 3., 4.])
y = tensor([5., 5., 5., 5.])
mse:
  7.5
```



# Introduction

## ❖ Tensor

### ❖ Important functions



```
1 # Concatenate tensors
2 import torch
3
4 # Create two tensors
5 x = torch.tensor([[1, 2],
6                   [3, 4]])
7 y = torch.tensor([[3, 4],
8                   [5, 6]])
9
10 # Concat tensors along the first dim
11 tensor1 = torch.cat((x, y), dim=0)
12 print(f'Tensor1:\n {tensor1}')
13
14 # Concat tensors along the second dim
15 tensor2 = torch.cat((x, y), dim=1)
16 print(f'Tensor2:\n {tensor2}')
```

```
Tensor1:
tensor([[1, 2],
        [3, 4],
        [3, 4],
        [5, 6]])
Tensor2:
tensor([[1, 2, 3, 4],
        [3, 4, 5, 6]])
```

# Introduction

## ❖ Tensor

### ❖ Important functions

3	4
4	2
1	1

 $\text{.argmax}(\text{axis}=0) =$ 

1	0
---	---

3	4
4	2
1	1

 $\text{.argmax}(\text{axis}=1) =$ 

1	0	0
---	---	---

```
1 # argmax
2 import torch
3
4 # Creates a 3x2 tensor
5 data = torch.randint(low=0, high=9, size=(3, 2))
6 print(f'data:\n {data}')
7
8 # Compute argmax across the rows (dimension 0)
9 argmax_dim0 = torch.argmax(data, dim=0)
10 print(f'argmax1:\n {argmax_dim0}')
11
12 # Compute argmax across the columns (dimension 1)
13 argmax_dim1 = torch.argmax(data, dim=1)
14 print(f'argmax1:\n {argmax_dim1}')
```


```
data:
  tensor([[3, 4],
          [4, 2],
          [1, 1]])
argmax1:
  tensor([1, 0])
argmax1:
  tensor([1, 0, 0])
```

# Introduction

## ❖ Gradient computation

$$y = x^2$$

$$z = 3y + 2$$


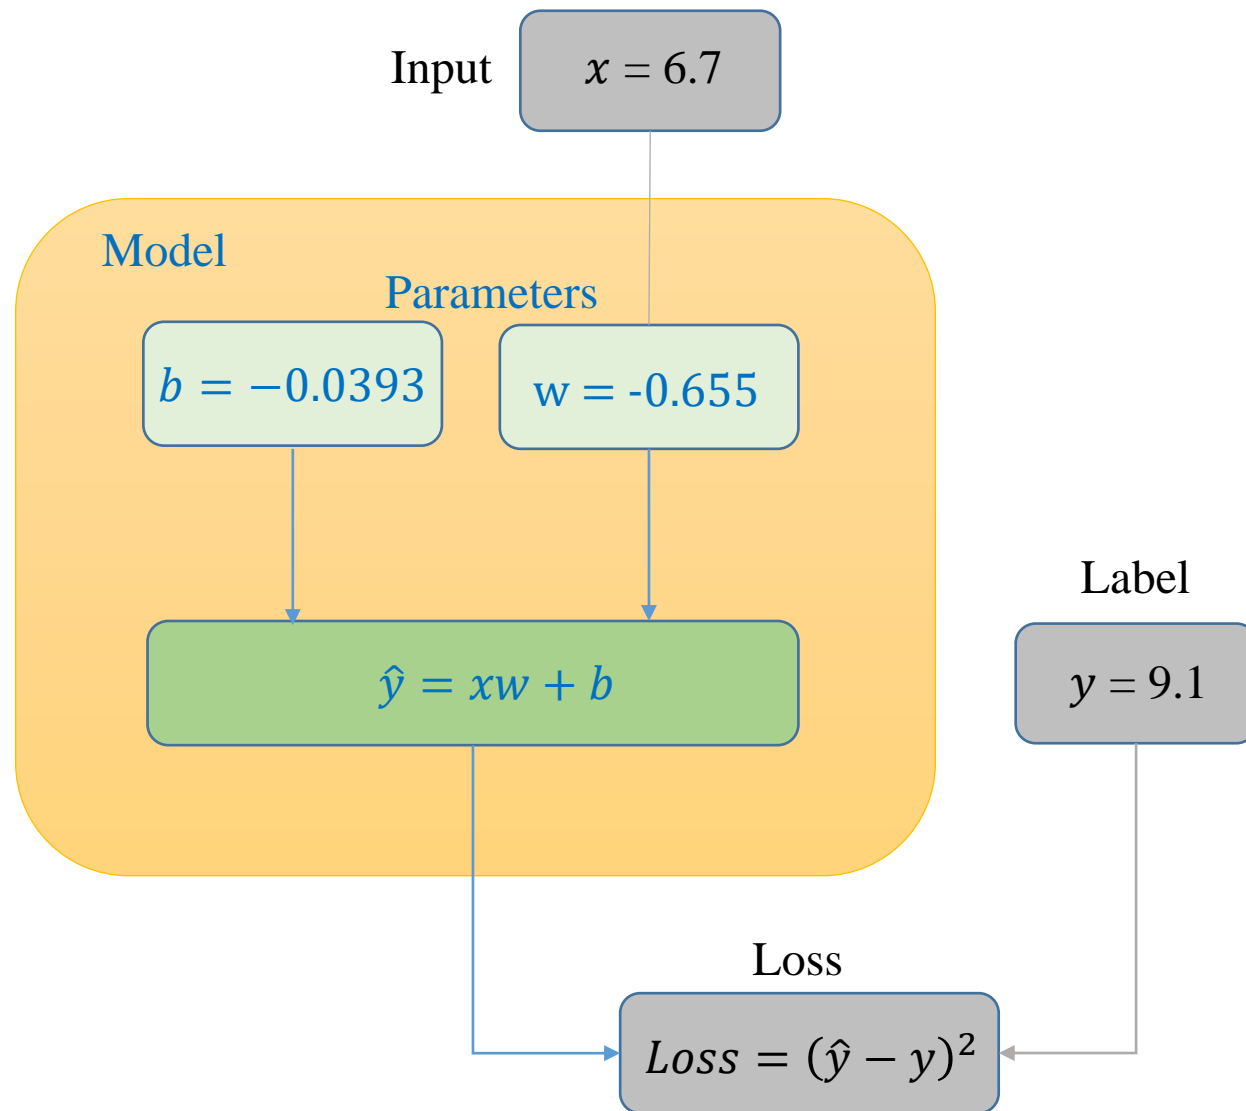


```
1 # autograd
2 import torch
3
4 # Create a tensor to compute gradients
5 x = torch.tensor(2.0, requires_grad=True)
6
7 # operation
8 y = x ** 2
9 z = 3*y + 2
10 print(f'z: {z}')
11
12 # Backpropagate to compute gradients
13 z.backward()
14
15 # Print the gradient. dz/dx at x=2.0
16 print(f"Gradient of z w.r.t. x: {x.grad}")
```

z: 14.0

Gradient of z with respect to x: 12.0

# Example 1




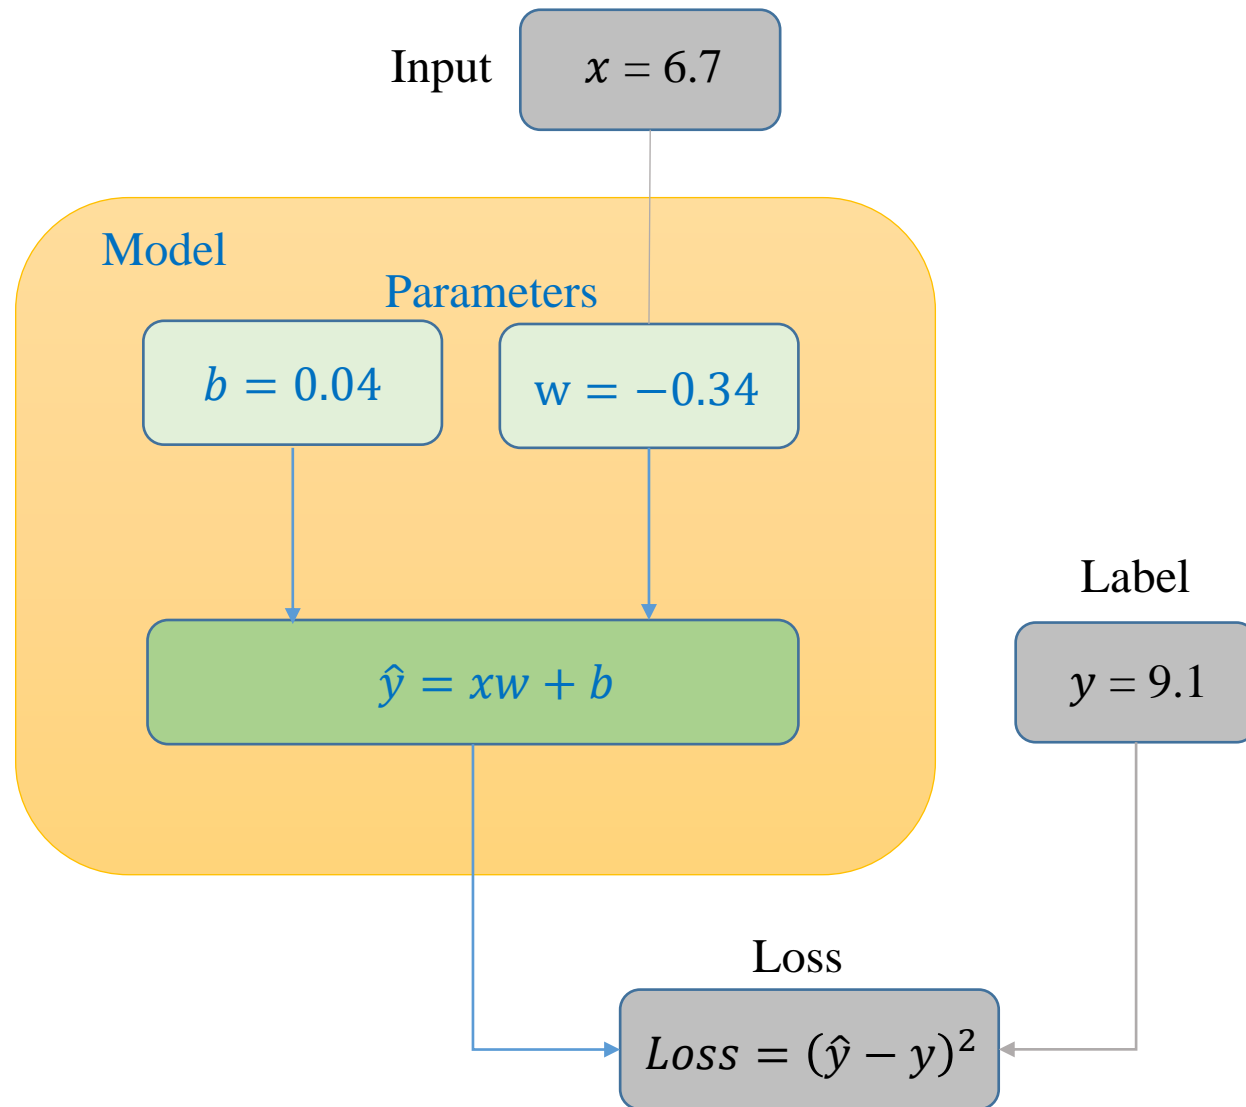
```
import torch.nn as nn

# Create a Linear layer
linear = nn.Linear(1, 1)

print(f'b - {linear.bias}')
print(f'w - {linear.weight}')
```

b - Parameter containing:  
tensor([-0.0393], requires\_grad=True)  
w - Parameter containing:  
tensor([[ -0.6550]], requires\_grad=True)

# Example 1



```
import torch.nn as nn
import torch

# Create a Linear Layer
linear = nn.Linear(1, 1)


# set values
linear.bias.data = torch.Tensor([0.04])
linear.weight.data = torch.Tensor([[ -0.34]])

print(f'b - {linear.bias}')
print(f'w - {linear.weight}')

b - Parameter containing:
tensor([0.0400], requires_grad=True)
w - Parameter containing:
tensor([[ -0.3400]], requires_grad=True)
```

# Example 1

$$L'_w \quad L'_b$$



Input

$x = 6.7$

Model

Parameters

$b = 0.04$

$w = -0.34$

$\hat{y} = xw + b$

Label

$y = 9.1$

Loss

$Loss = (\hat{y} - y)^2$



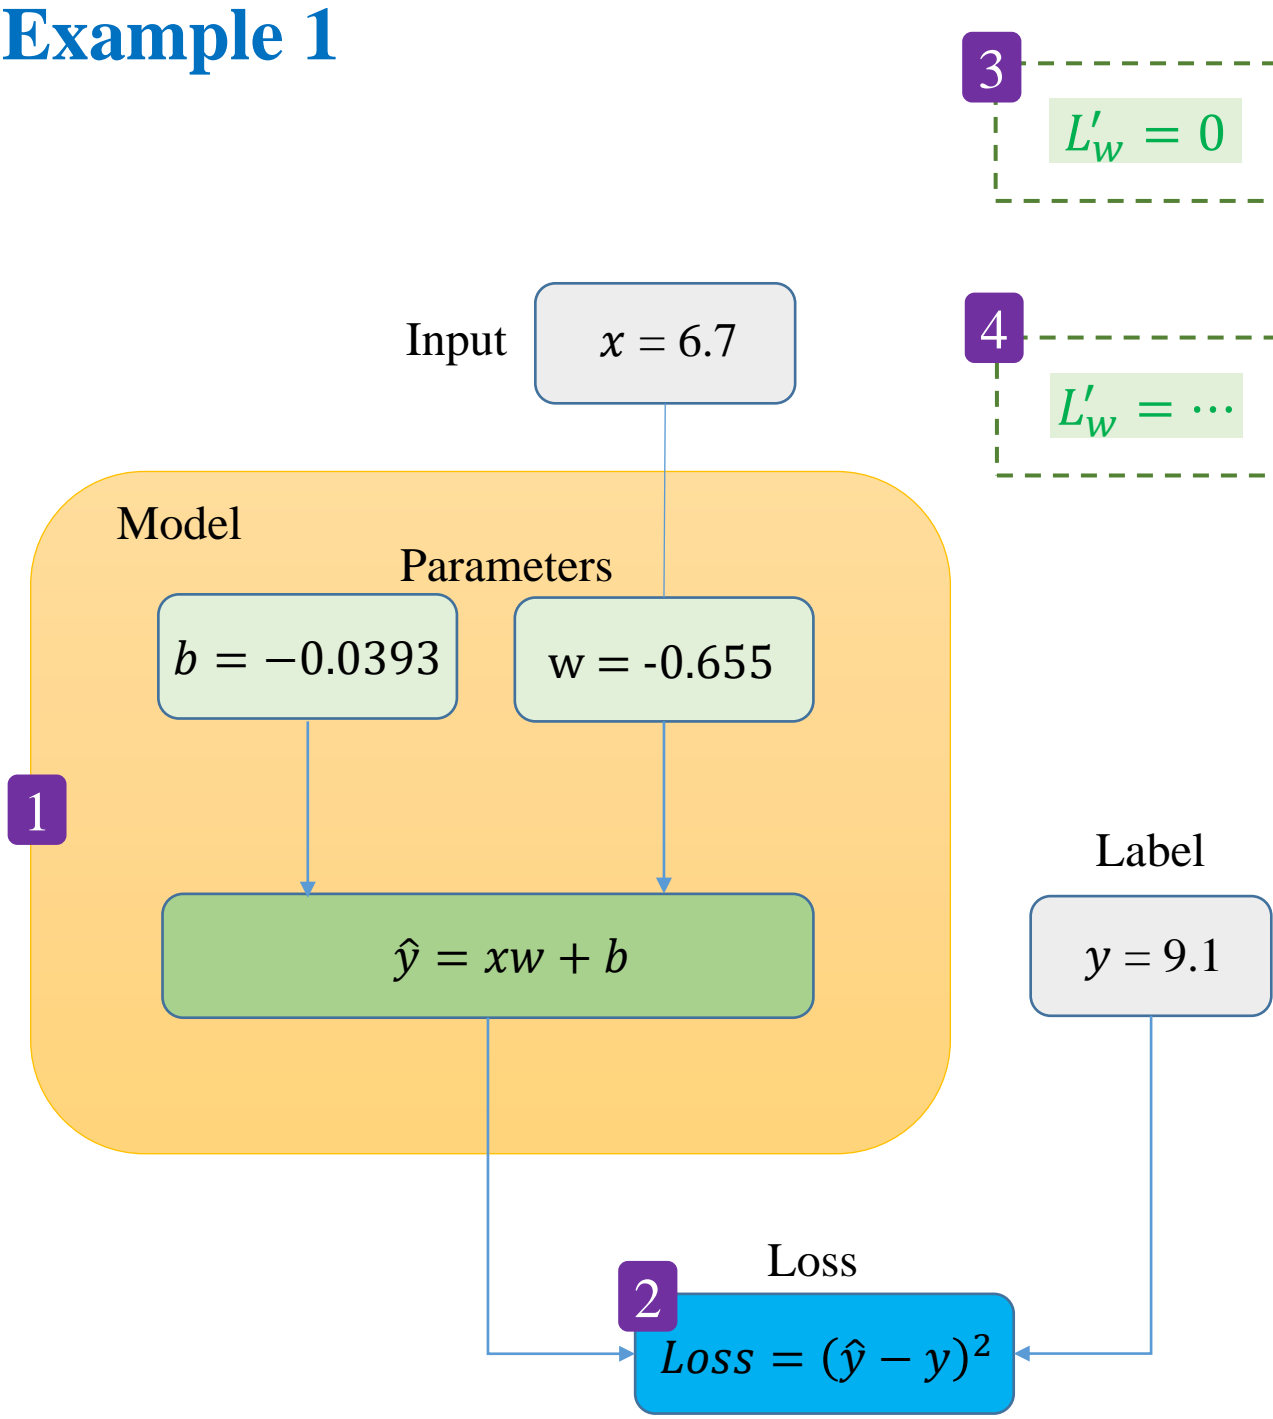
```
import torch.nn as nn
import torch

# Create a Linear Layer
linear = nn.Linear(1, 1)

# set values
linear.bias.data = torch.Tensor([0.04])
linear.weight.data = torch.Tensor([[ -0.34]])

# loss function and optimizer
loss_fn = torch.nn.MSELoss()
optimizer = torch.optim.SGD(linear.parameters(),
                             lr=0.01)
```

# Example 1



```
# y_hat
y_hat = linear(x)

# loss
loss = loss_fn(y_hat, y)

# compute gradient
optimizer.zero_grad()
loss.backward()

# update
optimizer.step()
```

Annotations for the code block:

- Purple box "1" points to the line `y_hat = linear(x)`.
- Purple box "2" points to the line `loss = loss_fn(y_hat, y)`.
- Purple box "3" points to the line `optimizer.zero_grad()`.
- Purple box "4" points to the line `loss.backward()`.
- Purple box "5" points to the line `optimizer.step()`.

# Example 1

Initialize  $b$  and  $w$  randomly

Input

$x = 6.7$

	area	price	
	6.7	9.1	
	4.6	5.9	
	3.5	4.6	
	5.5	6.7	

Model

Parameters

$b = 0.04$

$w = -0.34$

$$\hat{y} = xw + b = -2.238$$

Label

$y = 9.1$

Loss

$$(\hat{y} - y)^2 = 128.55$$

```
import numpy as np
import torch
import torch.nn as nn
```

### Data preparation

```
data = np.genfromtxt('data.csv', delimiter=',')
x_data = torch.from_numpy(data[:, 0:1]).float()
y_data = torch.from_numpy(data[:, 1:]).float()
```

# Create model, loss and optimizer

```
linear = nn.Linear(1, 1)
loss_fn = torch.nn.MSELoss()
optimizer = torch.optim.SGD(linear.parameters(),
                              lr=0.01)
```

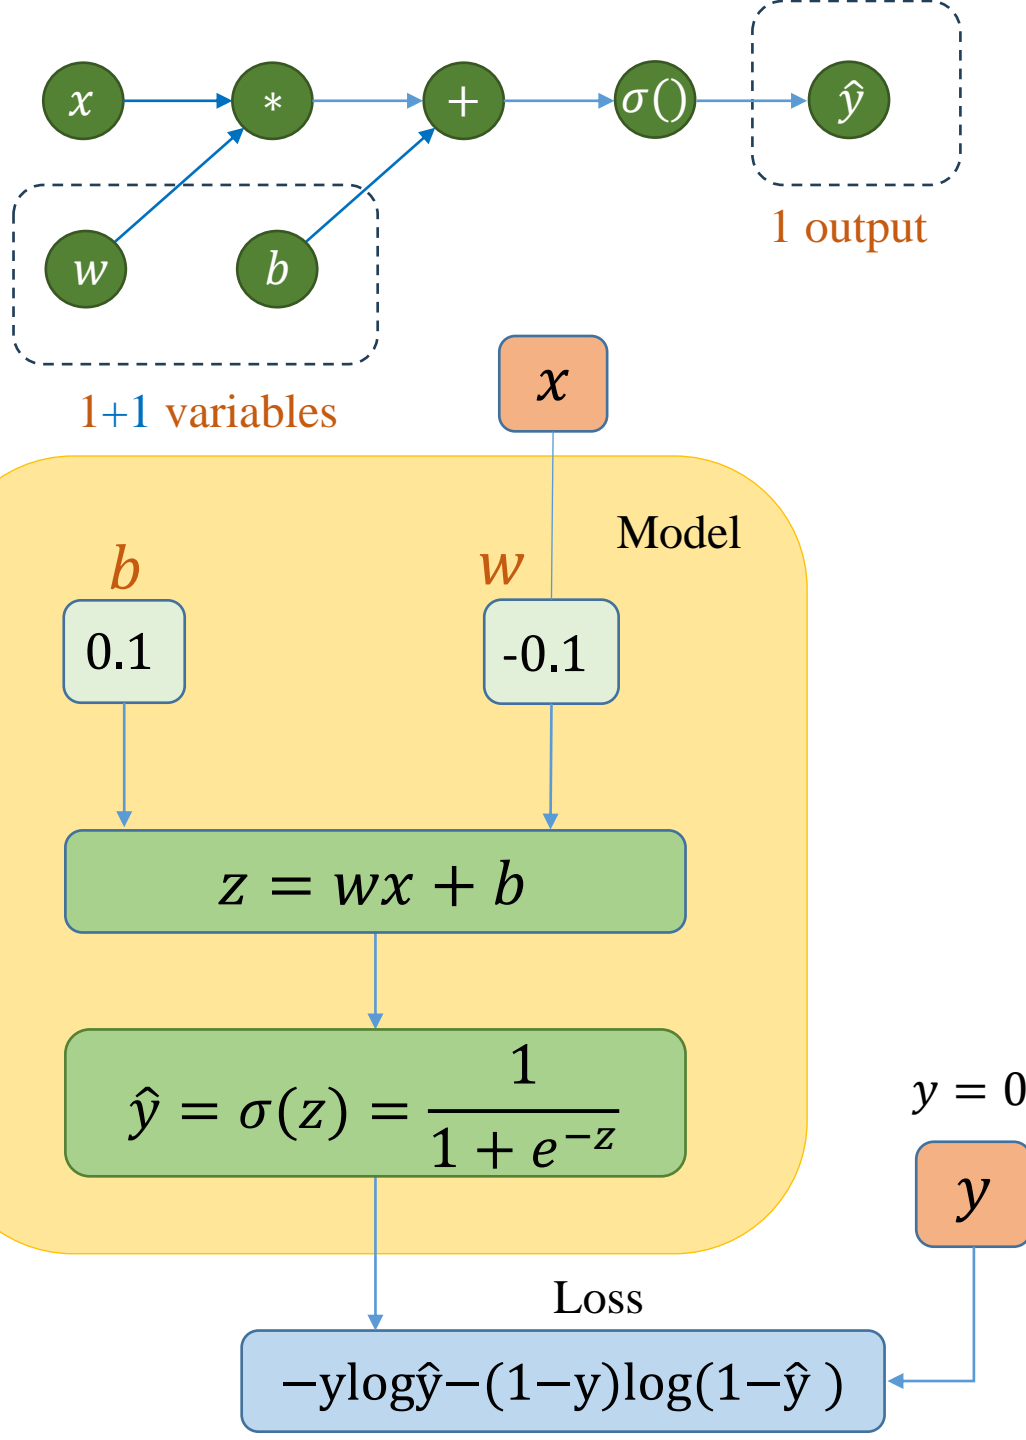
# training

```
epochs = 100
for epoch in range(epochs):
    y_hat = linear(x_data)
    loss = loss_fn(y_hat, y_data)
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```





# Logistic Regression

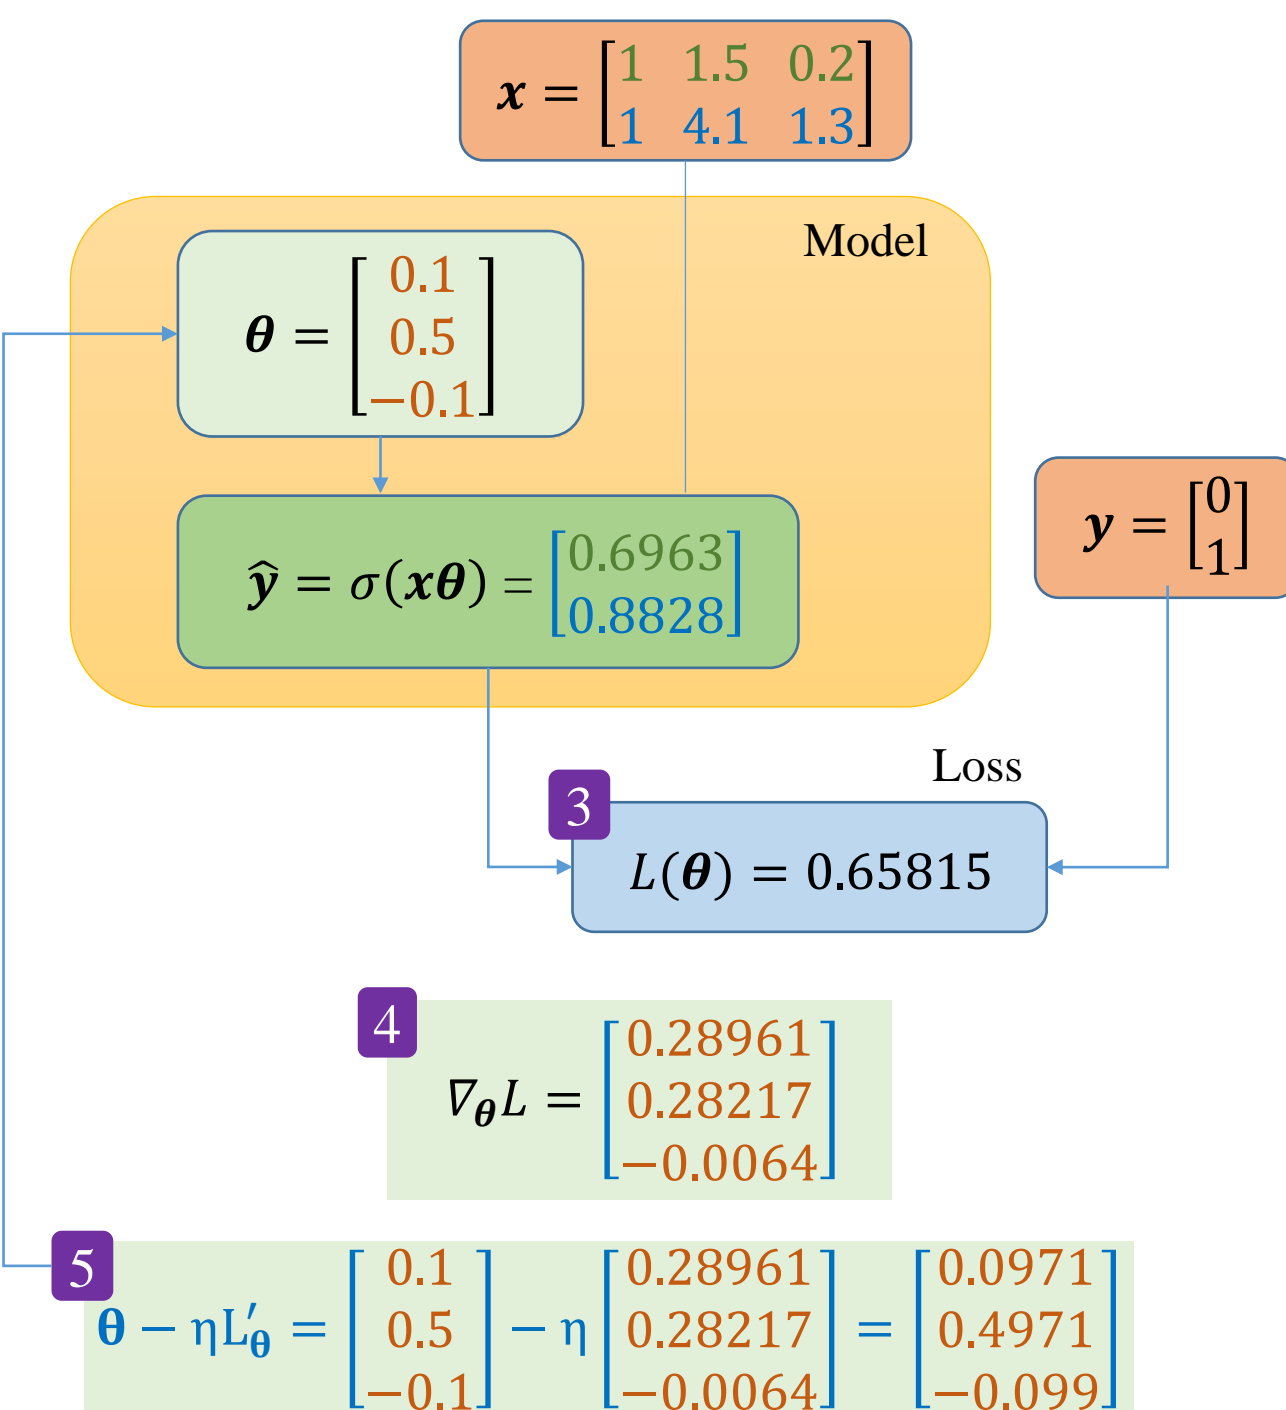


```
data = np.genfromtxt('iris_1D.csv',  
                    delimiter=',', skip_header=1)  
X = torch.from_numpy(data[:,0:1]).float()  
y = torch.from_numpy(data[:,1:]).float()
```

```
# Create a Linear Layer  
linear = nn.Linear(1, 1)  
  
# loss and optimizer  
loss_fn = torch.nn.BCELoss()  
optimizer = torch.optim.SGD(linear.parameters(),  
                             lr=0.01)
```

```
# training  
for epoch in range(epochs):  
    y_hat = torch.sigmoid(linear(X))  
    loss = loss_fn(y_hat, y)  
    optimizer.zero_grad()  
    loss.backward()  
    optimizer.step()
```

# Logistic Regression



```
data = np.genfromtxt('iris_2D_demo.csv',  
                    delimiter=',', skip_header=1)  
X = torch.from_numpy(data[:,0:2]).float()  
y = torch.from_numpy(data[:,2:]).float()
```

```
# Create a Linear Layer  
linear = nn.Linear(2, 1)  
  
# loss and optimizer  
loss_fn = torch.nn.BCELoss()  
optimizer = torch.optim.SGD(linear.parameters(),  
                             lr=0.01)
```

```
# training  
for epoch in range(epochs):  
    y_hat = torch.sigmoid(linear(X))  
    loss = loss_fn(y_hat, y)  
    optimizer.zero_grad()  
    loss.backward()  
    optimizer.step()
```



# Loss Functions

## ❖ Cross-entropy

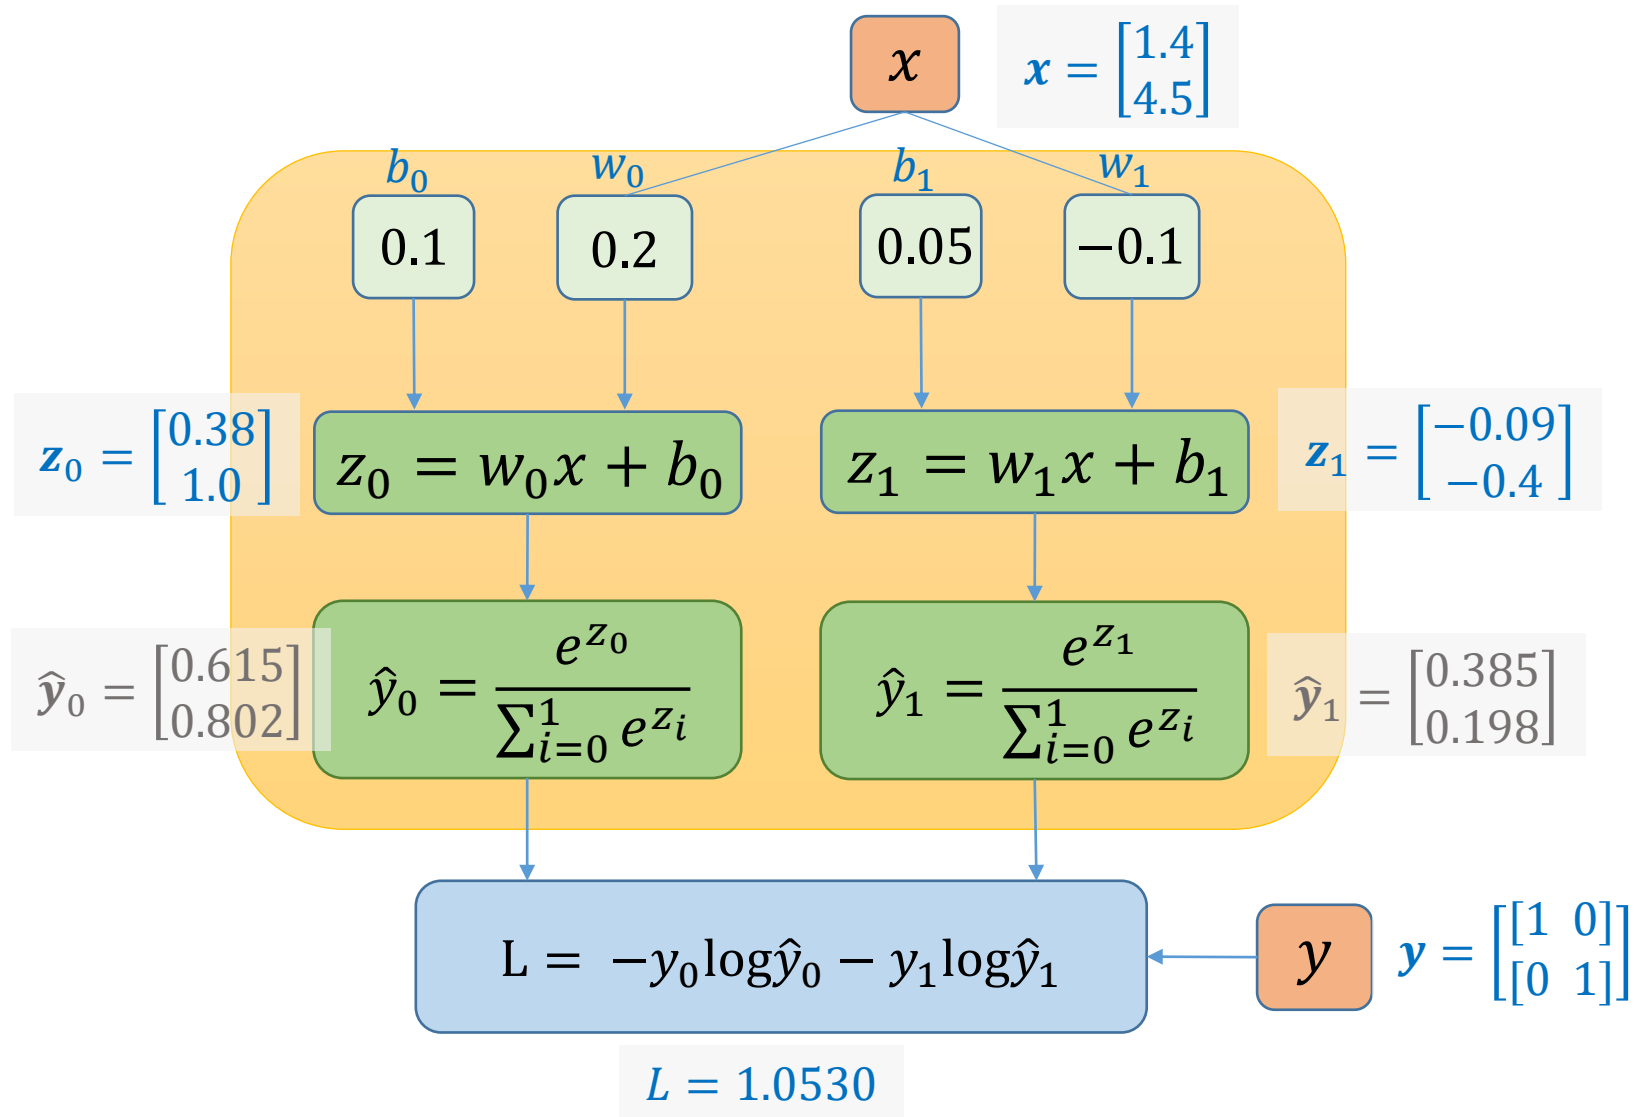
$$L(\hat{y}, y) = \sum_i -y^i \log(\hat{y}^i)$$

```
import torch.nn as nn
import torch

y_true = torch.Tensor([0, 1]).long()
logit_pred = torch.Tensor([[0.38, -0.09],
                           [1.0, -0.4]])


criterion = nn.CrossEntropyLoss()
loss = criterion(logit_pred, y_true)
print(loss)

tensor(1.0530)
```

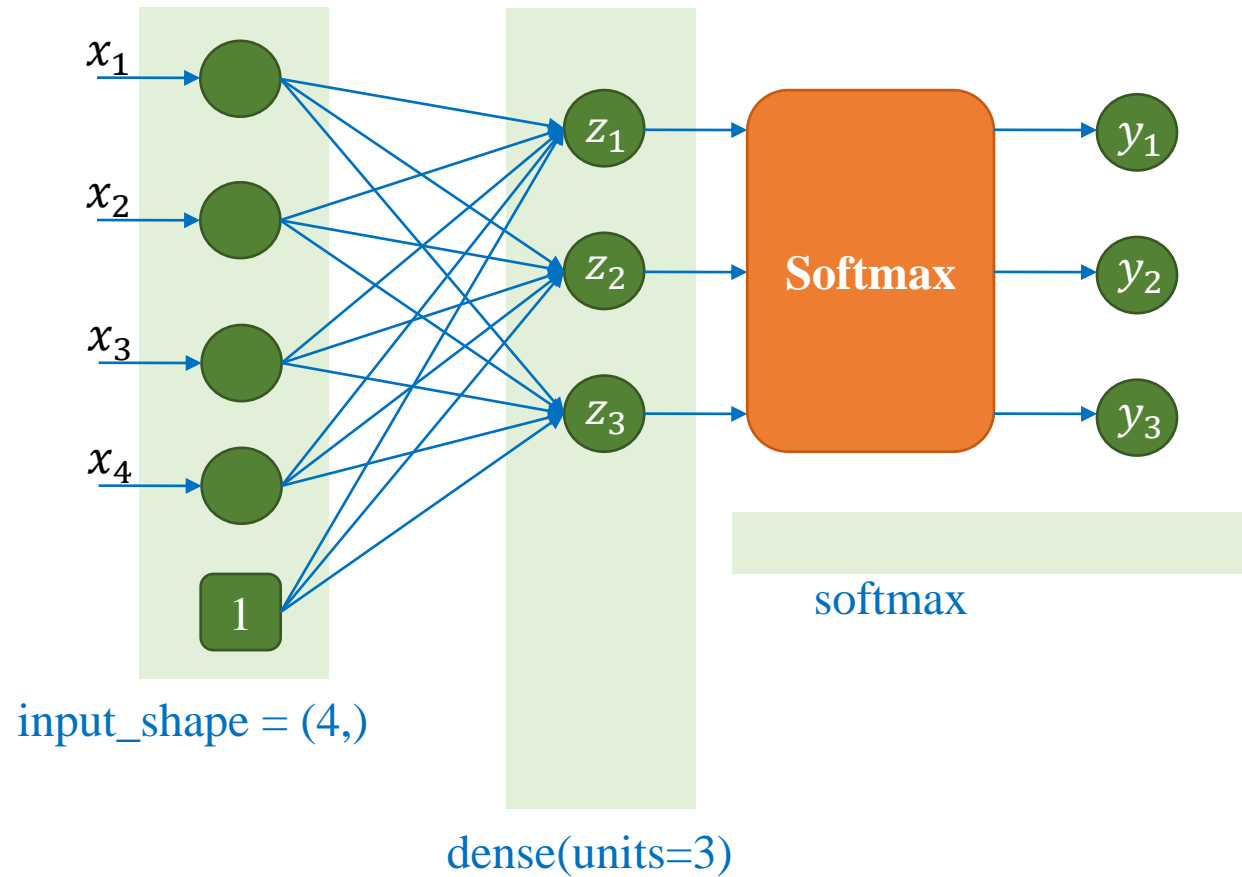


# nn.Sequential()

Model




Model





# nn.Sequential()

## Model



nn.Sequential()



```
import torch.nn as nn
model = nn.Sequential(
    nn.Linear(in_features=3, out_features=1)
)
summary(model, input_size=(3,))
```

Layer (type)

Output Shape

Param #

Linear-1

[-1, 1]

4

```
import torch.nn as nn
model = nn.Sequential(
    nn.Linear(in_features=4, out_features=3)
)
summary(model, input_size=(4,))
```

Layer (type)

Output Shape

Param #


Linear-1

[-1, 3]

15

# nn.Sequential()

## Model

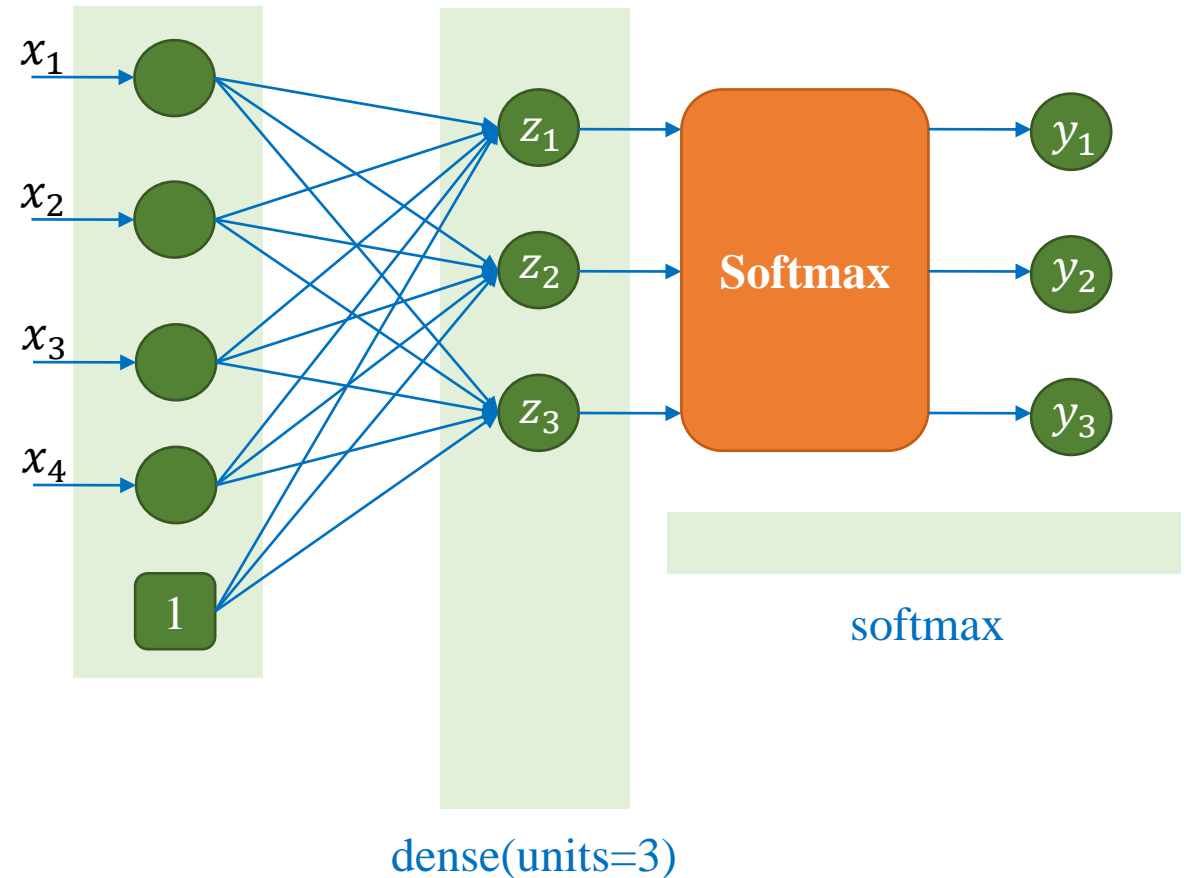


nn.Sequential()

```
import torch.nn as nn
model = nn.Sequential(
    nn.Linear(in_features=4, out_features=3)
    # don't include a softmax layer
)
summary(model, input_size=(4,))
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 3]	15

input\_shape = (4,)





# Outline

- **Introduction to PyTorch**
- **Model Construction**
- **Model Training and Inference**
- **Applying Softmax for Image Data**

# Model Construction

## ❖ Linear regression


Feature	Label
area	price
6.7	9.1
4.6	5.9
3.5	4.6
5.5	6.7

House price data

### Model

$$\text{price} = w * \text{area} + b$$

$$\hat{y} = wx + b$$



```
import torch.nn as nn
model = nn.Sequential(
    nn.Linear(in_features=1, out_features=1)
)
summary(model, input_size=(1,))
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 1]	2
Total params: 2		
Trainable params: 2		
Non-trainable params: 0		

# Model Construction

## ❖ Linear regression


Features			Label
TV	↕ Radio	↕ Newspaper	↕ Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9

Advertising-based sale data

### Model

$$\text{Sale} = w_1 * TV + w_2 * Radio + w_3 * Newspaper + b$$

$$\hat{y} = w_1x_1 + w_2x_2 + w_3x_3 + b$$



```
import torch.nn as nn
model = nn.Sequential(
    nn.Linear(in_features=3, out_features=1)
)
summary(model, input_size=(3,))
```

Layer (type)	Output Shape	Param #
Linear-1	<code>[-1, 1]</code>	4

# Model Construction

## ❖ Linear regression

Boston House  
Price Data

Features														Label													
crim	↕	zn	↕	indus	↕	chas	↕	nox	↕	rm	↕	age	↕	dis	↕	rad	↕	tax	↕	ptratio	↕	black	↕	lstat	↕	medv	↕
0.00632		18		2.31		0		0.538		6.575		65.2		4.09		1		296		15.3		396.9		4.98		24	
0.02731		0		7.07		0		0.469		6.421		78.9		4.9671		2		242		17.8		396.9		9.14		21.6	
0.03237		0		2.18		0		0.458		6.998		45.8		6.0622		3		222		18.7		394.63		2.94		33.4	
0.06905		0		2.18		0		0.458		7.147		54.2		6.0622		3		222		18.7		396.9		5.33		36.2	
0.08829		12.5		7.87		0		0.524		6.012		66.6		5.5605		5		311		15.2		395.6		12.43		22.9	

### Model


$$\text{medv} = w_1 * x_1 + \dots + w_{13} * x_{13} + b$$

# Model Construction

## ❖ Linear regression

### Model

$$\text{medv} = w_1 * x_1 + \dots + w_{13} * x_{13} + b$$



```
import torch.nn as nn
model = nn.Sequential(
    nn.Linear(in_features=13, out_features=1)
)
summary(model, input_size=(13,))
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 1]	14

Total params: 14  
Trainable params: 14  
Non-trainable params: 0

Input size (MB): 0.00  
Forward/backward pass size (MB): 0.00  
Params size (MB): 0.00  
Estimated Total Size (MB): 0.00

# Model Construction

Feature Label

Petal_Length	Category
1.4	0
1	0
1.5	0
3	1
3.8	1
4.1	1

Model

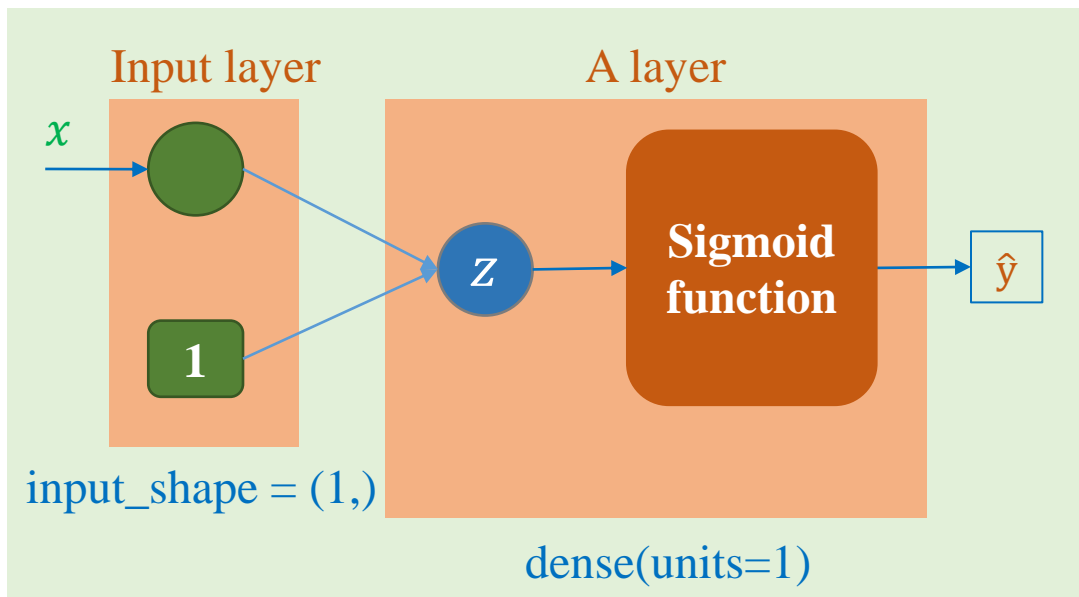
$$z = \theta^T x$$

$$\hat{y} = \frac{1}{1 + e^{-z}}$$

❖ Logistic regression

```
import torch.nn as nn
model = nn.Sequential(
    nn.Linear(in_features=1, out_features=1)
    # don't include a sigmoid layer
)
summary(model, input_size=(1,))
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 1]	2
Total params: 2		
Trainable params: 2		
Non-trainable params: 0		



criterion = nn.BCEWithLogitsLoss()

# Model Construction

## ❖ Logistic regression

Feature Label

Petal_Length	Petal_Width	Label
1.5	0.2	0
1.4	0.2	0
1.6	0.2	0
4.7	1.6	1
3.3	1.1	1
4.6	1.3	1

Model

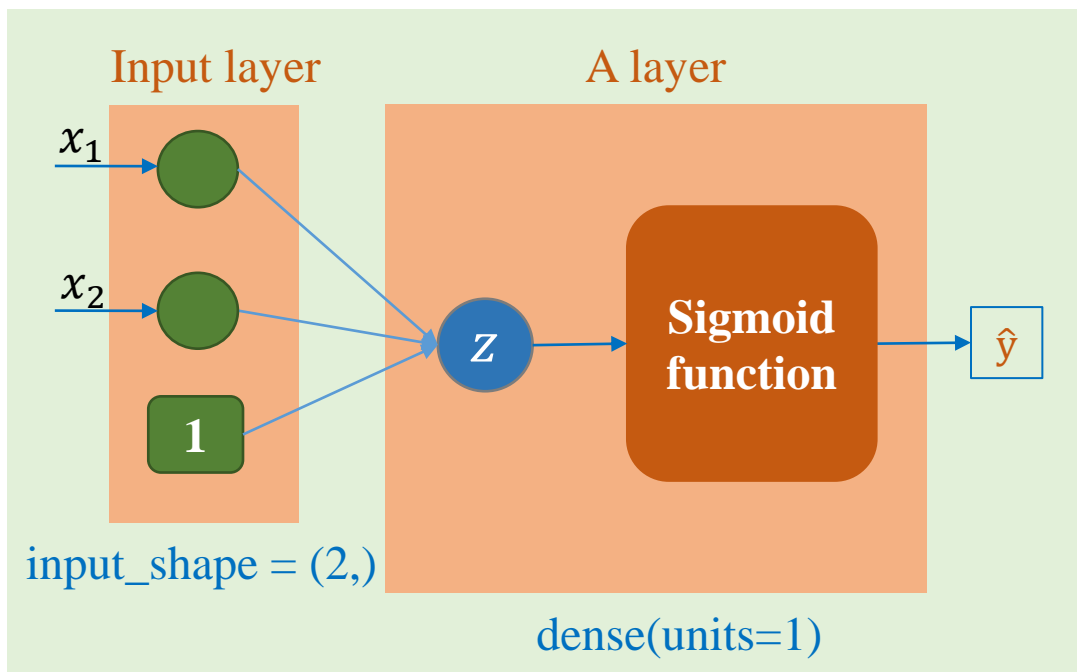
$$z = \theta^T x$$

$$\hat{y} = \frac{1}{1 + e^{-z}}$$

```
import torch.nn as nn
model = nn.Sequential(
    nn.Linear(in_features=2, out_features=1)
    # don't include a sigmoid layer
)
summary(model, input_size=(2,))
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 1]	3

Total params: 3  
Trainable params: 3  
Non-trainable params: 0



criterion = nn.BCEWithLogitsLoss()

# Model Construction

Feature

Label

Sepal_Length	Sepal_Width	Petal_Length	Petal_Width	Label
5.2	3.5	1.5	0.2	0
5.2	3.4	1.4	0.2	0
4.7	3.2	1.6	0.2	0
6.3	3.3	4.7	1.6	1
4.9	2.4	3.3	1.1	1
6.6	2.9	4.6	1.3	1

Model

$$z = \theta^T x$$

$$\hat{y} = \frac{1}{1 + e^{-z}}$$

❖ Logistic regression

```
import torch.nn as nn
model = nn.Sequential(
    nn.Linear(in_features=4, out_features=1)
    # don't include a sigmoid layer
)
summary(model, input_size=(4,))
```

Layer (type)

Output Shape

Param #

Linear-1

[-1, 1]

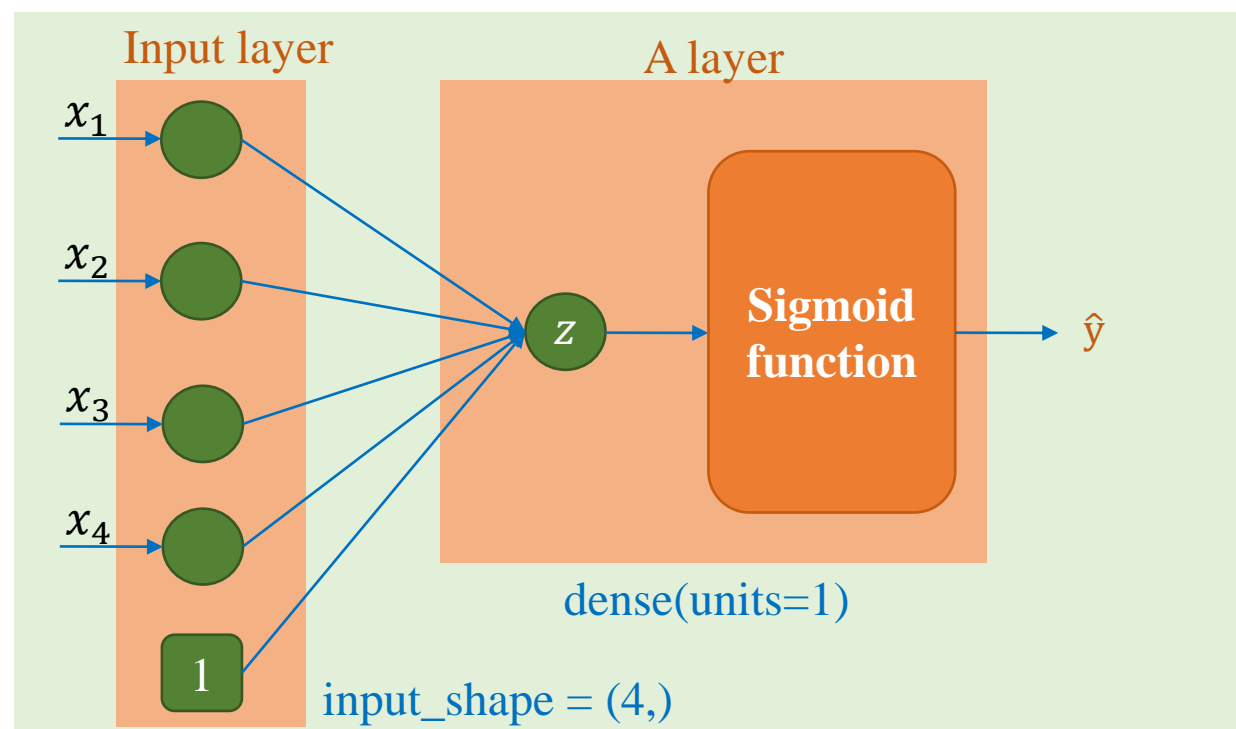
5

Total params: 5

Trainable params: 5

Non-trainable params: 0

criterion = nn.BCEWithLogitsLoss()





# Model Construction

Feature Label

Petal_Length	Label
1.4	1
1.3	1
1.5	1
4.5	2
4.1	2
4.6	2

Iris Classification Data

$$z_1 = xw_1 + b_1$$

$$z_2 = xw_2 + b_2$$

$$\hat{y}_1 = \frac{e^{z_1}}{\sum_{j=1}^2 e^{z_j}}$$

$$\hat{y}_2 = \frac{e^{z_2}}{\sum_{j=1}^2 e^{z_j}}$$

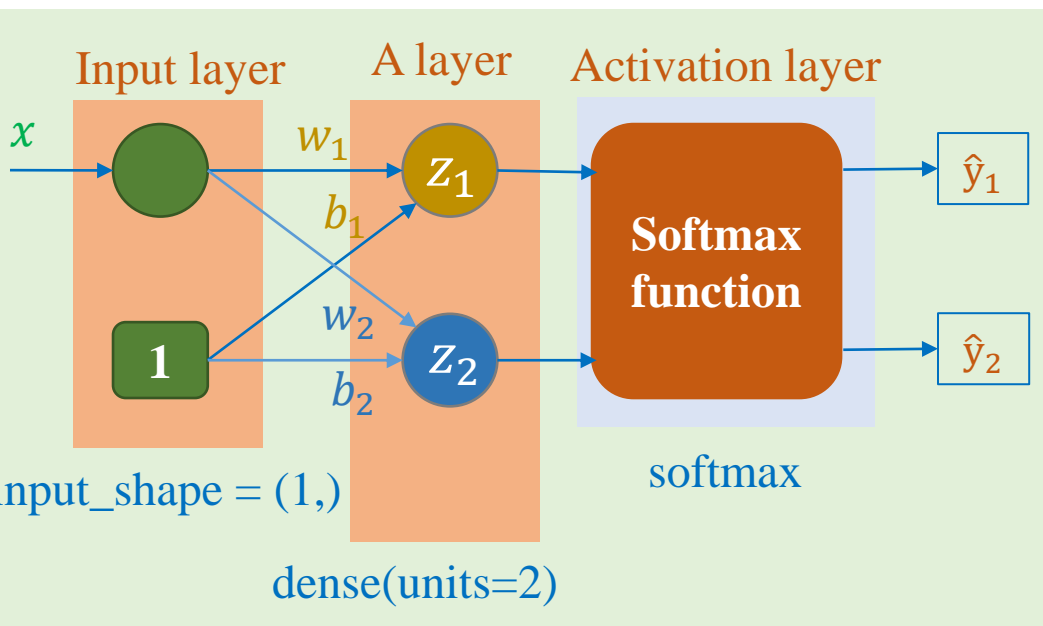
## ❖ Softmax regression

```
import torch.nn as nn
model = nn.Sequential(
    nn.Linear(in_features=1, out_features=2)
    # don't include a softmax layer
)
summary(model, input_size=(1,))
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 2]	4

Total params: 4  
Trainable params: 4  
Non-trainable params: 0

`criterion = nn.CrossEntropyLoss()`



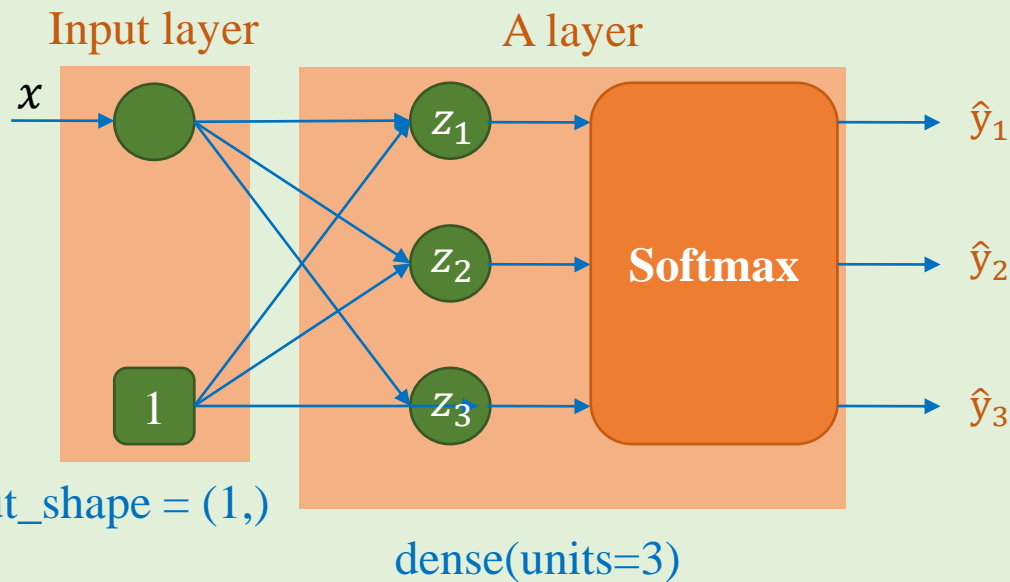
# Model Construction

## ❖ Softmax regression

Petal_Length	Label
1.4	1
1.3	1
1.5	1
4.5	2
4.1	2
4.6	2
5.2	3
5.6	3
5.9	3

```
import torch.nn as nn
model = nn.Sequential(
    nn.Linear(in_features=1, out_features=3)
    # don't include a softmax layer
)
summary(model, input_size=(1,))
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 3]	6
Total params: 6		
Trainable params: 6		
Non-trainable params: 0		



criterion = nn.CrossEntropyLoss()

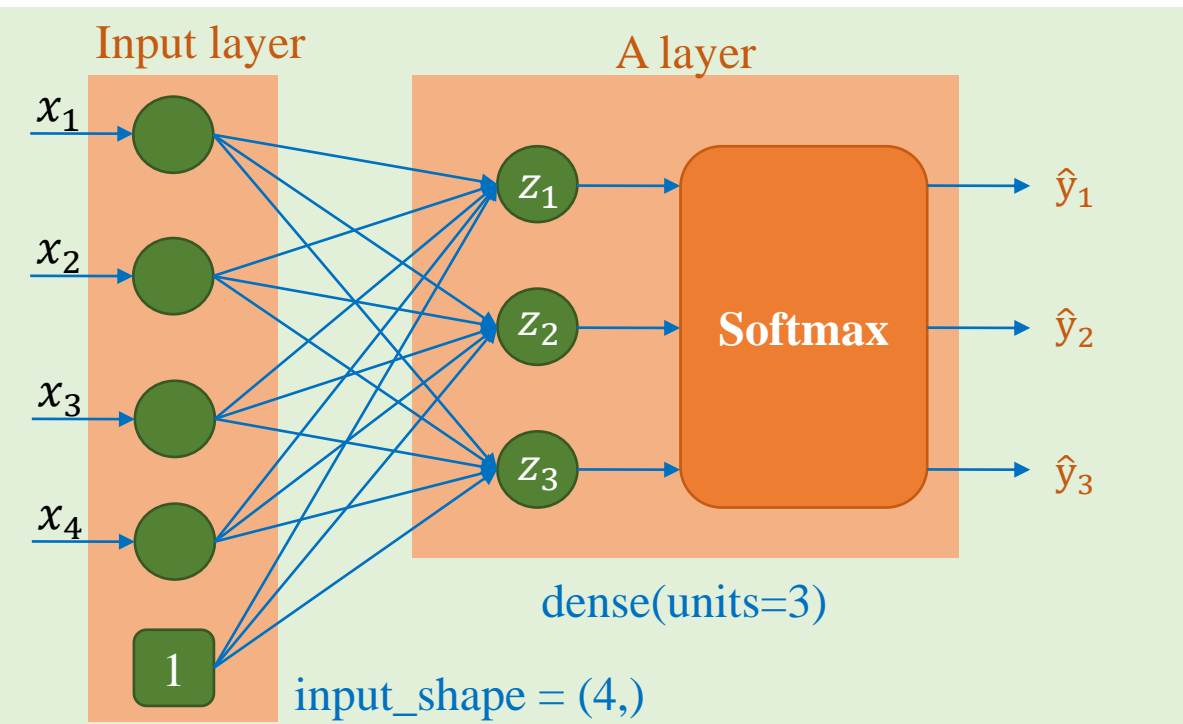
# Model Construction

Sepal_Length	Sepal_Width	Petal_Length	Petal_Width	Label
5.2	3.5	1.5	0.2	1
5.2	3.4	1.4	0.2	1
4.7	3.2	1.6	0.2	1
6.3	3.3	4.7	1.6	2
4.9	2.4	3.3	1.1	2
6.6	2.9	4.6	1.3	2
6.4	2.8	5.6	2.2	3
6.3	2.8	5.1	1.5	3
6.1	2.6	5.6	1.4	3

## ❖ Softmax regression

### Forward computation

$$z = \theta^T x \quad \hat{y} = \frac{e^z}{\sum_{i=1}^k e^{z_i}}$$



```
import torch.nn as nn
model = nn.Sequential(
    nn.Linear(in_features=4, out_features=3)
    # don't include a softmax layer
)
summary(model, input_size=(4,))
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 3]	15
Total params: 15		
Trainable params: 15		
Non-trainable params: 0		

# Outline

- **Introduction to PyTorch**
- **Model Construction**
- **Model Training and Inference**
- **Applying Softmax for Image Data**

# Training

## ❖ Logistic regression

→ Tính output  $\hat{y}$

$$z = \theta^T x$$
$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

→ Tính loss (binary cross-entropy)

$$L(\theta) = (-y^T \log \hat{y} - (1-y)^T \log(1-\hat{y}))$$

→ Tính đạo hàm

$$L'_\theta = x^T (\hat{y} - y)$$

→ Cập nhật tham số (Stochastic gradient descent)

$$\theta = \theta - \eta L'_\theta$$

Declare optimizer and loss function

```
criterion = nn.BCEWithLogitsLoss()
optimizer = optim.SGD(model.parameters(),
                        lr=learning_rate)
```

Start training

```
for epoch in range(max_epoch):
    # Zero the gradients
    optimizer.zero_grad()

    # Forward pass
    outputs = model(X)

    # Compute loss
    loss = criterion(outputs, y)

    # Backward pass and optimization
    loss.backward()
    optimizer.step()
```

# Training

## ❖ Softmax regression

→ Tính output  $\hat{y}$

$$z = \theta^T x \quad \hat{y} = \frac{e^z}{\sum_{i=1}^k e^{z_i}}$$

→ Tính loss (cross-entropy)

$$L(\theta) = - \sum_{i=1}^k y_i \log \hat{y}_i$$

→ Tính đạo hàm

$$\frac{\partial L}{\partial \theta_i} = x(\hat{y}_i - y_i)$$

→ Cập nhật tham số (Stochastic gradient descent)

$$\theta = \theta - \eta L'_\theta$$

### Declaration and Training

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(),
                        lr=learning_rate)
```

```
for epoch in range(max_epoch):
    # Zero the gradients
    optimizer.zero_grad()

    # Forward pass
    outputs = model(X)

    # Compute loss
    loss = criterion(outputs, y)

    # Backward pass and optimization
    loss.backward()
    optimizer.step()
```


# Training

## ❖ Softmax regression

Model

$$z = \theta^T x$$
$$\hat{y} = \frac{e^z}{\sum_{i=1}^k e^{z_i}}$$

Petal_Length	Label
1.4	1
1.3	1
1.5	1
4.5	2
4.1	2
4.6	2
5.2	3
5.6	3
5.9	3



```
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim

# Load data
iris = np.genfromtxt('iris_2D_3c.csv', dtype=None,
                    delimiter=',', skip_header=1)
X = torch.tensor(iris[:, 0:2], dtype=torch.float32)
y = torch.tensor(iris[:, 2], dtype=torch.int64)

# Define a simple Sequential model
input_dim = X.shape[1]
output_dim = len(torch.unique(y))
model = nn.Sequential(
    nn.Linear(in_features=input_dim, out_features=output_dim)
)

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
```


# Training

Petal_Length	Label
1.4	1
1.3	1
1.5	1
4.5	2
4.1	2
4.6	2
5.2	3
5.6	3
5.9	3

## ❖ Softmax regression

Model

$$z = \theta^T x$$
$$\hat{y} = \frac{e^z}{\sum_{i=1}^k e^{z_i}}$$



```
# Training loop
max_epoch = 100
losses = []

for epoch in range(max_epoch):
    # Zero the gradients
    optimizer.zero_grad()

    # Forward pass
    outputs = model(X)

    # Compute loss
    loss = criterion(outputs, y)
    losses.append(loss.item())

    # Backward pass and optimization
    loss.backward()
    optimizer.step()
```

```
import matplotlib.pyplot as plt
plt.plot(losses, 'g')
```




# Outline

- **Introduction to PyTorch**
- **Model Construction**
- **Model Training and Inference**
- **Applying Softmax for Image Data**

# Image Classification: Image Data

## ❖ Grayscale images



230	194	147	108	90	98	84	96	91	101
237	206	188	195	207	213	163	123	116	128
210	183	180	205	224	234	188	122	134	147
198	189	201	227	229	232	200	125	127	135
249	241	237	244	232	226	202	116	125	126
251	254	241	239	230	217	196	102	103	99
243	255	240	231	227	214	203	116	95	91
204	231	208	200	207	201	200	121	95	95
144	140	120	115	125	127	143	118	92	91
121	121	108	109	122	121	134	106	86	97

(Height, Width)

Pixel  $p$  = scalar

$$0 \leq p \leq 255$$

Resolution: #pixels

Resolution = Height $\times$ Width

## ❖ Color images



(Channel, Height, Width,)

		233	188	137	96	90	95	63	73	73	82
	237	202	159	120	105	110	88	107	112	121	109
226	191	147	110	101	112	98	123	110	119	142	131
221	191	176	182	203	214	169	144	133	145	155	122
185	160	161	184	205	223	186	137	147	161	140	115
181	174	189	207	206	215	194	136	142	151	133	87
246	237	237	231	208	206	192	122	143	144	111	74
254	254	241	224	199	192	181	99	122	117	107	74
239	248	232	207	187	182	184	110	114	110	113	74
193	215	193	167	158	164	181	114	112	111	105	82
113	119	110	111	113	123	135	120	108	106	113	
93	97	91	103	107	111	122	112	104	114		

## RGB color image

$$\text{Pixel } p = \begin{bmatrix} r \\ g \\ b \end{bmatrix}$$

$$0 \leq r, g, b \leq 255$$

## Resolution: #pixels

Resolution = HeightxWidth

# Important Packages

## ❖ Some functions

### To download a file

```
import urllib.request as req
req.urlretrieve(url, name)
```

### To open an image

```
from PIL import Image
img = Image.open(name)
```

### To show an image

```
import matplotlib.pyplot as plt
plt.imshow(img)
```

```
import urllib.request as req
from PIL import Image
import matplotlib.pyplot as plt

# download an image
req.urlretrieve('https://www.dropbox.com/s/zwy8ddkdm3thatr/nature.jpg?dl=1',
               'image.jpg')

# show the image
img = Image.open('image.jpg')
plt.imshow(img)
```

<matplotlib.image.AxesImage at 0x7f5088018b90>



# Image Data


## MNIST dataset

Grayscale images

Resolution=28x28

Training set: 60000 samples

Testing set: 10000 samples





# Image Data

## Fashion-MNIST dataset

Grayscale images

Resolution=28x28

Training set: 60000 samples

Testing set: 10000 samples


T-shirt



Trouser



Pullover



Dress




Coat



Sandal



Shirt



Sneaker




Bag



Ankle  
Boot



# Using Softmax Regression



```
X_train: (60000, 784)
y_train: (60000,)
X_test: (10000, 784)
y_test: (10000,)
```

Data Sets


```
class SoftmaxRegression(nn.Module):
    def __init__(self, input_dim, output_dim):
        super(SoftmaxRegression, self).__init__()
        self.linear = nn.Linear(input_dim, output_dim)
        self.flatten = nn.Flatten()

    def forward(self, x):
        x = self.flatten(x)
        return self.linear(x)
```

# Demo

```
# Testing the model
with torch.no_grad():
    correct = 0
    total = 0
    for images, labels in test_loader:
        outputs = model(images)
        predicted = torch.argmax(outputs.data, dim=1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
```

Training Loss over Epochs



Test Accuracy  
91.97%

```
# Load MNIST dataset
train_loader = ...
test_loader = ...

# Input and output dimensions
input_dim = 28 * 28 # MNIST images are 28x28
output_dim = 10      # 10 classes for MNIST

# Create a Sequential model
model = nn.Sequential(nn.Flatten(),
                      nn.Linear(input_dim, output_dim))

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.1)

# Training loop
num_epochs = 5
losses = [] # to store the average loss for each epoch

for epoch in range(num_epochs):
    for i, (images, labels) in enumerate(train_loader):
        outputs = model(images)
        loss = criterion(outputs, labels)
        losses.append(loss.item())

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
```



