

Multi-layer Perception


Quang-Vinh Dinh
Ph.D. in Computer Science

Outline

- **Image Data Loading Using Numpy&PyTorch**
- **Softmax+Normalization for Fashion-MNIST**
- **MLP and Examples**
- **Step-by-Step Implementation**
- **Training Strategy (optional)**

Image Classification: Image Data

❖ Grayscale images



(Height, Width)

Pixel p = scalar

$$0 \leq p \leq 255$$

Resolution: #pixels

Resolution = Height \times Width

❖ Color images



(Height, Width, channel)

		233	188	137	96	90	95	63	73	73	82
	237	202	159	120	105	110	88	107	112	121	109
226	191	147	110	101	112	98	123	110	119	142	131
221	191	176	182	203	214	169	144	133	145	155	122
185	160	161	184	205	223	186	137	147	161	140	115
181	174	189	207	206	215	194	136	142	151	133	87
246	237	237	231	208	206	192	122	143	144	111	74
254	254	241	224	199	192	181	99	122	117	107	74
239	248	232	207	187	182	184	110	114	110	113	74
193	215	193	167	158	164	181	114	112	111	105	82
113	119	110	111	113	123	135	120	108	106	113	
93	97	91	103	107	111	122	112	104	114		

RGB color image

$$\text{Pixel } p = \begin{bmatrix} r \\ g \\ b \end{bmatrix}$$

$$0 \leq r, g, b \leq 255$$

Resolution: #pixels

Resolution = HeightxWidth

Important Packages

❖ Some functions

To download a file

```
import urllib.request as req
req.urlretrieve(url, name)
```

To open an image

```
from PIL import Image
img = Image.open(name)
```

To show an image

```
import matplotlib.pyplot as plt
plt.imshow(img)
```

```
import urllib.request as req
from PIL import Image
import matplotlib.pyplot as plt

# download an image
req.urlretrieve('https://www.dropbox.com/s/zwy8ddkdm3thatr/nature.jpg?dl=1',
               'image.jpg')

# show the image
img = Image.open('image.jpg')
plt.imshow(img)
```

<matplotlib.image.AxesImage at 0x7f5088018b90>




Image Data

MNIST dataset

Grayscale images

Resolution=28x28

Training set: 60000 samples

Testing set: 10000 samples




Image Data

MNIST dataset

Grayscale images

Resolution=28x28

Training set: 60000 samples

Testing set: 10000 samples

TRAINING SET LABEL FILE (train-labels-idx1-ubyte):

[offset]	[type]	[value]	[description]
0000	32 bit integer	0x00000801(2049)	magic number (MSB first)
0004	32 bit integer	60000	number of items
0008	unsigned byte	??	label
0009	unsigned byte	??	label
.....			
xxxx	unsigned byte	??	label

The labels values are 0 to 9.

TRAINING SET IMAGE FILE (train-images-idx3-ubyte):

[offset]	[type]	[value]	[description]
0000	32 bit integer	0x00000803(2051)	magic number
0004	32 bit integer	60000	number of images
0008	32 bit integer	28	number of rows
0012	32 bit integer	28	number of columns
0016	unsigned byte	??	pixel
0017	unsigned byte	??	pixel
.....			
xxxx	unsigned byte	??	pixel

<http://yann.lecun.com/exdb/mnist/>

Image Data

Fashion-MNIST dataset

Grayscale images

Resolution=28x28


Training set: 60000 samples

Testing set: 10000 samples


T-shirt



Trouser




Pullover



Dress




Coat



Sandal



Shirt



Sneaker



Bag



Ankle
Boot



Image Classification

Fashion-MNIST dataset

Grayscale images

Resolution=28x28

Training set: 60000 samples

Testing set: 10000 samples

```
from urllib import request
```

```
filenames = ["train-images-idx3-ubyte.gz",  
             "train-labels-idx1-ubyte.gz",  
             "t10k-images-idx3-ubyte.gz",  
             "t10k-labels-idx1-ubyte.gz"]
```

```
folder = 'data_fashion_mnist/'
```

```
base_url = "http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/"
```

```
for name in filenames:
```





```
    print("Downloading " + name + "...")
```

```
    # Lưu vào folder data_fashion_mnist
```

```
    request.urlretrieve(base_url + name, folder + name)
```

```
print("Download complete.")
```

Download data

Name	Size
 t10k-images-idx3-ubyte.gz	4.4 MB
 t10k-labels-idx1-ubyte.gz	5.1 kB
 train-images-idx3-ubyte.gz	26.4 MB
 train-labels-idx1-ubyte.gz	29.5 kB









Image Data

Fashion-MNIST data

Download data

Name	Size
 t10k-images-idx3-ubyte.gz	4.4 MB
 t10k-labels-idx1-ubyte.gz	5.1 kB
 train-images-idx3-ubyte.gz	26.4 MB
 train-labels-idx1-ubyte.gz	29.5 kB



```
1 import numpy as np
2 import gzip
3
4 # load training images
5 with gzip.open('data_fashion_mnist/train-images-idx3-ubyte.gz', 'rb') as f:
6     X_train = np.frombuffer(f.read(), np.uint8, offset=16).reshape(-1, 28*28)
7
8 # load testing images
9 with gzip.open('data_fashion_mnist/t10k-images-idx3-ubyte.gz', 'rb') as f:
10     X_test = np.frombuffer(f.read(), np.uint8, offset=16).reshape(-1, 28*28)
11
12 # load training labels
13 with gzip.open('data_fashion_mnist/train-labels-idx1-ubyte.gz', 'rb') as f:
14     y_train = np.frombuffer(f.read(), np.uint8, offset=8)
15
16 # load testing labels
17 with gzip.open('data_fashion_mnist/t10k-labels-idx1-ubyte.gz', 'rb') as f:
18     y_test = np.frombuffer(f.read(), np.uint8, offset=8)
```

```
X_train: (60000, 784)
y_train: (60000,)
X_test: (10000, 784)
y_test: (10000,)
```

Demo

Image Data

❖ Using Pytorch

230	194	147	108	90	98	84	96	91	101
237	206	188	195	207	213	163	123	116	128
210	183	180	205	224	234	188	122	134	147
198	189	201	227	229	232	200	125	127	135
249	241	237	244	232	226	202	116	125	126
251	254	241	239	230	217	196	102	103	99
243	255	240	231	227	214	203	116	95	91
204	231	208	200	207	201	200	121	95	95
144	140	120	115	125	127	143	118	92	91
121	121	108	109	122	121	134	106	86	97

data
(ndarray, tensor)


Size					...				
Mode					...				
230	194	147	108	90	98	84	96	91	101
237	206	188	195	207	213	163	123	116	128
210	183	180	205	224	234	188	122	134	147
198	189	201	227	229	232	200	125	127	135
249	241	237	244	232	226	202	116	125	126
251	254	241	239	230	217	196	102	103	99
243	255	240	231	227	214	203	116	95	91
204	231	208	200	207	201	200	121	95	95
144	140	120	115	125	127	143	118	92	91
121	121	108	109	122	121	134	106	86	97

data
(ndarray, tensor)

Fashion MNIST

❖ Using Pytorch

Each sample is a tuple (PIL image, label)



```
from torchvision.datasets import FashionMNIST
trainset = FashionMNIST(root='data',
                        train=True,
                        download=True)
```

```
img, label = trainset[0]
print(type(img), label)
```

```
<class 'PIL.Image.Image'> 9
```

```
import matplotlib.pyplot as plt

img, _ = trainset[0]

plt.figure(figsize=(2,2))
plt.imshow(img, cmap='gray')
plt.axis('off') # Hide axis
plt.show()
```



Fashion MNIST

❖ Using Pytorch

Each sample is a tuple (image tensor, label)



```
from torch.utils.data import DataLoader
from torchvision import transforms

transform = transforms.Compose([transforms.ToTensor()])
trainset = FashionMNIST(root='data', train=True,
                        download=True, transform=transform)


img, label = trainset[0]
print(type(img), label)

<class 'torch.Tensor'> 9
```

```
import matplotlib.pyplot as plt

img, _ = trainset[0]
np_img = img.numpy()
np_img = np.transpose(np_img, (1, 2, 0))


plt.figure(figsize=(2,2))
plt.imshow(np_img, cmap='gray')
plt.axis('off')
plt.show()
```



Fashion MNIST

❖ Using Pytorch

Each sample is a tuple (image tensor, label)



```
from torchvision.datasets import FashionMNIST
from torch.utils.data import DataLoader
from torchvision import transforms

transform = transforms.Compose([transforms.ToTensor()])
trainset = FashionMNIST(root='data',
                        train=True,
                        download=True,
                        transform=transform)

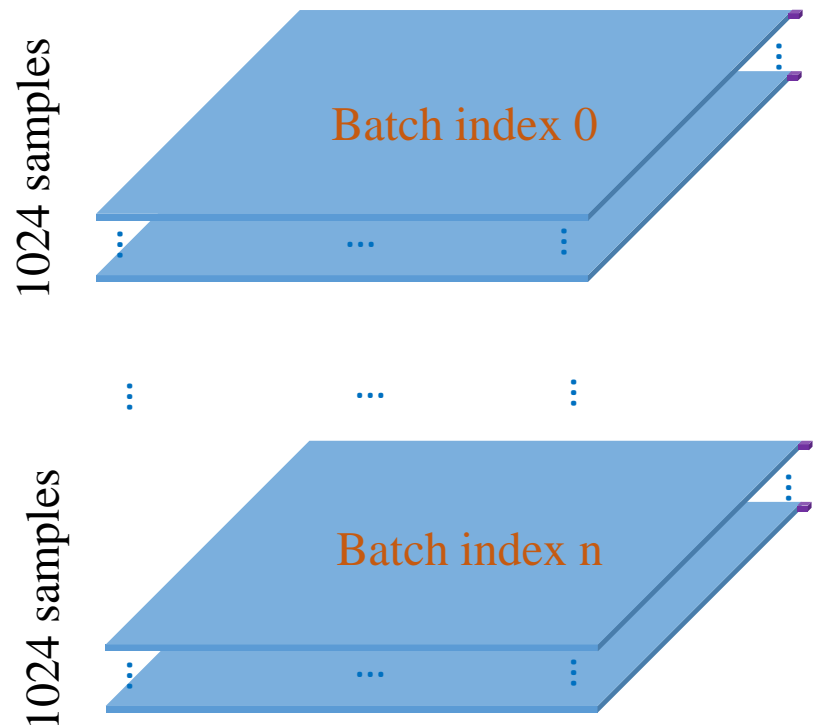
trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=2,
                        shuffle=True)

print(len(trainloader))
```

Fashion MNIST

❖ Using Pytorch

Each sample is a tuple
(image tensor, label)



```
# batch_size=3500
for i, (inputs, labels) in enumerate(trainloader, 0):
    print(f'Batch index {i} -- {inputs.shape} -- {labels.shape}')
```

```
Batch index 0 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 1 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 2 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 3 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 4 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 5 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 6 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 7 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 8 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 9 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 10 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 11 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 12 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 13 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 14 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 15 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 16 -- torch.Size([3500, 1, 28, 28]) -- torch.Size([3500])
Batch index 17 -- torch.Size([500, 1, 28, 28]) -- torch.Size([500])
```

Outline


- **Image Data Loading Using Numpy&PyTorch**
- **Softmax+Normalization for Fashion-MNIST**
- **MLP and Examples**
- **Step-by-Step Implementation**
- **Training Strategy (optional)**

Petal_Length	Label
1.4	0
1.3	0
1.5	0
4.5	1
4.1	1
4.6	1

$$x = \begin{bmatrix} 1 \\ x \end{bmatrix}$$

$$\theta = \begin{bmatrix} b_0 & b_1 \\ w_0 & w_1 \end{bmatrix}$$

One-hot encoding for label



$$z_0 = xw_0 + b_0$$

$$z_1 = xw_1 + b_1$$

$$\hat{y}_0 = \frac{e^{z_0}}{\sum_{j=0}^1 e^{z_j}}$$

$$\hat{y}_1 = \frac{e^{z_1}}{\sum_{j=0}^1 e^{z_j}}$$


Softmax Regression

$$\mathbf{z} = \begin{bmatrix} z_0 \\ z_1 \end{bmatrix} = \begin{bmatrix} b_0 & w_0 \\ b_1 & w_1 \end{bmatrix} \begin{bmatrix} 1 \\ x \end{bmatrix} = \begin{bmatrix} \theta_0^T \\ \theta_1^T \end{bmatrix} \begin{bmatrix} 1 \\ x \end{bmatrix} = \theta^T x$$

$$\hat{\mathbf{y}} = \begin{bmatrix} \hat{y}_0 \\ \hat{y}_1 \end{bmatrix} = \frac{1}{\sum_{j=0}^1 e^{z_j}} \begin{bmatrix} e^{z_0} \\ e^{z_1} \end{bmatrix} = \frac{e^{\mathbf{z}}}{\sum_{j=0}^1 e^{z_j}}$$

$$L(\theta) = - \sum_{i=0}^1 y_i \log \hat{y}_i = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

Model



Derivative

$$\frac{\partial L}{\partial \hat{y}_i} = -\frac{y_i}{\hat{y}_i}$$

$$\frac{\partial \hat{y}_i}{\partial z_j} = \begin{cases} \hat{y}_i(1 - \hat{y}_i) & \text{if } i = j \\ -\hat{y}_i \hat{y}_j & \text{if } i \neq j \end{cases}$$

$$\frac{\partial L}{\partial z_i} = \hat{y}_i - y_i$$

$$\frac{\partial L}{\partial w_i} = x(\hat{y}_i - y_i)$$

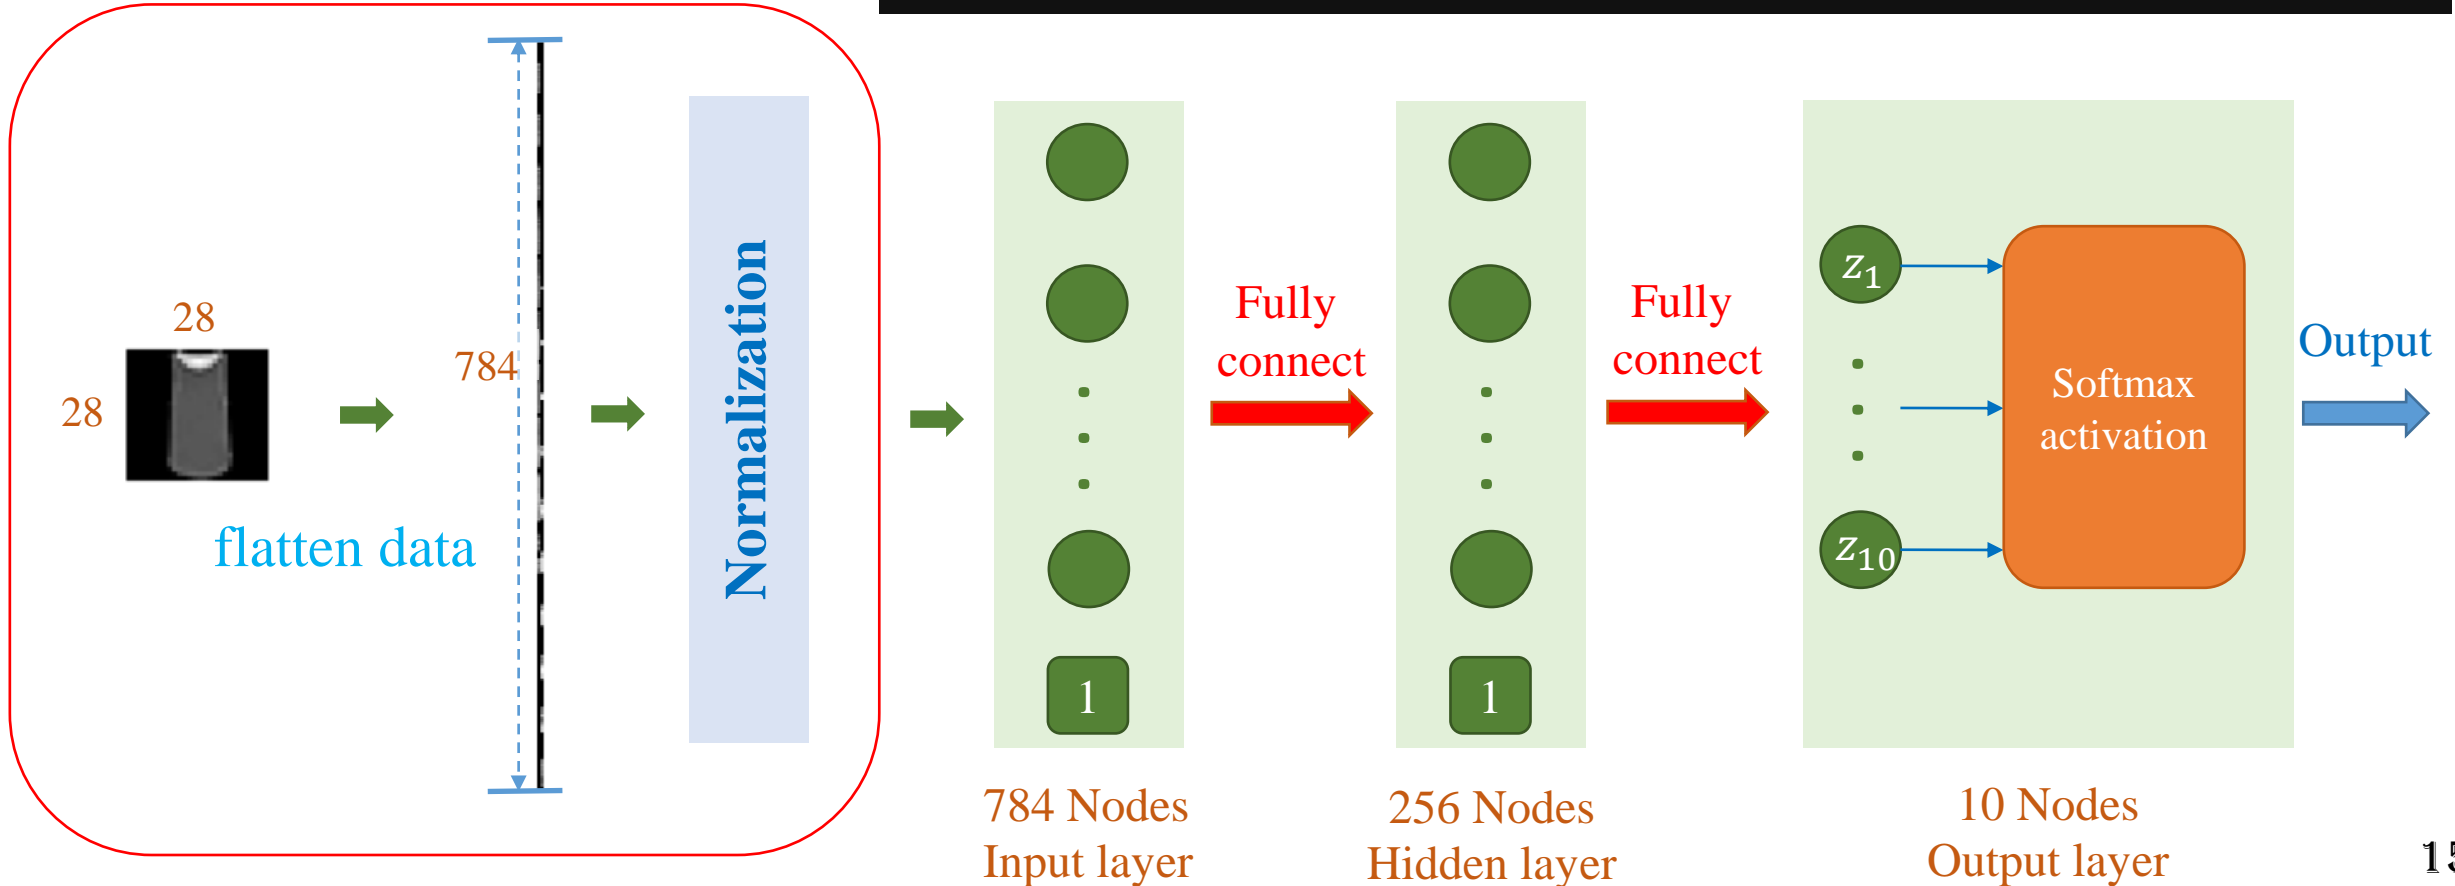
$$\frac{\partial L}{\partial b_i} = \hat{y}_i - y_i$$

Where to put Flatten

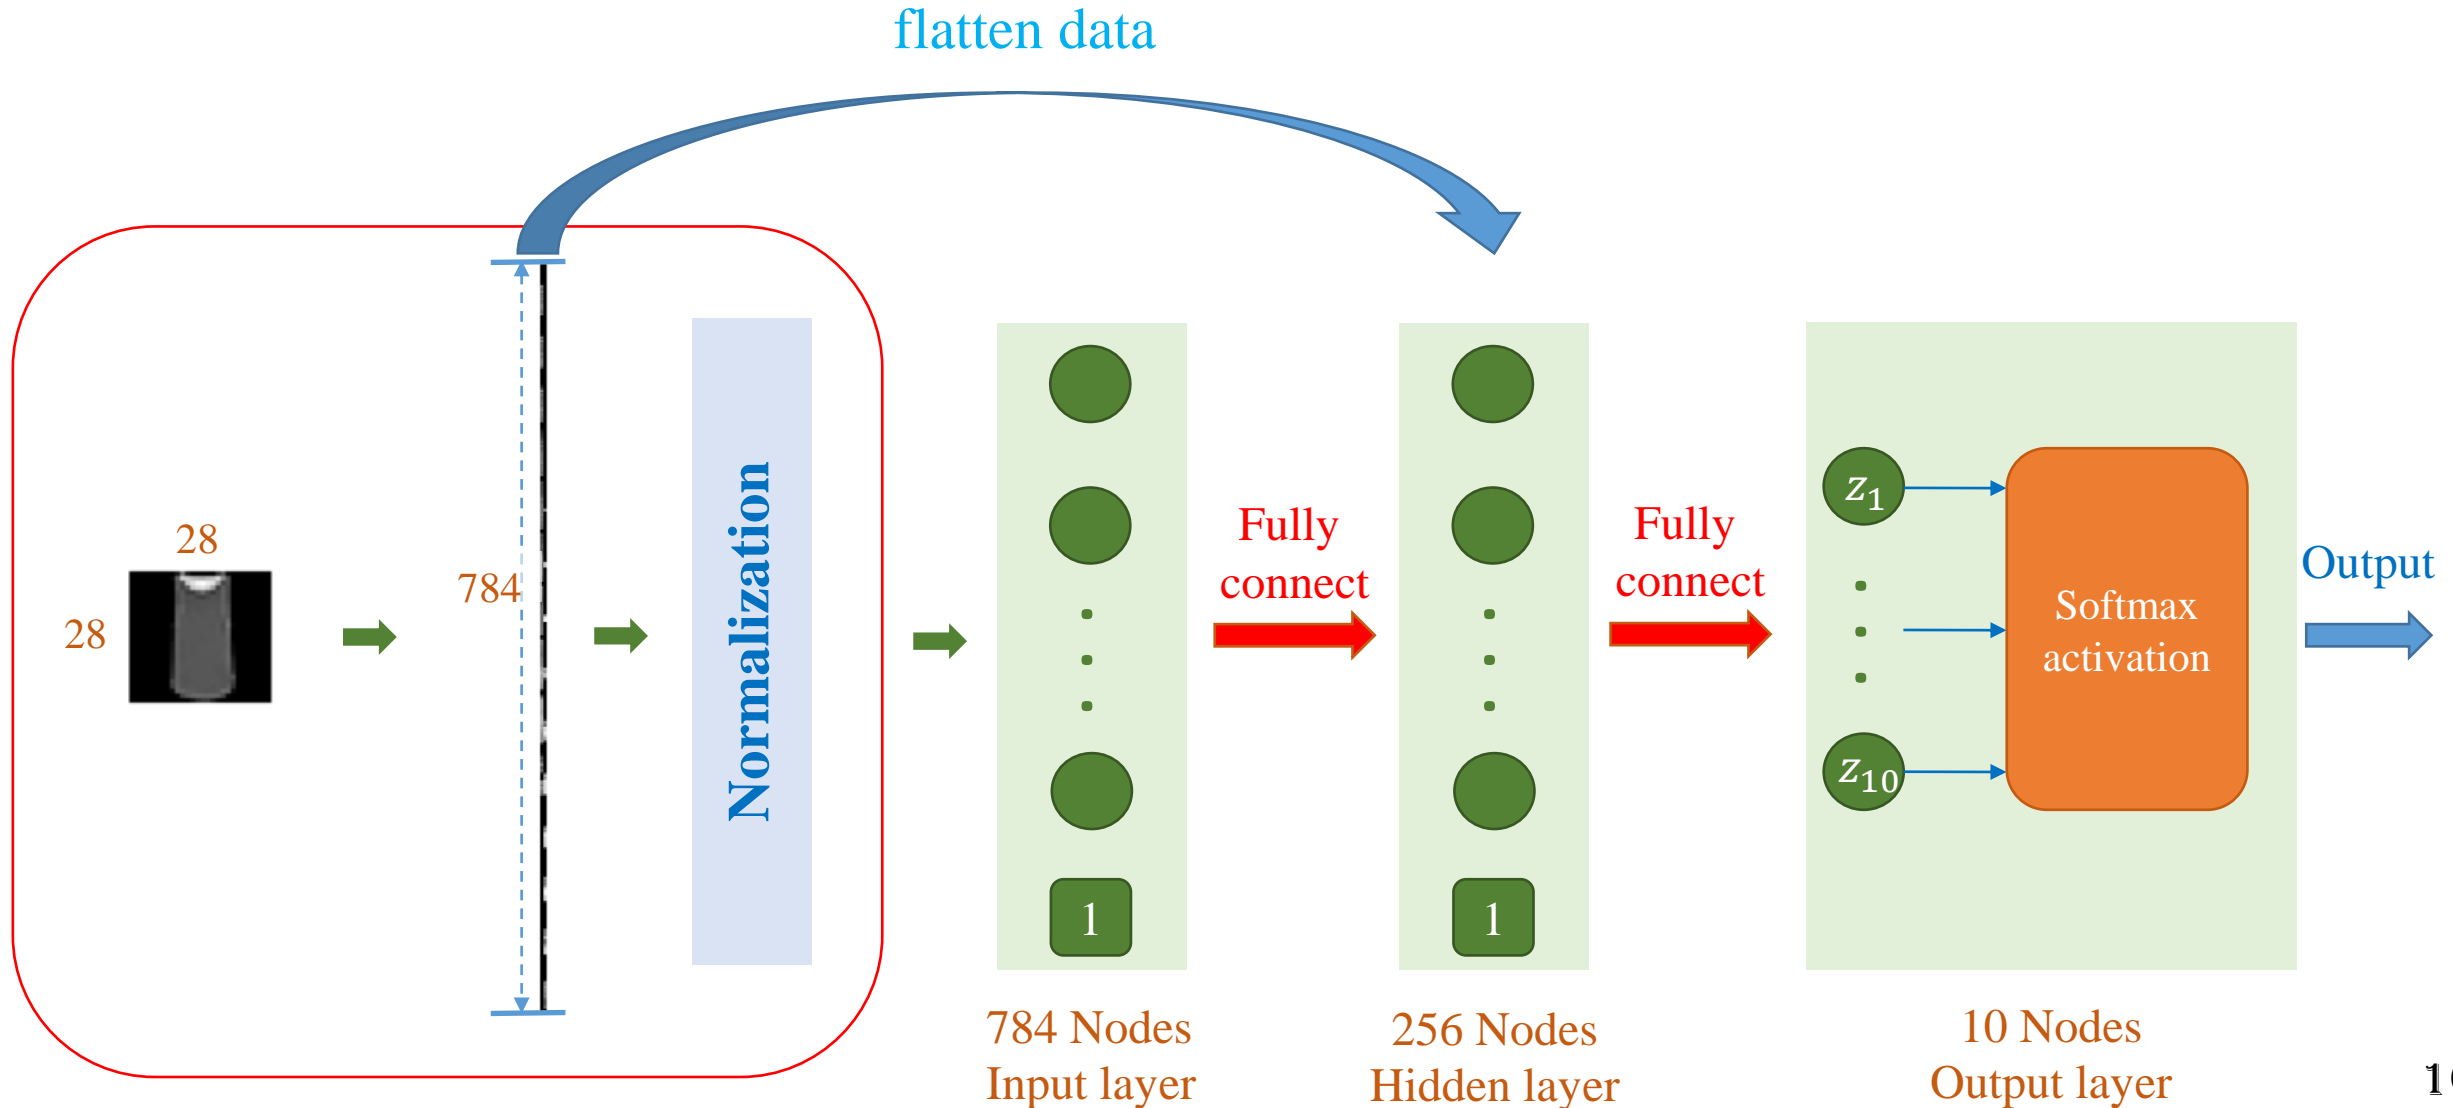
```
transform = transforms.Compose([transforms.ToTensor(),
                                transforms.Lambda(torch.flatten)])
trainset = FashionMNIST(root='data', train=True,
                        download=True, transform=transform)

img, _ = trainset[0]
print(img.shape)

torch.Size([784])
```



Where to put Flatten




Where to put Flatten

```
transform = transforms.Compose([transforms.ToTensor()])  
trainset = FashionMNIST(root='data', train=True,  
                        download=True, transform=transform)
```

```
img, label = trainset[0]  
print(img.shape)
```


```
torch.Size([1, 28, 28])
```



Softmax Regression

without
normalization

learning rate = 0.01



```
X_train: (60000, 784)
y_train: (60000,)
X_test: (10000, 784)
y_test: (10000,)
```

Data Sets

```
import torch.nn as nn



model = nn.Sequential(
    nn.Flatten(), nn.Linear(784, 10)
)
print(model)
```

```
Sequential(
  (0): Flatten(start_dim=1, end_dim=-1)
  (1): Linear(in_features=784, out_features=10, bias=True)
)
```

```
# Generating a random tensor
input_tensor = torch.rand(5, 28, 28)

# Feeding the tensor into the model
output = model(input_tensor)
print(output.shape)
```

```
torch.Size([5, 10])
```



```
model = nn.Sequential(nn.Flatten(), nn.Linear(784, 10))
model = model.to(device)
```

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
```

```
# train
for epoch in range(max_epoch):
    for i, (inputs, labels) in enumerate(trainloader, 0):
        # Move inputs and labels to the device
        inputs, labels = inputs.to(device), labels.to(device)

        # Zero the parameter gradients
        optimizer.zero_grad()

        # Forward pass
        outputs = model(inputs)
        loss = criterion(outputs, labels)


        # Determine class predictions and track accuracy
        _, predicted = torch.max(outputs.data, 1)
        correct += (predicted == labels).sum().item()

        # Backward pass and optimization
        loss.backward()
        optimizer.step()
```

Softmax Regression

without
normalization

learning rate = 0.00001



```
X_train: (60000, 784)
y_train: (60000,)
X_test: (10000, 784)
y_test: (10000,)
```

Data Sets

```
import torch.nn as nn



model = nn.Sequential(
    nn.Flatten(), nn.Linear(784, 10)
)
print(model)
```

```
Sequential(
  (0): Flatten(start_dim=1, end_dim=-1)
  (1): Linear(in_features=784, out_features=10, bias=True)
)
```

```
# Generating a random tensor
input_tensor = torch.rand(5, 28, 28)

# Feeding the tensor into the model
output = model(input_tensor)
print(output.shape)
```

```
torch.Size([5, 10])
```



```
model = nn.Sequential(nn.Flatten(), nn.Linear(784, 10))
model = model.to(device)
```

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.00001)
```

```
# train
for epoch in range(max_epoch):
    for i, (inputs, labels) in enumerate(trainloader, 0):
        # Move inputs and labels to the device
        inputs, labels = inputs.to(device), labels.to(device)

        # Zero the parameter gradients
        optimizer.zero_grad()

        # Forward pass
        outputs = model(inputs)
        loss = criterion(outputs, labels)

        # Determine class predictions and track accuracy
        _, predicted = torch.max(outputs.data, 1)
        correct += (predicted == labels).sum().item()

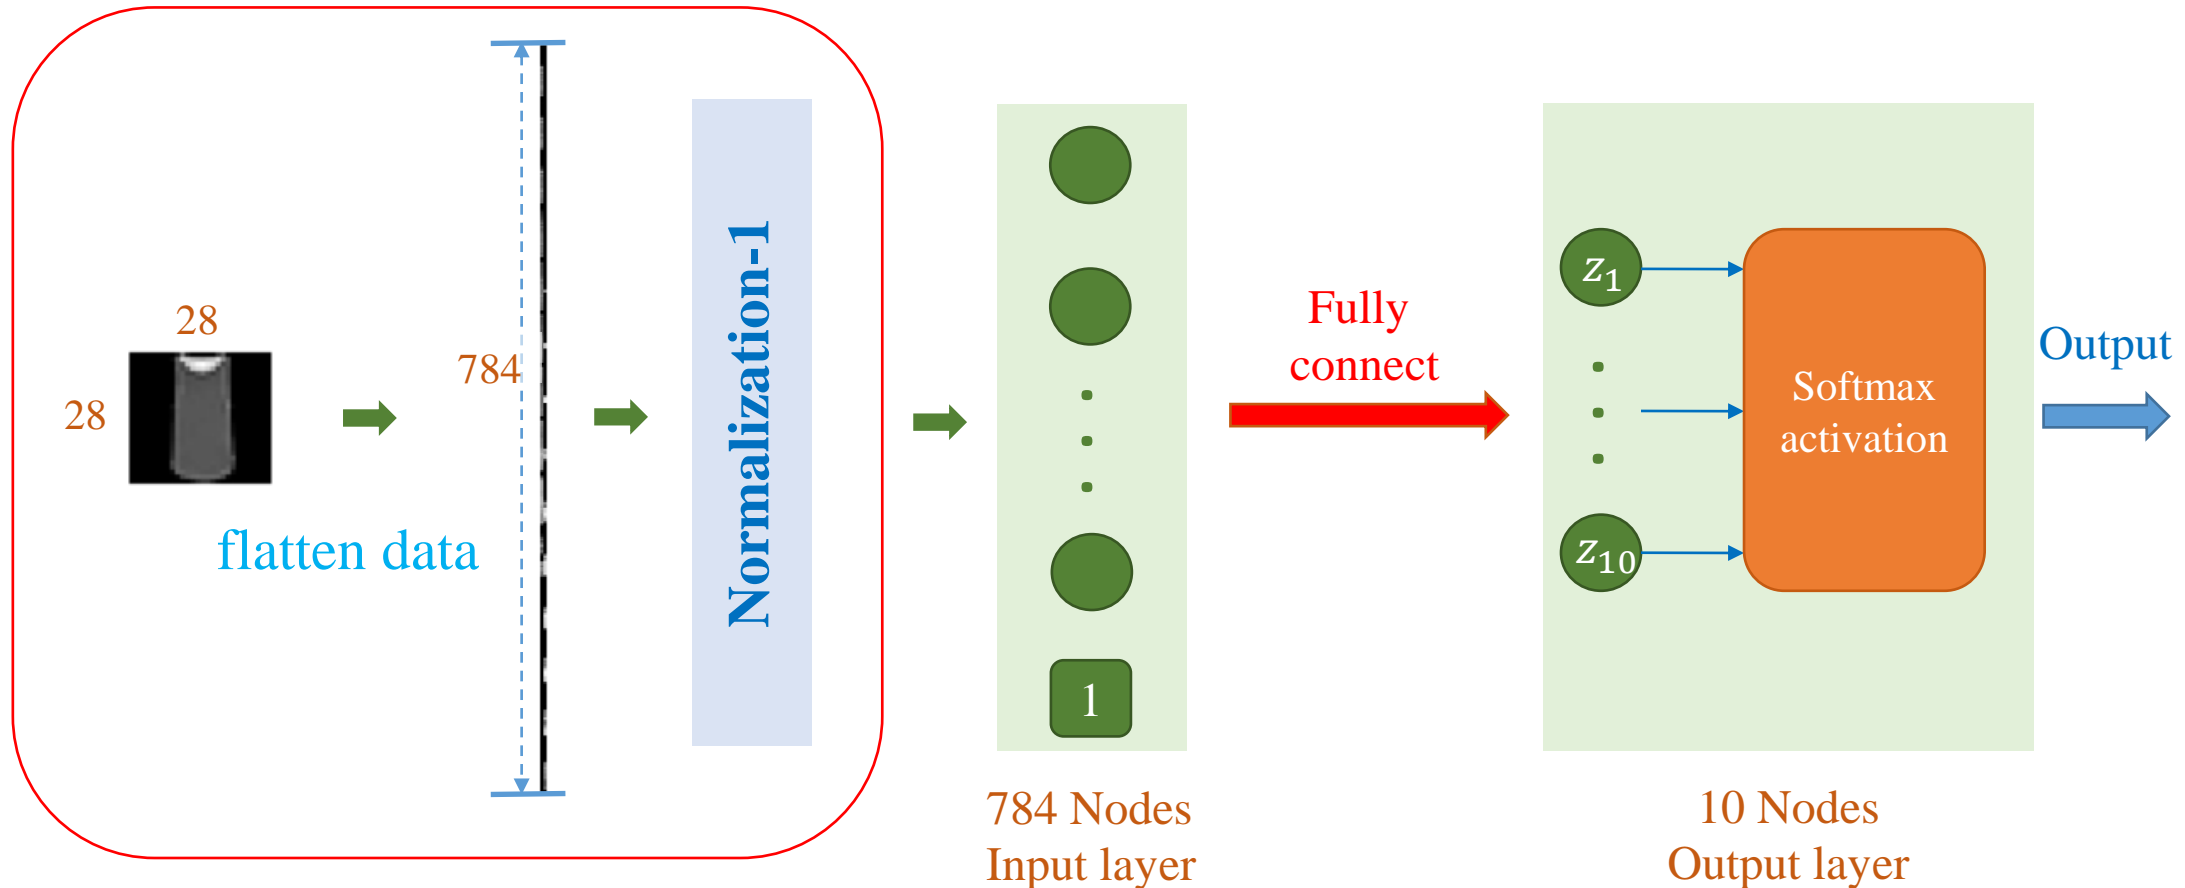
        # Backward pass and optimization
        loss.backward()
        optimizer.step()
```



Softmax Regression + Normalization

$$\text{Image} = \frac{\text{Image}}{255}$$

```
import torchvision.transforms as transforms

transform = transforms.Compose([transforms.ToTensor()])
trainset = torchvision.datasets.FashionMNIST(root='data',
                                             train=True,
                                             download=True,
                                             transform=transform)
```





```
model = nn.Sequential(nn.Flatten(), nn.Linear(784, 10))
model = model.to(device)
```

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
```

```
# train
for epoch in range(max_epoch):
    for i, (inputs, labels) in enumerate(trainloader, 0):
        # Move inputs and labels to the device
        inputs, labels = inputs.to(device), labels.to(device)

        # Zero the parameter gradients
        optimizer.zero_grad()


        # Forward pass
        outputs = model(inputs)
        loss = criterion(outputs, labels)



        # Determine class predictions and track accuracy
        _, predicted = torch.max(outputs.data, 1)
        correct += (predicted == labels).sum().item()

        # Backward pass and optimization
        loss.backward()
        optimizer.step()
```

Softmax Regression + Normalization

$$\text{Image} = \frac{\text{Image}}{127.5} - 1$$

[illegible]



```
model = nn.Sequential(nn.Flatten(), nn.Linear(784, 10))  
model = model.to(device)
```

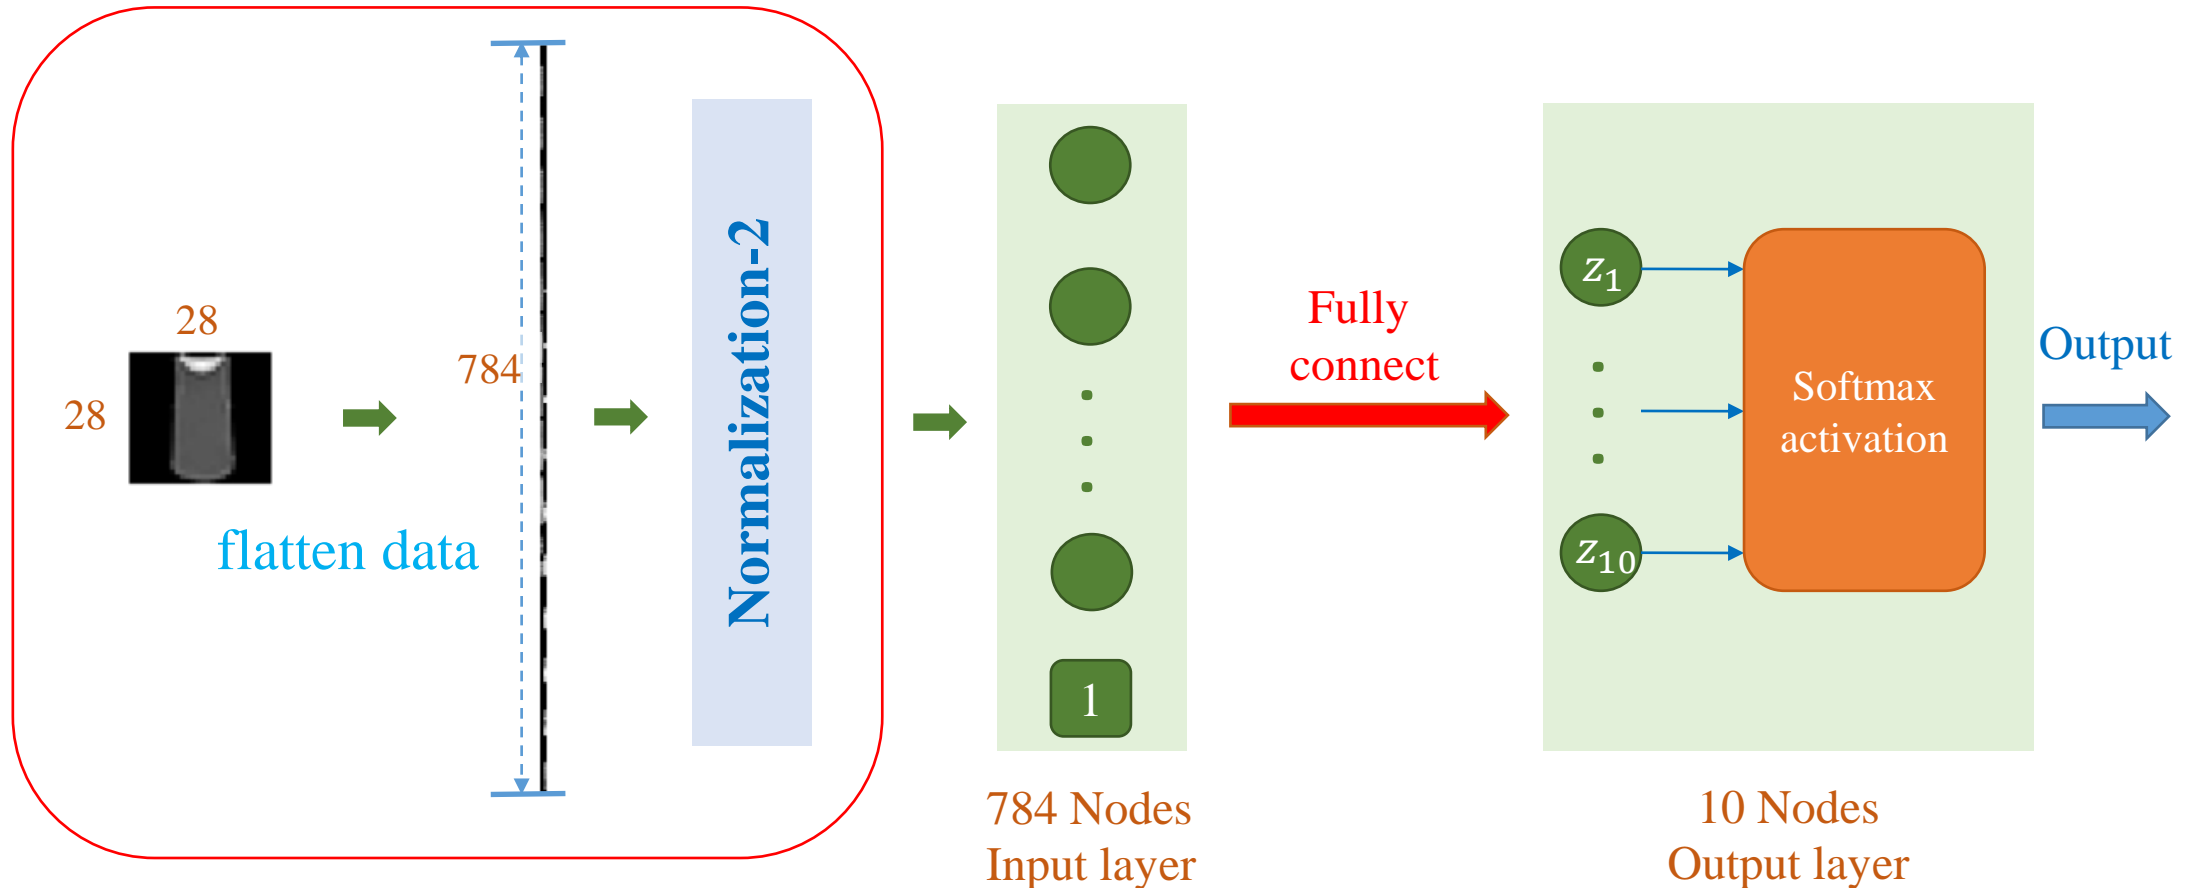
```
criterion = nn.CrossEntropyLoss()  
optimizer = optim.SGD(model.parameters(), lr=0.01)
```



```
# train  
for epoch in range(max_epoch):  
    for i, (inputs, labels) in enumerate(trainloader, 0):  
        # Move inputs and labels to the device  
        inputs, labels = inputs.to(device), labels.to(device)  
  
        # Zero the parameter gradients  
        optimizer.zero_grad()  
  
        # Forward pass  
        outputs = model(inputs)  
        loss = criterion(outputs, labels)  
  
        # Determine class predictions and track accuracy  
        _, predicted = torch.max(outputs.data, 1)  
        correct += (predicted == labels).sum().item()  
  
        # Backward pass and optimization  
        loss.backward()  
        optimizer.step()
```

Softmax Regression + Normalization

$$\text{Image} = \frac{\text{Image} - \mu}{\sigma}$$

```
transform = transforms.Compose([transforms.ToTensor(),  
                                transforms.Normalize((mean,),  
                                                    (std,))])  
  
trainset = torchvision.datasets.FashionMNIST(root='data',  
                                              train=True,  
                                              download=True,  
                                              transform=transform)
```





```
model = nn.Sequential(nn.Flatten(), nn.Linear(784, 10))
model = model.to(device)
```

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
```

```
# train
for epoch in range(max_epoch):
    for i, (inputs, labels) in enumerate(trainloader, 0):
        # Move inputs and labels to the device
        inputs, labels = inputs.to(device), labels.to(device)

        # Zero the parameter gradients
        optimizer.zero_grad()

        # Forward pass
        outputs = model(inputs)
        loss = criterion(outputs, labels)

        # Determine class predictions and track accuracy
        _, predicted = torch.max(outputs.data, 1)
        correct += (predicted == labels).sum().item()

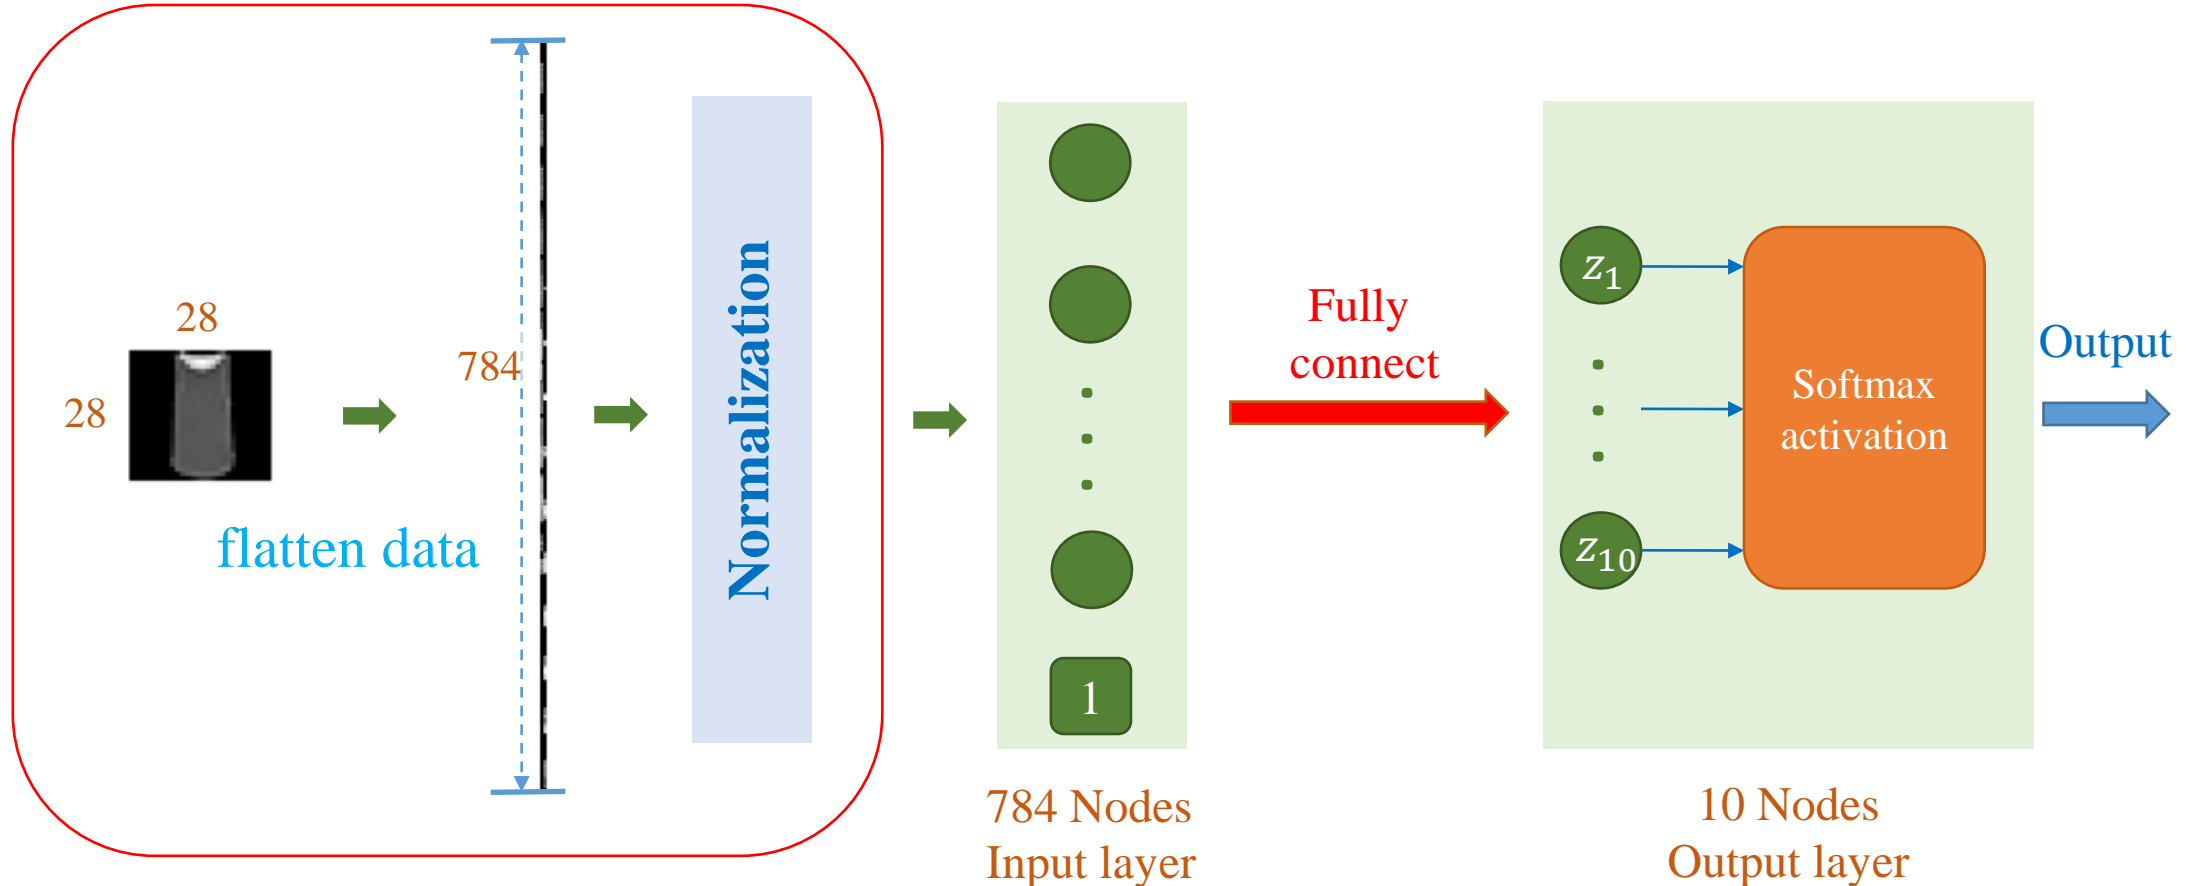
        # Backward pass and optimization
        loss.backward()
        optimizer.step()
```

Outline

- **Image Data Loading Using Numpy&PyTorch**
- **Softmax+Normalization for Fashion-MNIST**
- **MLP and Examples**
- **Step-by-Step Implementation**
- **Training Strategy (optional)**

MLP - Motivation

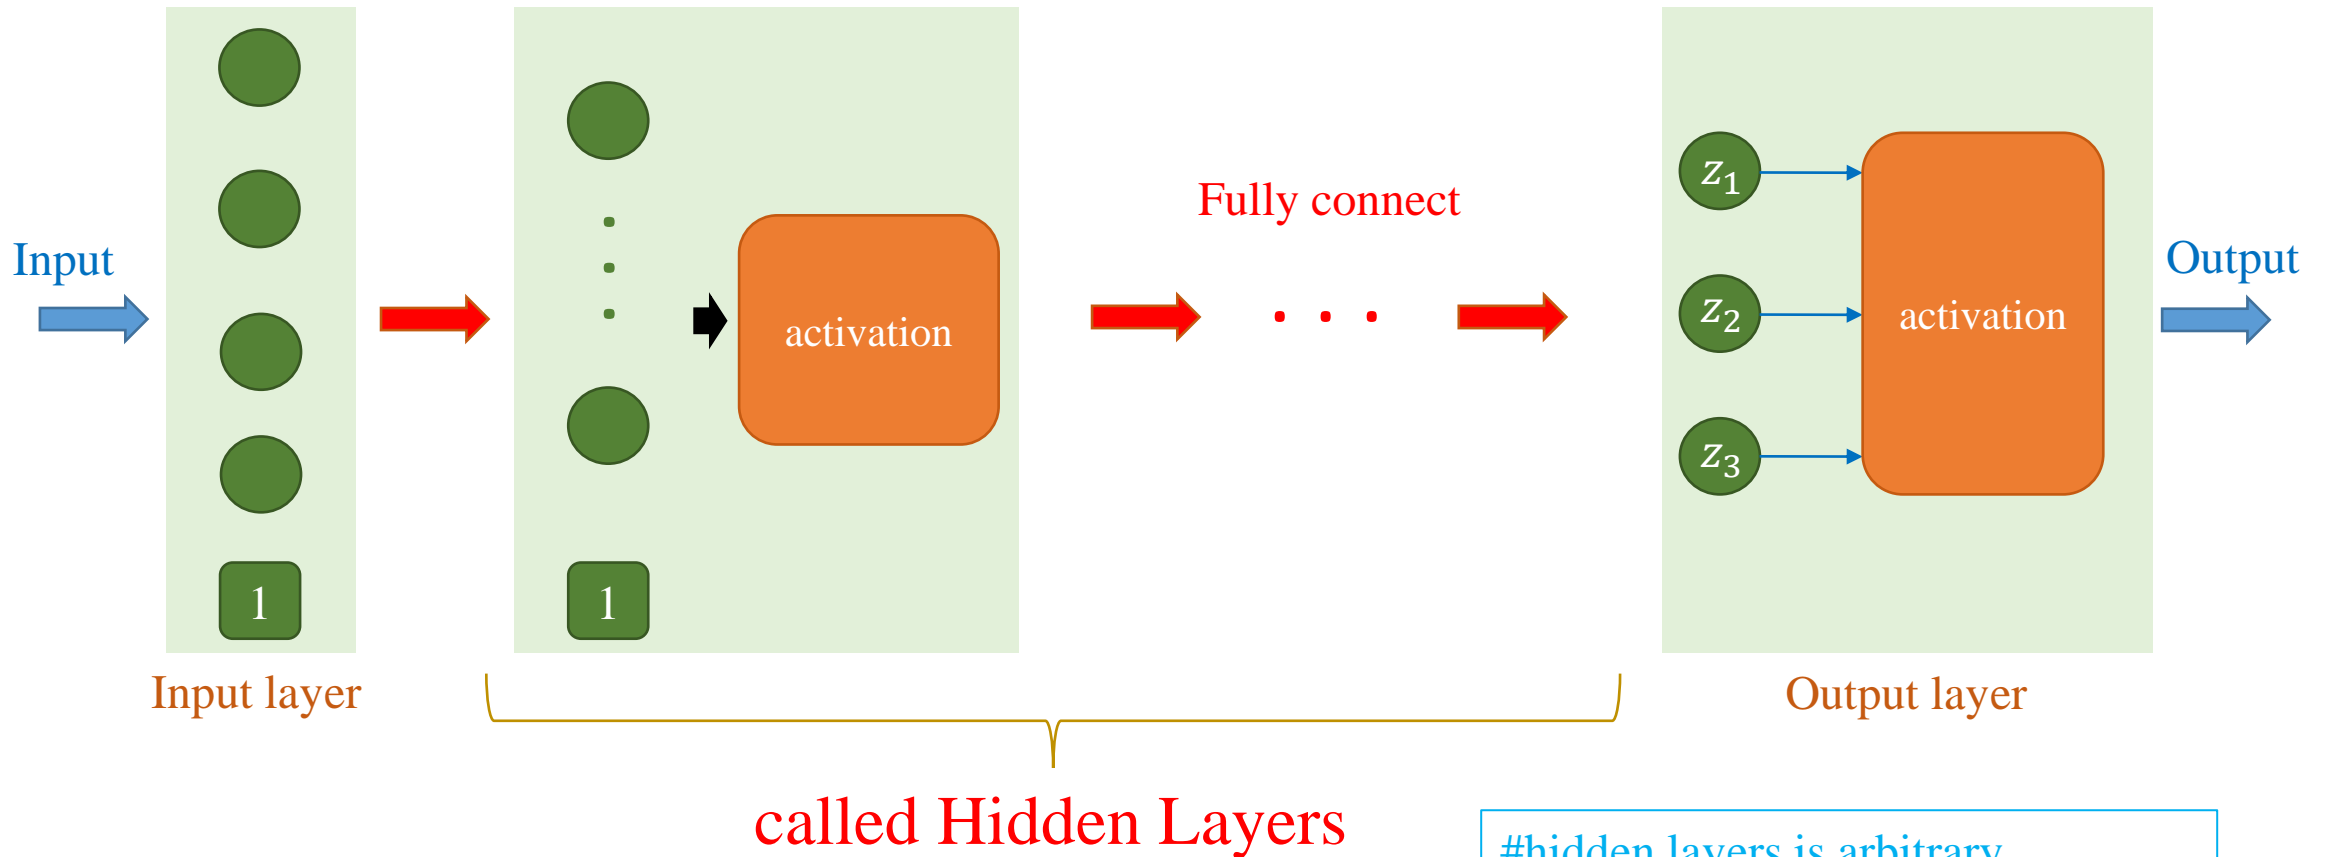
- ❖ John Von Neumann's quote “with four parameters I can fit an elephant, with five I can make him wiggle his trunk”
- ❖ More parameters \rightarrow better capacity (~stronger model)



Multi-layer Perceptron

❖ An idea: More parameters \rightarrow better capacity (~stronger model)

❖ Adding more layers



#hidden layers is arbitrary
#nodes in a hidden layer is arbitrary

Multi-layer Perceptron

❖ ReLU function


$$\text{ReLU}(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

data =



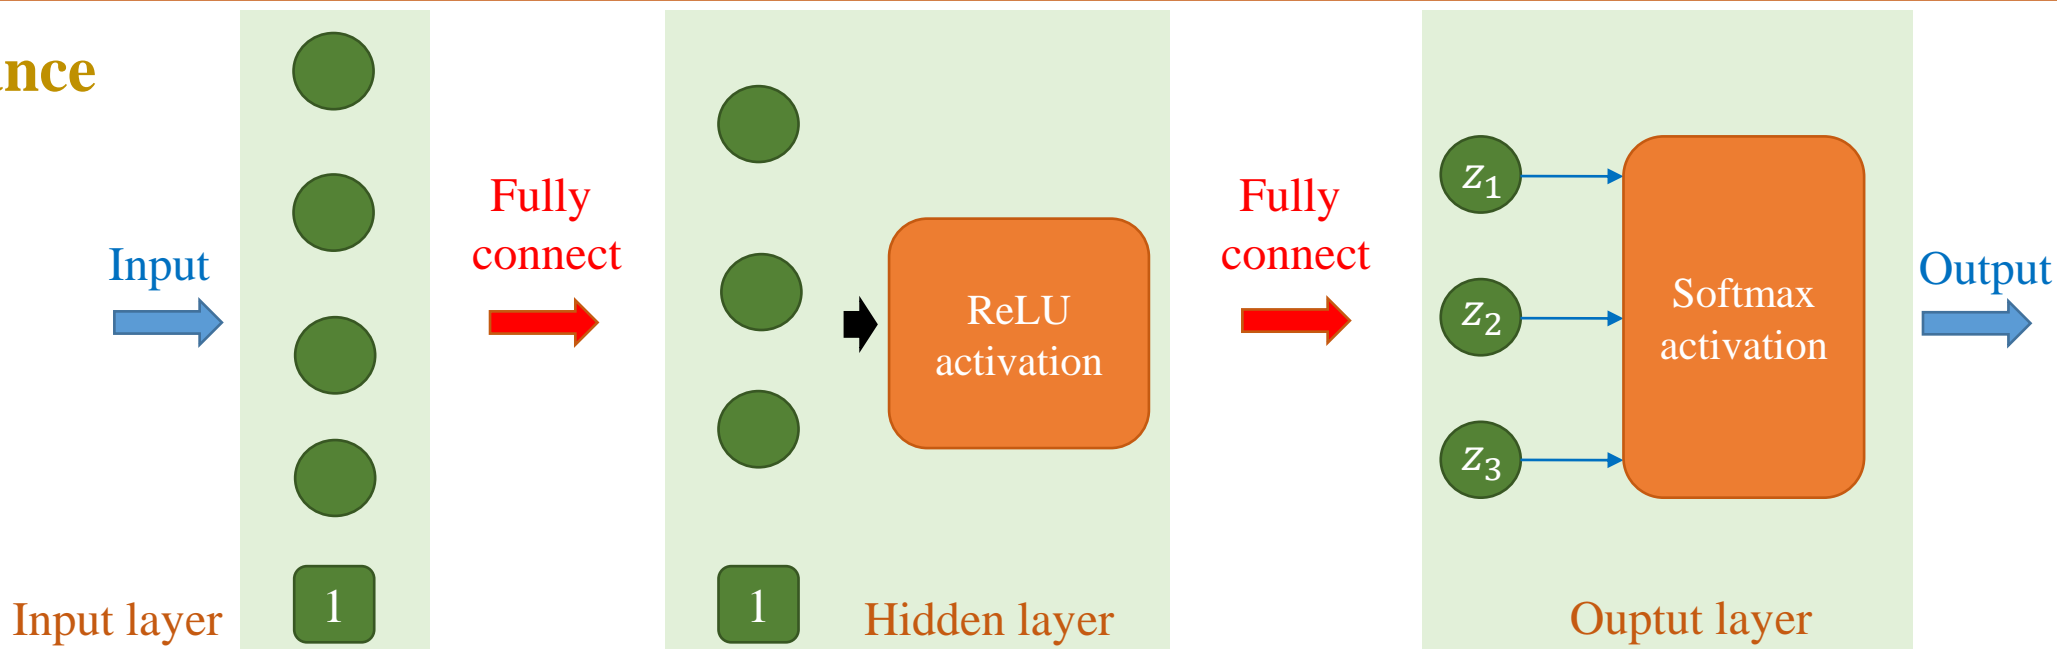
data_a = ReLU(data)

data_a =



Multi-layer Perceptron

An instance

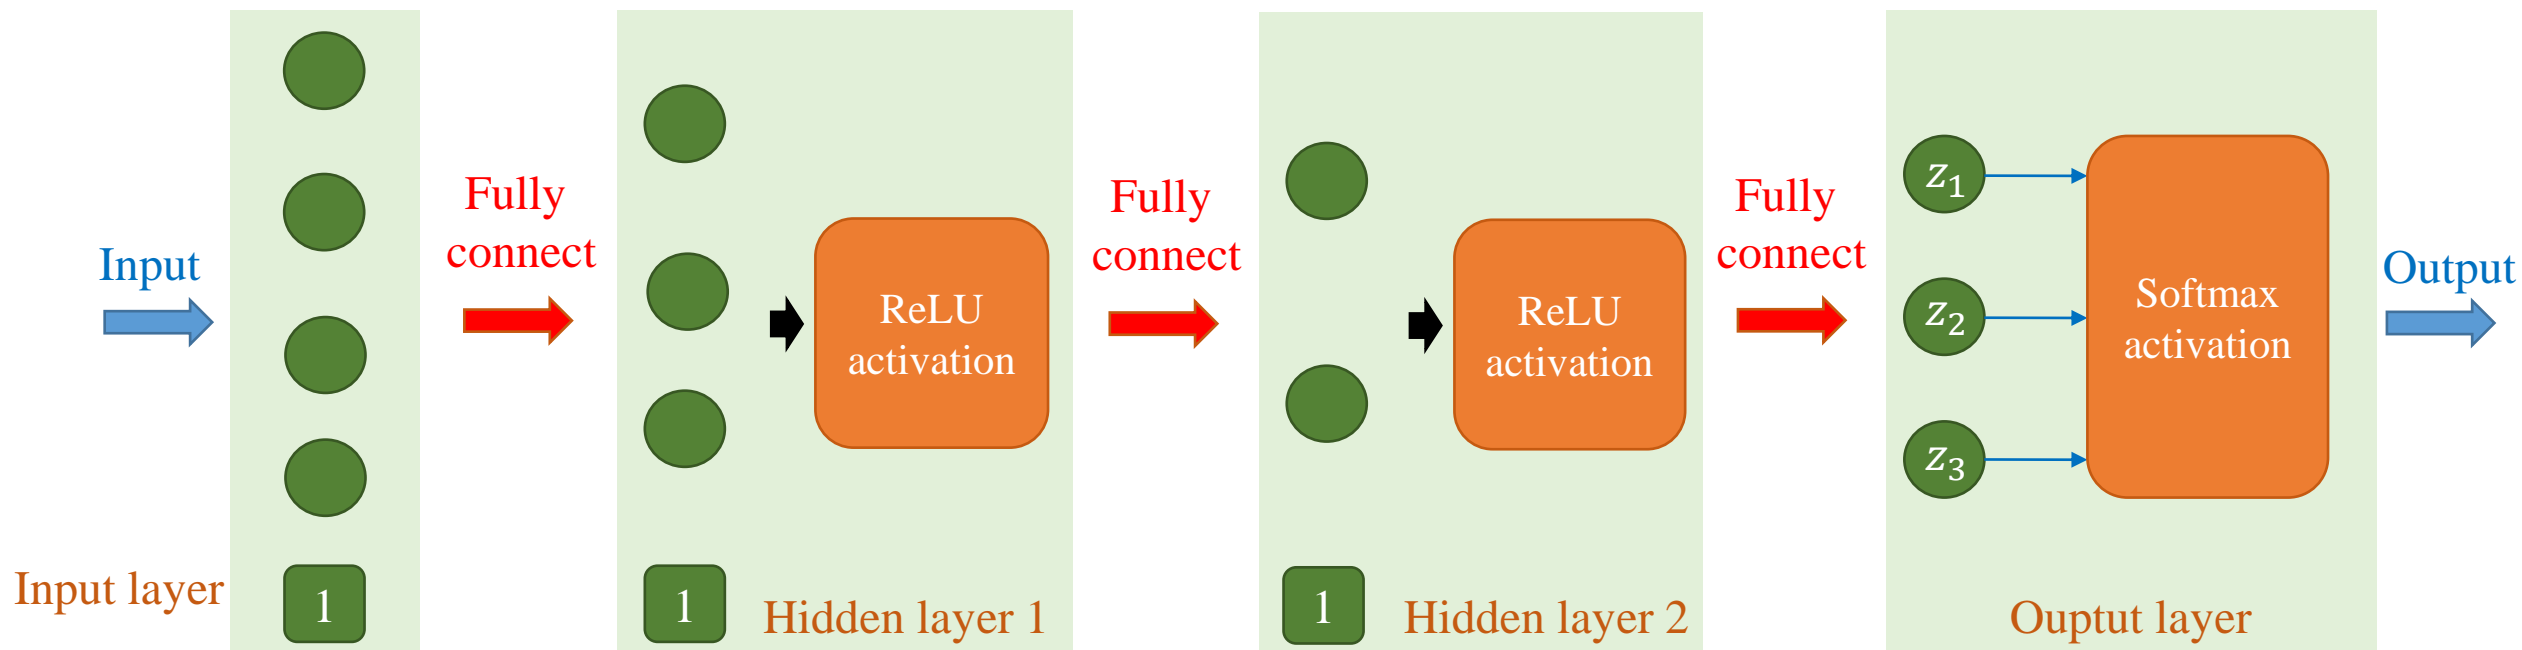


```
import torch.nn as nn

model = nn.Sequential(
    nn.Linear(4, 3),
    nn.ReLU(),
    nn.Linear(3, 3)
)
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 3]	15
ReLU-2	[-1, 3]	0
Linear-3	[-1, 3]	12
Total params: 27		
Trainable params: 27		
Non-trainable params: 0		

Multi-layer Perceptron



```
import torch.nn as nn

model = nn.Sequential(
    nn.Linear(4, 3),
    nn.ReLU(),
    nn.Linear(3, 2),
    nn.ReLU(),
    nn.Linear(2, 3)
)
```



Layer (type)	Output Shape	Param #
Linear-1	[-1, 3]	15
ReLU-2	[-1, 3]	0
Linear-3	[-1, 2]	8
ReLU-4	[-1, 2]	0
Linear-5	[-1, 3]	9

Total params: 32

Back to Fashion-MNIST

$$\text{Image} = \frac{\text{Image}}{255.0}$$

```
model = nn.Sequential(  
    nn.Linear(784, 256),  
    nn.ReLU(),  
    nn.Linear(256, 10)  
)  
print(model)  
  
Sequential(  
  (0): Linear(in_features=784, out_features=256, bias=True)  
  (1): ReLU()  
  (2): Linear(in_features=256, out_features=10, bias=True)  
)
```

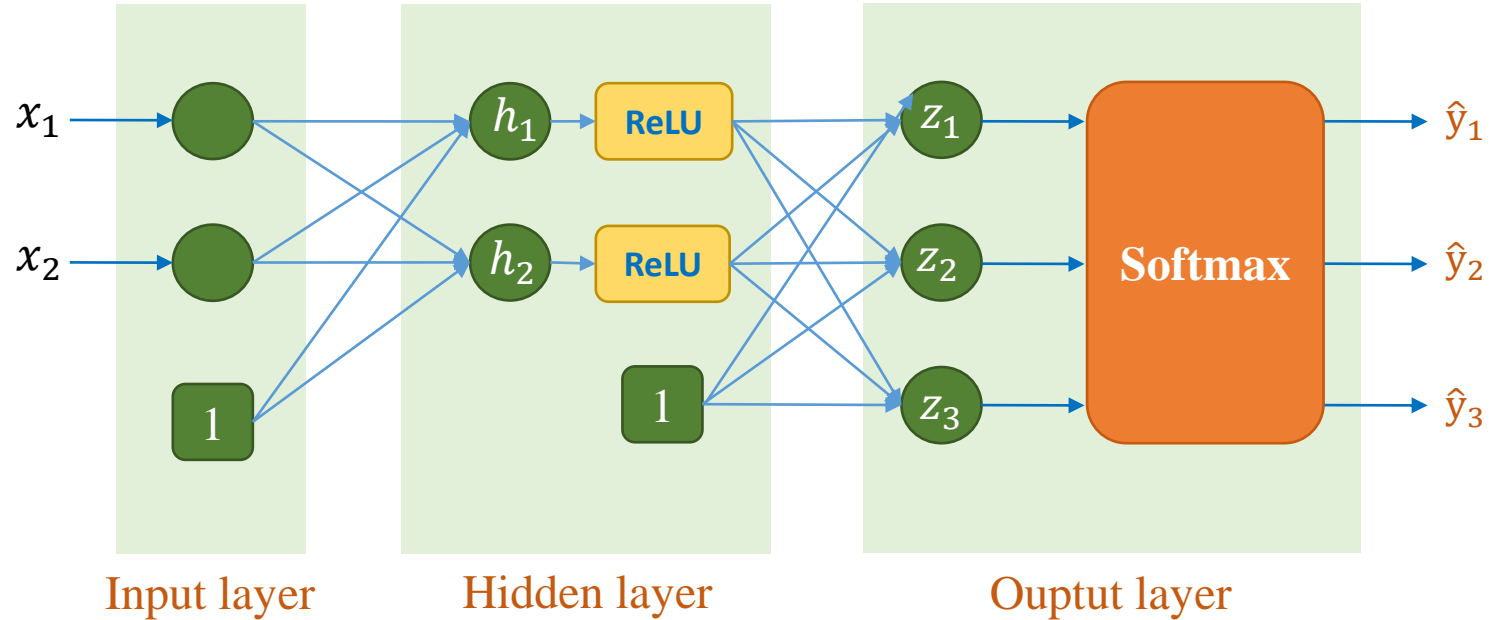


MLP Example

Feature		Label
Petal Length	Petal Width	Label
1.5	0.2	0
1.4	0.2	0
1.6	0.2	0
4.7	1.6	1
3.3	1.1	1
4.6	1.3	1
5.6	2.2	2
5.1	1.5	2
5.6	1.4	2

$$\mathbf{x} = \begin{bmatrix} \mathbf{x}^{(1)} \\ \mathbf{x}^{(2)} \\ \mathbf{x}^{(3)} \end{bmatrix} = \begin{bmatrix} 1.5 & 0.2 \\ 4.7 & 1.6 \\ 5.6 & 2.2 \end{bmatrix}$$

$$\mathbf{y} = \begin{bmatrix} 0 \\ 1 \\ 2 \end{bmatrix}$$



$$\mathbf{W}_h = [\mathbf{W}_{h1} \quad \mathbf{W}_{h2}]$$

$$= \begin{bmatrix} 0.0 & 0.0 \\ 0.86 & -1.04 \\ 0.41 & -0.65 \end{bmatrix}$$


$$\mathbf{W}_z = [\mathbf{W}_{z1} \quad \mathbf{W}_{z2} \quad \mathbf{W}_{z3}]$$

$$= \begin{bmatrix} 0.0 & 0.0 & 0.0 \\ 0.32 & 0.25 & 0.14 \\ -0.47 & -1.06 & 0.063 \end{bmatrix}$$

$$\mathbf{h} = \mathbf{x}\mathbf{W}_h = \begin{bmatrix} 1 & 1.5 & 0.2 \\ 1 & 4.7 & 1.6 \\ 1 & 5.6 & 2.2 \end{bmatrix} \begin{bmatrix} 0.0 & 0.0 \\ 0.86 & -1.04 \\ 0.41 & -0.65 \end{bmatrix} = \begin{bmatrix} 1.373 & -1.696 \\ 4.708 & -5.951 \\ 5.731 & -7.281 \end{bmatrix}$$

$$\text{ReLU}(\mathbf{h}) = \begin{bmatrix} 1.373 & 0 \\ 4.708 & 0 \\ 5.731 & 0 \end{bmatrix}$$

Feature		Label
Petal Length	Petal Width	Label
1.5	0.2	0
1.4	0.2	0
1.6	0.2	0
4.7	1.6	1
3.3	1.1	1
4.6	1.3	1
5.6	2.2	2
5.1	1.5	2
5.6	1.4	2



$$\mathbf{x} = \begin{bmatrix} \mathbf{x}^{(1)} \\ \mathbf{x}^{(2)} \\ \mathbf{x}^{(3)} \end{bmatrix} = \begin{bmatrix} 1 & 1.5 & 0.2 \\ 1 & 4.7 & 1.6 \\ 1 & 5.6 & 2.2 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 0 \\ 1 \\ 2 \end{bmatrix}$$

$$\begin{aligned} \mathbf{W}_h &= [\mathbf{W}_{h1} \quad \mathbf{W}_{h2}] & \mathbf{W}_z &= [\mathbf{W}_{z1} \quad \mathbf{W}_{z2} \quad \mathbf{W}_{z3}] \\ &= \begin{bmatrix} 0.0 & 0.0 \\ 0.86 & -1.04 \\ 0.41 & -0.65 \end{bmatrix} & &= \begin{bmatrix} 0.0 & 0.0 & 0.0 \\ 0.32 & 0.25 & 0.14 \\ -0.47 & -1.06 & 0.063 \end{bmatrix} \end{aligned}$$

$$\text{ReLU}(\mathbf{h}) = \begin{bmatrix} 1.373 & 0 \\ 4.708 & 0 \\ 5.731 & 0 \end{bmatrix}$$

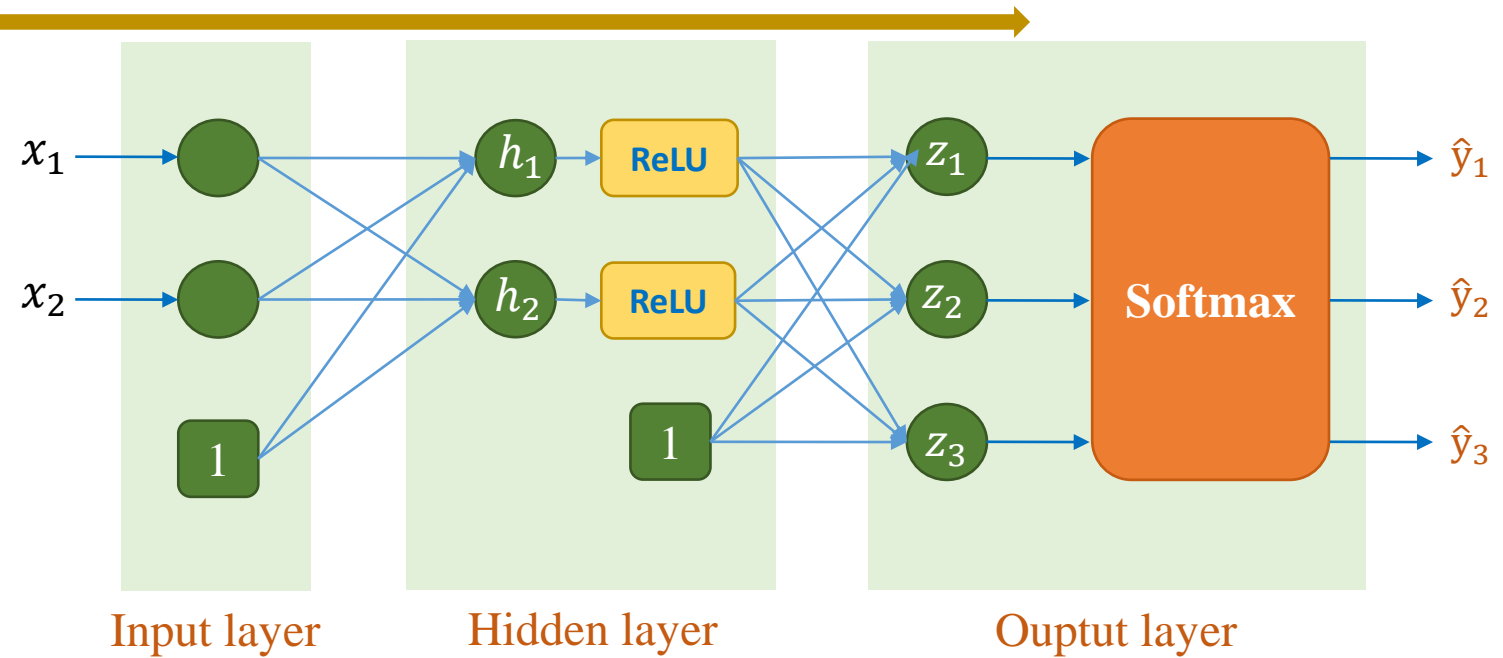
$$[\mathbf{1} \quad \text{ReLU}(\mathbf{h})] = \begin{bmatrix} 1 & 1.373 & 0 \\ 1 & 4.708 & 0 \\ 1 & 5.731 & 0 \end{bmatrix}$$

Feature			Label
Petal Length	Petal Width		Label
1.5	0.2		0
1.4	0.2		0
1.6	0.2		0
4.7	1.6		1
3.3	1.1		1
4.6	1.3		1
5.6	2.2		2
5.1	1.5		2
5.6	1.4		2

$$\mathbf{x} = \begin{bmatrix} \mathbf{x}^{(1)} \\ \mathbf{x}^{(2)} \\ \mathbf{x}^{(3)} \end{bmatrix} = \begin{bmatrix} 1 & 1.5 & 0.2 \\ 1 & 4.7 & 1.6 \\ 1 & 5.6 & 2.2 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 0 \\ 1 \\ 2 \end{bmatrix}$$

$$\mathbf{z} = [\mathbf{1} \quad \text{ReLU}(\mathbf{h})] \mathbf{W}_z = \begin{bmatrix} 1 & 1.373 & 0 \\ 1 & 4.708 & 0 \\ 1 & 5.731 & 0 \end{bmatrix} \begin{bmatrix} 0.0 & 0.0 & 0.0 \\ 0.32 & 0.25 & 0.14 \\ -0.47 & -1.06 & 0.063 \end{bmatrix}$$

$$= \begin{bmatrix} 0.439 & 0.356 & 0.195 \\ 1.507 & 1.220 & 0.670 \\ 1.835 & 1.485 & 0.816 \end{bmatrix}$$



$$\mathbf{W}_h = [\mathbf{W}_{h1} \quad \mathbf{W}_{h2}] \quad \mathbf{W}_z = [\mathbf{W}_{z1} \quad \mathbf{W}_{z2} \quad \mathbf{W}_{z3}]$$

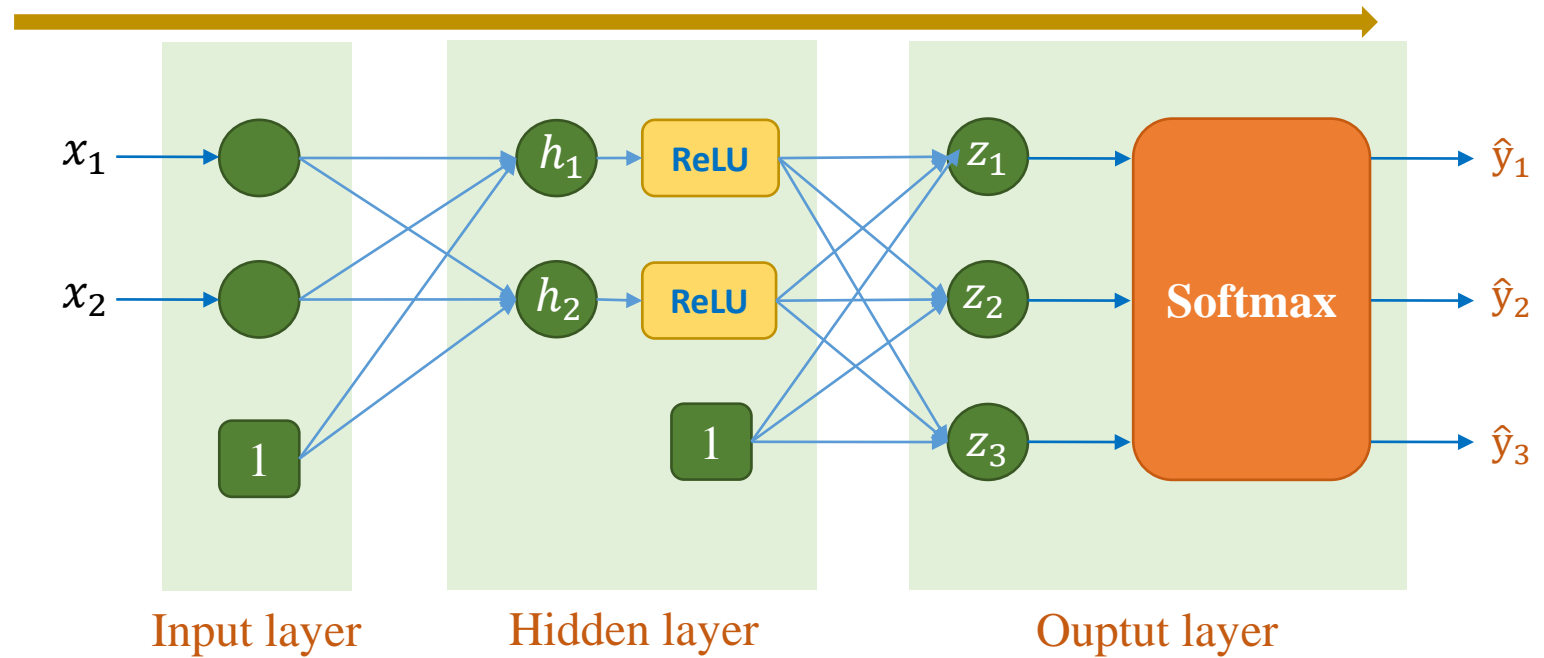
$$= \begin{bmatrix} 0.0 & 0.0 \\ 0.86 & -1.04 \\ 0.41 & -0.65 \end{bmatrix} \quad = \begin{bmatrix} 0.0 & 0.0 & 0.0 \\ 0.32 & 0.25 & 0.14 \\ -0.47 & -1.06 & 0.063 \end{bmatrix}$$

$$\mathbf{z} = \begin{bmatrix} 0.439 & 0.356 & 0.195 \\ 1.507 & 1.220 & 0.670 \\ 1.835 & 1.485 & 0.816 \end{bmatrix}$$

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{z}) = \begin{bmatrix} \hat{y}^{(1)} \\ \hat{y}^{(2)} \\ \hat{y}^{(3)} \end{bmatrix} = \begin{bmatrix} 0.369 & 0.340 & 0.289 \\ 0.458 & 0.343 & 0.198 \\ 0.484 & 0.341 & 0.174 \end{bmatrix}$$

$$\text{loss} = 1.269$$

Feature			Label
Petal Length	Petal Width		Label
1.5	0.2		0
1.4	0.2		0
1.6	0.2		0
4.7	1.6		1
3.3	1.1		1
4.6	1.3		1
5.6	2.2		2
5.1	1.5		2
5.6	1.4		2



$$\mathbf{x} = \begin{bmatrix} x^{(1)} \\ x^{(2)} \\ x^{(3)} \end{bmatrix} = \begin{bmatrix} 1 & 1.5 & 0.2 \\ 1 & 4.7 & 1.6 \\ 1 & 5.6 & 2.2 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 0 \\ 1 \\ 2 \end{bmatrix}$$

$$\mathbf{W}_h = [\mathbf{W}_{h1} \quad \mathbf{W}_{h2}]$$


$$= \begin{bmatrix} 0.0 & 0.0 \\ 0.86 & -1.04 \\ 0.41 & -0.65 \end{bmatrix}$$

$$\mathbf{W}_z = [\mathbf{W}_{z1} \quad \mathbf{W}_{z2} \quad \mathbf{W}_{z3}]$$


$$= \begin{bmatrix} 0.0 & 0.0 & 0.0 \\ 0.32 & 0.25 & 0.14 \\ -0.47 & -1.06 & 0.063 \end{bmatrix}$$

Softmax and MLP

```
model = nn.Sequential(  
    nn.Flatten(),  
    nn.Linear(784, 10)  
)  
model = model.to(device)
```



```
model = nn.Sequential(  
    nn.Flatten(), nn.Linear(784, 256),  
    nn.ReLU(), nn.Linear(256, 10)  
)  
model = model.to(device)
```




Outline

- **Image Data Loading Using Numpy&PyTorch**
- **Softmax+Normalization for Fashion-MNIST**
- **MLP and Examples**
- **Step-by-Step Implementation**
- **Training Strategy (optional)**

Step-by-Step Implementation

❖ 1. Data Preparation

Each sample is a tuple (image tensor, label)



```
transform = T.Compose([T.ToTensor(),
                       T.Normalize((0.5,),
                                   (0.5,))])

trainset = FashionMNIST(root='data',
                        train=True,
                        download=True,
                        transform=transform)

trainloader = DataLoader(trainset,
                        batch_size=64,
                        shuffle=True)

testset = FashionMNIST(root='data',
                      train=False,
                      download=True,
                      transform=transform)


testloader = DataLoader(testset,
                      batch_size=64,
                      shuffle=False)
```

Step-by-Step Implementation

❖ 2. Model, loss and optimizer

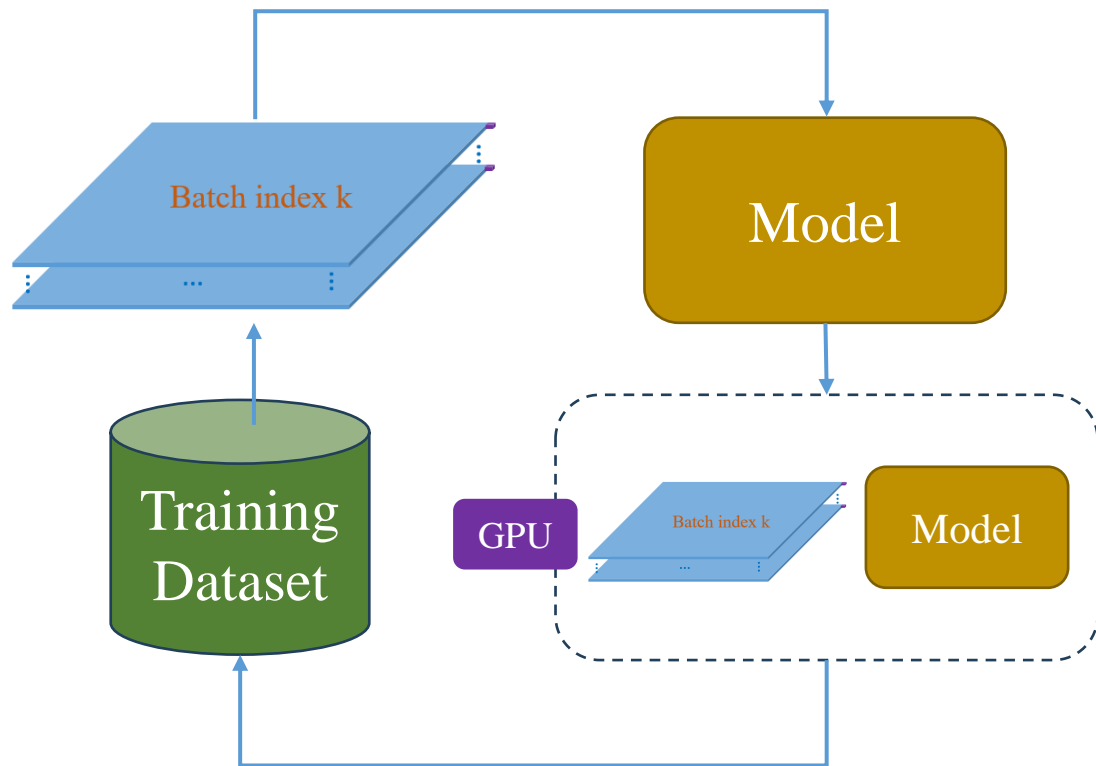
```
model = model.to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
```

```
# Define the MLP model
model = nn.Sequential(
    nn.Flatten(),
    nn.Linear(784, 256),
    nn.ReLU(),
    nn.Linear(256, 10)
)
```



Step-by-Step Implementation

❖ 3. Training



```
# Training the model
```

```
max_epoch = 5
```

```
for epoch in range(max_epoch):
```

```
    for i, (inputs, labels) in enumerate(trainloader, 0):
```

```
        # Move inputs and labels to the device
```

```
        inputs, labels = inputs.to(device), labels.to(device)
```

```
        # Zero the parameter gradients
```

```
        optimizer.zero_grad()
```

```
        # Forward pass
```

```
        outputs = model(inputs)
```

```
        loss = criterion(outputs, labels)
```

```
        # Backward pass and optimization
```

```
        loss.backward()
```

```
        optimizer.step()
```

```
    print(f"Epoch [{epoch + 1}/{max_epoch}]")
```

```
Epoch [1/5]
```

```
Epoch [2/5]
```


```
Epoch [3/5]
```

```
Epoch [4/5]
```

```
Epoch [5/5]
```


Step-by-Step Implementation

❖ 4. Inference



```
correct = 0
total = 0
with torch.no_grad():
    for images, labels in testloader:
        # Move inputs and labels to the device
        images = images.to(device)
        labels = labels.to(device)

        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)

        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = 100 * correct / total
print(f"accuracy: {accuracy}")
```

Step-by-Step Implementation

❖ Addition 1: Compute Training Loss and Accuracy

Epoch [1/5],	Loss: 0.7866,	Accuracy: 74.73%
Epoch [2/5],	Loss: 0.5205,	Accuracy: 81.59%
Epoch [3/5],	Loss: 0.4706,	Accuracy: 83.28%
Epoch [4/5],	Loss: 0.4432,	Accuracy: 84.25%
Epoch [5/5],	Loss: 0.4232,	Accuracy: 85.13%

```
1 for epoch in range(5):
2     running_loss = 0.0
3     correct = 0    # to track number of correct predictions
4     total = 0      # to track total number of samples
5
6     for i, (inputs, labels) in enumerate(trainloader, 0):
7         # see comments from the previous example
8         inputs, labels = inputs.to(device), labels.to(device)
9         optimizer.zero_grad()
10        outputs = model(inputs)
11        loss = criterion(outputs, labels)
12        loss.backward()
13        optimizer.step()
14
15        # Determine class predictions and track accuracy
16        _, predicted = torch.max(outputs.data, 1)
17        total += labels.size(0)
18        correct += (predicted == labels).sum().item()
19
20        # accumulate loss
21        running_loss += loss.item()
22
23    epoch_accuracy = 100 * correct / total
24    running_loss = running_loss / (i + 1)
```

Step-by-Step Implementation

❖ Addition 2: Compute Test Loss and Accuracy

```
def evaluate(model, testloader, criterion):
    model.eval()
    test_loss = 0.0
    correct = 0
    total = 0
    with torch.no_grad():
        for images, labels in testloader:
            # Move inputs and labels to the device
            images = images.to(device)
            labels = labels.to(device)

            outputs = model(images)
            loss = criterion(outputs, labels)
            test_loss += loss.item()
            _, predicted = torch.max(outputs.data, 1)

            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    accuracy = 100 * correct / total
    return test_loss / len(testloader), accuracy
```

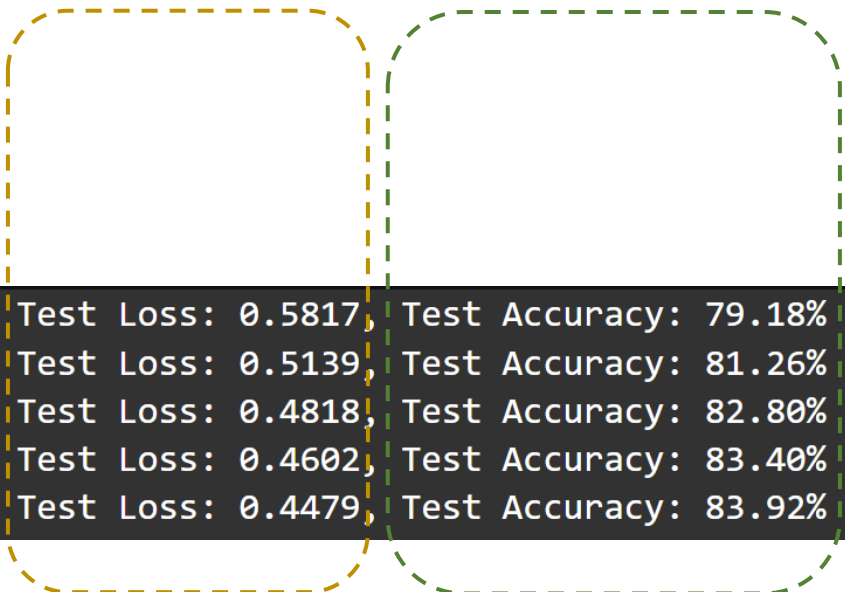
```
for epoch in range(5):
    running_loss = 0.0
    correct = 0    # to track number of correct predictions
    total = 0      # to track total number of samples

    for i, (inputs, labels) in enumerate(trainloader, 0):
        # see comments from the previous example
        inputs, labels = inputs.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        # Determine class predictions and track accuracy
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
        # accumulate loss
        running_loss += loss.item()

    epoch_accuracy = 100 * correct / total
    running_loss = running_loss / (i + 1)
    test_loss, test_accuracy = evaluate(model,
                                         testloader,
                                         criterion)
```

Step-by-Step Implementation

❖ Addition 2: Compute Test Loss and Accuracy



Epoch [1/5],	Test Loss: 0.5817,	Test Accuracy: 79.18%
Epoch [2/5],	Test Loss: 0.5139,	Test Accuracy: 81.26%
Epoch [3/5],	Test Loss: 0.4818,	Test Accuracy: 82.80%
Epoch [4/5],	Test Loss: 0.4602,	Test Accuracy: 83.40%
Epoch [5/5],	Test Loss: 0.4479,	Test Accuracy: 83.92%

```
for epoch in range(5):
    running_loss = 0.0
    correct = 0    # to track number of correct predictions
    total = 0      # to track total number of samples

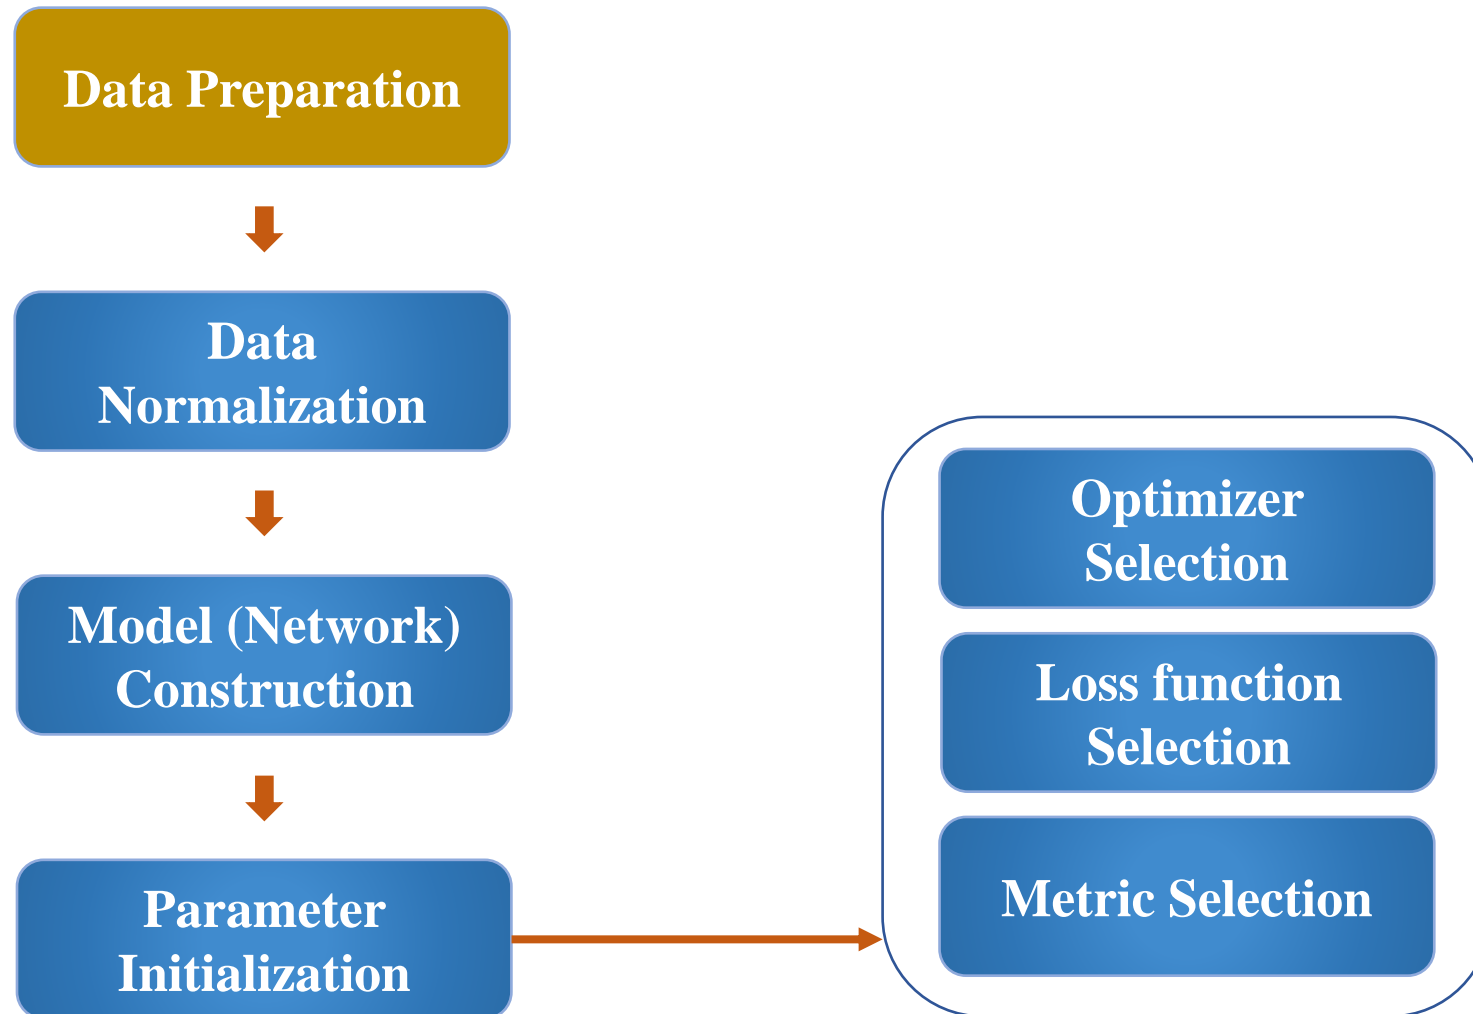
    for i, (inputs, labels) in enumerate(trainloader, 0):
        # see comments from the previous example
        inputs, labels = inputs.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        # Determine class predictions and track accuracy
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
        # accumulate loss
        running_loss += loss.item()

    epoch_accuracy = 100 * correct / total
    running_loss = running_loss / (i + 1)
    test_loss, test_accuracy = evaluate(model,
                                       testloader,
                                       criterion)
```

Outline

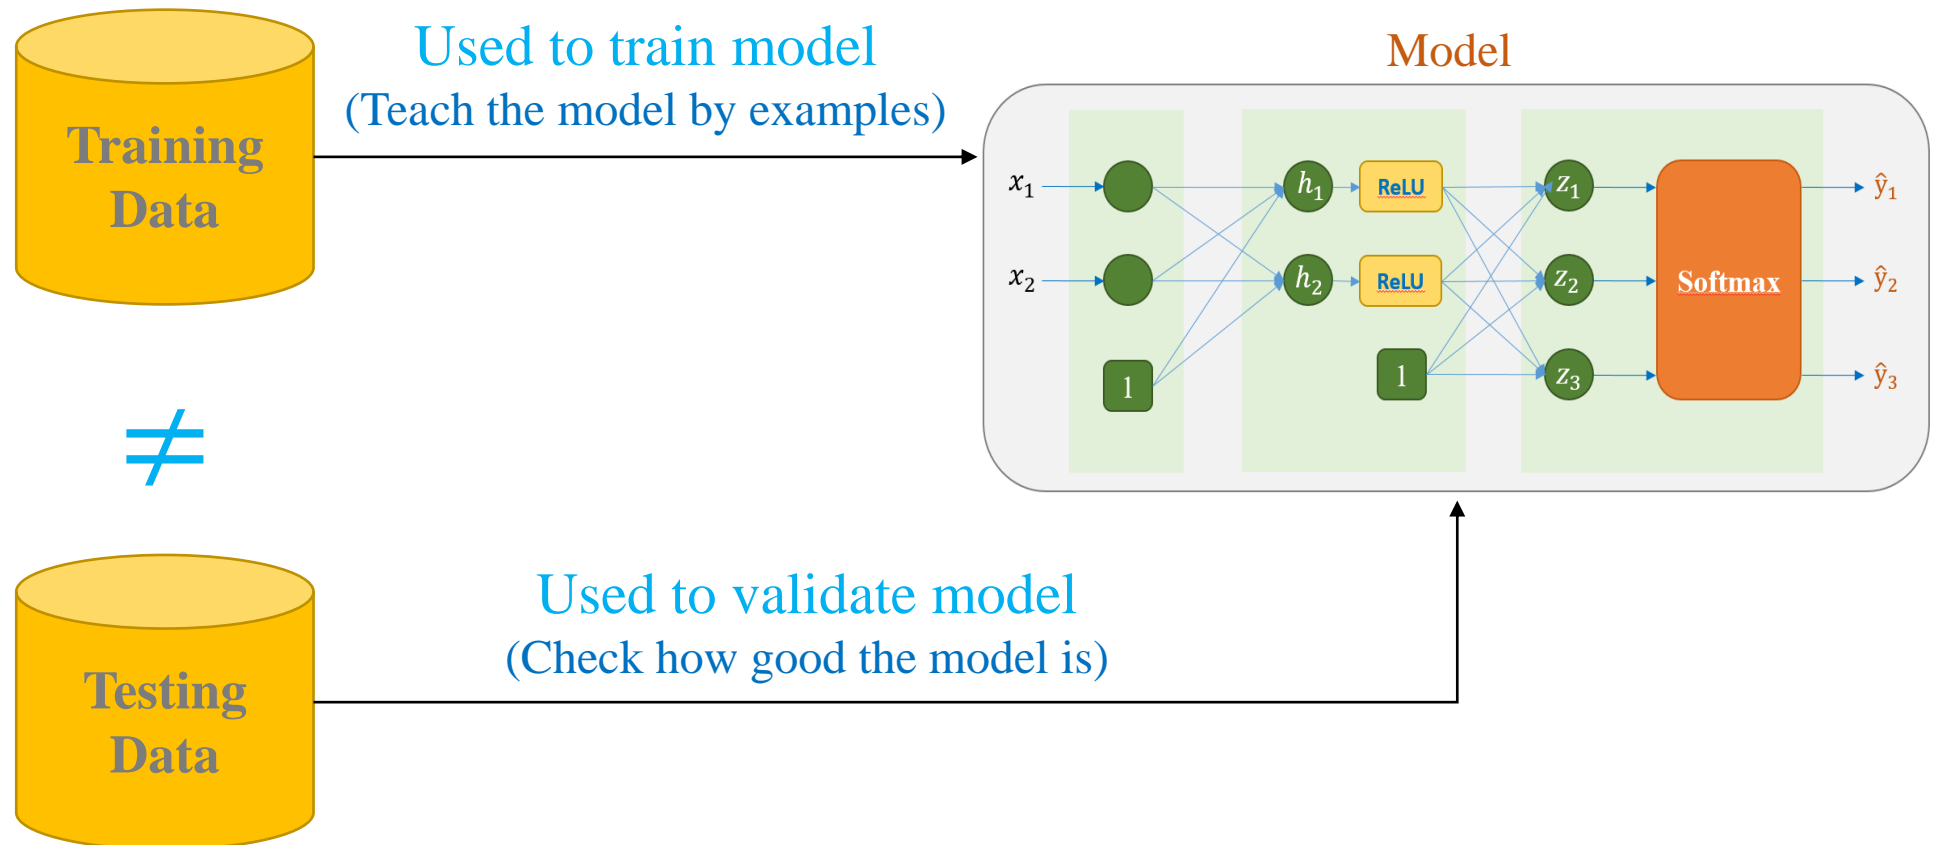
- **Image Data Loading Using Numpy&PyTorch**
- **Softmax+Normalization for Fashion-MNIST**
- **MLP and Examples**
- **Step-by-Step Implementation**
- **Training Strategy (optional)**

To-do List for Training

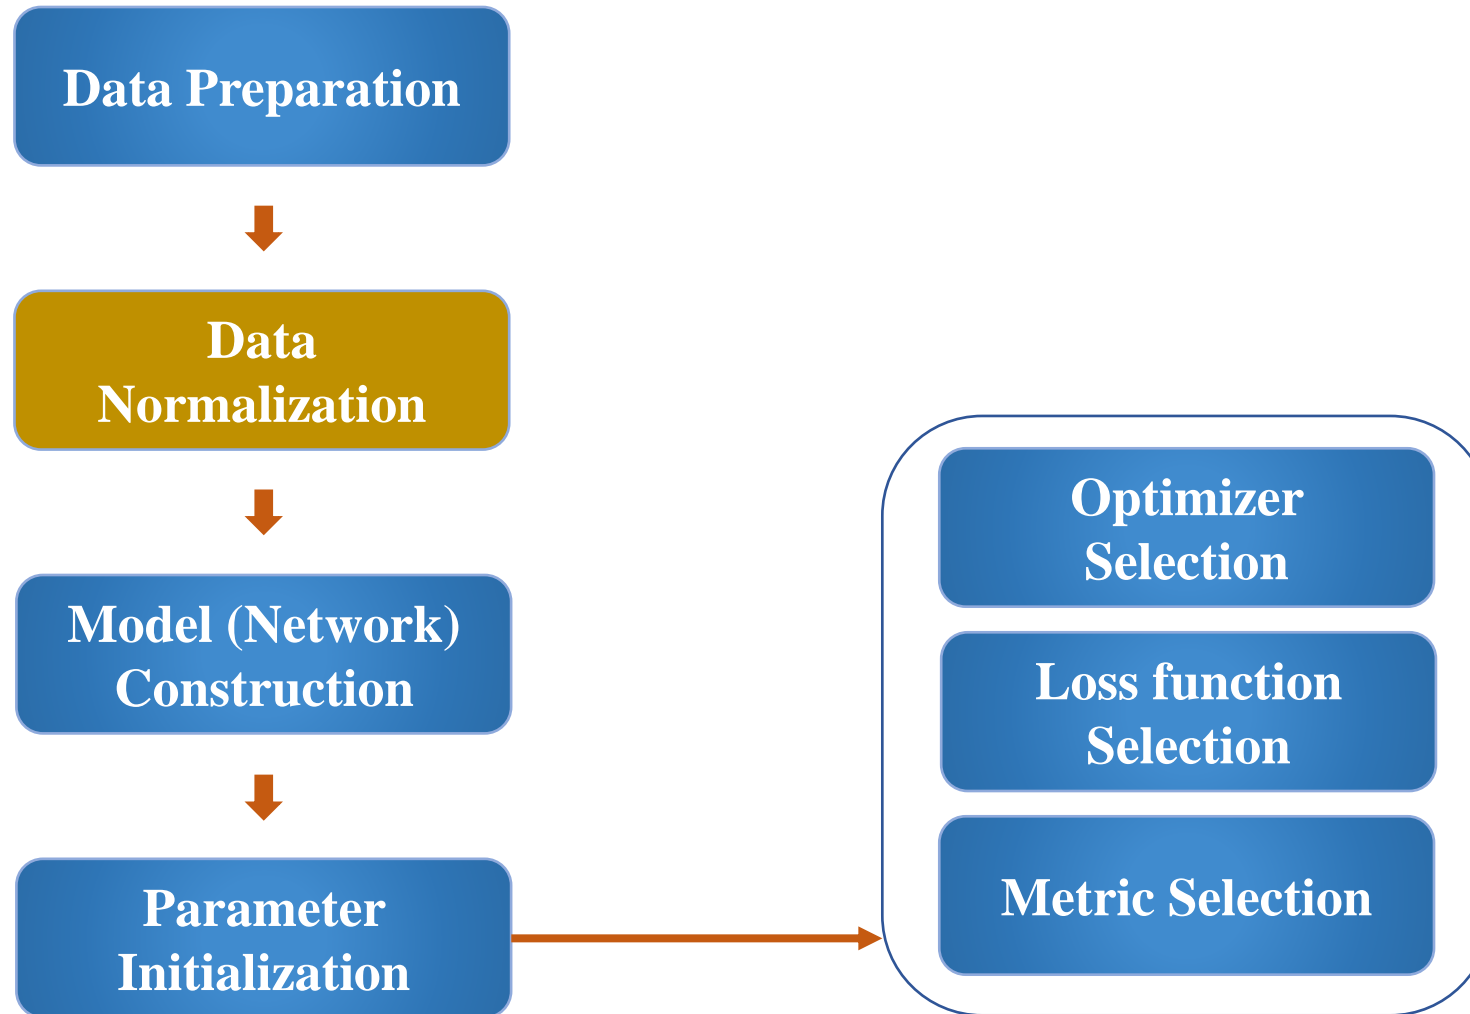


To-do List for Training

Data Preparation



To-do List for Training



To-do List for Training

Data Normalization



Convert to the range [0,1]

$$\text{Image} = \frac{\text{Image}}{255}$$

Convert to the range [-1,1]

$$\text{Image} = \frac{\text{Image}}{127.5} - 1$$

Z-score normalization

$$\text{Image} = \frac{\text{Image} - \mu}{\sigma}$$

μ is the mean of
the image (or training data)

σ is the standard deviation
of the image (or training data)

Implmentation

In Theory

$$X \in [0, 255]$$

Convert to the range [0,1]

$$\text{Image} = \frac{\text{Image}}{255}$$

Convert to the range [-1,1]

$$\text{Image} = \frac{\text{Image}}{127.5} - 1$$

Z-score normalization

$$\text{Image} = \frac{\text{Image} - \mu}{\sigma}$$

In Pytorch

$$X \in [0, 1]$$

Normalize(*mean*, *std*)

$$\text{Image} = \frac{\text{Image} - \textit{mean}}{\textit{std}}$$

[0,1]

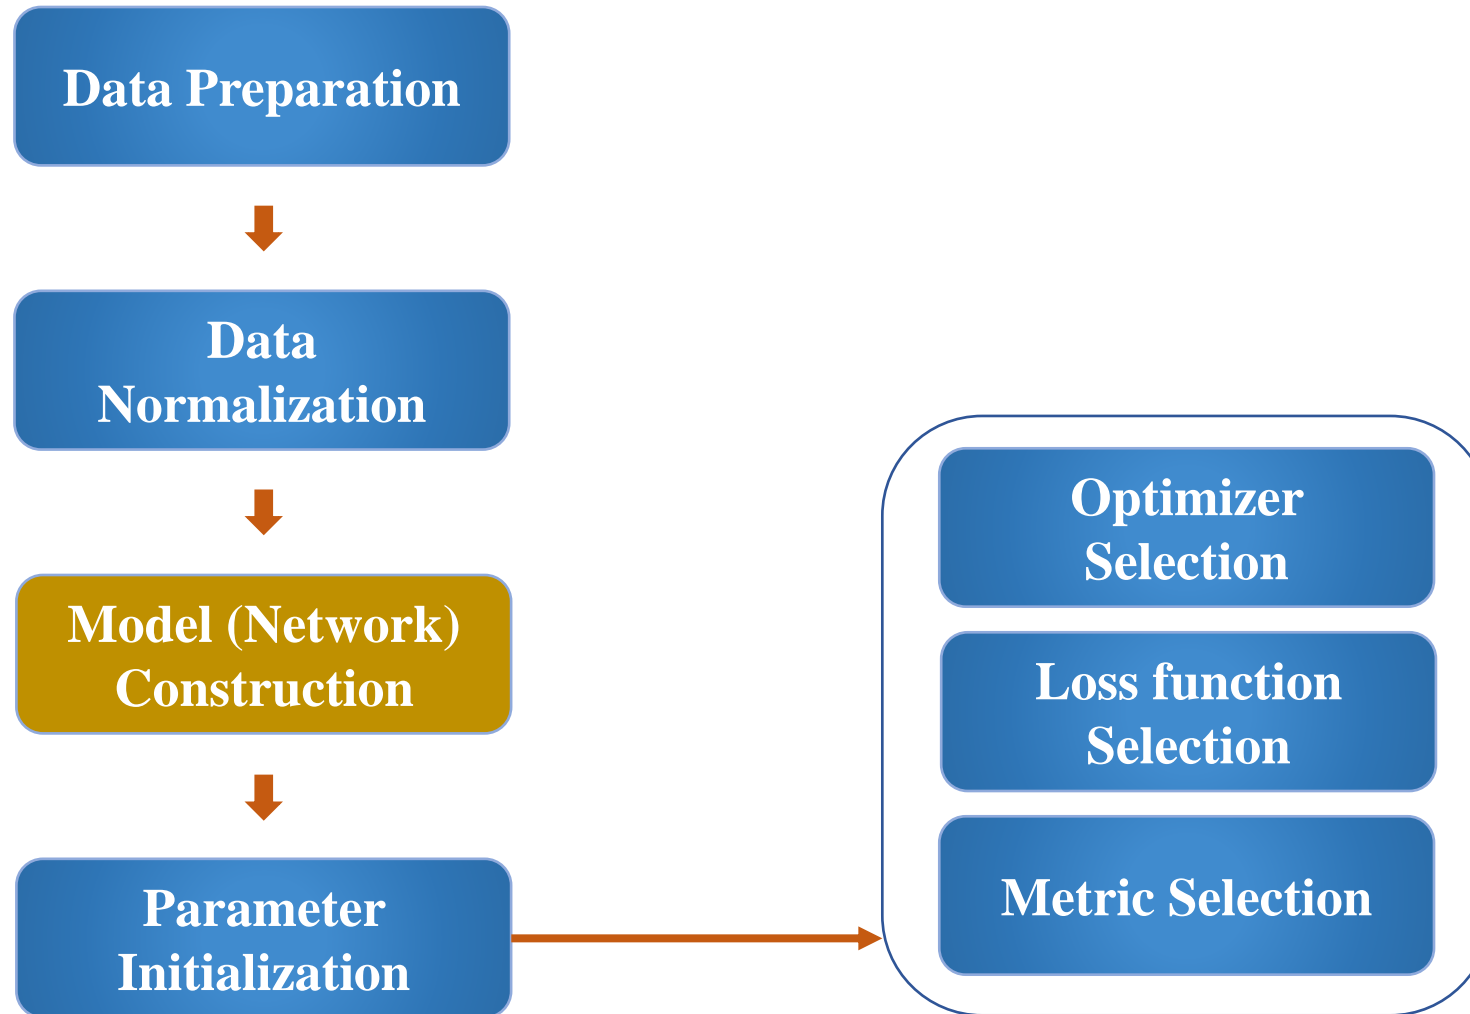
mean = 0 ; std = 1

[-1,1]

mean = 0.5; std = 0.5

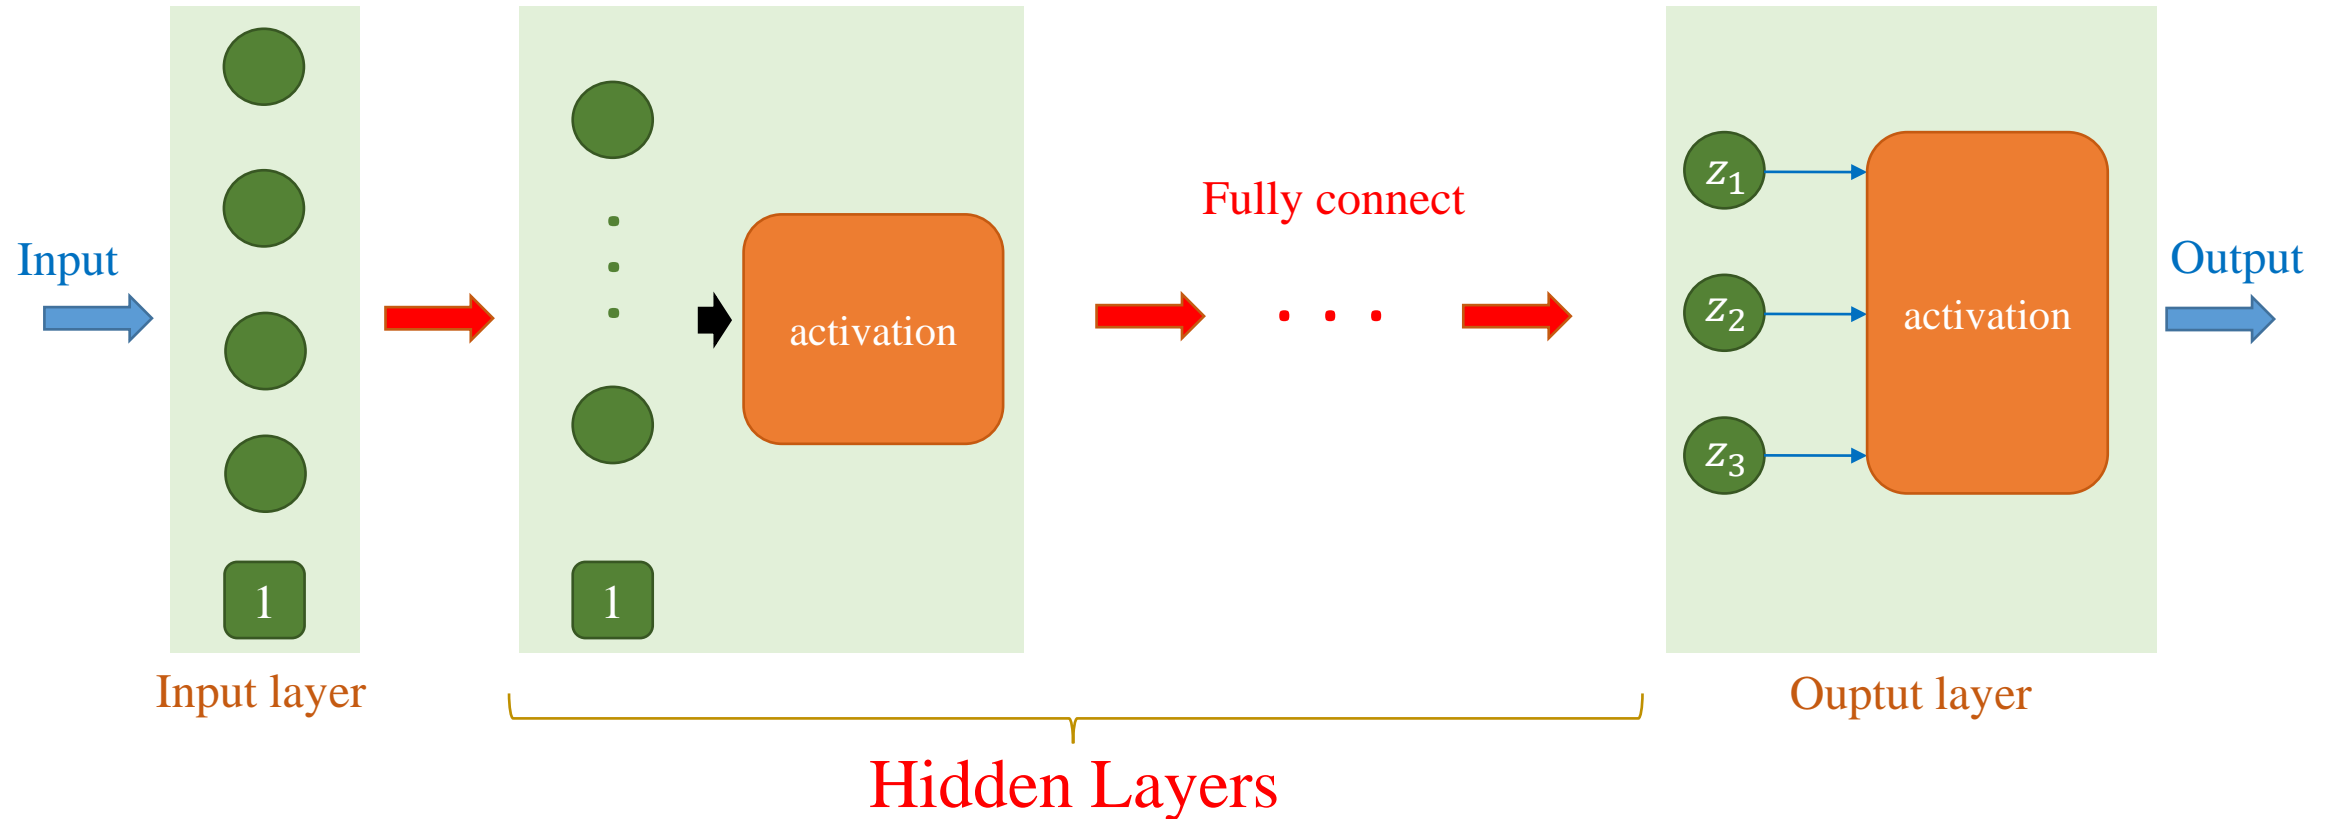
Compute mean and std
from data

To-do List for Training



To-do List for Training

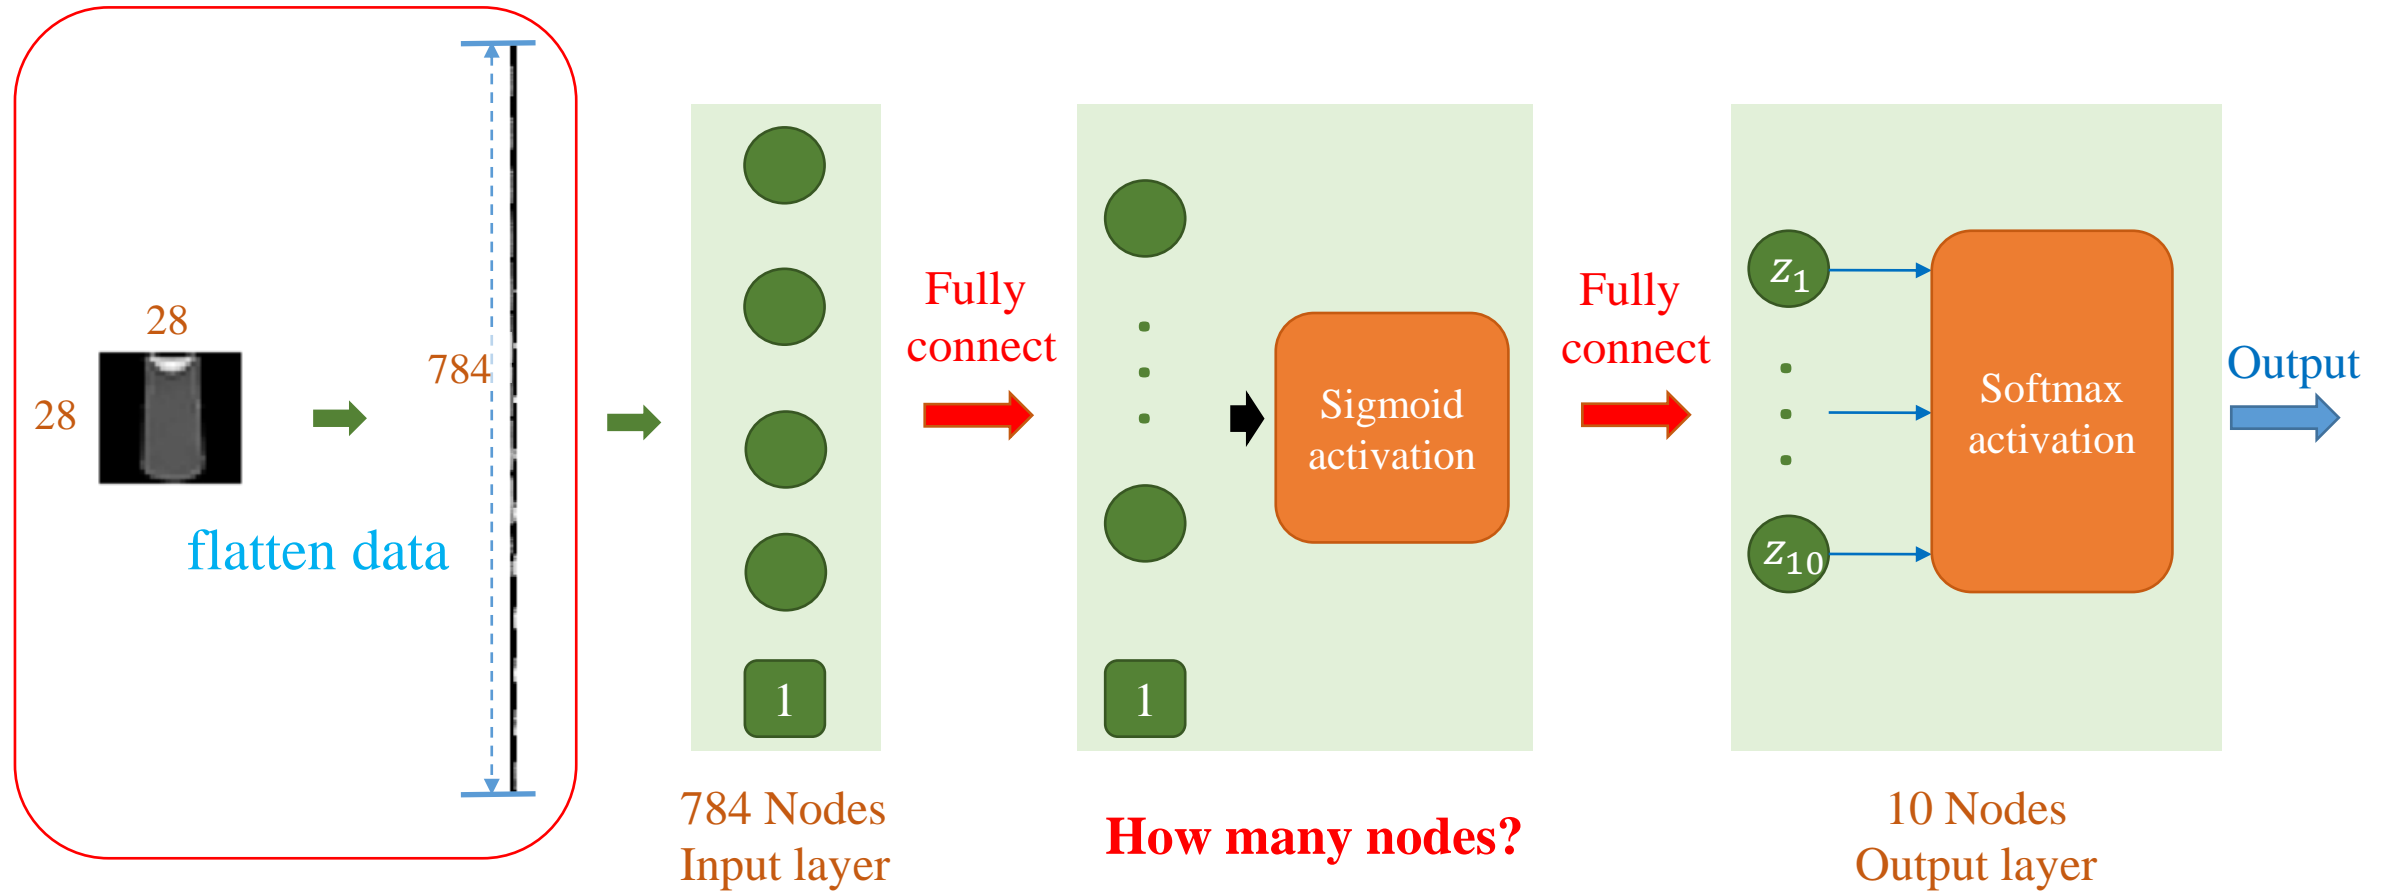
Model (Network) Construction



How many hidden layers?
How many nodes in a hidden layer?

Which activation function?
Which network components?

How many nodes?



How many nodes?

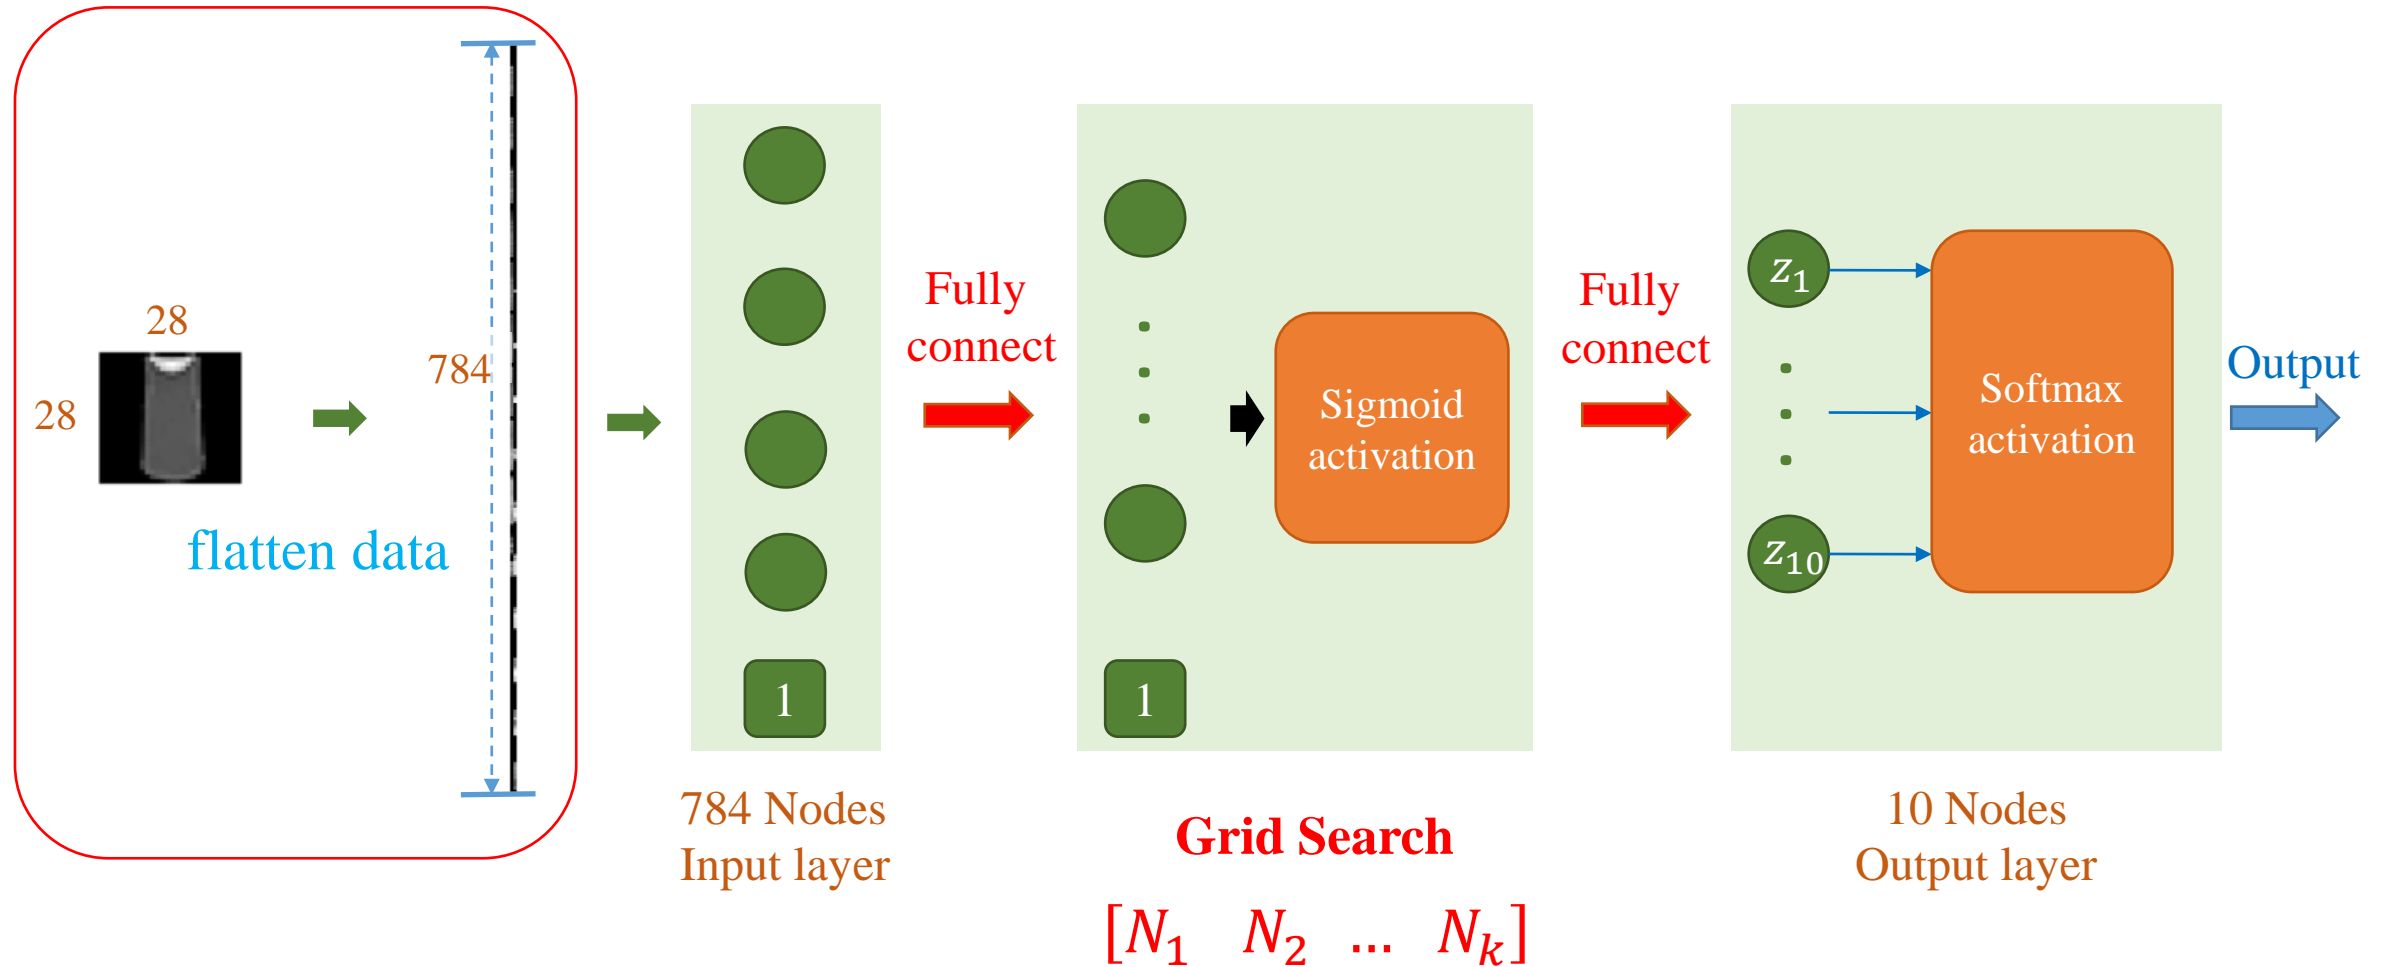


Image Classification

Cifar-10 dataset

Color images

Resolution=32x32

Training set: 50000 samples

Testing set: 10000 samples

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck



