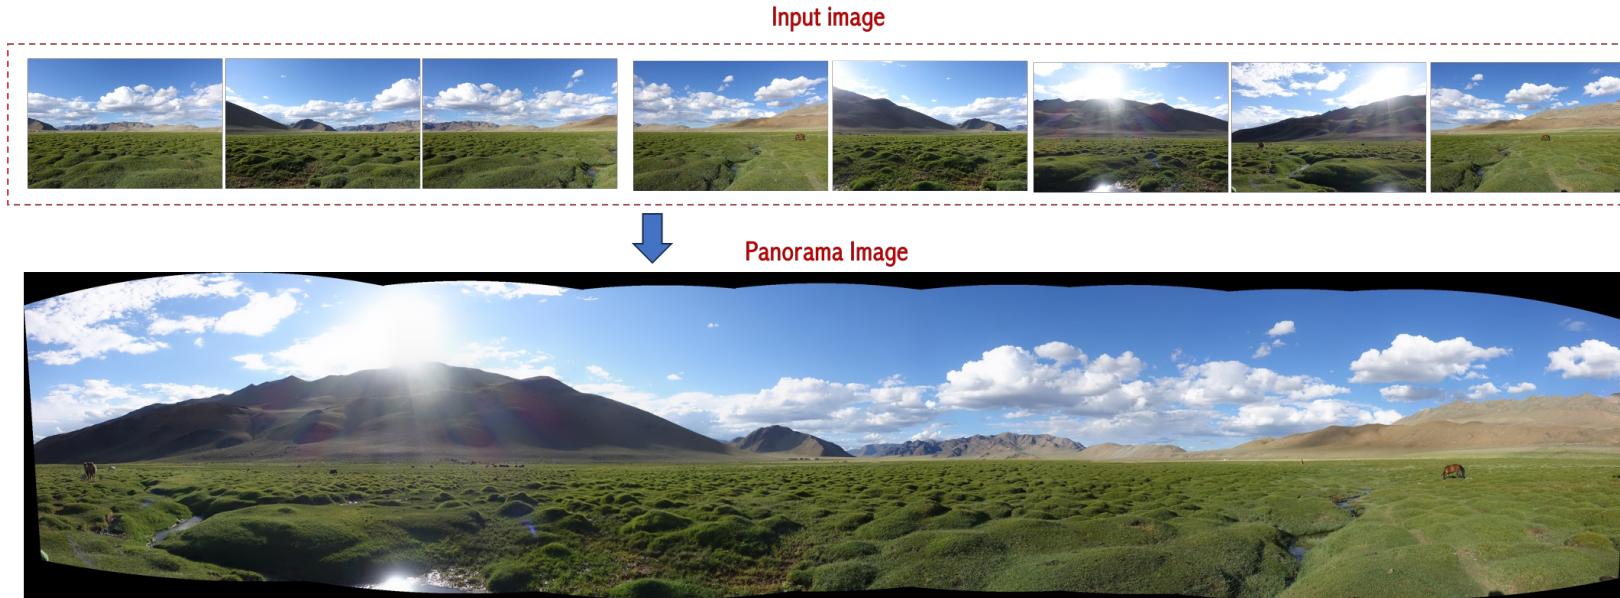


Image Stitching: Panorama Image

(Computer Vision Foundation)



Vinh Dinh Nguyen
PhD in Computer Science

Outline

- **Image Stitching/Panorama Image**
- **Edge Detector**
- **Blob Detector**
- **SIFT detector**
- **Image Transformation: 2D & 3D**
- **Image stitching/panorama Techniques**

Outline

- **Image Stitching/Panorama Image**
- **Edge Detector**
- **Blob Detector**
- **SIFT detector**
- **Image Transformation: 2D & 3D**
- **Image stitching/panorama Techniques**

How to build a Panaroma

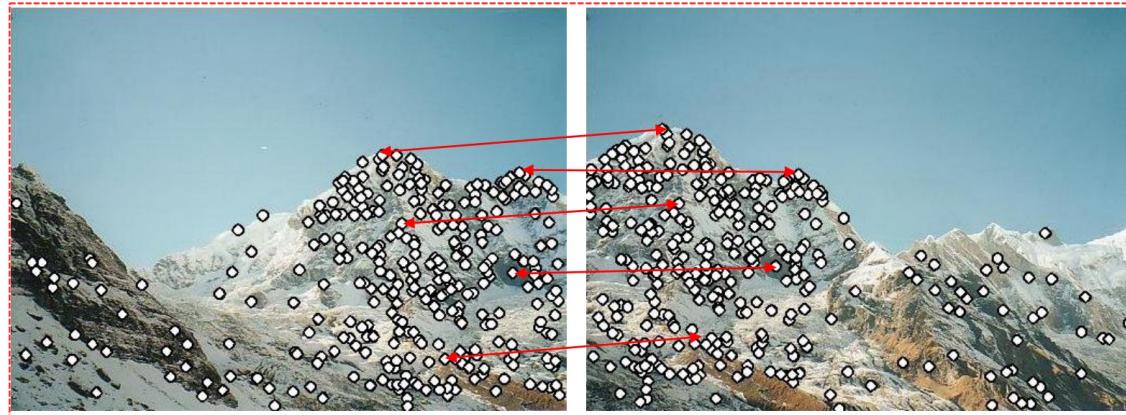


We need to match (align) images



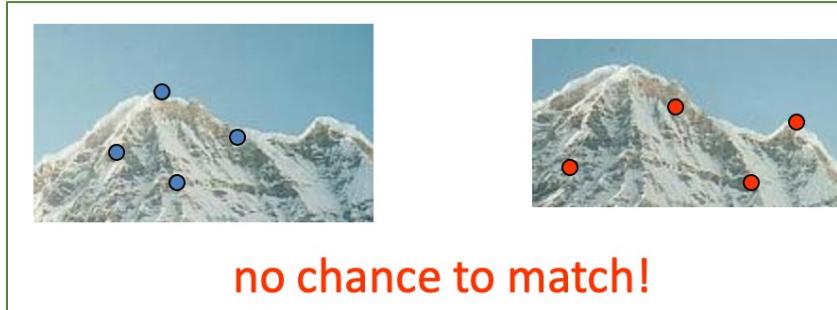
How to build a Panorama

- Detect feature points in both images
- Find corresponding pairs
- Use these pairs to align images

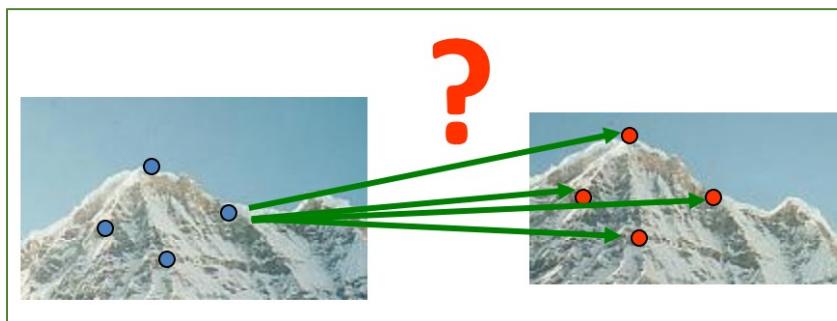


Problems

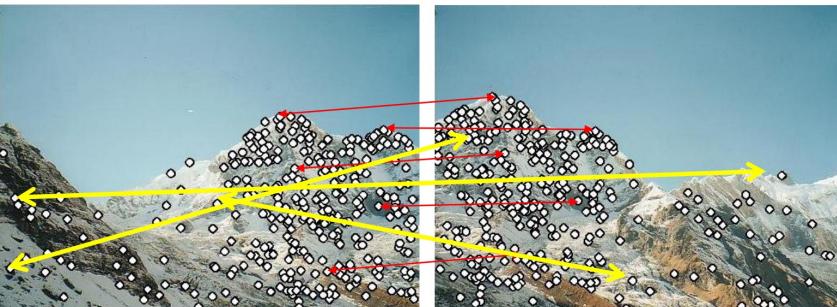
Detect the same point independently in both images



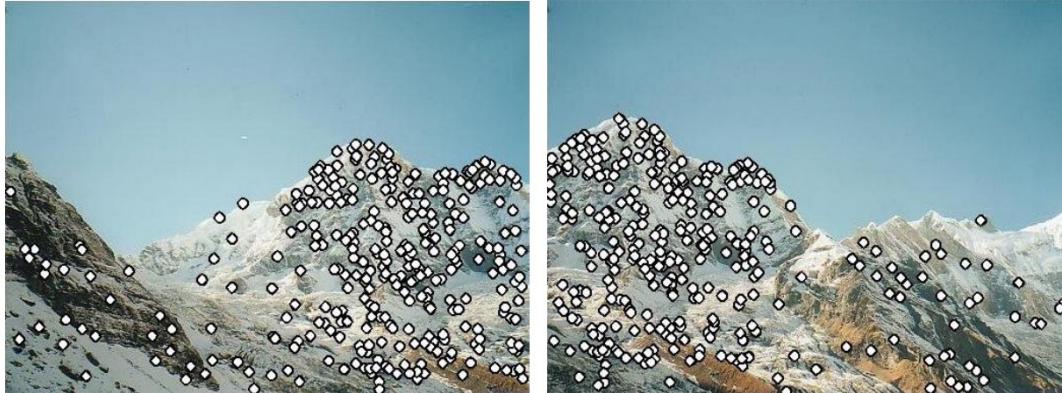
For each point correctly recognize the corresponding one



Need to estimate transformation between images, despite erroneous correspondences



Characteristics of Good Features

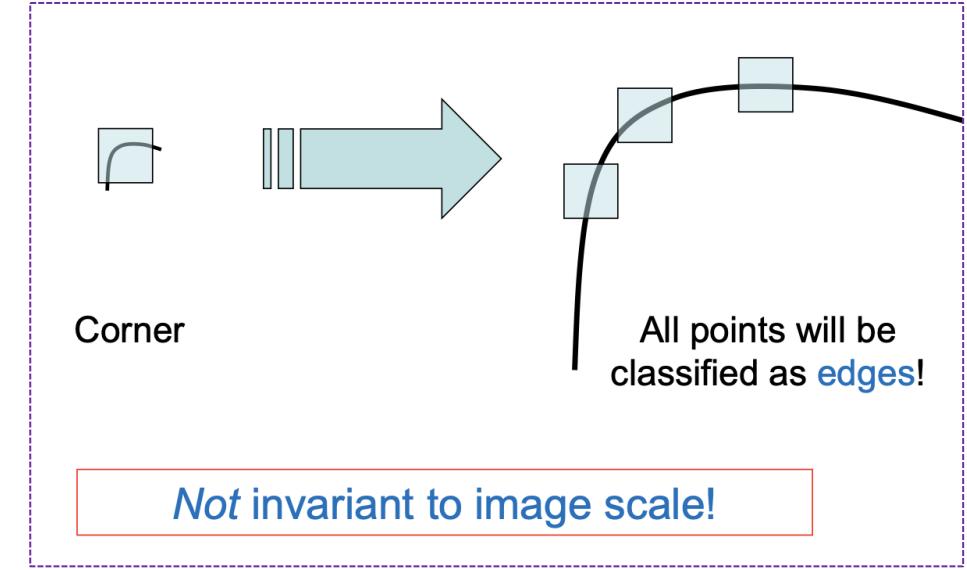
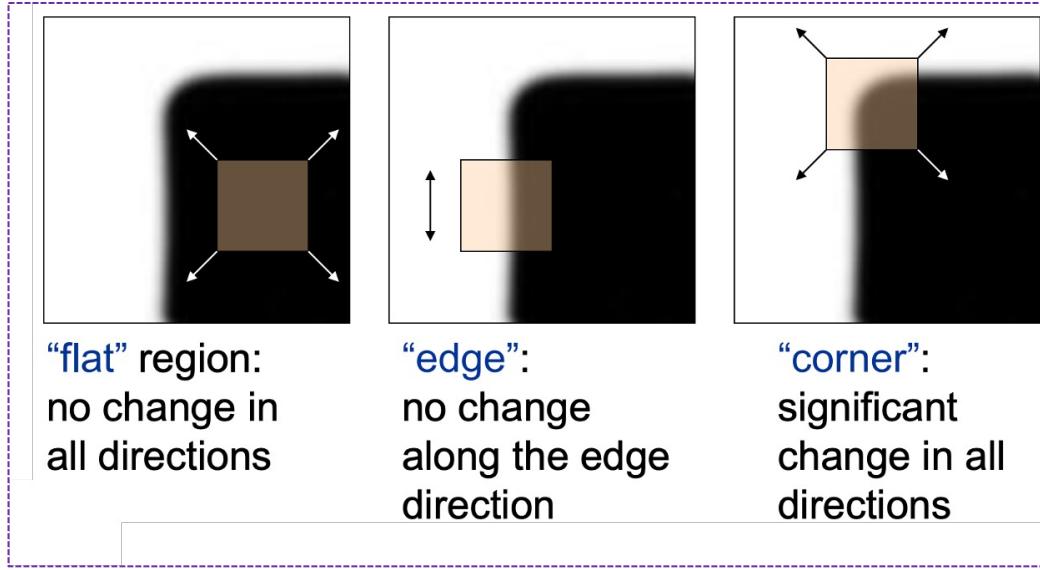


- **Repeatability**
 - The same feature can be found in several images despite geometric and photometric transformations
- **Saliency**
 - Each feature has a distinctive description
- **Compactness and efficiency**
 - Many fewer features than image pixels
- **Locality**
 - A feature occupies a relatively small area of the image; robust to clutter and occlusion

Feature points are used for:

- Motion tracking
- Image alignment
- 3D reconstruction
- Object recognition
- Indexing and database retrieval
- Robot navigation

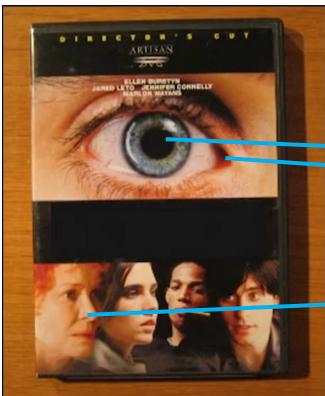
Corners



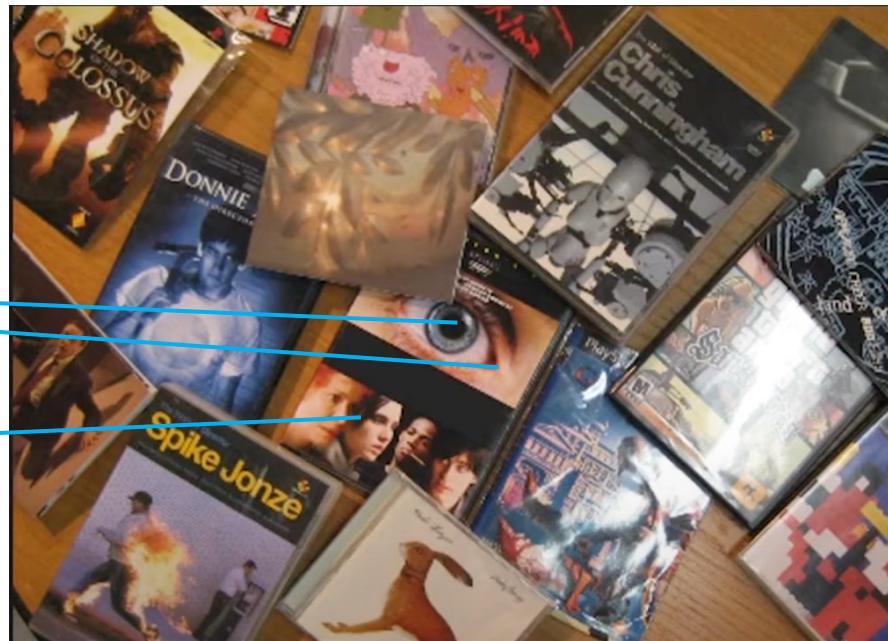
Limitations

SIFT Detector: Motivation

How would you recognize the following types of objects



Template



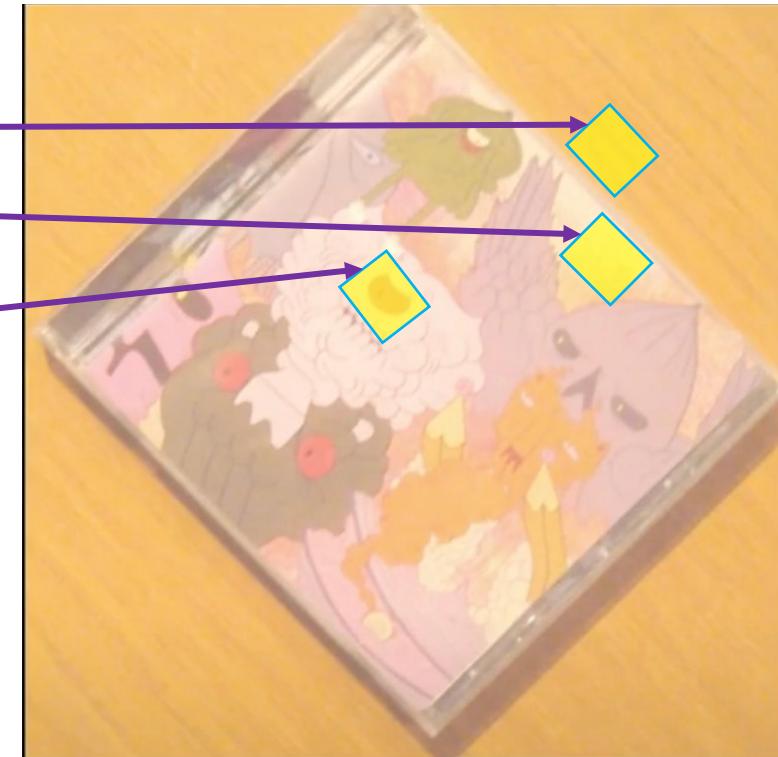
Find and Match “Interesting Points or Features”

Scale Invariant Feature Transformation (SIFT): Image alignment and 2D object recognition

Image credit: Professor Shree Nayar who is faculty in the Computer Science Department, School of Engineering and Applied Sciences, Columbia University.

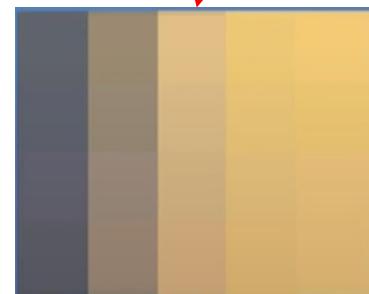
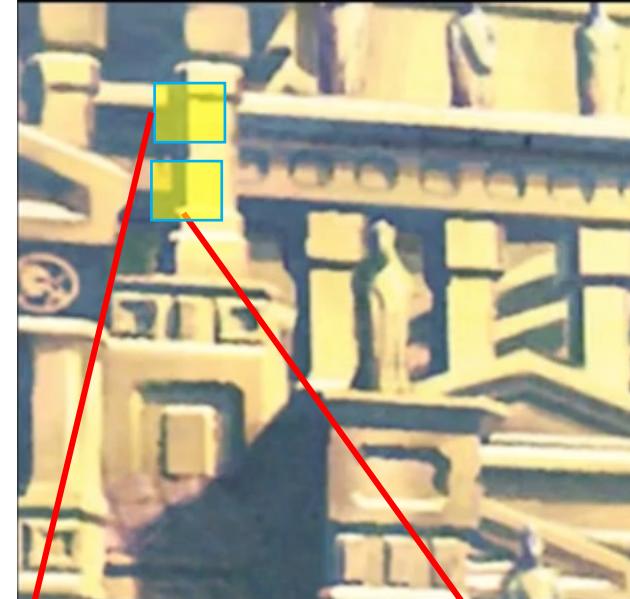
SIFT Detector: Interesting Point

Different size, orientation, lighting, brightness, etc,...

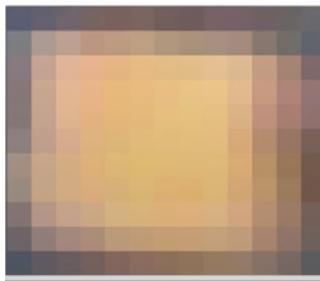
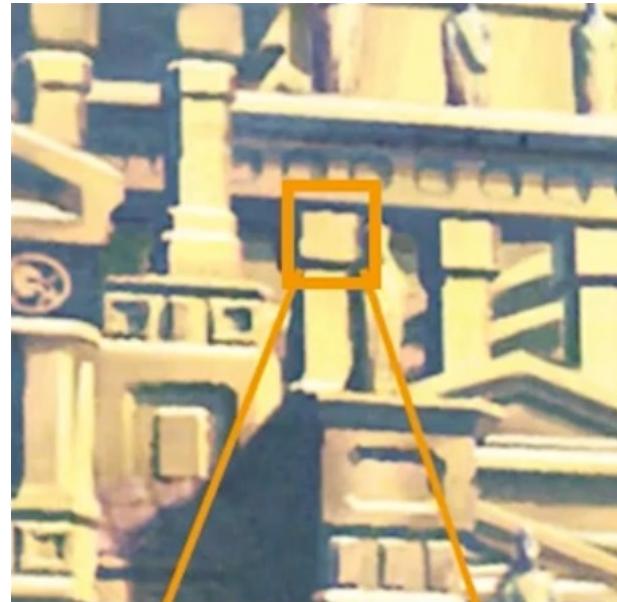


Interesting point: (1) Rich content within the local window. (2) Well-defined representation for matching/comparing with other points. (3) Well-defined position in the image. (4) Should be invariant to image rotation and scaling. (5) Should be insensitive to lighting changes

Are Line/Edges Interesting Point?



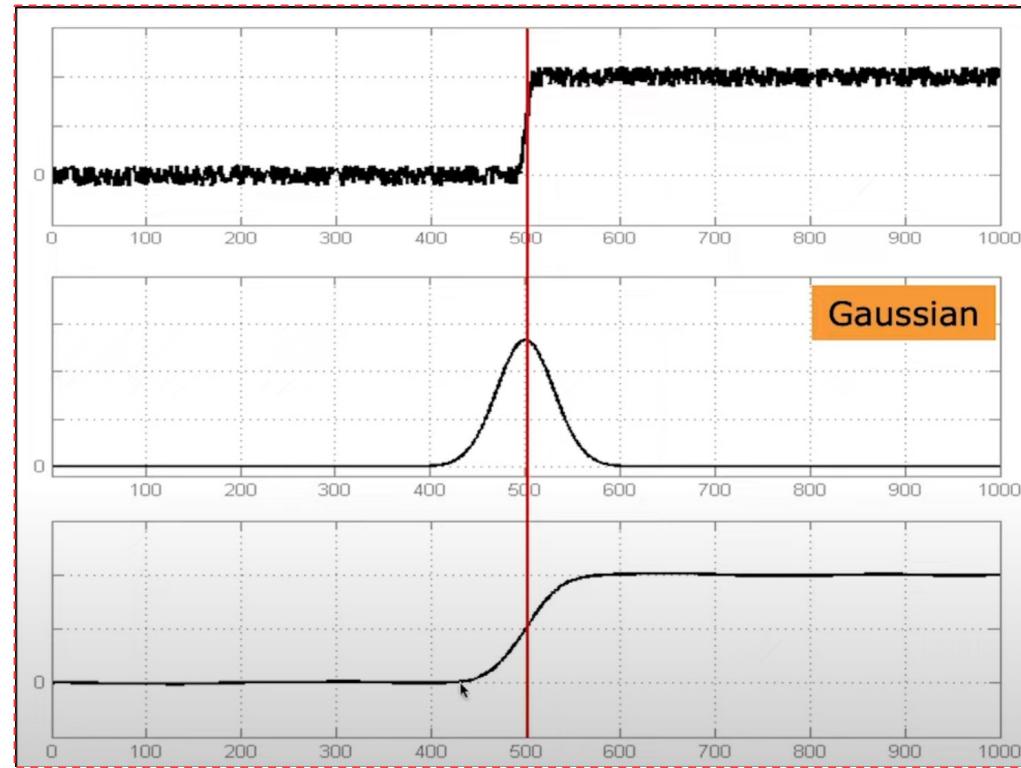
Are Blob Interesting?



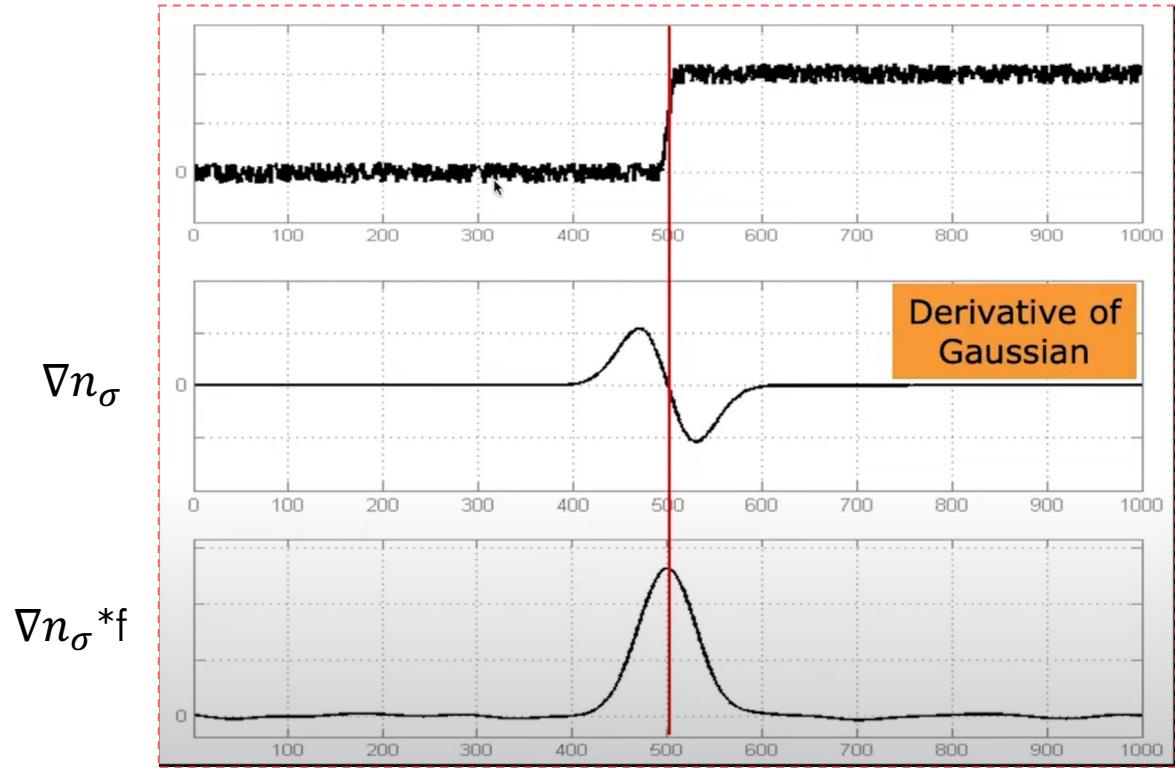
Blob has fixed position and definite size: location, size, and orientation

A Blob, in a sense, is anything that is considered a large object or anything bright in a dark background, in images, we can generalize it as a group of pixel values that forms a somewhat colony or a large object that is distinguishable from its background.

Gaussian Filter: Review

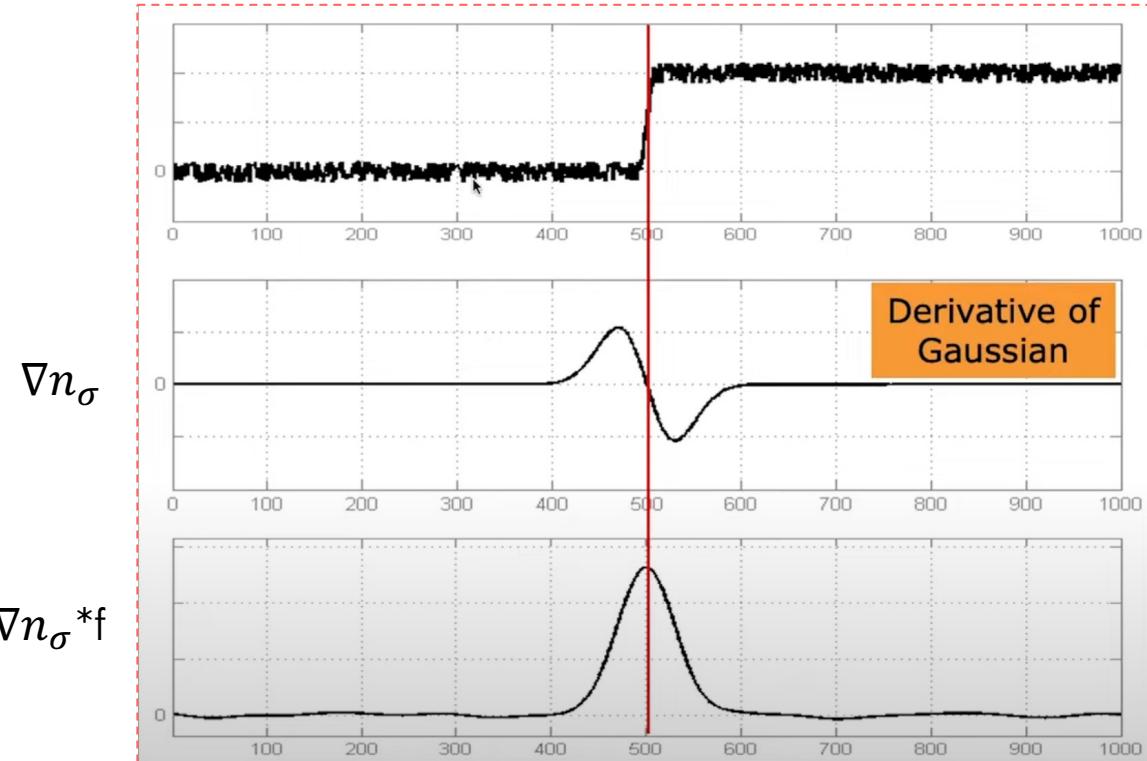


Gaussian Filter is used to remove noise by smoothing

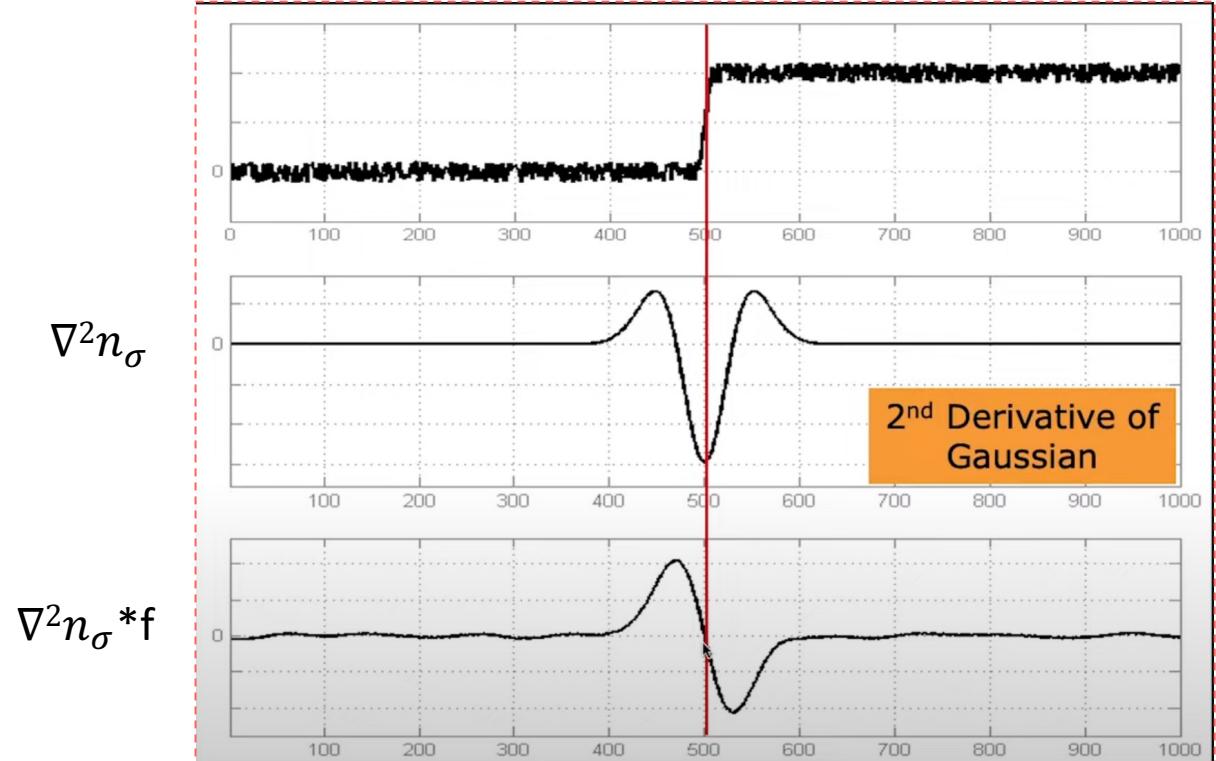


Extremum of Derivative of Gaussian denotes an edge.

Gaussian Filter: Review



Extremum of Derivative of Gaussian denotes an edge.

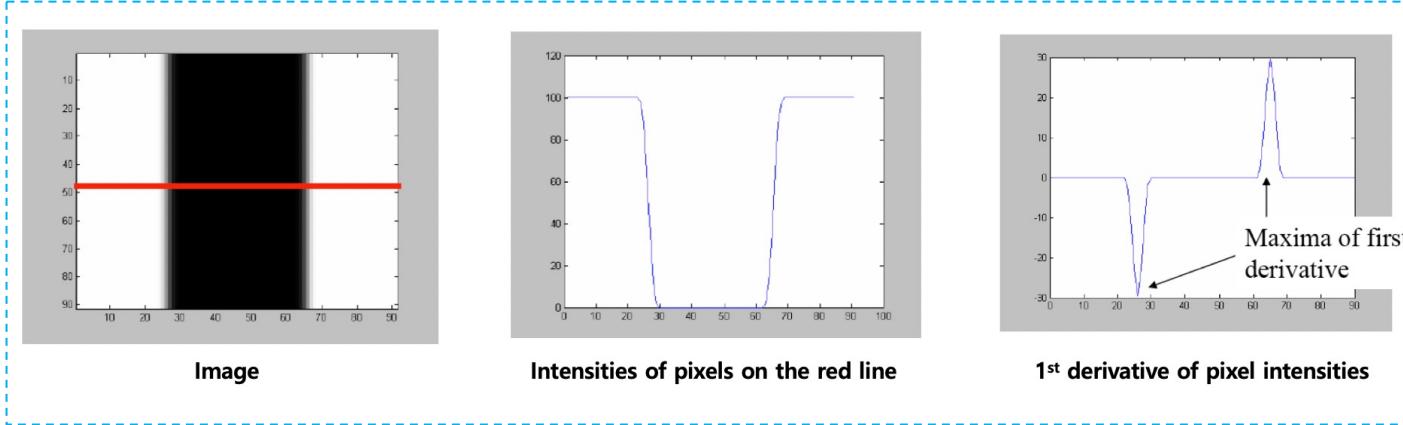


Zero crossing in 2nd derivative of Gaussian denote an Edge

Outline

- What is image stitching/panorama image
- Edge Detector
- Blob Detector
- SIFT detector
- Image Transformation: 2D & 3D
- Image stitching/panorama Techniques

Edge Detection



$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

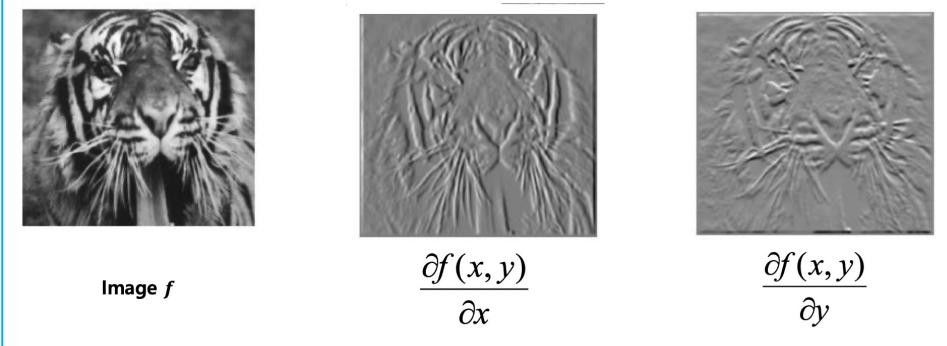
Continuous data

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x+1, y) - f(x, y)}{1}$$

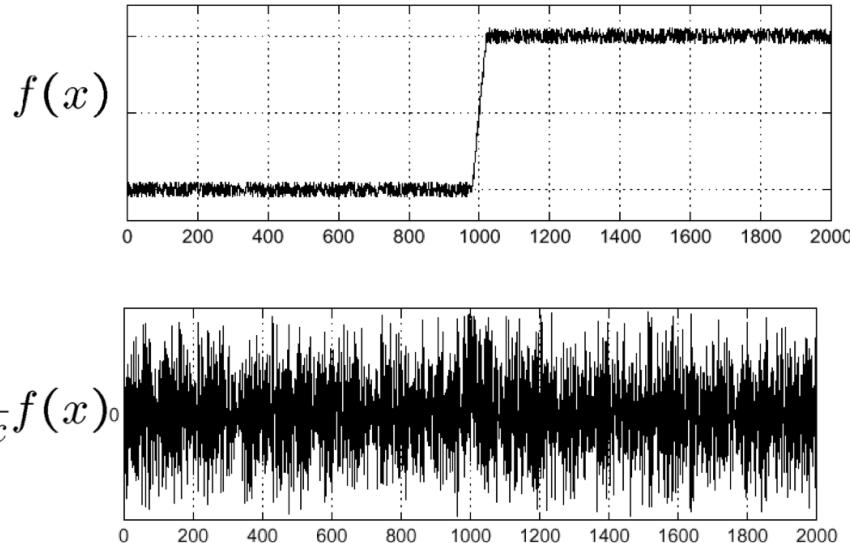
Discrete data

1	-1
---	----

Kernel

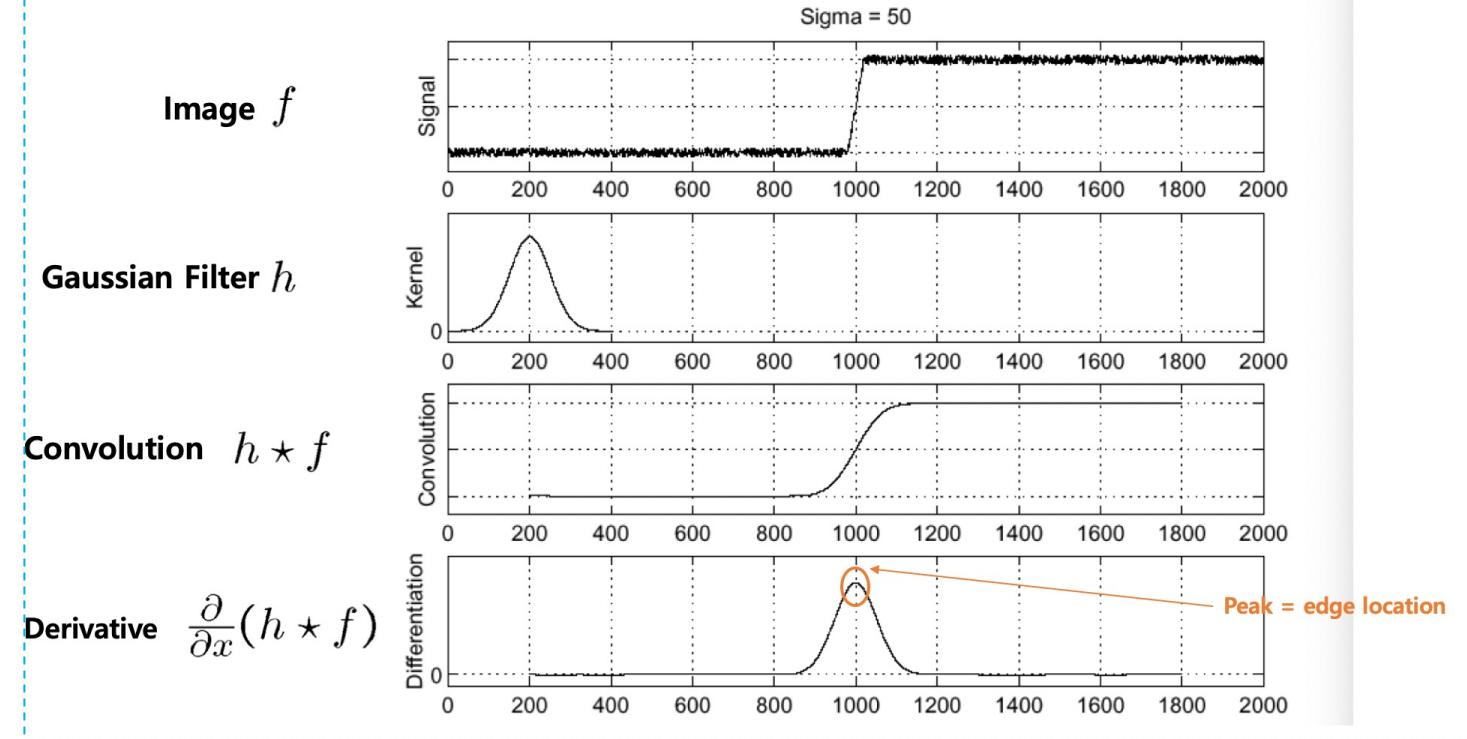


Edge Detection with Gaussian Filter



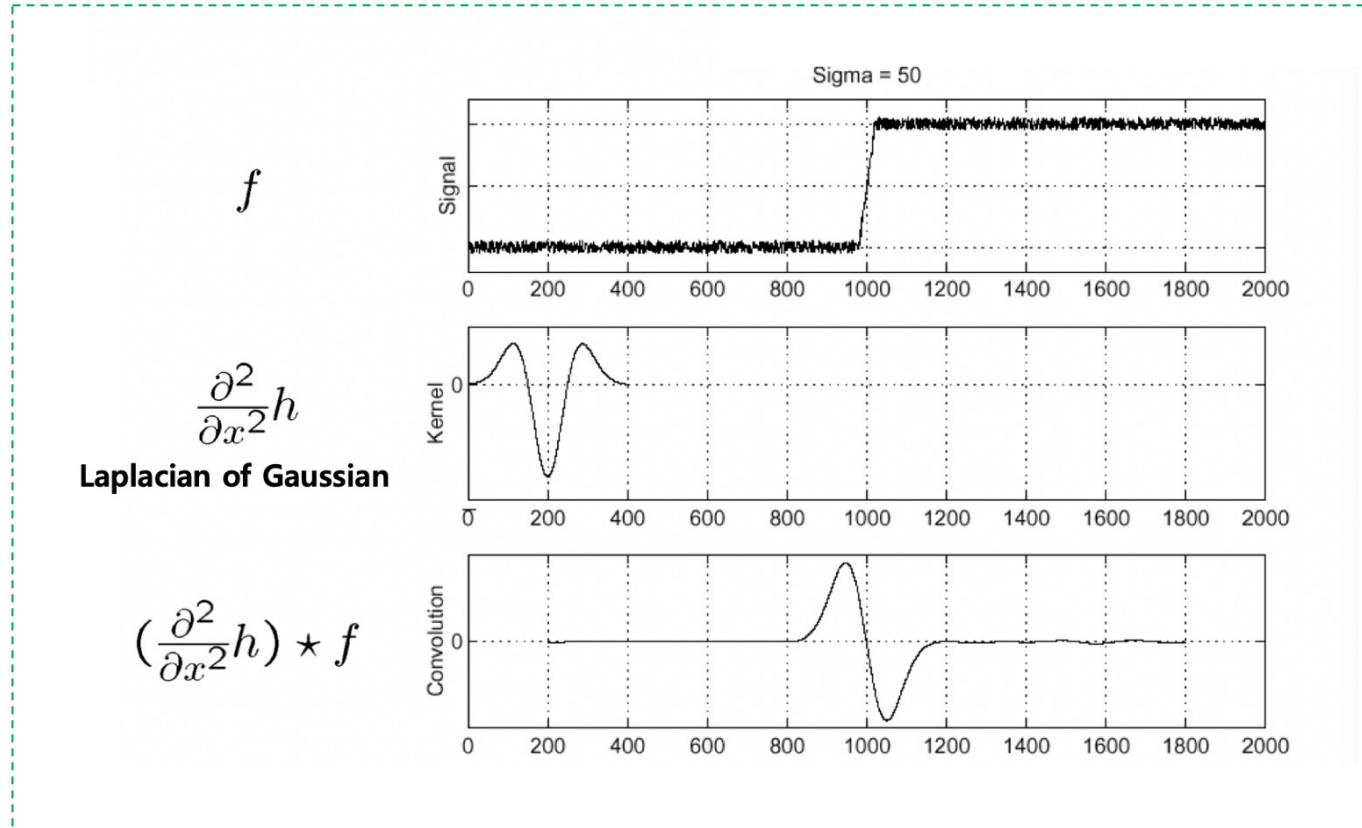
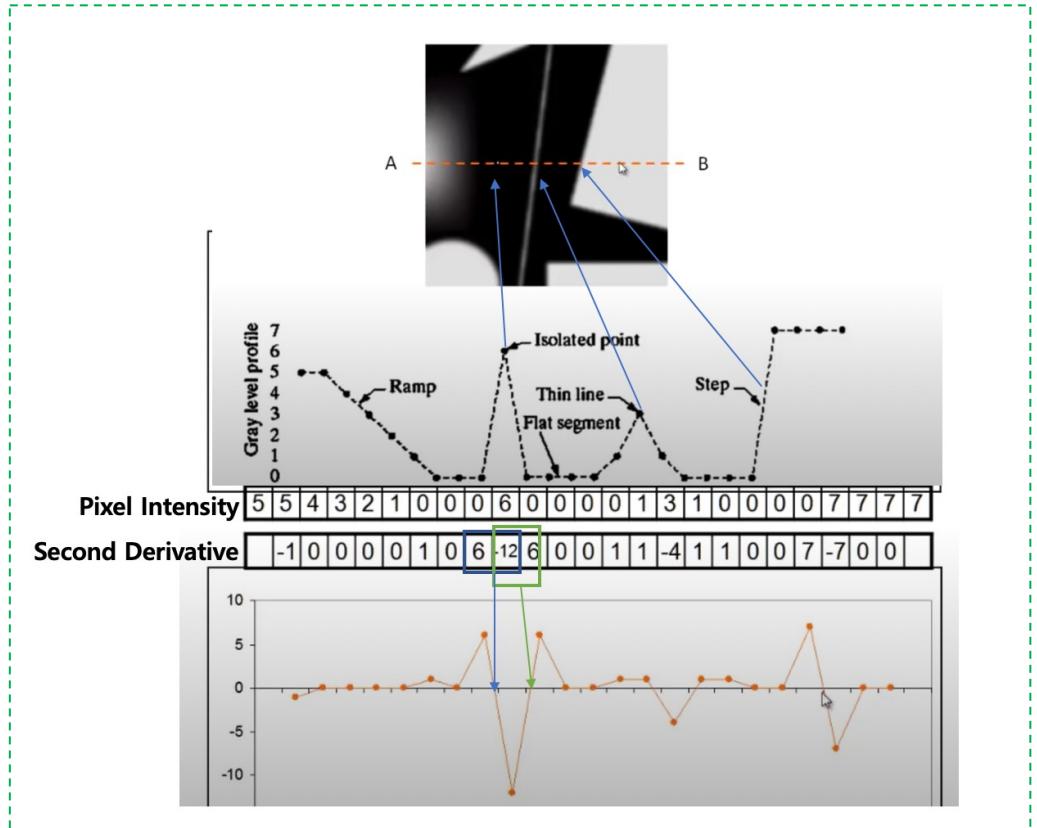
$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

$[1 \quad -1]$	\star	$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix}$
Derivative Filter		
Derivative of Gaussian		



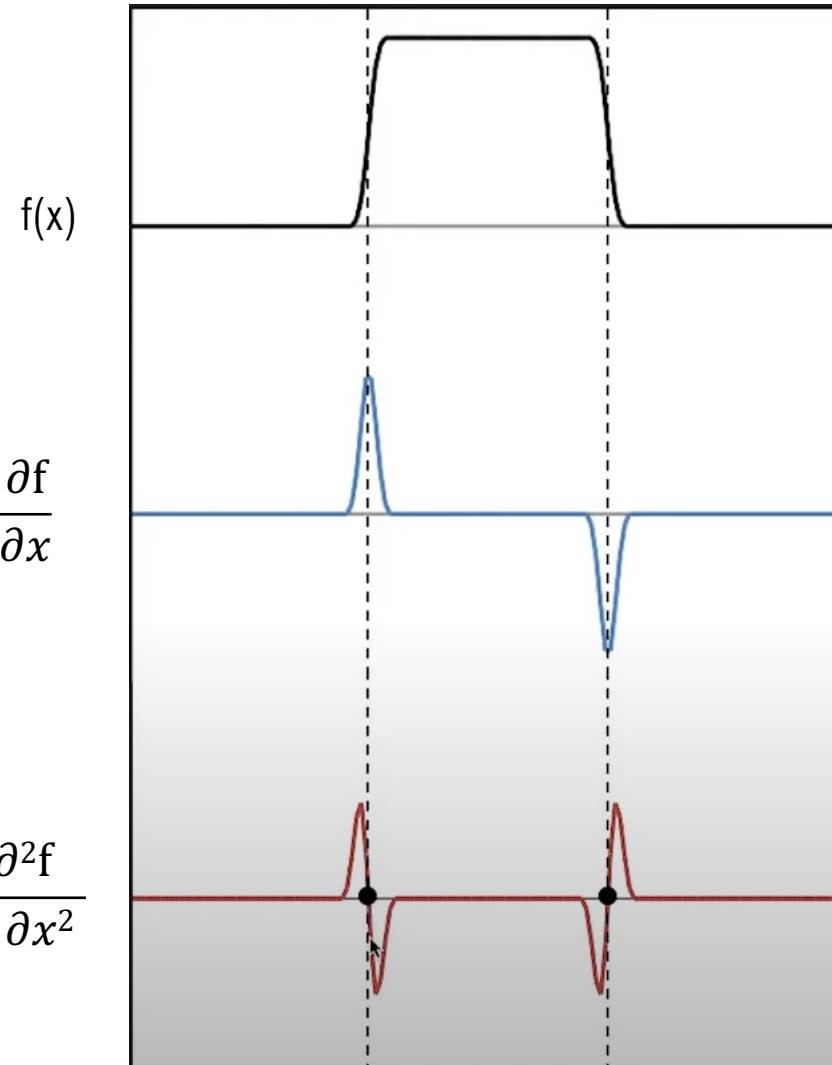
The result of this convolution (Derivative filter $\star h$) is called the Derivative of Gaussian

Laplacian as Edge Detector



$$\begin{aligned}
 \text{Laplacian } \nabla^2 f &= \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2} \\
 &= f(x+1, y) + f(x-1, y) - 2f(x, y) + f(x, y+1) + f(x, y-1) - 2f(x, y) \\
 &= f(x+1, y) + f(x-1, y) + f(x, y+1) + f(x, y-1) - 4f(x, y)
 \end{aligned}$$

Laplacian as Edge Detector



Local extrema
indicate Edges

Zero-crossing
indicate edges

Laplacian: Sum of Pure Second Derivatives

$$\nabla^2 I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$

Edges are “zero-crossing” in Laplacian of image
Laplacian does not provide directions of edges

0	1	0
1	-4	1
0	1	0

Kernel

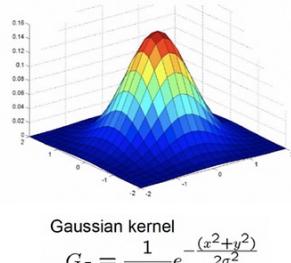
$$\begin{aligned}\text{Laplacian } \nabla^2 f &= \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2} \\ &= f(x+1, y) + f(x-1, y) - 2f(x, y) + f(x, y+1) + f(x, y-1) - 2f(x, y) \\ &= f(x+1, y) + f(x-1, y) + f(x, y+1) + f(x, y-1) - 4f(x, y)\end{aligned}$$

Laplacian of Gaussian

$$LoG = -\frac{1}{\pi\sigma^4} \left[1 - \frac{x^2 + y^2}{2\sigma^2} \right] e^{-\frac{x^2+y^2}{2\sigma^2}}$$

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

★



$$= \begin{array}{|cccccccccc|} \hline 0 & 1 & 1 & 2 & 2 & 2 & 1 & 1 & 0 \\ \hline 1 & 2 & 4 & 5 & 5 & 5 & 4 & 2 & 1 \\ 1 & 4 & 5 & 3 & 0 & 3 & 5 & 4 & 1 \\ 2 & 5 & 3 & -12 & -24 & -12 & 3 & 5 & 2 \\ 2 & 5 & 0 & -24 & -40 & -24 & 0 & 5 & 2 \\ 2 & 5 & 3 & -12 & -24 & -12 & 3 & 5 & 2 \\ 1 & 4 & 5 & 3 & 0 & 3 & 5 & 4 & 1 \\ 1 & 2 & 4 & 5 & 5 & 5 & 4 & 2 & 1 \\ 0 & 1 & 1 & 2 & 2 & 2 & 1 & 1 & 0 \\ \hline \end{array}$$

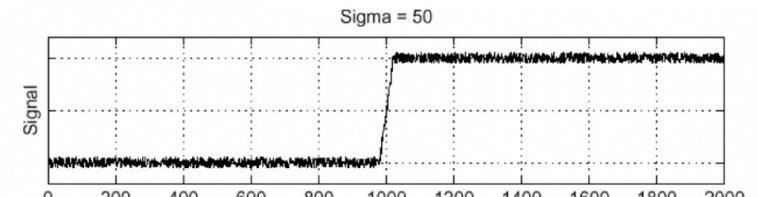
LoG Filter

Laplacian Filter

$$G(x, y; \sigma) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

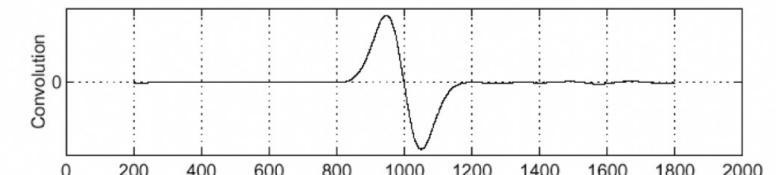
$$LoG(x, y; \sigma) = \Delta_{(x,y)} G(x, y; \sigma) = \frac{\partial^2 G(x, y; \sigma)}{\partial x^2} + \frac{\partial^2 G(x, y; \sigma)}{\partial y^2} = \frac{1}{\pi\sigma^4} \left(\frac{x^2 + y^2}{2\sigma^2} - 1 \right) e^{-\frac{x^2+y^2}{2\sigma^2}}$$

f

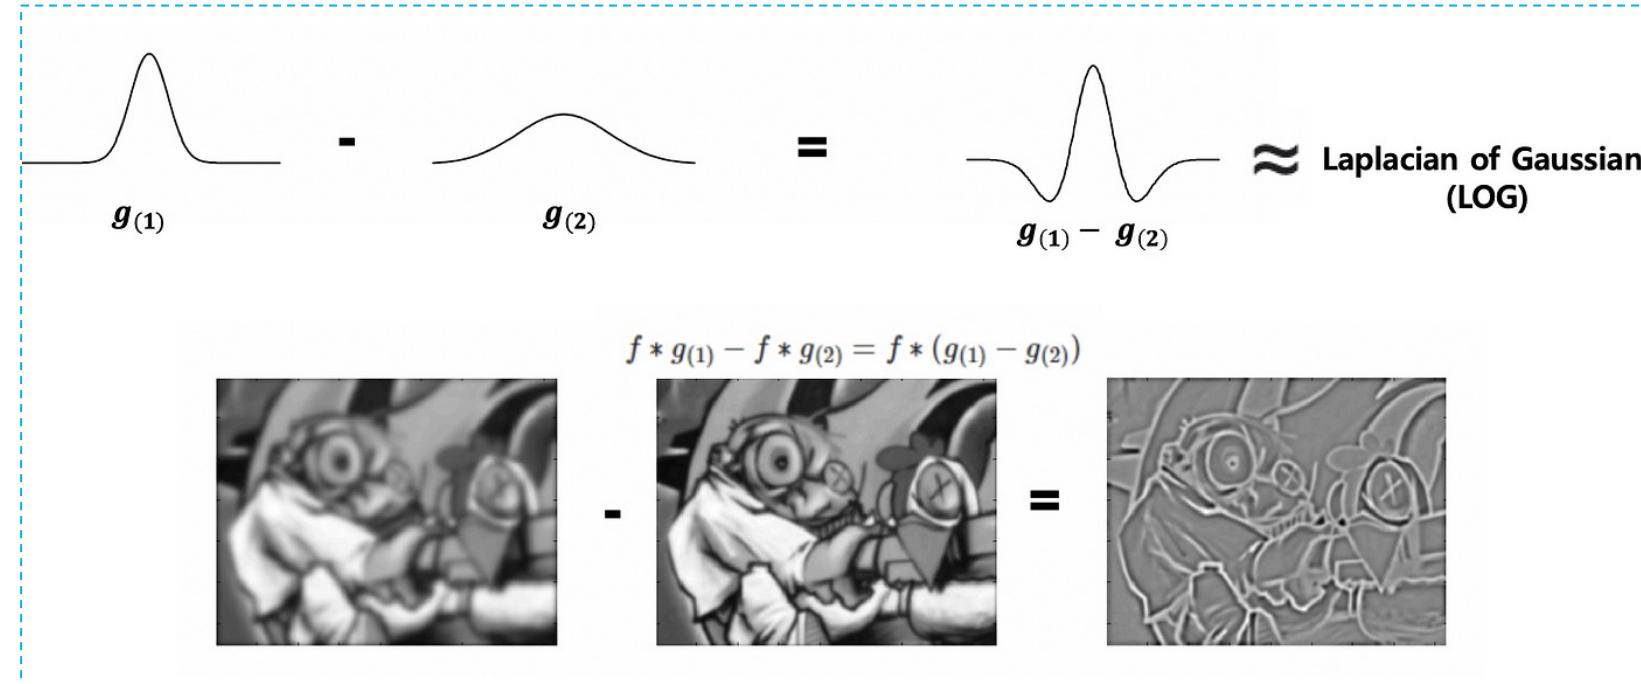


$$\frac{\partial^2 h}{\partial x^2}$$

Laplacian of Gaussian



Difference of Gaussians (DoG)



Subtracting one Gaussian by another approximates the Laplacian of Gaussian

This indicates while approximating the LoG, there is no actual derivative computation needed

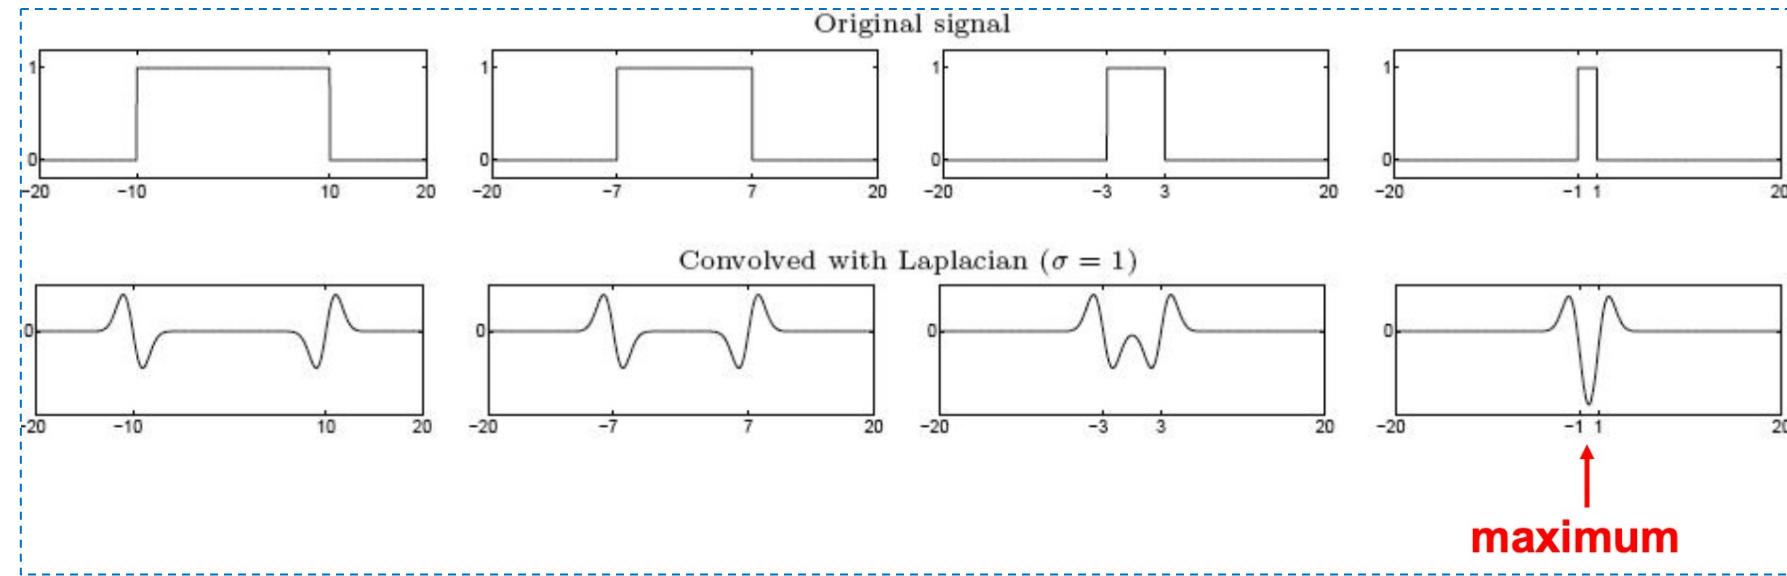
Outline

- What is image stitching/panorama image
- Edge Detector
- Blob Detector
- SIFT detector
- Image Transformation: 2D & 3D
- Image stitching/panorama Techniques

From Edges to Blobs

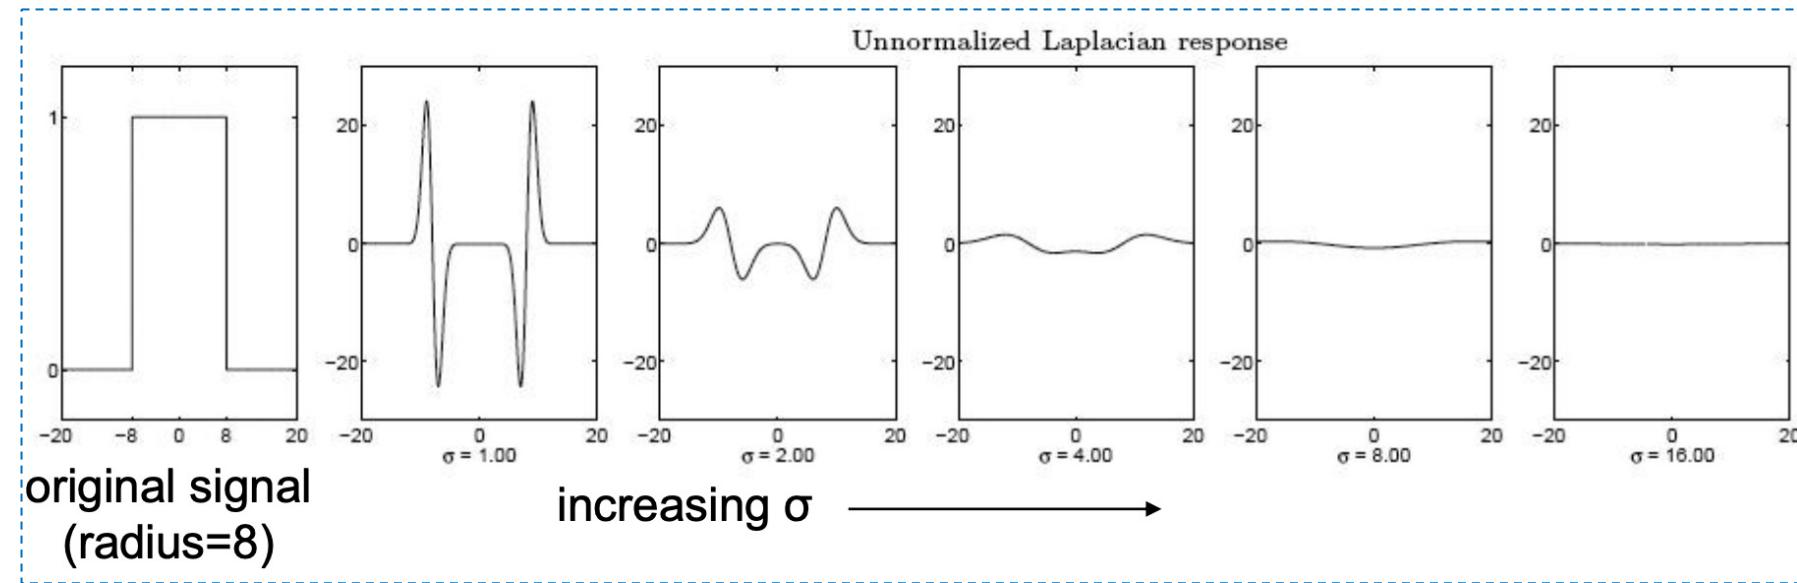
Edge = ripple

Blob = superposition of two ripples



From Edges to Blobs

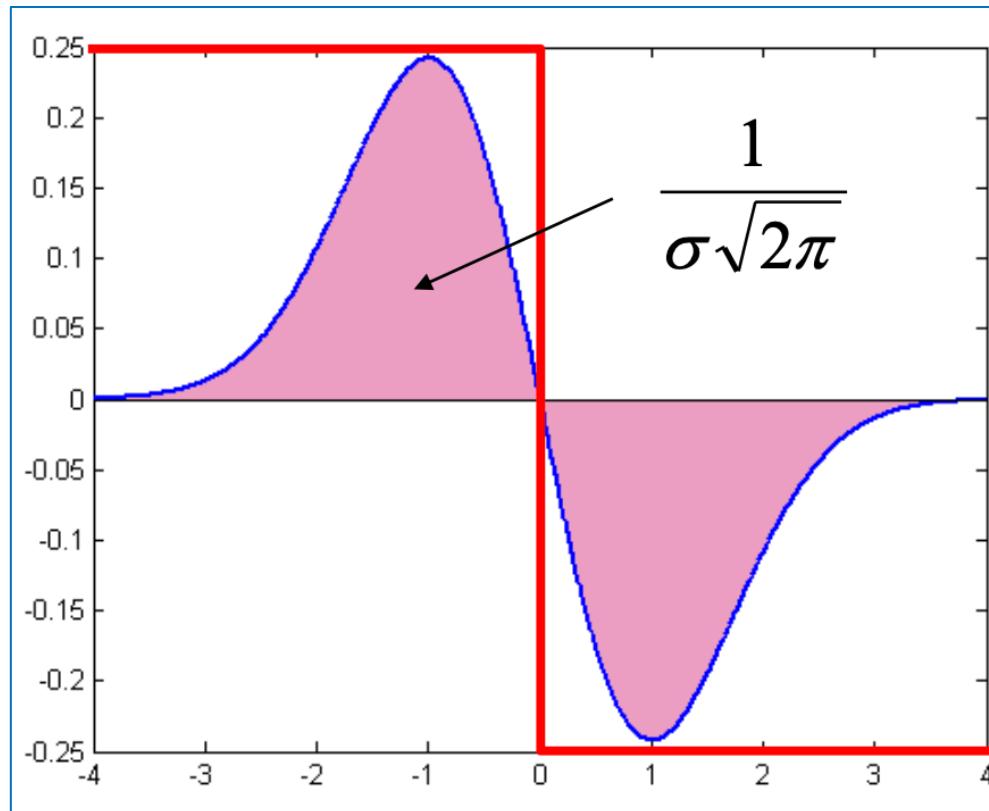
- We want to find the characteristic scale of the blob by convolving it with Laplacians at several scales and looking for the maximum response
- However, Laplacian response decays as scale increases



Why does this happen?

Scale normalization

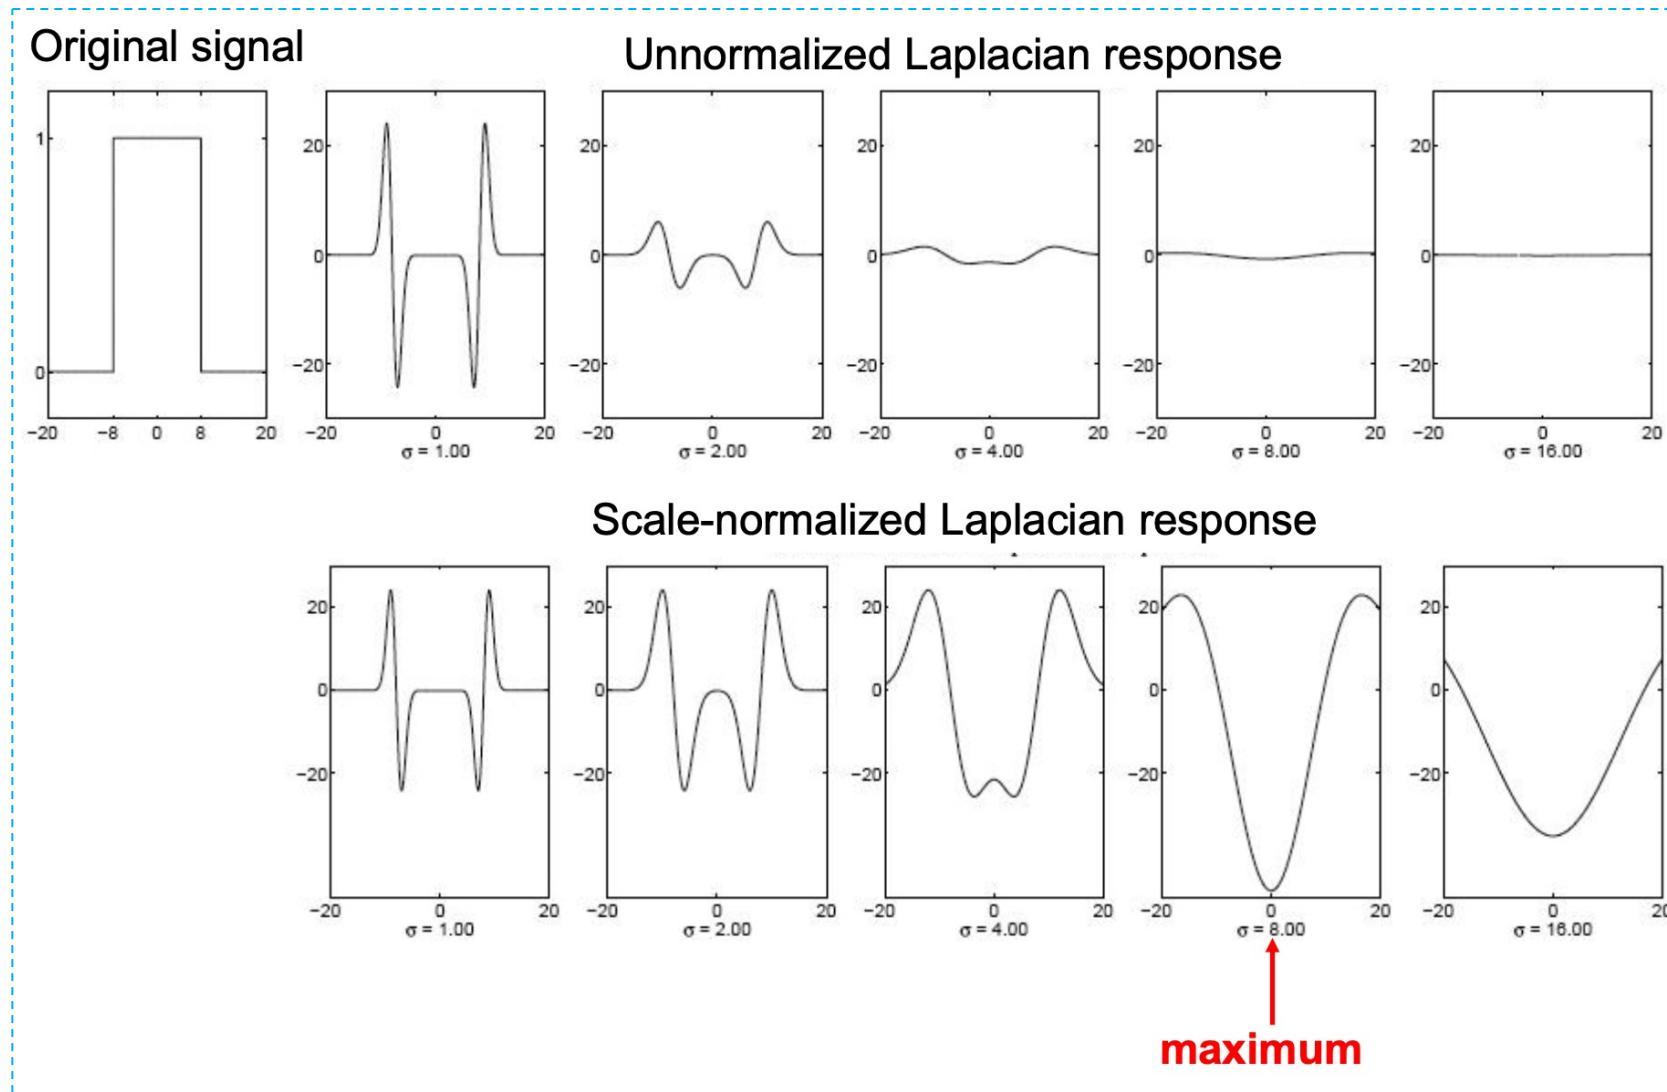
The response of a derivative of Gaussian filter to a perfect step edge decreases as σ increases



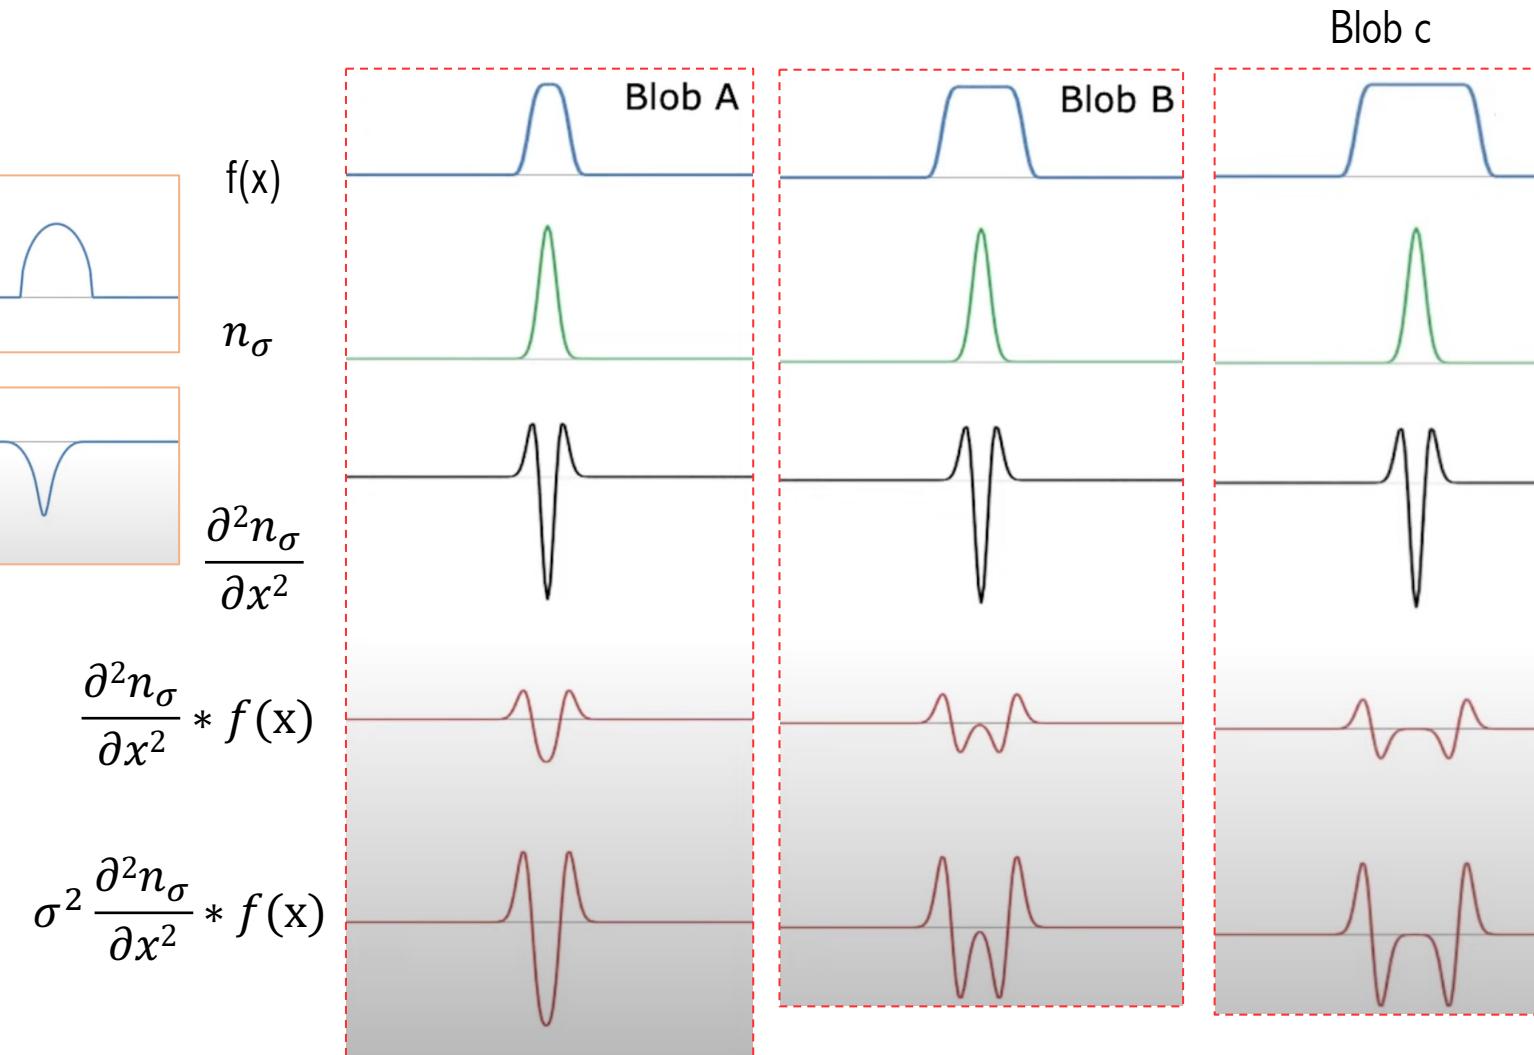
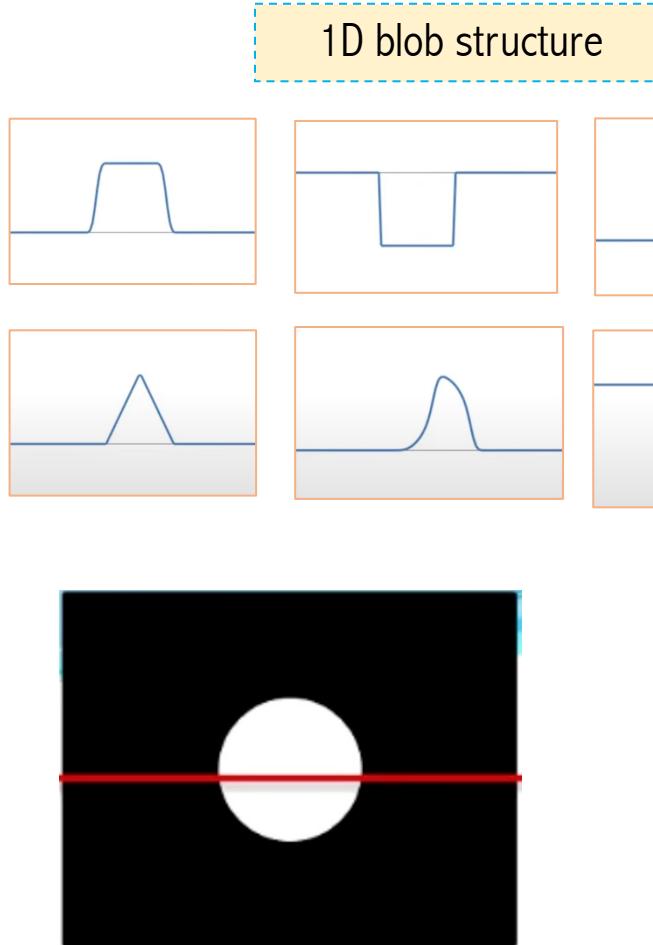
To keep response the same (scale-invariant), must multiply Gaussian derivative by σ

Laplacian is the second Gaussian derivative, so it must be multiplied by σ^2

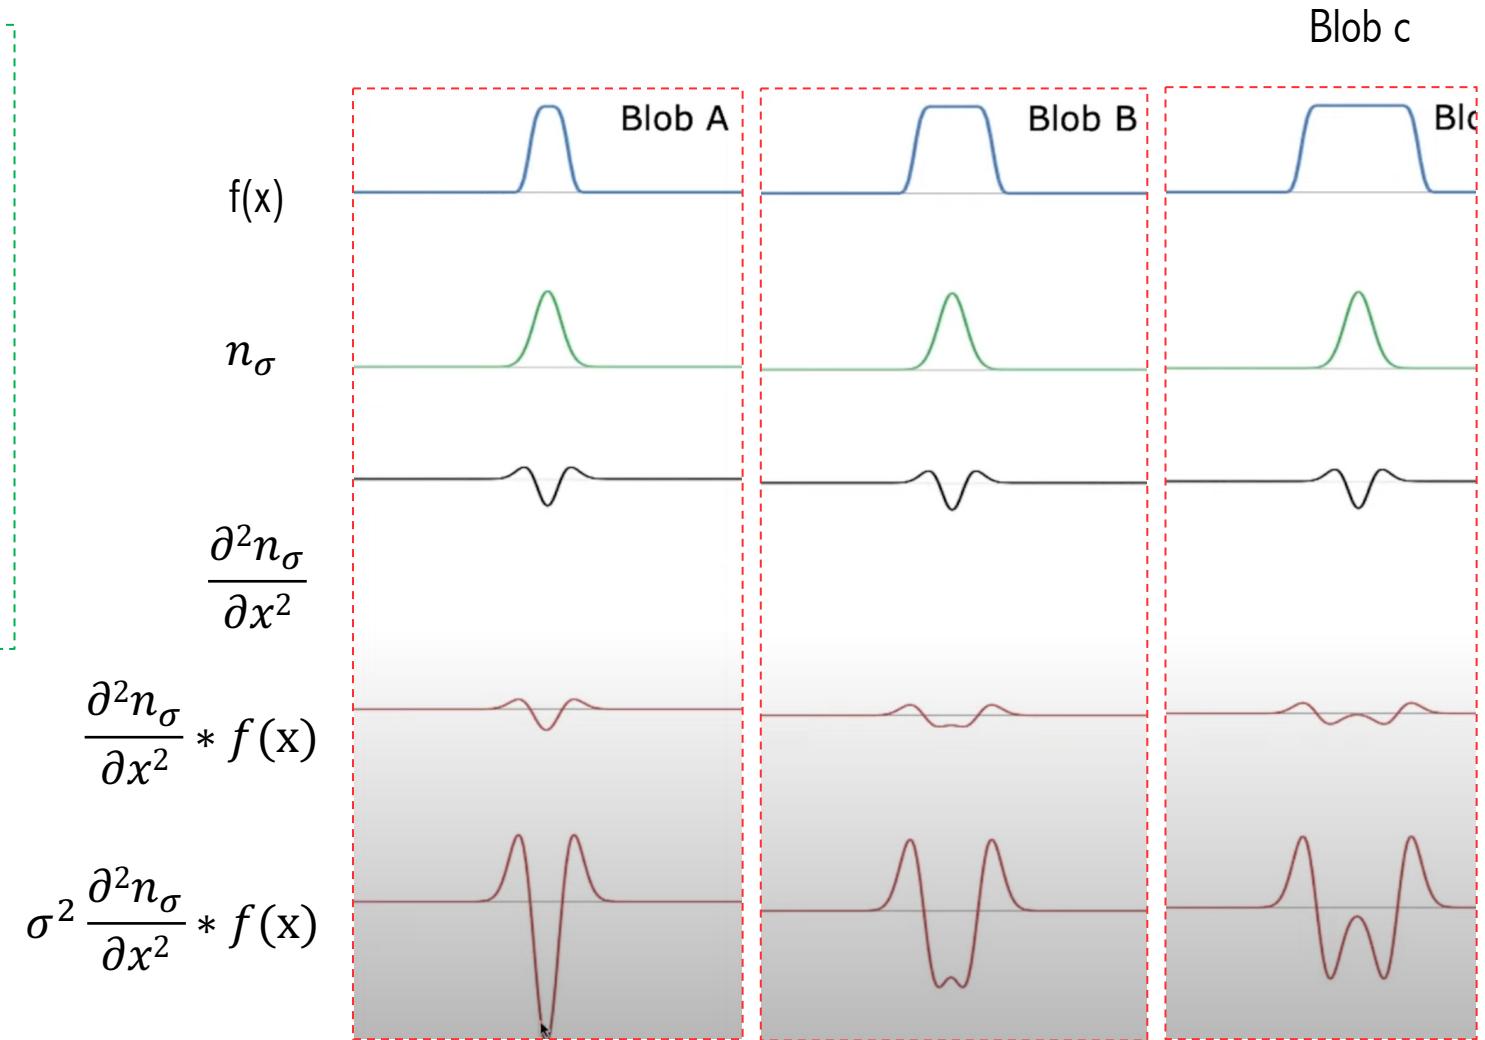
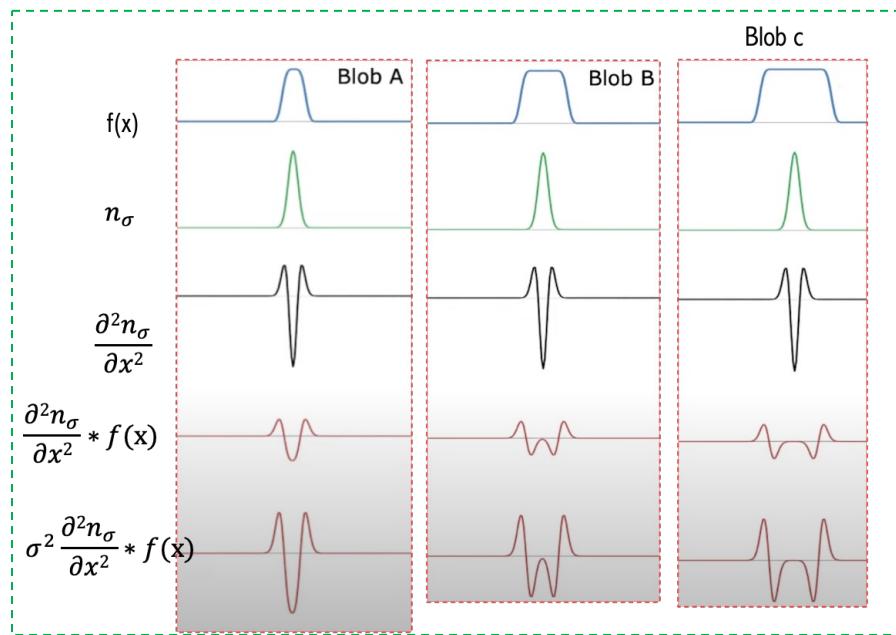
Effect of Scale Normalization



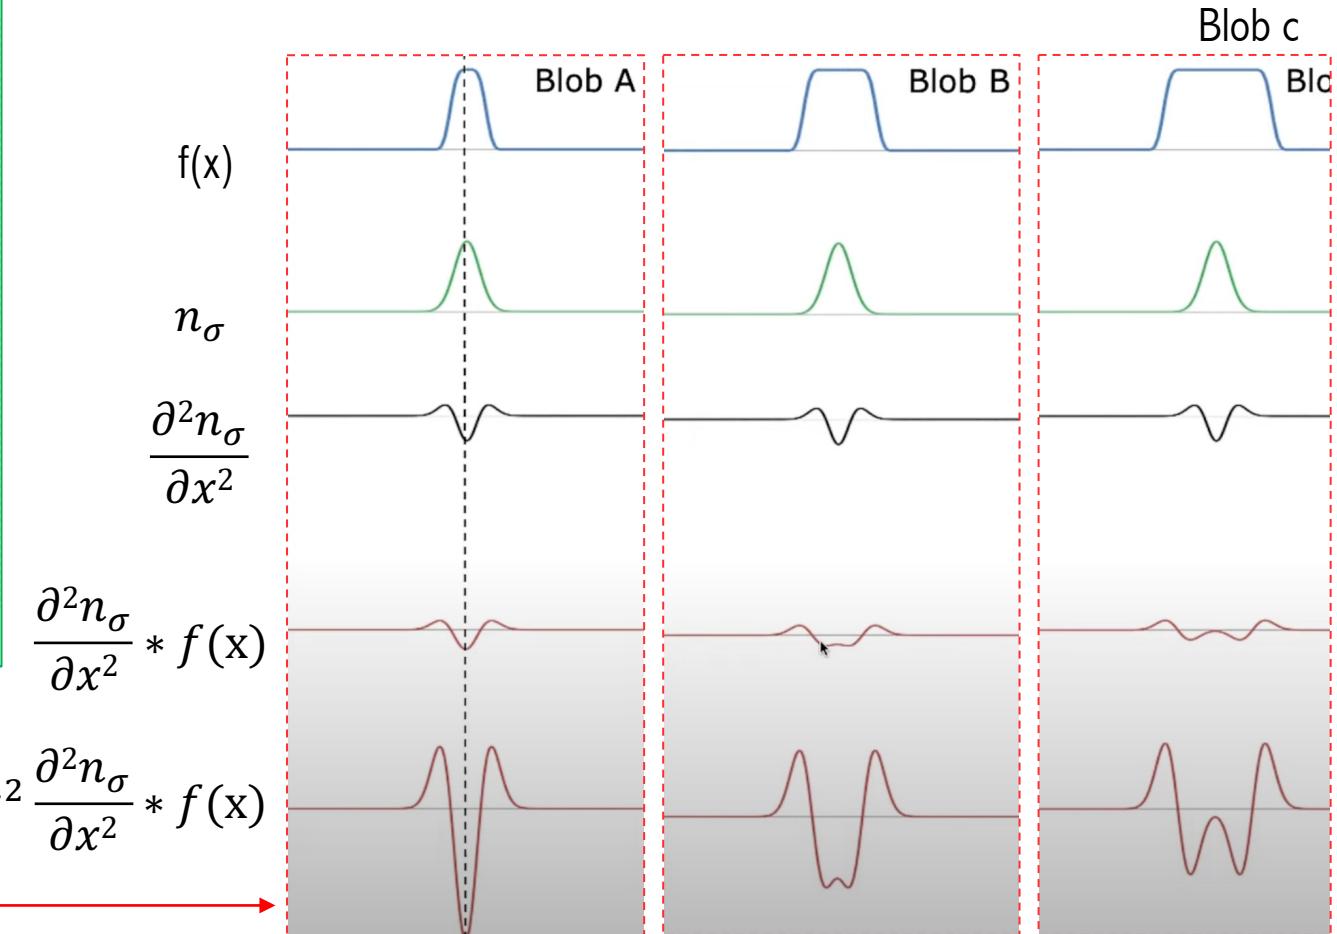
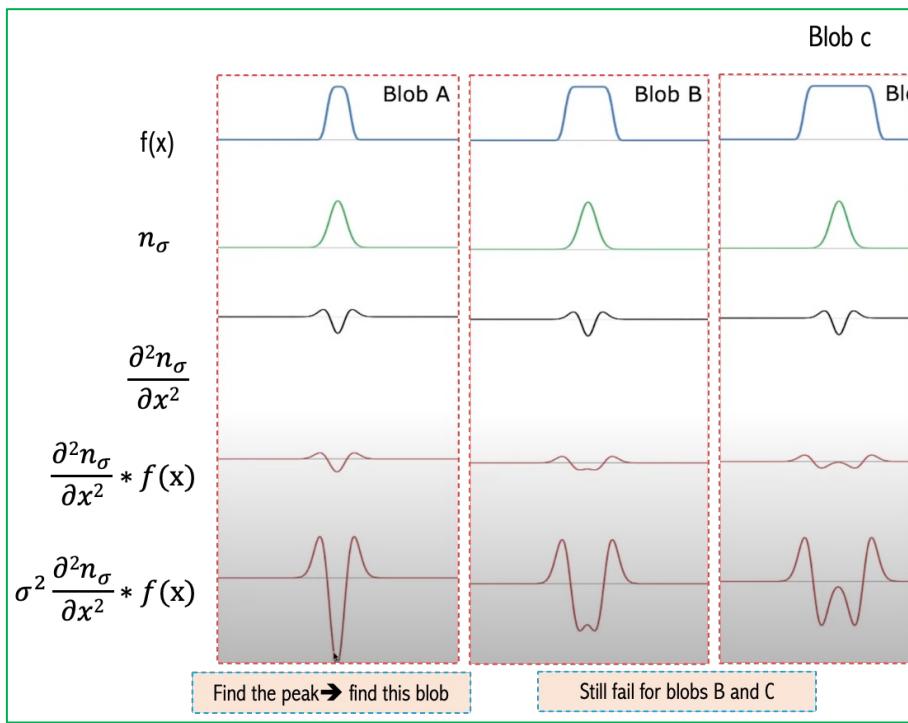
How to detect Blobs



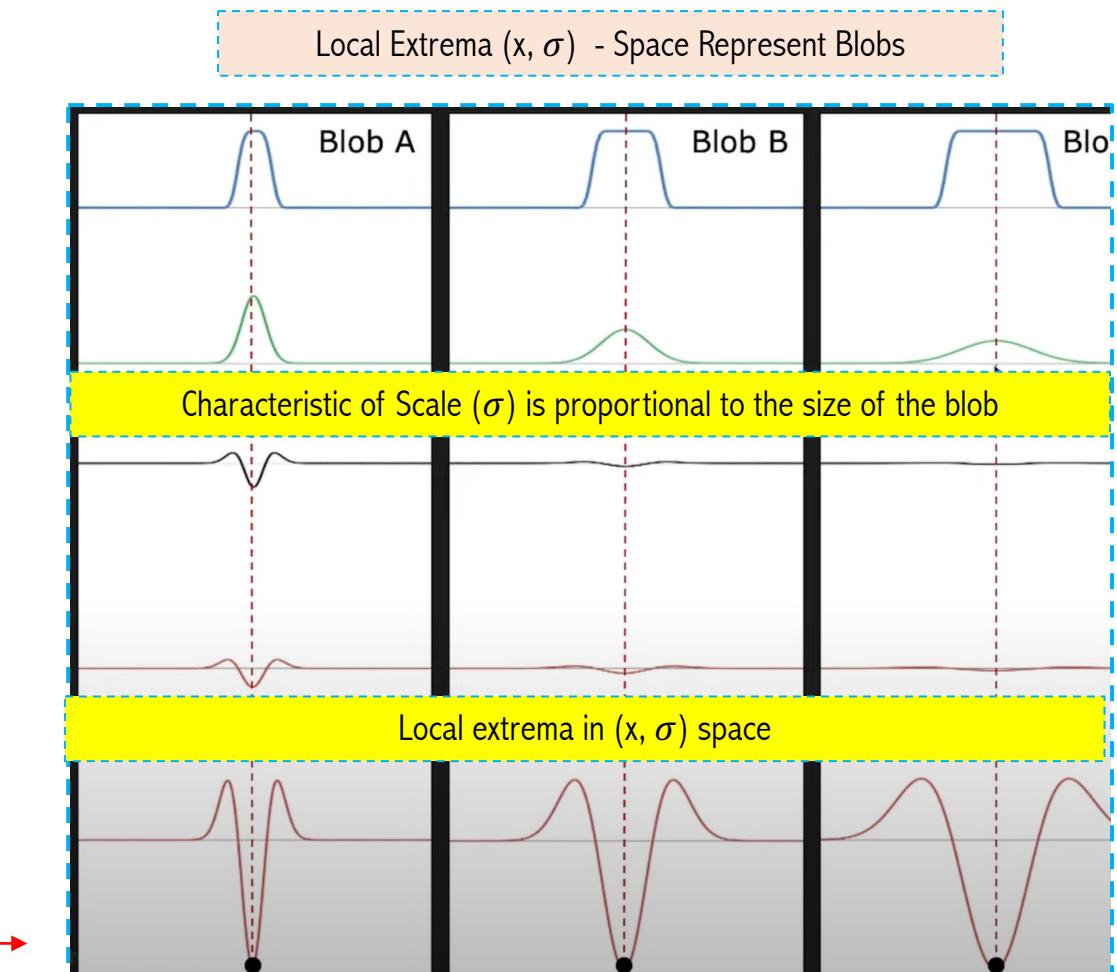
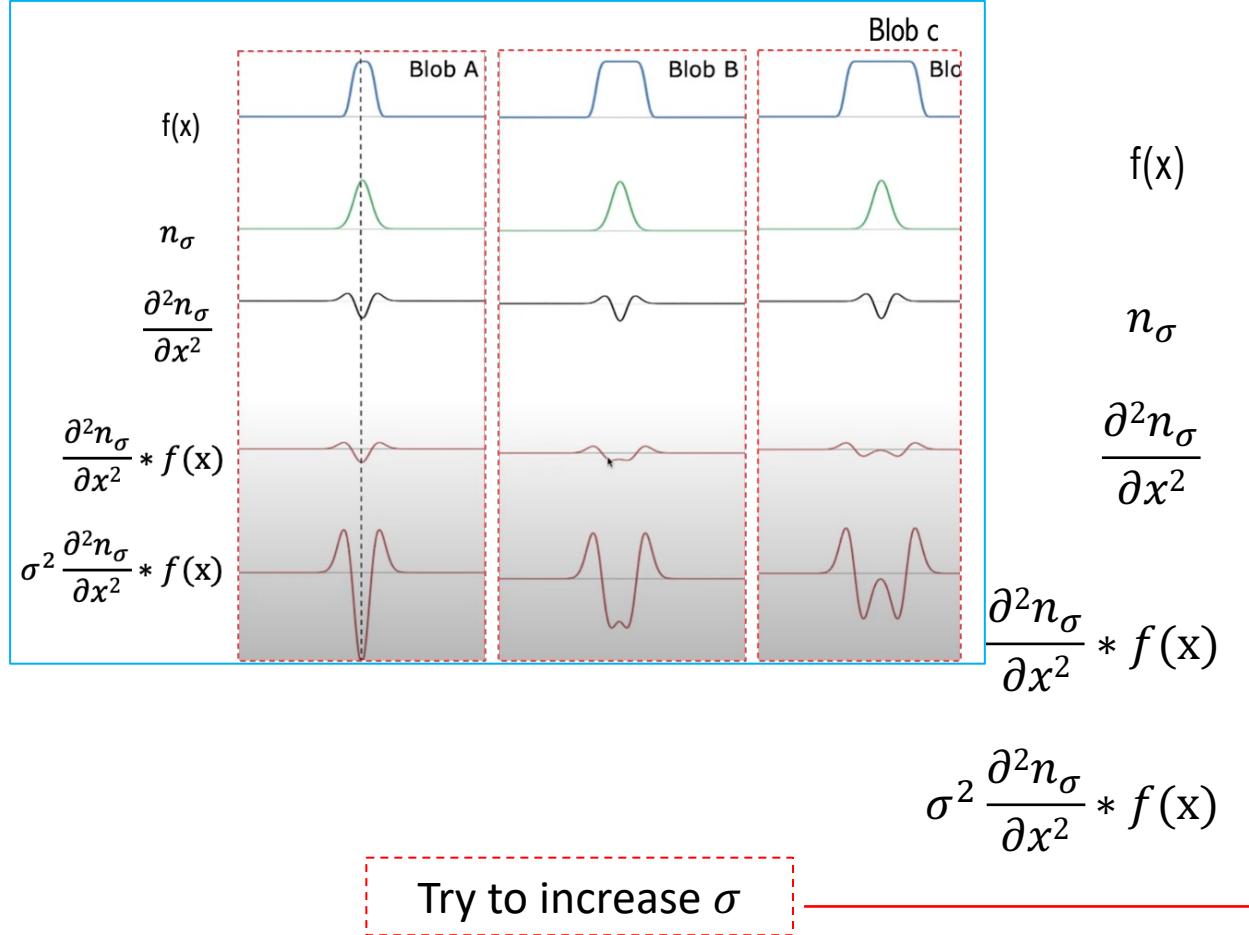
How to detect Blobs



How to detect Blobs



How to detect Blobs



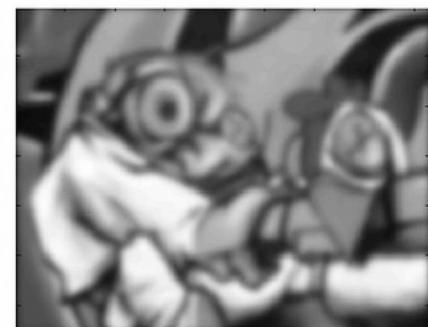
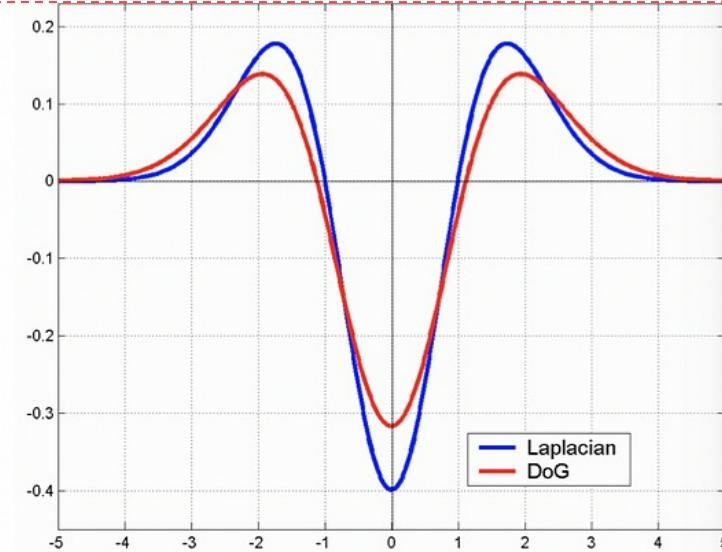
Comparison of Laplacian-of-Gaussian and Difference-of-Gaussian

$$L = \sigma^2 (G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma))$$

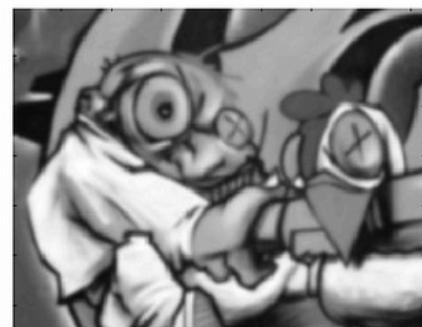
(Laplacian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

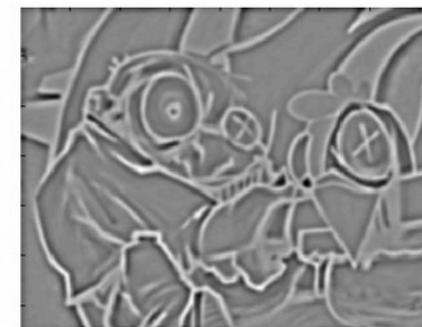
(Difference of Gaussians)



$G(x, y, k\sigma)$



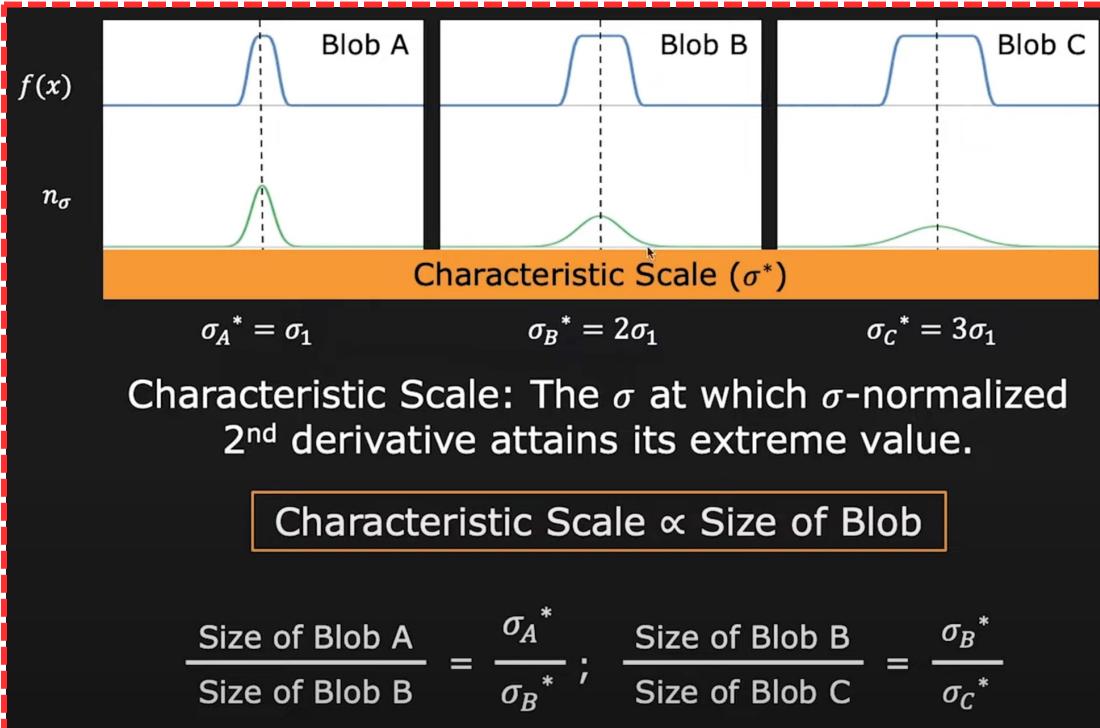
$G(x, y, \sigma)$



$G(x, y, k\sigma) - G(x, y, \sigma)$

$\approx LoG$

How to detect 1D Blobs



Given: 1D signal $f(x)$

Compute: $\sigma^2 \frac{\partial^2 n_\sigma}{\partial x^2} * f(x)$ at many scales $(\sigma_0, \sigma_1, \sigma_2, \dots, \sigma_k)$.

Find:
$$(x^*, \sigma^*) = \arg \max_{(x, \sigma)} \left| \sigma^2 \frac{\partial^2 n_\sigma}{\partial x^2} * f(x) \right|$$

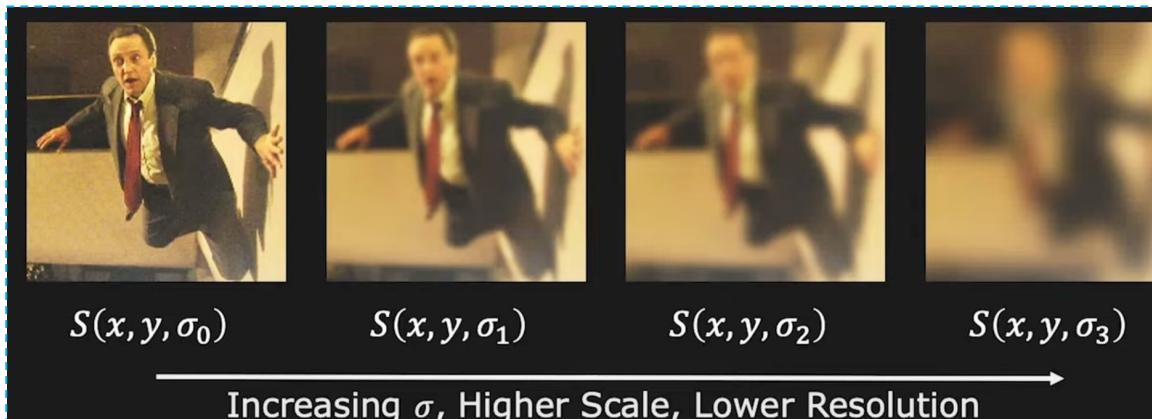
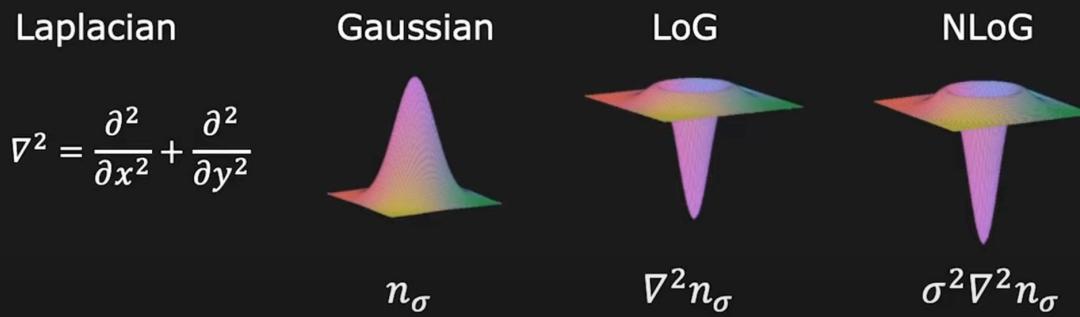
x^* : Blob Position

σ^* : Characteristic Scale (Blob Size)

From Prof. Shree Nayar Lecture

How to detect 2D Blobs: Scale-Space

Normalized Laplacian of Gaussian (NLoG) is used as the 2D equivalent for Blob Detection.



Selecting sigmas to generate the scale-space:

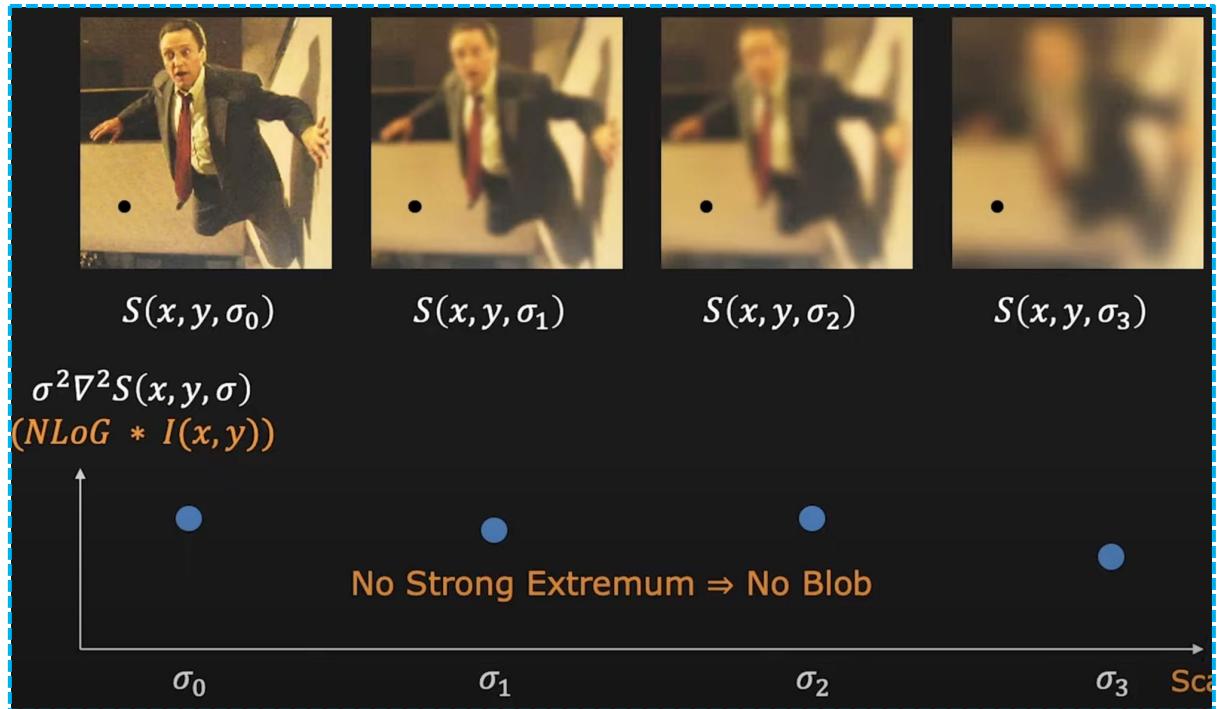
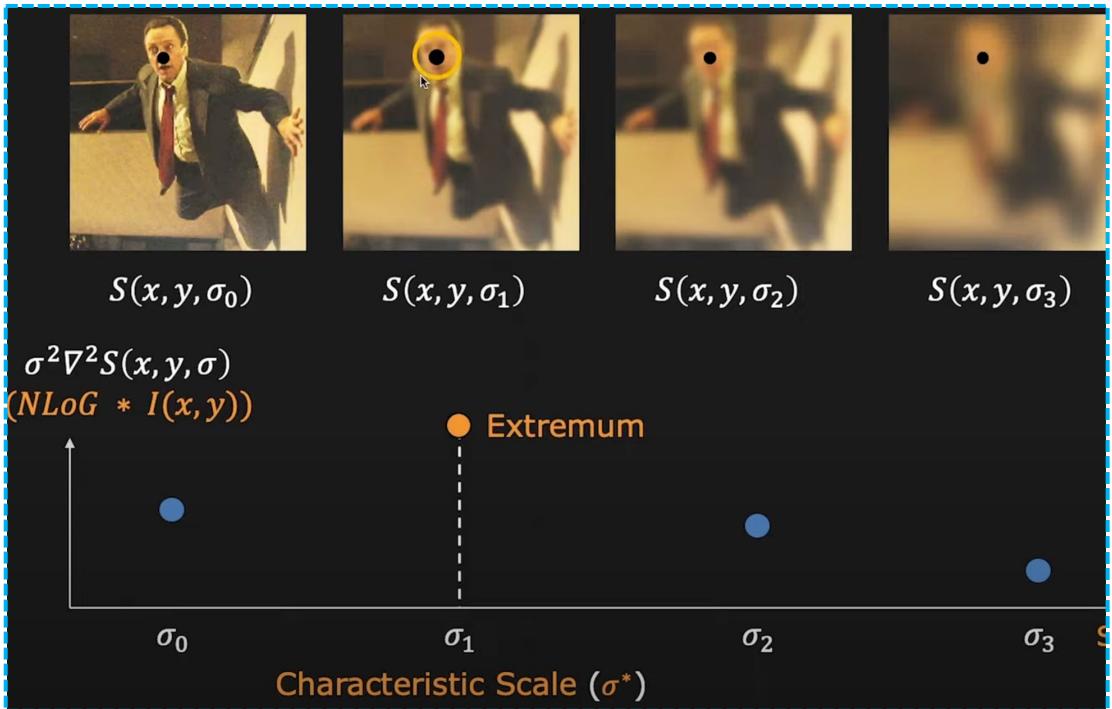
$$\sigma_k = \sigma_0 s^k \quad k = 0, 1, 2, 3, \dots$$

s : Constant multiplier

σ_0 : Initial Scale

From Prof. Shree Nayar Lecture

How to detect 2D Blobs



2D Blob Detection: Summary

Given an image $I(x, y)$

Convolve the image using NLoG at many scales σ

Find:

$$(x^*, y^*, \sigma^*) = \arg \max_{(x,y,\sigma)} |\sigma^2 \nabla^2 n_\sigma * I(x, y)|$$

(x^*, y^*) : Position of the blob

σ^* : Size of the blob

Outline

- What is image stitching/panorama image
- Edge Detector
- Blob Detector
- SIFT detector
- Image Transformation: 2D & 3D
- Image stitching/panorama Techniques

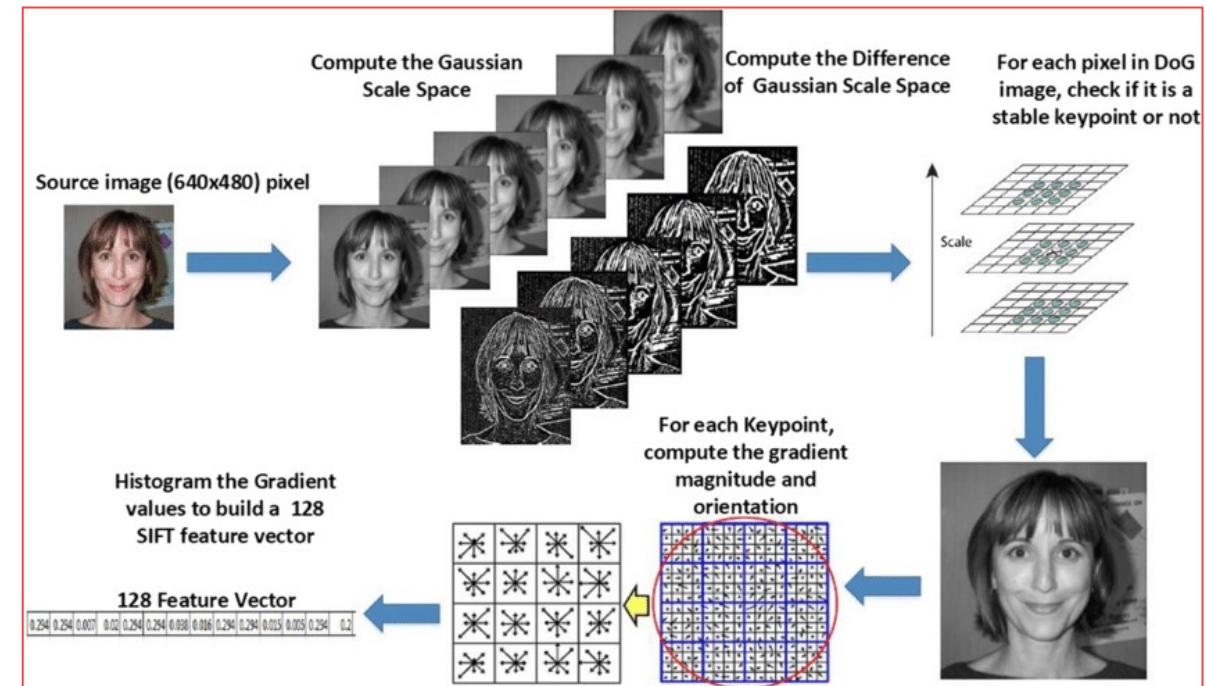
Bird's-eye view of SIFT

Scale-space extrema detection

Keypoint localization and filtering

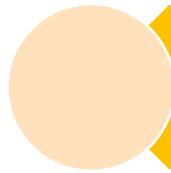
Orientation assignment

Generation of vector representation for interest points



https://www.researchgate.net/publication/331185020_A_novel_SIFT_architecture_and ASIC_implementation_for_real_time_SOC_application/figures?lo=1

Bird's-eye view of SIFT



Scale-space extrema detection



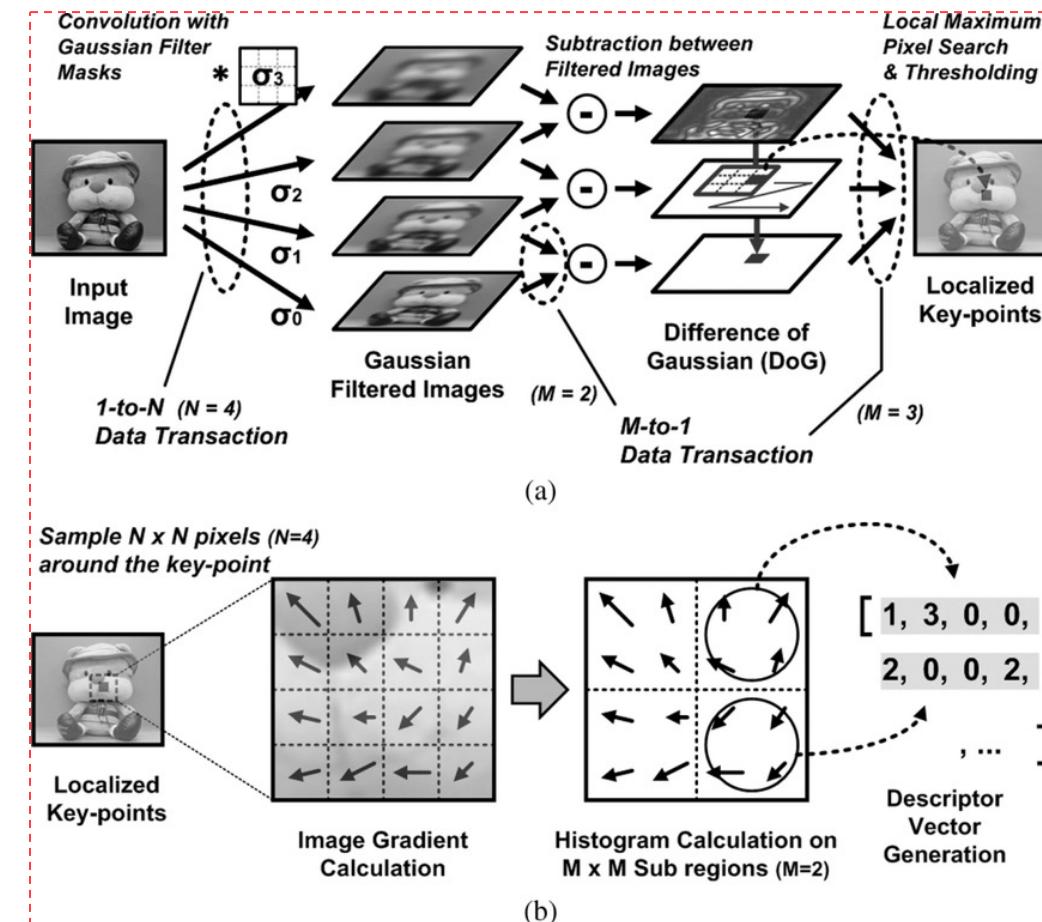
Keypoint localization and filtering



Orientation assignment

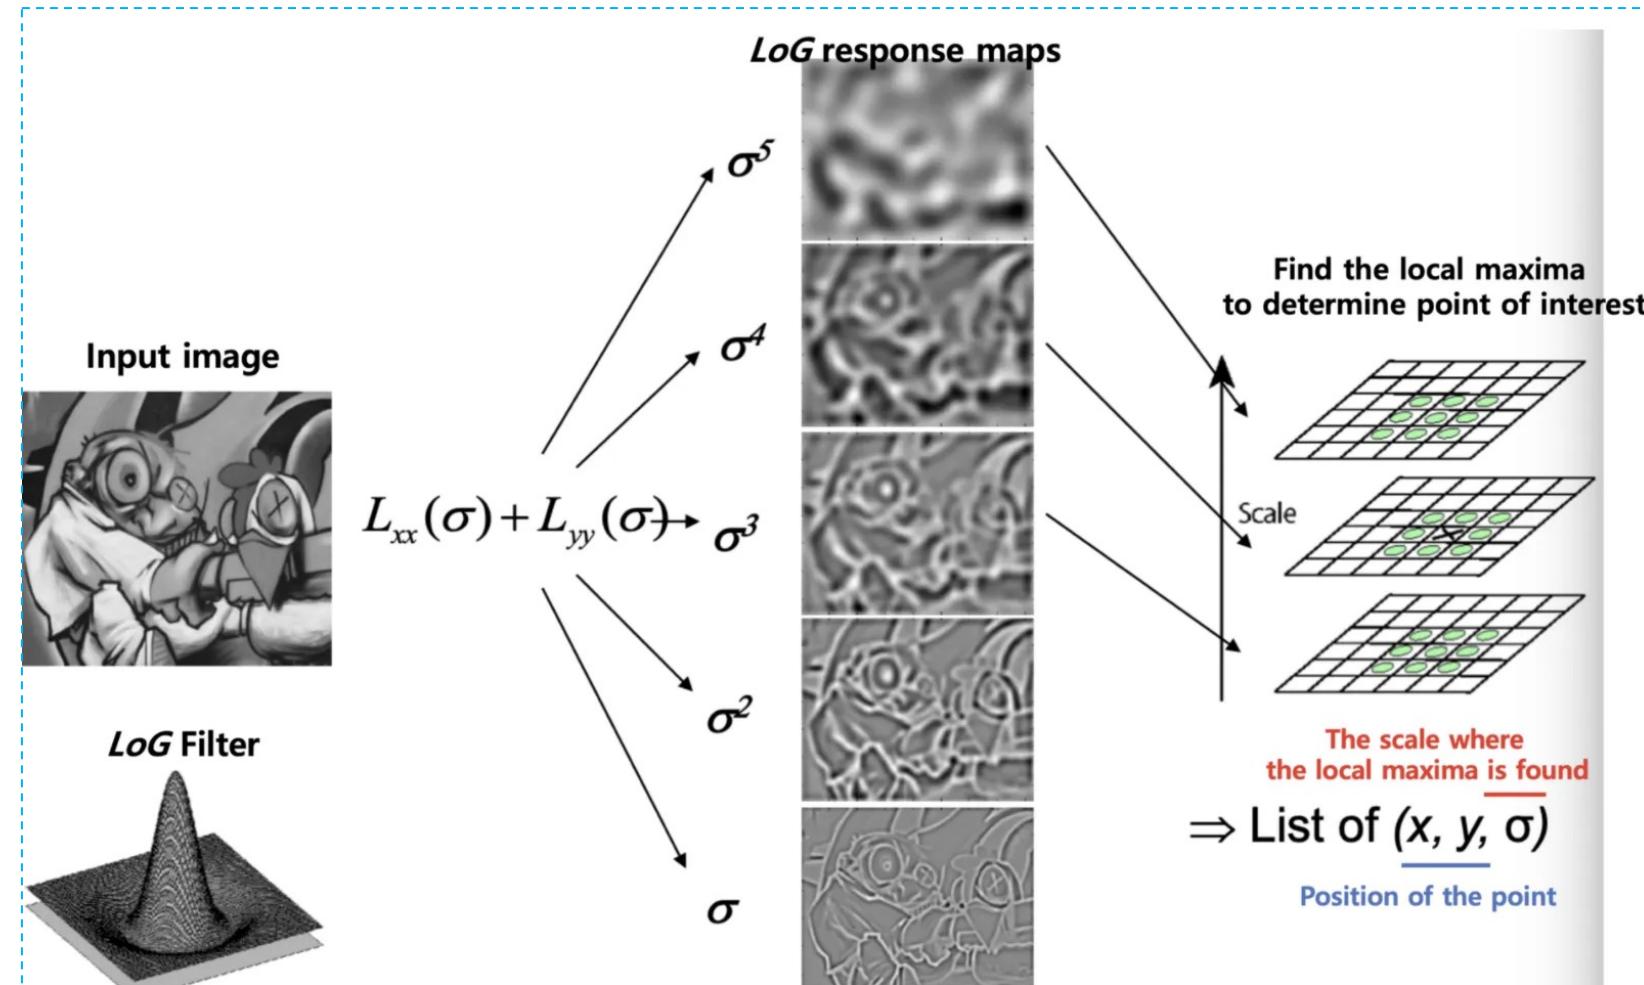


Generation of vector representation for interest points



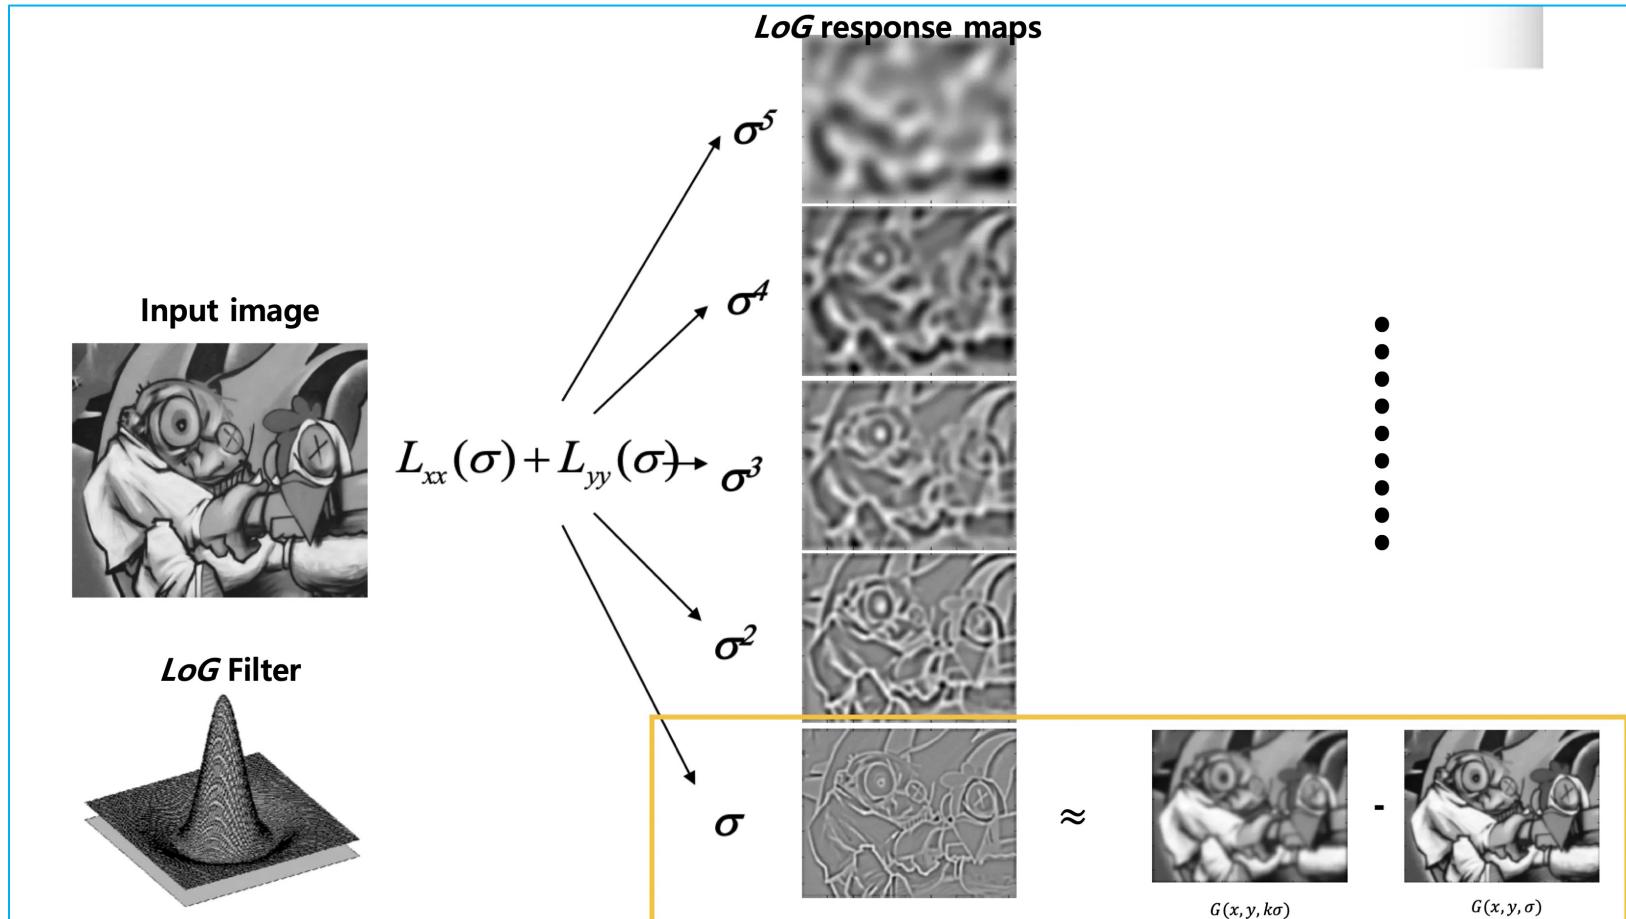
https://www.researchgate.net/publication/331185020_A_novel_SIFT_architecture_and ASIC_implementation_for_real_time_SOC_application/figures?lo=1

Scale-space extrema detection



Automatic scale selection with Laplacian of Gaussian

Scale-space extrema detection

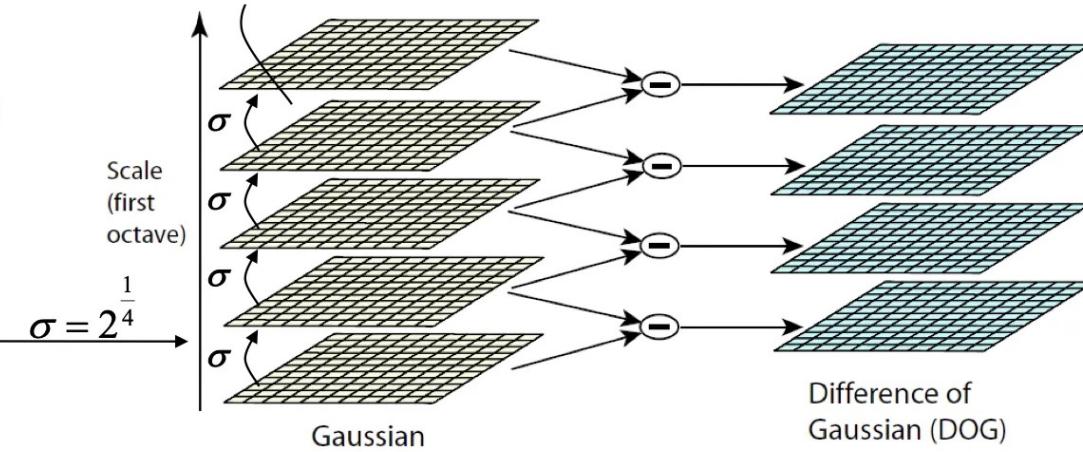


Applying DoG instead of computing LoG response map in each scale space

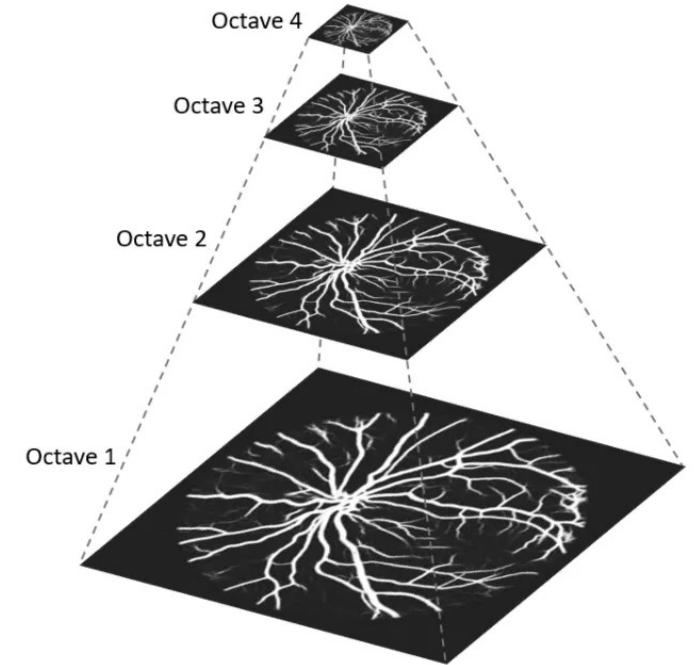
Scale-space extrema detection



Original image

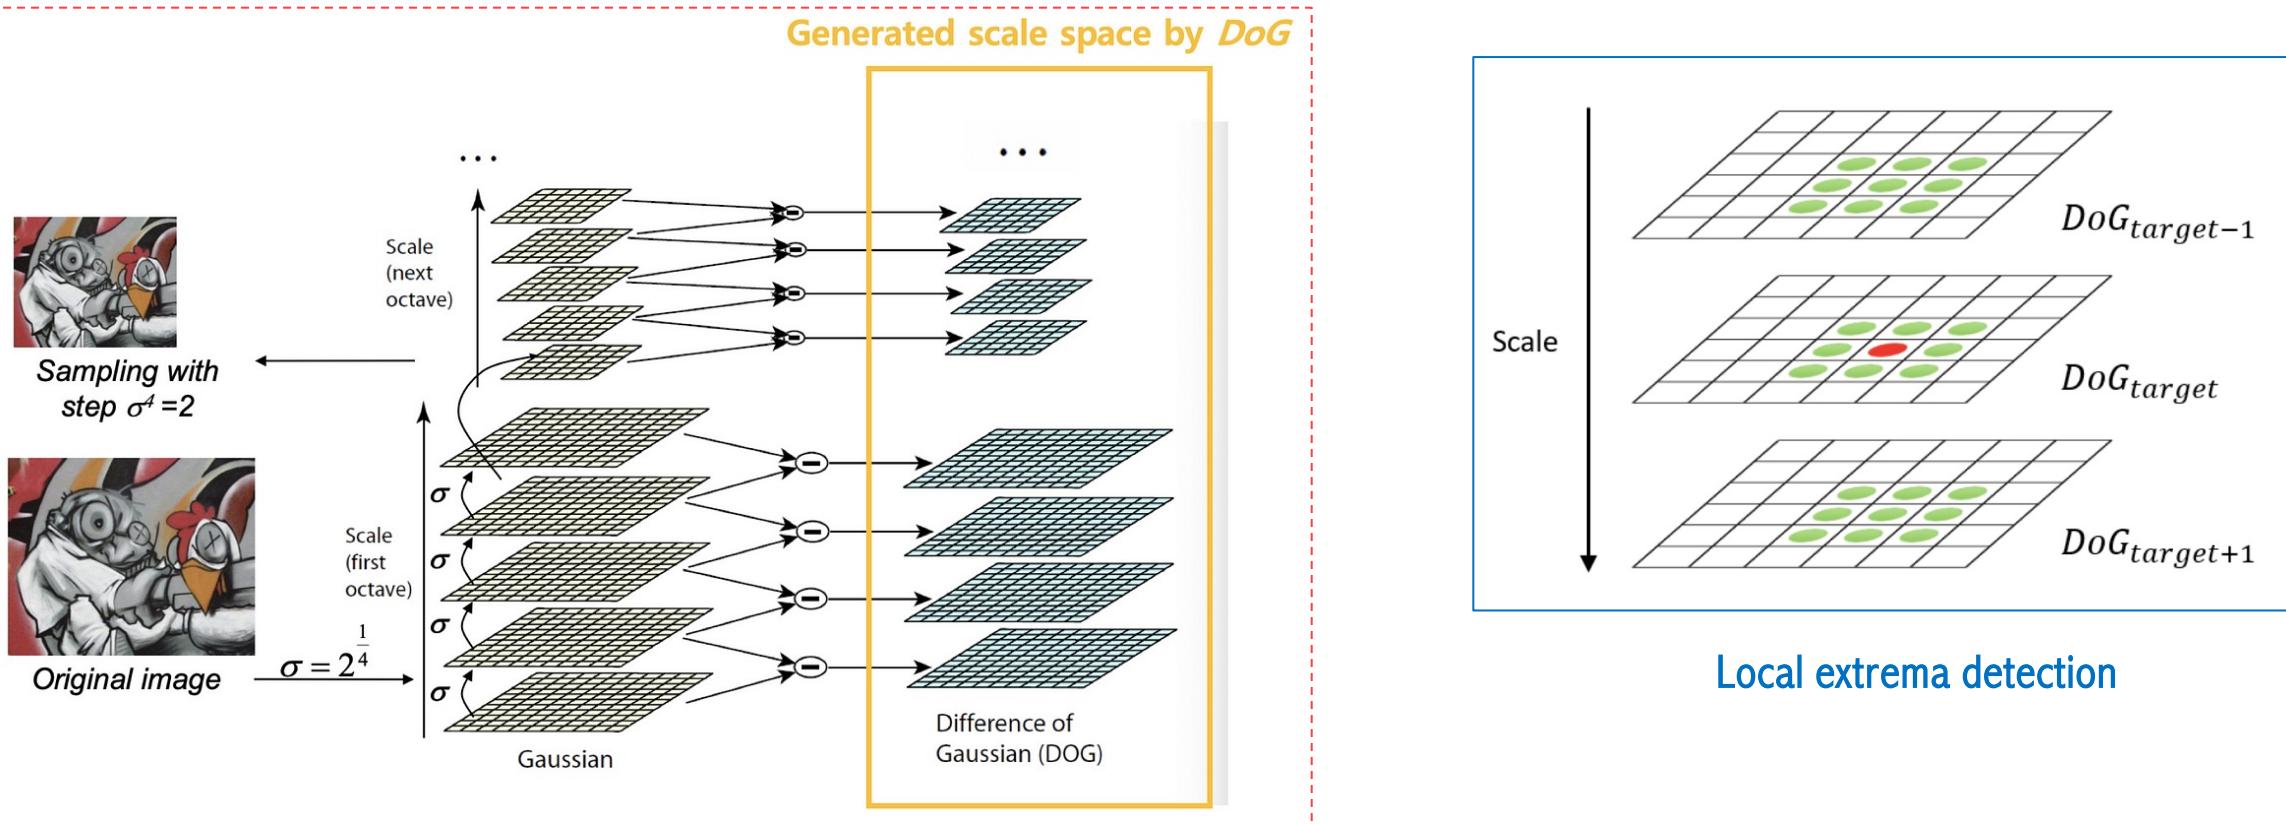


However, just doing this does not guarantee that SIFT is scale-invariant because the algorithm currently only looks at the image in one fixed scale. In order to observe the image from many different scales, SIFT integrates the idea of *Gaussian Pyramid* into this.



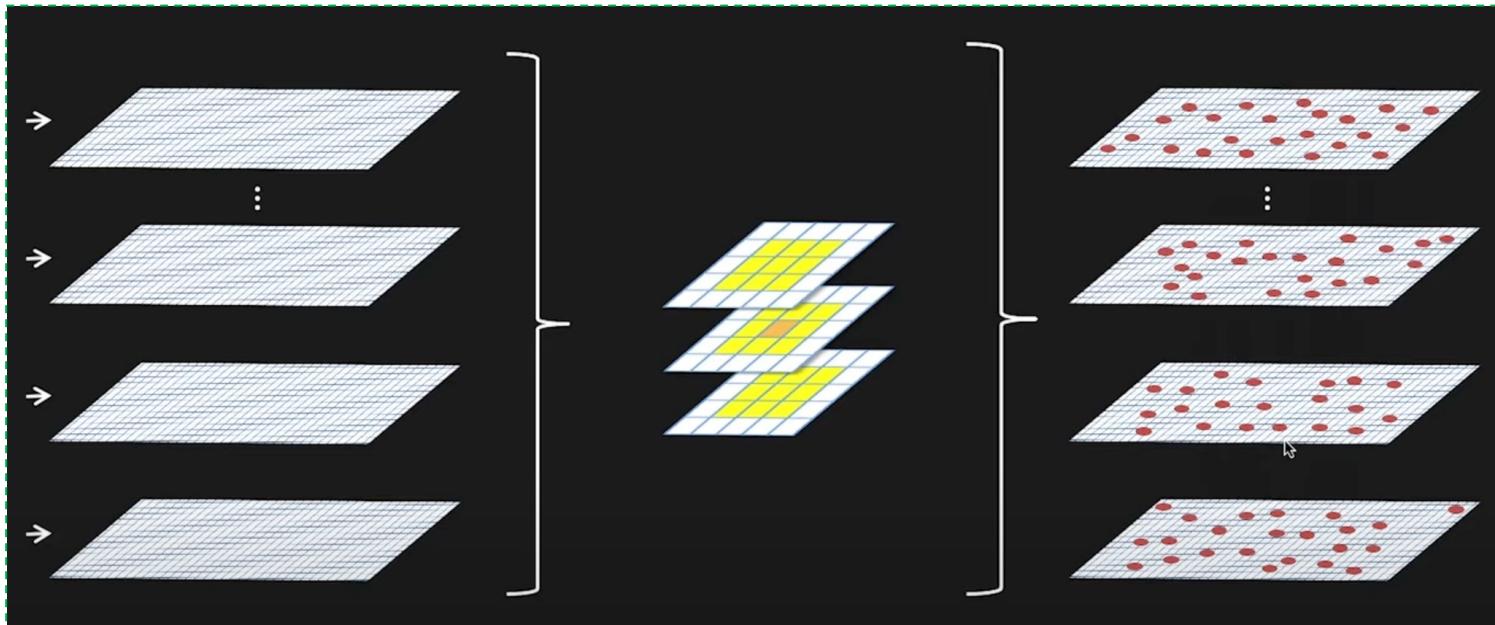
Gaussian pyramid is a multi-scale representation of an image containing a properly re-sized(re-scaled) version of the original image by sampling with Gaussian blurring in order to avoid the aliasing issue

Scale-space extrema detection



Local extrema (either maxima or minima) are detected by comparing a pixel (*red circle*) to its 26 neighbors (*green*) in 3×3 window at the current and adjacent (above and below) scales. More specifically, only a point $DoG(x, y, \sigma)$ is selected if it is larger or smaller than all of its neighbors. The list of detected extremum (x, y, σ) is the **keypoint candidates** which will be processed later on in *SIFT* algorithm.

Local Non-Maximal Suppression



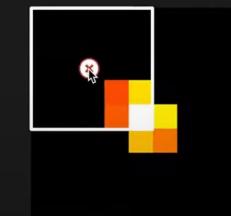
Differences of Gaussian

Find extremum in very
3x3x3 grid

Interest Point Candidates
(includes weak extrema)

Non-maximal suppression

1. Slide a window of size k over the image.
2. At each position, if the pixel at the center is the maximum value within the window, label it as positive (retain it). Else label it as negative (suppress it).



Suppress



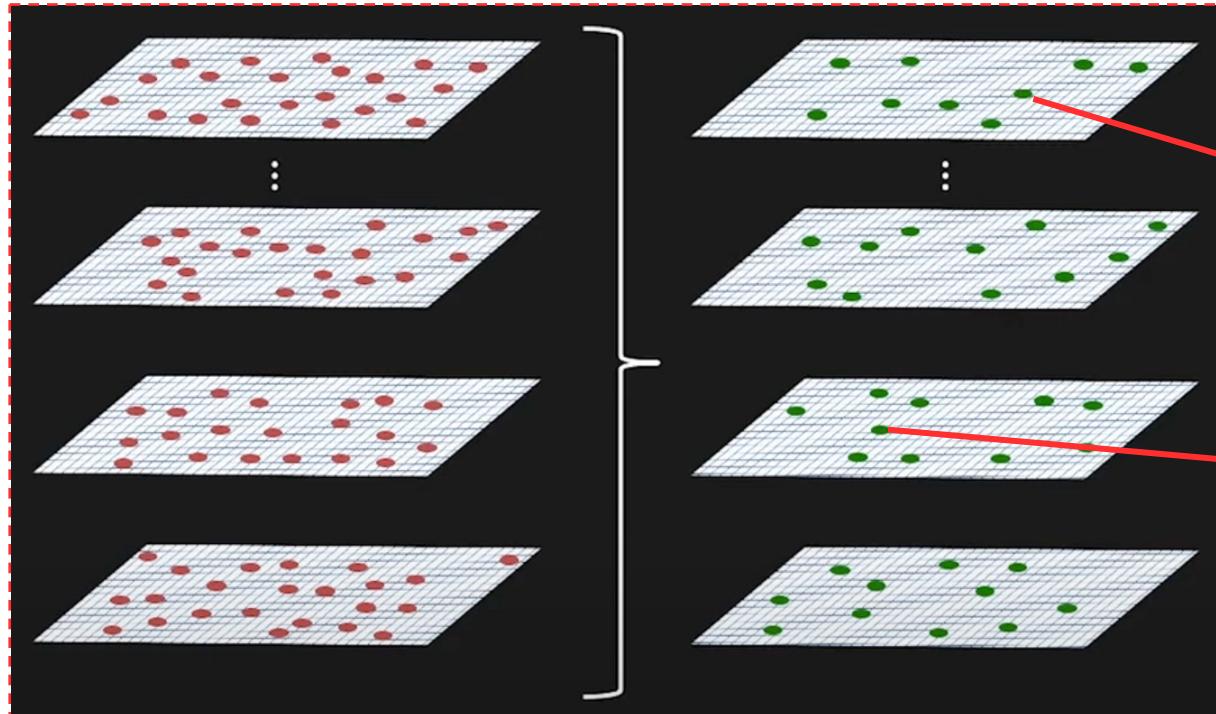
Suppress



Retain

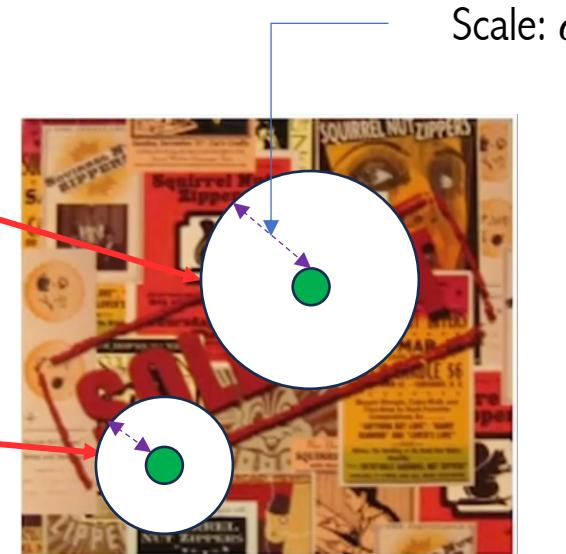
Used for finding Local Extrema (Maxima/Minima)

Keypoint localization and filtering

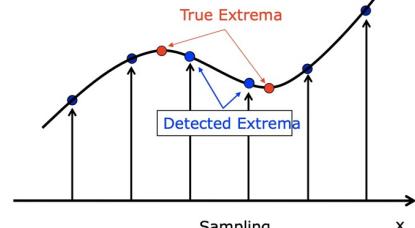


Interest Point Candidates
(includes weak extrema)

Interest Point Candidates
(remove weak extrema)

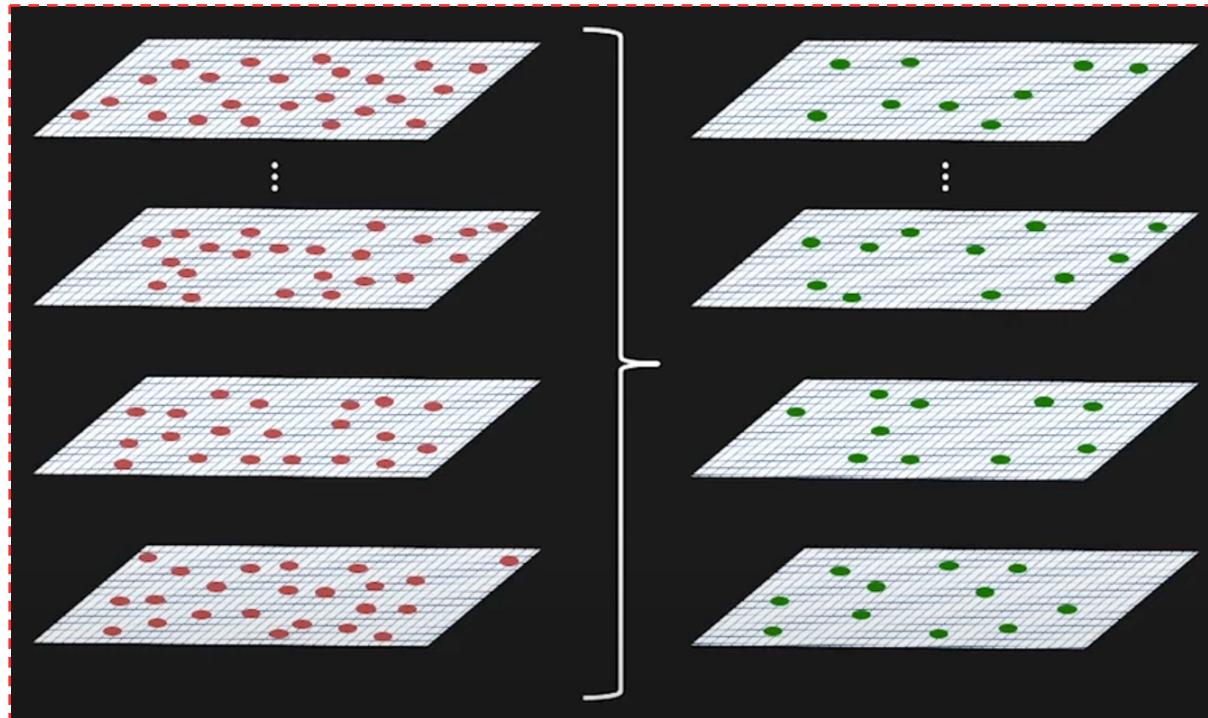


■ The problem:



Taylor series expansion of scale space to get a more accurate location of extrema, and if the intensity at this extrema is less than a threshold value (0.03 as per the paper), it is rejected

Keypoint localization and filtering



Interest Point Candidates
(includes weak
extrema)

Interest Point Candidates
(remove weak extrema)



Interest Point Visualization

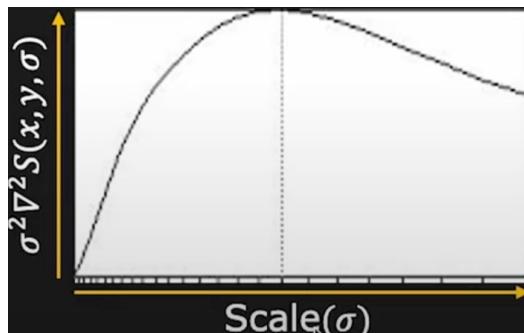
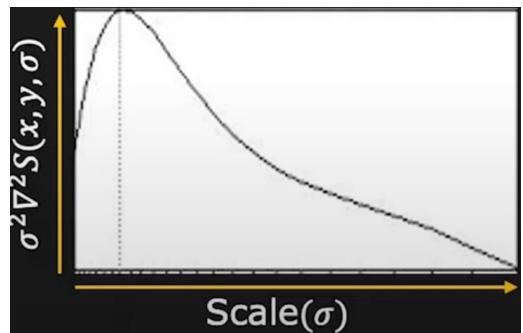
Sift Detector Example



```
import numpy as np
import cv2
import imageio
img = imageio.imread('http://www.ic.unicamp.br\
/~helio/imagens_registro/foto1A.jpg')

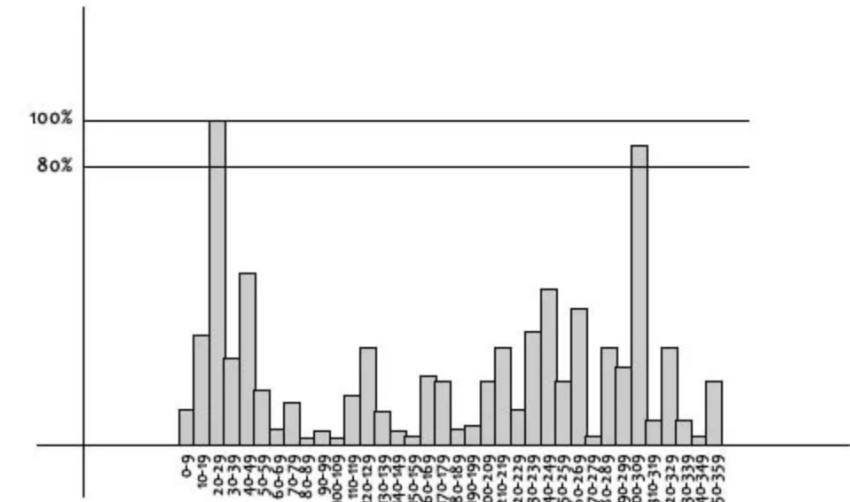
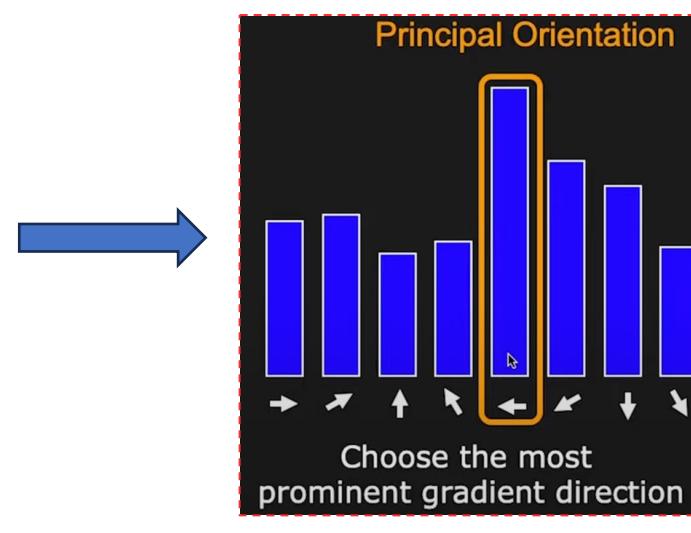
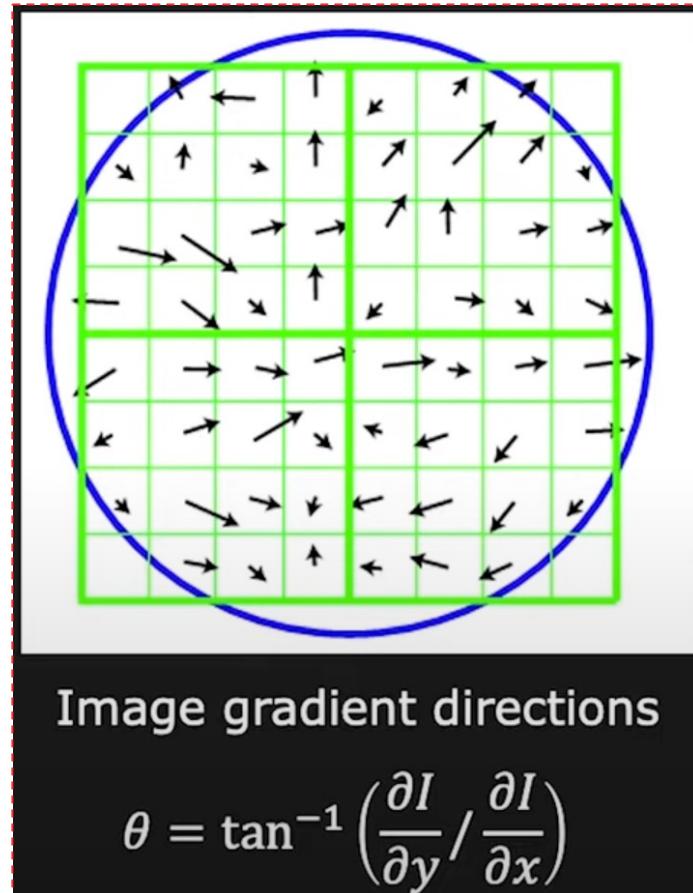
gray= cv2.cvtColor(img,cv2.COLOR_BGR2GRAY)
sift = cv2.SIFT_create()
kp = sift.detect(gray,None)
img=cv2.drawKeypoints(gray,kp,img, flags=cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS)
cv2.imwrite('sift_keypoints.jpg',img)
```

SIFT Scale Invariance



$\frac{\sigma_1^*}{\sigma_2^*}$: Ratio of Blob Sizes

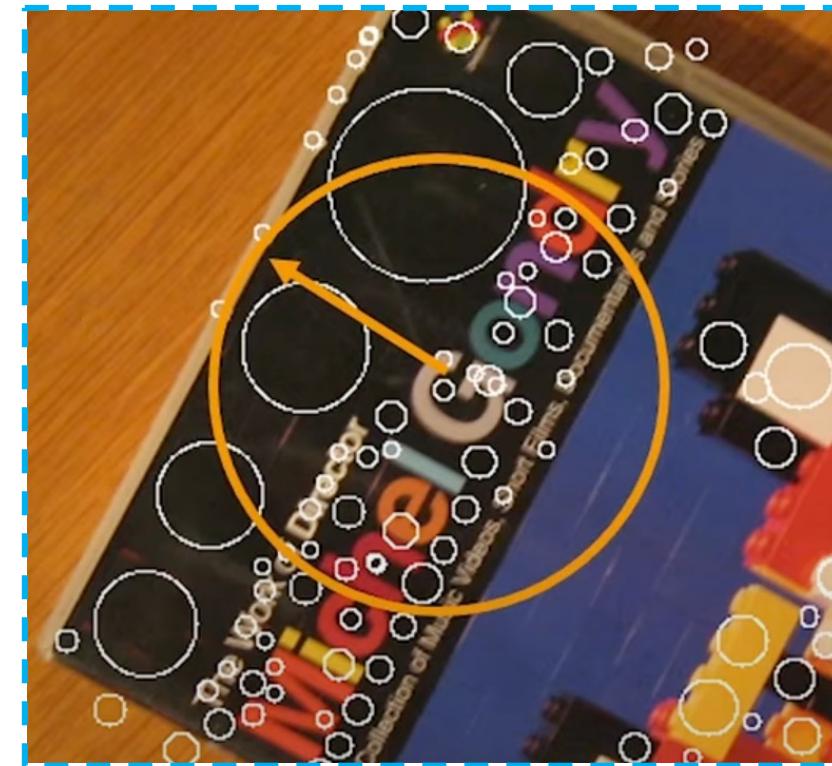
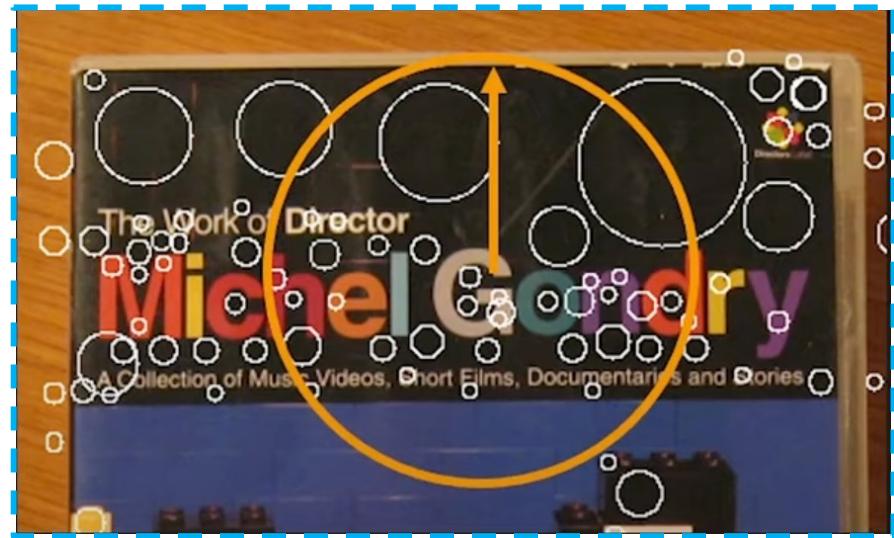
Orientation Assignment



Histogram of oriented gradient

An orientation histogram with 36 bins covering 360 degrees is created

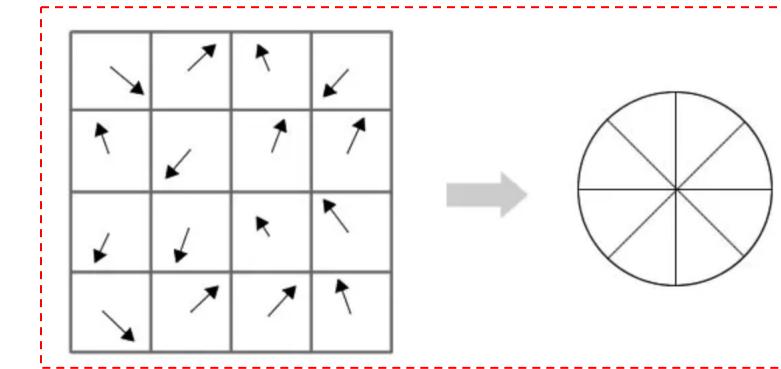
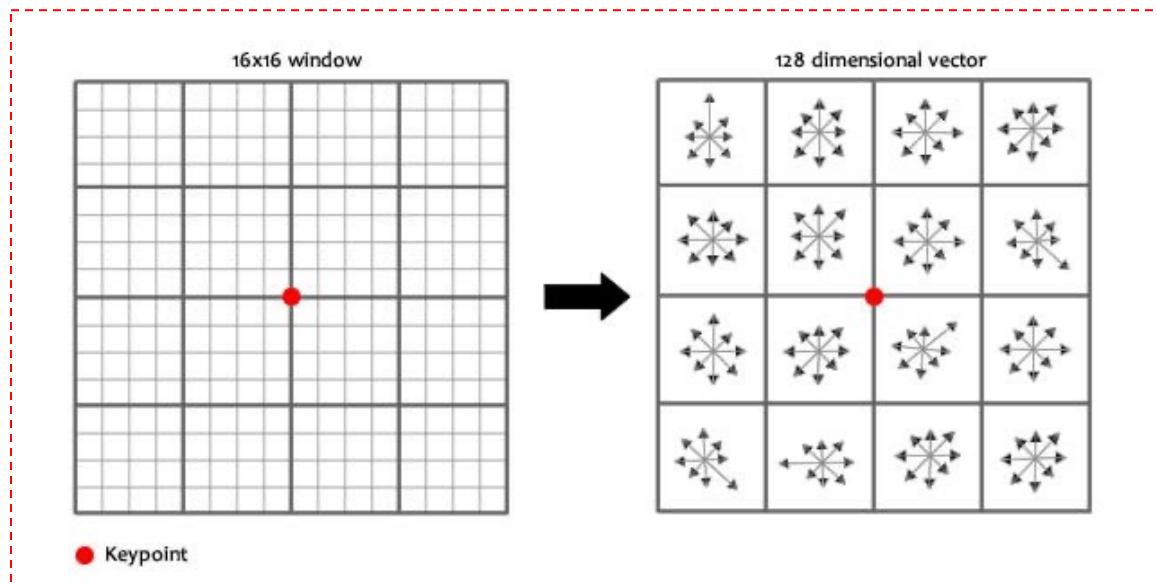
Orientation Assignment



Use the principal orientation to undo the rotation

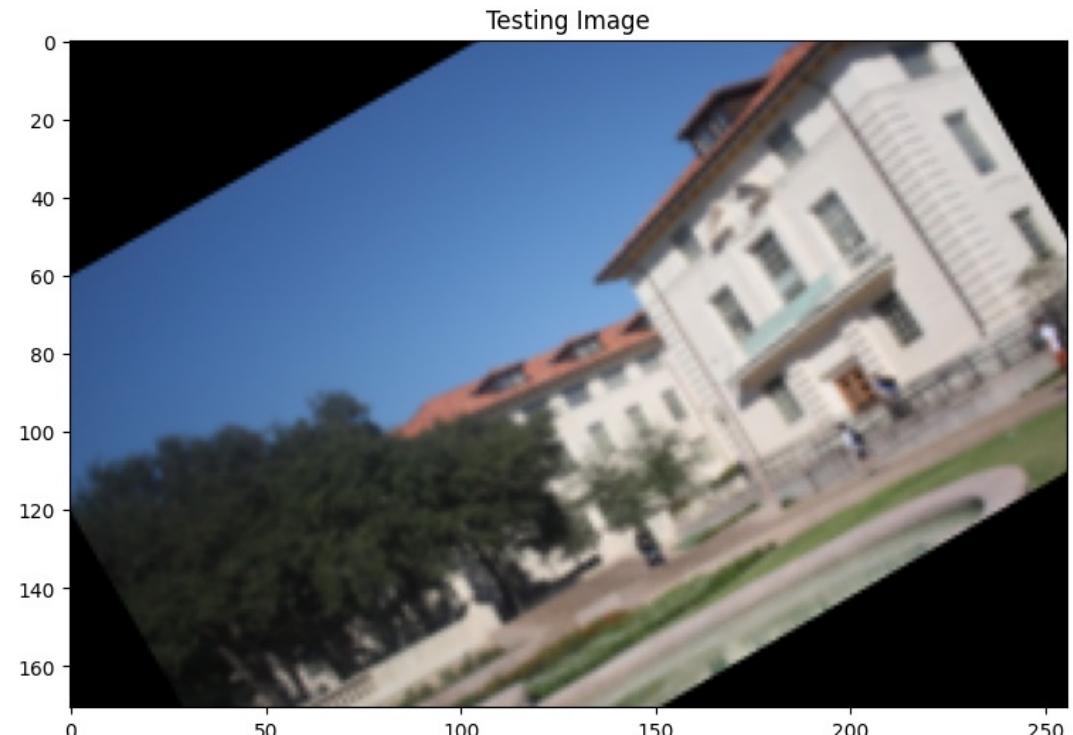
Keypoint descriptor

a 16x16 window around the keypoint is taken. It is divided into 16 sub-blocks of 4x4 size.

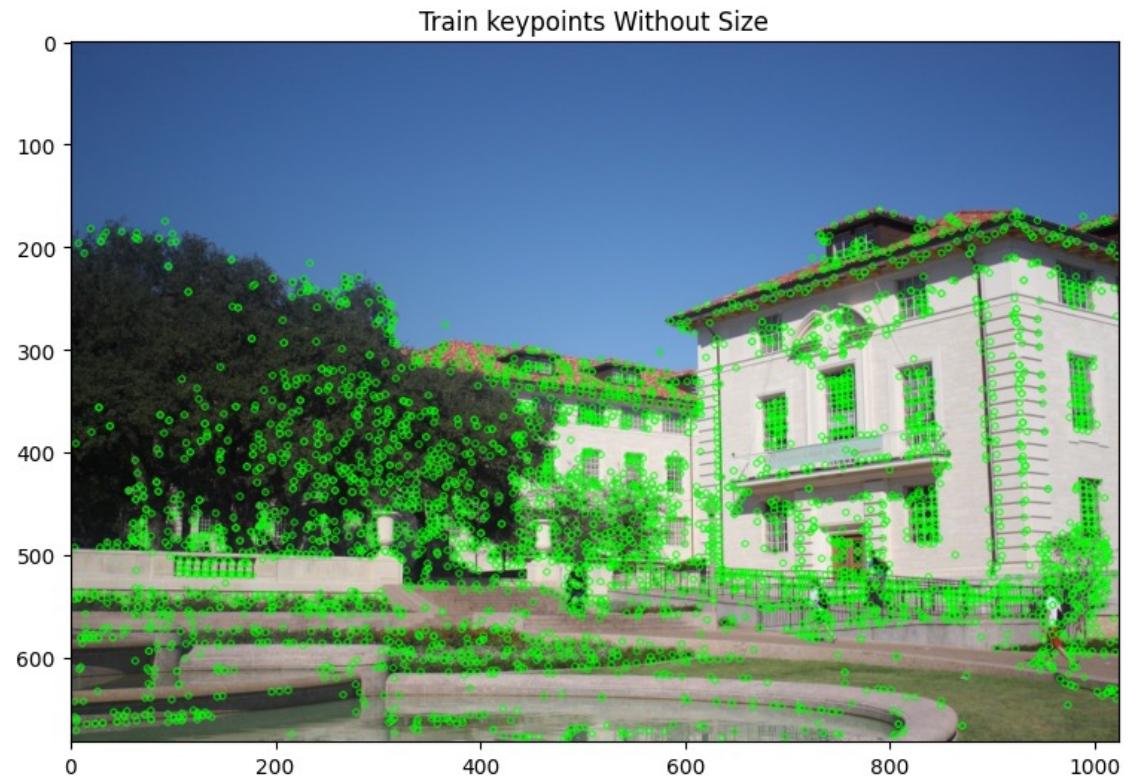
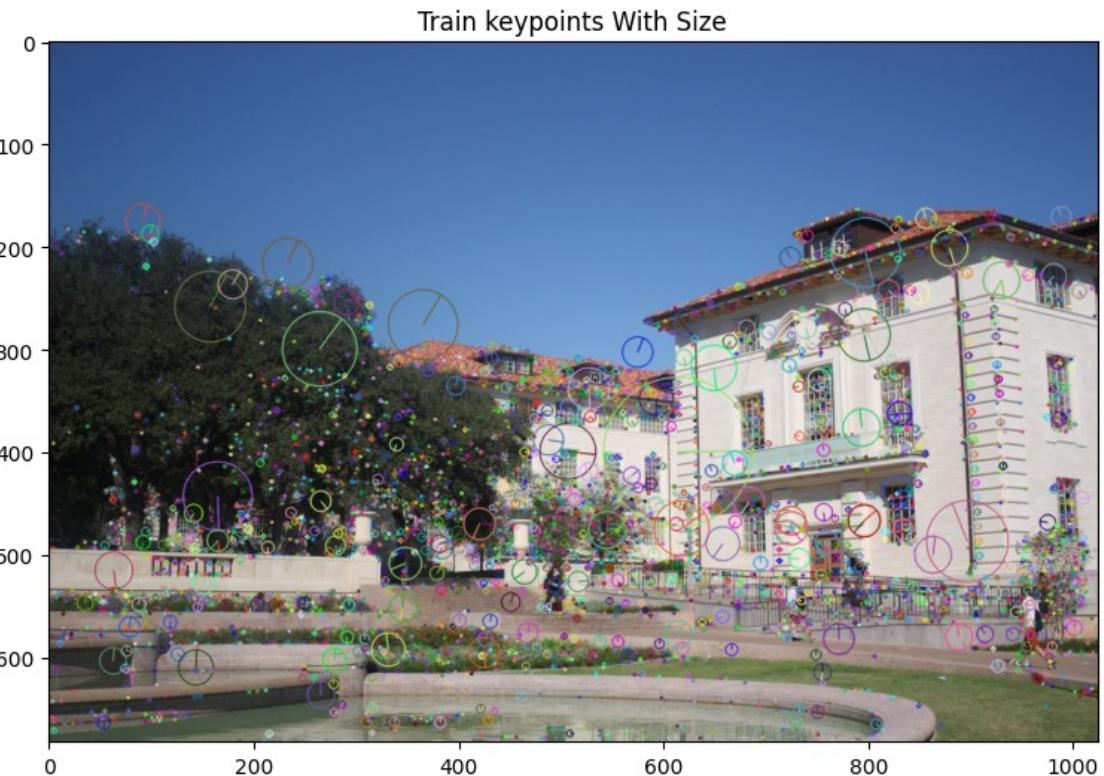


For each sub-block, 8 bin orientation histogram is created.

SIFT Matching



SIFT Matching



SIFT Matching



Let $H_1(k)$ and $H_2(k)$ be two arrays of data of length N .

L2 Distance:

$$d(H_1, H_2) = \sqrt{\sum_k (H_1(k) - H_2(k))^2}$$

Smaller the distance metric, better the match.

Perfect match when $d(H_1, H_2) = 0$

Image Stitching



2x2 Image Transformation

3x3 Image Transformation

Compute Homography

Dealing with Outlier: RANSAC

Warping and Blending Images

Outline

- What is image stitching/panorama image
- Edge Detector
- Blob Detector
- SIFT detector
- Image Transformation: 2D & 3D
- Image stitching/panorama Techniques

Image Manipulation



$$G(x, y) = Tr(f(x, y))$$

Image filtering: change range(brightness)

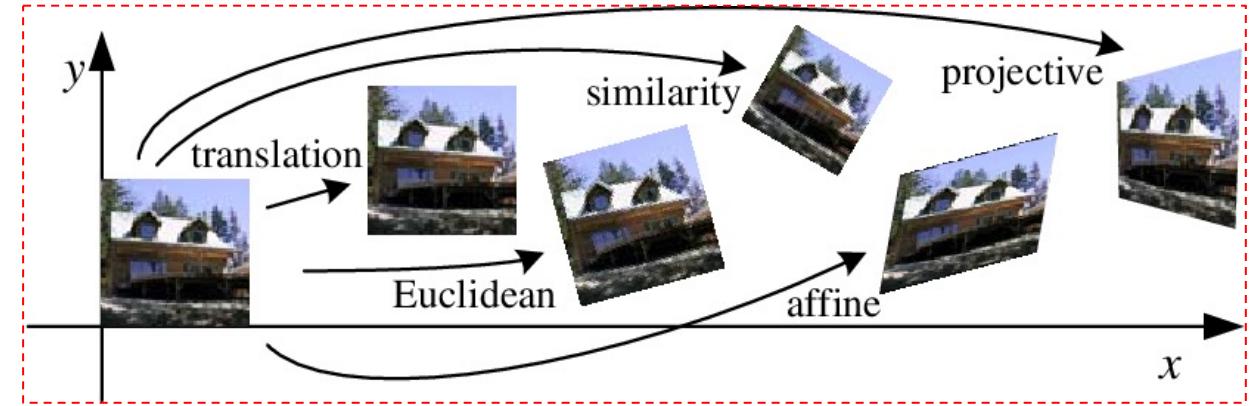
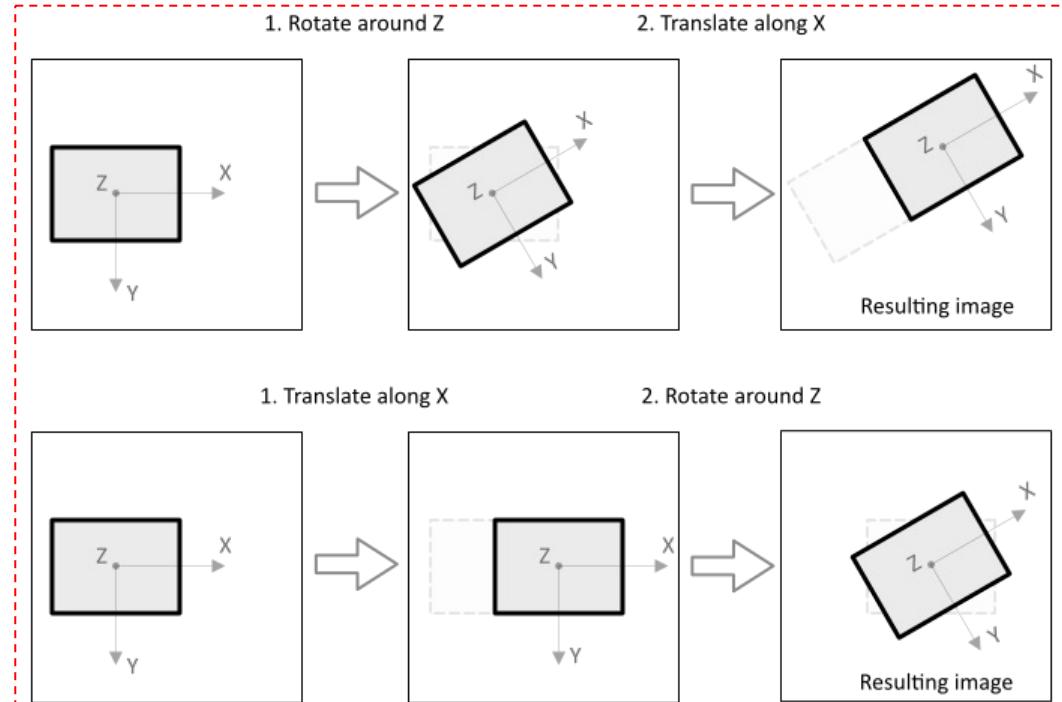


$$G(x, y) = Td(f(x, y))$$

Image filtering: change domain(location)



Global Warping/Transformation

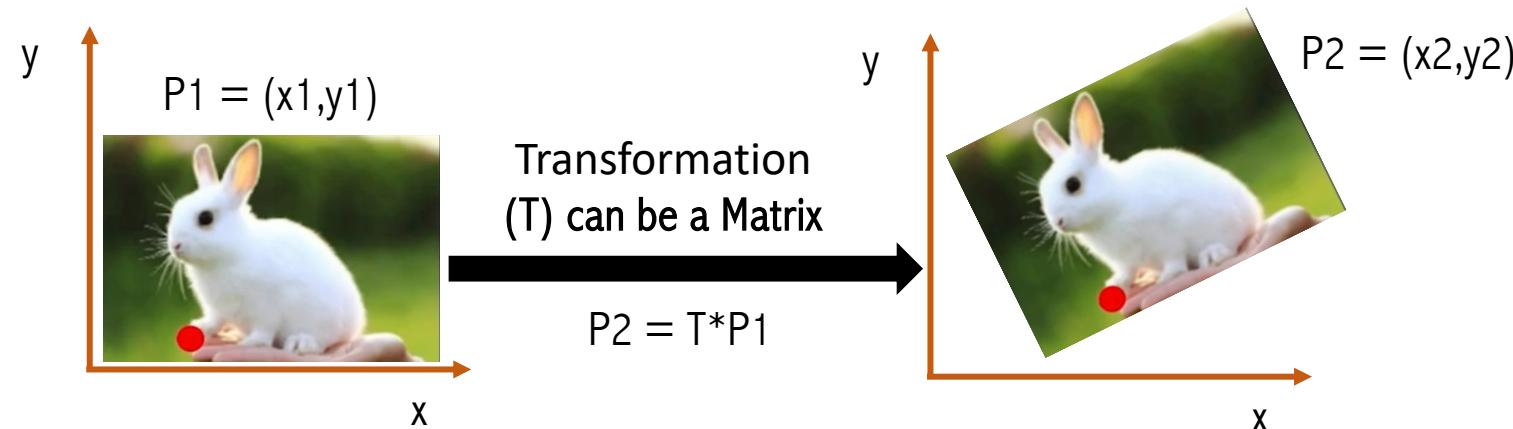


$$\begin{bmatrix} \mathbf{X}_{\text{scaled}} \\ \mathbf{Y}_{\text{scaled}} \\ 1 \end{bmatrix} = \begin{bmatrix} \text{Scale}_x & 0 & 0 \\ 0 & \text{Scale}_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} \mathbf{X} \\ \mathbf{Y} \\ 1 \end{bmatrix}$$

$$\begin{bmatrix} \mathbf{X}_{\text{translated}} \\ \mathbf{Y}_{\text{translated}} \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & D_x \\ 0 & 1 & D_y \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} \mathbf{X} \\ \mathbf{Y} \\ 1 \end{bmatrix}$$

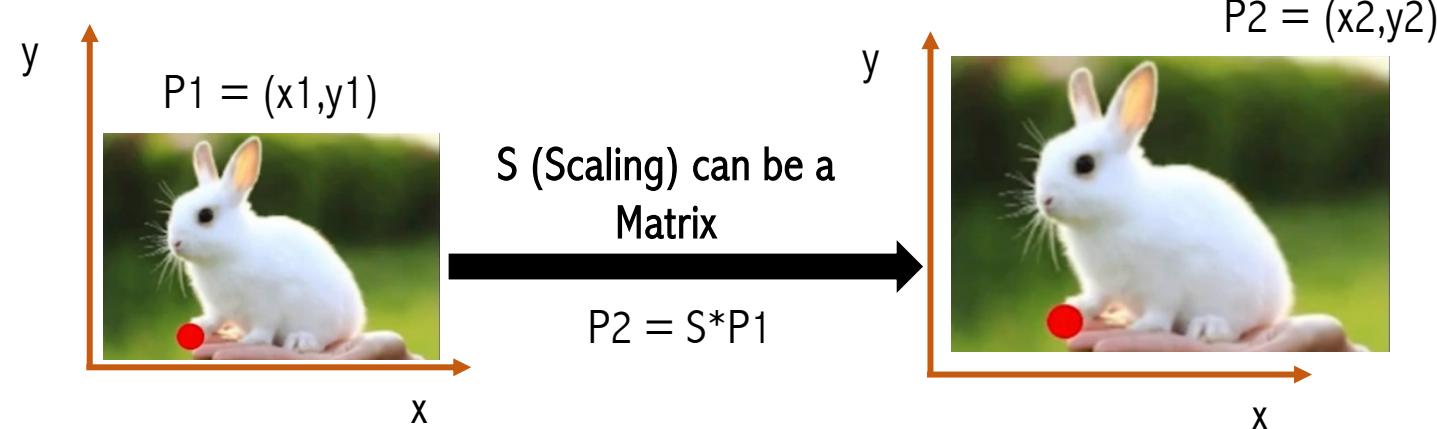
$$\begin{bmatrix} \mathbf{X}_{\text{rotated}} \\ \mathbf{Y}_{\text{rotated}} \\ 1 \end{bmatrix} = \begin{bmatrix} \cos(\theta) & -\sin(\theta) & 0 \\ \sin(\theta) & \cos(\theta) & 0 \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} \mathbf{X} \\ \mathbf{Y} \\ 1 \end{bmatrix}$$

Transformation and Scaling (2D)



$$\begin{bmatrix} x_2 \\ y_2 \end{bmatrix} = T \begin{bmatrix} x_1 \\ y_1 \end{bmatrix}$$

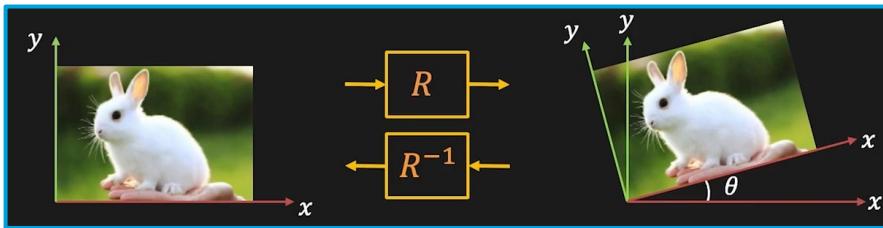
$$\begin{bmatrix} x_2 \\ y_2 \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \end{bmatrix}$$



$$\begin{bmatrix} x_2 \\ y_2 \end{bmatrix} = S \begin{bmatrix} x_1 \\ y_1 \end{bmatrix}$$

$$\begin{bmatrix} x_2 \\ y_2 \end{bmatrix} = \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \end{bmatrix}$$

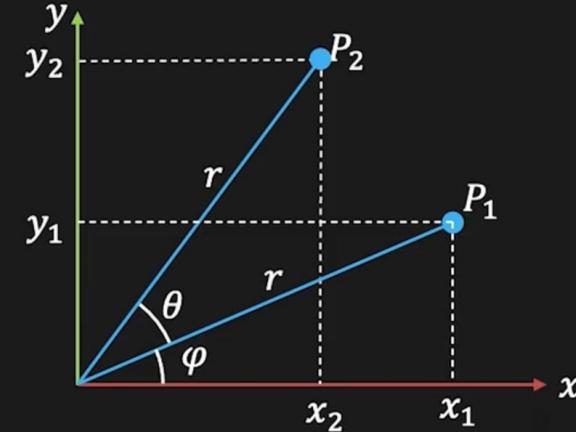
2D Rotation



$$x_2 = x_1 \cos\theta - y_1 \sin\theta$$

$$y_2 = x_1 \sin\theta + y_1 \cos\theta$$

$$\begin{bmatrix} x_2 \\ y_2 \end{bmatrix} = R \begin{bmatrix} x_1 \\ y_1 \end{bmatrix} = \begin{bmatrix} \cos\theta & -\sin\theta \\ \sin\theta & \cos\theta \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \end{bmatrix}$$



$$x_1 = r \cos(\varphi)$$

$$y_1 = r \sin(\varphi)$$

$$x_2 = r \cos(\varphi + \theta)$$

$$x_2 = r \cos \varphi \cos \theta - r \sin \varphi \sin \theta$$

$$x_2 = x_1 \cos \theta - y_1 \sin \theta$$

$$y_2 = r \sin(\varphi + \theta)$$

$$y_2 = r \cos \varphi \sin \theta + r \sin \varphi \cos \theta$$

$$y_2 = x_1 \sin \theta + y_1 \cos \theta$$

Skew and Mirror



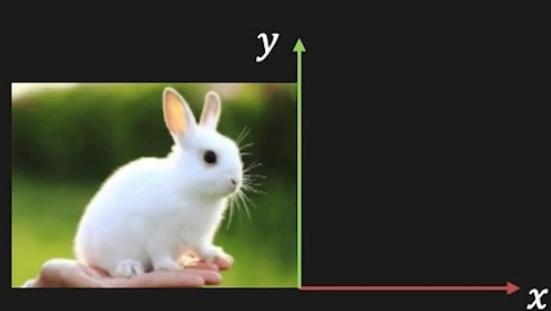
Horizontal Skew:

$$x_2 = x_1 + m_x y_1$$
$$y_2 = y_1$$



Vertical Skew:

$$x_2 = x_1$$
$$y_2 = m_y x_1 + y_1$$



Mirror about Y-axis:

$$x_2 = -x_1$$

$$y_2 = y_1$$

$$M_y = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$$



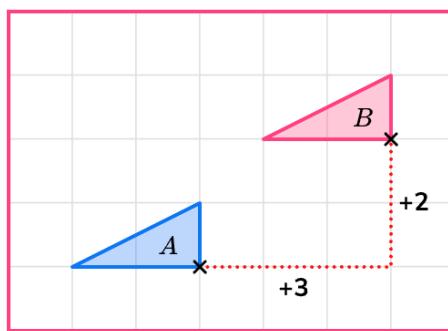
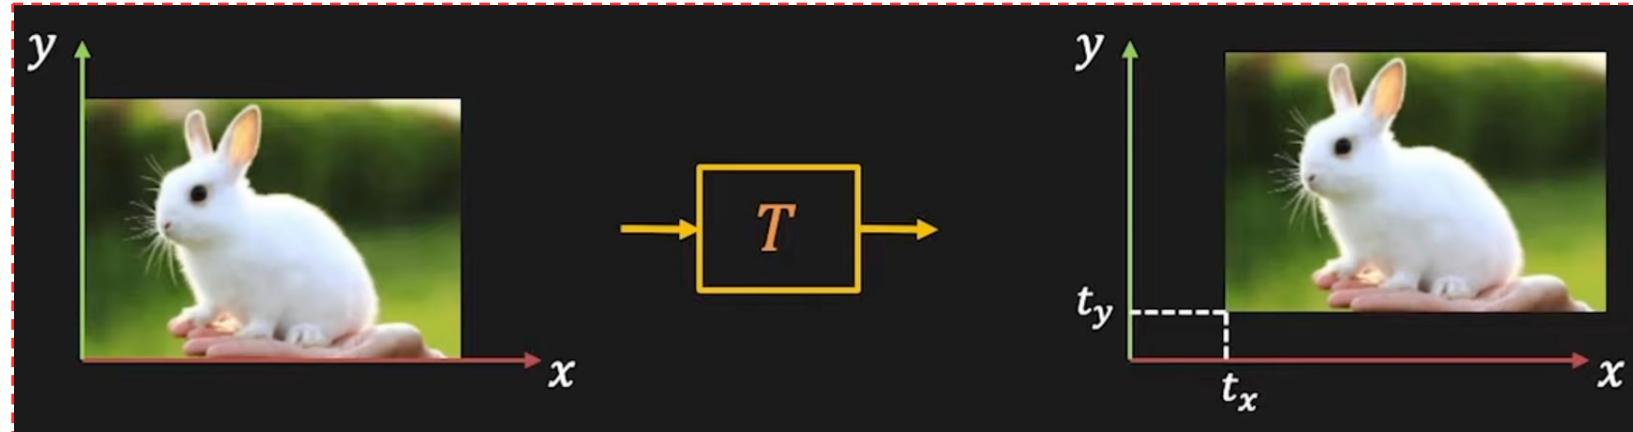
Mirror about line $y = x$:

$$x_2 = y_1$$

$$y_2 = x_1$$

$$M_{xy} = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$$

Translation



$$x_2 = x_1 + t_x$$

$$y_2 = y_1 + t_y$$

Can translation be expressed as 2x2 matrix?



Homogenous Coordinates

$$\begin{bmatrix} 2 \\ 4 \end{bmatrix} \xrightarrow{\text{Euclidean}} \begin{bmatrix} 2 \\ 4 \\ 1 \end{bmatrix} \xrightarrow{\text{homogeneous}}$$

$$\begin{bmatrix} 4 \\ 8 \\ 2 \end{bmatrix} = \begin{bmatrix} 4/2 \\ 8/2 \\ 1 \end{bmatrix} \xrightarrow{\text{homogeneous}} \begin{bmatrix} 4/2 \\ 8/2 \end{bmatrix} = \begin{bmatrix} 2 \\ 4 \end{bmatrix} \xrightarrow{\text{Euclidean}}$$

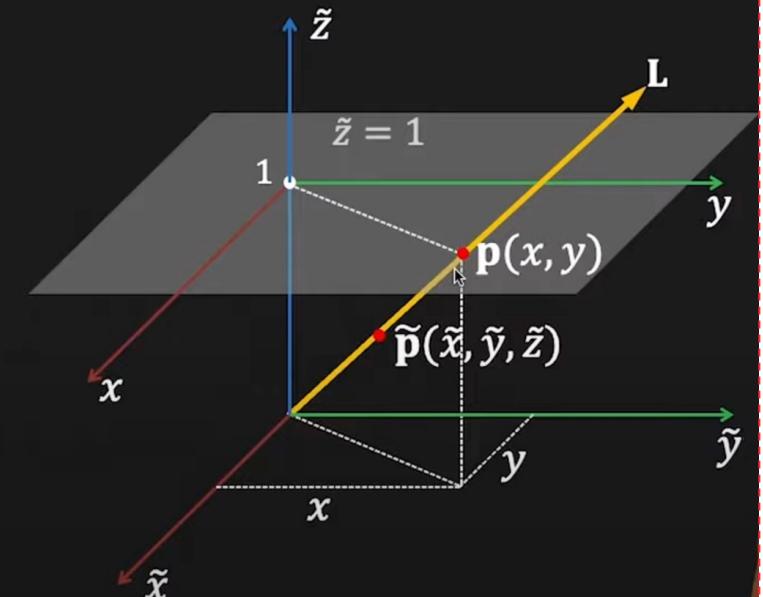
$$\begin{bmatrix} x \\ y \end{bmatrix} \xrightarrow{\text{Euclidean}} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \xrightarrow{\text{homogeneous}}$$

$$\begin{bmatrix} u \\ v \\ w \end{bmatrix} = \begin{bmatrix} u/w \\ v/w \\ 1 \end{bmatrix} \xrightarrow{\text{homogeneous}} \begin{bmatrix} u/w \\ v/w \end{bmatrix} = \begin{bmatrix} x \\ y \end{bmatrix} \xrightarrow{\text{Euclidean}}$$

The homogenous representation of a 2D point $\mathbf{p} = (x, y)$ is a 3D point $\tilde{\mathbf{p}} = (\tilde{x}, \tilde{y}, \tilde{z})$. The third coordinate $\tilde{z} \neq 0$ is fictitious such that:

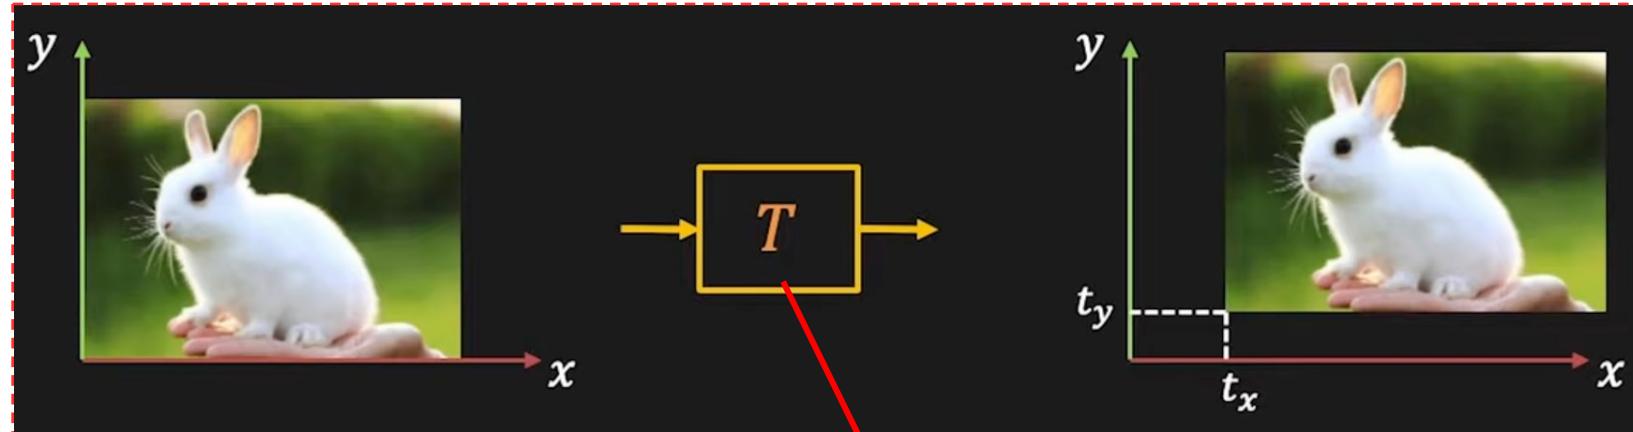
$$x = \frac{\tilde{x}}{\tilde{z}} \quad y = \frac{\tilde{y}}{\tilde{z}}$$

$$\mathbf{p} \equiv \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \equiv \begin{bmatrix} \tilde{z}x \\ \tilde{z}y \\ \tilde{z} \end{bmatrix} \equiv \begin{bmatrix} \tilde{x} \\ \tilde{y} \\ \tilde{z} \end{bmatrix} = \tilde{\mathbf{p}}$$



Every point on line L (except origin) represent the homogeneous coordinate of $p(x,y)$

Translation (3D)



$$x_2 = x_1 + t_x$$

$$y_2 = y_1 + t_y$$

$$\begin{bmatrix} x_2 \\ y_2 \\ 1 \end{bmatrix} \equiv \begin{bmatrix} \tilde{x}_2 \\ \tilde{y}_2 \\ \tilde{z}_2 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix}$$

Scaling, Translation, Rotation, Skew

$$\begin{bmatrix} \tilde{x}_2 \\ \tilde{y}_2 \\ \tilde{z}_2 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix}$$

Scaling

$$\begin{bmatrix} \tilde{x}_2 \\ \tilde{y}_2 \\ \tilde{z}_2 \end{bmatrix} = \begin{bmatrix} 1 & m_x & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix}$$

Skew

$$\begin{bmatrix} \tilde{x}_2 \\ \tilde{y}_2 \\ \tilde{z}_2 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix}$$

Translation

$$\begin{bmatrix} \tilde{x}_2 \\ \tilde{y}_2 \\ \tilde{z}_2 \end{bmatrix} = \begin{bmatrix} \cos\theta & -\sin\theta & 0 \\ \sin\theta & \cos\theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix}$$

Rotation

Composition of these transformation (scaling + skew + translation + rotation)?

Affine Transformation

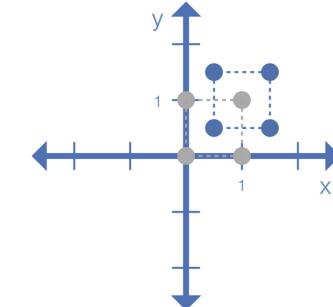
Any transformation of the form:

$$\begin{bmatrix} x_2 \\ y_2 \\ 1 \end{bmatrix} \equiv \begin{bmatrix} \tilde{x}_2 \\ \tilde{y}_2 \\ \tilde{z}_2 \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \tilde{x}_1 \\ \tilde{y}_1 \\ \tilde{z}_1 \end{bmatrix}$$



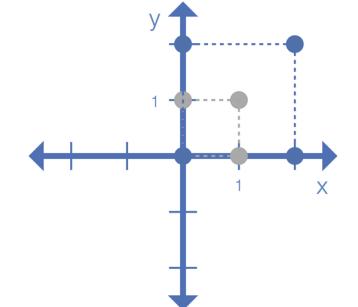
Translate

$$\begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$



Scale

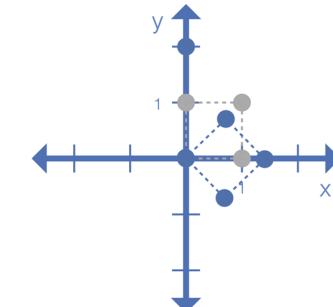
$$\begin{bmatrix} 2 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



Rotate

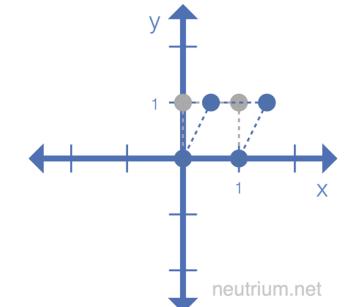
$$\begin{bmatrix} c & s & 0 \\ -s & c & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$$c = s = \sin(45^\circ)$$



Shear

$$\begin{bmatrix} 1 & 0.5 & 0 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$



neutrium.net

Outline

- What is image stitching/panorama image
- Edge Detector
- Blob Detector
- SIFT detector
- Image Transformation: 2D & 3D
- Image stitching/panorama Techniques

Projective Transformation

Mapping of one plane to another through a point

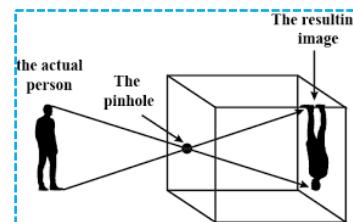
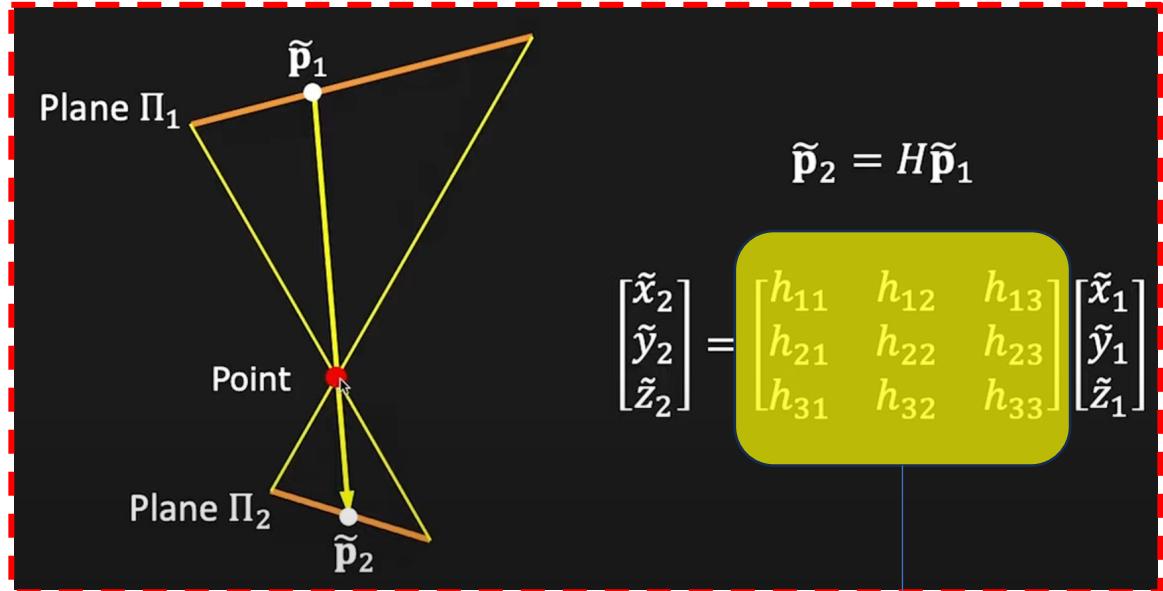
Any transformation of the form:

$$\begin{bmatrix} \tilde{x}_2 \\ \tilde{y}_2 \\ \tilde{z}_2 \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} \tilde{x}_1 \\ \tilde{y}_1 \\ \tilde{z}_1 \end{bmatrix} \quad \tilde{\mathbf{p}}_2 = H\tilde{\mathbf{p}}_1$$



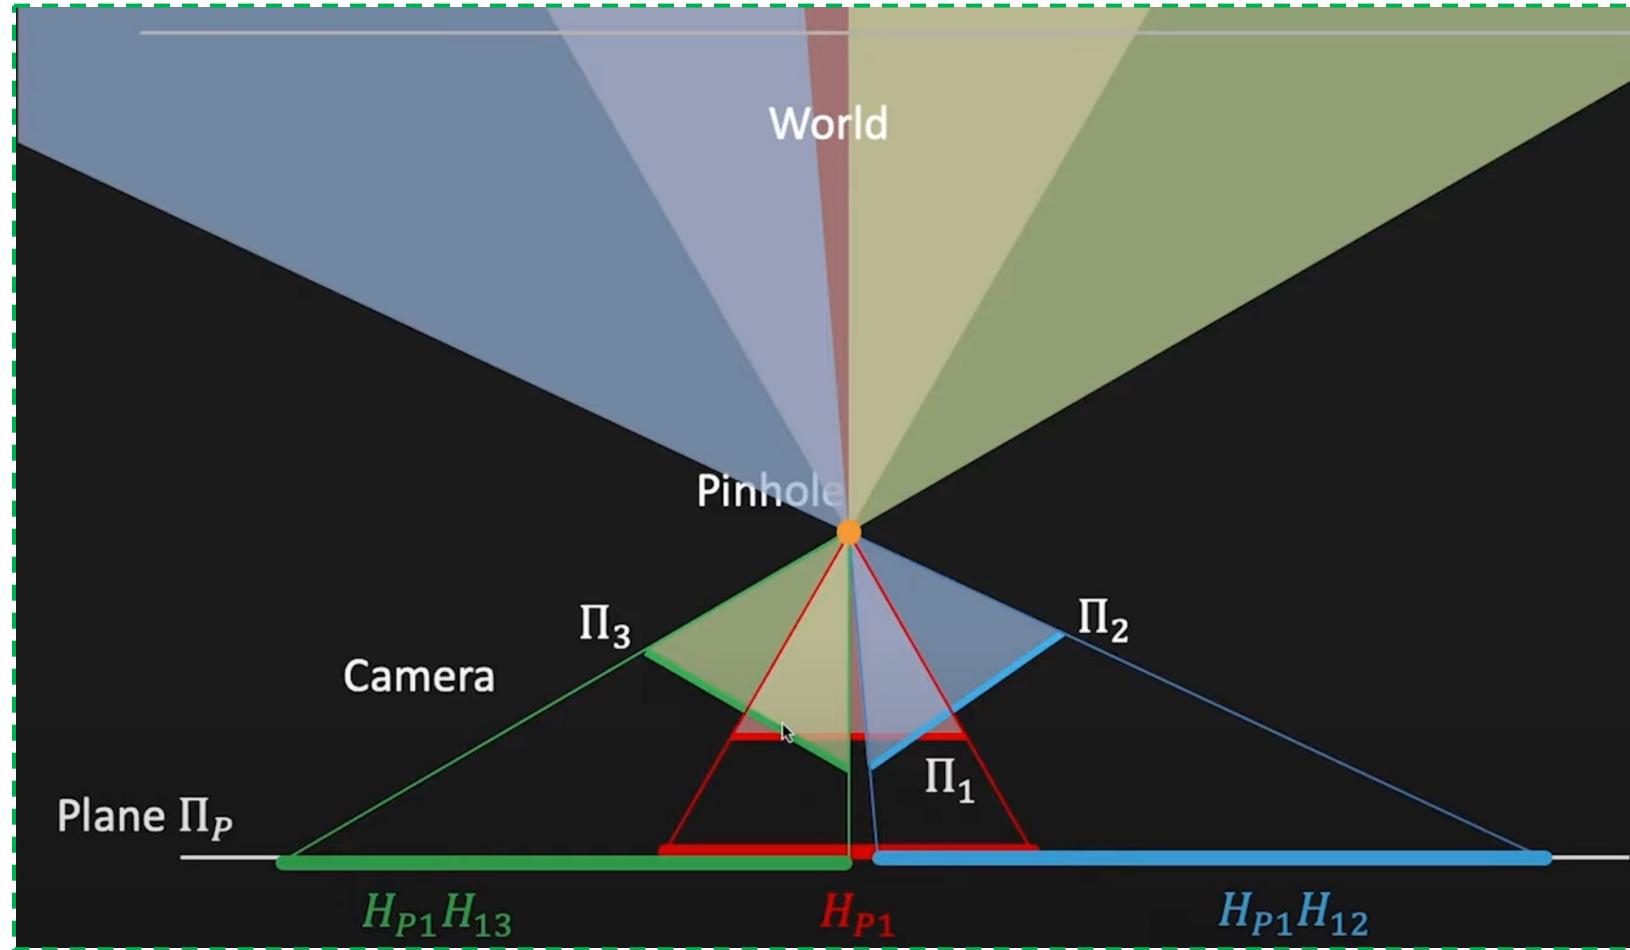
Homography can only be defined up to a scale.

$$\begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} \tilde{x}_1 \\ \tilde{y}_1 \\ \tilde{z}_1 \end{bmatrix} \equiv \begin{bmatrix} \tilde{x}_2 \\ \tilde{y}_2 \\ \tilde{z}_2 \end{bmatrix} \equiv k \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} \tilde{x}_1 \\ \tilde{y}_1 \\ \tilde{z}_1 \end{bmatrix}$$

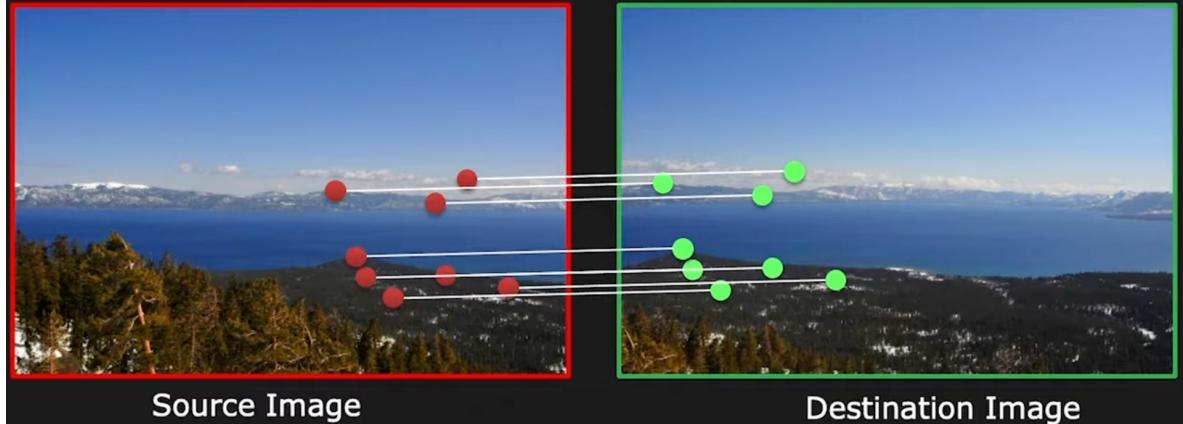


A homography is a projective transformation between two planes. How to compute it?

Homography Composition

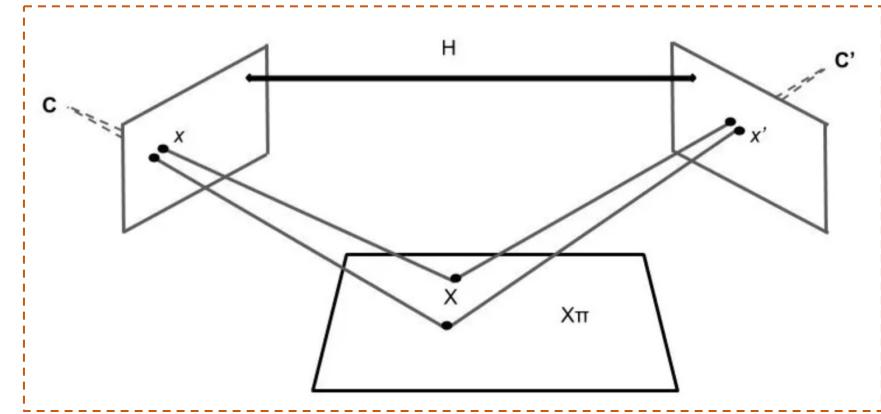


Computing Homography



Given a set of matching features/points between image 1 and image 2, find the homography H that best “agrees” with the matches. How many minimum pairs of matching points?

$$\begin{bmatrix} x_d \\ y_d \\ 1 \end{bmatrix} \equiv \begin{bmatrix} \tilde{x}_d \\ \tilde{y}_d \\ \tilde{z}_d \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x_s \\ y_s \\ 1 \end{bmatrix}$$



For a given pair i of corresponding points

$$x_d^{(i)} = \frac{\tilde{x}_d^{(i)}}{\tilde{z}_d^{(i)}} = \frac{h_{11}x_s^{(i)} + h_{12}y_s^{(i)} + h_{13}}{h_{31}x_s^{(i)} + h_{32}y_s^{(i)} + h_{33}}$$

$$y_d^{(i)} = \frac{\tilde{y}_d^{(i)}}{\tilde{z}_d^{(i)}} = \frac{h_{21}x_s^{(i)} + h_{22}y_s^{(i)} + h_{23}}{h_{31}x_s^{(i)} + h_{32}y_s^{(i)} + h_{33}}$$

Computing Homography

$$x_d^{(i)} = \frac{\tilde{x}_d^{(i)}}{\tilde{z}_d^{(i)}} = \frac{h_{11}x_s^{(i)} + h_{12}y_s^{(i)} + h_{13}}{h_{31}x_s^{(i)} + h_{32}y_s^{(i)} + h_{33}}$$

$$y_d^{(i)} = \frac{\tilde{y}_d^{(i)}}{\tilde{z}_d^{(i)}} = \frac{h_{21}x_s^{(i)} + h_{22}y_s^{(i)} + h_{23}}{h_{31}x_s^{(i)} + h_{32}y_s^{(i)} + h_{33}}$$



$$x_d^{(i)} (h_{31}x_s^{(i)} + h_{32}y_s^{(i)} + h_{33}) = h_{11}x_s^{(i)} + h_{12}y_s^{(i)} + h_{13}$$

$$y_d^{(i)} (h_{31}x_s^{(i)} + h_{32}y_s^{(i)} + h_{33}) = h_{21}x_s^{(i)} + h_{22}y_s^{(i)} + h_{23}$$



$$\begin{bmatrix} x_s^{(1)} & y_s^{(1)} & 1 & 0 & 0 & 0 & -x_d^{(1)}x_s^{(1)} & -x_d^{(1)}y_s^{(1)} & -x_d^{(1)} \\ 0 & 0 & 0 & x_s^{(1)} & y_s^{(1)} & 1 & -y_d^{(1)}x_s^{(1)} & -y_d^{(1)}y_s^{(1)} & -y_d^{(1)} \\ & & & & & \vdots & & & \\ x_s^{(i)} & y_s^{(i)} & 1 & 0 & 0 & 0 & -x_d^{(i)}x_s^{(i)} & -x_d^{(i)}y_s^{(i)} & -x_d^{(i)} \\ 0 & 0 & 0 & x_s^{(i)} & y_s^{(i)} & 1 & -y_d^{(i)}x_s^{(i)} & -y_d^{(i)}y_s^{(i)} & -y_d^{(i)} \\ & & & & & \vdots & & & \\ x_s^{(n)} & y_s^{(n)} & 1 & 0 & 0 & 0 & -x_d^{(n)}x_s^{(n)} & -x_d^{(n)}y_s^{(n)} & -x_d^{(n)} \\ 0 & 0 & 0 & x_s^{(n)} & y_s^{(n)} & 1 & -y_d^{(n)}x_s^{(n)} & -y_d^{(n)}y_s^{(n)} & -y_d^{(n)} \end{bmatrix} \begin{bmatrix} h_{11} \\ h_{12} \\ h_{13} \\ h_{21} \\ h_{22} \\ h_{23} \\ h_{31} \\ h_{32} \\ h_{33} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \\ 0 \\ \vdots \\ 0 \\ 0 \end{bmatrix}$$

A
(Known) \mathbf{h}
(Unknown)

$$\begin{bmatrix} x_s^{(i)} & y_s^{(i)} & 1 & 0 & 0 & 0 & -x_d^{(i)}x_s^{(i)} & -x_d^{(i)}y_s^{(i)} & -x_d^{(i)} \\ 0 & 0 & 0 & x_s^{(i)} & y_s^{(i)} & 1 & -y_d^{(i)}x_s^{(i)} & -y_d^{(i)}y_s^{(i)} & -y_d^{(i)} \end{bmatrix} \begin{bmatrix} h_{11} \\ h_{12} \\ h_{13} \\ h_{21} \\ h_{22} \\ h_{23} \\ h_{31} \\ h_{32} \\ h_{33} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \\ 0 \\ \vdots \\ 0 \\ 0 \end{bmatrix}$$

(Known) \mathbf{h}
(Unknown)

Nayar

All corresponding points

One corresponding point

Constraint Least Square

$$\begin{bmatrix} x_s^{(1)} & y_s^{(1)} & 1 & 0 & 0 & 0 & -x_d^{(1)}x_s^{(1)} & -x_d^{(1)}y_s^{(1)} & -x_d^{(1)} \\ 0 & 0 & 0 & x_s^{(1)} & y_s^{(1)} & 1 & -y_d^{(1)}x_s^{(1)} & -y_d^{(1)}y_s^{(1)} & -y_d^{(1)} \\ & & & & & \vdots & & & \\ x_s^{(i)} & y_s^{(i)} & 1 & 0 & 0 & 0 & -x_d^{(i)}x_s^{(i)} & -x_d^{(i)}y_s^{(i)} & -x_d^{(i)} \\ 0 & 0 & 0 & x_s^{(i)} & y_s^{(i)} & 1 & -y_d^{(i)}x_s^{(i)} & -y_d^{(i)}y_s^{(i)} & -y_d^{(i)} \\ & & & & & \vdots & & & \\ x_s^{(n)} & y_s^{(n)} & 1 & 0 & 0 & 0 & -x_d^{(n)}x_s^{(n)} & -x_d^{(n)}y_s^{(n)} & -x_d^{(n)} \\ 0 & 0 & 0 & x_s^{(n)} & y_s^{(n)} & 1 & -y_d^{(n)}x_s^{(n)} & -y_d^{(n)}y_s^{(n)} & -y_d^{(n)} \end{bmatrix} = \begin{bmatrix} h_{11} \\ h_{12} \\ h_{13} \\ h_{21} \\ h_{22} \\ h_{23} \\ h_{31} \\ h_{32} \\ h_{33} \end{bmatrix}$$

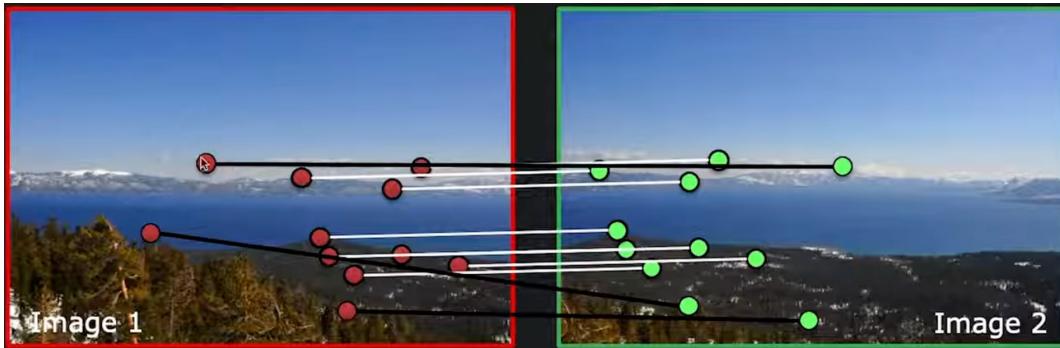
Solve for \mathbf{h} : $A \mathbf{h} = \mathbf{0}$ such that $\|\mathbf{h}\|^2 = 1$

Define least squares problem:

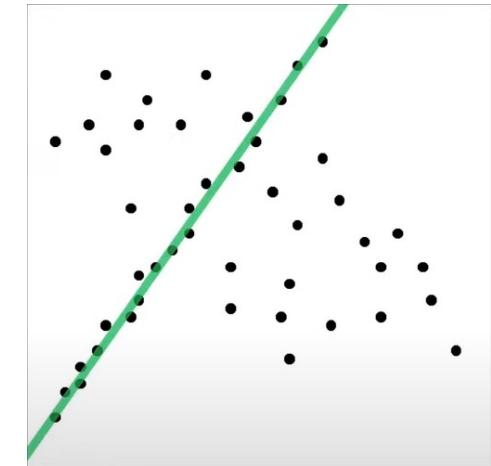
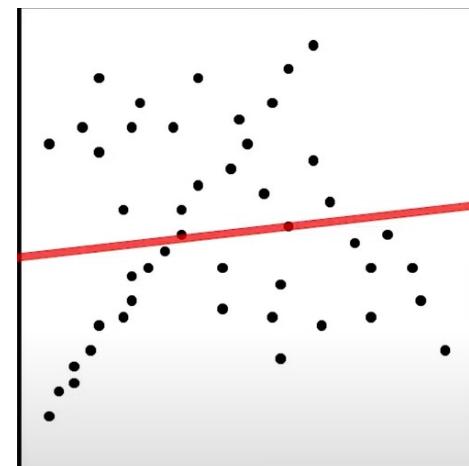
$$\min_{\mathbf{h}} \|\mathbf{A}\mathbf{h}\|^2 \text{ such that } \|\mathbf{h}\|^2 = 1$$

Deal with outlier: RANSAC

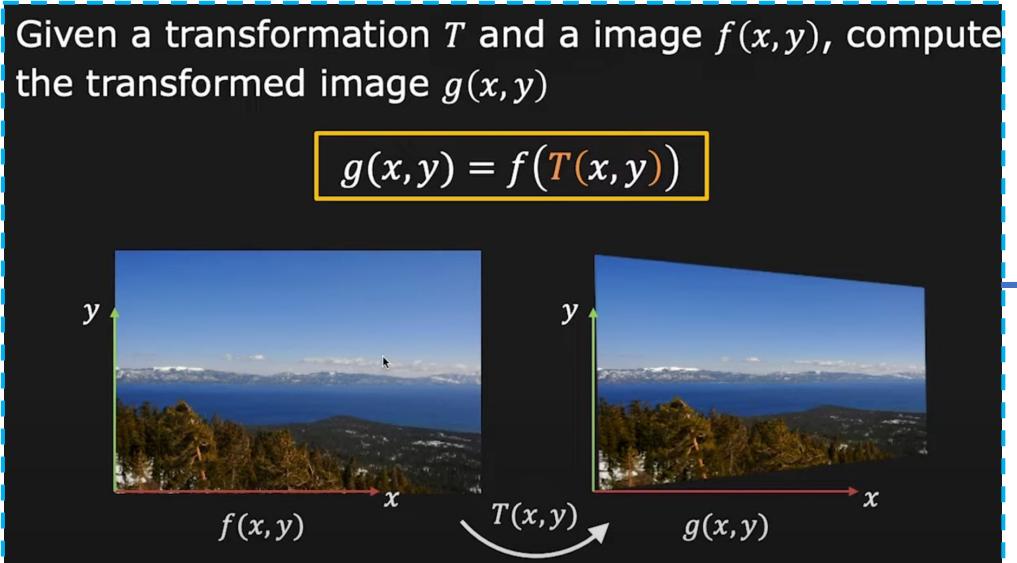
How to compute the homography in the presence of outliers



1. Randomly choose s samples. Typically s is the minimum samples to fit a model
2. Fit the model to the randomly chosen samples
3. Count the number M of data points (inliers) that fit the model within a measure of error ε
4. Repeat steps 1-3 N times
5. Choose the model that has the largest number M of inliers



Warping Images

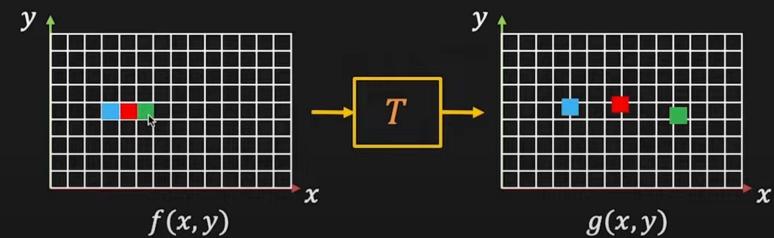


Result in holes

Forward Warping

Send each pixel (x, y) in $f(x, y)$ to its corresponding location $T(x, y)$ in $g(x, y)$

$$g(x, y) = f(T(x, y))$$



Get each pixel (x, y) in $g(x, y)$ from its corresponding location $T^{-1}(x, y)$ in $f(x, y)$

$$g(x, y) = f(T^{-1}(x, y))$$

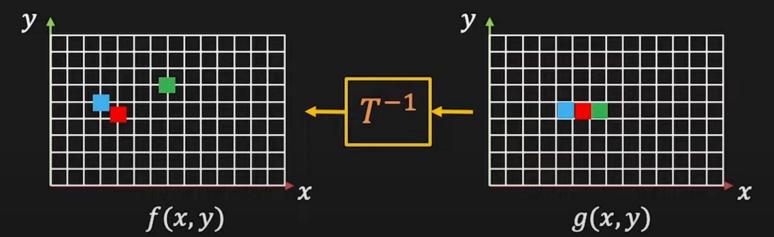
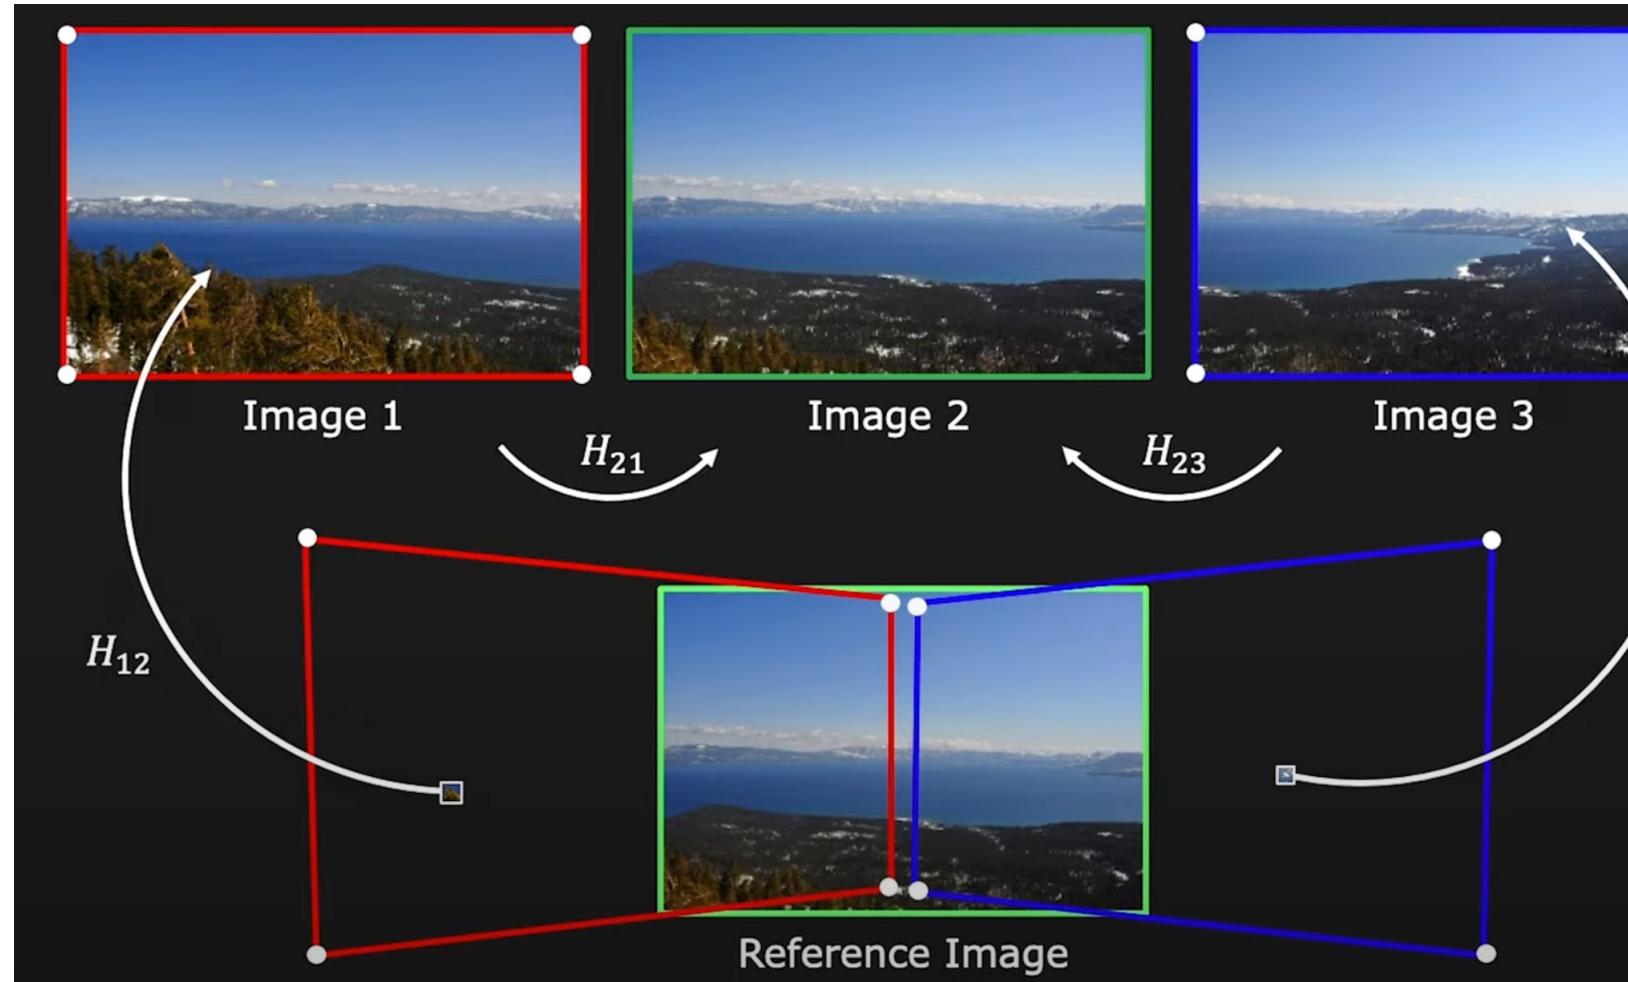
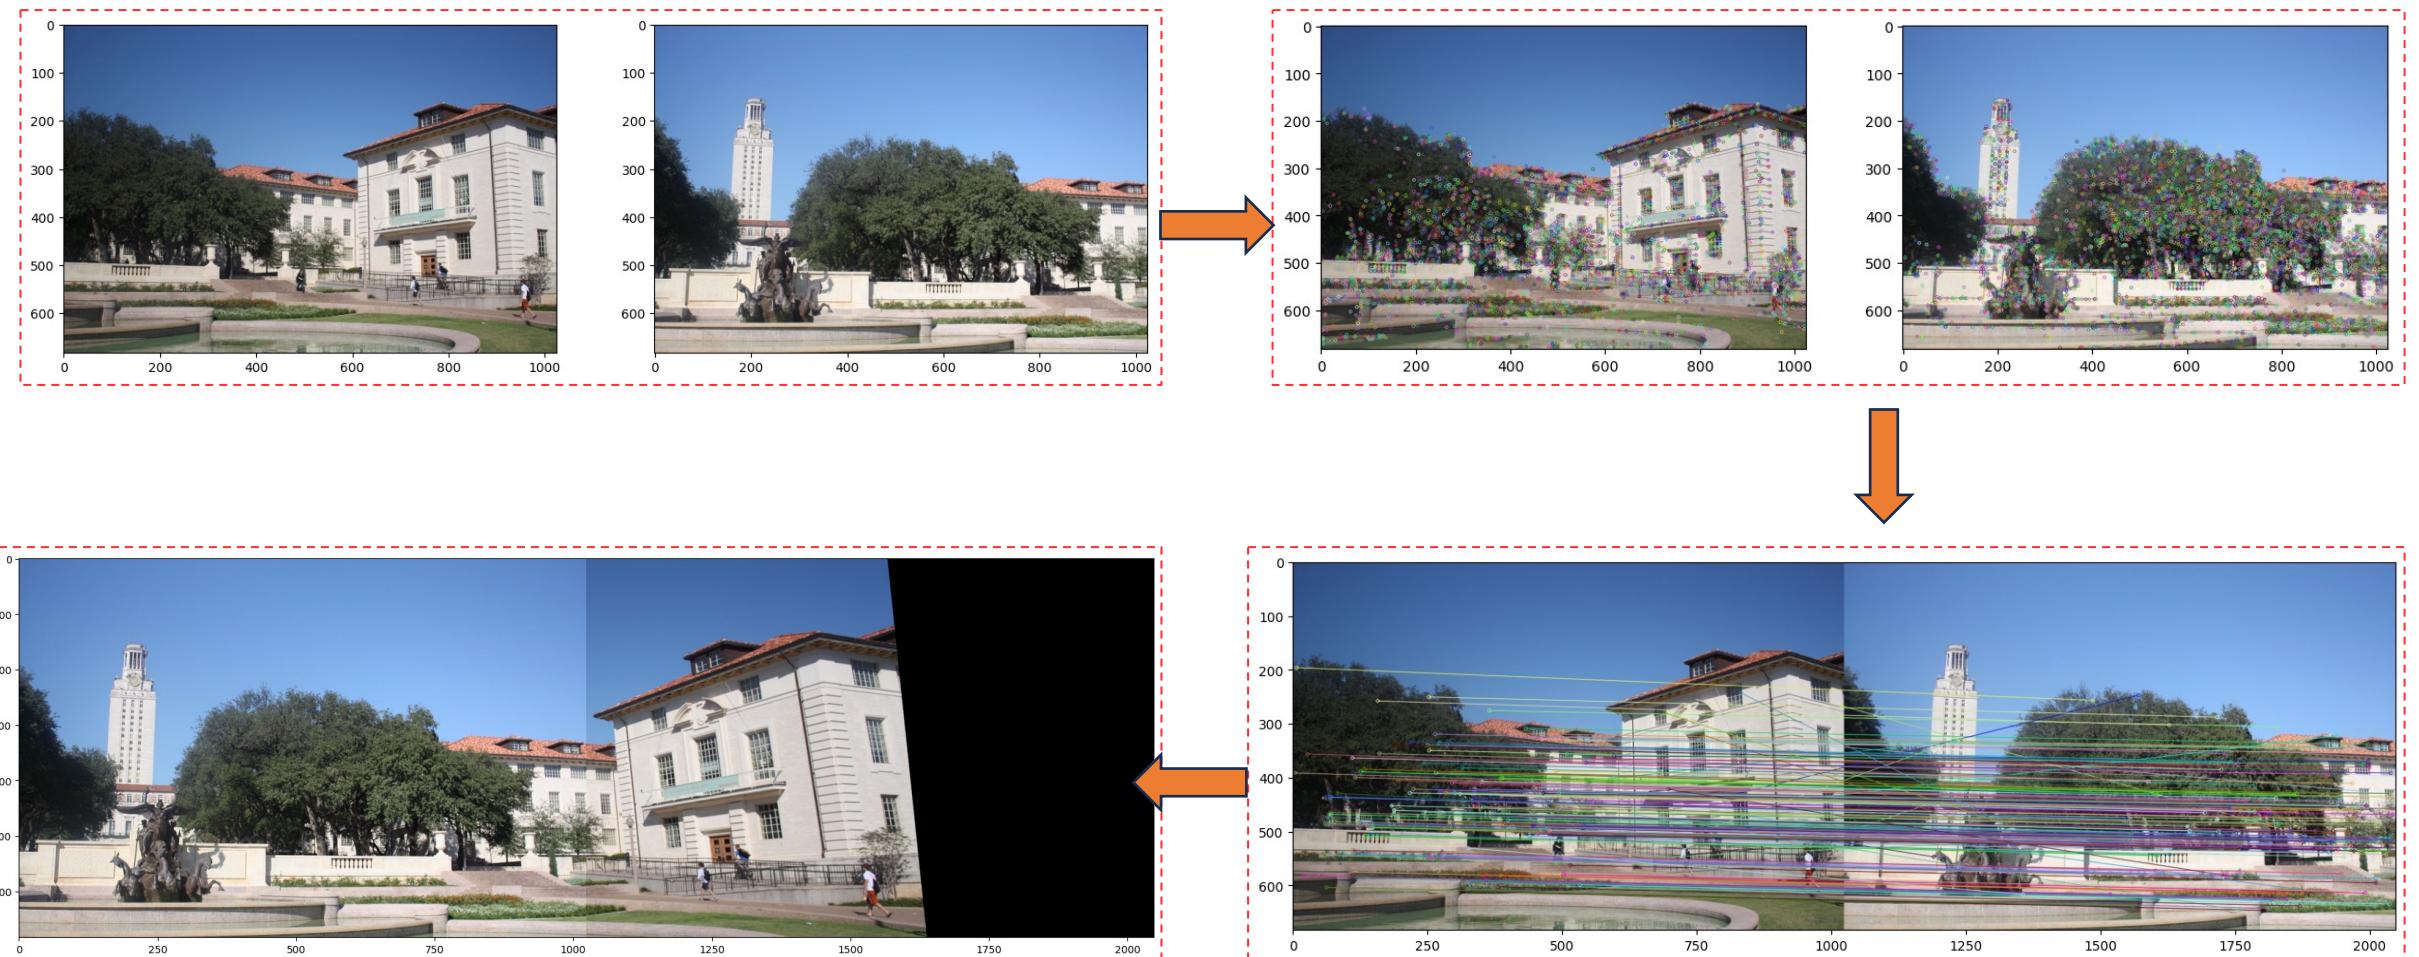


Image Alignment Process



Blending Images

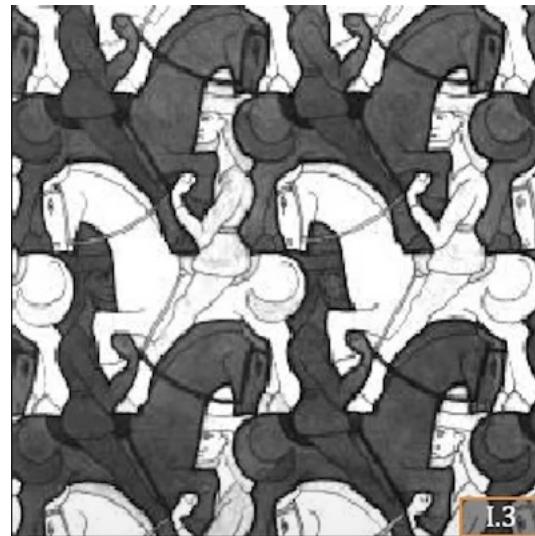


Seams still visible

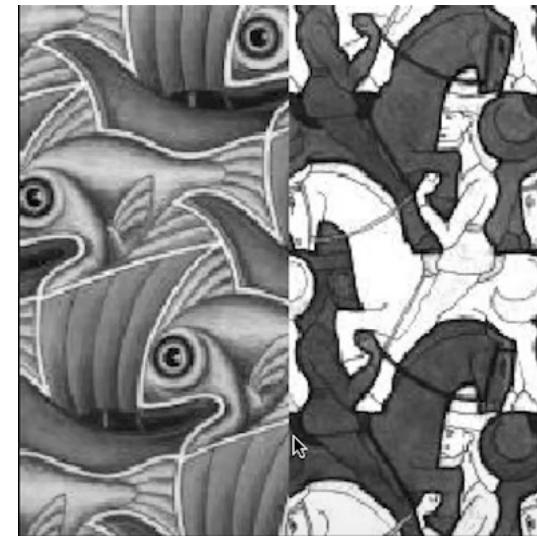
Blending Images



+



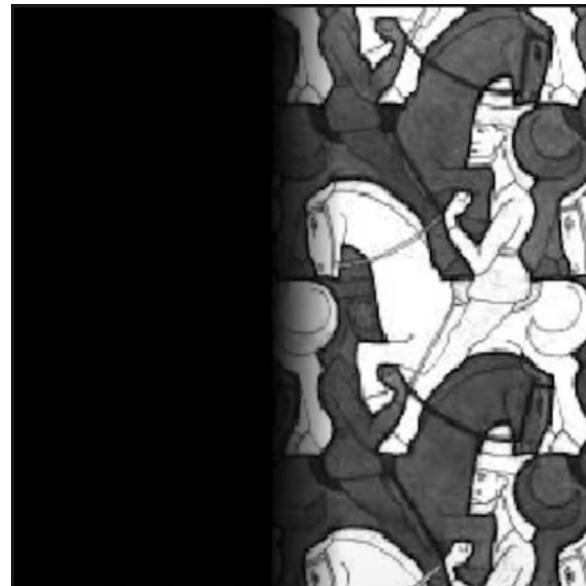
=



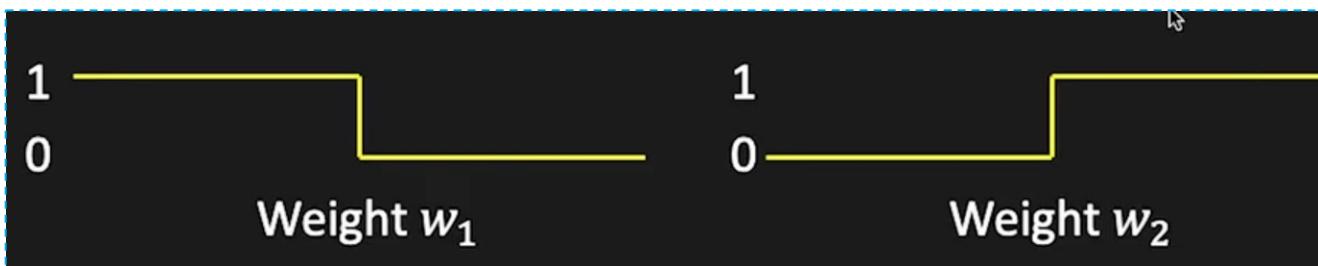
Blending Images



+

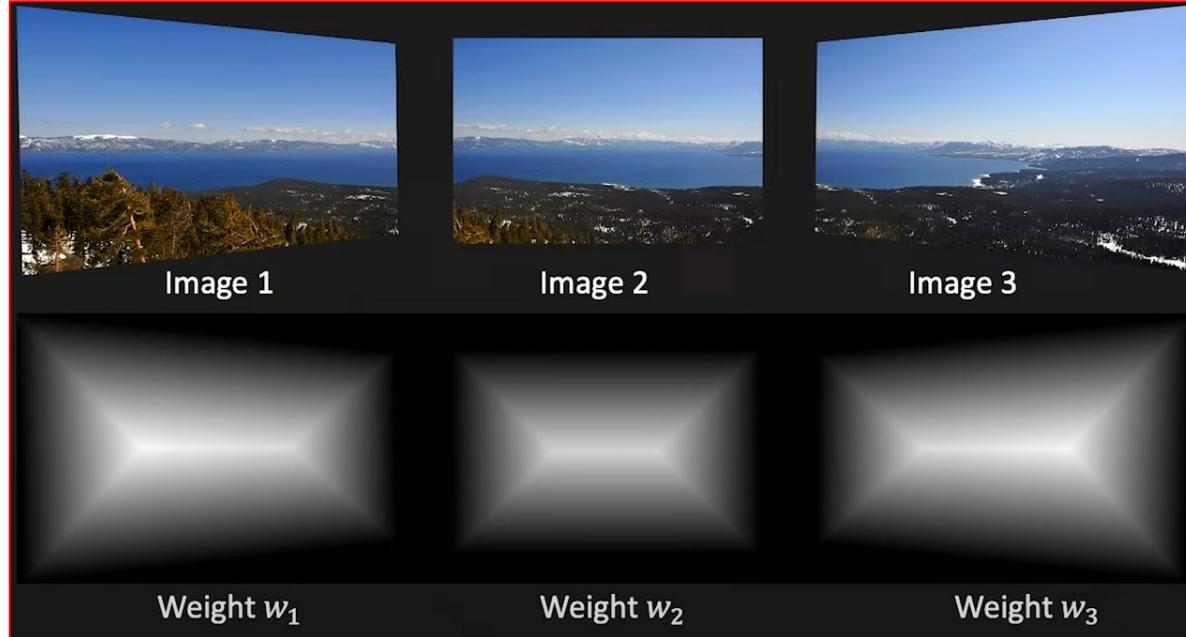


=

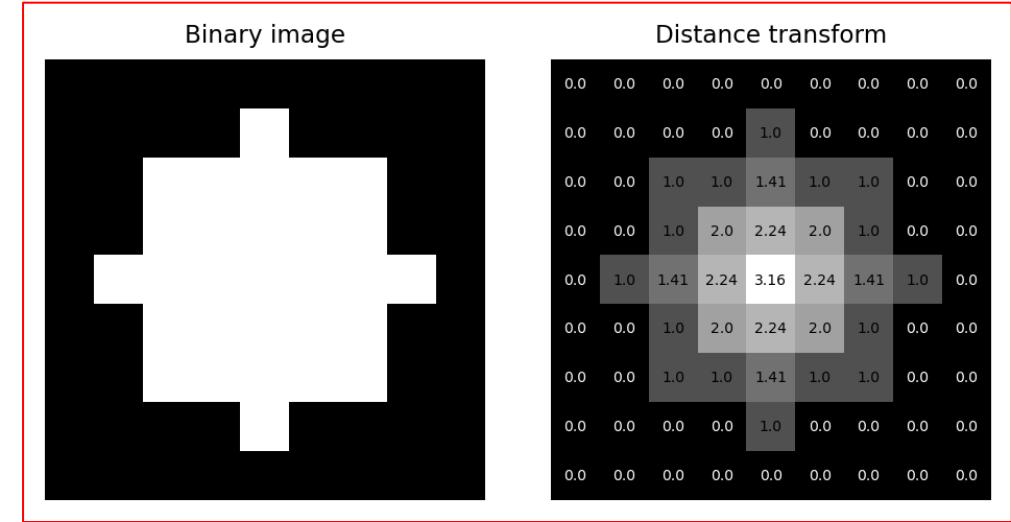


$$I_{\text{blend}} = \frac{w_1 I_1 + w_2 I_2}{w_1 + w_2}$$

Blending Images (Distance Transform)



Pixels closer to the edge get lower weight



The distance transform is an operator normally only applied to binary images. The result of the transform is a graylevel image that looks similar to the input image, except that the graylevel intensities of points inside foreground regions are changed to show the distance to the closest boundary from each point

Distance Transform

Euclidean

0	0	0
0	1	0
0	0	0

Image

1.41	1.0	1.41
1.0	0.0	1.0
1.41	1.0	1.41

Distance Transform

Chessboard

0	0	0
0	1	0
0	0	0

Image

1	1	1
1	0	1
1	1	1

Distance Transform

City Block

0	0	0
0	1	0
0	0	0

Image

2	1	2
1	0	1
2	1	2

Distance Transform

Quasi-Euclidean

0	0	0	0	0
0	0	0	0	0
0	0	1	0	0
0	0	0	0	0
0	0	0	0	0

Image

2.8	2.4	2.0	2.4	2.8
2.4	1.4	1.0	1.4	2.4
2.0	1.0	0	1.0	2.0
2.4	1.4	1.0	1.4	2.4
2.8	2.4	2.0	2.4	2.8

Distance Transform

Blending Images (Distance Transform)

