

Imbalanced Data

Quang-Vinh Dinh
Ph.D. in Computer Science

Outline


- **Introduction**
- **Examples and Discussion**
- **Metrics**
- **Case Study**

Introduction

❖ Imbalanced Data vs. Balanced Data


Imbalanced Data

Negative	100000
Positive	200




Balanced Data

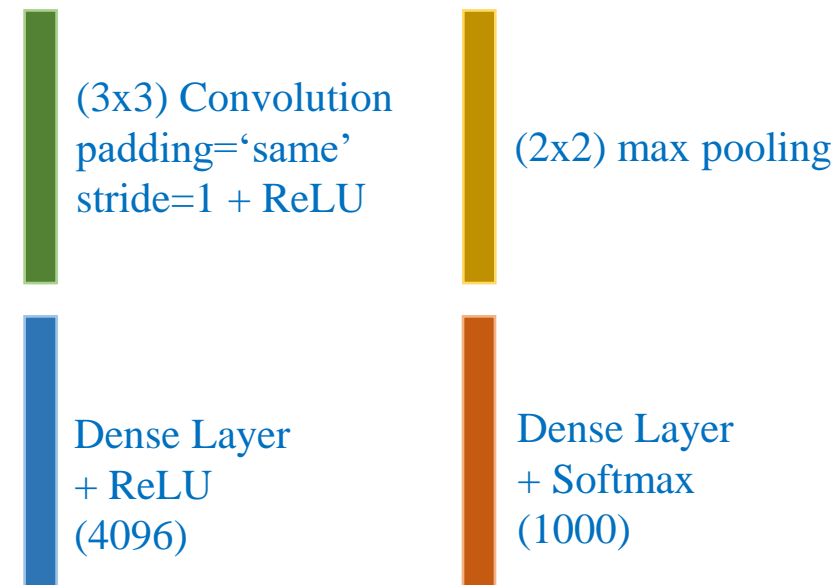
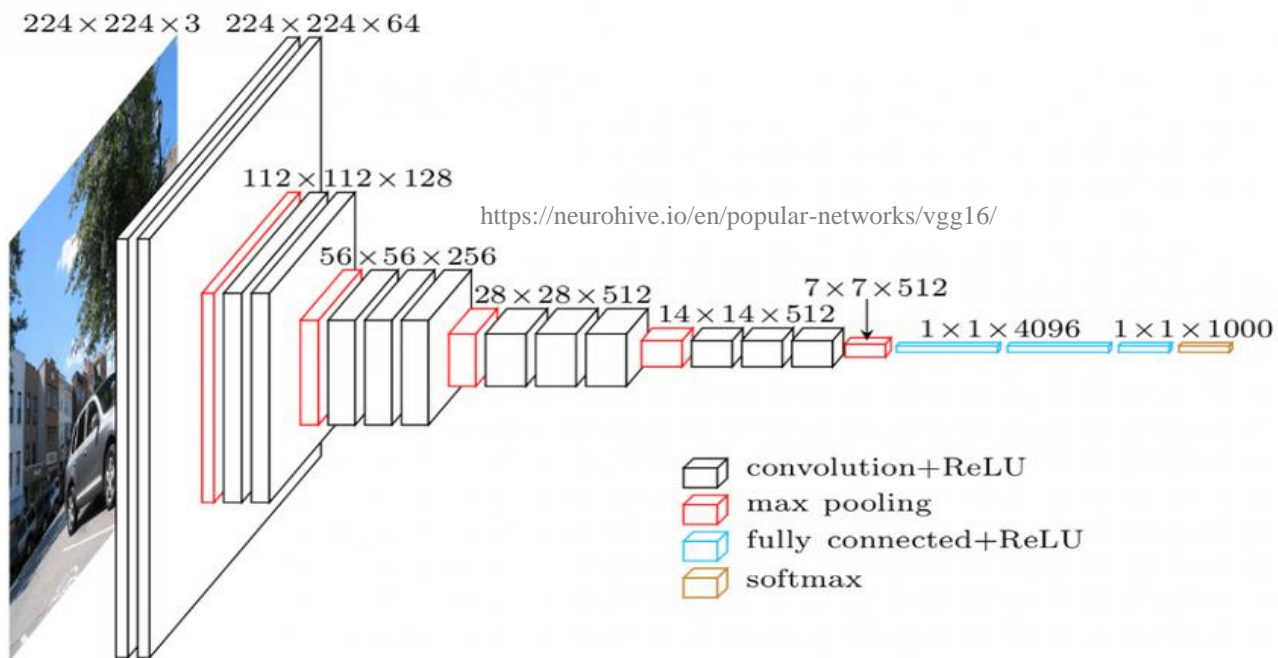
Negative	50100
Positive	50100



Introduction


❖ Cat-Dog dataset



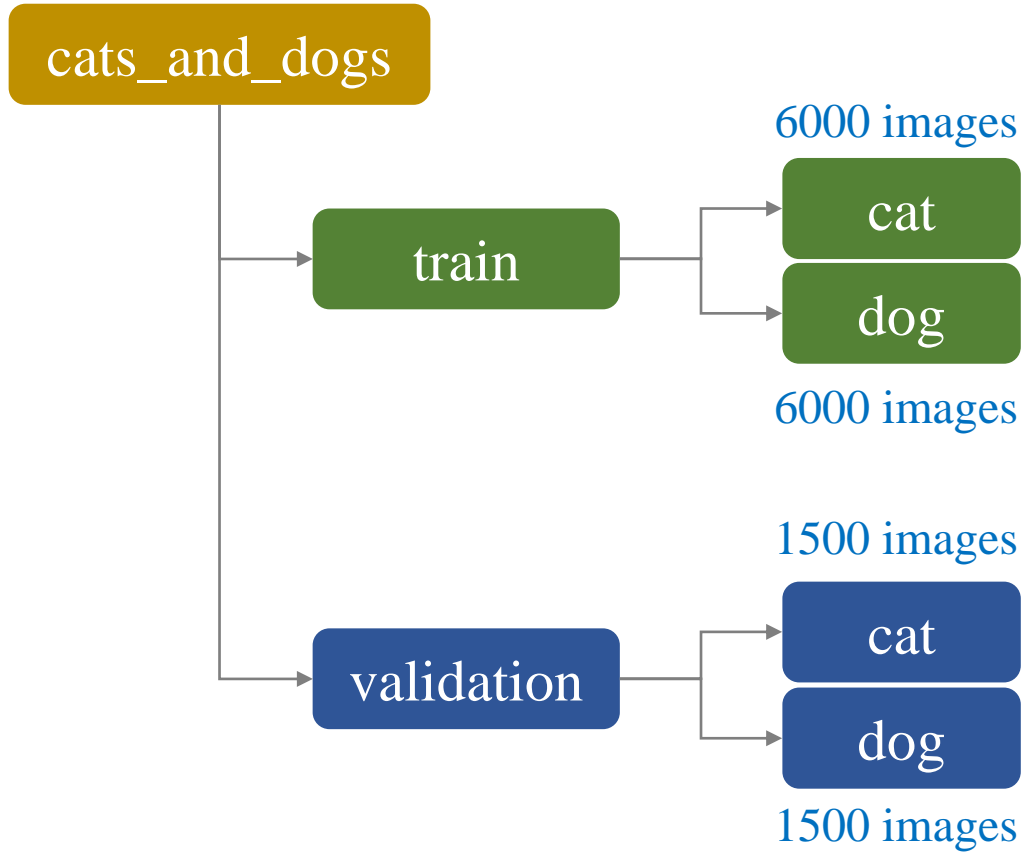


VGG16

VGG-16



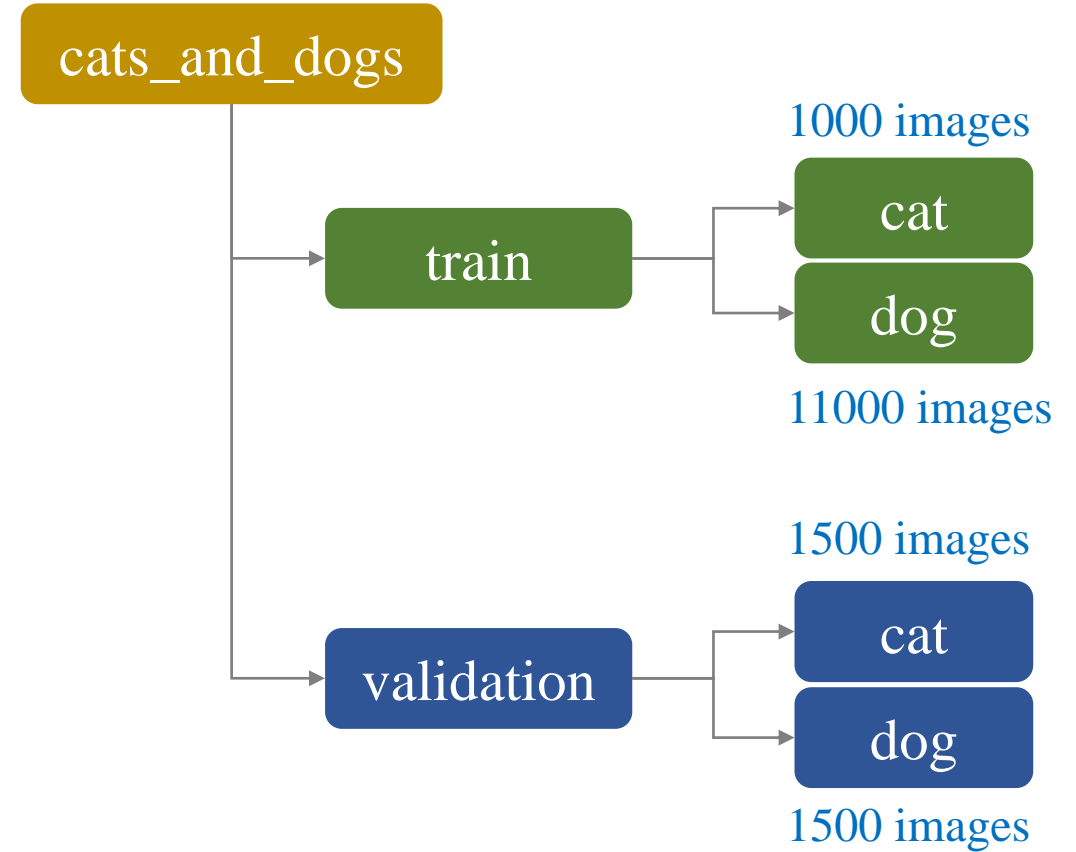
Introduction



Correct prediction

$$\#_{cat} = 1245$$

$$\#_{dog} = 1236$$



Correct prediction

$$\#_{cat} = 1052$$

$$\#_{dog} = 1443$$

Introduction

❖ Why?

Correct prediction (1)

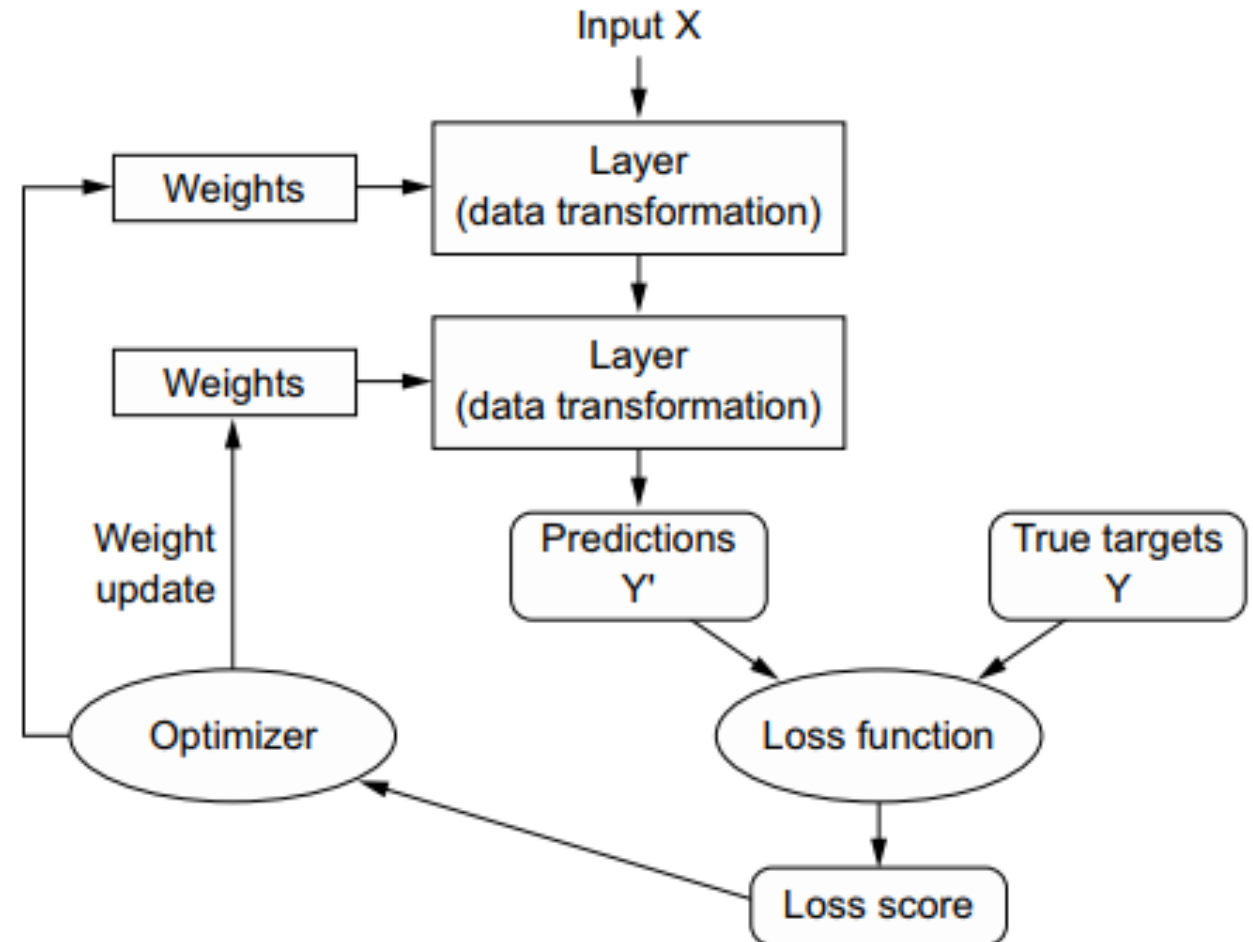
$\#_{cat} = 1245$

$\#_{dog} = 1236$

Correct prediction (2)

$\#_{cat} = 1052$

$\#_{dog} = 1443$



Outline

- **Introduction**
- **Examples and Discussion**
- **Metrics**
- **Case Study**

Imbalanced Data

❖ Lie/Truth classification

Feature Output Label

Input	Output	Label
...	0.3	0
...	0.8	0
...	0.7	0
...	0.4	0
...	0.6	1
...	0.8	1
...	0.9	1
...	0.2	1

if $y = 0$



$$P(\hat{y}) = -\log(1 - \hat{y})$$

if $y = 1$

$$P(\hat{y}) = -\log(\hat{y})$$

Binary cross-entropy

$$L(.) = -y\log\hat{y} - (1 - y)\log(1 - \hat{y})$$



Imbalanced Data

❖ Lie/Truth classification

Feature Output Label

Input	Output	Label	Loss
...	0.3	0	0.356
...	0.8	0	1.609
...	0.7	0	1.203
...	0.4	0	0.511
...	0.6	1	0.511
...	0.8	1	0.223
...	0.9	1	0.105
...	0.2	1	1.609

if $y = 0$



$$L(\hat{y}) = -\log(1 - \hat{y})$$

if $y = 1$

$$L(\hat{y}) = -\log(\hat{y})$$

Binary cross-entropy

$$L(y, \hat{y}) = -y\log\hat{y} - (1 - y)\log(1 - \hat{y})$$



Imbalanced Data

❖ Lie/Truth classification

Feature	Output	Label
Input	Output	Label
...	0.2	0
...	0.8	0
...	0.3	0
...	0.1	0
...	0.8	1
...	0.8	1
...	0.9	1
...	0.9	1

After a period of time

if $y = 0$



$$P(\hat{y}) = -\log(1 - \hat{y})$$

if $y = 1$

$$P(\hat{y}) = -\log(\hat{y})$$

Binary cross-entropy

$$L(.) = -y\log\hat{y} - (1 - y)\log(1 - \hat{y})$$



Imbalanced Data

❖ Lie/Truth classification

Feature	Output	Label
Input	Output	Label
...	0.2	0
...	0.7	0
...	0.7	1
...	0.8	1
...	0.8	1
...	0.8	1
...	0.9	1
...	0.9	1
...	0.8	1
...	0.7	1

A special context

if $y = 0$



$$P(\hat{y}) = -\log(1 - \hat{y})$$

if $y = 1$

$$P(\hat{y}) = -\log(\hat{y})$$

Binary cross-entropy

$$L(.) = -y\log\hat{y} - (1 - y)\log(1 - \hat{y})$$



Imbalanced Data

❖ Lie/Truth classification

Feature	Output	Label
Input	Output	Label
...	0.7	0
...	0.8	1
...	0.7	1
...	0.8	1
...	0.8	1
...	0.8	1
...	0.9	1
...	0.9	1
...	0.8	1
...	0.7	1

a more severe context!

$$loss_{y=0} = 1.204$$

$$loss_{y=1} = 2.039$$

if $y = 0$



$$P(\hat{y}) = -\log(1 - \hat{y})$$

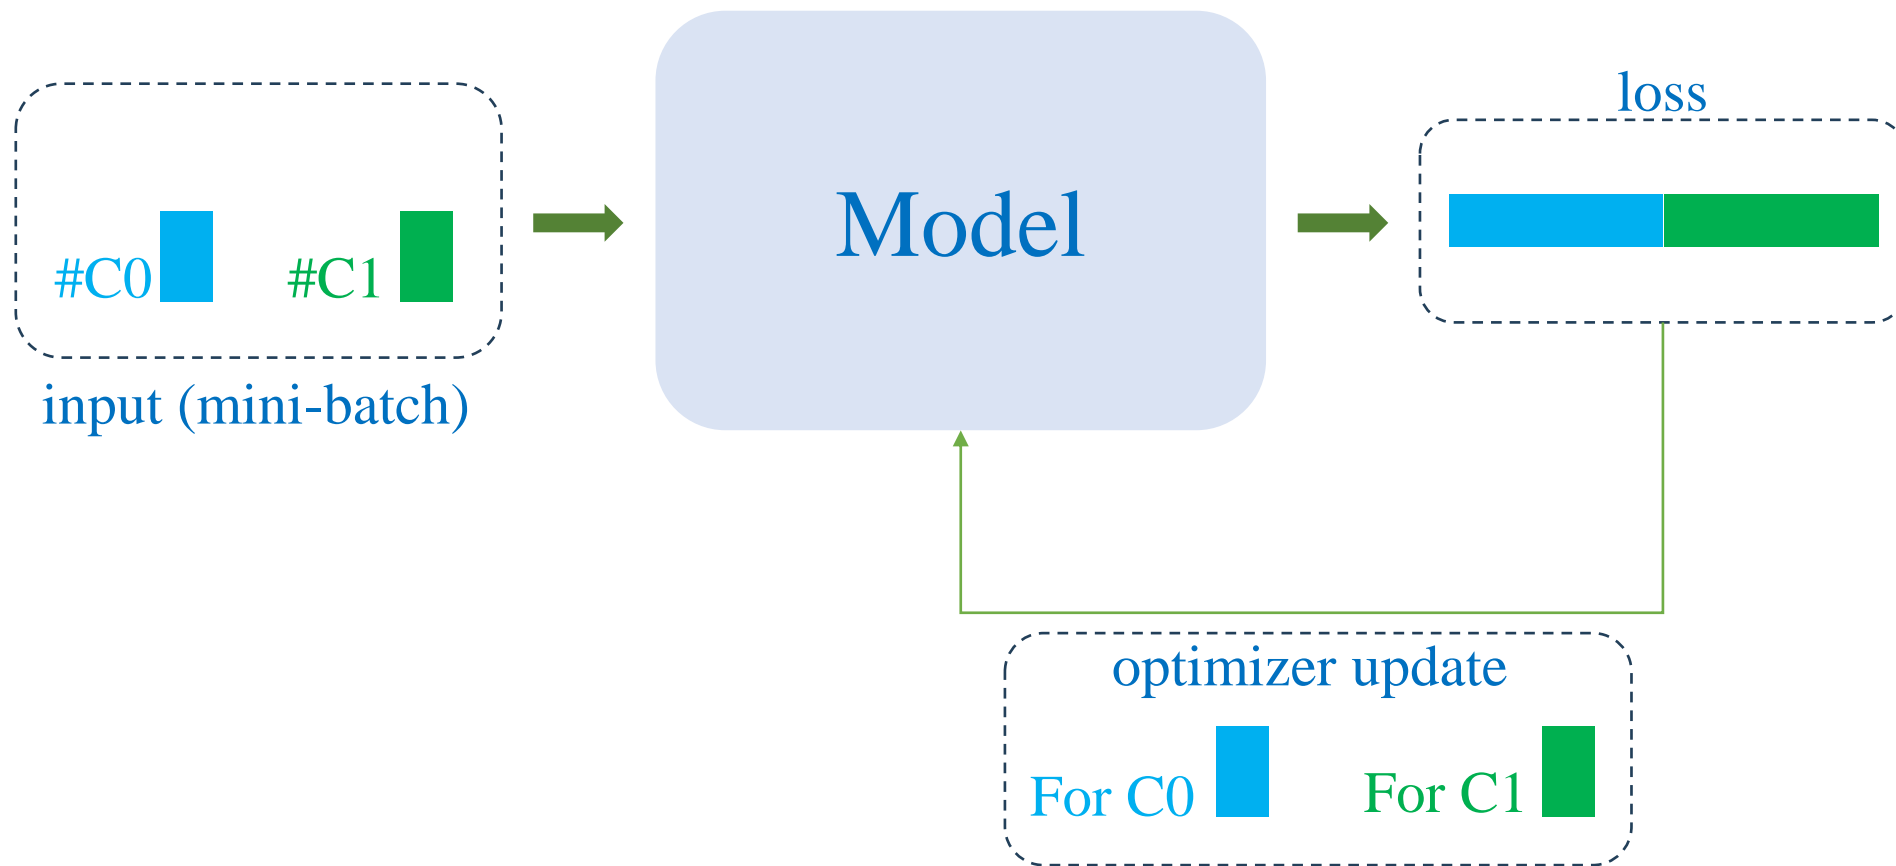
if $y = 1$

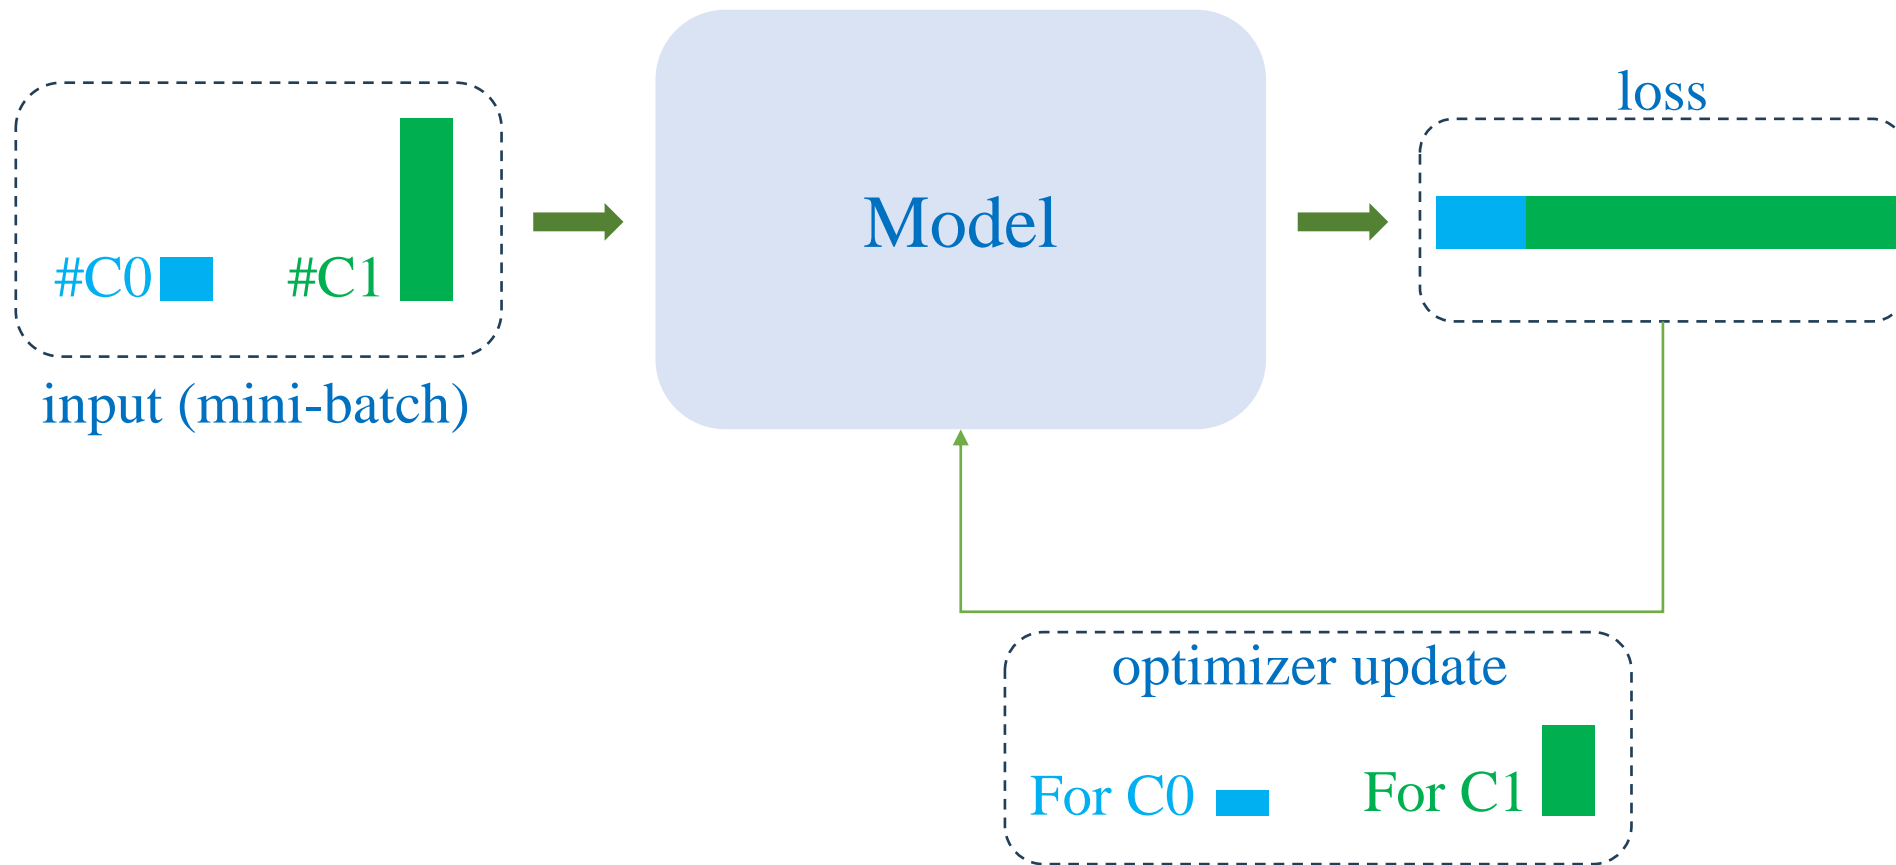
$$P(\hat{y}) = -\log(\hat{y})$$

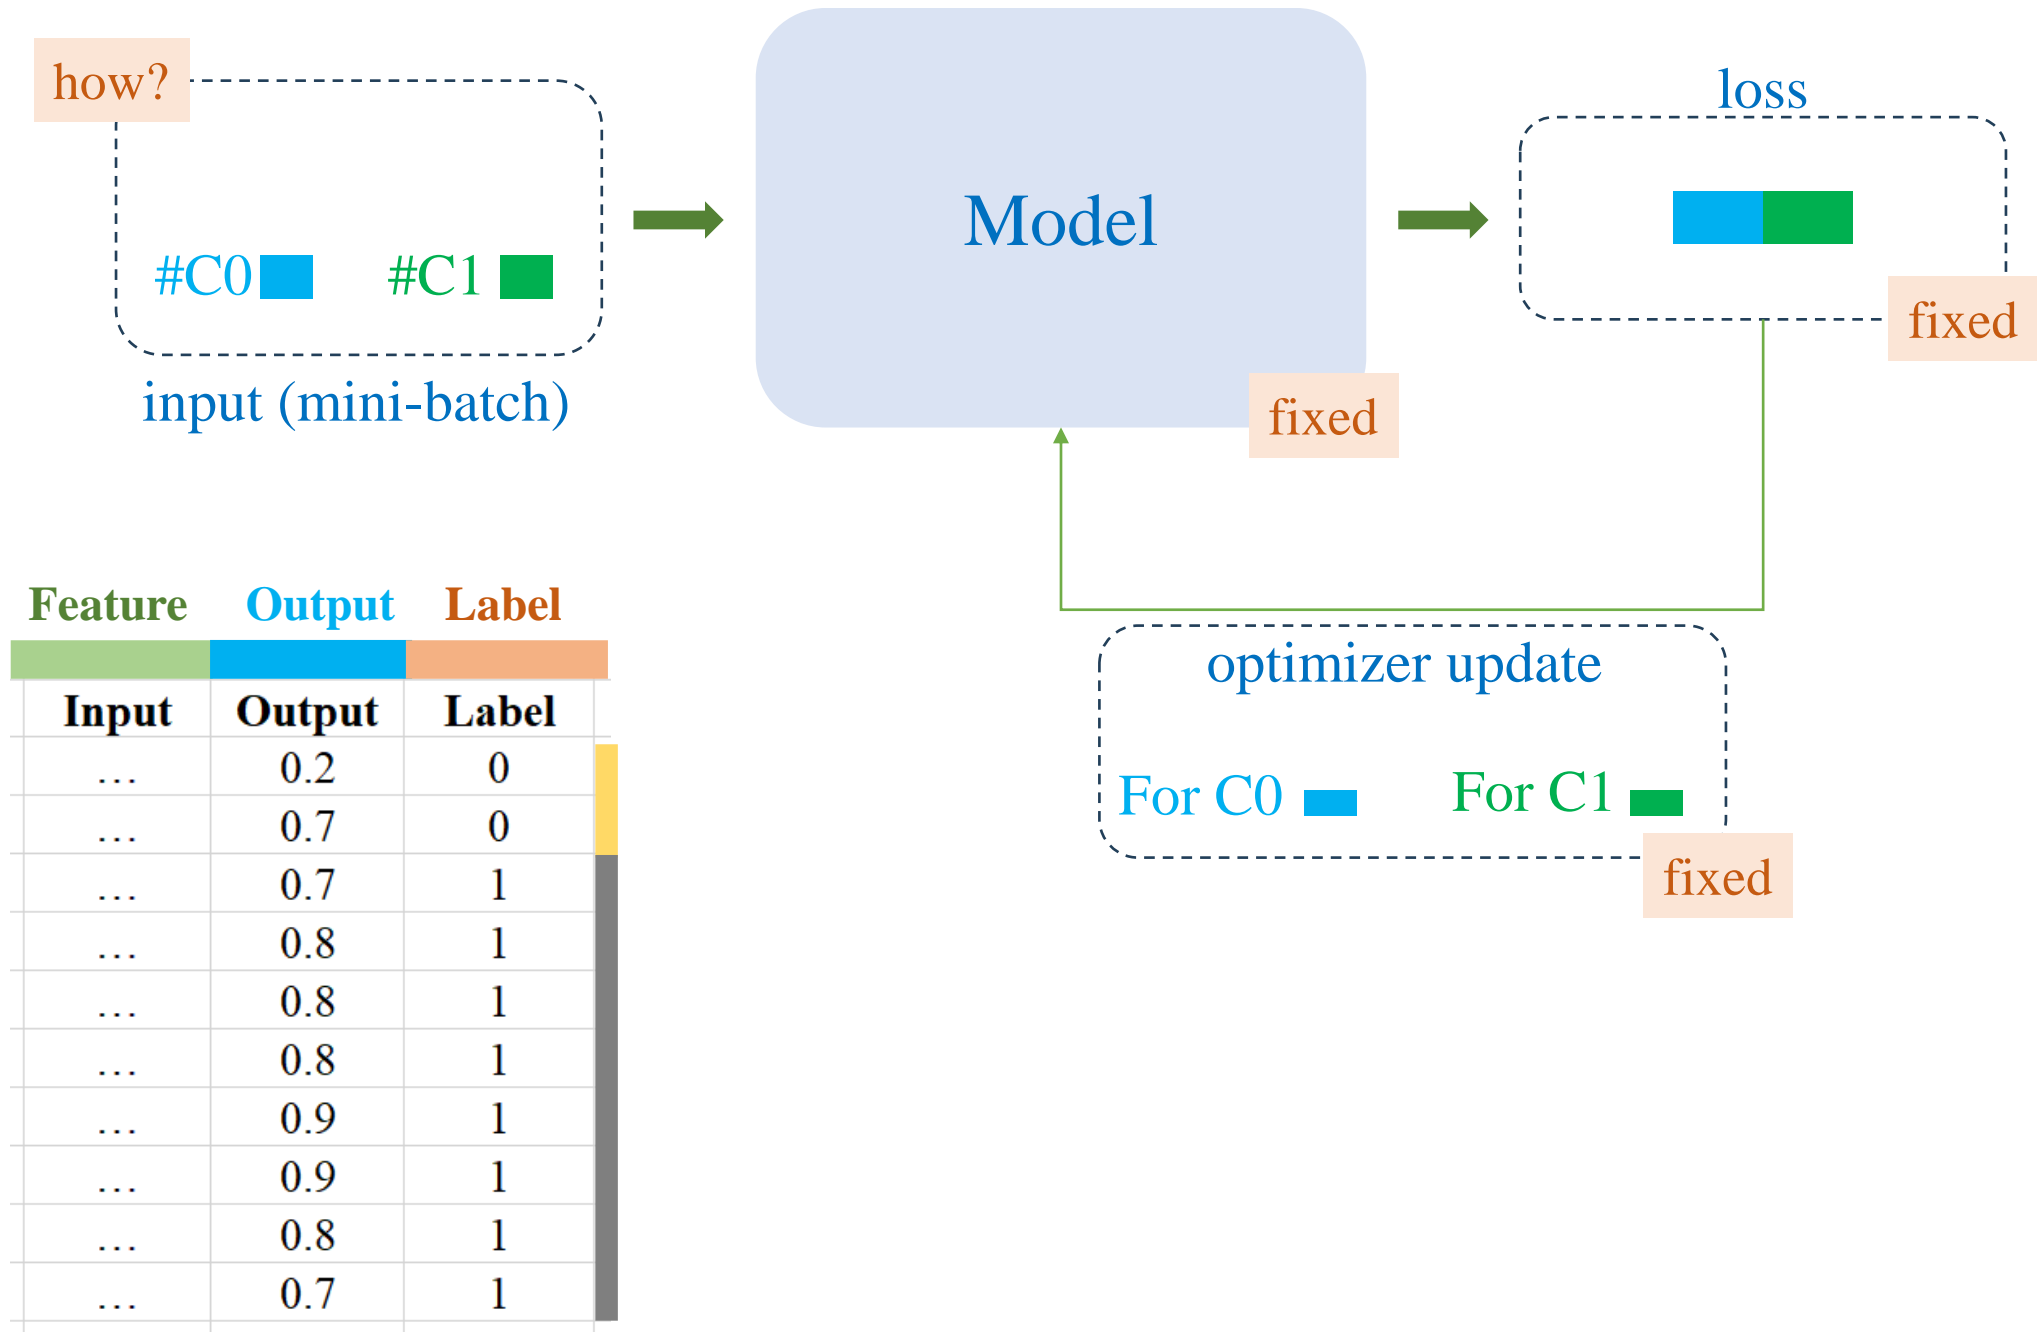
Binary cross-entropy

$$P = -y\log\hat{y} - (1 - y)\log(1 - \hat{y})$$

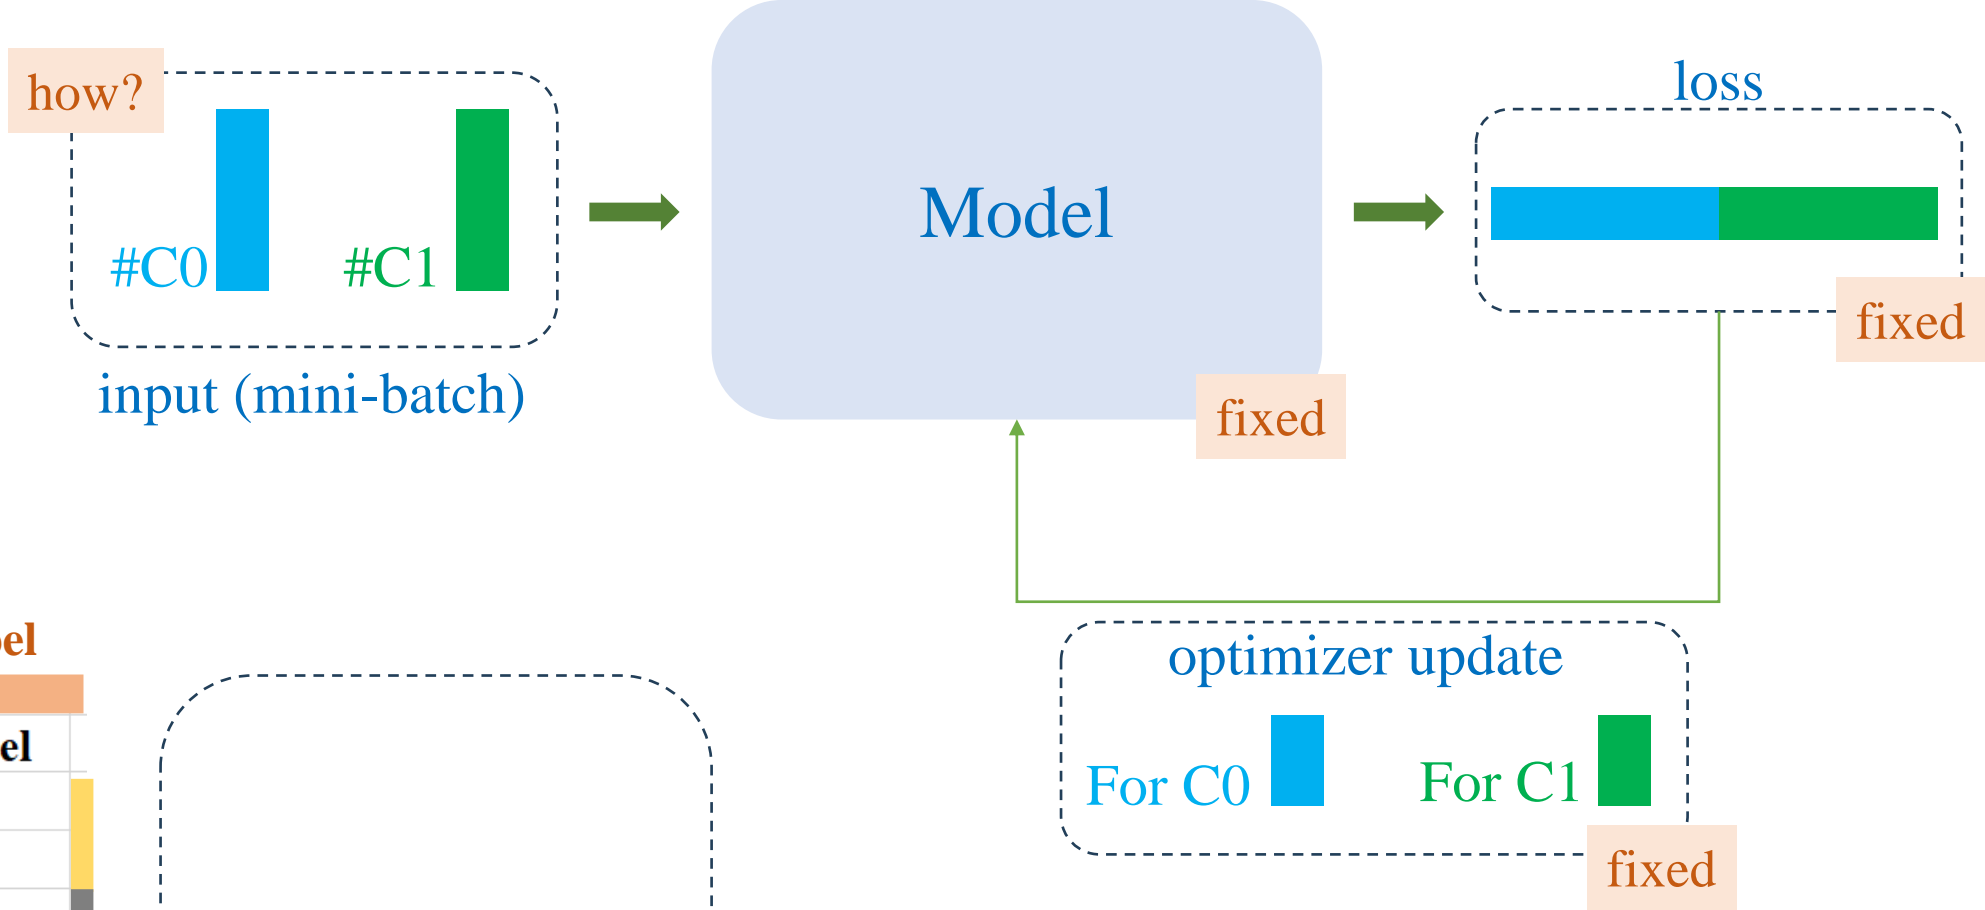









solution 2




Feature	Output	Label
Input	Output	Label
...	0.2	0
...	0.7	0
...	0.7	1
...	0.8	1
...	0.8	1
...	0.8	1
...	0.9	1
...	0.9	1
...	0.8	1
...	0.7	1

Imbalanced Data

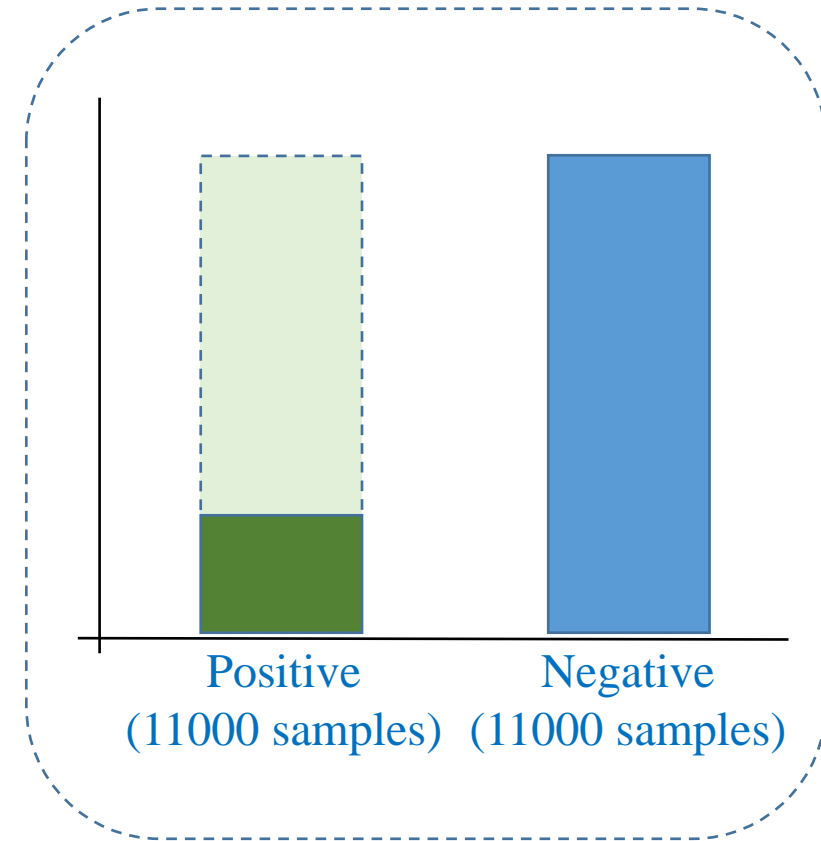
❖ Approach 1: Data manipulation



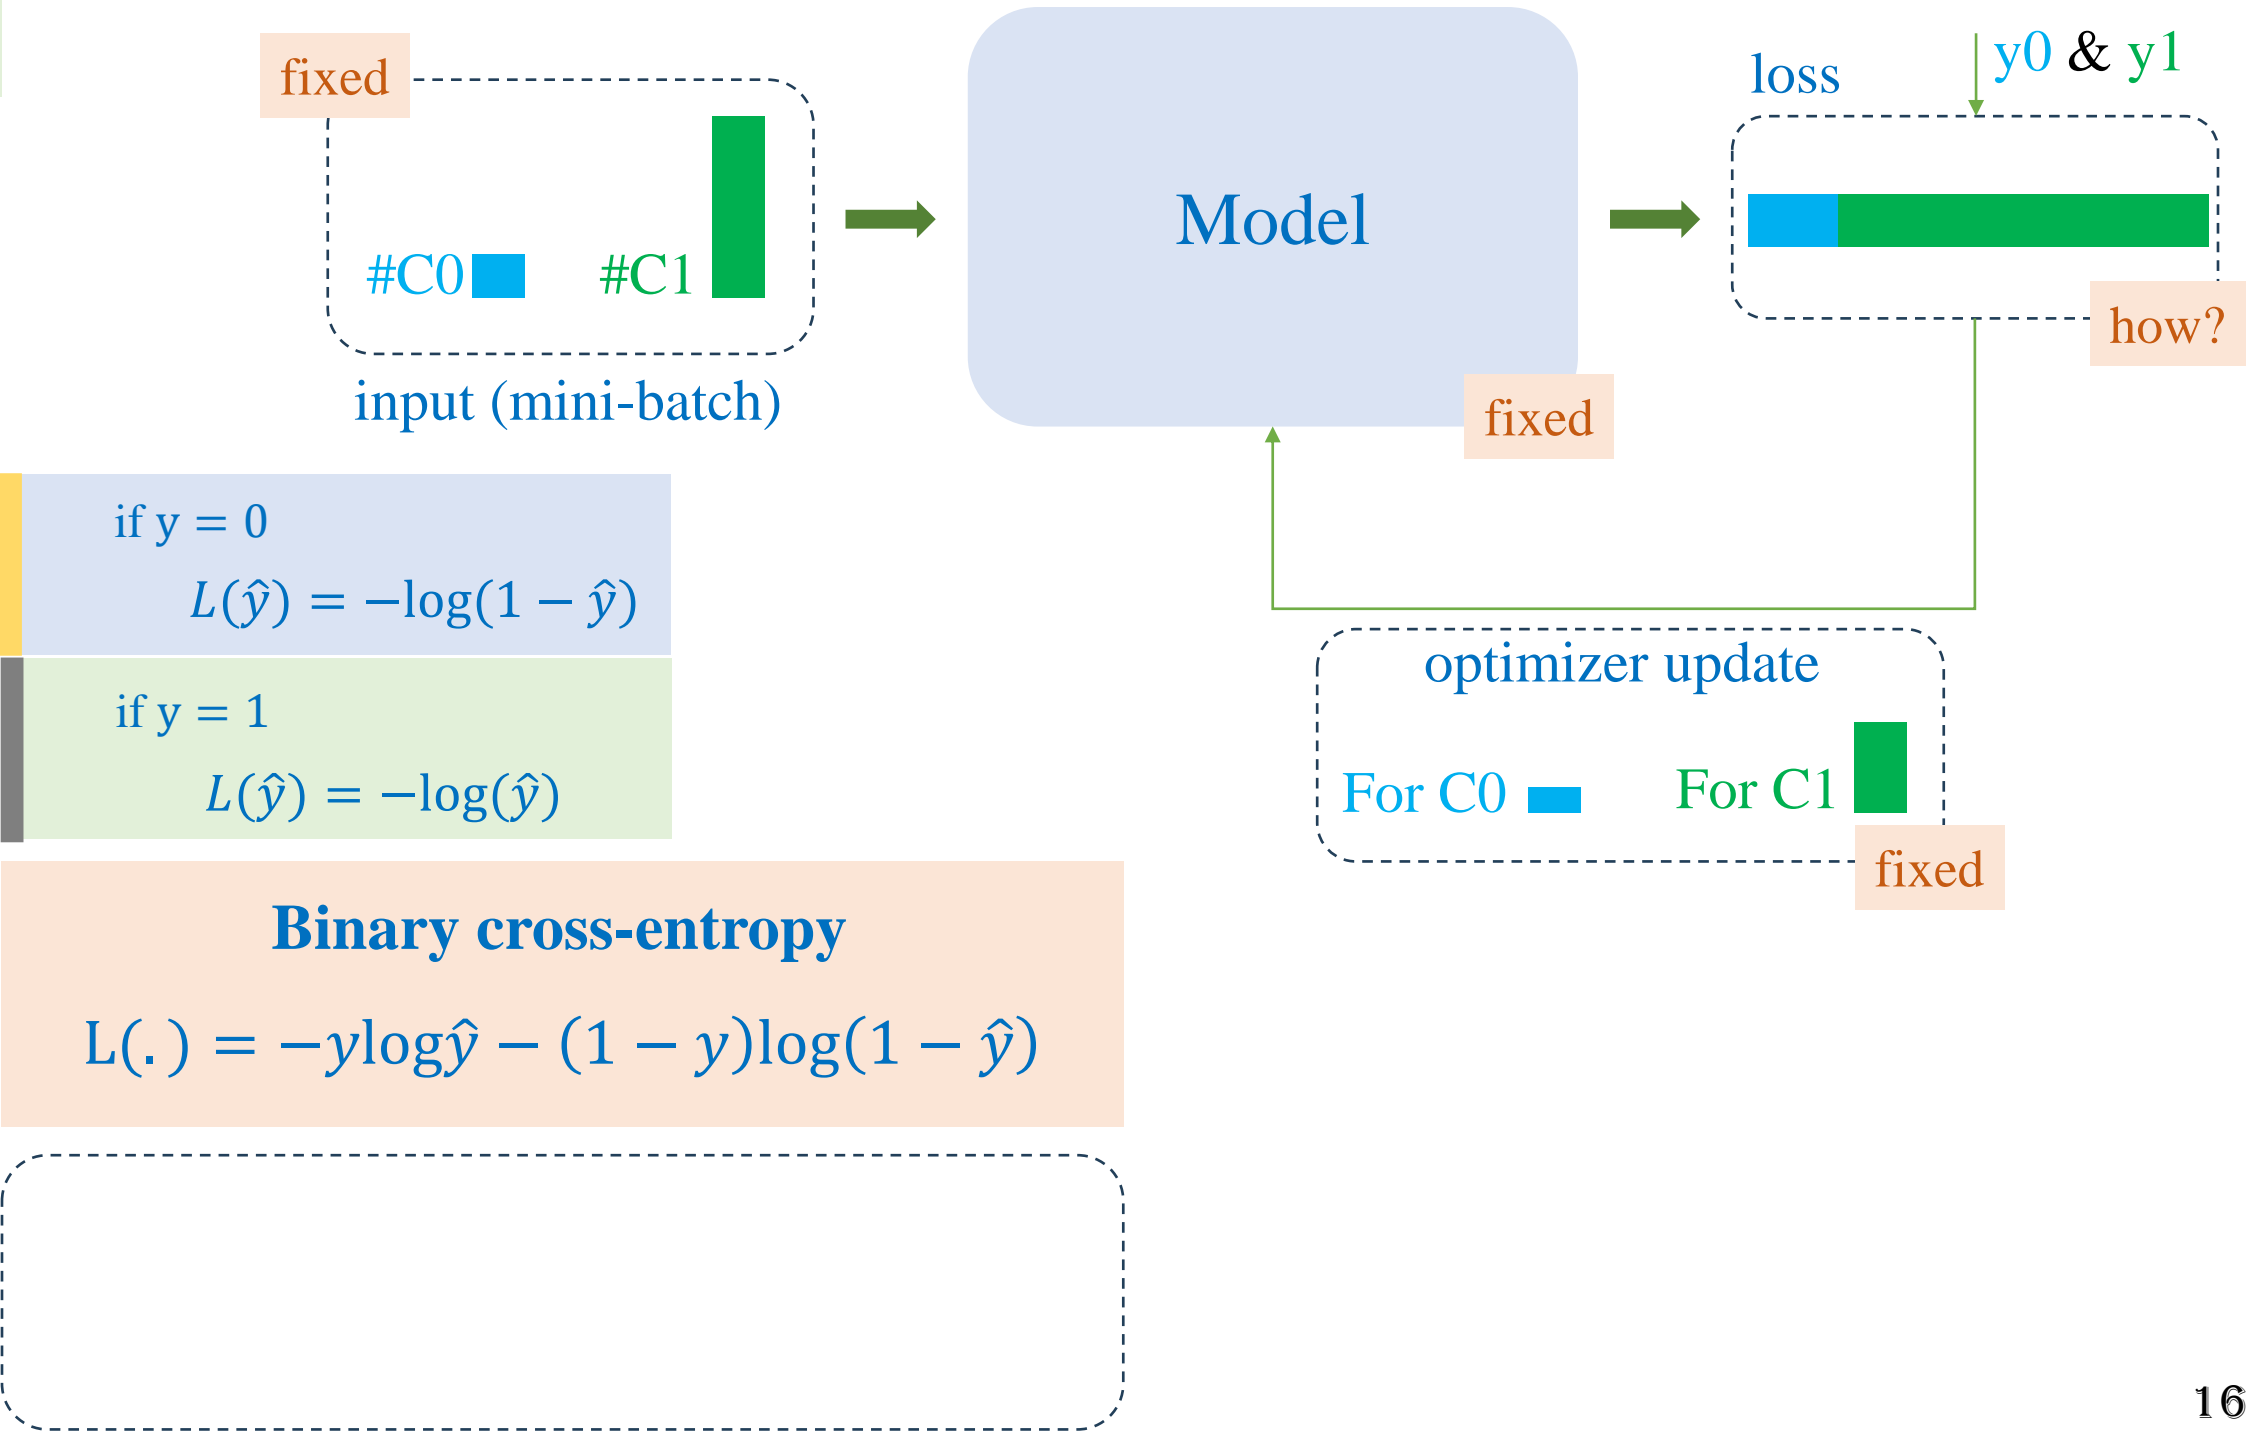
Original Data



Undersampling Data





Oversampling Data



Imbalanced Data

❖ Lie/Truth classification

Feature	Output	Label
Input	Output	Label
...	0.7	0
...	0.8	1
...	0.7	1
...	0.8	1
...	0.8	1
...	0.8	1
...	0.9	1
...	0.9	1
...	0.8	1
...	0.7	1

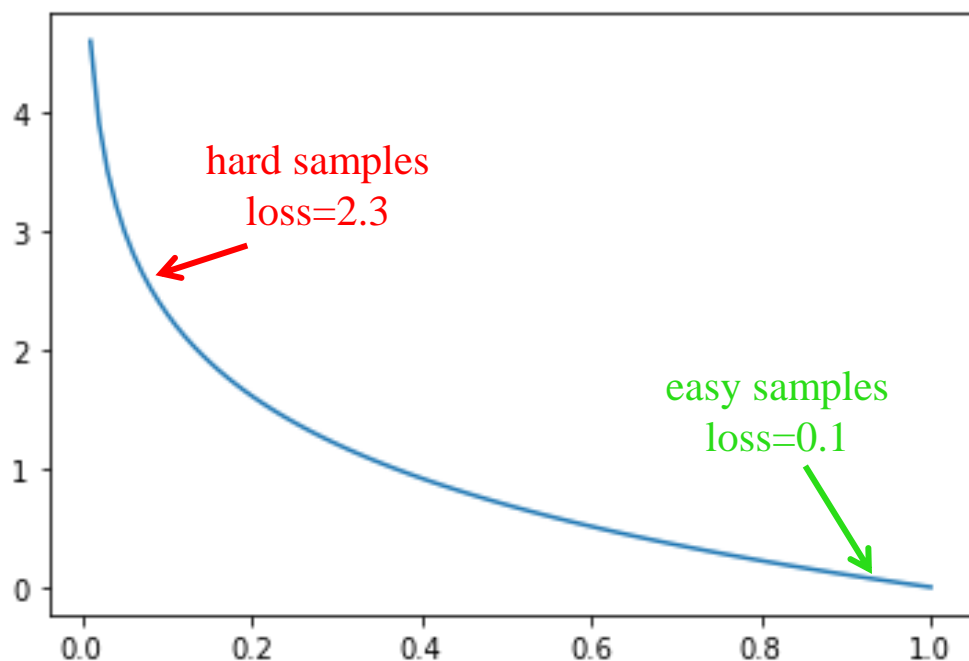


Class-weighted Binary cross-entropy

$$L(.) = -\alpha_1 y \log \hat{y} - \alpha_2 (1 - y) \log(1 - \hat{y})$$

Imbalanced Data

❖ Imbalance data



$$\text{Easy samples loss} = 100000 * 0.1 = 10000$$

$$\text{Hard samples loss} = 100 * 2.3 = 230$$

$$\text{Loss} = \text{Easy samples loss} + \text{Hard samples loss}$$







$$\text{Easy samples loss} : \text{Hard samples loss} = 10000 : 230 \approx 43$$

BCE không tốt cho trường hợp data bị imbalance nặng


Imbalance Case:

- 100000 easy samples vs 100 hard samples

How to solve it!!!




Designing a Function





Given $0 \leq k \leq 1$


if $f(\hat{y}) * k$
where k approaches 1
 $\rightarrow f(\hat{y}) * k$ reduces slightly



if $f(\hat{y}) * k$
where k approaches 0
 $\rightarrow f(\hat{y}) * k$ reduces
significantly




Designing a Function



Given $0 \leq k \leq 1$



if $f(\hat{y}) * k$
where k approaches 1
 $\rightarrow f(\hat{y}) * k$ reduces slightly







if $f(\hat{y}) * k$
where k approaches 0
 $\rightarrow f(\hat{y}) * k$ reduces significantly



Reducing significantly for
the correct part

Reducing slightly for the
incorrect part







Applying to our Problem

$$L(y, \hat{y}, \gamma) = -y(1 - \hat{y})^\gamma \log \hat{y} - (1 - y)\hat{y}^\gamma \log(1 - \hat{y})$$

Input	Output	Label	Gamma=0	Gamma=1	Gamma=2	Gamma=3	Gamma=4
...	0.7	0	1.204	0.842	0.589	0.412	0.289
...	0.8	1	0.223	0.044	0.008	0.001	0.0003
...	0.7	1	0.356	0.107	0.032	0.009	0.002
...	0.8	1	0.223	0.044	0.008	0.001	0.0003
...	0.8	1	0.223	0.044	0.008	0.001	0.0003
...	0.8	1	0.223	0.044	0.008	0.001	0.0003
...	0.9	1	0.105	0.011	0.001	0.0001	0.00001
...	0.9	1	0.105	0.011	0.001	0.0001	0.00001
...	0.8	1	0.223	0.044	0.008	0.001	0.0003
...	0.7	1	0.356	0.107	0.032	0.009	0.002

$loss_{y=1}$: 2.039 0.458 0.111 0.028 0.007





Combine with Class Weight

$$L(.) = -\alpha_1 y(1 - \hat{y})^\gamma \log \hat{y} - \alpha_2 (1 - y) \hat{y}^\gamma \log(1 - \hat{y})$$


Input	Output	Label	Gamma=0	Gamma=1	Gamma=2	Gamma=3	Gamma=4
...	0.7	0	1.204	0.842	0.589	0.412	0.289
...	0.8	1	0.223	0.044	0.008	0.001	0.0003
...	0.7	1	0.356	0.107	0.032	0.009	0.002
...	0.8	1	0.223	0.044	0.008	0.001	0.0003
...	0.8	1	0.223	0.044	0.008	0.001	0.0003
...	0.8	1	0.223	0.044	0.008	0.001	0.0003
...	0.9	1	0.105	0.011	0.001	0.0001	0.00001
...	0.9	1	0.105	0.011	0.001	0.0001	0.00001
...	0.8	1	0.223	0.044	0.008	0.001	0.0003
...	0.7	1	0.356	0.107	0.032	0.009	0.002

$loss_{y=1}$: 2.039 0.458 0.111 0.028 0.007




solution 5

Parameter
Space




solution 5

convergence 1




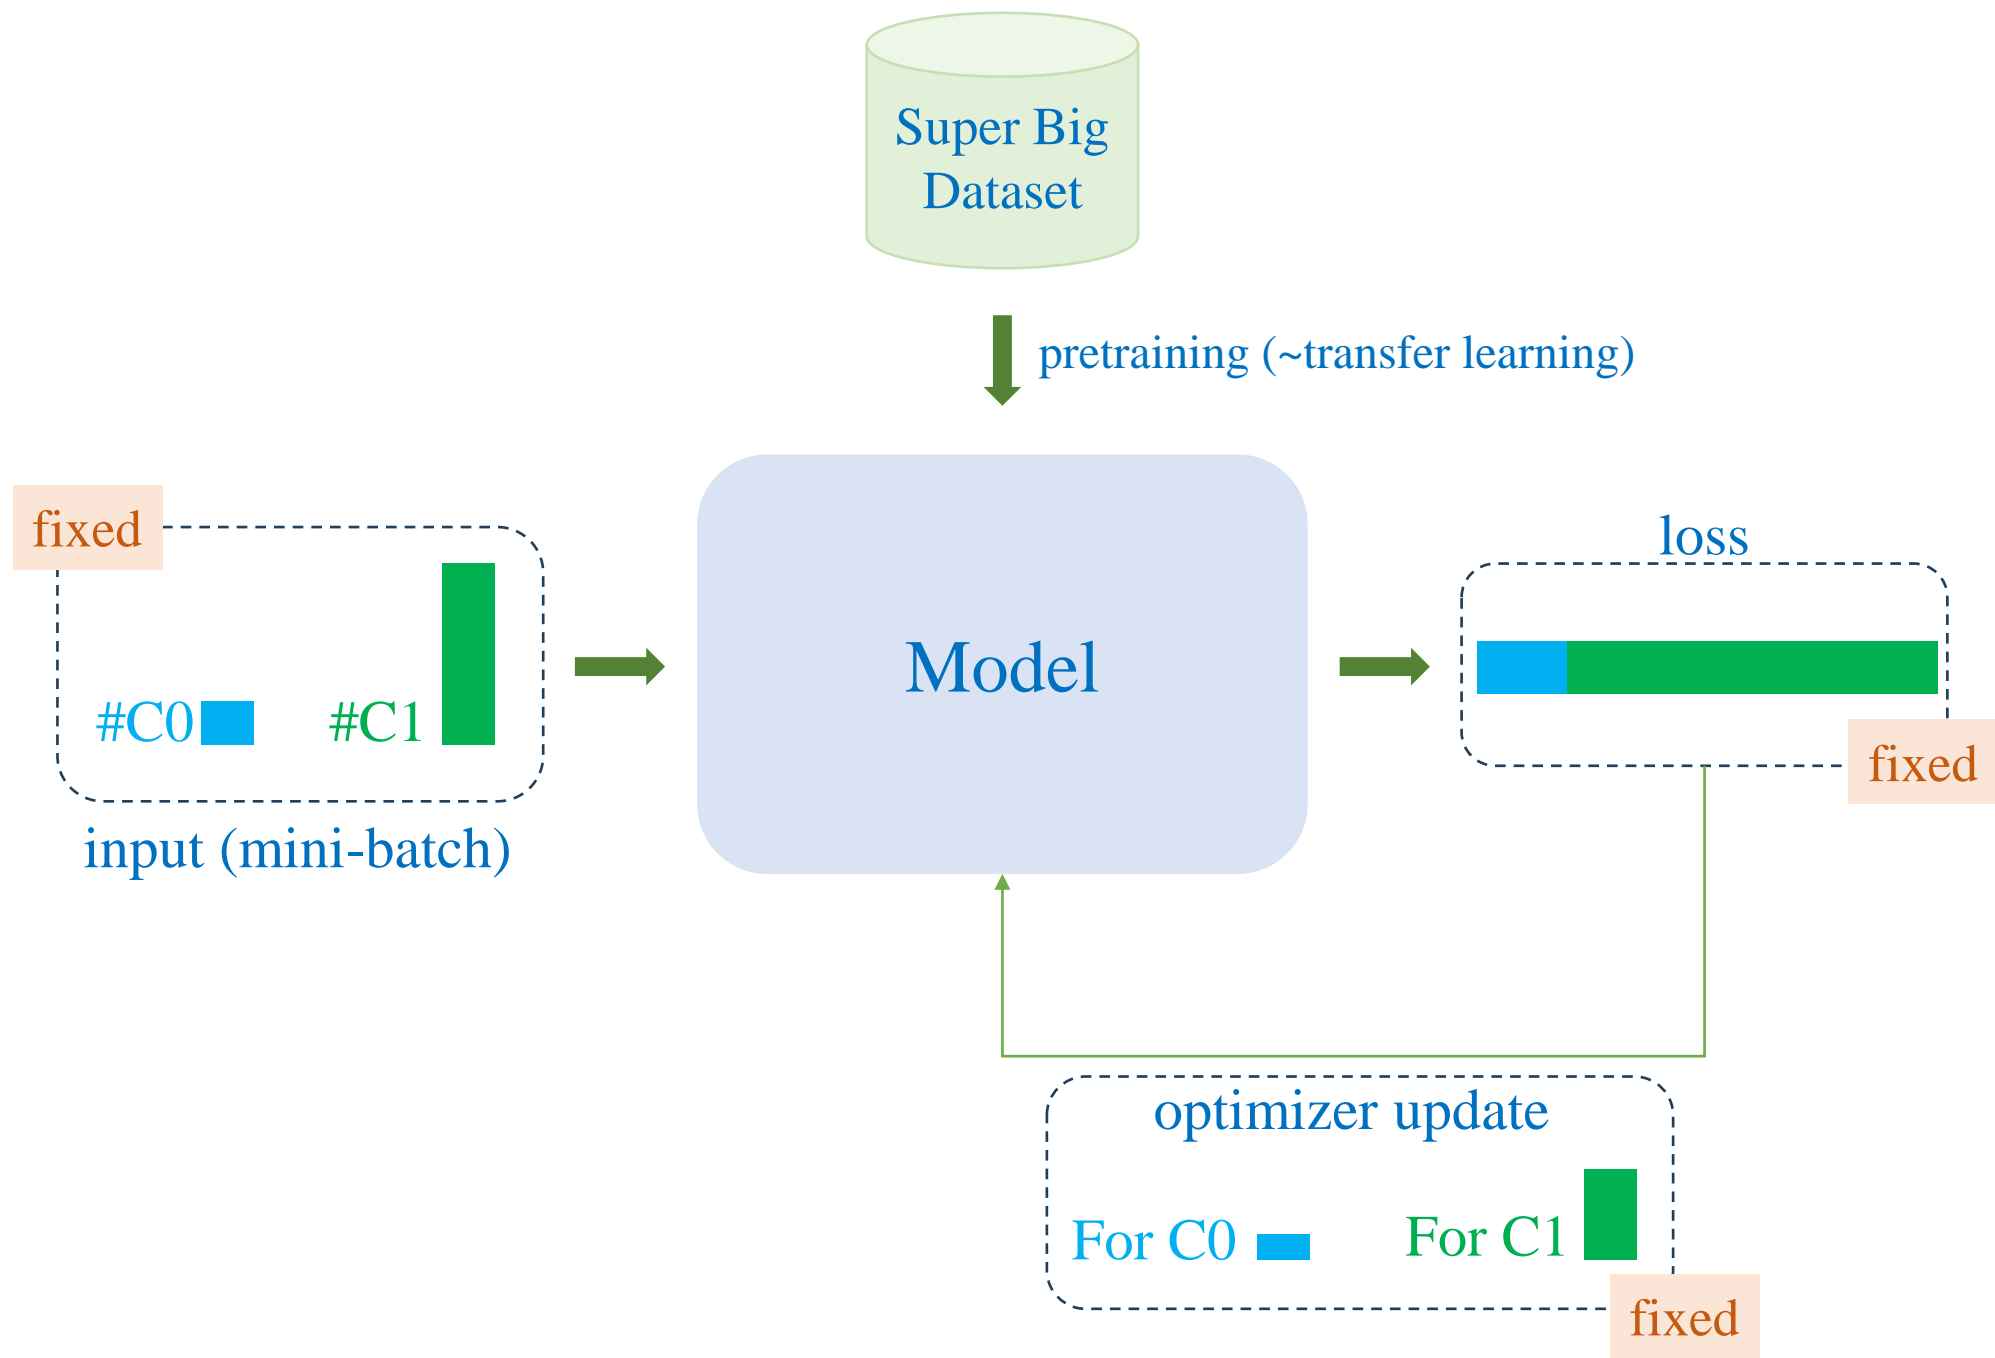
solution 5

convergence 2



Using
Pretrained
Model

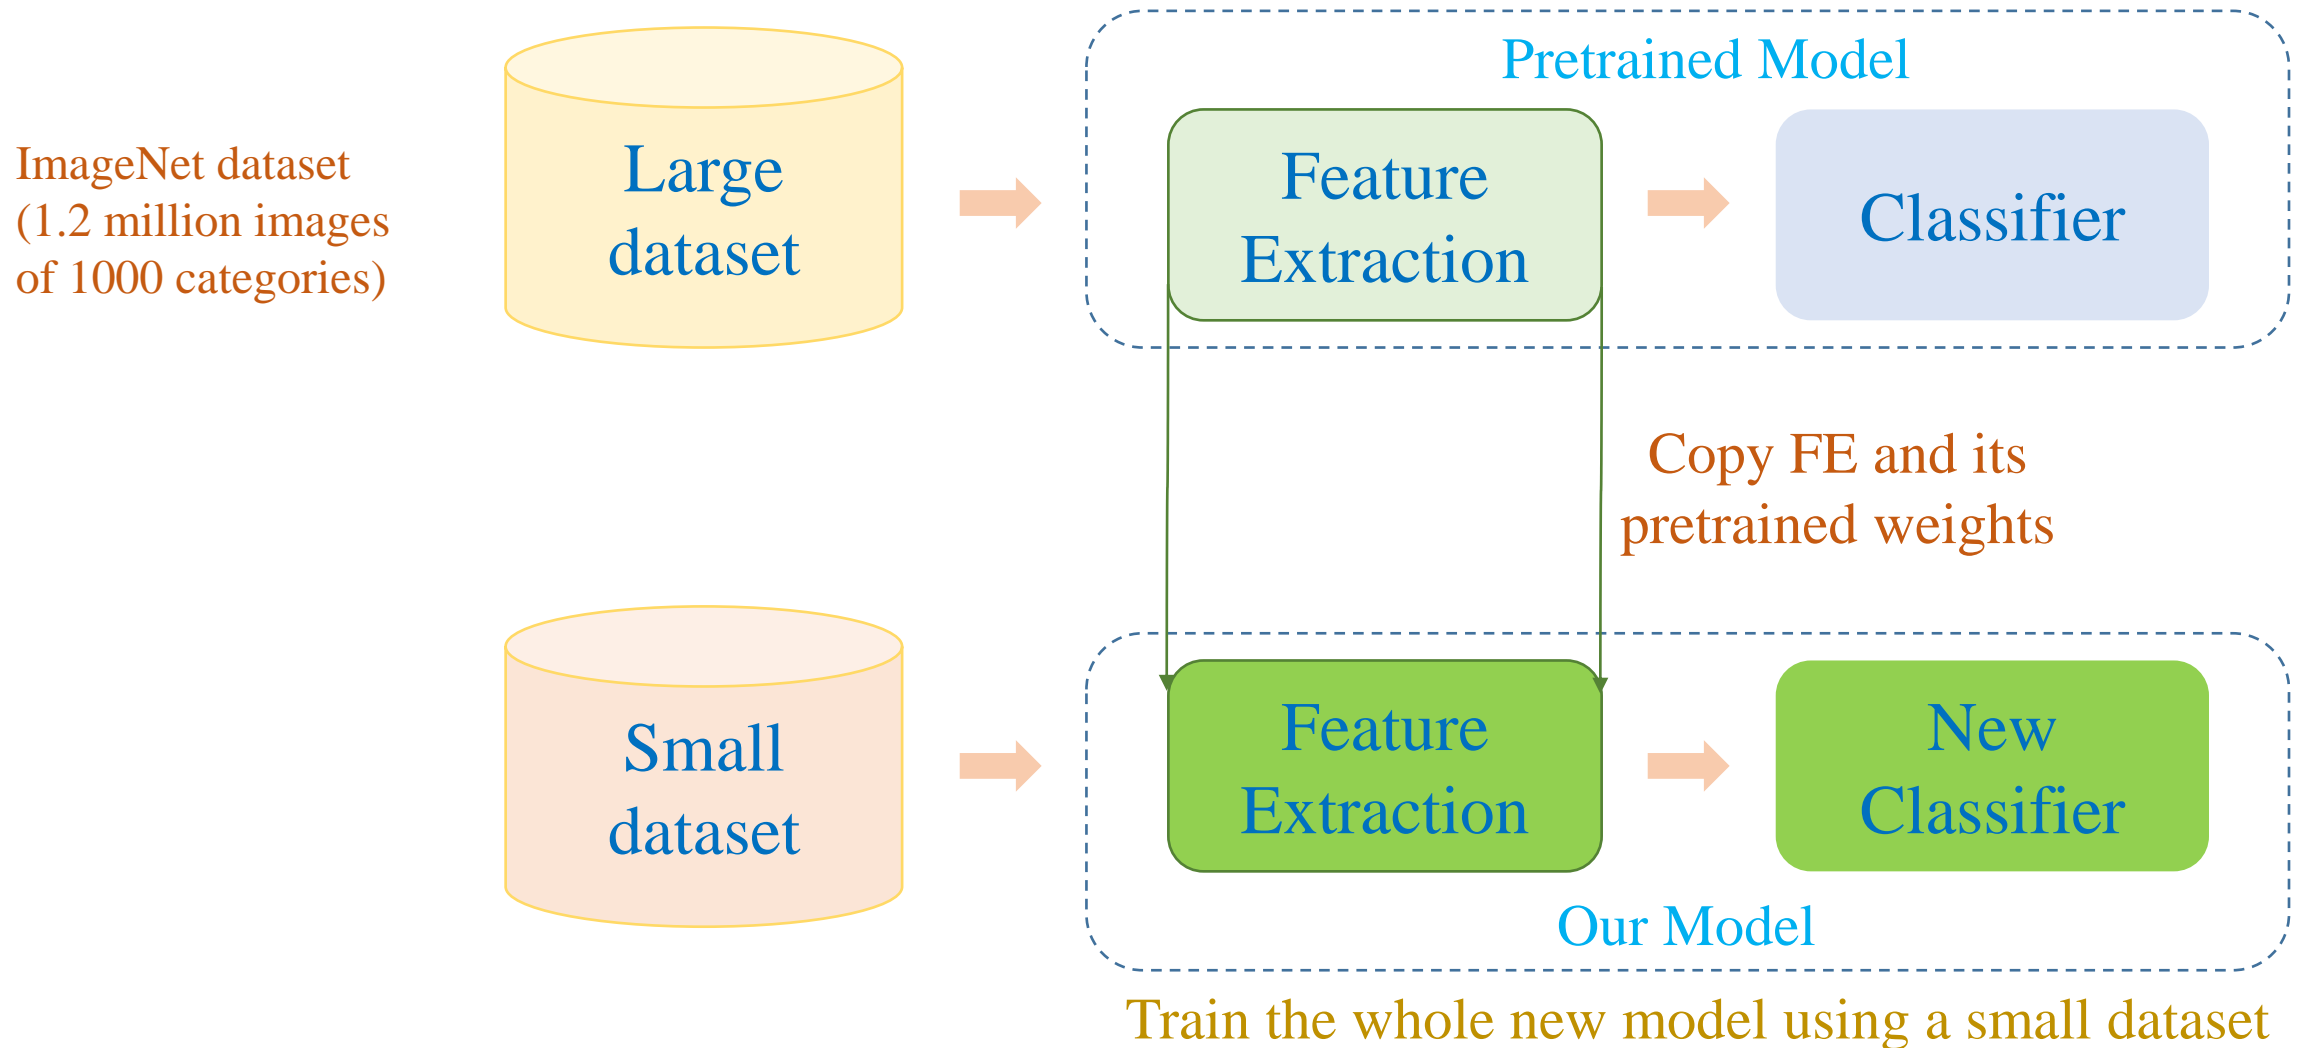




Exploitation of Pretrained Models

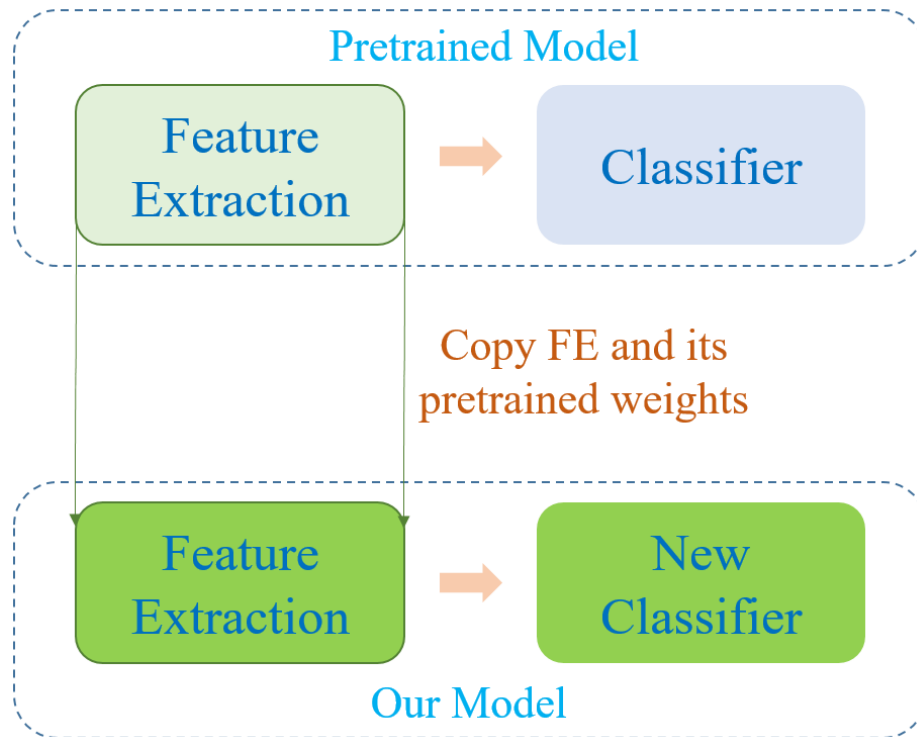
❖ As Initialization

■ Will be trained with the small dataset



Exploitation of Pretrained Models


❖ Use the pretrained weights as an initialization



```
1 # Load the pretrained VGG16 model
2 vgg16 = models.vgg16(weights=models.VGG16_Weights.DEFAULT)
3 f_extractor = vgg16.features
4
5 model = nn.Sequential(f_extractor,
6                       nn.Flatten(),
7                       nn.Dropout(0.3),
8                       nn.Linear(512*7*7, 512),
9                       nn.ReLU(),
10                      nn.Dropout(0.3),
11                      nn.Linear(512, 2))
12
13 import torchvision.models as models
14 import torch.nn as nn
15
16 # Load Resnet18 model and pre-train weight
17 model = models.resnet18(weights="IMAGENET1K_V1")
18 model.fc = nn.Linear(512, 2, bias=True)
```

Imbalanced Data

❖ Approach 2: Loss Functions



Class weight

$$w_c = \frac{N}{2N_c}$$

Pay more attention to samples from an under-represented class

A higher loss \rightarrow higher optimization

Focal loss

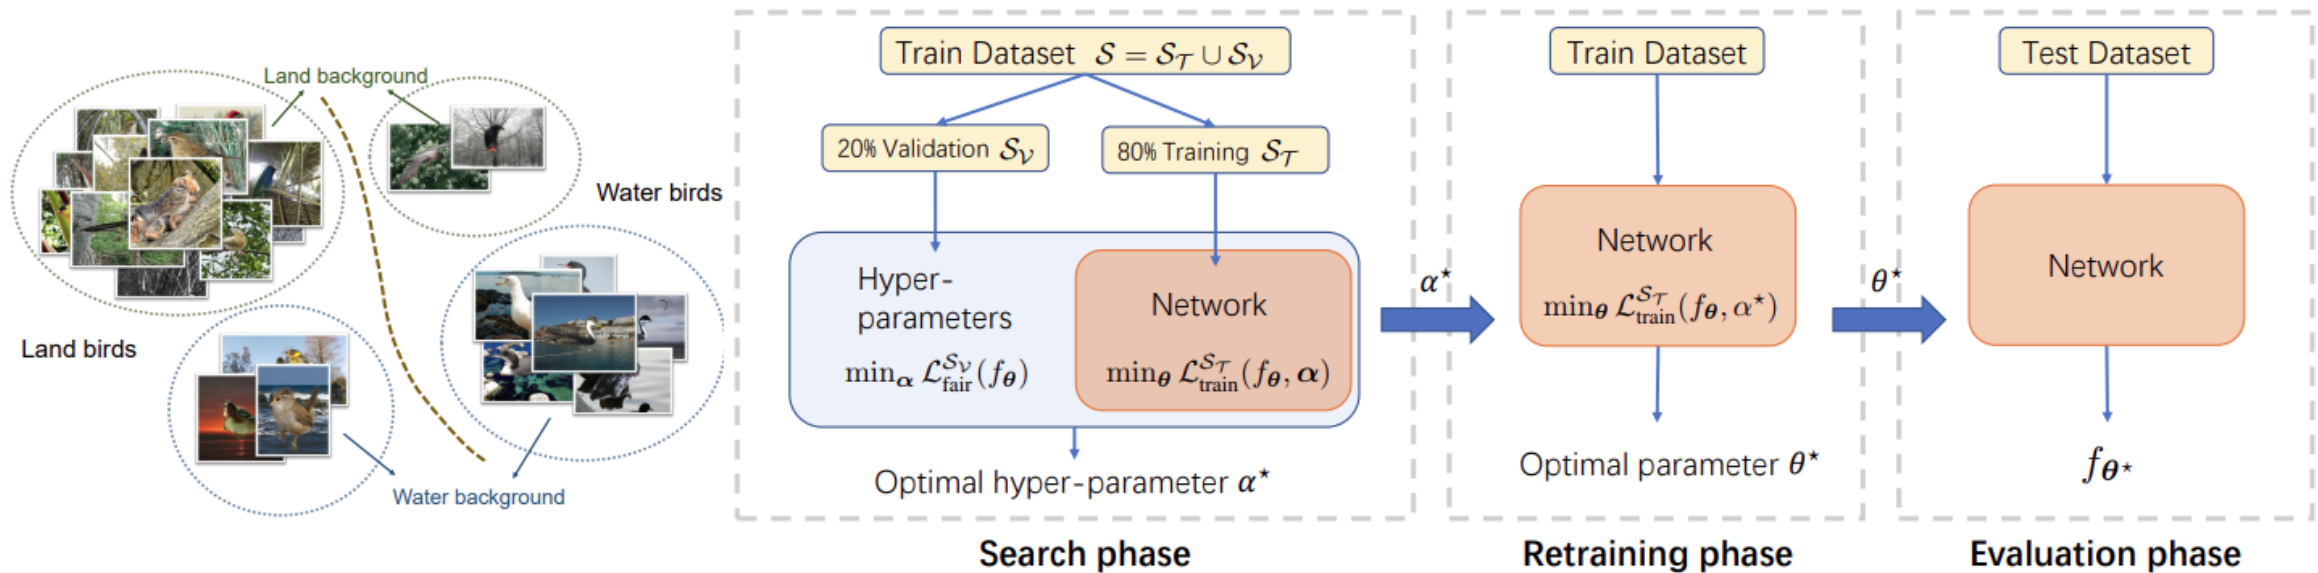
$$\text{FL}(p_t) = -\alpha_t(1 - p_t)^\gamma \log(p_t)$$

γ helps the loss function focus on 'hard' samples

α_t balances losses using a number of samples in a class

Imbalanced Data

❖ Solution 6



<https://arxiv.org/pdf/2201.01212.pdf>


AutoBalance: Optimized Loss Functions for Imbalanced Data, 2022

Outline


- **Introduction**
- **Examples and Discussion**
- **Metrics**
- **Case Study**

Metrics

❖ Confusion matrix




Model



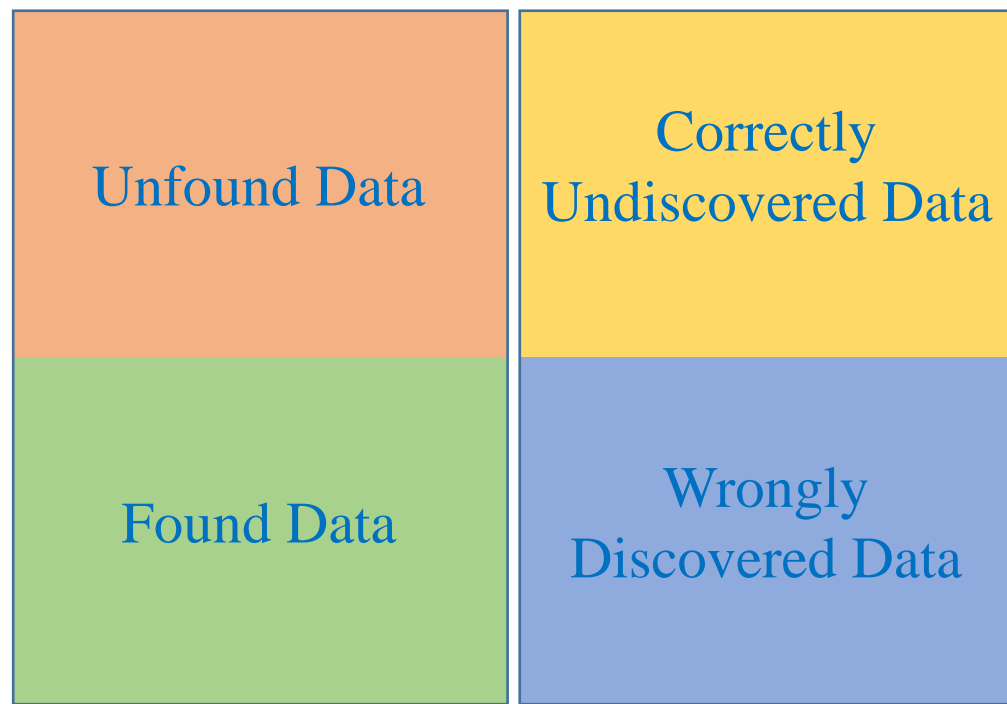
Positive data

Negative data

1

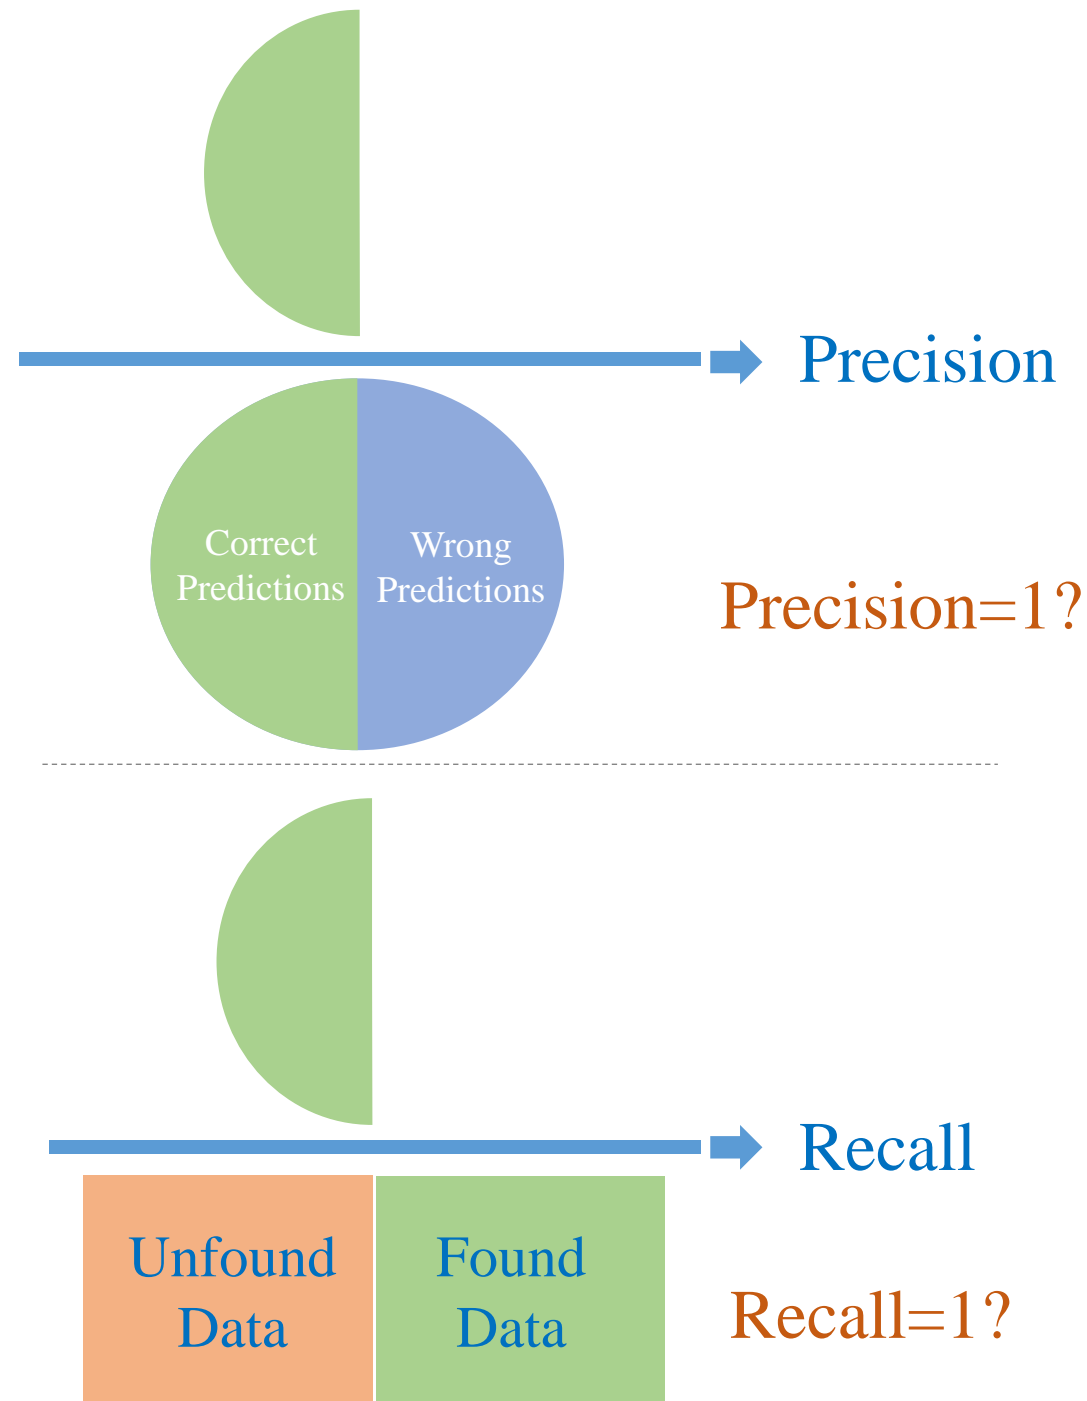


2



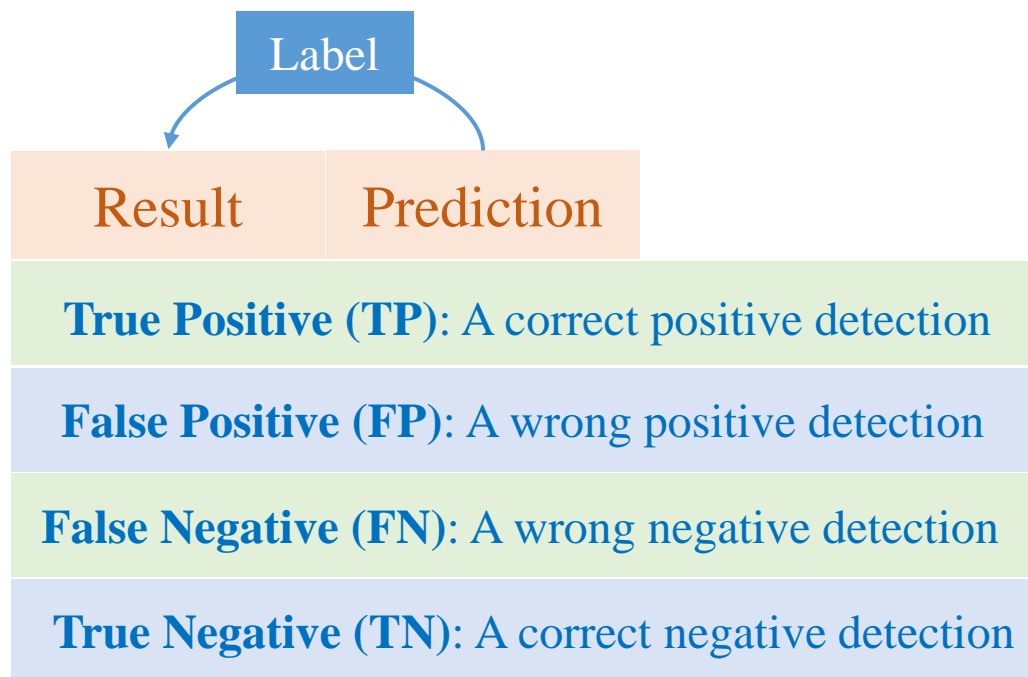
Positive data

Negative data



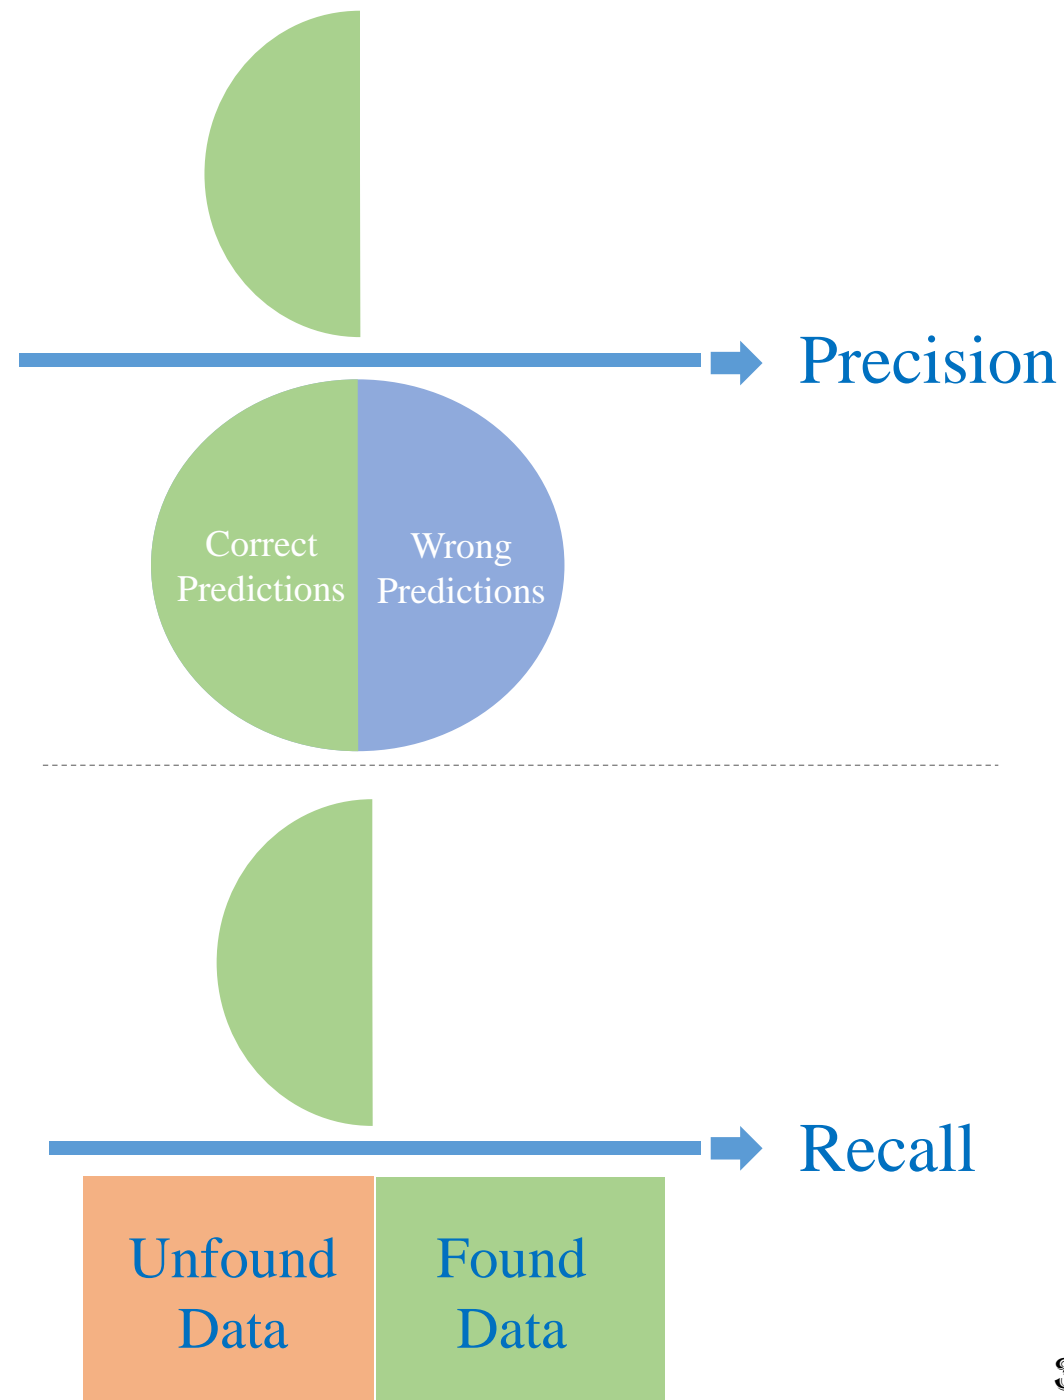
Metrics

❖ Precision and recall



$$\text{Precision} = \frac{TP}{TP + FP} = \frac{\text{Số dự đoán chính xác}}{\text{Tổng số lần dự đoán positive}}$$

$$\text{Recall} = \frac{TP}{TP + FN} = \frac{\text{Số dự đoán chính xác}}{\text{Tổng số ground truth cho positive}}$$



Metrics

❖ Quiz

Fill TP, FP, TN, FN into appropriate cells

	Actual: NEGATIVE	Actual: POSITIVE
Predicted: NEGATIVE		
Predicted: POSITIVE		



Metrics

❖ Accuracy


$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

❖ Precision

$$\text{Precision} = \frac{TP}{TP + FP} = \frac{TP}{\text{all positive decisions}}$$

❖ Recall

$$\text{Recall} = \frac{TP}{TP + FN} = \frac{TP}{\text{all positive samples}}$$



Metrics

❖ Accuracy

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

❖ Precision

$$\text{Precision} = \frac{TP}{TP + FP} = \frac{TP}{\text{all positive decisions}}$$

❖ Recall

$$\text{Recall} = \frac{TP}{TP + FN} = \frac{TP}{\text{all positive samples}}$$

n = 165	Predicted:		
	NEGATIVE	POSITIVE	
Actual:	NEGATIVE	TN=50	60
	POSITIVE	FN=10	105
		60	105

Accuracy = 145/165

Precision: When it predicts yes, how often is it correct?

$$\text{TP/predicted yes} = 95/105 = 0.9$$

Recall: When it's actually yes, how often does it predict yes?

$$\text{TP/actual yes} = 95/105 = 0.90$$

Metrics

❖ Accuracy

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

n = 165

	Predicted: NEGATIVE	Predicted: POSITIVE	
Actual: NEGATIVE	TN=60	FP=0	60
Actual: POSITIVE	FN=20	TP=85	105
	80	85	

Accuracy = 145/165

Precision: When it predicts yes, how often is it correct?

$$TP/\text{predicted yes} = 85/85 = 1.0$$

Recall: When it's actually yes, how often does it predict yes?

$$TP/\text{actual yes} = 85/105 = 0.81$$

❖ Precision

$$\text{Precision} = \frac{TP}{TP + FP} = \frac{TP}{\text{all positive decisions}}$$

❖ Recall

$$\text{Recall} = \frac{TP}{TP + FN} = \frac{TP}{\text{all positive samples}}$$

Metrics

❖ Accuracy

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$n = 165$

	Predicted: NEGATIVE	Predicted: POSITIVE	
Actual: NEGATIVE	TN=40	FP=20	60
Actual: POSITIVE	FN=0	TP=105	105
	60	105	

Accuracy = 145/165

Precision: When it predicts yes, how often is it correct?

$TP/\text{predicted yes} = 105/125 = 0.84$

Recall: When it's actually yes, how often does it predict yes?

$TP/\text{actual yes} = 105/105 = 1.0$

❖ Precision

$$\text{Precision} = \frac{TP}{TP + FP} = \frac{TP}{\text{all positive decisions}}$$

❖ Recall



$$\text{Recall} = \frac{TP}{TP + FN} = \frac{TP}{\text{all positive samples}}$$

Metrics

❖ Combine precision and recall

$$F_1 = \frac{\text{precision} + \text{recall}}{2}$$

$$\frac{2}{F_1} = \frac{1}{\text{precision}} + \frac{1}{\text{recall}}$$



Metrics

❖ F1 Score

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$


$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

$$\text{F1} = \frac{2 * \text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

$$\frac{2}{F_1} = \frac{1}{\text{precision}} + \frac{1}{\text{recall}}$$


$$F_1 = 2 \frac{1}{\frac{1}{\text{precision}} + \frac{1}{\text{recall}}} = 2 \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

Precision	Recall	F1
1	1	1
0.1	0.1	0.1
0.5	0.5	0.5
1	0.1	0.182
0.3	0.8	0.36



Example

❖ Cat-Dog dataset



Cat is of the positive class

Correct prediction

$$\#_{cat} = 821$$

$$\#_{dog} = 1489$$

Validation data (3000 samples)

n=3000	Predicted: NEGATIVE	Predicted: POSITIVE	
	Actual: NEGATIVE	Actual: POSITIVE	
	TN=1489	FP=11	1500
	FN=679	TP=821	1500
	2168	832	

$$\text{Recall} = \frac{TP}{TP + FN} = \frac{821}{821 + 679} \approx 0.547$$

$$\text{Precision} = \frac{TP}{TP + FP} = \frac{821}{821 + 11} \approx 0.987$$

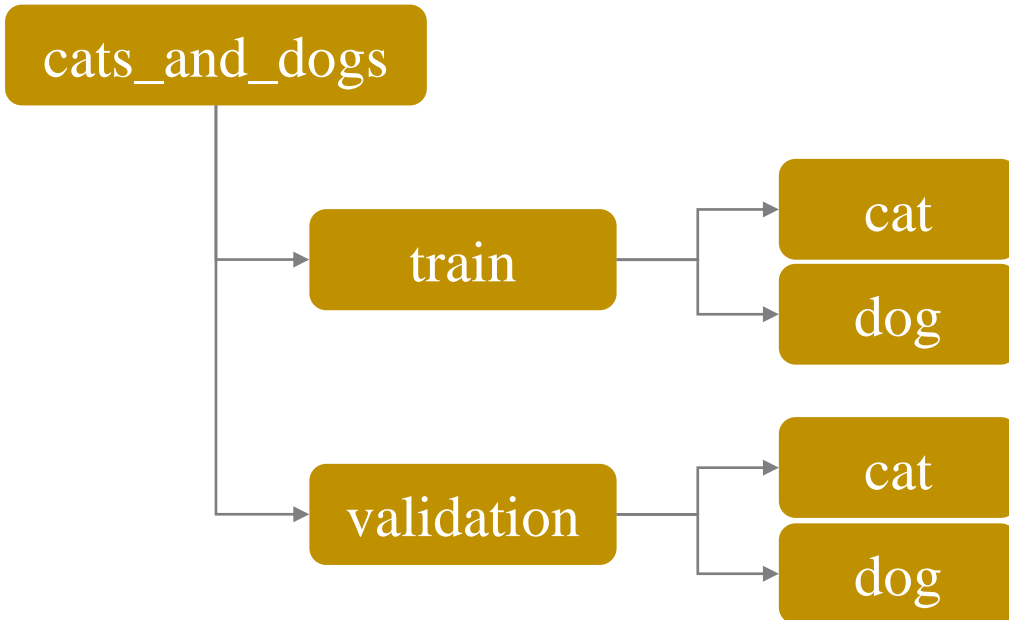
$$F1 = \frac{2 * \text{Recall} * \text{Precision}}{\text{Recall} + \text{Precision}} = \frac{2 * 0.547 * 0.987}{0.547 + 0.987} \approx 0.704$$

Outline

- **Introduction**
- **Examples and Discussion**
- **Metrics**
- **Case Study**

Experiments

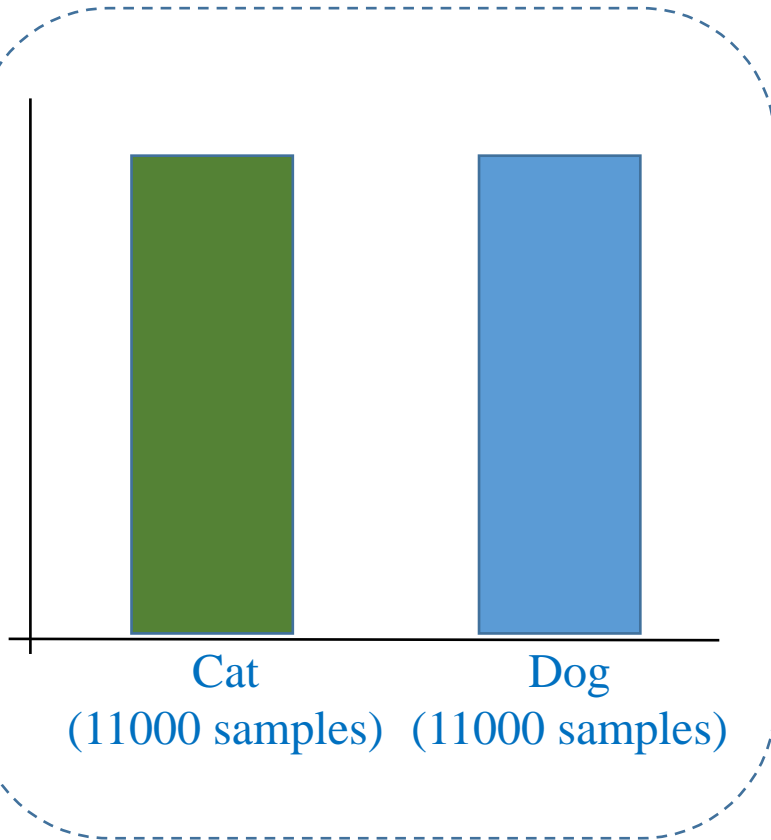
❖ Cat-Dog dataset




Experiments

❖ Cat-Dog dataset


■ Validation data (3000 samples)



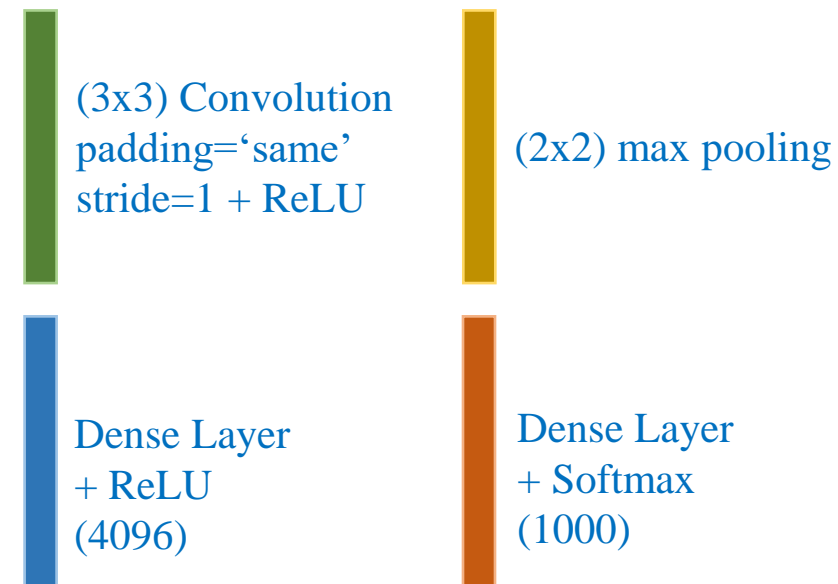

Balanced Data



Imbalanced Data 1




Imbalanced Data 2



VGG16


VGG-16



Experiments

❖ Cat-Dog dataset: Results from the TF codes

Validation data (3000 samples)



Balanced Loss

$$L_b = L_c + L_d$$

Correct prediction

$$\#_{cat} = 1445$$

$$\#_{dog} = 1436$$

$$F_1 = 0.96$$

Imbalanced Loss 1

$$L_b = L_c + 100 \times L_d$$

Correct prediction

$$\#_{cat} = 1076$$

$$\#_{dog} = 1499$$

$$F_1 = 0.835$$

Imbalanced Loss 2

$$L_b = L_c + 1000 \times L_d$$

Correct prediction

$$\#_{cat} = 670$$


$$\#_{dog} = 1498$$

$$F_1 = 0.617$$

Experiments

❖ Cat-Dog dataset

Results from the TF codes



Imbalanced Data 1

■ Validation data (3000 samples)

Balanced Loss

$$L_b = L_c + L_d$$

Correct prediction

$$\#_{cat} = 1082$$

$$\#_{dog} = 1483$$

$$F_1 = 0.833$$

Imbalanced Loss

$$L_b = 6 \times L_c + 0.55 \times L_d$$

Correct prediction

$$\#_{cat} = 1163$$

$$\#_{dog} = 1379$$

$$F_1 = 0.835$$

Oversampling Data

$$\#_{cat} = 1167$$

$$\#_{dog} = 1438$$

$$F_1 = 0.855$$

Focal loss

$$\#_{cat} = 1210$$

$$\#_{dog} = 1447$$


$$F_1 = 0.876$$

$$FL(p_t) = -\alpha_t(1 - p_t)^\gamma \log(p_t) \quad \gamma = 2.0 \quad \alpha_t = 0.5$$

Experiments

❖ Cat-Dog dataset

Results from the TF codes



Imbalanced Data 2

Validation data (3000 samples)

Balanced Loss

$$L_b = L_c + L_d$$

Correct prediction

$$\#_{cat} = 821$$

$$\#_{dog} = 1489$$

$$F_1 = 0.704$$

Imbalanced Loss

$$L_b = 11.5 \times L_c + 0.52 \times L_d$$


Correct prediction

$$\#_{cat} = 1123$$

$$\#_{dog} = 1309$$

$$F_1 = 0.798$$

Oversampling



Correct prediction

$$\#_{cat} = 1159$$

$$\#_{dog} = 1386$$

$$F_1 = 0.836$$

Focal Loss

$$FL(p_t) = -\alpha_t(1 - p_t)^\gamma \log(p_t)$$

$$\gamma = 2.0 \quad \alpha_t = 0.5$$

Correct prediction


$$\#_{cat} = 1210$$

$$\#_{dog} = 1447$$

$$F_1 = 0.876$$

Experiments

❖ Cat-Dog dataset: Using Resnet and Pytorch



Normal cross-entropy

$$F_1 = 0.67$$

Class weight

$$F_1 = 0.75$$

Focal Loss

$$F_1 = ???$$

Using the pretrained model

$$F_1 = 0.96$$

