AI VIETNAM
All-in-One Course

Generative Adversarial Networks

Quang-Vinh Dinh
Ph.D. in Computer Science

Year 2024

Objectives

- ✓ Study Generative Adversarial Network
- ✓ Study Deep Convolutional GAN

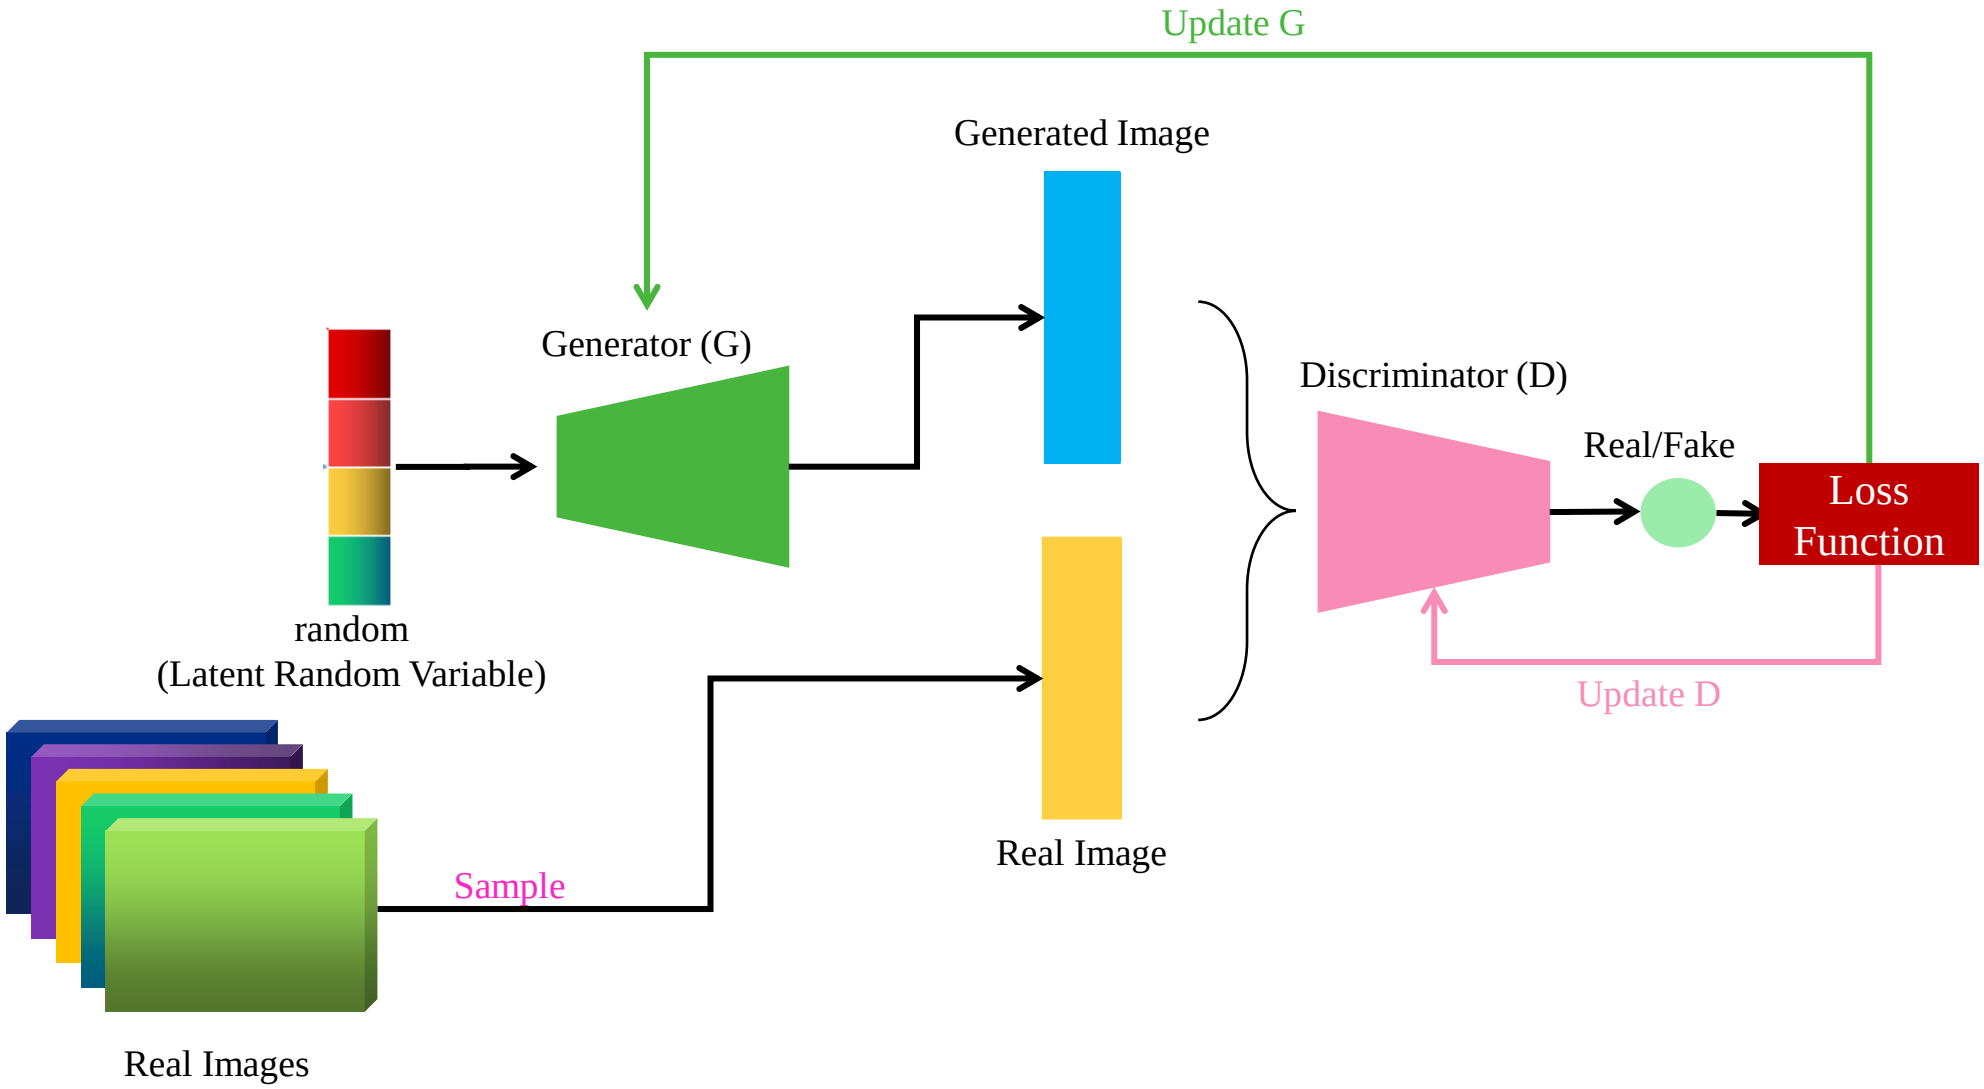

Outline

- Introduction
- GAN
- DCGAN
- Implementation

Introduction

❖ Introduction

A normal case

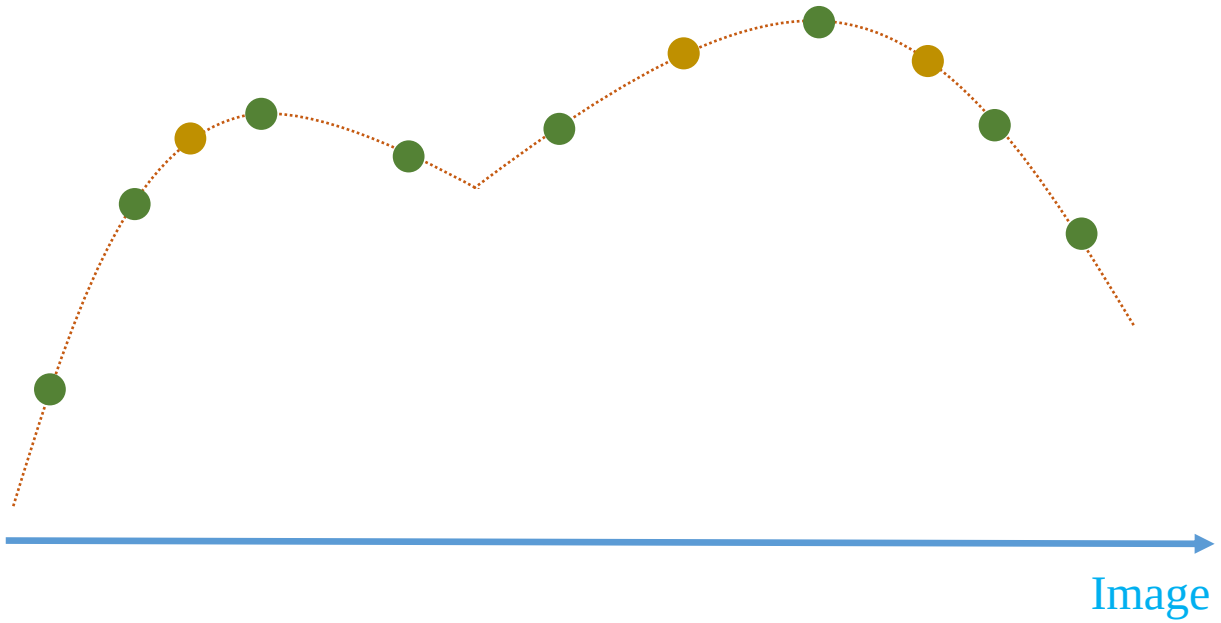

A perfect case: Have unlimited training

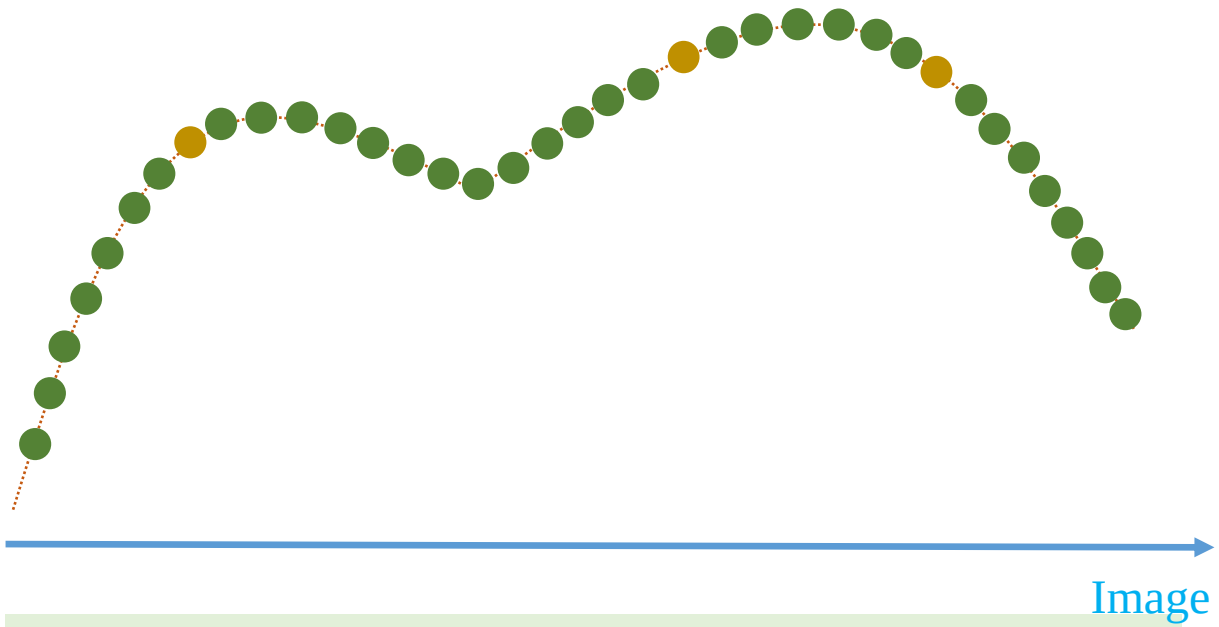

Training data cover the whole distribution

But, impractical!!!

Introduction

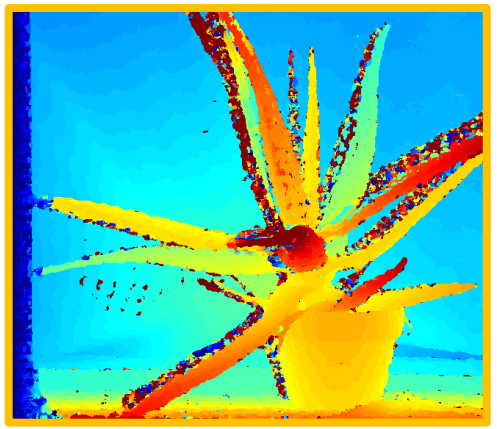

Images

Input Space

Network (Unet)

Images

Output Space

AI VIETNAM
All-in-One Course

Introduction

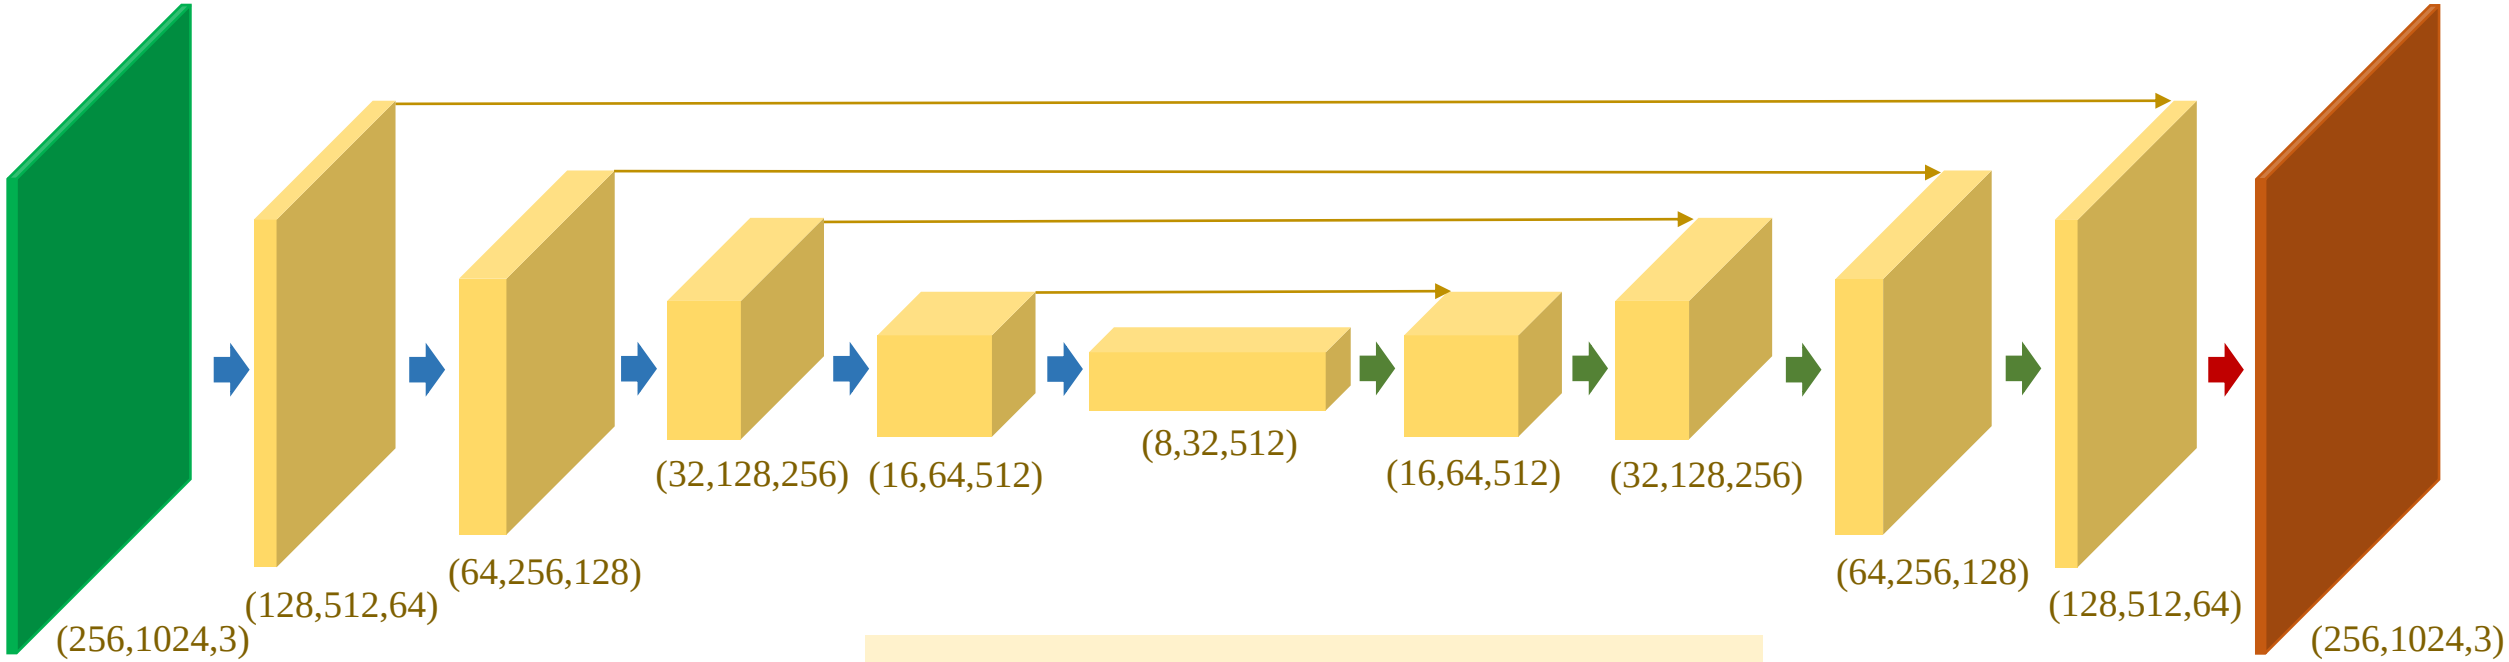

Input Image
[-1,1]

Skip connection (UNet)

Output Image
[-1,1]

(3x3) Convolution padding = 'same' stride = 1 + ReLU + Batch Norm + (2x2) max pooling

Feature Map Upsamping + (3x3) Convolution padding = 'same' stride = 1 + ReLU + Batch Norm

Feature Map Upsamping + (3x3) Convolution padding = 'same' stride = 1 + Tanh

Introduction

MNIST dataset

- Grayscale images
- Resolution=28x28
- Training set: 60000 samples
- Testing set: 10000 samples

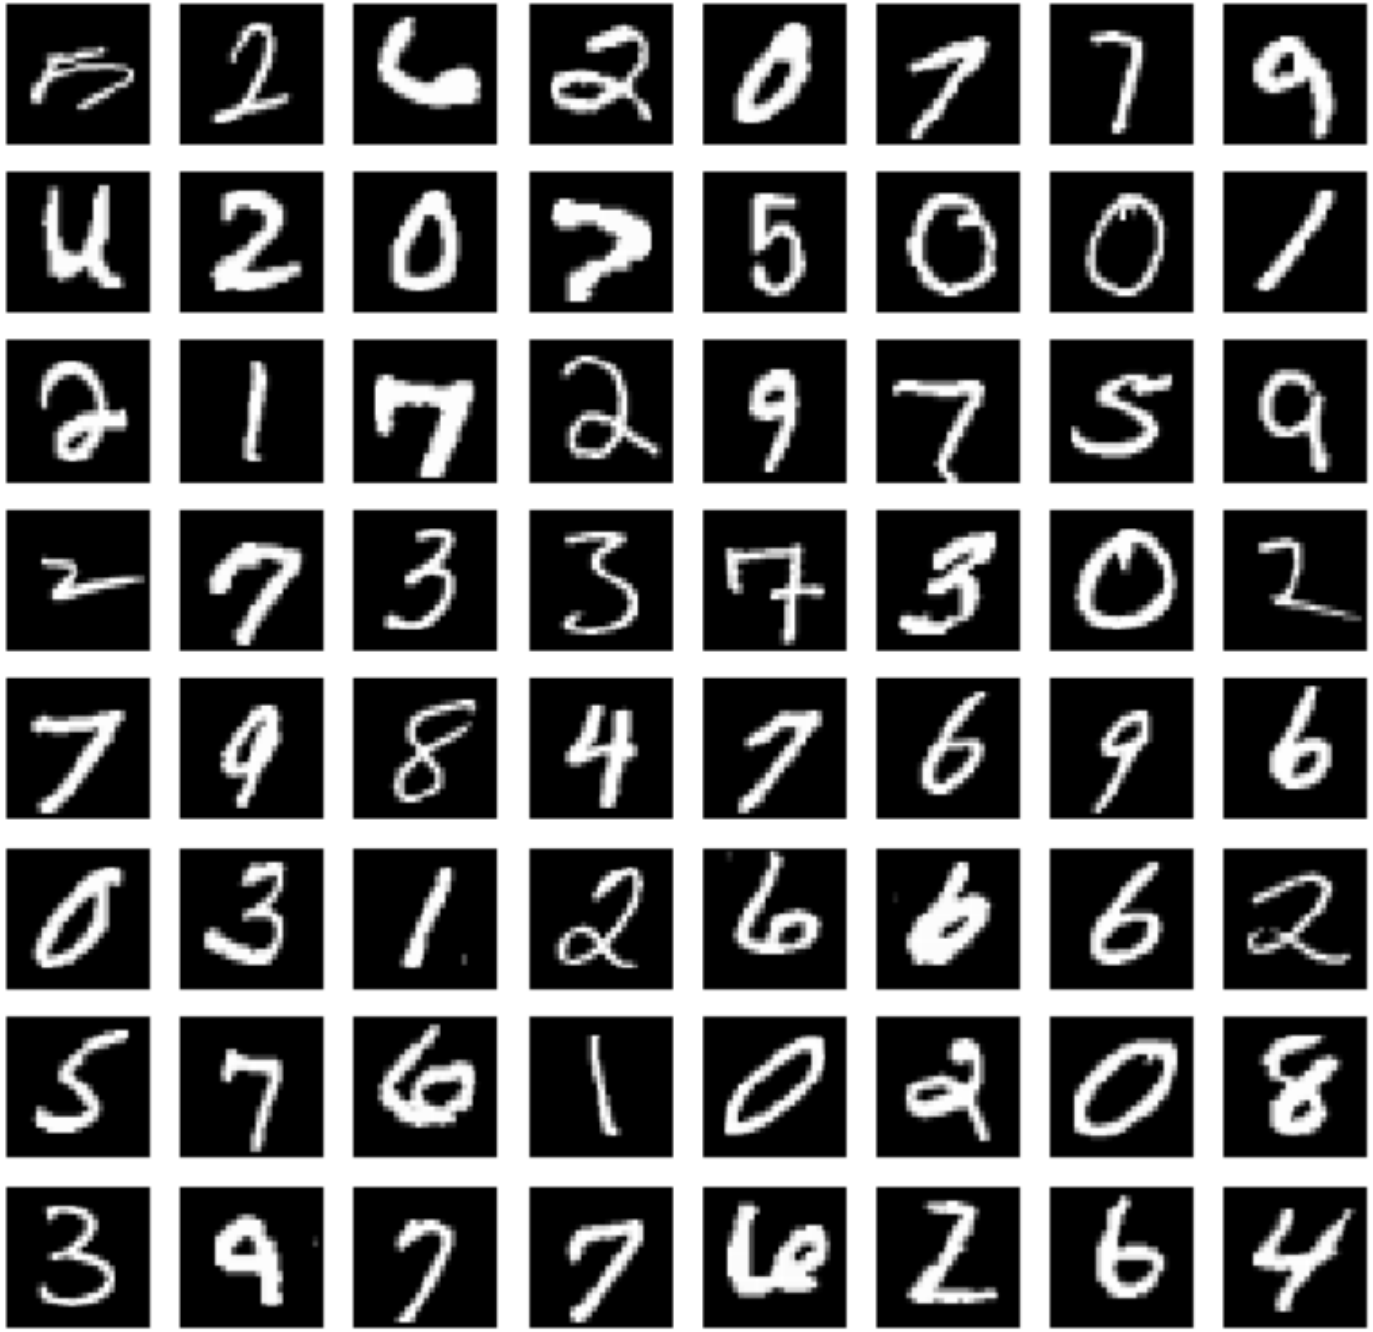

Introduction

❖ Input?

A perfect case: Have unlimited training

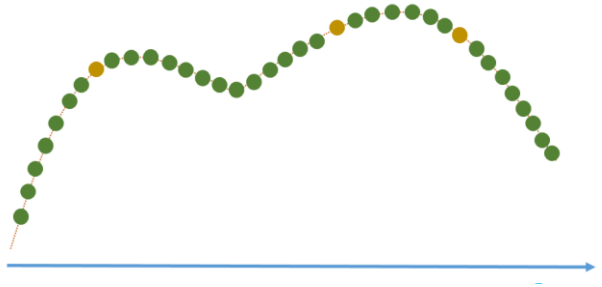

?

➡

Data Generation Network

➡

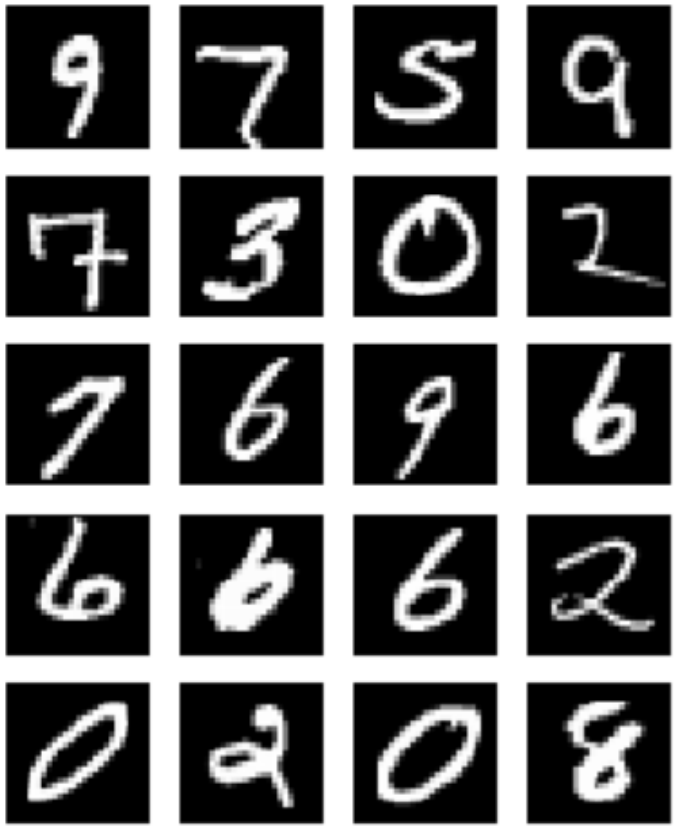

Generative Adversarial Networks

❖ Input?

A perfect case: Have unlimited training

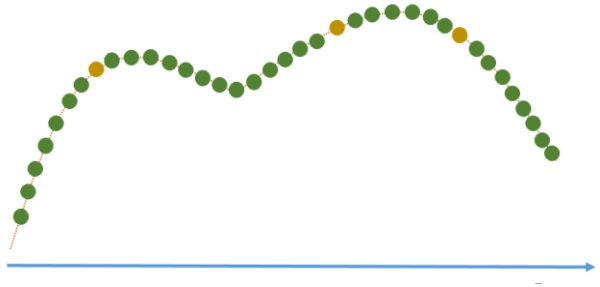

Random samples

➡

Data Generation Network

➡

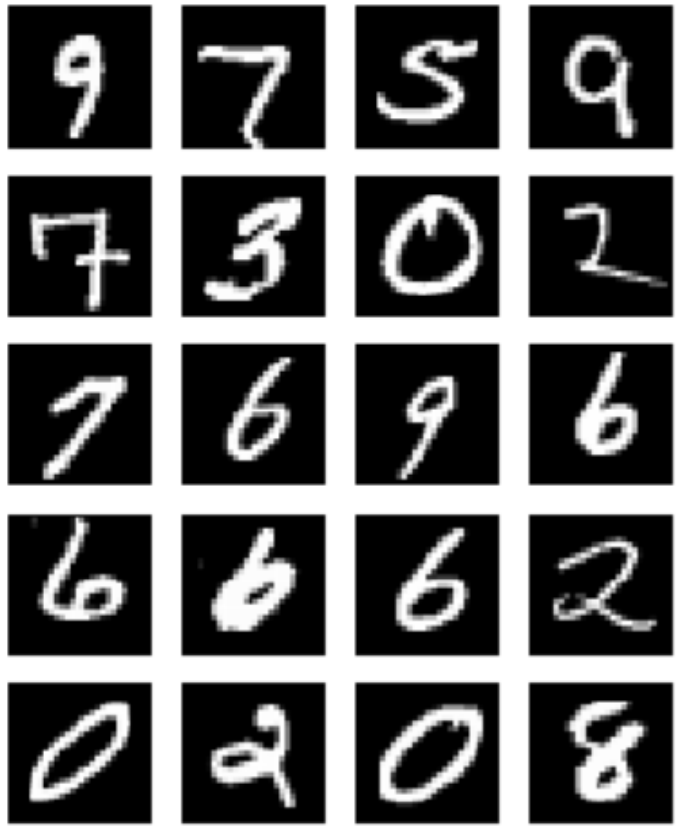

Generative Adversarial Networks

❖ **Loss function?**

Demo

Generative Adversarial Networks

❖ Loss function: A network

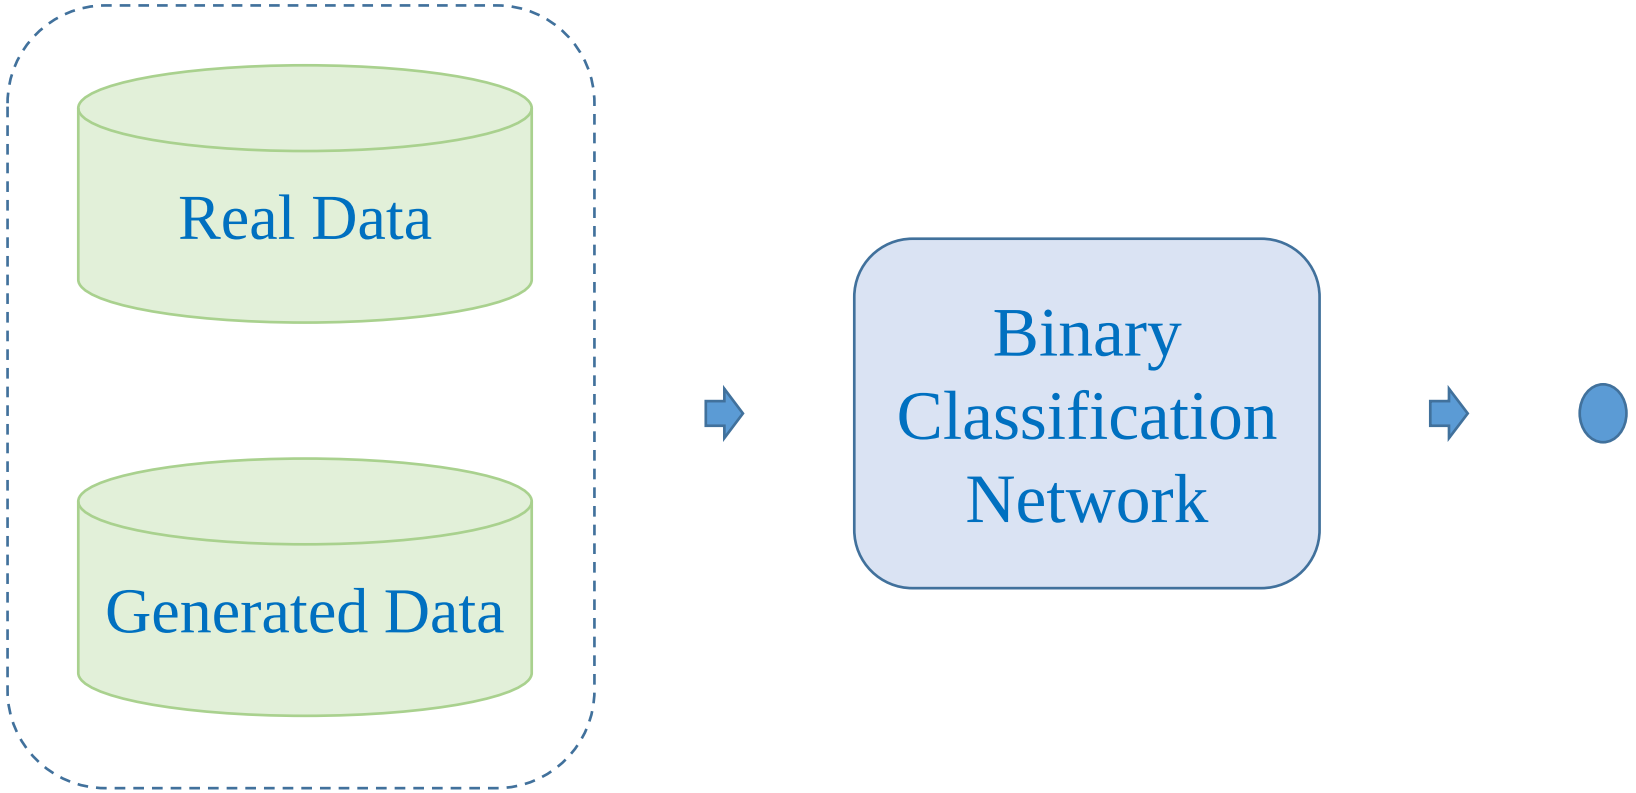

Generative Adversarial Networks

Model so far!

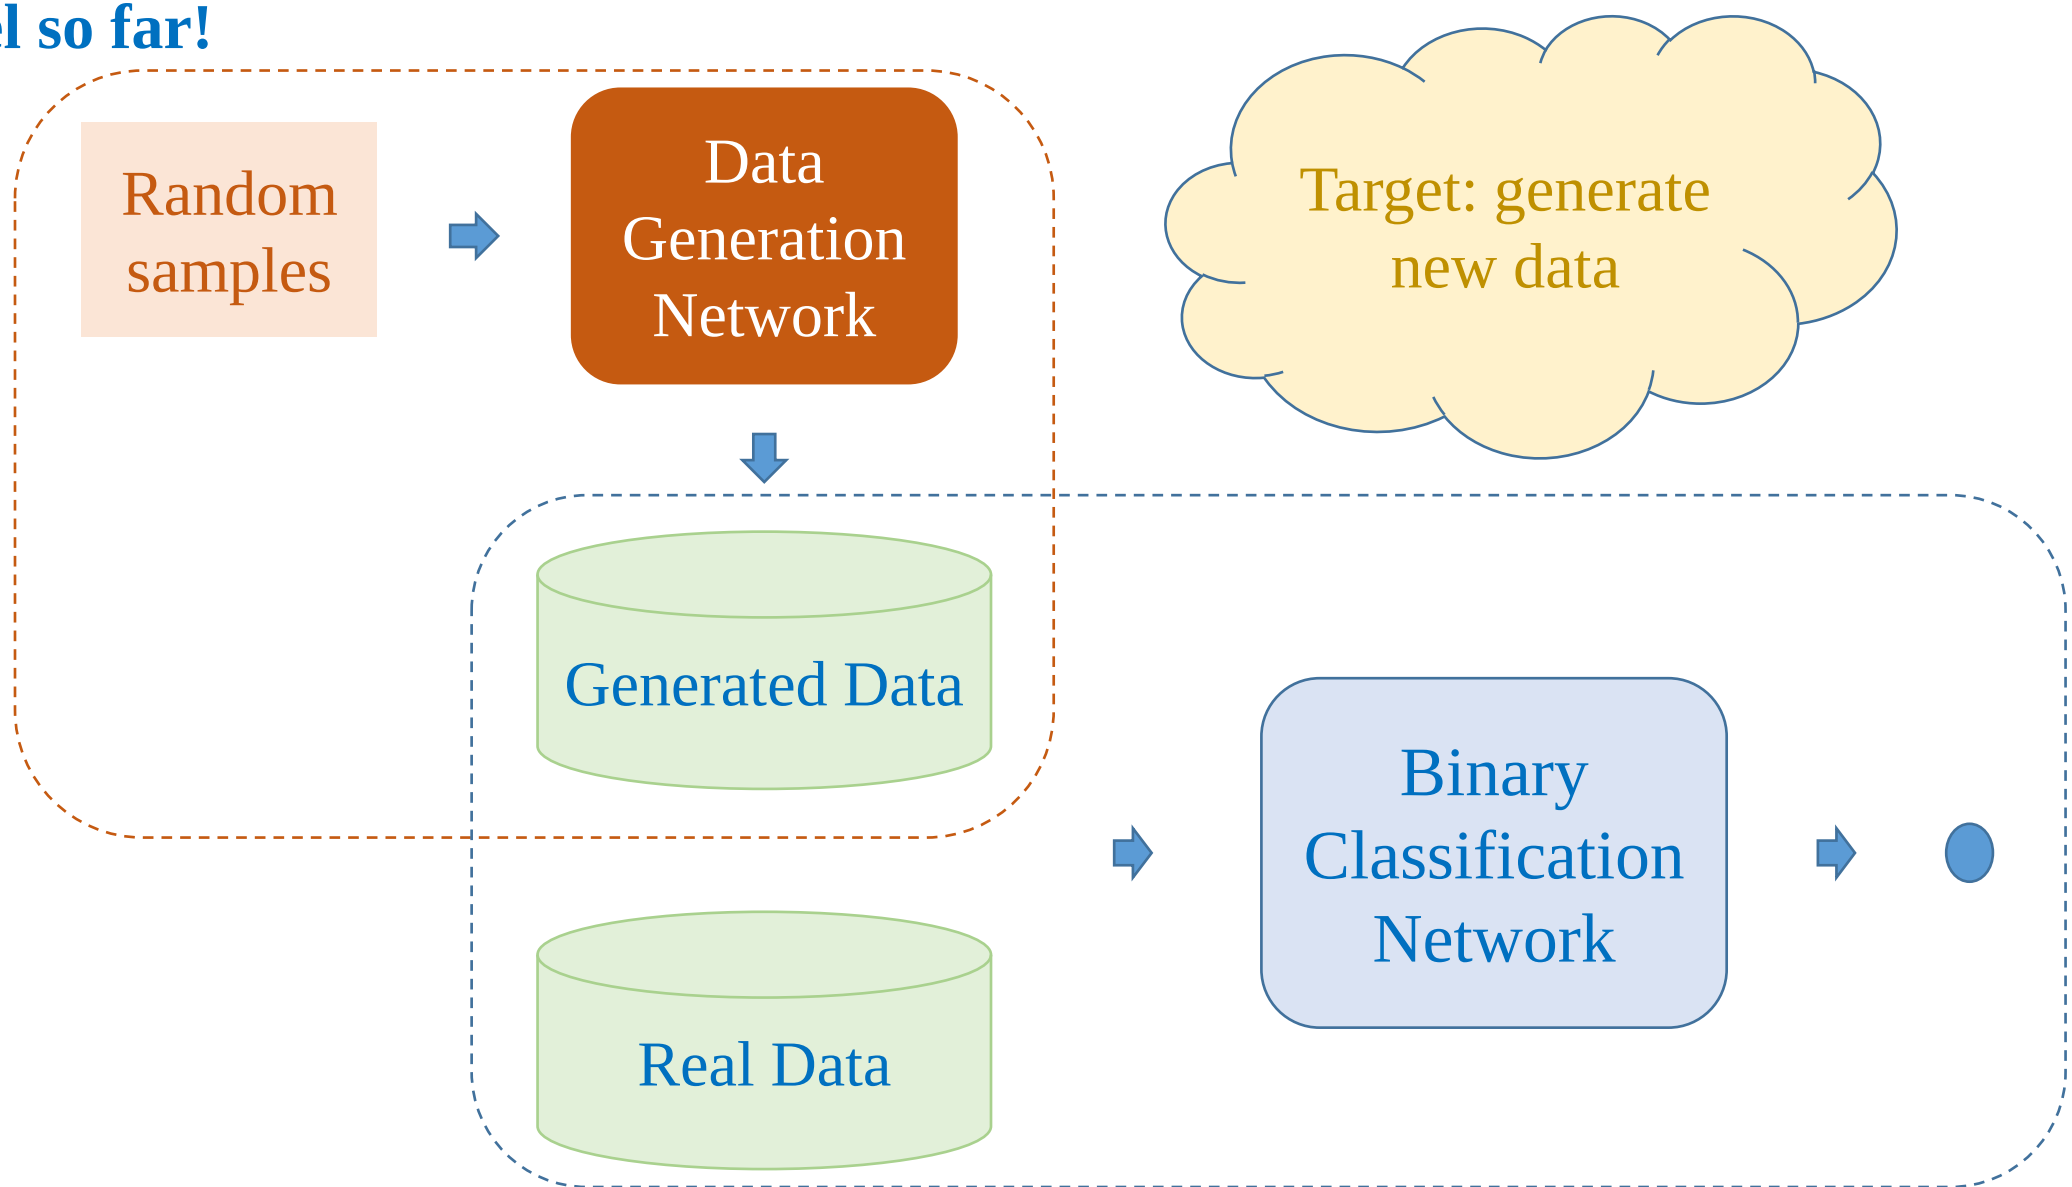

Outline

- Introduction
- GAN
- DCGAN
- Implementation

Generator Loss

```
criterion = nn.BCELoss()
real_labels = torch.ones(imgs.shape[0], 1)

z = torch.randn((imgs.shape[0], latent_dim))
gen_imgs = generator(z)
G_loss = criterion(discriminator(gen_imgs),
                   real_labels)
```

Model so far!

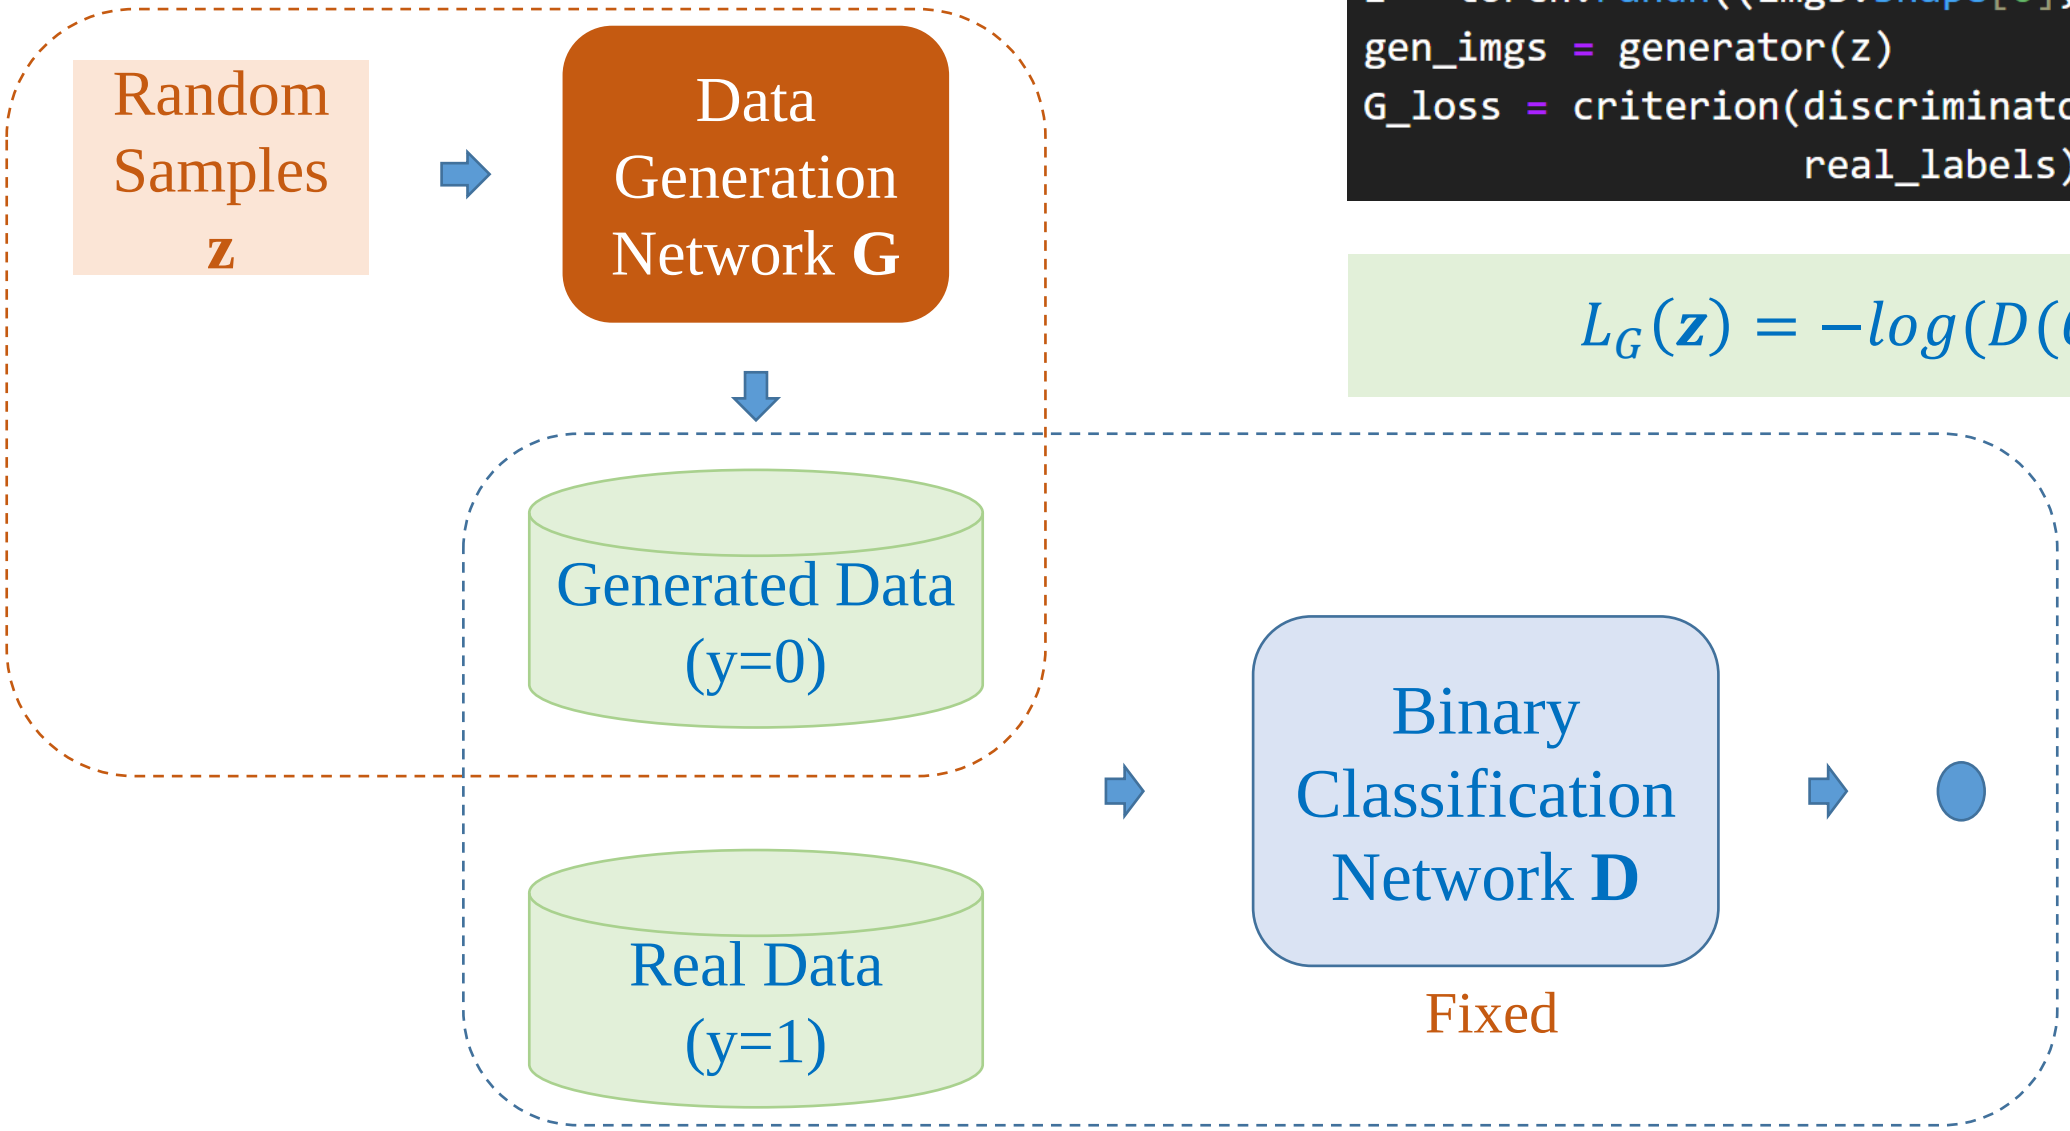

$$L_G(\boldsymbol{z}) = -log(D(G(\boldsymbol{z})))$$

Discriminator Loss

```
criterion = nn.BCELoss()
real_labels = torch.ones(imgs.shape[0], 1)
fake_labels = torch.zeros(imgs.shape[0], 1)

z = torch.randn((imgs.shape[0], latent_dim))
gen_imgs = generator(z)
real_loss = criterion(discriminator(real_imgs), real_labels)
fake_loss = criterion(discriminator(gen_imgs.detach()), fake)
D_loss = (real_loss + fake_loss) / 2
```

Model so far!

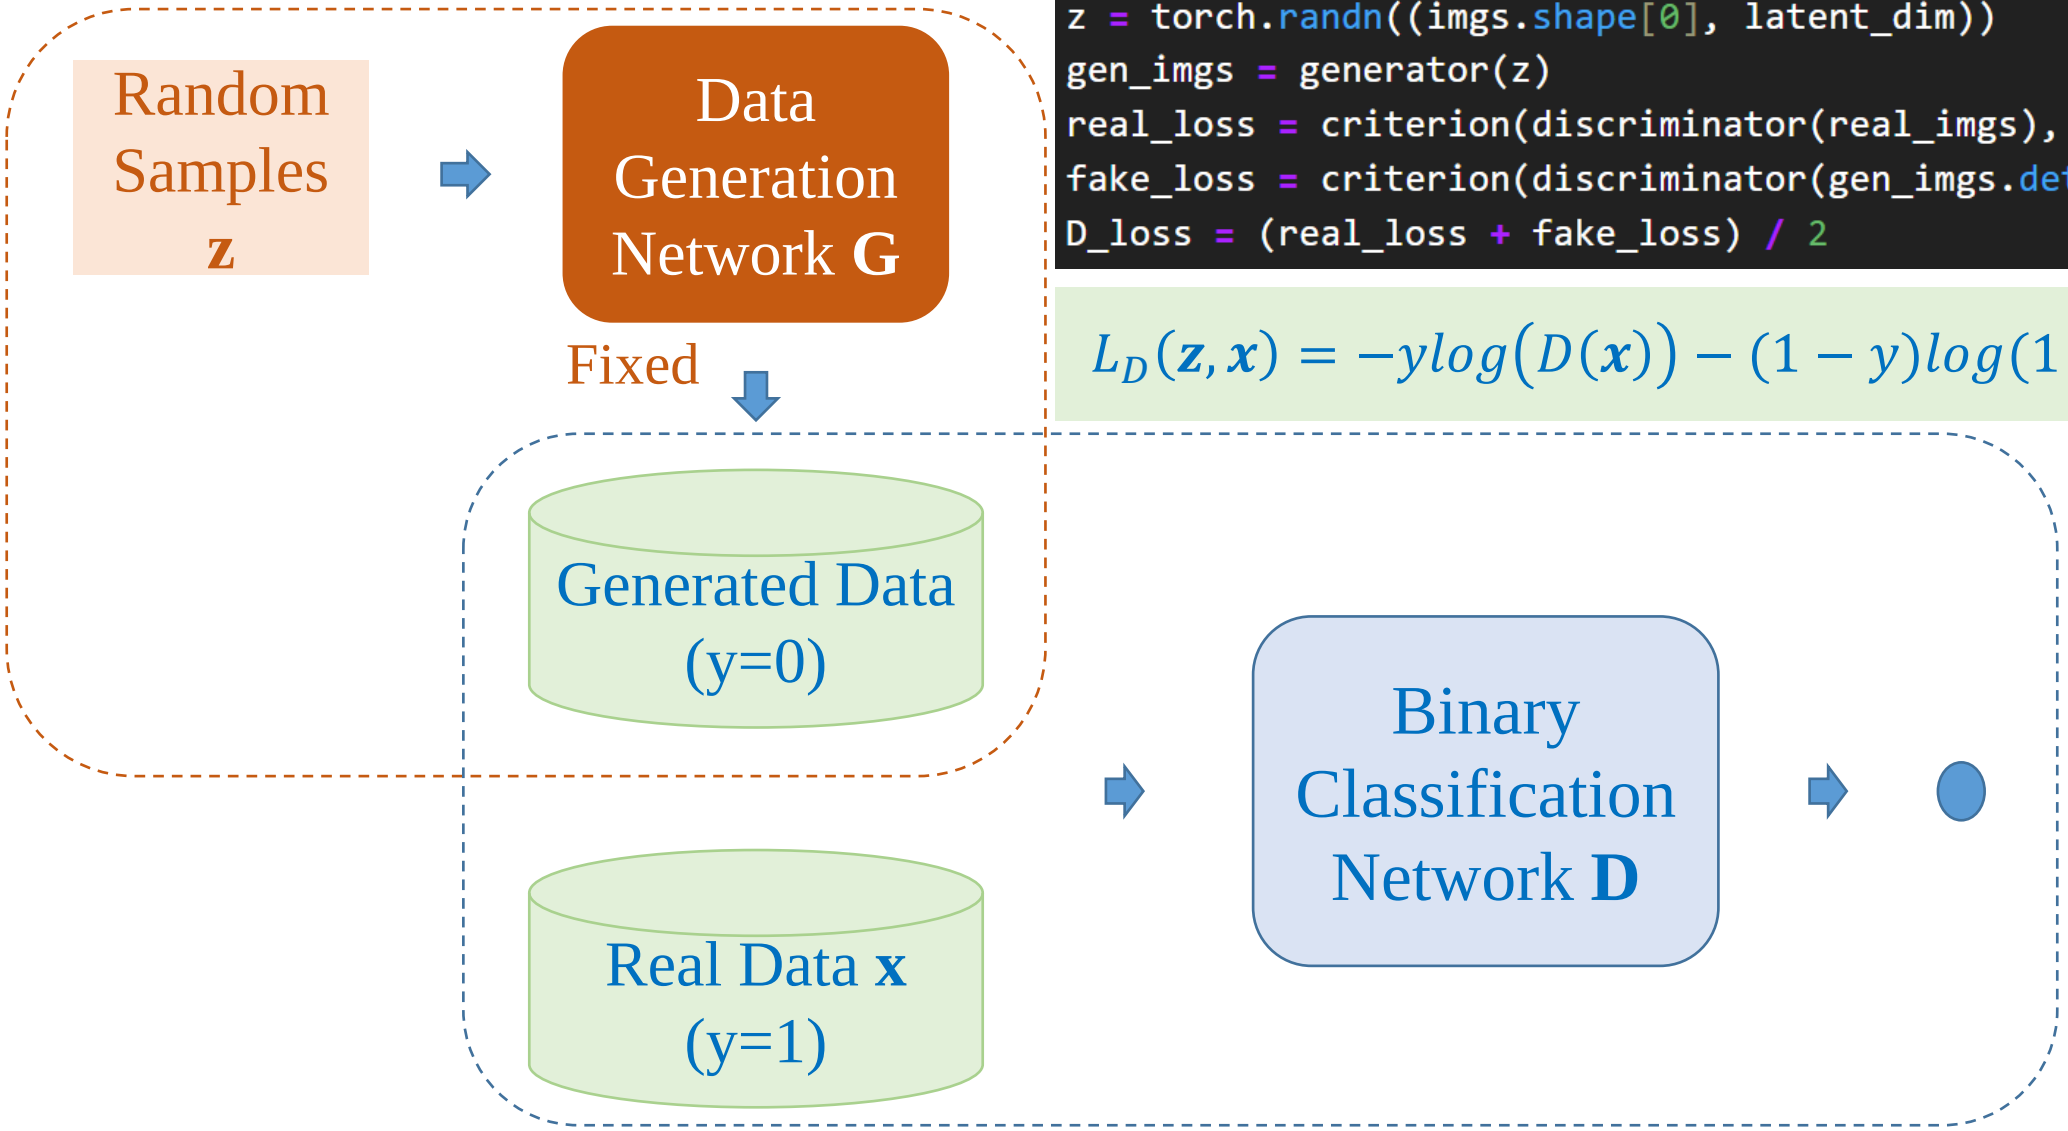

$$L_D(\boldsymbol{z}, \boldsymbol{x}) = -y log\big(D(\boldsymbol{x})\big) - (1-y) log(1 - D(G(\boldsymbol{z})))$$

Standard GAN

❖ Architecture

```
self.model = nn.Sequential(
    nn.Linear(int(np.prod(img_shape)), 512),
    nn.LeakyReLU(0.2, inplace=True),

    nn.Linear(512, 256),
    nn.LeakyReLU(0.2, inplace=True),

    nn.Linear(256, 1),
    nn.Sigmoid()
)
```

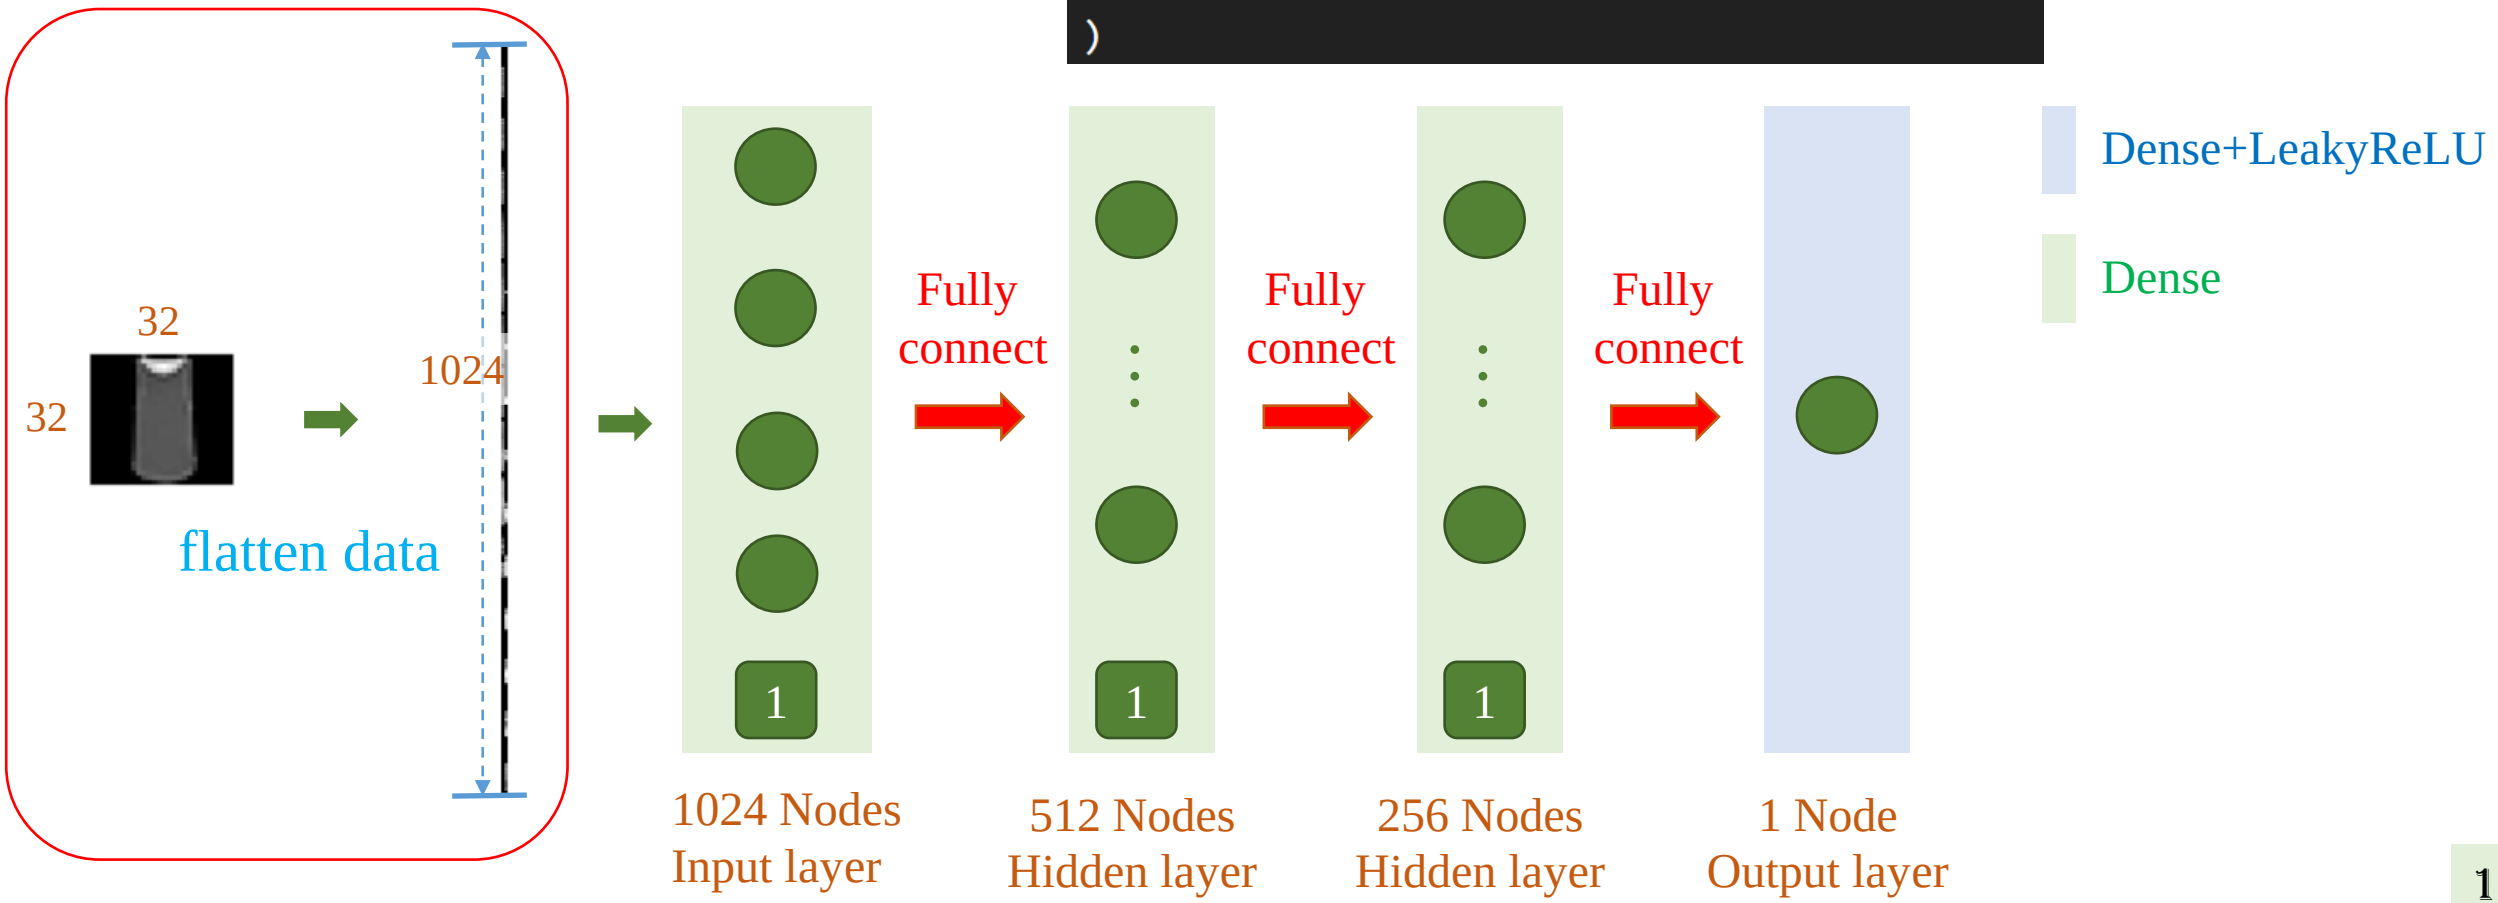

Standard GAN

❖ Architecture

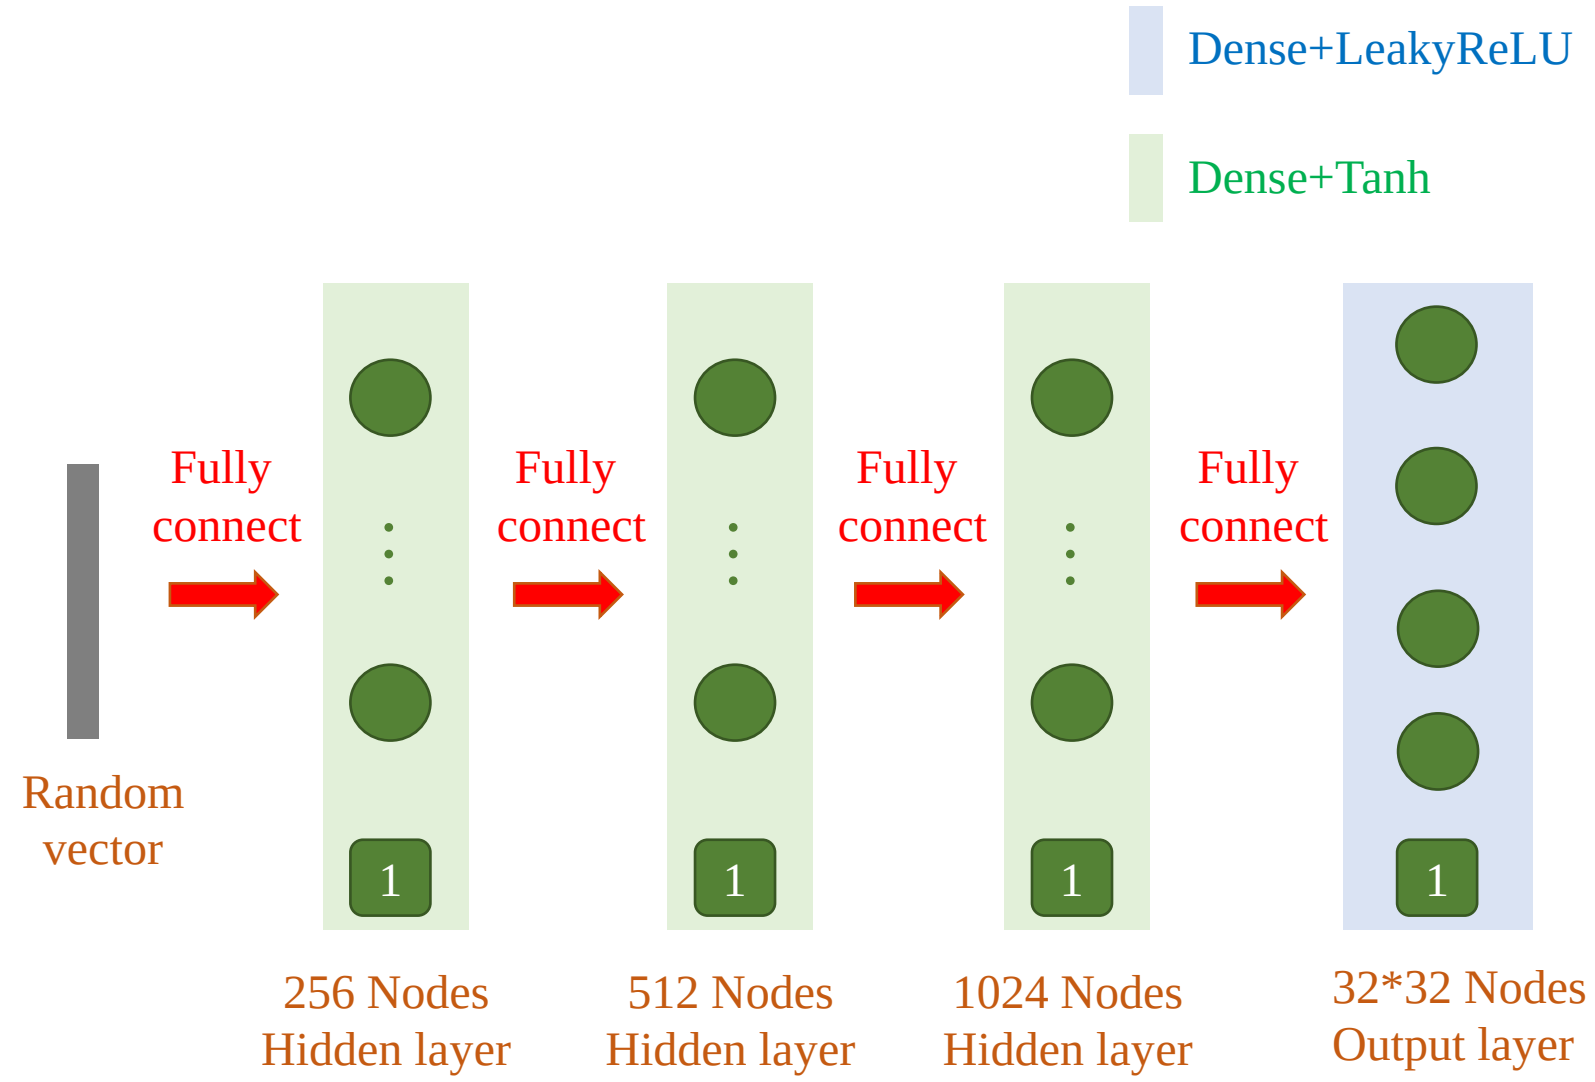


```python
class Generator(nn.Module):
    def __init__(self):
        super().__init__()
        self.model = nn.Sequential(
            nn.Linear(latent_dim, 256),
            nn.BatchNorm1d(256),
            nn.LeakyReLU(0.2, inplace=True),

            nn.Linear(256, 512),
            nn.BatchNorm1d(512),
            nn.LeakyReLU(0.2, inplace=True),

            nn.Linear(512, 1024),
            nn.BatchNorm1d(1024),
            nn.LeakyReLU(0.2, inplace=True),

            nn.Linear(1024, 32*32)),
            nn.Tanh()
        )

    def forward(self, z):
        img = self.model(z)
        img = img.view(img.size(0),
                       1, 32, 32)
        return img
```

Generative Adversarial Networks

Model so far!

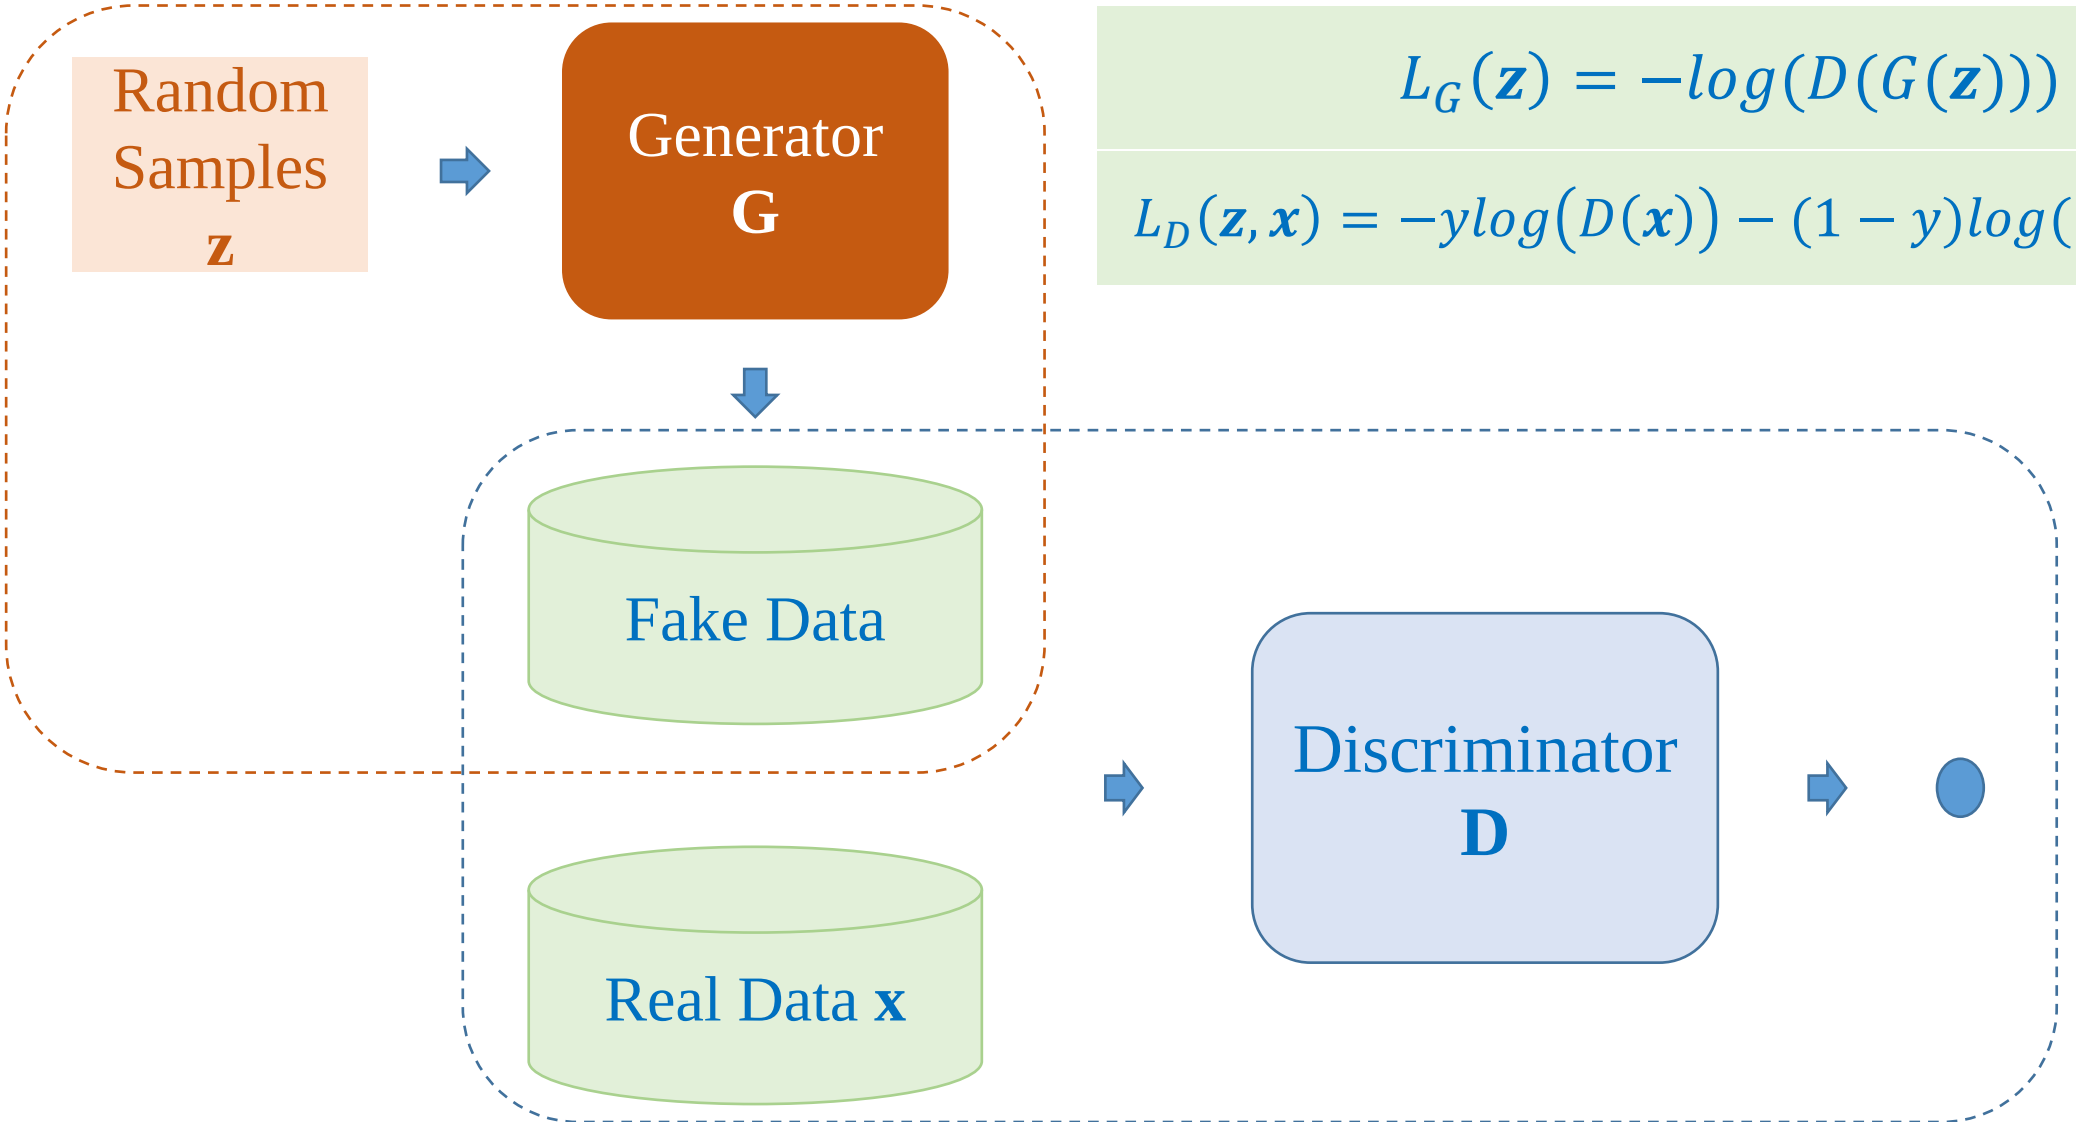

$$L_G(\boldsymbol{z}) = -log(D(G(\boldsymbol{z})))$$

$$L_D(\boldsymbol{z}, \boldsymbol{x}) = -ylog\big(D(\boldsymbol{x})\big) - (1-y)log(1 - D(G(\boldsymbol{z})))$$

Generative Adversarial Networks

Workflow

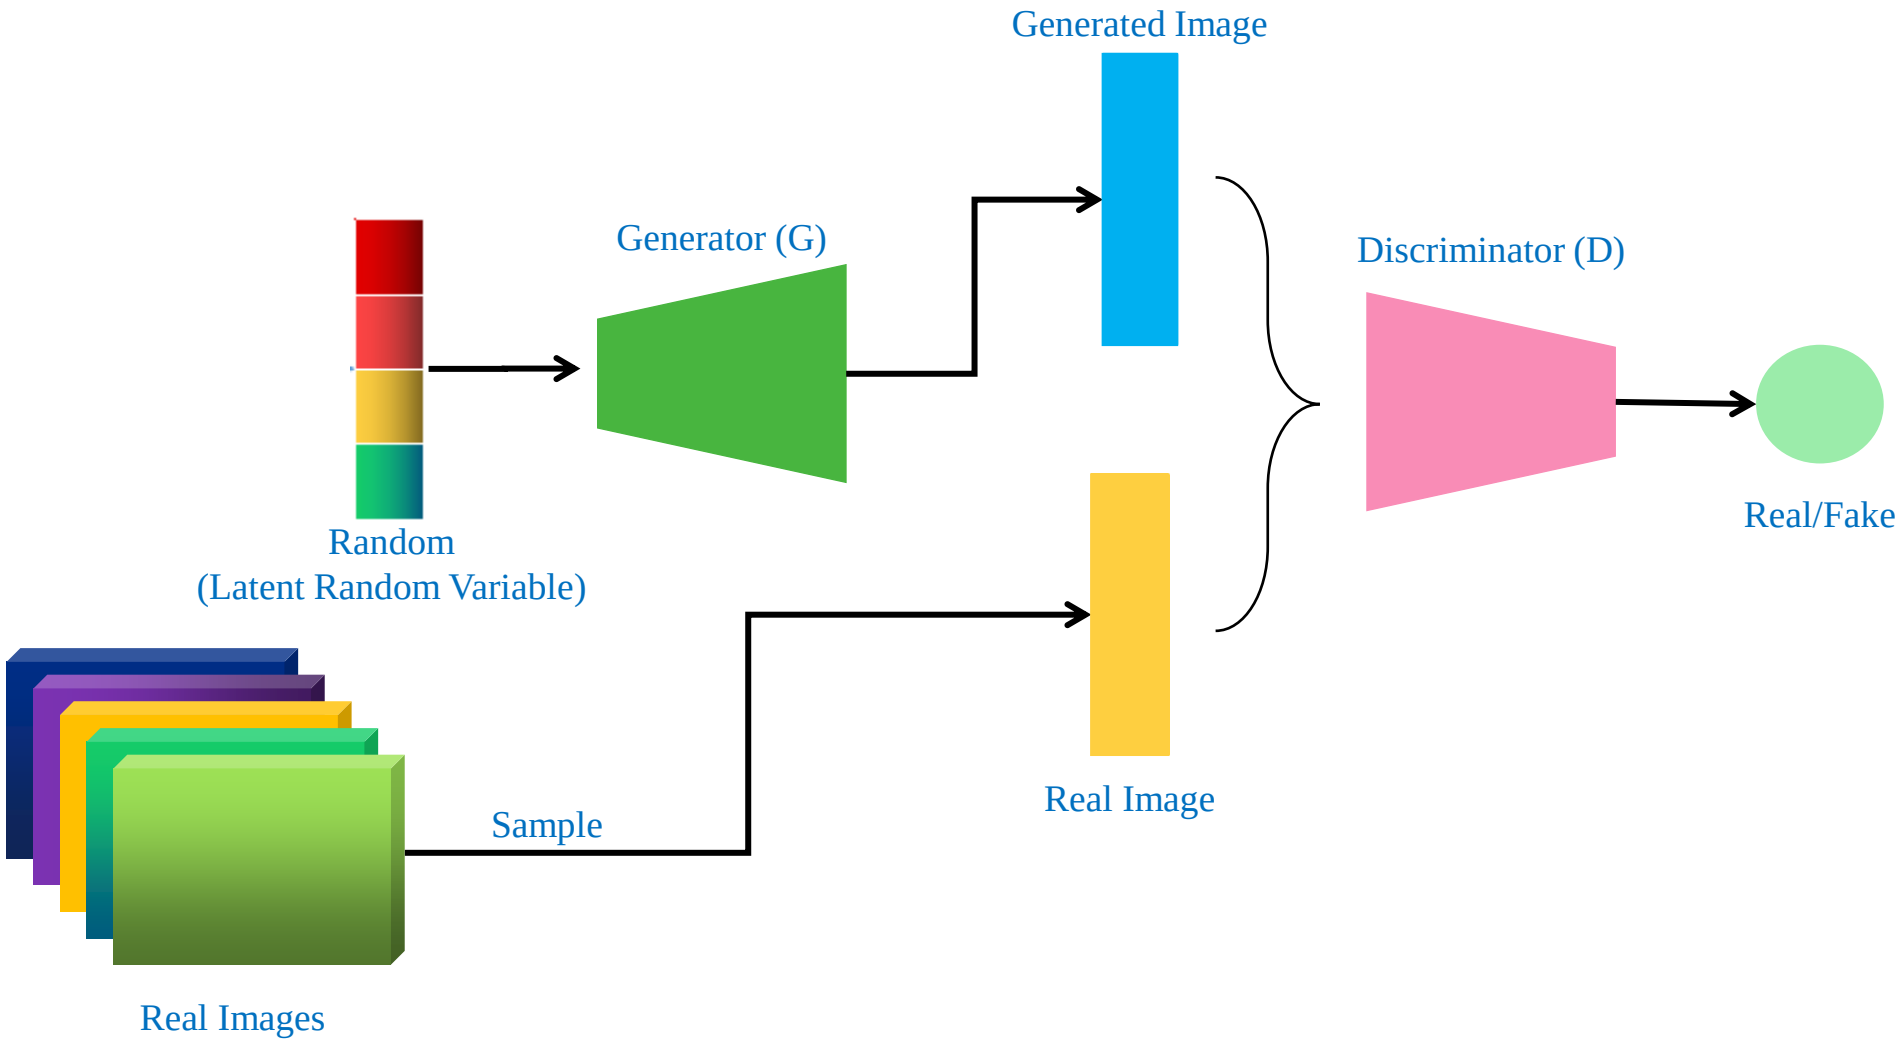

Generative Adversarial Networks

GAN Training

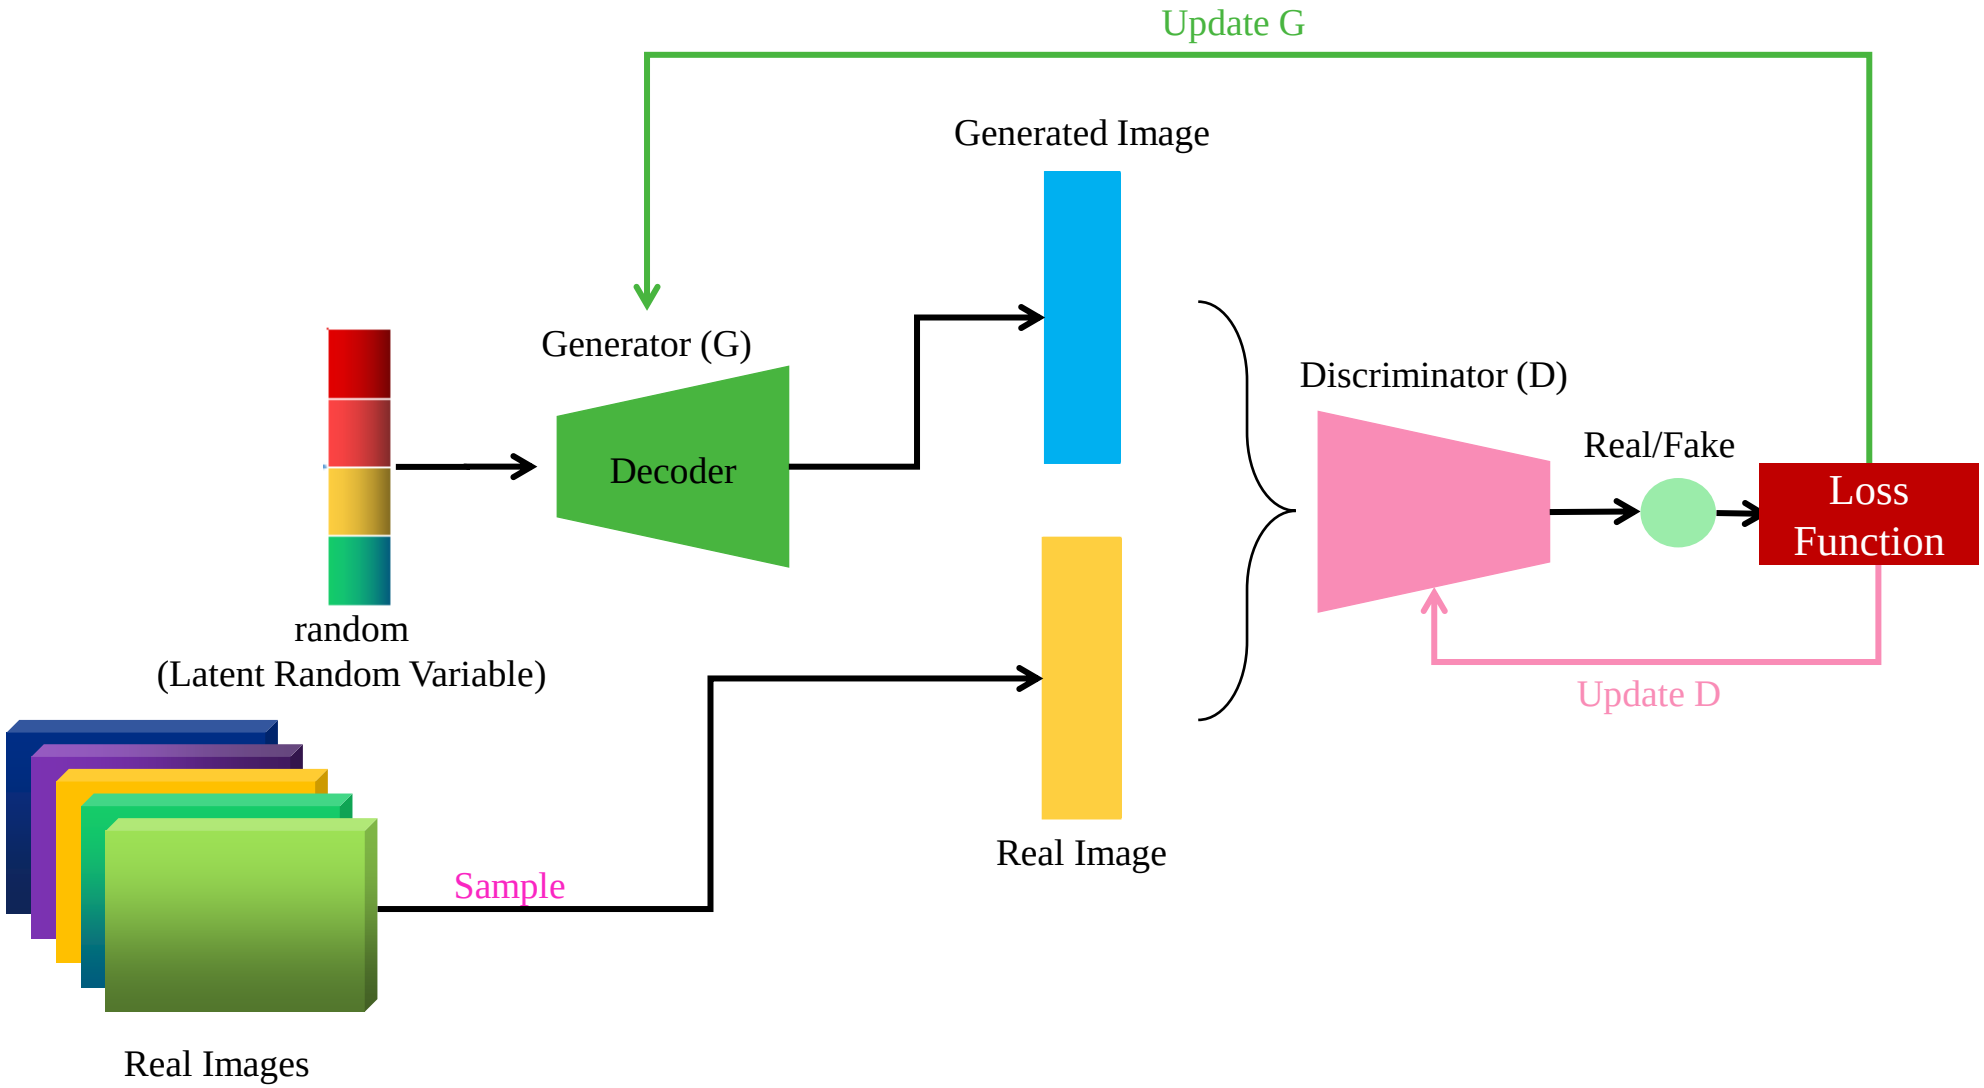

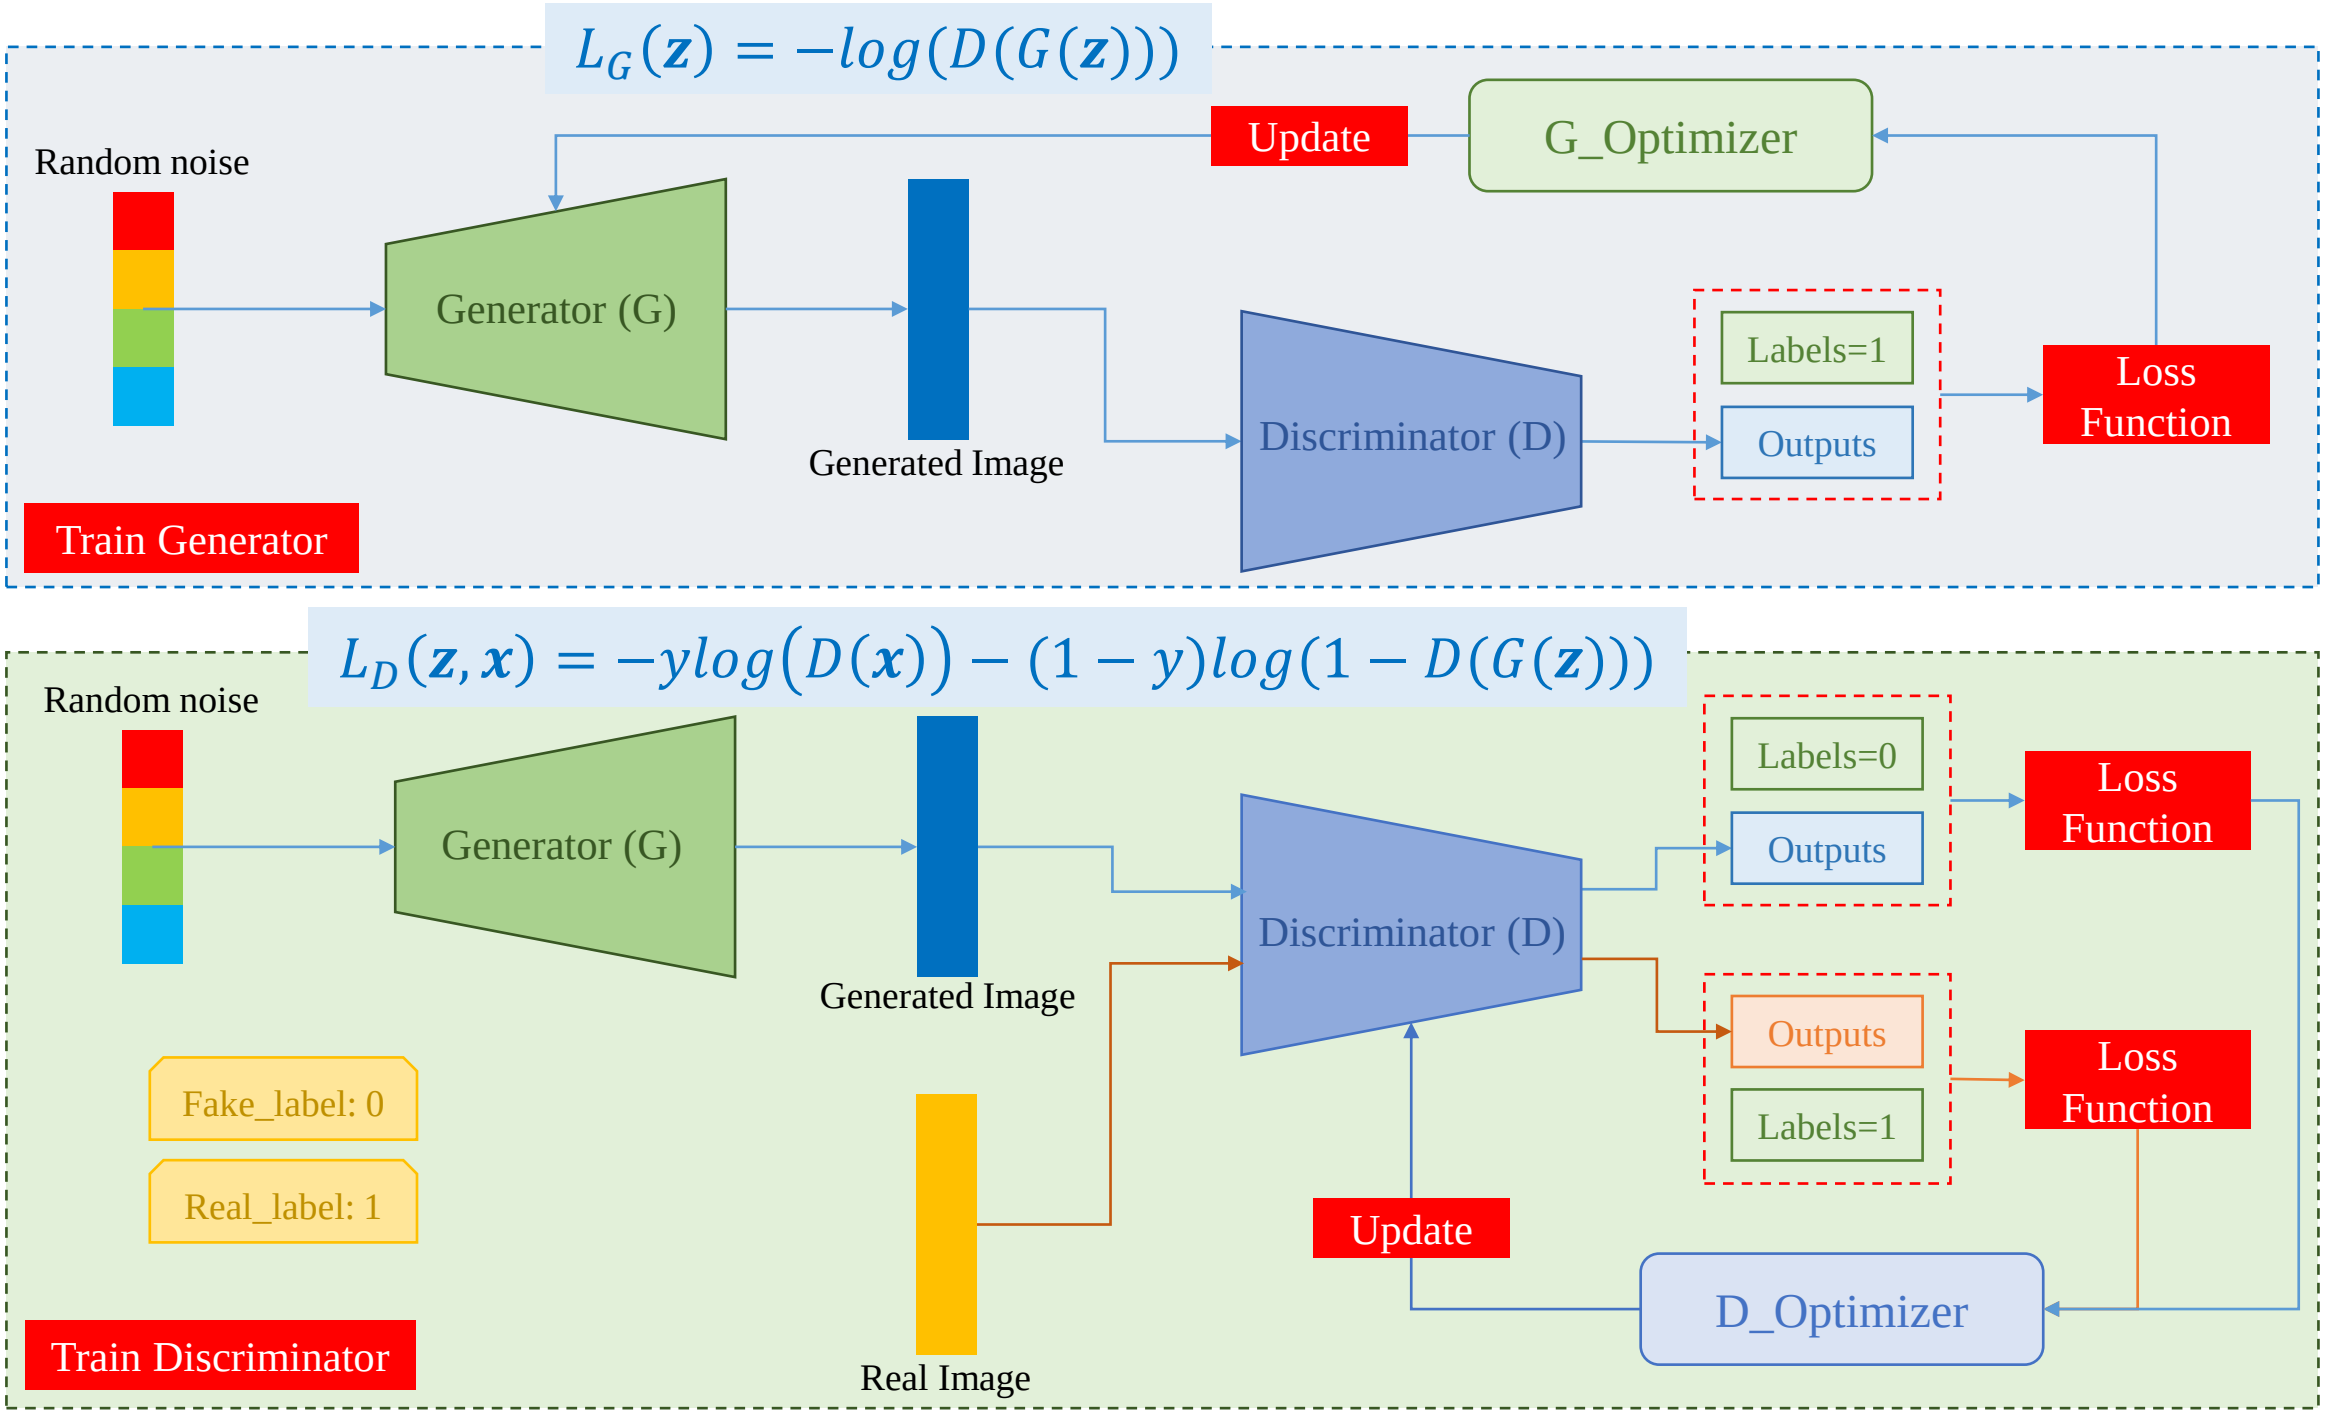
$L_G(z) = -log(D(G(z)))$
Random noise
Update
G_Optimizer
Generator (G)
Generated Image
Discriminator (D)
Labels=1
Outputs
Loss Function
Train Generator
$L_D(z, x) = -ylog(D(x)) - (1-y)log(1-D(G(z)))$
Random noise
Generator (G)
Generated Image
Discriminator (D)
Labels=0
Outputs
Loss Function
Outputs
Loss Function
Labels=1
Fake_label: 0
Real_label: 1
Update
D_Optimizer
Real Image
Train Discriminator

Standard GAN

❖ Train Generator

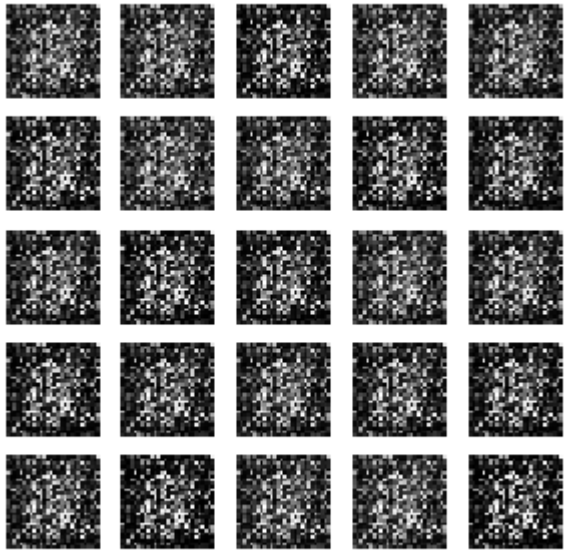

Epoch 1

```
criterion = nn.BCELoss()
optimizer_G = torch.optim.Adam(generator.parameters(),
                               lr=0.0001)

for i, (imgs, _) in enumerate(dataloader)
    # prepare labels
    real_labels = torch.ones(imgs.shape[0], 1)
    fake_labels = torch.zeros(imgs.shape[0], 1)

    # zero_grad
    optimizer_G.zero_grad()

    # Noise input for Generator
    z = torch.randn((imgs.shape[0], latent_dim))

    # generate images
    gen_imgs = generator(z)

    # comupte loss
    G_loss = criterion(discriminator(gen_imgs),
                       real_labels)

    # update
    G_loss.backward()
    optimizer_G.step()
```

Standard GAN

❖ Train Discriminator

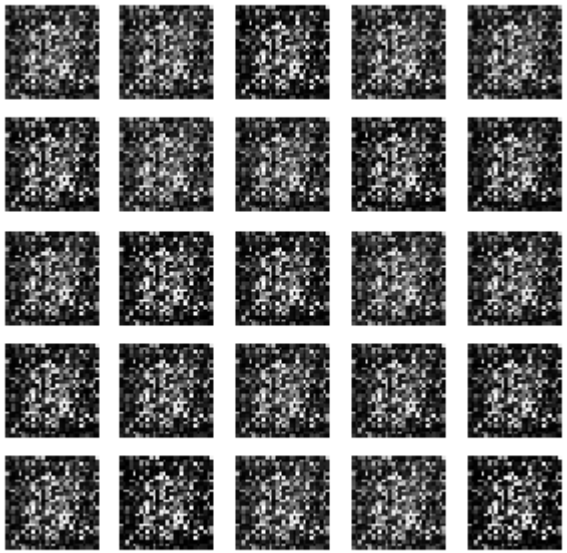

Epoch 1

```
criterion = nn.BCELoss()
optimizer_D = torch.optim.Adam(discriminator.parameters(),
                               lr=0.0002)

for i, (imgs, _) in enumerate(dataloader)
    # prepare labels
    real_labels = torch.ones(imgs.shape[0], 1)
    fake_labels = torch.zeros(imgs.shape[0], 1)

    # zero_grad
    optimizer_D.zero_grad()

    # Noise input for Generator
    z = torch.randn((imgs.shape[0], latent_dim))

    # generate images
    gen_imgs = generator(z)

    # comupte loss
    real_loss = criterion(discriminator(real_imgs),
                          real_labels)
    fake_loss = criterion(discriminator(gen_imgs.detach()),
                          fake_labels)
    D_loss = (real_loss + fake_loss) / 2

    # update
    D_loss.backward()
    optimizer_D.step()
```

Outline

- Introduction
- GAN
- DCGAN
- Implementation

Generative Adversarial Networks

❖ How to improve?

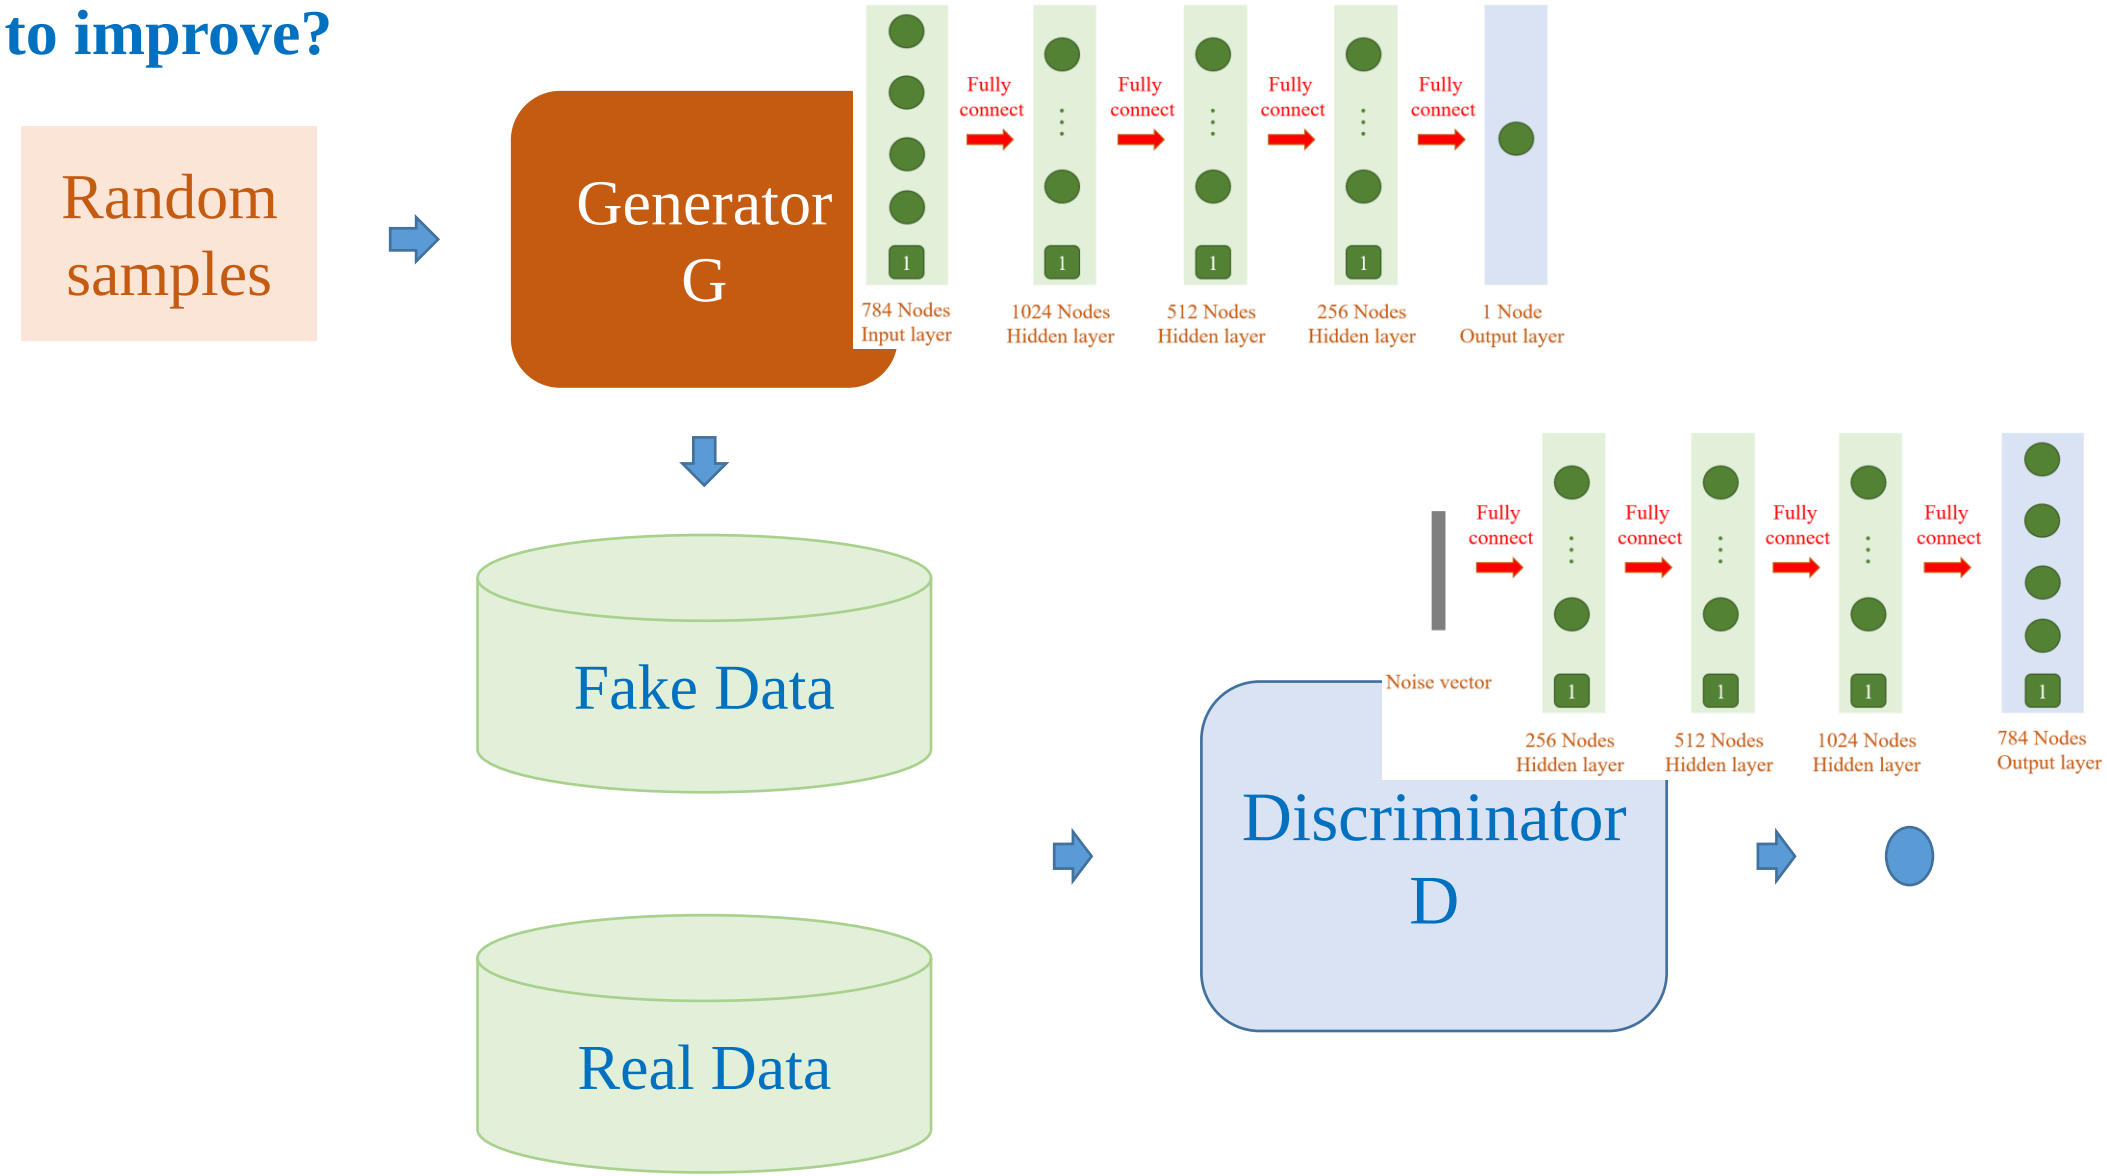

Generative Adversarial Networks

Network architecture

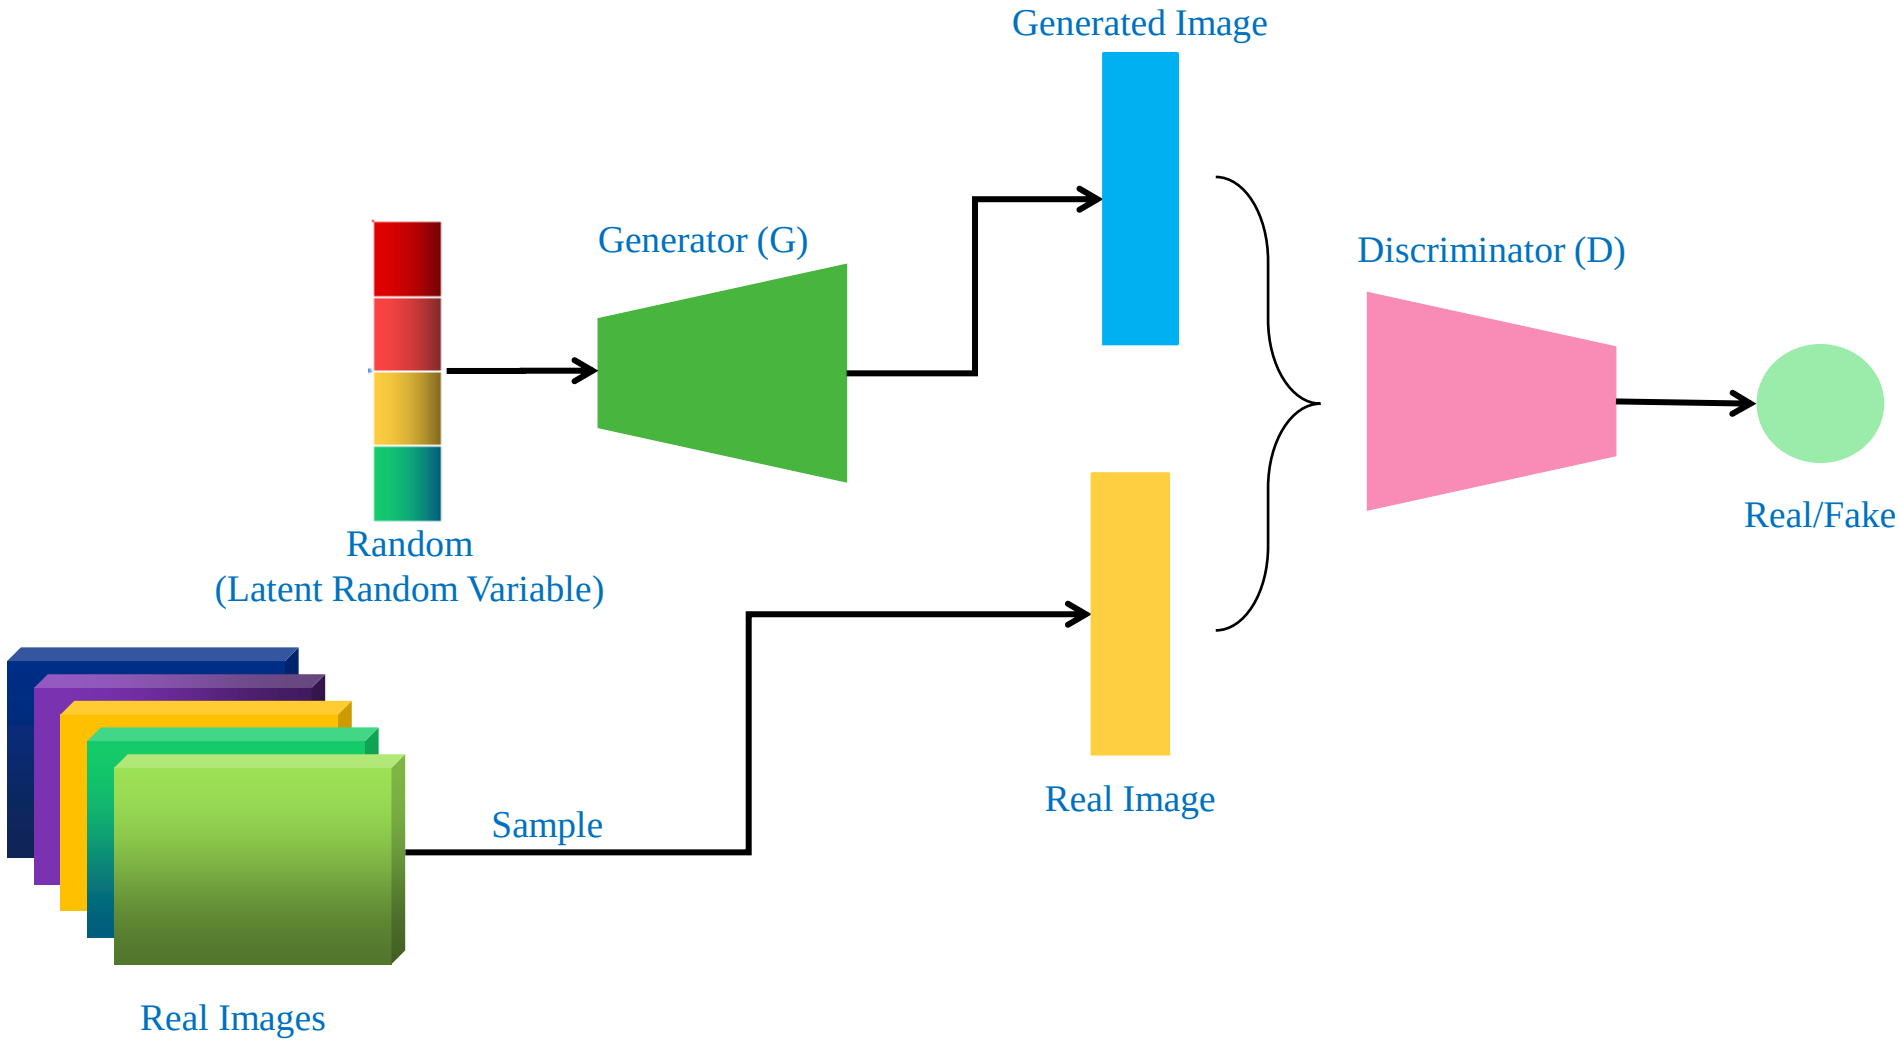

Generative Adversarial Networks

❖ Architecture

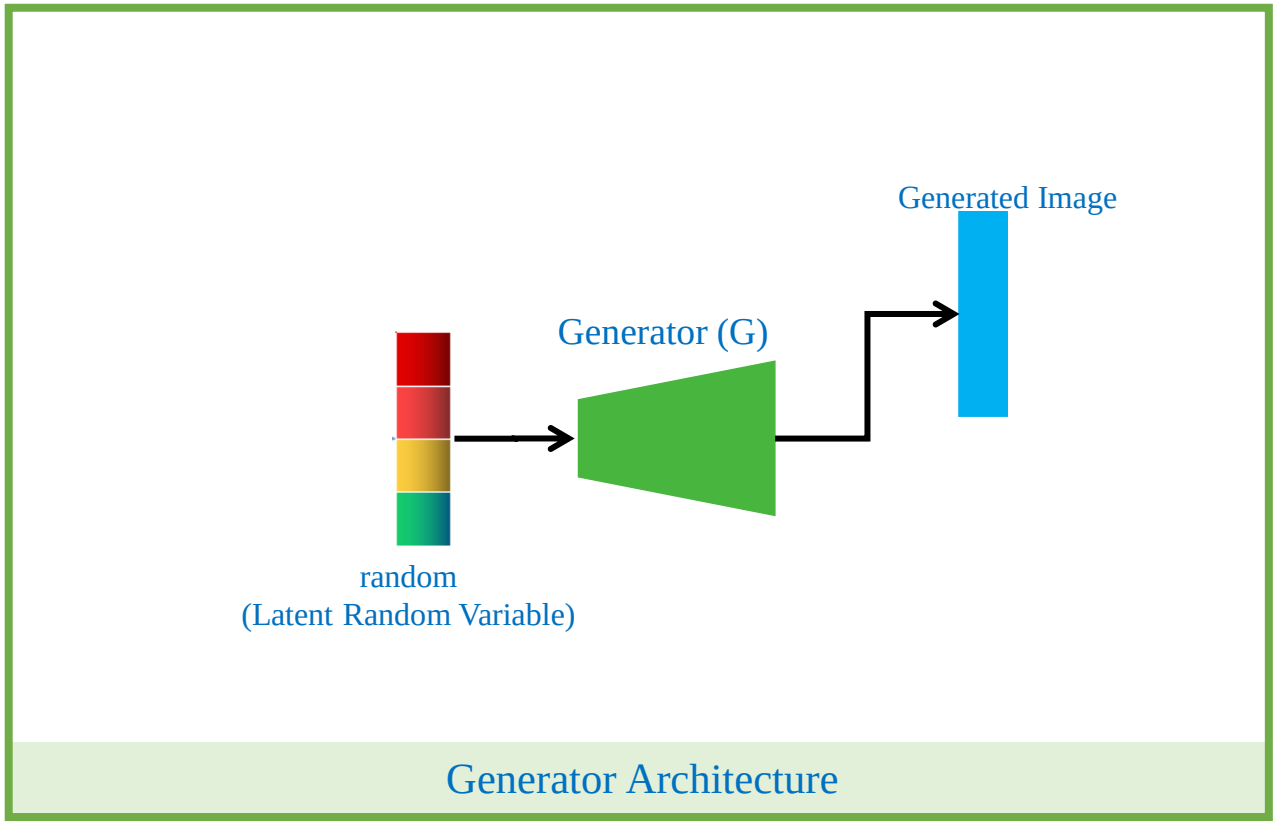

Generator Architecture

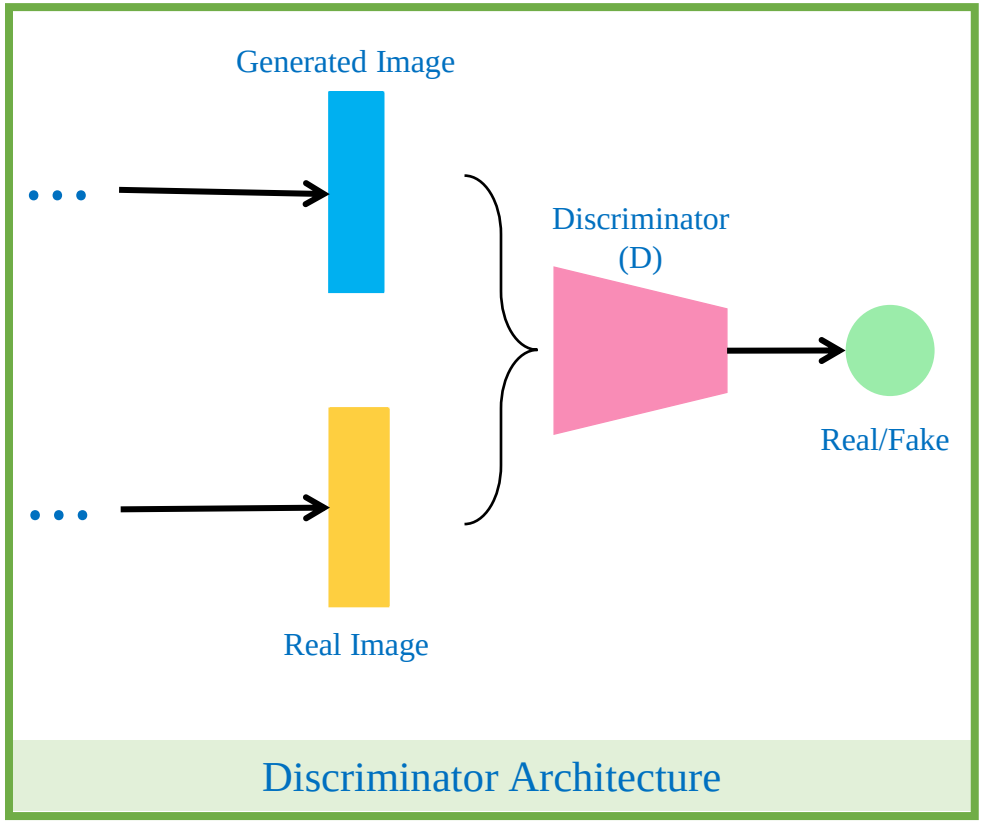

Discriminator Architecture

Generative Adversarial Networks

GAN Training

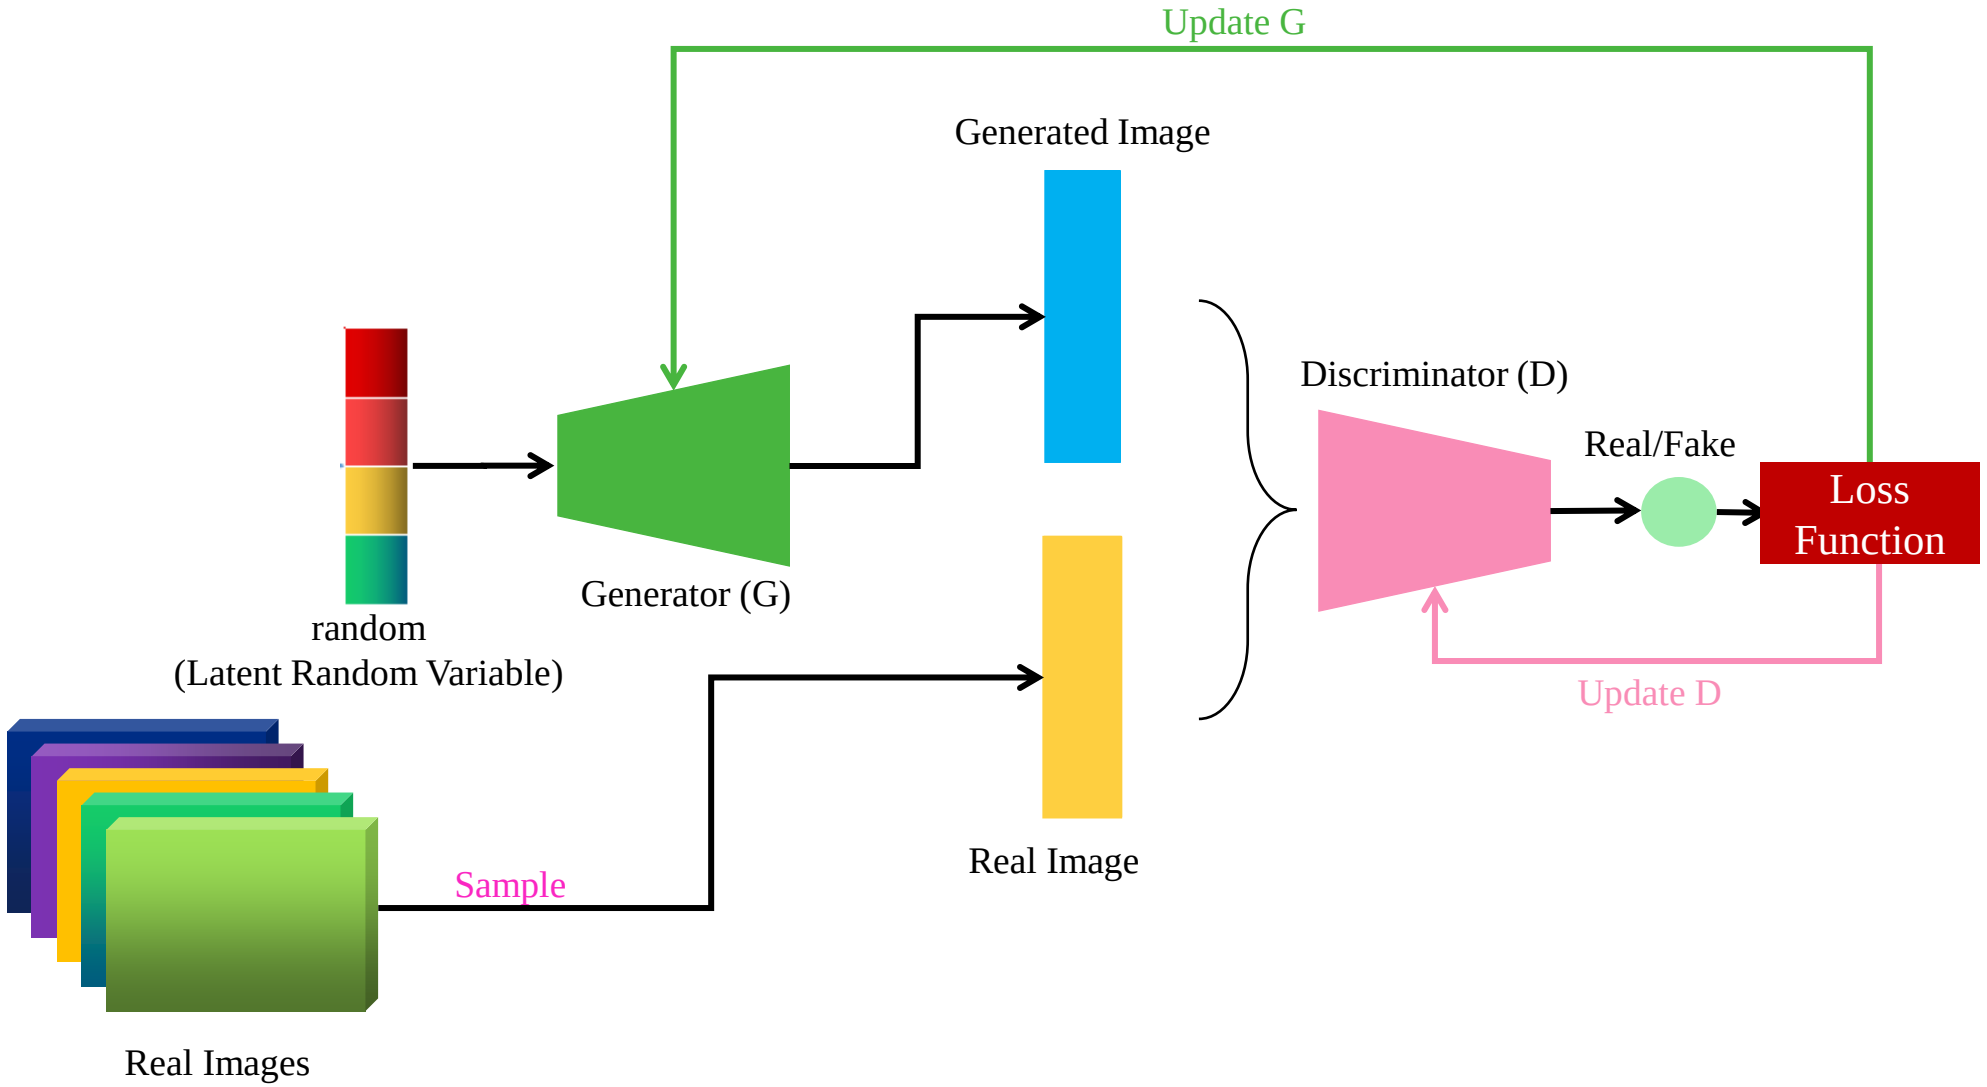

DCGAN

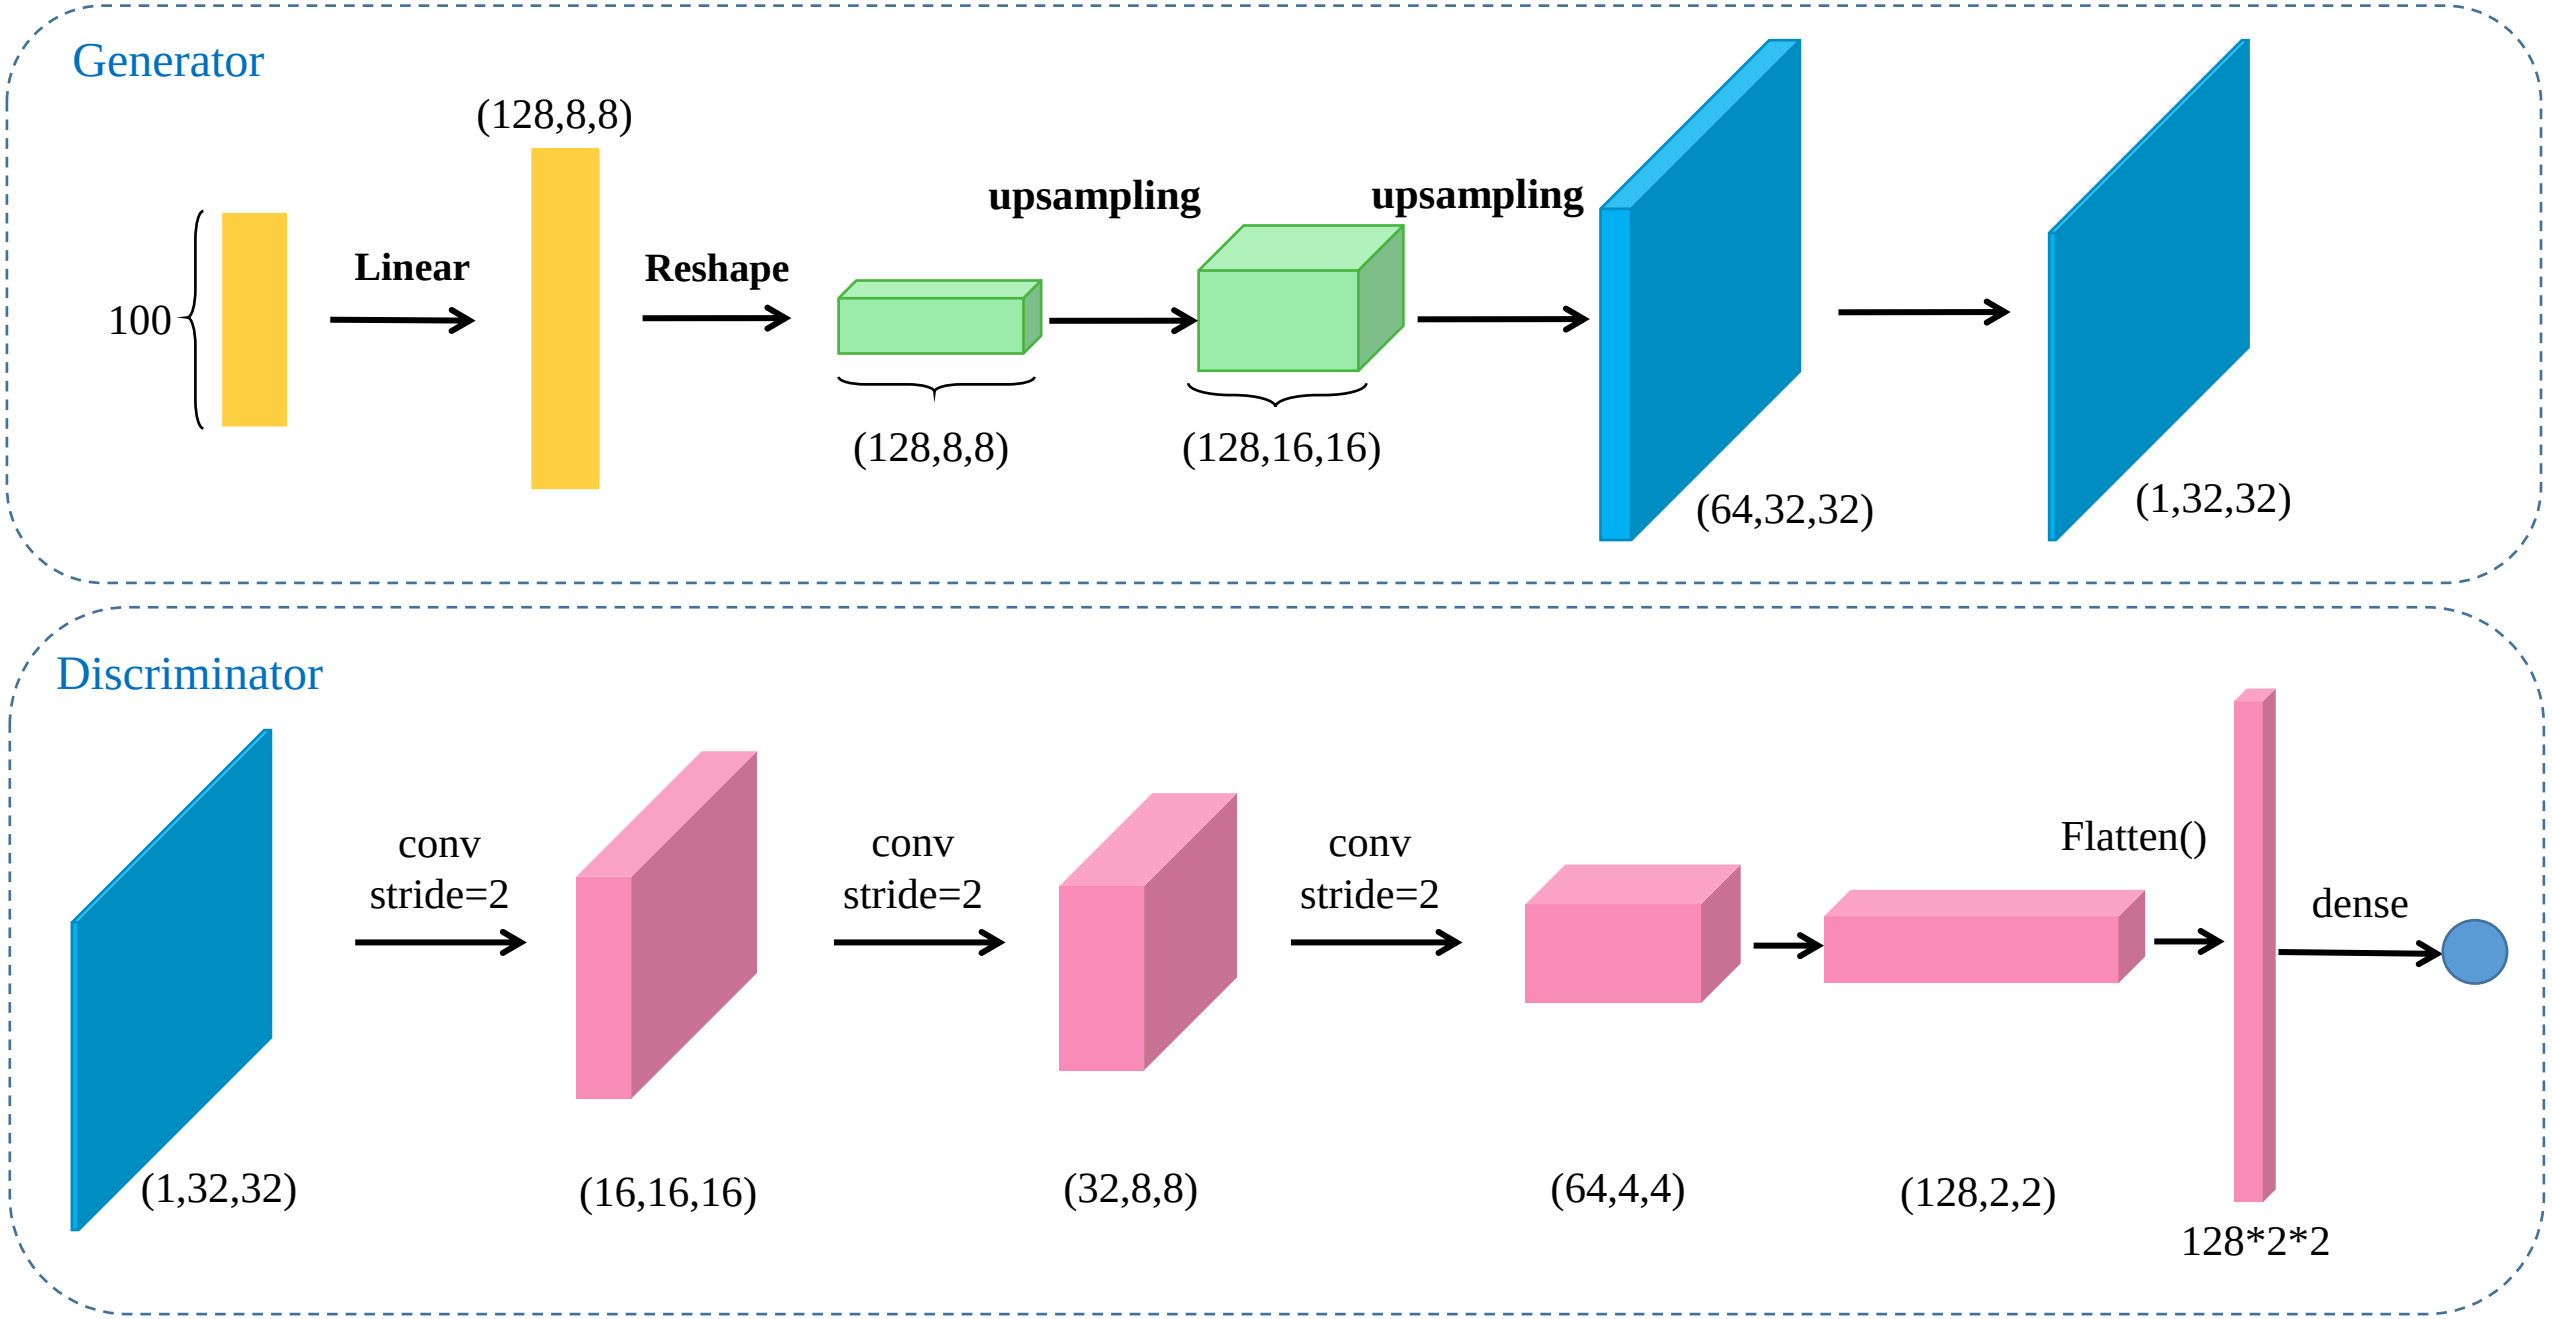
Generator
100
Linear
(128,8,8)
Reshape
upsampling
upsampling
(128,8,8)
(128,16,16)
(64,32,32)
(1,32,32)
Discriminator
(1,32,32)
conv
stride=2
(16,16,16)
conv
stride=2
(32,8,8)
conv
stride=2
(64,4,4)
(128,2,2)
Flatten()
128*2*2
dense


QUIZ TIME

```python
class Generator(nn.Module):
    def __init__(self):
        super().__init__()
        self.init_size = 8
        self.fc = nn.Linear(latent_dim, 128*8*8)
        self.conv_blocks = nn.Sequential(
            nn.BatchNorm2d(128),
            nn.Upsample(scale_factor=2),
            nn.Conv2d(128, 128, 3, padding=1),

            nn.BatchNorm2d(128),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Upsample(scale_factor=2),
            nn.Conv2d(128, 64, 3, padding=1),

            nn.BatchNorm2d(64,),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(64, channels,
                      kernel_size=3, padding=1),
            nn.Tanh()
        )

    def forward(self, z):
        x = self.fc(z)
        x = x.view(x.shape[0], 128,
                   self.init_size, self.init_size)
        img = self.conv_blocks(x)
        return img
```

```python
class Descriminator(nn.Module):
    def __init__(self):
        super().__init__()
        self.model = nn.Sequential(
            nn.Conv2d(channels, 16, kernel_size=3, stride=2, padding=1),
            nn.LeakyReLU(0.2, inplace=True)

            nn.Conv2d(16, 32, kernel_size=3, stride=2, padding=1),
            nn.LeakyReLU(0.2, inplace=True),

            nn.Conv2d(32, 64, kernel_size=3, stride=2, padding=1),
            nn.LeakyReLU(0.2, inplace=True),

            nn.Conv2d(64, 128, kernel_size=3, stride=2, padding=1),
            nn.LeakyReLU(0.2, inplace=True),
        )

        self.adv_layer = nn.Sequential(
            nn.Linear(128*2*2, 1),
            nn.Sigmoid()
        )

    def forward(self, img):
        x = self.model(img)
        x = x.view(x.shape[0], -1)
        validity = self.adv_layer(x)
        return validity
```

DCGAN

❖ Train Generator

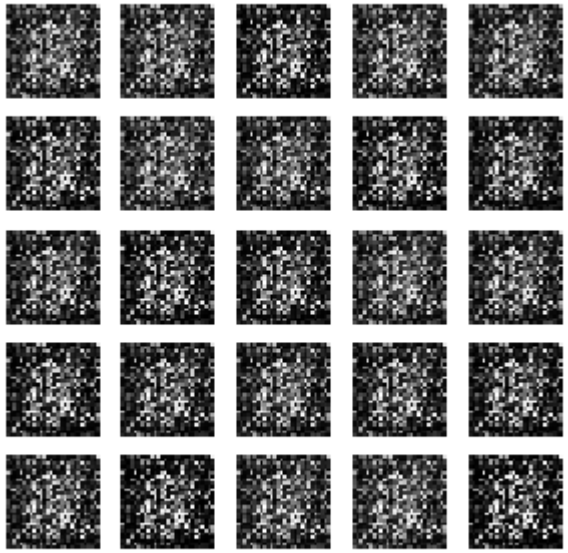

Epoch 1

```
criterion = nn.BCELoss()
optimizer_G = torch.optim.Adam(generator.parameters(),
                               lr=0.0001)

for i, (imgs, _) in enumerate(dataloader)
    # prepare labels
    real_labels = torch.ones(imgs.shape[0], 1)
    fake_labels = torch.zeros(imgs.shape[0], 1)

    # zero_grad
    optimizer_G.zero_grad()

    # Noise input for Generator
    z = torch.randn((imgs.shape[0], latent_dim))

    # generate images
    gen_imgs = generator(z)

    # comupte loss
    G_loss = criterion(discriminator(gen_imgs),
                       real_labels)

    # update
    G_loss.backward()
    optimizer_G.step()
```

DCGAN

❖ Train Discriminator

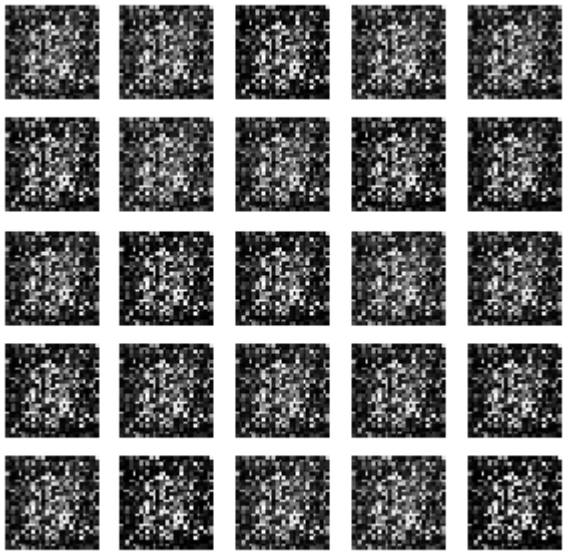

Epoch 1

```
criterion = nn.BCELoss()
optimizer_D = torch.optim.Adam(discriminator.parameters(),
                               lr=0.0002)

for i, (imgs, _) in enumerate(dataloader)
    # prepare labels
    real_labels = torch.ones(imgs.shape[0], 1)
    fake_labels = torch.zeros(imgs.shape[0], 1)

    # zero_grad
    optimizer_D.zero_grad()

    # Noise input for Generator
    z = torch.randn((imgs.shape[0], latent_dim))

    # generate images
    gen_imgs = generator(z)

    # comupte loss
    real_loss = criterion(discriminator(real_imgs),
                          real_labels)
    fake_loss = criterion(discriminator(gen_imgs.detach()),
                          fake_labels)
    D_loss = (real_loss + fake_loss) / 2

    # update
    D_loss.backward()
    optimizer_D.step()
```

AI VIETNAM
All-in-One Course

DCGAN

❖ Implementation

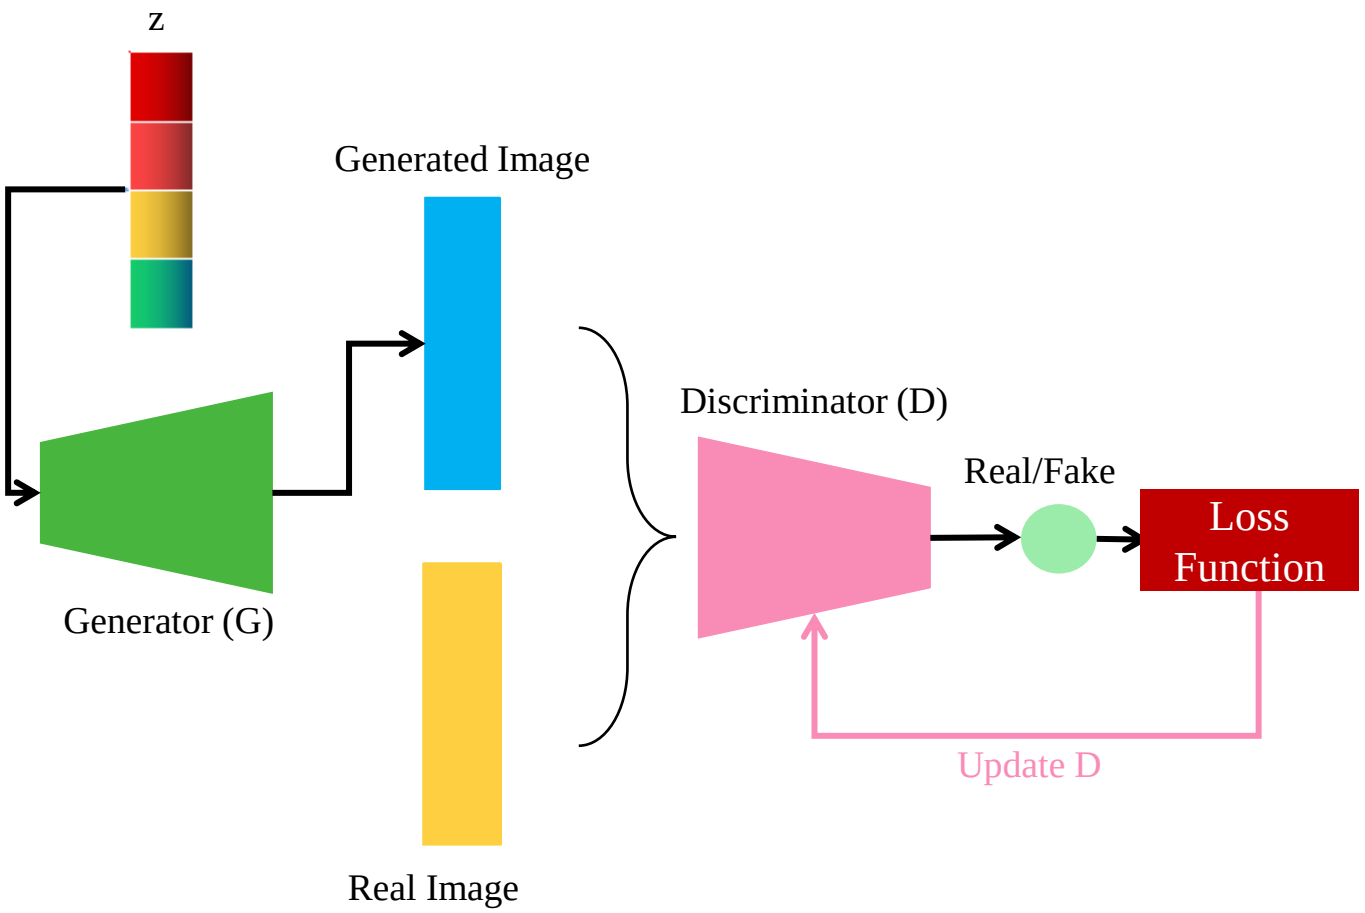


```python
criterion = nn.BCELoss()
optimizer_D = torch.optim.Adam(discriminator.parameters(),
                               lr=0.0002)

for i, (imgs, _) in enumerate(dataloader)
    # prepare labels
    real_labels = torch.ones(imgs.shape[0], 1)
    fake_labels = torch.zeros(imgs.shape[0], 1)

    # zero_grad
    optimizer_D.zero_grad()

    # Noise input for Generator
    z = torch.randn((imgs.shape[0], latent_dim))

    # generate images
    gen_imgs = generator(z)

    # comupte loss
    real_loss = criterion(discriminator(real_imgs),
                          real_labels)
    fake_loss = criterion(discriminator(gen_imgs.detach()),
                          fake_labels)
    D_loss = (real_loss + fake_loss) / 2

    # update
    D_loss.backward()
    optimizer_D.step()
```

DCGAN

Implementation

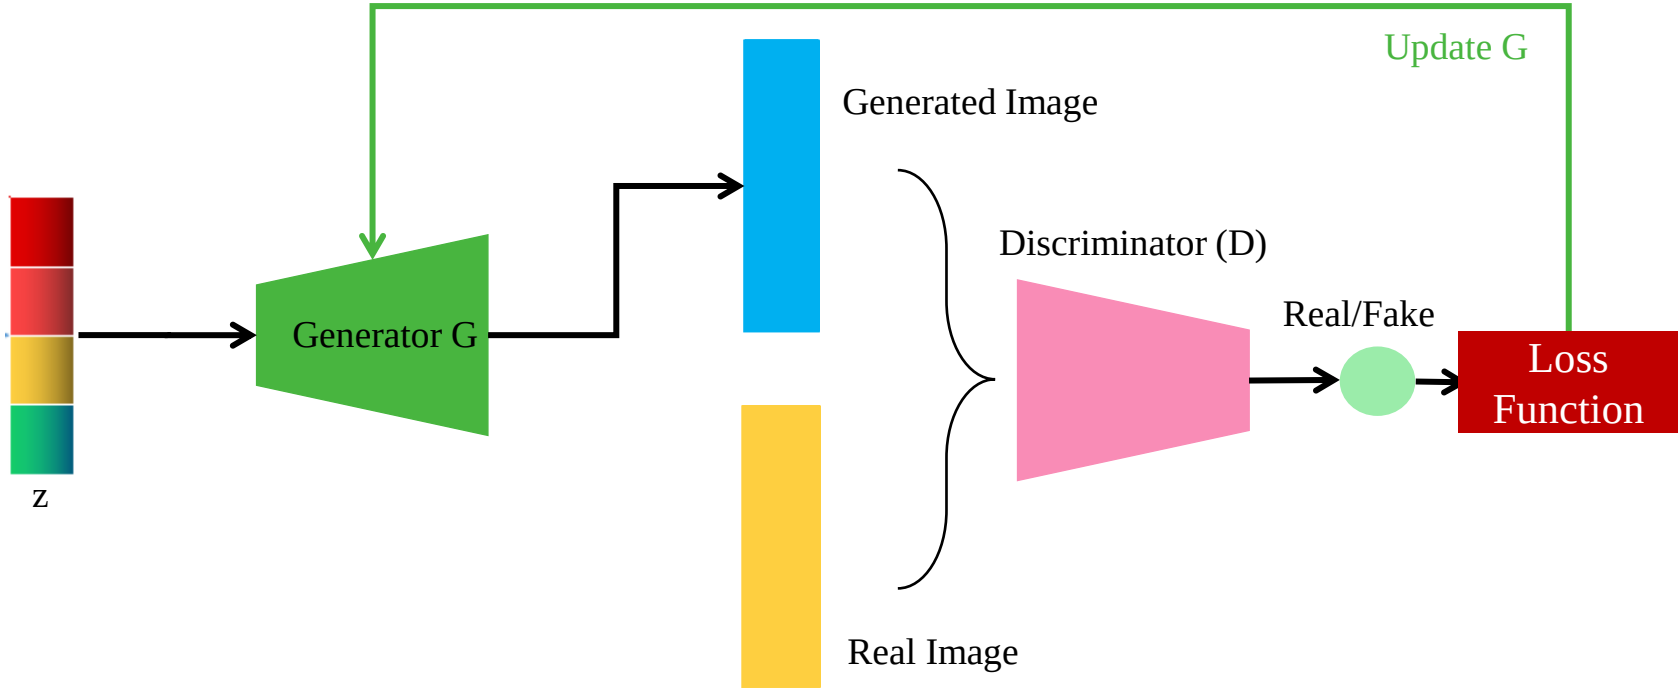

1

```
criterion = nn.BCELoss()
optimizer_G = torch.optim.Adam(generator.parameters(),
                               lr=0.0001)

for i, (imgs, _) in enumerate(dataloader)
    # prepare labels
    real_labels = torch.ones(imgs.shape[0], 1)
    fake_labels = torch.zeros(imgs.shape[0], 1)

    # zero_grad
    optimizer_G.zero_grad()
```

2

```
    z = torch.randn((imgs.shape[0], latent_dim))

    # generate images
    gen_imgs = generator(z)

    # comupte loss
    G_loss = criterion(discriminator(gen_imgs),
                       real_labels)

    # update
    G_loss.backward()
    optimizer_G.step()
```

Outline

- Introduction
- GAN
- DCGAN
- Implementation

Summary

✓ Studied Generative Adversarial Network

✓ Studied Deep Convolutional GAN

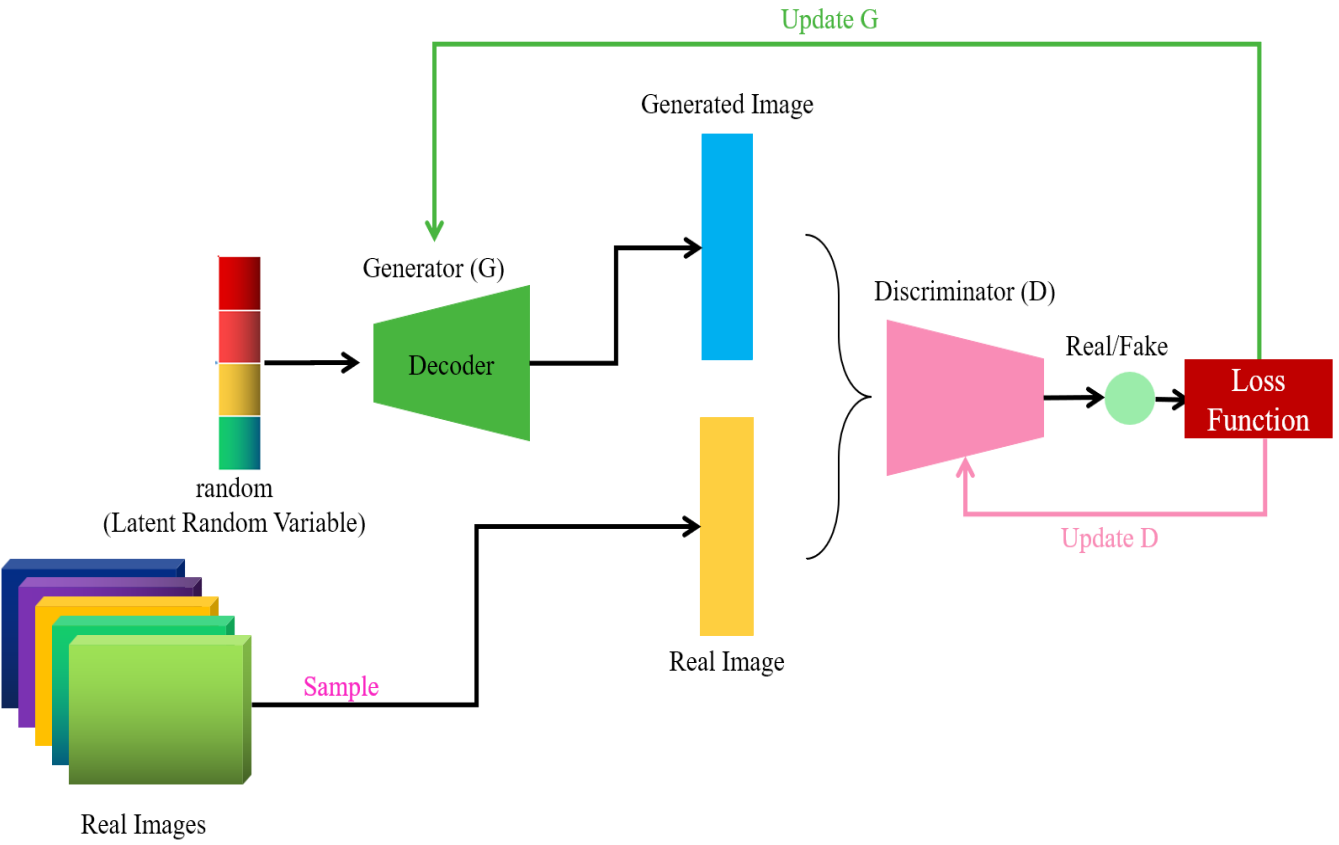

✓ GAN Extensions